



Stephen Prata

Sixth Edition

C++ Primer Plus

Developer's Library



Paul DuBois

ISBN-13: 978-0-672-32938-8

Linux Kernel Development

Robert Love

ISBN-13: 978-0-672-32946-3

Developer's Library books are available at r
as by subscription from Safari Books Online

**Developer's
Library**

informit.com/devlibrary

Stephe

◆◆ Addisc

Upper Saddle River, NJ • Boston
New York • Toronto • Montreal •
Cape Town • Sydney • Tokyo

duction, storage in a retrieval system, or transmission in any form, electronic, mechanical, photocopying, recording, or likewise. To obtain permission to reproduce material from this work, please submit a written request to Pearson Education, Inc., Department, One Lake Street, Upper Saddle River, New Jersey 07633. For a request to (201) 236-3290.

ISBN-13: 978-0-321-77640-2

ISBN-10: 0-321-77640-2

Text printed in the United States on recycled paper at R.R. Donnelley, Indianapolis, Indiana.

Second printing: January 2012

- 9 Memory Models and Namespaces**
- 10 Objects and Classes 505**
- 11 Working with Classes 563**
- 12 Classes and Dynamic Memory Allocation**
- 13 Class Inheritance 707**
- 14 Reusing Code in C++ 785**
- 15 Friends, Exceptions, and More 877**
- 16 The `string` Class and the Standard Template Library 951**
- 17 Input, Output, and Files 1061**
- 18 Visiting with the New C++ Standard**

Simple Variables	66
The <code>const</code> Qualifier	90
Floating-Point Numbers	92
C++ Arithmetic Operators	97
Summary	109
Chapter Review	110
Programming Exercises	111

4 Compound Types 115

Introducing Arrays	116
Strings	120
Introducing the <code>string</code> Class	131
Introducing Structures	140
Unions	149
Enumerations	150
Pointers and the Free Store	153
Pointers, Arrays, and Pointer Arithmetic	
Combinations of Types	184
Array Alternatives	186
Summary	190
Chapter Review	191
Programming Exercises	192

7 Functions: C

Function Review

Function Arguments

Functions and Arrays

Functions and Pointers

Functions and Structures

Functions and Unions

Functions and Enumerations

Functions and Preprocessor

Recursion 3

Pointers to Functions

Summary 3

Chapter Review

Programming

8 Adventures in

C++ Inline Functions

Reference Variables

Default Arguments

Function Overloading

Function Templates

11 Working with Classes 563

Operator Overloading 564

Time on Our Hands: Developing an Operator

Overloading Example 565

Introducing Friends 578

Overloaded Operators: Member Versus

Functions 587

More Overloading: A Vector Class 588

Automatic Conversions and Type Casts

Summary 621

Chapter Review 623

Programming Exercises 623

12 Classes and Dynamic Memory Allocation

Dynamic Memory and Classes 628

The New, Improved `String` Class 64

Things to Remember When Using `new`

in Constructors 659

Observations About Returning Objects

Using Pointers to Objects 665

Reviewing Techniques 676

A Queue Simulation 678

Summary 8
Chapter Review
Programming

15 Friends, Exceptions, and Runtime Type Casting
Friends 877
Nested Class
Exceptions
Runtime Type
Type Cast Operator
Summary 9
Chapter Review
Programming

16 The string and Template Libraries
The string
Smart Pointer
The Standard
Generic Programming
Function Objects
Algorithms
Other Libraries

Language Change	1205
What Now?	1207
Summary	1208
Chapter Review	1209
Programming Exercises	1212

Appendixes

A	Number Bases	1215
B	C++ Reserved Words	1221
C	The ASCII Character Set	1225
D	Operator Precedence	1231
E	Other Operators	1235
F	The <code>string</code> Template Class	1249
G	The Standard Template Library Methods Functions	1271
H	Selected Readings and Internet Resources	
I	Converting to ISO Standard C++	1301
J	Answers to Chapter Reviews	1335
	Index	1367

Acknowledgments for the Third Edition

I'd like to thank the editors from Macmillan who played in putting this book together: Tracy L. Rosenberg. Thanks, too, to Russ Jacobs for his help at Metrowerks, I'd like to thank Dave Mark, Al, for their help and cooperation.

Acknowledgments for the Second Edition

I'd like to thank Mitchell Waite and Scott C. Joel Fugazzotto and Joanne Miller for guiding me. Michael Marcotty of Metrowerks for dealing with the CodeWarrior compiler. I'd also like to thank you for the time to give us feedback on the first edition. Thanks to Holland, Andy Yao, Larry Sanders, Shahin M. I'd like to thank Heidi Brumbaugh for her helpful comments.

Acknowledgments for the First Edition

Many people have contributed to this book. I'd like to thank Harry Henderson for his work in developing, shaping, and reshaping the manuscript. I appreciate Harry Henderson's work.

As the reader of this book, *you* are our most important person. We value your opinion and want to know what we're doing well at, what areas you'd like to see us publish in, and any other comments that pass our way.

You can email or write directly to let us know what you think of the book—as well as what we can do to make our books better.

Please note that we cannot help you with technical questions, as that due to the high volume of mail we receive, we cannot respond to individual questions.

When you write, please be sure to include the author's name, email address, and phone number.

Email: feedback@developers-library.info
Mail: Reader Feedback
Addison-Wesley Developer's Library
800 East 96th Street
Indianapolis, IN 46240 USA

Reader Services

Visit our website and register this book at www.wiley.com/it-ebooks to get access to any updates, downloads, or errata that we post.

C++ Primer Plus discusses the basic C language, and this book self-contained. It presents C++ fundamentals with to-the-point programs that are easy to copy and paste. It covers input/output (I/O), how to make programs work with files, the many ways to handle data, and how to use the standard library features C++ has added to C, including the following:

- Classes and objects
- Inheritance
- Polymorphism, virtual functions, and runtime polymorphism
- Function overloading
- Reference variables
- Generic, or type-independent, programming with the Standard Template Library (STL)
- The exception mechanism for handling errors
- Namespaces for managing names of functions and variables

- It provides a variety of sidebars, including

The author and editors of this book do our best to make it as simple, and fun. Our goal is that by the end of the book, you will have effective programs and enjoy yourself doing so.

Sample Code Used in This Book

This book provides an abundance of sample code and programs. Like the previous editions, this book is not tied to any particular kind of computer, operating system, or compiler. We were tested on a Windows 7 system, a Macintosh, and Linux. Those programs using C++11 features require a C++11 compiler. The remaining programs should work with any C++ compiler.

The sample code for the complete programs is available on the book's website. See the registration link given in the preface.

How This Book Is Organized

This book is divided into 18 chapters and 10 appendices.

- **Chapter 1: Getting Started with C++**—This chapter introduces the C++ programming language and the C++ standard library.

values of a single type; structures, which group related data; pointers, which identify locations in memory; strings, which store text strings and to handle text I/O; the C++ `string` class. Finally, you'll learn about memory allocation, including using the `new` and `delete` operators explicitly.

- **Chapter 5: Loops and Relational Expressions** Loops are repetitive actions, and C++ provides three types: the `for` loop, the `while` loop, and the `do-while` loop. You'll learn how they should terminate, and the C++ `break` and `continue` statements to guide such loops. In Chapter 5 you'll learn how to process input character-by-character. You'll also learn about relational arrays and how to use nested loops.
- **Chapter 6: Branching Statements and the Conditional Operator** You'll learn how to control program flow by using branching statements and the conditional operator. You'll learn how to help express decision-making tests. Also, you'll learn how to help for evaluating character relations, such as whether a character is a nonprinting character. Finally, you'll get

at different methods of managing memory, which determine what parts of a program

- **Chapter 10: Objects and Classes**—An object (such as a variable) is an instance of a class. This chapter introduces object-oriented programming and to class design. It covers how data is organized as information stored in a class object and also how objects are created and destroyed. Class objects. Some parts of an object are public (the portion), and some are hidden (the private portion). Constructors and destructors come into play here. You will learn about all this and other concepts. This chapter shows how classes can be used to implement *Abstraction*.
- **Chapter 11: Working with Classes**—This chapter covers the understanding of classes. First, you'll learn about how to define and use classes. You'll define how operators such as + will work with classes. You'll see how friend functions, which can access class members, are used. You'll see how certain constructors and destructors can be used to manage conversion to and from other types.
- **Chapter 12: Classes and Dynamic Memory**—This chapter shows how to have a class member point to dynamically allocated memory. You'll see how a class constructor to allocate dynamic memory and how to provide an appropriate destructor, of defining

automobile has a motor. You also can use the model such relationships. This chapter discusses differences among the different approaches. You can let you define a class in terms of some other class. A template to create specific classes in terms of a template enables you to create a stack of integers. You learn about multiple public inheritance and how to use one class.

- **Chapter 15: Friends, Exceptions, and `final` and `override`** of friends to include friend classes and functions. This chapter discusses several new developments in C++, beginning with the `final` and `override` mechanism for dealing with unusual polymorphism. It also discusses function argument values and running-time type identification (RTTI), a mechanism for identifying objects. It also discusses alternatives to unrestricted typecasting.
- **Chapter 16: The `string` Class and the `string` Library** discusses some useful class libraries recommended by the C++ standard. `string` is a convenient and powerful alternative to C-style strings. The `string` class helps manage dynamically allocated memory. It also provides containers, including template representations. It also provides an efficient library of general string operations.

- **Appendix C: The ASCII Character Set**—The ASCII character set, along with decimal, octal, hexadecimal, and hexadecimal character representations.
- **Appendix D: Operator Precedence**—The order of decreasing precedence.
- **Appendix E: Other Operators**—Appendix E covers operators such as the bitwise operators, not covered in the main text.
- **Appendix F: The `string` Template Class**—The `string` class methods and functions.
- **Appendix G: The Standard Template Library**—Appendix G summarizes the STL containers and algorithm functions.
- **Appendix H: Selected Readings and References**—A list of some books that can further your understanding of C++.
- **Appendix I: Converting to ISO Standard C++**—Guidelines for moving from C and older C++ implementations to ISO C++.
- **Appendix J: Answers to Chapter Review Questions**—The review questions posed at the end of each chapter.

quer strategy.

- Most chapters are short enough to cover a single topic.
- The book discusses *when* to use certain features. For example, it links public inheritance to private inheritance to *has-a* relationships, and when not to.

Conventions Used in This Book

This book uses several typographic conventions:

- Code lines, commands, statements, variables, and file names are in a computer typeface:

```
#include <iostream>
int main()
{
    using namespace std;
    cout << "What's up, Doc!\n";
    return 0;
}
```

Systems Used to Develop Programming Examples

For the record, the C++11 examples in this book were compiled using g++ 4.8.1 under C++ 2010 and Cygwin with Gnu g++ 4.5.0. The remaining examples were tested with the g++ 4.2.1 under OS X 10.6.8 and on an Ubuntu 10.04 LTS. The pre-C++11 examples were originally developed using Metrowerks CodeWarrior Development Studio 5.5.1, which was tested on Windows Professional and checked using the Borland C++ 5.5.2 on the same system, using Comeau 4.3.2 on Linux, and using Metrowerks Development Studio 5.5.1 on Windows. C++ offers a lot to the programmer; learn

- The mechanics of creating a program

Welcome to C++! This exciting language for object-oriented programming and for general-purpose programming is one of the most important programming languages of the 1990s. Its ancestry brings to C++ the tradition of an excellent language. Its object-oriented heritage brings C++ a framework to cope with the escalating complexity of modern programming. Its ability to bring yet another new programming methodology to the table is both a blessing and a bane. It makes the learning process there's a lot to learn.

This chapter explores C++'s background and the ground rules for creating C++ programs. The C++ language, going from the modest basics of object-oriented programming (OOP) and its supporting concepts: encapsulation, data hiding, polymorphism, and inheritance, to generic programming. (Of course, as you learn more, you'll find from buzzwords to the necessary vocabulary

particular computer processor. So if you want a particular kind of computer, you may have to complete the program in assembly language. It was a bit as if each time a new car designers decided to change where the controls were, you had to relearn how to drive.

But Unix was intended to work on a variety of computers, which suggested using a high-level language. A *high-level language* is solving instead of toward specific hardware. So you write a high-level language to the internal language of the computer. The same high-level language program on different computers would run for each platform. Ritchie wanted a language that gave you hardware access with high-level generality and portability. In languages, he created C.

C Programming Philosophy

Because C++ grafts a new programming philosophy on top of at the older philosophy that C follows. In general, C focuses on concepts—data and algorithms. The *data* concepts are the processes. The *algorithms* are the methods the programmer uses. In mainstream languages when C was created, C++ emphasizes the algorithm side of programming.

Earlier procedural languages, such as FORTRAN, had problems as programs grew larger. For example, they had `goto` statements which route execution to one or another set of instructions on some sort of test. Many older programs had such "spaghetti programming") that it was virtually impossible to modify such a program was an invitation to disaster. C developed a more disciplined style of programming. C includes features to facilitate this approach. For example, it has branching (choosing which instruction to do next) and loops (repeating instructions). C incorporates these constructions into its syntax (the `while` loop, and the `if else` statement) into its control flow.

Top-down design was another of the new programming techniques. You divide a program into smaller, more manageable tasks. You then divide it into yet smaller tasks. You continue to divide until the task is compartmentalized into small, easily programmed tasks. (For example, to organize your desk, your table top, your filing cabinet, you start with the desk and organize each drawer, then the table top, perhaps I can manage that task.) C's design facilitates this approach.

height and width, the color and style of the line to fill the rectangle. The operations part of the class describes moving the rectangle, resizing it, rotating it, and moving the rectangle to another location. If you then use the `draw` method to create an object according to the class specifications, you pass values describing the rectangle, and you could draw the rectangle. If you drew two rectangles, the program would draw two rectangles.

The OOP approach to program design is to represent the things with which the program deals as objects. You define classes to represent rectangles, lines, circles, etc. The definitions, recall, include a description of the operations you want: moving a circle or rotating a line. Then you create objects of those classes. The process of going from low-level classes, to a higher level, such as program design, is the essence of OOP.

There's more to OOP than the binding of data to operations. For example, OOP facilitates creating reusable code. Information hiding safeguards data from unauthorized access. You can create multiple definitions for operators and methods, and then determine which definition is used. Inheritance allows you to reuse code. As you can see, OOP introduces many new concepts to programming than does procedural programming.

type. Generic programming involves extending the `function` for a generic (that is, an unspecified) type. C++ templates provide a mechanism for

The Genesis of C++

Like C, C++ began its life at Bell Labs, where, in the early 1980s. In Stroustrup's own words, "my friends and I would not have to program in assembly languages. Its main purpose was to make writing for the individual programmer" (Bjarne Stroustrup, 1997, 1st Edition. Reading, MA: Addison-Wesley, 1997).

Bjarne Stroustrup's Home Page

Bjarne Stroustrup designed and implemented C++. He is the author of the definitive reference manuals *The C++ Programming Language* and *Evolution of C++*. His personal website at www.research.att.com/~bs is a bookmark, or favorite, you create:

www.research.att.com/~bs

This site includes an interesting historical perspective on the language, Stroustrup's biographical material, and answers to frequently asked questions. One frequently asked question may be how to pronounce C++. Stroustrup's website and download the `.WAV` file.

OOP onto C, you can ignore C++'s object-orientation; that's all you do.

Only after C++ achieved some success did the merits of template programming. And only after the template programming became apparent that templates were perhaps more significant, some would argue. The fact that generic programming, as well as the more traditional programming that C++ emphasizes the utilitarian over the artistic, are the reasons for the language's success.

Portability and Standards

Say you've written a handy C++ program for Windows 2000 at work, but management decides to get a new computer using a different operating system, a different processor design, such as a SPARC processor, or a new platform? Of course you'll have to recompile the program designed for the new platform. But will you have to rewrite it? If you can recompile the program without any major hitch, we say the program is *portable*.

There are a couple obstacles to portability, that is hardware specific is not likely to be portable. An IBM PC video board, for example, speaks gibberish. You can minimize portability problems by localizing the hardware modules; then you just have to rewrite those sections of programming in this book.

The second obstacle to portability is language. It is a problem with spoken languages. A Yorkshireman's English is not portable to Brooklyn, even though English is a world language, too, can develop dialects. Although C++ has their versions of C++ compatible with others, there is no standard describing exactly how the language is implemented. The Standards Institute (ANSI) created a committee to develop a standard for C++. (ANSI had already developed a standard for Standardization (ISO) soon joined the effort (ISO-WG-21), creating a joint ANSI/ISO effort.

and the corresponding rules for C++, but the some features first introduced in C++, such as a qualifier.

Prior to the emergence of ANSI C, the C was based on the book *The C Programming Language* (Prentice-Hall Publishing Company, Reading, MA, 1978). The emergence of ANSI C, the simpler K&R.

The ANSI C Standard not only defines the standard library that ANSI C implementations must support, but also refers to it as the *standard C library* or the *standard C library*. The standard provides a standard library of C++ classes.

The C Standard was last revised as C99, which was adopted by ANSI in 2000. This standard adds some features that some C++ compilers support.

Language Growth

Originally, the de facto standard for C++ was the 328-page *The C++ Programming Language*, by Bjarne Stroustrup.

The next major published de facto standard was the *C++ Reference* by Ellis and Stroustrup (Addison-Wesley, 1990), which provided substantial commentary in addition to reference.

to creating programs.

The Mechanics of Creating

Suppose you've written a C++ program. How it runs depends on your computer environment and the compiler you use. However, the steps they should resemble the following steps (see Figure 1-1):

1. Use a text editor of some sort to write the program. The file you create constitutes the *source code* for your program.
2. Compile the source code. This means run the compiler, which translates the source code to the internal language, called *machine code*. The file containing the translated program is called the *object code*.
3. Link the object code with additional code from other programs, called *libraries*. A C++ library contains object code for common routines, called *functions*, to perform tasks such as calculating the square root of a number. The linker combines the object code for the functions you use and the object code for your program to produce a runtime version of your program called the *executable code*.

You will encounter the term *source code* throughout this book. It refers to the code in your personal random-access memory.

Most of the programs in this book are generic and support C++98. However, some, particularly those in Chapter 1, require C++11 support. At the time of this writing, some compilers require the `-std=c++11` flag when compiling a source code file.

```
g++ -std=c++11 use_auto.cpp
```

The steps for putting together a program may vary. Let's look at some examples.

Creating the Source Code File

The rest of the book deals with what goes into a source code file and the mechanics of creating one. Some C++ implementations, such as Embarcadero C++ Builder, Apple Xcode, Open Watcom, and Freescale CodeWarrior, provide *integrated development environments* that manage all steps of program development, including editing, compiling, and linking. Other implementations, such as GNU C++ on Unix and Linux, and the free versions of the Borland 5.5 (distributed by Embarcadero) compilers, just handle the compilation and linking stages and leave the rest on the system command line. In such cases, you can use a text editor to edit and modify source code. On a Unix system, for example, you can use `emacs`. On a Windows system running in the Command

The extension you use depends on the C++ compiler. There are a number of common choices. For example, `spiffy.C` is a common choice for Unix. Note that Unix is case sensitive, meaning you must use uppercase `C`. Usually, a lowercase `c` extension also works, but standardizing on uppercase `C` avoids confusion on Unix systems, you should use `c` extensions on other systems. If you don't mind typing an extra character, you can use `cc` extensions with some Unix systems. DOS, being case insensitive, doesn't distinguish between uppercase and lowercase letters, so you can use either optional letters, as shown in Table 1.1, to distinguish between the two.

Table 1.1 Source Code Extensions

C++ Implementation	Source Code Extensions
Unix	C, cc, cxx, Cpp
GNU C++	C, cc, cxx, Cpp
Digital Mars	cpp, cxx
Borland C++	cpp
Watcom	cpp
Microsoft Visual C++	cpp, cxx, cc
Freestyle CodeWarrior	cpp, cp, cc, Cpp

compiler being invoked differing from system compiler. If the `cc` compiler is available, but realize that you might have to use a different compiler for the following discussion.

You use the `cc` command to compile your program. To distinguish it from the standard Unix C compiler, meaning you type compilation command.

For example, to compile the C++ source file `spiffy.C`, type the following command at the Unix prompt:

```
cc spiffy.C
```

If, through skill, dedication, or luck, your program compiles successfully, you get an object code file with an `o` extension. In this case, the object code file is `spiffy.o`.

Next, the compiler automatically passes the object code file to the linker program that combines your code with library code. By default, the executable file is called `a.out`. If you run the linker, it deletes the `spiffy.o` file because it's no longer needed. The default name of the executable file:

a.out

Note that if you compile a new program, you get a new `a.out` file, overwriting the previous `a.out`. (That's because executable file name `a.out` helps reduce storage demands.) If you want to save the executable file, you can use the `-o` option to specify a different name.

be installed. The `g++` compiler works much like `gcc`, so the following produces an executable file called `spiffy`:

```
g++ spiffy.cxx
```

Some versions might require that you link with the C++ standard library:

```
g++ spiffy.cxx -lg++
```

To compile multiple source files, you just list them all:

```
g++ my.cxx precious.cxx
```

This produces an executable file called `a.out`. If you subsequently modify just `my.cxx`, you can recompile by using `my.cxx` and the previously compiled `precious.o`:

```
g++ my.cxx precious.o
```

The GNU compiler is available for many platforms, including a native Windows mode for Windows-based PCs as well as for Unix-like systems.

Command-Line Compilers for Windows C++

An inexpensive route for compiling C++ programs on Windows is to use a free command-line compiler that runs in Windows. The compiler opens an MS-DOS-like window. Free Windows C++ compilers are Cygwin and MinGW; they use `g++` and `gcc` under the hood.

cation. Some of these may be available in both character-based and graphical user interface (GUI) modes. Because the programs in this book are generally platform-specific code, such as Windows applications, you may need to run them in a character-based mode. The choice depends on the application. For example, if there is an option labeled Console, character-based mode is preferred. For instance, in Microsoft Visual C++ 2010, select File > Project Settings > click Application Settings, and select the Empty Project > Console Application under C++ Builder Projects.

After you have the project set up, you have several options. The IDE typically gives you several choices, such as Compile, Run, and Debug (but not necessarily all these).

- *Compile* typically means compile the code.
- *Build* or *Make* typically means compile the project. This is often an incremental process; you change just one, and then just that one.
- *Build All* typically means compile all the code.
- As described earlier, *Link* means combine the code with necessary library code.
- *Run* or *Execute* means run the program. If you have done the earlier steps, Run does them before trying to run the program.
- *Debug* means run the program with the debugger.

Occasionally, compilers get confused after including `<iostream.h>`, giving meaningless error messages that cannot be fixed up by selecting Build All to restart the process. I can't distinguish this situation from the more common situation that seems to be meaningless.

Usually, the IDE lets you run the program and keep the window open. In the window, the program finishes execution, the compiler closes the window, you'll have a hard time seeing it with quick eyes and a photographic memory. To see the output of the code at the end of the program:

```
        cin.get(); // add this statement
        cin.get(); // and maybe this, too
        return 0;
    }
```

The `cin.get()` statement reads the next character from the input stream, making the program to wait until you press the Enter key. (Note that you must press Enter, so there's no point in pressing any other key.) If the program otherwise leaves an unprocessed character in the input stream, for example, if you enter a number, you type the number and the program reads the number but leaves the Enter keystroke unprocessed. To fix this, use `cin.get()`.

programming and generic programming. This facilitates the creation of reusable code, which saves

The popularity of C++ has resulted in a large number of computing platforms; the C++ ISO standard committee is responsible for keeping these many implementations mutually compatible. The features the language should have, the base standard library of functions, classes, and templates, and the portable language across different computing platforms are the language.

To create a C++ program, you create one or more source files as expressed in the C++ language. These are then compiled to produce the machine-language files that can be linked to produce the executable. These steps are often accomplished in an IDE that provides a compiler and a linker for producing executables. IDEs also provide management and debugging capabilities. But you can also use a command-line environment by invoking the

- Placing comments in a C++ program
- How and when to use `endl`
- Declaring and using variables
- Using the `cin` object for input
- Defining and using simple functions

When you construct a simple home, you know how to work. If you don't have a solid structure from which to begin in the details, such as windows, door frames, and ballrooms. Similarly, when you learn a computer language, you learn the basic structure for a program. Only then can you begin to add and objects. This chapter gives you an overview of the language and previews some topics—notably functions and objects. The idea is to introduce you to the language in detail in later chapters. (The idea is to introduce you to the language en route to the great awakenings that come later.)

C++ Initiation

Let's begin with a simple C++ program that uses the `cout` (pronounced “see-out”) facility to produce output. The program includes several comments to the reader; these are ignored by the compiler. C++ is *case sensitive*; that is, it distinguishes between uppercase and lowercase letters.

You might find that you must alter the example. The most common reason is a matter of the program environment: some environments run the program in a separate window, and the window closes when the program finishes. As discussed in Chapter 1, you can keep the window open until you strike a key by adding the following line:

```
cin.get();
```

For some programs you must add two of these lines to make the program wait for you to press a key. You'll learn more about `cin.get()` in Chapter 1. If you have a very old system, it may not support `cin.get()`. Some programs require a compiler with some extensions. The extensions will be clearly identified and, if possible, altered to work on older compilers.

After you use your editor of choice to copy the source code from the files available online from this book's web page (see the back cover for more information), you can use your favorite compiler to compile the code, as Chapter 1 outlines. Here is the output of the program in Listing 2.1:

```
Come up and C++ me some time.  
You won't regret it!
```

place to start because some of the features that are new in C99, and the new directives, are simpler to understand after you

Features of the `main()` Function

Stripped of the trimmings, the sample program shows the fundamental structure:

```
int main()
{
    statements
    return 0;
}
```

These lines state that there is a function called `main()` and how the function behaves. Together they constitute a *function definition*. It has two parts: the first line, `int main()`, which is called the *function signature*, and the code enclosed in braces (`{` and `}`), which is the *function body*. Braces also go by other names, including “curly braces,” “brackets,” and “chicken lips.” However, the ISO Standard uses the term *function definition*. The function header shows the `main()` function. The function header defines the interface with the rest of the program, and the function body tells the computer about what the function should do. Each line is called a *statement*. You must terminate each statement with a semicolon when you type the examples.

ment separator. FORTRAN, for example, uses from the next. Pascal uses a semicolon to se you can omit the semicolon in certain cases, when you aren't actually separating two state agree about whether *can* implies *should*.) But rather than as a separator. The difference is t of the statement rather than a marker *between* C++ you should never omit the semicolon.

The Function Header as an Interface

Right now the main point to remember is that the definition of the `main()` function with this header function header syntax in more detail later, in [Chapter 10](#). I can't put their curiosity on hold, here's a preview:

In general, a C++ function is activated, or header describes the interface between a function and the code that calls it. The code preceding the function name is called the *function signature*. The code following the function name from a function back to the function that calls it is called the *argument list*. The flow from the calling function to the called function is called the *control flow*. When you apply it to `main()` because you are using it as the entry point of your program. Typically, however, `main()` is a function that the compiler adds to your program to mediate between the operating system and your program.

int. However, C++ has phased out that usage.

You can also use this variant:

```
int main(void)           // very explicit style
```

Using the keyword `void` in the parentheses of `main` indicates that the function takes no arguments. Under C++ (but not C), this is the same as using `void` in the parentheses. (In C, the compiler is remaining silent about whether there are arguments.)

Some programmers use this header and one of the following:

```
void main()
```

This is logically consistent because a `void` function does not return a value. However, although this variant is allowed by the C++ Standard, on other systems it fails to conform to the C++ Standard form; it doesn't require that `main` return a value.

Finally, the ISO C++ Standard makes a considerable effort to remove the tiresome necessity of having to place a `return` statement at the end of `main()`. The compiler reaches the end of `main()` without encountering a `return` statement, and it is the same as if you ended `main()` with this:

```
return 0;
```

This implicit return is provided only for `main()`.

grammer to the reader that usually identifies a particular aspect of the code. The compiler ignores comments as well as you do, and, in any case, it's incapable of doing what the compiler is concerned, Listing 2.1 looks as if it

```
#include <iostream>
int main()
{
    using namespace std;
    cout << "Come up and C++ me some time."
    cout << endl;
    cout << "You won't regret it!" << endl;
    return 0;
}
```

C++ comments run from the `//` to the end of the line, or it can be on the same line as code. In our case, `// myfirst.cpp -- displays a message`

In this book all programs begin with a comment line of code and a brief program summary. As mentioned earlier, the location of source code depends on your C++ system. For example, `myfirst.cxx` for names.

There are some alternatives to using the `std` namespace now. (If your compiler doesn't like these lines, there are many other problems with the examples in the book that make your programs work, but now let's take a look at C++.

C++, like C, uses a *preprocessor*. This is a program that runs before the main compilation takes place. (Some C++ compilers, as discussed in Chapter 1, use a translator program to convert C++ to C; this translator is also a form of preprocessor, we're not discussing the one that handles directives whose purpose is to do anything special to invoke this preprocessor. In this book, we'll use the program.

Listing 2.1 uses the `#include` directive:

```
#include <iostream>    // a PREPROCESSOR
```

This directive causes the preprocessor to include the `iostream` header program. This is a typical preprocessor action that occurs before it's compiled.

This raises the question of why you should use `iostream` in the program. The answer concerns communication with the outside world. The `io` in `iostream` refers to *input*, which is information sent to the program, and to *output*, which is information sent from the program. This scheme involves several definitions found in the `iostream` header.

the C++ version of `math.h` is the `cmath` header. The names of C header files are identical, whereas in other cases, the C++ header files have a different name. This is a cosmetic change, for the h-free header files also have the same name. Table 2.1 summarizes the naming conventions in this chapter.

Table 2.1 Header File Naming Conventions

Kind of Header	Convention	Example
C++ old style	Ends in <code>.h</code>	<code>iostream.h</code>
C old style	Ends in <code>.h</code>	<code>math.h</code>
C++ new style	No extension	<code>iostream</code>
Converted C	<code>c</code> prefix, no extension	<code>cmath</code>

In view of the C tradition of using different header file types, it appears reasonable to have some way to indicate C++ header files. The ANSI/ISO committee is working on which extension to use, so eventually the

In this spirit, the classes, functions, and variables in the C++ standard library are now placed in a namespace called `std`. This means, for example, that the `cout` variable in the `<iostream>` header is really called `std::cout` and that you must use the `using` directive and, instead, code in the program:

```
std::cout << "Come up and C++ me some time";  
std::cout << std::endl;
```

However, many users don't feel like converting all the `<iostream.h>` and `cout`, to namespace code, so the C++ standard allows them to do so without a lot of hassle. This is done by the following line means you can use names defined in the `std` namespace with the `std::` prefix:

```
using namespace std;
```

This `using` directive makes all the names in the `std` namespace available. Some people regard this as a bit lazy and potentially a bad practice. Two other approaches are to use the `std::` qualifier or to use the `using` directive to make just particular names available:

```
using std::cout;    // make cout available  
using std::endl;    // make endl available  
using std::cin;     // make cin available
```

as you might remember from Chapter 1, is a `cout` object that defines how data is stored and used.)

Well, using objects so soon is a bit awkward, but we'll wait until several more chapters. Actually, this reveals one of the things you need to know the innards of an object in order to use it, that is, how to use it. The `cout` object has a `string` member, so you can do the following to display it:

```
cout << string;
```

This is all you must know to display a string. The `cout` object's conceptual view represents the process. In this view, the stream of characters flowing from the program. The `cout` object is in the `iostream` file, represents that stream. The `<<` insertion operator (`<<`) that inserts the information into the stream. The following statement (note the terminating semicolon):

```
cout << "Come up and C++ me some time.";
```

It inserts the string “Come up and C++ me some time.” into the output stream. Somehow, that sounds like a good idea.

A First Look at Operator Overloading

If you're coming to C++ from C, you probably just like the bitwise left-shift operator (<<). In which the same operator symbol can have different text to figure out which meaning is intended. For example, the & symbol represents both the address-of operator and the bitwise AND operator. The * symbol represents both multiplication and pointer dereferencing. The point here is not the exact function of these operators, but that more than one meaning, with the compiler determining the meaning. (You do much the same when you determine the "sound financial basis.") C++ extends the operator overloading to fine operator meanings for the user-defined types.

The Manipulator endl

Now let's examine an odd-looking notation in Listing 2.1:

```
cout << endl;
```

endl is a special C++ notation that represents a new line. Inserting endl into the output stream causes the beginning of the next line. Special notations

The Newline Character

C++ has another, more ancient, way to indicate a new line:

```
cout << "What's next?\n";    // \n means start a new line
```

The `\n` combination is considered to be a single character.

If you are displaying a string, you need less code than to tag an `endl` onto the end:

```
cout << "Pluto is a dwarf planet.\n";  
cout << "Pluto is a dwarf planet." << endl;
```

On the other hand, if you want to generate a new line without the extra code, you can use `endl` for the same amount of typing, but most people find it more comfortable:

```
cout << "\n";    // start a new line  
cout << endl;    // start a new line
```

Typically, this book uses an embedded new line in strings and the `endl` manipulator otherwise. C++ guarantees that the output will be *flushed* (in, this case, immediately) when `endl` moves on. You don't get that guarantee with `\n`.

```
"Come up and C++ me some time."  
;    cout <<  
endl; cout <<  
"You won't regret it!" <<  
endl;return 0; }
```

This is visually ugly but valid code. You do not know C and C++ you can't put a space, tab, or carriage return as a name, nor can you place a carriage return in the middle of what you can't do:

```
int main()          // INVALID -- space in name  
{  
    return 0; // INVALID -- carriage return in middle  
    cout << "Behold the Beans  
        of Beauty!"; // INVALID -- carriage return in middle
```

(However, the *raw* string, added by C++11, allows including a carriage return in a string.)

Tokens and White Space in Source Code

The indivisible elements in a line of code are called *tokens*. To separate one token from the next with a space, tab, or carriage return. These characters are termed *white space*. Some single characters

```
return0;           // INVALID, must be ret
return(0);         // VALID, white space c
return (0);        // VALID, white space u
intmain();         // INVALID, white space
int main()         // VALID, white space c
int main ( )      // ALSO VALID, white sp
```

C++ Source Code Style

Although C++ gives you much formatting freedom, it is best to follow a sensible style. Having valid but inconsistent code is a waste of time. Most programmers use styles similar to that of

- One statement per line
- An opening brace and a closing brace for each block
- Statements in a function indented from the function definition
- No whitespace around the parentheses of function calls

The first three rules have the simple intent of making code easier to read. The fourth helps to differentiate functions from so-called macros, which also use parentheses. This book alerts you to

```
    cout << carrots;           // display the number of carrots
    cout << " carrots.";
    cout << endl;
    carrots = carrots - 1;    // modify the number of carrots
    cout << "Crunch, crunch. Now I have " << carrots << endl;
    return 0;
}
```

A blank line separates the declaration from the rest of the program, as is usual C convention, but it's somewhat less common in C++ programs. For Listing 2.2:

```
I have 25 carrots.
Crunch, crunch. Now I have 24 carrots.
```

The next few pages examine this program in more detail.

Declaration Statements and Variables

Computers are precise, orderly machines. To store data in memory, you must identify both the storage location and the type of information requires. One relatively painless way to do this is with a *declaration statement* to indicate the type of storage and the name of the variable. For example, the program in Listing 2.2 has this declaration:

```
int carrots;
```


short term: the problem is that if you misspell a variable, you can create a new variable without realizing it. That's the problem in the following:

```
CastleDark = 34
...
CastleDark = CastleDark + MoreGhosts
...
PRINT CastleDark
```

Because `CastleDark` is misspelled (the *r* was dropped), the program will leave `CastleDark` unchanged. This kind of error is common in BASIC. However, in C++, `CastleDark` would not be declared. Therefore, the compiler would catch the error about the need to declare a variable for you to avoid the potential bug.

In general, then, a declaration indicates the type of data the program will use for the data that's stored there. In this example, a variable called `carrots` in which it can store an integer.

The declaration statement in the program is a *definition*, for short. This means that its presence defines the space for the variable. In more complex situations, you might have multiple declarations. These tell the computer to use a variable that has already been defined. In general, a declaration need not be a definition, but a definition is always a declaration.

Tip

The C++ style for declaring variables is to do it as early as possible.

Assignment Statements

An assignment statement assigns a value to a variable. The following statement assigns the integer 25 to the location `carrots`:

```
carrots = 25;
```

The `=` symbol is called the *assignment operator*. Because you can use the assignment operator several times, you can write:

```
int steinway;  
int baldwin;  
int yamaha;  
yamaha = baldwin = steinway = 88;
```

The assignment works from right to left. In the previous example, the value of `steinway`, which is now 88, is assigned to `baldwin`, which is then assigned to `yamaha`. (C++ follows C's penchance for this.)

The second assignment statement in Listing 10-10 shows the value of a variable:

```
carrots = carrots - 1; // modify the variable
```

Bases, discusses this representation.) The main problem is that C is not smart enough to recognize that carrots is an integer.

Perhaps the contrast with old C will indicate the difference. For the string "25" and the integer 25 in C, you could use C's `printf()` to print them:

```
printf("Printing a string: %s\n", "25");  
printf("Printing an integer: %d\n", 25);
```

Without going into the intricacies of `printf()`, you can use `%s` and `%d` to indicate whether you are going to print a string or an integer. In C++, you can tell `printf()` to print a string but give it an integer, and the compiler will be notified to notice your mistake. It just goes ahead and prints the integer.

The intelligent way in which `cout` behaves is a result of the C++ design. In essence, the C++ insertion operator (`<<`) and `cout` are designed to follow it. This is an example of operator overloading, which is a function overloading and operator overloading design that you design yourself.

`cout` and `printf()`

If you are used to C and `printf()`, you might be tempted to cling to your hard-won mastery of `printf()`. But `cout` is more powerful than `printf()`, with all its conversion and formatting significant advantages. Its capability to recognize and handle different data types is a major advantage.

```
    cout << "How many carrots do you have ";  
    cin >> carrots; // C++  
    cout << "Here are two more. ";  
    carrots = carrots + 2;  
    // the next line concatenates output  
    cout << "Now you have " << carrots << "  
    return 0;  
}
```

Program Adjustments

If you found that you had to add a `cin.get()` to add two `cin.get()` statements to this list onscreen. The first one will read the newline key after typing a number, and the second will Return or Enter again.

Here is an example of output from the program:

```
How many carrots do you have?
```

```
12
```

```
Here are two more. Now you have 14 carrots
```

The program has two new features: using four output statements into one. Let's take a look

The second new feature of `getinfo.cpp` is `cout`. The `iostream` file defines the `<<` operator so `cout` output as follows:

```
cout << "Now you have " << carrots << " carrots."
```

This allows you to combine string output and integer output. The resulting output is the same as what the following code produces:

```
cout << "Now you have ";
cout << carrots;
cout << " carrots.";
cout << endl;
```

While you're still in the mood for `cout` and `<<`, you can use another version this way, spreading the single statement over four lines:

```
cout << "Now you have "
     << carrots
     << " carrots."
     << endl;
```

That's because C++'s free format rules treat `<<` as a left-associative operator. This last technique is convenient when you have a long string of `<<` operators.

Another point to note is that `cout` always inserts a space before the first argument.

```
Now you have 14 carrots.
```

you've been exposed to different OOP terminology. In C++, a class corresponds to what some languages term a type, and an object corresponds to an object instance or instance variable.

Now let's get a little more specific. Recall

```
int carrots;
```

This creates a particular variable (`carrots`). If the variable is, `carrots` can store an integer and can be used in an arithmetic operation, for example. Now consider `cout`. It is an object of the `ostream` class. The `ostream` class definition describes the sort of data an `ostream` object may accept in input form with and to it, such as inserting a number. `cin` is an object created with the properties of the `istream` class.

Note

The class describes all the properties of a data type. A class is formed with it, and an object is an entity created from it.

You have learned that classes are user-defined. You can design the `ostream` and `istream` classes. Just as user-defined classes can come in class libraries. That's the case with C++ standard library. Technically, they are not built in to the C++ language.

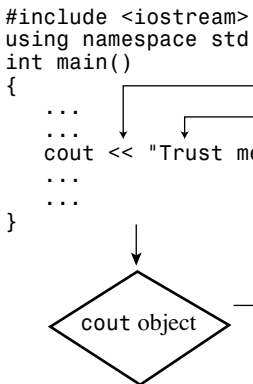


Figure 2.5 Sending a

Functions

Because functions are the modules from which are essential to C++ OOP definitions, you should understand them. Some aspects of functions are advanced and come later, in Chapter 7, “Functions: C++’s

Calling Function

```
int main()  
{
```

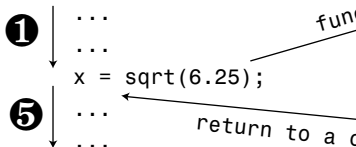


Figure 2.6 C

The value in the parentheses (6.25, in this case) is the value *passed* to the function. This value is called an *argument* or *parameter* (see Figure 2.7). The function `sqrt` calculates the square root of 6.25, which is 2.5, and sends that value back to the caller. This value is called the *return value* of the function. Think of the function call in the statement after the function definition as sending the return value to the variable `x`. In short, a function call sends a value to a function, and the return value is a value sent back to the caller.

is to use a function prototype statement.

Note

A C++ program should provide a prototype for

A function prototype does for functions what a variable declaration does for variables. It tells what types are involved. For example, to take a number with (potentially) a fractional part and return a number of the same type. Some languages use the name `double` for this type, but the name C++ uses for this type is `double`. The function prototype for `sqrt()` looks like this:

```
double sqrt(double);    // function prototype
```

The initial `double` means `sqrt()` returns a `double`. The parentheses means `sqrt()` requires a `double` argument. `sqrt()` is used exactly as used in the following code:

```
double x;                // declare x as a type double
x = sqrt(6.25);
```

The terminating semicolon in the prototype makes it a prototype instead of a function header. If you omit the semicolon, the compiler interprets the line as a function header and expects you to define the function.

usual practice is to place prototypes just before Listing 2.4 demonstrates the use of the library, including the `cmath` file.

Listing 2.4 **sqrt.cpp**

```
// sqrt.cpp -- using the sqrt() function

#include <iostream>
#include <cmath>      // or math.h

int main()
{
    using namespace std;

    double area;
    cout << "Enter the floor area, in square feet: ";
    cin >> area;
    double side;
    side = sqrt(area);
    cout << "That's the equivalent of a square with a side of "
         << " feet to the side." << endl;
    cout << "How fascinating!" << endl;
    return 0;
}
```

That's the equivalent of a square 39.1918.
How fascinating!

Because `sqrt()` works with type `double` variable type. Note that you declare a type `double` variable when you declare a type `int` variable:

```
type-name variable-name;
```

Type `double` allows the variables `area` and `side`, such as 1536.0 and 39.1918. An apparent integer with a decimal fraction part of .0 when stored in `double`. Chapter 3, type `double` encompasses a much larger range.

C++ allows you to declare new variables anywhere you declare `side` until just before using it. C++ allows you to declare `side` when you create it, so you could also have done this:

```
double side = sqrt(area);
```

You'll learn more about this process, called `variable declaration`.

Note that `cin` knows how to convert information to `double`, and `cout` knows how to insert type `double` data. In the `main` function, these objects are smart.

the function can even if there are no arguments. There also are functions that have no return value. For example, you could write a function that displayed a number in dollars-and-cents format. If the argument of, say, 23.5, and it would display \$23.50. You would pass the value to the screen instead of to the calling process. You could indicate this in the prototype by using the keyword `void`:

```
void bucks(double); // prototype for function
```

Because `bucks()` doesn't return a value, you can't use it in an assignment statement or of some other expression. Here's an example statement:

```
bucks(1234.56); // function call, no return value
```

Some languages reserve the term *function* for functions that return a value. Other languages use the terms *procedure* or *subroutine* for those without a return value. C++ uses the term *function* for both variations.

User-Defined Functions

The standard C library provides more than 100 functions. You should use them when you need them, by all means use it. But often you have to design your own functions. Anyway, it's fun to design your own functions. You've already used several user-defined functions: `main()`. Every C++ program must have a `main()` function.

```

    }

void simon(int n)    // define the simon()
{
    using namespace std;
    cout << "Simon says touch your toes "
}                  // void functions don't

```

The `main()` function calls the `simon()` function once with a variable argument `count`. In `main()`, `count` is used to set the value of `count`. The example displays a prompting message. This results in the user inputting a value in response to the prompt. Here is a sample run of the program:

```

Simon says touch your toes 3 times.
Pick an integer: 512
Simon says touch your toes 512 times.
Done!

```

Function Form

The definition for the `simon()` function in Listing 10.1 is similar to the definition for `main()`. First, there is a function signature, followed by the function body. You can generalize the form of a function definition as follows:

function #3

{ do
{
}

Figure 2.8 Funct
sequential

Function Headers

The `simon()` function in Listing 2.5 has this

```
void simon(int n)
```

The initial `void` means that `simon()` has no return value. It produces a number that you can assign to a variable. The function looks like this:

```
simon(3);                               // ok for void function
```

Because poor `simon()` lacks a return value, the following is not allowed for `simple`:

```
simple = simon(3);   // not allowed for void function
```

The `int n` within the parentheses means that the function takes a single argument of type `int`. The `n` is a new variable.

an, nowhere in any of your programs have you

```
squeeze = main();    // absent from our pro
```

The answer is that you can think of your computer (or Windows) as calling your program. So `main()` is not part of the program but to the operating system, it's the program's return value. For example, Unix shell scripts and batch files can be designed to run programs and return *values*. The normal convention is that an exit value of 0 means success, whereas a nonzero value means there was a failure. A program to return a nonzero value if, say, it fails to find a script or batch file to run that program and to indicate a signals failure.

Keywords

Keywords are the vocabulary of a computer language. In C++, the keywords are: `int`, `void`, `return`, and `double`. Because these words are keywords, you can't use them for other purposes. That is, you can't use `int` or `double` as the name of a function. But you can use `painter` (with its hidden `int`) or `return_area`. Appendix A provides a complete list of C++ keywords. Incidentally, `return` is not part of the language. Instead, it is the name of a function, not a variable name. (That can cause a problem in C++ because `return` is a keyword and because it is confusing in any case, you'd better

```

        using namespace std;
        int stone;
        cout << "Enter the weight in stone: ";
        cin >> stone;
        int pounds = stonetolb(stone);
        cout << stone << " stone = ";
        cout << pounds << " pounds." << endl;
        return 0;
    }

    int stonetolb(int sts)
    {
        return 14 * sts;
    }

```

Here's a sample run of the program in Listing 10.1:

```

Enter the weight in stone: 15
15 stone = 210 pounds.

```

In `main()`, the program uses `cin` to provide the user with a prompt. The value entered by the user is passed to the `stonetolb()` function. In `stonetolb()`, the value `sts` is passed to the function. `stonetolb()` then uses `14 * sts` to calculate the number of pounds. This illustrates that you can use a simple number. Here, by using a more complex

sort of information goes into the function, and what is returned. Programmers sometimes describe functions (specified by the flow of information into and out of the function) perfectly portrays that point of view (see Figure 2.9).

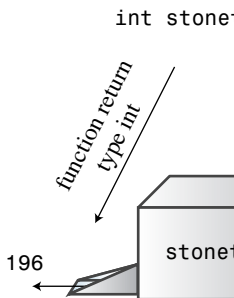


Figure 2.9 The function as a black box

The `stonetolb()` function is short and simple, with a few optional features:

- It has a header and a body.
- It accepts an argument.

```
int main()
{
    simon(3);
    cout << "Pick an integer: ";
    int count;
    cin >> count;
    simon(count);
    cout << "Done!" << endl;
    return 0;
}

void simon(int n)
{
    cout << "Simon says touch your toes "
}
```

The current prevalent philosophy is that it is best to limit access to the `std` namespace to only those functions in Listing 2.6, only `main()` uses `cout`, so therefore `cout` is available to the `stonetolb()` function. Thus, `main()` function only, limiting `std` namespace

Naming Conventions

C++ programmers are blessed (or cursed) with many naming conventions for classes, and variables. Programmers have strong opinions about these often surface as holy wars in public forums. When choosing a function name, a programmer might select an

```
MyFunction( )  
myfunction( )  
myFunction( )  
my_function( )  
my_funct( )
```

The choice will depend on the development team, the libraries used, and the tastes and preferences of the programmer. As long as any style consistent with the C++ rules is chosen, the C++ language is concerned, and it can be used.

Language allowances aside, it is worth noting that consistency and precision—is well as a personal naming convention is a hallmark of good programming throughout your programming career.

- along with the number and type of arguments.
- **Return statement**—A return statement is used to return a value to the calling function.

A class is a user-defined specification for a data type. The information is to be represented and also the operations that can be performed on the data. An object is an entity created according to a class. A variable is an entity created according to a data type.

C++ provides two predefined objects (`cin` and `cout`). They are examples of the `istream` and `ostream` classes. The `iostream` file. These classes view input and output streams. The `<<` operator, which is defined for the `ostream` class, is used to write to the output stream, and the extraction operator (`>>`), which is defined for the `istream` class, is used to extract information from the input stream. Both classes provide automatic conversion of information from one type to another in program context.

C++ can use the extensive set of C library functions. Every program should include the header file that provides the declarations for the functions.

Now that you have an overall view of the C++ language, the next chapters to fill in details and expand on the concepts.

9. What do the following function prototypes declare?
- ```
int froop(double t);
void rattle(int n);
int prune(void);
```
10. When do you not have to use the keyword `const`?
11. Suppose your `main()` function has the following code:
- ```
cout << "Please enter your PIN: ";
```

And suppose the compiler complains that you have a syntax error. What is the likely cause of this complaint, and what is the fix?

Programming Exercises

1. Write a C++ program that displays your name (without your privacy, a fictitious name and address).
2. Write a C++ program that asks for a distance in furlongs and displays it in yards. (One furlong is 220 yards.)

For reference, here is the formula for converting Celsius to Fahrenheit:

$$\text{Fahrenheit} = 1.8 \times \text{degrees Celsius} + 32$$

6. Write a program that has `main()` call a function that takes a number of light years as an argument and then returns the number of astronomical units. The program should request the light years from the user and then print the result, as shown in the following code:

```
Enter the number of light years: 4.2
4.2 light years = 265608 astronomical units
```

An astronomical unit is the average distance from the Earth to the Sun (about 150,000,000 km or 93,000,000 miles), a year (about 10 trillion kilometers or 6 trillion miles). The distance to the nearest star other than the sun is about 4.2 light years away.) Use the following conversion factor:

$$1 \text{ light year} = 63,240 \text{ astronomical units}$$

7. Write a program that asks the user to enter a number of hours and minutes. The `main()` function should then pass these values to a function that calculates the total number of minutes. The program should then play the two values in the format shown in the following code:

```
Enter the number of hours: 9
Enter the number of minutes: 28
Time: 9:28
```



```
int fooDLE; // valid and even more di
Int terrier; // invalid -- has to be in
int my_stars3 // valid
int _Mystars3; // valid but reserved -- s
int 4ever; // invalid because starts
int double; // invalid -- double is a
int begin; // valid -- begin is a Pas
int __fools; // valid but reserved -- s
int the_very_best_variable_i_can_be_versi
int honky-tonk; // invalid -- no h
```

If you want to form a name from two or more words, you can use the words with an underscore character, as in `myEye`, or you can use the first character of each word after the first, as in `myEye`. (This is the underscore method in the C tradition, whereas the C++ tradition uses the capitalization approach.) Either form makes it easy to distinguish between, say, `carDrip` and `cardR`.

Naming Schemes

Schemes for naming variables, like schemes for naming functions, have been a fervid discussion. Indeed, this topic produces more discussion than any other in programming. Again, as with function names, there are many possible variable names as long as they satisfy the rules. A consistent naming convention will serve you well.

computer memory can represent as possible. C++ provides several choices. This gives you the best that meets a program's particular requirements. This presages the designed data types of OOP.

The various C++ integer types differ in the amount of memory they use to represent an integer. A larger block of memory can represent a larger range of values. Some types (signed types) can represent both positive and negative values; others (unsigned types) can't represent negative values. The amount of memory used for an integer is *width*. The more memory, the larger the range of values. C++'s basic integer types, in order of increasing width, are `short`, `int`, `long`, and `long long`. Each comes in both signed and unsigned versions. You have a choice of ten different integer types! (Because the `char` type has some special properties, and because `wchar_t` represents characters instead of numbers), this chapter covers

The short, int, long, and long long

Computer memory consists of units called *bits*. (We'll discuss bits in more detail in this chapter.) By using different numbers of bits, the types `short`, `int`, `long`, and `long long` can represent up to different ranges of values. It's convenient if each type were always some particular size. If `short` were always 16 bits, `int` were always 32 bits, and `long` were always 64 bits, you could

A byte usually means an 8-bit unit of memory, but that describes the amount of memory in a computer, and a megabyte equal to 1,024 kilobytes. However, a byte consists of at least enough adjacent bits to hold the implementation. That is, the number of possible values or number of distinct characters. In the United States, the ASCII and EBCDIC sets, each of which can be stored in a typically 8 bits on systems using those character sets, can require much larger character sets, such as UTF-16, which use a 16-bit byte or even a 32-bit byte. Some

Many systems currently use the minimum of 32 bits. This still leaves several choices open for implementation and meet the standard. It could even be 64 bits, but at least that wide. Typically, `int` is 16 bits (the same as `short`) and 32 bits (the same as `long`) for Windows, macOS, and many other minicomputers. Some systems give you a choice of how to handle `int`. (Which system's example shows you how to determine the line number in an open a manual.) The differences between implementations can cause problems when you move a C++ program from one system to another using a different compiler on the same system. This chapter, can minimize those problems.

variable.

Listing 3.1 `limits.cpp`

```
// limits.cpp -- some integer limits
#include <iostream>
#include <climits>                // use limits
int main()
{
    using namespace std;
    int n_int = INT_MAX;          // initial
    short n_short = SHRT_MAX;    // symbols
    long n_long = LONG_MAX;
    long long n_llong = LLONG_MAX;

    // sizeof operator yields size of type
    cout << "int is " << sizeof (int) << "
    cout << "short is " << sizeof n_short << "
    cout << "long is " << sizeof n_long << "
    cout << "long long is " << sizeof n_llong << "
    cout << endl;

    cout << "Maximum values:" << endl;
    cout << "int: " << n_int << endl;
    cout << "short: " << n_short << endl;
```

```
long: 2147483647
```

```
long long: 9223372036854775807
```

```
Minimum int value = -2147483648
```

```
Bits per byte = 8
```

These particular values came from a system running

The following sections look at the chief programming

The sizeof Operator and the `climits` Header

The `sizeof` operator reports that `int` is 4 bytes on the
byte. You can apply the `sizeof` operator to a type name
use the `sizeof` operator with a type name, such as `int`
theses. But when you use the operator with the name of
parentheses are optional:

```
cout << "int is " << sizeof (int) << " bytes.\n";  
cout << "short is " << sizeof n_short << " bytes.\n";
```

The `climits` header file defines symbolic constants
stants the Preprocessor Way,” later in this chapter) to rep
previously, `INT_MAX` represents the largest value type in
2,147,483,647 for our Windows 7 system. The compiler
file that reflects the values appropriate to that compiler.
some older systems that used a 16-bit `int`, defines `INT_`

<code>UINT_MAX</code>	Maximum unsigned int
<code>LONG_MAX</code>	Maximum long value
<code>LONG_MIN</code>	Minimum long value
<code>ULONG_MAX</code>	Maximum unsigned long
<code>LLONG_MAX</code>	Maximum long long value
<code>LLONG_MIN</code>	Minimum long long value
<code>ULLONG_MAX</code>	Maximum unsigned long long

Symbolic Constants the Preprocessor V

The `climits` file contains lines similar to the

```
#define INT_MAX 32767
```

Recall that the C++ compilation process first uses the preprocessor. Here `#define`, like `#include`, is a preprocessor command. It tells the preprocessor is this: Look through the source file and replace each occurrence with `32767`. So the `#define` is like the find-and-replace command in a text editor or word processor. It does so after these replacements occur. The preprocessor also ignores comments (words) and skips embedded words. That is, the

time the program tries to initialize the other

The initialization syntax shown previously is C++ syntax that is not shared with C:

```
int owls = 101;    // traditional C initialization
int wrens(432);    // alternative C++ syntax
```

Caution

If you don't initialize a variable that is defined before use, its value is *indeterminate*. That means the value is whatever was in that location prior to the creation of the variable.

If you know what the initial value of a variable should be, you can do the declaring of a variable from assigning it a value.

```
short year;        // what could it be?
year = 1492;        // oh
```

But initializing the variable when you declare it is the preferred way to set the value later.

C++ added the parentheses form of initialization more like initializing class variables. C++11 now allows (with or without the =) with all types—a university may introduce you to initialization using the book of historical oddities retained for backward compatibility.

Unsigned Types

Each of the four integer types you just learned about can't hold negative values. This has the advantage that each type can hold a larger range of values. For example, if `short` represents the range -32,768 to 32,767, its unsigned version can represent the range 0 to 65,535. Use unsigned types only for quantities that are never negative, such as happy face manifestations. To create unsigned types, use the keyword `unsigned` to modify the declaration.

```
unsigned short change;           // unsigned short
unsigned int rovert;             // unsigned int
unsigned quarterback;            // also unsigned
unsigned long gone;              // unsigned long
unsigned long long lang_lang;    // unsigned long long
```

Note that `unsigned` by itself is short for `unsigned int`.

```
    sam = ZERO;
    sue = ZERO;
    cout << "Sam has " << sam << " dollar
    cout << " dollars deposited." << endl;
    cout << "Take $1 from each account."
    sam = sam - 1;
    sue = sue - 1;
    cout << "Sam has " << sam << " dollar
    cout << " dollars deposited." << endl;
    return 0;
}
```

Here's the output from the program in Listing 13.1:

```
Sam has 32767 dollars and Sue has 32767 dollars
Add $1 to each account.
Now Sam has -32768 dollars and Sue has 32768 dollars
Poor Sam!
Sam has 0 dollars and Sue has 0 dollars deposited
Take $1 from each account.
Now Sam has -1 dollars and Sue has 65535 dollars
Lucky Sue!
```

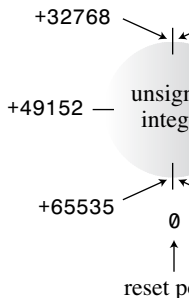
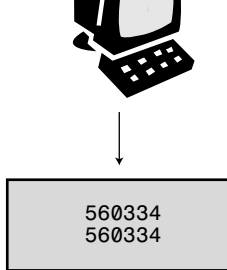


Figure 3.1 Typical overf

Choosing an Integer Type

With the richness of C++ integer types, which the most “natural” integer size for the target c form that the computer handles most efficiently choose another type, you should use `int`.

Now look at reasons why you might use an thing that is never negative, such as the number unsigned type; that way the variable can repre



Type `int` worked on this computer.

Figure 3.2 For portability

Using `short` can conserve memory if `short` is important only if you have a large array of integers. If you have several values of the same type sequentially in memory space, you should use `short` instead of `int`, for example, that you move your program from a 32-bit system to a 16-bit system. That doubles the amount of memory needed for a `short` array. Remember, a `short` is 2 bytes.

If you need only a single byte, you can use `char`.

```
int waist = 0x42;    // hexadecimal int  
int inseam = 042;    // octal integer 1  
  
cout << "Monsieur cuts a striking figure!" << endl;  
cout << "chest = " << chest << " (42 in decimal)" << endl;  
cout << "waist = " << waist << " (0x42 in hex)" << endl;  
cout << "inseam = " << inseam << " (042 in octal)" << endl;  
return 0;  
}
```

By default, `cout` displays integers in decimal. In this example, `cout` is used in a program, as the following output shows:

```
Monsieur cuts a striking figure!  
chest = 42 (42 in decimal)  
waist = 66 (0x42 in hex)  
inseam = 34 (042 in octal)
```

Keep in mind that these notations are merely for readability. If you are a member of a vintage PC club and read that a value is stored in hexadecimal, you don't have to convert the value to decimal. In your program, instead, you can simply use `0x10`, `012`, or `0xA`, it's stored the same way in memory.

```
    cout << "waist = " << waist << " (hex  
    cout << oct;          // manipulator for  
    cout << "inseam = " << inseam << " (o  
    return 0;  
}
```

Here's the program output for Listing 3.4:

```
Monsieur cuts a striking figure!  
chest = 42 (decimal for 42)  
waist = 2a (hexadecimal for 42)  
inseam = 52 (octal for 42)
```

Note that code like the following doesn't

```
cout << hex;
```

Instead, it changes the way `cout` displays its message to `cout` that tells it how to behave. A part of the `std` namespace and the program uses `hex` as the name of a variable. However, if you used `std::cout`, `std::endl`, `std::hex`, and `std::` name for a variable.

base 16; the term *decimal* does not necessarily without a suffix is represented by the smallest long, or long long. On a computer system unsigned int is represented as type int, 40000 is represented as long long. A hexadecimal or octal integer with one of the following types that can hold it: int, unsigned int, long, or unsigned long long. The same constant 0x9C40 represents the hexadecimal equivalent 0x9C40. The decimal is frequently used to express memory addresses. So unsigned int is more appropriate than long.

The char Type: Characters and Strings

It's time to turn to the final integer type: char. The char type is designed to store characters, such as letters and digits. Storing numbers is no big deal for computers, but programming languages take the easy way out by using a type that is another integer type. It's guaranteed to hold a range of basic symbols—all the letters, digits, punctuation, and so on—on any computer system. In practice, many systems support a single byte can represent the whole range. To handle characters, you can also use it as an integer.

```
cout << "Hola: ",  
cout << "Thank you for the " << ch <<  
return 0;  
}
```

Here's the output from the program in Listing 3.6:

Enter a character:

M

Hola! Thank you for the M character.

The interesting thing is that you type an **M**, but the program prints an **M**, not **77**. Yet if you look at the value stored in the `ch` variable, it's **77**. The magic, such as `cin` and `cout`, these worthy facilities make converting the keystroke input **M** to the value **77**. `Cin` displays character **M**; `cin` and `cout` are guided by the same value **77** into an `int` variable, `cout` displays the character **M**. (Listing 3.6 illustrates this point. In Listing 3.6, `cin` in C++: Enclose the character within two single quotation marks. In the example doesn't use double quotation marks for a string character and double quotation marks for a string. As Chapter 4 discusses, the two are quite different. In Listing 3.6, introduces a `cout` feature, the `cout.put()` function.


```
    cout << endl << "Done" << endl;  
    return 0;  
}
```

Here is the output from the program in Listing 3.6:

```
The ASCII code for M is 77
```

```
Add one to the character code:
```

```
The ASCII code for N is 78
```

```
Displaying char ch using cout.put(ch): N!
```

```
Done
```

Program Notes

In the program in Listing 3.6, the notation `'M'` initializes `ch` to the character `M`, so initializing the `char` variable `ch` to `'M'` is the same as `ch = 'M'`. Then `i` is assigned the identical value to the `int` variable `i`. Next, `cout` displays `ch` as `M` and `i` as `77`. As previously noted, `cout` chooses how to display that value—just another example of `cout`’s flexibility.

Because `ch` is really an integer, you can apply the `++` operator to it. This changes the value of `ch` to `78`. The program then displays `ch` as `N!`. (Alternatively, you can simply add `1` to `i`.) Again, `cout` displays the character and the `int` version as a number.

Remember that a class defines how to represent an object. A member function belongs to a class and describes a member function of that class. The `ostream` class, for example, has a `put()` member function that prints characters. You can use a member function only with an object of the class. The `cout` object, in this case. To use a class member function, you use a period to combine the object name with the member function name. The period is called the *membership operator*. To use the `put()` class member function with the `cout` object, you write `cout.put()`. We will detail when you reach classes in Chapter 10, but the classes you have are the `istream` and `ostream` classes. You will use member functions to get more comfortable with the classes.

The `cout.put()` member function provides a way to display a character. At this point you might wonder why not `cout.put()`. Much of the answer is historical. In C, character *variables* as characters but displayed as numbers. The problem was that earlier versions of C used `int` for characters. That is, the code 77 for 'M'. Meanwhile, `char` variables typically occupied only 8 bits (the important 8 bits) from the constant 32 bits of an `int`.

```
char ch = 'M';
```

- 'a' is 97, the ASCII code for a.
- '5' is 53, the ASCII code for the digit 5.
- ' ' is 32, the ASCII code for the space.
- '!' is 33, the ASCII code for the exclamation mark.

Using this notation is better than using the ASCII code because it doesn't assume a particular code. If a system uses a different code but 'A' still represents the character.

There are some characters that you can't enter on a keyboard. For example, you can't make the new line character. Enter key; instead, the program editor interprets the new line as a new line in your source code file. Other characters in a programming language imbue them with special significance. The backslash character delimits string literals, so you can't just use it in a string. C++ has special notations, called *escape sequences*, listed in Table 3.2. For example, \a represents the alert bell character, which makes a speaker or rings its bell. The escape sequence \" represents the double quotation mark as an ordinary character. You can use these notations in strings or in character constants.

The last line produces the following output:

```
Ben "Bugsie" Hacker  
was here!
```

Note that you treat an escape sequence, such as `"\n"`. That is, you enclose it in single quotes to create a string. You also use double quotes when including it as part of a string.

The escape sequence concept dates back to the days of teletypes using the teletype, an electromechanical typewriter. Teletypes always honor the complete set of escape sequences. The alarm character is silent for the alarm character.

The newline character provides an alternative way to insert a new line in the output. You can use the newline character in a character in a string (`"\n"`). All three of the following examples produce the same output of the next line:

```
cout << endl;      // using the endl manipulator  
cout << '\n';      // using a character constant  
cout << "\n";      // using a string
```

You can embed the newline character in a string. This is often easier than using `endl`. For example, the following two examples produce the same output:

```
cout << endl << endl << "What next?" << endl;  
cout << "\n\nWhat next?\nEnter a number:" << endl;
```

giant step for cursorkind), and the backspace character moves the cursor one column to the left. (Houdini once painted a picture of the Houdini escape artist, and he was, of course, a great escape artist.)

Listing 3.7 **bondini.cpp**

```
// bondini.cpp -- using escape sequences
#include <iostream>
int main()
{
    using namespace std;
    cout << "\aOperation \"HyperHype\" is\n";
    cout << "Enter your agent code:_____";
    long code;
    cin >> code;
    cout << "\aYou entered " << code << ".\n";
    cout << "\aCode verified! Proceed with\n";
    return 0;
}
```

Note

Some systems might behave differently, displaying the backspace character as a space, for example, or perhaps erasing the character to the left.

allows an implementation to offer extended character sets. Furthermore, those additional characters are a part of the name of an identifier. Thus, a German implementation supports umlauted vowels, and a French implementation supports accented vowels. The C++ standard provides a mechanism for representing such international characters on a particular keyboard: the use of *universal character names*.

Using universal character names is similar to using character names. A universal character name begins either with `\u` or `\U`. The `\u` form is followed by 4 hexadecimal digits, and the `\U` form by 16 hexadecimal digits. The hexadecimal digits are the code point for the character. (ISO 10646 is an international standard that provides numeric codes for a wide range of characters; see later in this chapter.)

If your implementation supports extended character sets, you can use universal character names in identifiers, as character constants, and in string literals. The following code:

```
int k\u00F6rper;  
cout << "Let them eat g\u00E2teau.\n";
```

The ISO 10646 code point for ö is 00F6, and the ISO 10646 code point for â is 00E2. A C++ code would set the variable name to `körper` and the string literal to `"Let them eat gâteau."`

also incorporates other Latin characters, such as characters from other alphabets, including Greek, Bengali; and ideographs, such as those used in Chinese. Unicode represents more than 109,000 symbols and more languages in the world. If you want to know more, you can check out unicode.org.

Unicode assigns a number, called a `code point`, to each character. The notation for Unicode code points looks like this: `U+0022` for the double quote character, and the `222B` is the hexadecimal representation of the code point in this case.

The International Organization for Standardization (ISO) developed ISO 10646, also a standard for coding characters. The Unicode group have worked together since 1990, and with one another.

signed char and unsigned char

Unlike `int`, `char` is not signed by default. Note that this is due to the C++ implementation in order to allow for portability to the hardware properties. If it is vital to you that your `char` be signed, use `signed char` or `unsigned char` explicitly.

```
char fodo;           // may be signed,
unsigned char bar;    // definitely unsigned
signed char snark;    // definitely signed
```

The `cin` and `cout` family consider input and output as streams of `wchar_t` characters, so they are not suitable for handling the wide-character input and output. There are parallel facilities in the form of `wcin` and `wcout`. The `L` prefix can indicate a wide-character constant or string literal. The `w` prefix in the code stores a `wchar_t` version of the letter `P`. The `wcout` stream outputs a `wchar_t` version of the word `tall`:

```
wchar_t bob = L'P';           // a wide-character constant
wcout << L"tall" << endl;     // outputting a wide-character string
```

On a system with a 2-byte `wchar_t`, this code uses twice as much memory. This book doesn't use the wide-character facilities, particularly if you become involved in internationalization (ISO 10646).

New C++11 Types: `char16_t` and `char32_t`

As the programming community gained more experience with Unicode, it became clear that the `wchar_t` type wasn't enough. It turns out that the number of characters on a computer system is more complicated than the number of values (called code points). In particular, it's useful to have a type of definite size and signedness. But there is no standard from one implementation to another. So C++11 introduced `char16_t`, which is unsigned and 16 bits, and `char32_t`, which is unsigned and 32 bits. The `u` prefix for `char16_t` character and string constants, and the `U` prefix for `char32_t` character and string constants, are new.

and `false`, and the predefined literals `true` and `false` can make statements like the following:

```
bool is_ready = true;
```

The literals `true` and `false` can be converted to `1` and `false` to `0`:

```
int ans = true;           // ans assigned 1
int promise = false;      // promise assigned 0
```

Also any numeric or pointer value can be converted (via an explicit type cast) to a `bool` value. Any nonzero value converts to `true`:

```
bool start = -100;        // start assigned true
bool stop = 0;            // stop assigned false
```

After the book introduces `if` statements (in the chapter “Logical Operators”), the `bool` type will become more useful.

The `const` Qualifier

Now let’s return to the topic of symbolic names. The `const` keyword suggests what the constant represents. Also if the value is constant and you need to change the value, you can just use the `const` keyword.

The general form for creating a constant is

```
const type name = value;
```

Note that you initialize a `const` in the declaration.

```
const int toes;      // value of toes undefined
toes = 10;           // too late!
```

If you don't provide a value when you declare a `const`, you get an error: "constant variable has no defined value that you cannot modify."

If your background is in C, you might feel that the `const` keyword discussed earlier, already does the job adequately. In C++, you can use `const` to specify the type explicitly. Second, you can use `const` to limit the scope of a variable to particular functions or files. (Scoping is to different modules; you'll learn about this in Chapter 11, "Models and Namespaces.") Third, you can use `const` to declare arrays and structures, as discussed in Chapter 12.

Tip

If you are coming to C++ from C and you are used to using `volatile` to declare a constant, use `const` instead.

ANSI C also uses the `const` qualifier, which is different from C++. With the ANSI C version, you should be aware

except it's based on binary numbers, so the scale is of 10. Fortunately, you don't have to know much about the main points are that floating-point numbers can represent very small values, and they have internal representations as integers.

Writing Floating-Point Numbers

C++ has two ways of writing floating-point numbers. The decimal-point notation you've been using much of the time is

```
12.34                // floating-point
939001.32            // floating-point
0.00023              // floating-point
8.0                  // still floating-point
```

Even if the fractional part is 0, as in 8.0, the number is represented in floating-point format and not as an integer. (For implementations to represent different local conventions for using the European method of using a comma as a decimal point. However, these choices govern how the number is stored, not in code.)

The second method for representing floating-point numbers looks like this: 3.45E6. This means that the value is 3.45 times 10 to the 6th power, which is 1 followed by six zeros.


```

#define DBL_DIG 15           // double
#define FLT_DIG 6           // float
#define LDBL_DIG 18         // long double

// the following are the number of bits used
#define DBL_MANT_DIG 53
#define FLT_MANT_DIG 24
#define LDBL_MANT_DIG 64

// the following are the maximum and minimum
#define DBL_MAX_10_EXP +308
#define FLT_MAX_10_EXP +38
#define LDBL_MAX_10_EXP +4932

#define DBL_MIN_10_EXP -307
#define FLT_MIN_10_EXP -37
#define LDBL_MIN_10_EXP -4931

```

Listing 3.8 examines types `float` and `double` to which they represent numbers (that's the significance). It views an `ostream` method called `setf()` from `<iomanip>`. This particular call forces output to stay in fixed notation and the precision. It prevents the program from switching to scientific notation. It causes the program to display six digits to the right of the decimal point.

```
}  
  
Here is the output from the program in Listing 3.8:  
  
tub = 3.333333, a million tubs = 3333333.  
and ten million tubs = 33333332.000000  
mint = 3.333333 and a million mints = 333
```

Program Notes

Normally `cout` drops trailing zeros. For example, `3.33333325`. The call to `cout.setf()` overrides this behavior. The main thing to note in Listing 3.8 is that `tub` and `mint` are initialized to $10.0 / 3$, which is $3.3333333333333333\ldots$ (etc.). Because `cout` is set to `fixed`, you can see that both `tub` and `mint` are accurate to seven decimal places. Each number is multiplied by a million, so you see that `tub` displays `3333333` and `mint` displays `33333332`. The `precision` of `cout` is set to 15, so you see 15 threes. `tub` is good to seven significant figures for `float`, but that's the worst-case scenario. `double` shows 13 threes, so it's good to at least 13 significant figures. If `precision` were set to 15, this shouldn't surprise you. Also note that `mint` doesn't quite result in the correct answer; this is due to floating-point precision.

When you write a floating-point constant in a program, how does the program store it? By default, floating-point constants are of type `double`. If you want a constant to be type `float`, you use an `f` or `F` suffix. (Because the lowercase `f` is a better choice.) Here are some examples:

```
1.234f           // a float constant
2.45E20F         // a float constant
2.345324E28      // a double constant
2.2L             // a long double constant
```

Advantages and Disadvantages of Floating-Point Numbers

Floating-point numbers have two advantages over integers. First, they can represent a much larger range of values. Second, because of the scaling, they can represent a much wider range of values. On the other hand, floating-point operations are slower than integer operations, and you can lose precision.

Listing 3.9 `fltadd.cpp`

```
// fltadd.cpp -- precision problems with floats
#include <iostream>

int main()
{
```

C++ brings some order to its basic types by grouping `char`, `short`, `int`, and `long` are termed *signed integer* types. The unsigned versions are termed *unsigned integer* types. C++11 adds `wchar_t`, signed integer, and unsigned integer types. C++11 adds `char16_t` and `char32_t` types. `float`, `double`, and `long double` types are termed *floating-point* types. `bool` is collectively termed *arithmetic* types.

C++ Arithmetic Operator

Perhaps you have warm memories of doing arithmetic calculations that same pleasure to your computer. C++ uses arithmetic operators for five basic arithmetic calculations: addition, subtraction, multiplication, division, and taking the modulus. Each of these operations calculate a final answer. Together, the operators are collectively termed *arithmetic* operators. For example, consider the following statement:

```
int wheels = 4 + 2;
```

The values 4 and 2 are operands, the + symbol is the operator, and the entire expression whose value is 6.


```
// arith.cpp      Some C++ arithmetic
#include <iostream>
int main()
{
    using namespace std;
    float hats, heads;

    cout.setf(ios_base::fixed, ios_base::floatfield);
    cout << "Enter a number: ";
    cin >> hats;
    cout << "Enter another number: ";
    cin >> heads;

    cout << "hats = " << hats << "; heads = " << heads << "\n";
    cout << "hats + heads = " << hats + heads << "\n";
    cout << "hats - heads = " << hats - heads << "\n";
    cout << "hats * heads = " << hats * heads << "\n";
    cout << "hats / heads = " << hats / heads << "\n";
    return 0;
}
```

As you can see in the following sample output, you can trust C++ to do simple arithmetic:

operator can be applied to the same operand. The addition operator is used first. The arithmetic operators are addition, multiplication, division, and the taking of the remainder. Thus $3 + 4 * 5$ means $3 + (4 * 5)$, not $(3 + 4) * 5$. Of course, you can use parentheses to enforce precedence. “Operator Precedence,” shows precedence for all the C++ operators in the same row in Appendix D. That means the addition and subtraction share a lower precedence.

Sometimes the precedence list is not enough.

```
float logs = 120 / 4 * 5;    // 150 or 6?
```

Once again, 4 is an operand for two operators. Because of precedence, so precedence alone doesn't tell you whether to divide 120 by 4 or multiply 4 by 5. Because the first choice leads to a result of 6, the choice is an important one. When operators have the same precedence, C++ looks at whether the operators have left associativity. Left-to-right associativity means that if two operands have the same precedence, you apply the leftmost operator first. If right associativity, you apply the right-hand operator first. Appendix D shows that multiplication and division have left associativity. That means you use 4 with the leftmost operand, 120, as a result, and then multiply the result by 5 to get 150.

Listing 3.11 **divide.cpp**

```
// divide.cpp -- integer and floating-point
#include <iostream>
int main()
{
    using namespace std;
    cout.setf(ios_base::fixed, ios_base::floatfield);
    cout << "Integer division: 9/5 = " << 9/5 << endl;
    cout << "Floating-point division: 9.0/5.0 = " << 9.0/5.0 << endl;
    cout << 9.0 / 5.0 << endl;
    cout << "Mixed division: 9.0/5 = " << 9.0/5 << endl;
    cout << "double constants: 1e7/9.0 = " << 1e7/9.0 << endl;
    cout << 1.e7 / 9.0 << endl;
    cout << "float constants: 1e7f/9.0f = " << 1e7f/9.0f << endl;
    cout << 1.e7f / 9.0f << endl;
    return 0;
}
```

Here is the output from the program in Listing 3.11:

```
Integer division: 9/5 = 1
Floating-point division: 9.0/5.0 = 1.80000
Mixed division: 9.0/5 = 1.800000
```

type int /type int

9 / 5

operator performs
int division

type double /type double

9.0 / 5.0

operator performs
double division

Figure 3.4 DI

The Modulus Operator

Most people are more familiar with addition than with the modulus operation, so let's take it in action. The modulus operator returns the remainder of an integer division. For example, if you divide 10 by 3, you get 3 with a remainder of 1. In the modulus operation with integer division, the modulus operator requires dividing a quantity into different integer units, such as converting feet to yards and inches or converting dollars to quarters,

```
}
```

Here is a sample run of the program in Listing 10.1:

```
Enter your weight in pounds: 181
181 pounds are 12 stone, 13 pound(s).
```

In the expression `lbs / Lbs_per_stn`, `lbs` performs integer division. With a `lbs` value of 168, the product of 12 and 14 is 168, so the remainder is 0. The value of `lbs % Lbs_per_stn` is 13. Now you are prepared to respond to questions about your weight when prompted.

Type Conversions

C++'s profusion of types lets you match the type of the data to the type of the computer. For example, adding two short integers requires fewer instructions than adding two long values. With many different types, the computer can have a lot of different ways to deal with integer types. To help deal with this potential mismatch, C++ automatically converts values when you assign a value of one type to a variable of another arithmetic type.

- C++ converts values when you assign a value of one type to a variable of another arithmetic type.

`float` can have just six significant figures, the while some conversions are safe, some may present possible conversion problems.

Table 3.3 **Potential Numeric Conversion Problems**

Conversion Type	Potential Problem
Bigger floating-point type to smaller floating-point type, such as <code>double</code> to <code>float</code>	Loss of precision
Floating-point type to integer type	Loss of precision; if value is too big, it is truncated
Bigger integer type to smaller integer type, such as <code>long</code> to <code>short</code>	Truncation

A zero value assigned to a `bool` variable is converted to `true`.

Assigning floating-point values to integer variables (converting floating-point to integer results in truncation of the fractional part). Second, a `float` value might be too big

Here is the output from the program in Listing 10.1:

```
tree = 3.000000  
guess = 3  
debt = 1634811904
```

In this case, `tree` is assigned the floating-point value 3.0, and the integer variable `guess` causes the value to be truncated (dropping the fractional part) and not rounding (finding the nearest integer). Finally, `debt` holds the value 7.2E12. This creates a situation where the debt is enormous! On this system, `debt` ends up with the value 1634811904, a novel way to reduce massive indebtedness!

Some compilers issue warnings of possible conversions from integer variables to floating-point values. Also, the results can vary from compiler to compiler. For example, running the program on a second system produced a value of 2147483648 for `debt`.

Initialization Conversions When { } Are Used

C++11 calls an initialization that uses braces *brace initialization*. It can be used more generally to provide lists of initializers. It is more restrictive in type conversions than the *direct initialization*. *Direct initialization* doesn't permit *narrowing*, which is a conversion that is not able to represent the assigned value. For example,

unsigned operand.

7. Otherwise, if the signed type can represent the unsigned operand is converted to the type.
8. Otherwise, both operands are converted to the unsigned type.

ANSI C follows the same rules as ISO C99. In C++, the preceding rules, and classic K&R C has yet to be defined. C always promotes float to double, even if both operands are float.

This list introduces the concept of ranking of types. In C++, expect, the basic ranking for signed integer types is `int`, `short`, and `signed char`. Unsigned types are ranked as `unsigned int`, `unsigned short`, and `unsigned char`. Signed type has the highest rank. The three types `char`, `signed char`, and `unsigned char` have the same rank. The `bool` type has the lowest rank. The `bool` type has the same types as their underlying types.

Conversions in Passing Arguments

Normally, C++ function prototyping controls the conversions of arguments, as you'll learn in Chapter 7, "Function Arguments". It is possible, although usually unwise, to waive the conversions. In that case, C++ applies the integral promotion rules.

The first form is straight C. The second form is C++ and the third form is to make a type cast look like a function call. The third form makes types look like the type conversions you can do in C.

C++ also introduces four type cast operators that can be used. Chapter 15, “Friends, Exceptions, and Namespaces,” discusses the `static_cast<>` operator, can be used for converting one type to another. For example, using it to convert the `thorn` variable to a `long`:

```
static_cast<long> (thorn)      // returns 1000000000
```

More generally, you can do the following:

```
static_cast<typeName> (value)  // convert value to typeName
```

As Chapter 15 discusses further, Stroustrup’s C++ is dangerously unlimited in its possibilities. The type cast is more powerful than the traditional type cast.

Listing 3.14 briefly illustrates both the basic and advanced uses of the type cast. Imagine that the first section of this listing is a program that does floating-point calculations that involve birds and animals. The results you get depend on whether you first add the floating-point values and then convert the results to `int`. But the calculations for `bats` and `coots` first convert the floating-point values to `int` and then sum the values. The first section uses a type cast to display the ASCII code for the character `'@'`.

```
        cout << static_cast<int>(cn) << endl;  
        return 0;  
    }
```

Here is the result of the program in Listing

```
auks = 31, bats = 30, coots = 30  
The code for Z is 90  
Yes, the code is 90
```

First, adding 19.99 to 11.99 yields 31.98. When you use the variable `auks`, it's truncated to 31. But using type `double` for `bats` and 11 before addition, making 30 the result for `coots`. The first two statements use type casts to convert a type `double` to `int`. These conversions cause `cout` to print the value 31.

This program illustrates two reasons to use `double`. First, values that are stored as type `double` but are used to calculate `int` values might be fitting a position to a grid or modeling a physical object. Floating-point numbers. You might want the code to cast floating-point numbers. Casting enables you to do so directly. Notice that the code for `Z` prints these values, when you convert to `int` and add them. The code for `Z` convert to `int`.

```
auto z = 0; // oops, z is int because
```

Using 0 instead of 0.0 doesn't cause problems with automatic type conversion.

Automatic type deduction becomes much more complicated types, such as those in the STL (Standard Template Library). Code might have this:

```
std::vector<double> scores;  
std::vector<double>::iterator pv = scores.begin();
```

C++11 allows you to write this instead:

```
std::vector<double> scores;  
auto pv = scores.begin();
```

We'll mention this new meaning of `auto` in more detail in the topics at hand.

Summary

C++'s basic types fall into two groups. One group consists of integers. The second group consists of values that are not integers. Integer types differ from each other in the amount of bits they use, whether they are signed or unsigned. From s

which operation takes place first.

C++ converts values from one type to another mix types in arithmetic, and use type casts to ensure conversions are “safe,” meaning you can make them safe. For example, you can convert an `int` value to a `long` value. Conversions of floating-point types to integer types are not safe.

At first, you might find the large number of conversions a bit cumbersome, but especially when you take into account the various uses of these types, you will eventually find occasions when one of the type conversions is useful. You’ll thank C++ for having it.

Chapter Review

1. Why does C++ have more than one integer type?
2. Declare variables matching the following:
 - a. A short integer with the value 8
 - b. An unsigned `int` integer with the value 10
 - c. An integer with the value 3,000,000

- e. `15 % 4`
- 9. Suppose `x1` and `x2` are two type `double` and assign to an integer variable. Construct an expression that gives you what you want to add them as type `double`.
- 10. What is the variable type for each of the following?
 - a. `auto cars = 15;`
 - b. `auto iou = 150.37f;`
 - c. `auto level = 'B';`
 - d. `auto crat = U'/U00002155';`
 - e. `auto fract = 8.25f/2.5;`

Programming Exercises

1. Write a short program that asks for your height in feet and inches. Have the program convert the height to centimeters. Indicate where to type the response. Also, print out the conversion factor.

lent time in days, hours, minutes, and seconds. The program should also calculate the number of hours in the day, the number of minutes in an hour, and the number of seconds in a minute. The output should look like this:

```
Enter the number of seconds: 31600000
31600000 seconds = 365 days, 17 hours, 46 minutes, 40 seconds
```

5. Write a program that requests the user to enter the current population of the U.S. (or other nation) and the world's population. The information is stored in variables of type long. The program should calculate the percentage that the U.S. (or other nation's) population is of the world's population. The output should look something like this:

```
Enter the world's population: 689875000
Enter the population of the US: 310700000
The population of the US is 4.50492% of the world's population
```

You can use the Internet to get more recent population figures.

6. Write a program that asks how many miles of gasoline you have used and then reports the result. Or, if you prefer, the program can request the number of gallons of gasoline used and then report the result European style.

- Creating and using structures
- Creating and using unions
- Creating and using enumerations
- Creating and using pointers
- Managing dynamic memory with new
- Creating dynamic arrays
- Creating dynamic structures
- Automatic, static, and dynamic storage
- The vector and array classes (an intro

Say you've developed a computer game called *Wits* with a cryptic and abusive computer intelligence. *Wits* keeps track of your monthly game sales for a year, your accumulation of hacker-hero trading cards, and so on. Nothing more than C++'s simple basic types to do this. C++ offers something more—compound types. The `string` type, `int`, and floating-point types. The most far-reaching extension is the use of OOP toward which we are progressing. But first, the compound types taken from C. The array, for example, is a compound type. A particular kind of array can hold a string, for example. An array can hold several values of differing types. The `new` operator tells a computer where data is placed. You'll ex-

short value.

```
short months[12];    // creates array of
```

Each element, in essence, is a variable that

This is the general form for declaring an a

```
typeName arrayName[arraySize];
```

The expression *arraySize*, which is the n
stant, such as 10 or a `const` value, or a constan
for which all values are known at the time co
arraySize cannot be a variable whose value
ever, later in this chapter you'll learn how to r
restriction.

The Array as Compound Type

An array is called a *compound type* because i
term *derived type*, but because C++ uses the
come up with a new term.) You can't simply d
to be an array of some particular type. There
many specific array types, such as array of ch
declaration:

```
float loans[20];
```

The type for `loans` is not “array”; rather, it is
`loans` array is built from the `float` type.

The Importance of Valid Subscript Values

The compiler does not check to see if you use valid subscript values. It won't complain if you assign a value to the memory address. The assignment could cause problems when the program runs, possibly causing the program to abort. So it is important that a program uses only valid subscript values.

The yam analysis program in Listing 4.1 demonstrates this, including declaring an array, assigning values to array elements,

Listing 4.1 **arrayone.cpp**

```
// arrayone.cpp -- small arrays of integers
#include <iostream>
int main()
{
    using namespace std;
    int yams[3];    // creates array with 3 elements
    yams[0] = 7;    // assign value to first element
```

Here is the output from the program in Listing 4.1:

```
Total yams = 21
The package with 8 yams costs 30 cents per yam.
The total yam expense is 410 cents.
```

```
Size of yams array = 12 bytes.
Size of one element = 4 bytes.
```

Program Notes

First, the program in Listing 4.1 creates a three-element array. The array has three elements, the elements are numbered 0 through 2, and the index values of 0 through 2 are used to assign values to the individual yam element. The individual yam element is an `int` with all the right attributes. The program `arrayone.cpp` can, and does, assign values to the array and display elements.

The program uses the long way to assign values to the array. It initializes array elements within the declaration and then assigns values to the `yamcosts` array:

```
int yamcosts[3] = {20, 30, 5};
```

```
int cards[4] = {3, 6, 8, 10};           // ok
int hand[4];                             // ok
hand[4] = {5, 6, 7, 9};                 // no
hand = cards;                           // no
```

However, you can use subscripts and assign

When initializing an array, you can provide an example, the following statement initializes o

```
float hotelTips[5] = {5.0, 2.5};
```

If you partially initialize an array, the compiler initializes the rest to zero. Thus, it's easy to initialize all the elements of an array by setting the first element to zero and then let the compiler

```
long totals[500] = {0};
```

Note that if you initialize to {1} instead of {0}, the rest still get set to 0.

If you leave the square brackets ([]) empty, the compiler counts the elements for you. Suppose, for example,

```
short things[] = {1, 5, 3, 8};
```

The compiler makes things an array of four elements.

Third, as discussed in Chapter 3, list-initialization

```
long plifs[] = {25, 92, 3.0};           /  
char slifs[4] {'h', 'i', 1122011, '\0'}; /  
char tlifs[4] {'h', 'i', 112, '\0'};    /
```

The first initialization fails because conversion from `double` to `long` type is narrowing, even if the floating-point value is an integer. The second initialization fails because `1122011` is not a `char`; we have an 8-bit `char`. The third succeeds because `112` is in the range of a `char`.

The C++ Standard Template Library (STL) provides a `vector` template class, and C++11 adds an `array` template class, more sophisticated and flexible than the built-in `array`. We will discuss them briefly later, and Chapter 16, “The Standard Template Library,” discusses them more fully.

Strings

A *string* is a series of characters stored in consecutive memory. This chapter discusses some of dealing with strings. The first, taken from C++11, is the `std::string` one this chapter examines. Later, this chapter discusses the `std::string_view` string class library.

The cat array example makes initializing a single quotes and then having to remember the other way to initialize a character array to a string *constant* or *string literal*, as in the following:

```
char bird[11] = "Mr. Cheeps";      // the
char fish[] = "Bubbles";           // let
```

Quoted strings always include the terminating character, so you have to spell it out (see Figure 4.2). Also the string from keyboard input into a char array is a character for you. (If, when you run the program, you have to use the keyword `static` to initialize arrays, too.)

Of course, you should make sure the array is big enough for the string, including the null character. Initialization is one case where it may be safer to let the compiler do it for you. There is no harm, other than wasted space, in that. That's because functions that work with strings work with the character, not by the size of the array. C++ in

Caution

When determining the minimum array size needed for a string, don't forget the terminating null character in your count.

address to shirt_size.

```
char shirt_size = "S";           // illegal
```

Because an address is a separate type in C++, this is nonsense. (We'll return to this point later in the book.)

Concatenating String Literals

Sometimes a string may be too long to convey a message, so you concatenate string literals—that is, to concatenate them. Indeed, any two string constants separated only by a space are automatically joined into one. Thus, all the following are equivalent to each other:

```
cout << "I'd give my right arm to be" " a great violinist." << endl;
cout << "I'd give my right arm to be a great violinist." << endl;
cout << "I'd give my right arm to be a great violinist." << endl;
cout << "I'd give my right arm to be a great violinist.\n";
```

Note that the join doesn't add any spaces to the second string immediately follows the last character of the first string. The `\0` character from the first string is replaced by the second string.

```
    cout << "Howdy! I'm " << name2;
    cout << "! What's your name?\n";
    cin >> name1;
    cout << "Well, " << name1 << ", your
    cout << strlen(name1) << " letters are
    cout << "in an array of " << sizeof(name1) << "
    cout << "Your initial is " << name1[0];
    name2[3] = '\0'; // set to null
    cout << "Here are the first 3 characters of my name: ";
    cout << name2 << endl;
    return 0;
}
```

Here is a sample run of the program in Listing 10.1:

```
Howdy! I'm C++owboy! What's your name?
```

```
Basicman
```

```
Well, Basicman, your name has 8 letters and 8 bytes stored
in an array of 15 bytes.
```

```
Your initial is B.
```

```
Here are the first 3 characters of my name: C++
```

```
name2[3] = '\\0';
```

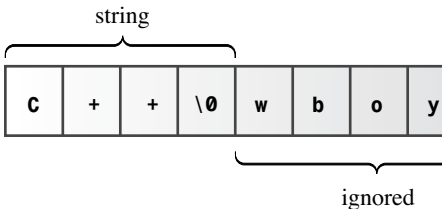


Figure 4.3 Shorter

Note that the program in Listing 4.2 uses `sizeof`. The size of an array appears in several statements. Using `sizeof` to represent the size of an array simplifies revising the program's size; you just have to change the value once, and

Adventures in String Input

The `strings.cpp` program has a blemish that is a by-product of the technique of carefully selected sample input. Listing 4.3 shows that string input can be tricky.

Enter your name:

Alistair Dreeb

Enter your favorite dessert:

I have some delicious Dreeb for you, Alis

We didn't even get a chance to respond to the prompt, and then immediately moved on to display the name.

The problem lies with how `cin` determines the end of a string. You can't enter the null character from the keyboard, so `cin` is locating the end of a string. The `cin` technique is to use newlines—to delineate a string. This means `cin` stops at the first newline character for a character array. After it reads this word, it stops at the first character when it places the string into the array.

The practical result in this example is that `cin` reads "Alistair" and puts it into the `name` array. This leaves `pointer` empty. When `cin` searches the input queue for the next string, it finds `Dreeb` still there. Then `cin` gobbles up `Dreeb` (see Figure 4.4).

Another problem, which didn't surface in this example, can turn out to be longer than the destination array. This is a protection against placing a 30-character string into a 20-character array.

Sao Paulo. You would want the program to read `cin` and `sao`. To be able to enter whole phrases instead of a different approach to string input. Specifically, you want a word-oriented method. You are in luck, for `std::string`, for example, has some line-oriented class member functions. `getline()` reads an entire input line—that is, up until a newline character. `get()` leaves the newline character, whereas `getline()` leaves it. `getline()` begins with `getline()`.

Line-Oriented Input with `getline()`

The `getline()` function reads a whole line, up until the Enter key to mark the end of input. You use it as a function call. The function takes two arguments: the target (that is, the array destined to hold the line) and a limit on the number of characters to be read. The limit is no more than 19 characters, leaving room to store the null character at the end. The `getline()` member function stops reading when it reaches the limit or when it reads a newline character, whichever comes first.

For example, suppose you want to use `getline()` to read a name array. You would use this call:

```
cin.getline(name, 20);
```

```
        return 0;  
    }
```

Here is some sample output for Listing 4.

Enter your name:

Dirk Hammernose

Enter your favorite dessert:

Radish Torte

I have some delicious Radish Torte for you.

The program now reads complete names. The `getline()` function conveniently gets a line of input, including the newline character marking the end of the line, but it doesn't remove it. It replaces it with a null character when storing the line in the array.

Line-Oriented Input with `get()`

Let's try another approach. The `istream` class has a `get()` member function which comes in several variations. One variation, `get(char*, rsize_t, char)`, takes the same arguments, interprets them the same way, and discards the newline character, leaving the rest of the line in the queue. Suppose you use two calls to `get()` in a row:

```
cin.get(name, ArSize);  
cin.get(dessert, ArSize);    // a problem
```

Because the first call leaves the newline character as the first character the second call reaches the end of line without having found anything past that newline character.

Fortunately, there is help in the form of a `getline()` (no arguments) reads the single next character, disposes of the newline character and prepares the sequence works:

```
cin.get(name, ArSize);           // read first
cin.get();                       // read newline
cin.get(dessert, ArSize);        // read second
```

Another way to use `get()` is to *concatenate*, as follows:

```
cin.get(name, ArSize).get(); // concatenate
```

What makes this possible is that `cin.get()` returns an object which is then used as the object that invokes

Mixing numeric input with line-oriented strings is a simple program in Listing 4.6.

Listing 4.6 **numstr.cpp**

```
// numstr.cpp -- following number input with strings
#include <iostream>
int main()
{
    using namespace std;
    cout << "What year was your house built? ";
    int year;
    cin >> year;
    cout << "What is its street address?\n";
    char address[80];
    cin.getline(address, 80);
    cout << "Year built: " << year << endl;
    cout << "Address: " << address << endl;
    cout << "Done!\n";
    return 0;
}
```

What year was your house built?

1966

What is its street address?

43821 Unsigned Short Street

Year built: 1966

Address: 43821 Unsigned Short Street

Done!

C++ programs frequently use pointers in that aspect of strings after talking a bit about more recent way to handle strings: the C++

Introducing the `string`

The ISO/ANSI C++98 Standard expanded now, instead of using a character array to hold (or object, to use C++ terminology). As you the array and also provides a truer representa

To use the `string` class, a program has to class is part of the `std` namespace, so you have or else refer to the class as `std::string`. The `string` and lets you treat a string much like an some of the similarities and differences between

```
        cout << "The third letter in " << str2 << endl;
        << str2[2] << endl;    // use array notation
    }

    return 0;
}
```

Here is a sample run of the program in Listing 10.1:

```
Enter a kind of feline: ocelot
Enter another kind of feline: tiger
Here are some felines:
ocelot jaguar tiger panther
The third letter in jaguar is g
The third letter in panther is n
```

You should learn from this example that, in the same manner as a character array:

- You can initialize a string object to a C++ string.
- You can use `cin` to store keyboard input in a string object.
- You can use `cout` to display a string object.
- You can use array notation to access individual characters in a string object.

```
string fourth_date {"Hank's Fine Eats"};
```

Assignment, Concatenation, and

The `string` class makes some operations simpler than you can't simply assign one array to another. Here are some examples:

```
char charr1[20];           // create an array
char charr2[20] = "jaguar"; // create an array
string str1;               // create a string
string str2 = "panther";   // create a string
charr1 = charr2;           // INVALID, no implicit conversion
str1 = str2;               // VALID, object assignment
```

The `string` class simplifies combining strings. You can concatenate two `string` objects together and the `+=` operator can be used to add a `string` object to another `string` object. Continuing with the preceding example:

```
string str3;
str3 = str1 + str2;        // assign str1 + str2 to str3
str1 += str2;              // add str2 to str1
```

Listing 4.8 illustrates these usages. Note that you can also assign a `string` object to a `string` object, as well as `string` objects to a `string` object.

```
    cout << "s1 += s2 yields s1 = " << s1 << endl;
    s2 += " for a day";
    cout << "s2 += \" for a day\" yields s2 = " << s2 << endl;

    return 0;
}
```

Recall that the escape sequence `\"` represents a literal character rather than as marking the line as a comment in the program in Listing 4.8:

```
You can assign one string object to another:
s1: penguin, s2: penguin
You can assign a C-style string to a string object:
s2 = "buzzard"
s2: buzzard
You can concatenate strings: s3 = s1 + s2
s3: penguinbuzzard
You can append strings.
s1 += s2 yields s1 = penguinbuzzard
s2 += " for a day" yields s2 = buzzard for a day
```

```
char charr2[20] = "jaguar";
string str1;
string str2 = "panther";

// assignment for string objects and
str1 = str2;                // copy s
strcpy(charr1, charr2);     // copy c

// appending for string objects and c
str1 += " paste";           // add pa
strcat(charr1, " juice");   // add ju

// finding the length of a string obj
int len1 = str1.size();     // obtain
int len2 = strlen(charr1); // obtain

cout << "The string " << str1 << " co
    << len1 << " characters.\n";
cout << "The string " << charr1 << "
    << len2 << " characters.\n";

return 0;
}
```

to `strcat()` and `strcpy()`, called `strncat()` and `strncpy()`, respectively. Taking a third argument to indicate the maximum number of characters to copy, using them adds another layer of complexity.

Notice the different syntax used to obtain the length of a string:

```
int len1 = str1.size();    // obtain length of str1
int len2 = strlen(charr1); // obtain length of charr1
```

The `strlen()` function is a regular function, and that returns the number of characters in the string. The `str1.size()` method does the same thing, but the syntax for it is different. In the first case, the argument, `str1` precedes the function name and in the second case, with the `put()` method in Chapter 3, this syntax is reversed. The `size()` with `str1` with the `put()` method in Chapter 3, this syntax is reversed. The `size()` is a class method. A method is a function belonging to the same class as the method. In the first case, `strlen()` and `size()` is a string method. In short, the `strlen()` identifies which string to use, and the C++ `strlen()` uses the dot operator to indicate which string to use.

More on string Class I/O

As you've seen, you can use `cin` with the `>>` operator and `cout` with the `<<` operator to display a `string` object.

```
        getline(cin, str);           // cin no longer has any data
        cout << "You entered: " << str << endl;
        cout << "Length of string in charr after input: "
              << strlen(charr) << endl;
        cout << "Length of string in str after input: "
              << str.size() << endl;

        return 0;
    }
}
```

Here's a sample run of the program in Listing 13.1:

```
Length of string in charr before input: 2
Length of string in str before input: 0
Enter a line of text:
peanut butter
You entered: peanut butter
Enter another line of text:
blueberry jam
You entered: blueberry jam
Length of string in charr after input: 13
Length of string in str after input: 13
```


So it takes `cin` as an argument that tells it what to read. This is an argument for the size of the string because the `string` class needs to know the size of the string.

So why is one `getline()` an `istream` class method and the other `string::getline()`? The `istream` class was part of C++ long before the `string` class. The C++ design recognizes basic C++ types such as `double`, `int`, `char`, and `string` type. Therefore, there are `istream` class methods for the other basic types, but there are no `istream` class methods for `string` objects.

Because there are no `istream` class methods for `string`, you might wonder why code like this works:

```
cin >> str; // read a word into the string str
```

It turns out that code like the following does work because of a function of the `istream` class:

```
cin >> x; // read a value into a basic C++ variable x
```

But the `string` class equivalent uses a friend function. This is a friend function of the `string` class. You'll have to wait until Chapter 11 to see how this technique works. In the meantime, you can use `string` objects and not worry about the inner workings of the `string` class.

strings.

```
cout << R"(Jim "King" Tutt uses "\n" instead of endl.)"
```

This would display the following:

```
Jim "King" Tutt uses \n instead of endl.
```

The standard string literal equivalent would be this:

```
cout << "Jim \"King\" Tutt uses \" \\n\" instead of endl."
```

Here we had to use `\\` to display `\` because a single backslash is the first character of an escape sequence.

If you press the Enter or Return key while typing a line of text, the cursor moves to the next line onscreen, it also places a carriage return character in the text.

What if you want to display the combination `)` in a string? Won't the compiler interpret the first occurrence of `)` in the string as the end of the statement? No, it won't. But the raw string syntax allows you to place additional characters in a string without escaping them. You just add an opening `"` and `(`. This implies that the same additional characters must be added to the closing `)` and `"`. So a raw string beginning with `R"+"` and ending with `)"` is a statement. The statement

```
cout << R" (*(Who wouldn't?), she whispered.) +"
```

would display the following:

```
*(Who wouldn't?), she whispered.
```

representation by storing all the related basket information in a single structure. If you want to keep track of a whole team's statistics, a structure type is also a stepping stone to that. In this chapter, a little about structures now takes you that much further.

A structure is a user-definable type, with a name and a set of data properties. After you define the type, you can create structure variables, or, more generally, objects. Thus, creating a structure is a two-part process: first, you describe that describes and labels the different types of data; then, you can create structure variables, or, more generally, objects. This description's plan.

For example, suppose that Bloataire, Inc., wants to manage its product line of designer inflatables. In particular, it needs to track the item, its volume in cubic feet, and its selling price. It meets those needs:

```
struct inflatable    // structure declaration
{
    char name[20];
    float volume;
    double price;
};
```

The keyword `struct` indicates that the code defines a new data type. The identifier `inflatable` is the name, or *tag*, for the new type.

After you have defined the structure, you

```
inflatable hat;           // hat is a
inflatable whoopee_cushion; // type inflatable
inflatable mainframe;     // type inflatable
```

If you're familiar with C structures, you'll find that C++ allows you to drop the keyword `struct` when declaring a structure.

```
struct inflatable goose; // keyword struct
inflatable vincent;      // keyword omitted
```

In C++, the structure tag is used just like a variable name. This emphasizes that a structure declaration defines a new type, not a `struct` from the list of curse-inducing errors.

Given that `hat` is type `inflatable`, you can access its individual members. For example, `hat.volume` refers to the volume member of the `hat` variable, and `hat.price` refers to the price member of the `vincent` variable. In short, the members of a structure much as indices enable you to access members. If the price member is declared as type `double`, `hat.price` is equivalent to type `double` variables and can be used as a `double` variable can be used. In short, `hat` is

```

using namespace std;
inflatable guest =
{
    "Glorious Gloria", // name value
    1.88,               // volume value
    29.99               // price value
}; // guest is a structure variable of type inflatable
// It's initialized to the indicated value
inflatable pal =
{
    "Audacious Arthur",
    3.12,
    32.99
}; // pal is a second variable of type inflatable
// NOTE: some implementations require using volatile
// static inflatable guest =

    cout << "Expand your guest list with " << guest.name << " and " << pal.name << "!\n";
// pal.name is the name member of the pal inflatable
    cout << "You can have both for $";
    cout << guest.price + pal.price << "!\n";
    return 0;
}

```

local declaration—can be
used only in this function

type parts variable
type perks variable

type parts variable
can't declare a type
perks variable here

Figure 4.7 Local and ext

Variables, too, can be defined internally or
among functions. (Chapter 9, “Memory Mod
topic.) C++ practices discourage the use of e

structure, not an array.

C++11 Structure Initialization

As with arrays, C++11 extends the features of

```
inflatable duck {"Daphne", 0.12, 9.98};
```

Next, empty braces result in the individual
following declaration results in `mayor.volume`
bytes in `mayor.name` being set to 0:

```
inflatable mayor {};
```

Finally, narrowing is not allowed.

Can a Structure Use a string Class

Can you use a `string` class object instead of a
is, can you declare a structure like this:

```
#include <string>
struct inflatable // structure definition
{
    std::string name;
    float volume;
    double price;
};
```

```
struct Inflatable
{
    char name[20];
    float volume;
    double price;
};

int main()
{
    using namespace std;
    Inflatable bouquet =
    {
        "sunflowers",
        0.20,
        12.49
    };
    Inflatable choice;
    cout << "bouquet: " << bouquet.name << endl;
    cout << bouquet.price << endl;

    choice = bouquet; // assign one struct to another
    cout << "choice: " << choice.name << endl;
    cout << choice.price << endl;
    return 0;
}
```

```
        "Packard"           // value for mr_glit  
};
```

However, keeping the structure definition locally makes a program easier to read and follow.

Another thing you can do with structures is do this by omitting a tag name while simultaneously declaring a structure variable:

```
struct           // no tag  
{  
    int x;       // 2 members  
    int y;  
} position;     // a structure variable
```

This creates one structure variable called `position`. You can use the membership operator, as in `position.x`, to access the members. You can't subsequently create other variables of this limited form of structure.

Aside from the fact that a C++ program can use structures, structures have all the features discussed so far. The only changes. But C++ structures go further. Unlike C structures, they can have member functions in addition to member variables. These features most typically are used with classes rather than structures. When we cover classes, beginning with Chapter 11, we'll see how they are used.

```

inflatable guests[2] =           // in
{
    {"Bambi", 0.5, 21.99},       // fi
    {"Godzilla", 2000, 565.99}   // ne
};

```

As usual, you can format this the way you want: all members on the same line, or each separate structure member on a new line.

Listing 4.13 shows a short example that uses an array of `inflatable` structures. `guest[0]` is the first element of the `guests` array, and `guest[0].name` is the `name` member of the `inflatable` structure. You use the dot operator to access a member of the structure.

Listing 4.13 **arrstruc.cpp**

```

// arrstruc.cpp -- an array of structures
#include <iostream>
using namespace std;

struct inflatable
{
    char name[20];
    float volume;
    double price;
};

int main()
{

```

merations are discussed later in this chapter), and the actual number of bits to be used. You can use a member variable to indicate the number of bits used. This member is termed a *bit field*. Here's an example:

```
struct toggle_register
{
    unsigned int SN : 4;    // 4 bits for serial number
    unsigned int : 4;       // 4 bits unused
    bool goodIn : 1;        // valid input
    bool goodToggle : 1;    // successful toggle
};
```

You can initialize the fields in the usual manner. The following example shows how to use the struct definition to access bit fields:

```
toggle_register tr = { 14, true, false };
...
if (tr.goodIn)    // if statement covered by goodIn
...

```

Bit fields are typically used in low-level programming and the bitwise operators listed in Appendix B are used in the approach.

another time. The member name identifies the largest member. Hence, the size of the union

One use for a union is to save space when but never simultaneously. For example, suppose sets, some of which have an integer ID, and so you could use the following:

```
struct widget
{
    char brand[20];
    int type;
    union id          // format depends on
    {
        long id_num;    // type 1 widgets
        char id_char[20]; // other widgets
    } id_val;
};

...
widget prize;

...
if (prize.type == 1)          // if-
    cin >> prize.id_val.id_num;    // use
else
    cin >> prize.id_val.id_char;
```


...

However, some implementations do not handle the type limits. For example, if `band` is of type `enum spectrum`, `++band`, if valid, increments `band` to 8, which is not a valid value. Again, for maximum portability, you should use `int`.

Enumerators are of integer type and can be converted to `int`. But `int` cannot be converted automatically to the enumeration type.

```
int color = blue;           // valid, spectrum
band = 3;                   // invalid, int
color = 3 + red;            // valid, red color
...
```

Note that in this example, even though 3 is a valid value for `band`, assigning 3 to `band` is a type error. But assigning `blue` to `color` is both type `spectrum`. Again, some implementations may not handle the expression `3 + red`, addition isn't defined for `int` and `enum spectrum`. Because `3` is of type `int`, and the result is type `int`. Because of this situation, you can use enumerations in place of `int` with ordinary integers, even though arithmetic operations are not defined.

The earlier example

```
band = orange + red;        // not valid, but
```

variables of the enumeration type, you can on example:

```
enum {red, orange, yellow, green, blue, v
```

Setting Enumerator Values

You can set enumerator values explicitly by u

```
enum bits{one = 1, two = 2, four = 4, eigh
```

The assigned values must be integers. You a explicitly:

```
enum bigstep{first, second = 100, third};
```

In this case, `first` is 0 by default. Subsequ one than their predecessors. So, `third` would

Finally, you can create more than one enum

```
enum {zero, null = 0, one, numero_uno = 1}
```

Here, both `zero` and `null` are 0, and both sions of C++, you could assign only `int` val merators, but that restriction has been remove long `long` values.

in a minus sign. (For example, if the smallest [times a minus sign] is -8, and the lower limit

The idea is that the compiler can choose a representation. It might use 1 byte or less for an enumeration with type `long` values.

C++11 extends enumerations with a form that discusses this form briefly in the section “Class

Pointers and the Free Store

The beginning of Chapter 3 mentions three things a program must keep track of when it stores data. Thumbing back to that chapter, here are those things:

- Where the information is stored
- What value is kept there
- What kind of information is stored

You’ve used one strategy for accomplishing this: a declaration statement provides the type and a `new` statement tells the program to allocate memory for the value.


```
        return 0;  
    }
```

Here is the output from the program in Listing 1.1:

donuts value = 6 and donuts address = 0x0065fd44
cups value = 4.5 and cups address = 0x0065fd40

The particular implementation of `cout` shows the variables displaying address values because that is the way the system address. (Some implementations use base 10 notation.) `donuts` is at a lower memory location than `cups`, so its address is 0x0065fd44 - 0x0065fd40, or 4. This makes sense because `donuts` uses 4 bytes. Different systems, of course, will have different addresses; some may store `cups` first, then `donuts`, giving `donuts` a higher address. And some may not even use adjacent memory locations.

Using ordinary variables, then, treats the variable as a derived quantity. Now let's look at the pointer programming philosophy of memory management (see "Pointers and the C++ Philosophy.")

Making runtime decisions is not unique to O, but it is more straightforward than does C.

The new strategy for handling stored data is to treat the variable as the named quantity and the value as a separate entity. The *pointer*—holds the address of a value. Thus, the *pointer* holds the location. Applying the `*` operator, called the *dereference*, yields the value at the location. (Yes, this is the same as the `*` operator in C++ uses the context to determine whether it is a pointer. Suppose, for example, that `manly` is a pointer. Then `*manly` represents the value at that address. This is different from an ordinary `int` variable. Listing 4.15 shows how to declare a pointer.

Listing 4.15 **pointer.cpp**

```
// pointer.cpp -- our first pointer variable
#include <iostream>
int main()
{
    using namespace std;
    int updates = 6;           // declare a variable
    int * p_updates;          // declare a pointer
```

two sides of the same coin. The `updates` variable is the primary and uses the `&` operator to get the address, whereas the `*p_updates` points to `updates`, `*p_updates` and `updates` are the same. `*p_updates` exactly as you would use a type `int`. As Figure 4.8 shows, you can even assign values to `*p_updates` and `updates` pointed-to value, `updates`.

```
int jumbo = 23;  
int * pe = &jumbo;
```

These are
the same.

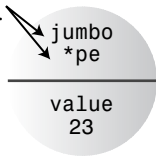


Figure 4.8 Two

1006	1000
1008	
1010	
1012	
1014	
1016	

```
int ducks = 12;
```

creates ducks variable, stores
the value 12 in the variable

Figure 4.9 Pointer

Incidentally, the use of spaces around the *
programmers have used this form:

```
int *ptr;
```

This accentuates the idea that the combin
programmers, on the other hand, use this form:

```
int* ptr;
```

both are pointers, they are pointers of two different types, and they point to different types of data. This is true on other types.

Note that whereas `tax_ptr` and `str` point to variables `tax_ptr` and `str` themselves are typed. The size of a `char` is the same size as the address of a `double`. The address for a department store, whereas 1024 could be the size or value of an address doesn't really tell you what store or building you find at that address. Usually, the address is relative to the computer system. (Some systems might have different address sizes for different types.)

You can use a declaration statement to initialize the pointed-to value, is initialized. That is, the value of the variable `higgins`:

```
int higgins = 5;  
int * pt = &higgins;
```

Listing 4.16 demonstrates how to initialize a pointer.

Listing 4.16 **`init_ptr.cpp`**

```
// init_ptr.cpp -- initialize a pointer  
#include <iostream>  
int main()  
{
```

when you create a pointer in C++, the compiler does not allocate memory to hold the data. The space for the data involves a separate step. On the surface, this is an invitation to disaster:

```
long * fellow;           // create a pointer  
*fellow = 223323;        // place a value
```

Sure, `fellow` is a pointer. But where does it point to? To `fellow`. So where is the value 223323 placed? If `fellow` is uninitialized, it could have any value. Whatever that value is, it's the address at which to store 223323. If `fellow` holds a random address, the computer attempts to place the data at address `fellow` in the middle of your program code. Chances are good that you don't want where you want to put the number 223323.7 This is one of the most insidious and hard-to-trace bugs.

Caution

Pointer Golden Rule: *Always* initialize a pointer before you apply the dereferencing operator (*) to it.

Now both sides of the assignment statement are valid. Note that just because it is valid doesn't mean that `pt` itself is type `int`. For example, `0` is an `int` is a 2-byte value and the addresses are 4-byte values.

Pointers have some other interesting properties. Meanwhile, let's look at how pointers can be used to manage memory space.

Allocating Memory with `new`

Now that you have a feel for how pointers work, let's look at an important technique of allocating memory as opposed to pointers to the addresses of variables; the variable is created at compile time, and each pointer merely provides a way to access directly by name anyway. The true worth of pointers is in *unnamed* memory during runtime to hold values that you need access to that memory. In C, you can allocate memory with `malloc()`. You can still do so in C++, but C++ has a better way.

Let's try out this new technique by creating a variable of type `int` and accessing the value with a pointer. To do this, you `new` for what data type you want memory; new

The general form for obtaining and assigning a new type can be a structure as well as a fundamental type:

```
typeName * pointer_name = new typeName;
```

You use the data type twice: once to specify the type to declare a suitable pointer. Of course, if you already have a type, you can use it rather than declare a new type. Listing 4.17 shows two different types.

Listing 4.17 **use_new.cpp**

```
// use_new.cpp -- using the new operator
#include <iostream>

int main()
{
    using namespace std;
    int nights = 1001;
    int * pt = new int;           // allocate memory
    *pt = 1001;                   // store value

    cout << "nights value = ";
    cout << nights << ": location " << &nights << endl;
    cout << "int ";
    cout << "value = " << *pt << ": location " << &pt << endl;
}
```


Program Notes

The program in Listing 4.17 uses `new` to allocate double data objects. This occurs while the program points to these two data objects. Without them. With them, you can use `*pt` and `*pd` just as you would `pt` and `pd` to assign values to the new data objects and display those values.

The program in Listing 4.17 also demonstrates how to use the type a pointer points to. An address in itself is an object stored, not its type or the number of bytes it contains. Values. They are just numbers with no type or meaning. The pointer-to-`int` is the same as the size of a pointer because `use_new.cpp` declares the pointer type `pt` to have a value of 8 bytes, whereas `*pt` is an `int` value of 4 bytes. The value of `*pd`, `cout` can tell how many bytes to print.

Another point to note is that typically `new` and `delete` are used for the ordinary variable definitions that we have seen. `pt` and `pd` have their values stored in a memory region allocated by `new` is in a region called the *heap* or *free store*.

This removes the memory to which `ps` points. You can reuse `ps`, for example, to point to another memory location. Unlike a use of `new` with a use of `delete`; otherwise, that is, memory that has been allocated but not freed, if it grows too large, it can bring a program seeking memory down.

You should not attempt to free a block of memory that was not allocated by `new`. The C++ Standard says the result of such an attempt is undefined. Sequences could be anything. Also you cannot `delete` a pointer to a variable or to a declaring ordinary variables:

```
int * ps = new int;    // ok
delete ps;             // ok
delete ps;             // not ok now
int jugs = 5;          // ok
int * pi = &jugs;      // ok
delete pi;             // not allowed, mem
```

Caution

You should use `delete` only to free memory that was allocated by `new`. Do not use `delete` to a null pointer.

that the array is built in to the program at compile time. With dynamic binding, you can create an array during runtime if you need it and skip creating it if you don't. You can select an array size after the program is running. With static binding, the array is created while the program is running. With dynamic binding, you must specify the array size at compile time. With static binding, you must specify the array size at compile time. With dynamic binding, the program can decide the array size at runtime.

For now, we'll look at two basic matters concerning arrays: how to use the `new` operator to create an array and how to use the `delete` operator to free an array.

Creating a Dynamic Array with `new`

It's easy to create a dynamic array in C++; you just specify the number of elements you want. The syntax requires you to specify the number of elements, in brackets. For example, to create an array of 10 integers, use this:

```
int * psome = new int [10]; // get a block of memory for 10 integers
```

The `new` operator returns the address of the first element of the array that value is assigned to the pointer `psome`.

As always, you should balance the call to `new` with a call to `delete` that finishes using that block of memory. However, `delete` requires using an alternative form of `delete` when you're deleting an array.

```
delete [] psome; // free the array
```

Now let's return to the dynamic array. Now, the first element of the block. It's your responsibility to free it in the block. That is, because the compiler doesn't know to the first of 10 integers, you have to write your own number of elements.

Actually, the program does keep track of the memory. It can be correctly freed at a later time when you no longer need the information isn't publicly available; you can't free it. The number of bytes in a dynamically allocated block.

The general form for allocating and assigning a block is:

```
type_name * pointer_name = new type_name [num_elements];
```

Invoking the new operator secures a block of memory. The *num_elements* elements of type *type_name*, you can use *pointer_name*. As you're about to see, you can use *pointer_name* to access the block. You can use an array name.

Using a Dynamic Array

After you create a dynamic array, how do you use it? Conceptually. The following statement creates a pointer to a block of 10 int values:

```
int * psome = new int [10]; // get a block of 10 int values
```

```
#include <iostream>
int main()
{
    using namespace std;
    double * p3 = new double [3]; // space
    p3[0] = 0.2;                    // treat
    p3[1] = 0.5;
    p3[2] = 0.8;
    cout << "p3[1] is " << p3[1] << ".\n";
    p3 = p3 + 1;                    // incre
    cout << "Now p3[0] is " << p3[0] << "
    cout << "p3[1] is " << p3[1] << ".\n";
    p3 = p3 - 1;                    // point
    delete [] p3;                   // free
    return 0;
}
```

Here is the output from the program in Listing 11.1:

p3[1] is 0.5.

Now p3[0] is 0.5 and p3[1] is 0.8.

double adds 8 to the numeric value on system to a pointer-to-short adds two to the pointer demonstrates this amazing point. It also shows the array name as an address.

Listing 4.19 **addpntrs.cpp**

```
// addpntrs.cpp -- pointer addition
#include <iostream>
int main()
{
    using namespace std;
    double wages[3] = {10000.0, 20000.0,
    short stacks[3] = {3, 2, 1};

    // Here are two ways to get the address of
    double * pw = wages;      // name of array
    short * ps = &stacks[0]; // or use address
    // with array element
    cout << "pw = " << pw << ", *pw = " <<
    pw = pw + 1;
    cout << "add 1 to the pw pointer:\n";
    cout << "pw = " << pw << ", *pw = " <<
```

```
ps = 0x28ccea, *ps = 3  
add 1 to the ps pointer:  
ps = 0x28ccec, *ps = 2
```

```
access two elements with array notation  
stacks[0] = 3, stacks[1] = 2  
access two elements with pointer notation  
*stacks = 3, *(stacks + 1) = 2  
24 = size of wages array  
4 = size of pw pointer
```

Program Notes

In most contexts, C++ interprets the name of an array as a pointer to its first element. Thus, the following statement makes `pw` a pointer to `wages`, which is the address of the first element of `wages`:

```
double * pw = wages;
```

For `wages`, as with any array, we have the following:
`wages = &wages[0]` = address of first element of `wages`

Just to show that this is no jive, the program initializes `pw` with the address of the first element of `stacks` by the expression `&stacks[0]` to initialize the `ps` pointer.

adding 1 to *pw* changes its value by 8 bytes.

Figure 4.10

After this, the program goes through similar logic to type `short` and because `short` is 2 bytes, adding 1 to *pw* changes its value by 2 (0x28ccea + 2 = 0x28cced in hexadecimal), and the program points to the next element of the array.

Note

Adding one to a pointer variable increases its value by the size of the type which it points.

Now consider the array expression `stacks[1]`. This is equivalent to `*(stacks + 1)`, which is the address of the second element of the array. The end result is precisely what `stacks[1]` means. The parentheses are necessary. Without them, 1 would be added to the value of `stacks`.

The program output demonstrates that `*(stacks + 1)` is the same as `stacks[1]`. Similarly, `*(stacks + 2)` is the same as `stacks[2]`. In pointer notation, C++ makes the following conversion: `arrayname[i]` becomes `*(arrayname + i)`.

array? Not quite—the name of the array is interpreted as an array, whereas applying the address operator to a variable is interpreted as a variable.

```
short tell[10];           // tell an array
cout << tell << endl;    // displays &tell
cout << &tell << endl;   // displays address of tell
```

Numerically, these two addresses are the same. The address of `tell`, is the address of a 2-byte block of memory. The address of a 2-byte block of memory. So the expression `&tell + 1` adds 20 to the address value. And `*(&tell + 1)` is type pointer-to-short, or `short *`, and `short (*)[20]`.

Now you might be wondering about the general case: how you could declare and initialize a pointer to an array of shorts.

```
short (*pas)[20] = &tell; // pas points to tell
```

If you omit the parentheses, precedence rules make `pas` an array of 20 pointers-to-short, so the type of `pas` is `short **`. To describe the type of a variable, you can use the `decltype` operator. To remove the variable name. Thus, the type of `pas` is `short (*)[20]` because `pas` is set to `&tell`, `*pas` is equivalent to `short`, and `pas` is an array of 20 elements of the type `short`.

You should assign a memory address to a pointer variable name to get an address of named memory or unnamed memory.

Here are some examples:

```
double * pn;           // pn can point to a double
double * pa;           // so can pa
char * pc;             // pc can point to a char
double bubble = 3.2;
pn = &bubble;          // assign address of bubble to pn
pc = new char;          // assign address of new char to pc
pa = new double[30];    // assign address of new double array to pa
```

Dereferencing Pointers

Dereferencing a pointer means referring to the value stored at the memory address pointed to by the pointer, or indirect value, operator (*) to a pointer variable. For example, dereferencing the pointer to bubble, as in the preceding example, would give the value 3.2, in this case.

Here are some examples:

```
cout << *pn; // print the value of bubble
*pc = 'S';   // place 'S' into the memory
```

Pointer Arithmetic

C++ allows you to add an integer to a pointer to get a new memory address value plus a value equal to the number of elements. You can also subtract an integer from a pointer to get a new memory address. The last operation, which yields an integer, is subtracting one pointer from another. This yields the number of elements into the same array (pointing to one position before the other). This is the separation between the two elements.

Here are some examples:

```
int tacos[10] = {5,2,8,4,1,2,2,4,6,8};
int * pt = tacos;           // suppose pt and
pt = pt + 1;                // now pt is 3004
int *pe = &tacos[9];        // pe is 3036 if a
pe = pe - 1;                // now pe is 3032,
int diff = pe - pt;          // diff is 7, the
                             // tacos[8] and ta
```

Dynamic Binding and Static Binding for Arrays

You can use an array declaration to create an array whose size is set during the compilation process.

```
int tacos[10]; // static binding, size fixed
```

```
int coats[10];  
*(coats + 4) = 12;           // set coats[4]
```

Pointers and Strings

The special relationship between arrays and pointers is illustrated by the following code:

```
char flower[10] = "rose";  
cout << flower << "s are red\n";
```

The name of an array is the address of its first element. In this case, the address of the first element of `flower` is the address of the `char` element containing the first character of the string "rose". `cout` assumes that the address of a `char` is the address of a string. It prints the character at that address and then continues printing characters until it reaches the null character (`\0`). In short, if you give `cout` the address of a `char`, it prints the character and then the character to the first null character that follows.

The crucial element here is not that `flower` is the address of a `char`. This implies that you can pass the address of a `char` element to `cout` also because it, too, is the address of a string. It's the address that points to the beginning of a string. We'll check this in a moment.

But what about the final part of the preceding code? If you pass the address of the first character of a string, what happens? It's consistent with `cout`'s handling of string output.

```

int main()
{
    using namespace std;
    char animal[20] = "bear";    // animal
    const char * bird = "wren"; // bird ho
    char * ps;                   // uniniti

    cout << animal << " and "; // display
    cout << bird << "\n";      // display
    // cout << ps << "\n";      //may disp

    cout << "Enter a kind of animal: ";
    cin >> animal;               // ok if i
    // cin >> ps; Too horrible a blunder t
    //                               point to allocated space

    ps = animal;                // set ps
    cout << ps << "!\n";        // ok, same
    cout << "Before using strcpy():\n";
    cout << animal << " at " << (int *) an
    cout << ps << " at " << (int *) ps <<

    ps = new char[strlen(animal) + 1]; //
    strcpy(ps, animal);               // copy st

```

"bear" string, just as you've initialized arrays new. It initializes a pointer-to-char to a string:

```
const char * bird = "wren"; // bird holds
```

Remember, "wren" actually represents the address of the string in memory. The `bird` pointer assigns the address of "wren" to the `bird` pointer in memory to hold all the quoted strings used in the program. (The compiler stores each stored string with its address.) This means that `bird` would use the string "wren", as in this example:

```
cout << "A concerned " << bird << " speak
```

String literals are constants, which is why you can't change them. Using `const` in this fashion means you can't change it. Chapter 7 takes up the topic of pointers. The pointer `ps` remains uninitialized, so it doesn't point to anything. It's usually a bad idea, and this example is no exception.

Next, the program illustrates that you can't use `bird` equivalently with `cout`. Both, after all, are pointers. The two strings ("bear" and "wren") stored at those addresses makes the error of attempting to display `ps`, yet the program displays garbage, and you might get a program crash, a bit like distributing a blank signed check: You

lems: You just use a sufficiently large `char` array to receive input or uninitialized pointers. To avoid these issues and use `std::string` objects instead.

Caution

When you read a string into a program-style string, you must use previously allocated memory. This address cannot be a pointer that has been initialized using `new`.

Next, notice what the following code accomplishes:

```
ps = animal;                // set ps to point to animal
...
cout << animal << " at " << (int *) animal << endl;
cout << ps << " at " << (int *) ps << endl;
```

It produces the following output:

```
fox at 0x0065fd30
fox at 0x0065fd30
```

Normally, if you give `cout` a pointer, it prints the address. If you give `cout` a `char *`, `cout` displays the pointed-to string. If you give `cout` a pointer to a pointer, you have to type cast the pointer to another pointer.

here by using `strlen()` to find the correct size:

```
fox at 0x0065fd30
fox at 0x004301c8
```

Also note that `new` located the new storage at the address that is the address of the array `animal`.

Often you encounter the need to place a string in memory when you initialize an array; otherwise, you use the `strcpy()` function; it works like this:

```
char food[20] = "carrots"; // initialize with "carrots"
strcpy(food, "flan");      // otherwise
```

Note that something like the following can happen if the string is smaller than the string:

```
strcpy(food, "a picnic basket filled with fruit");
```

In this case, the function copies the rest of the string immediately following the array, which can overwrite other data. To avoid that problem, you should use `strncpy()`, which takes a minimum number of characters to be copied. Be

many structures as a program needs during a program's execution. `new` is a tool to use. With it, you can create dynamic structures that are allocated during runtime, not at compile time. For dynamic structures, you are able to use the techniques you use for arrays and classes, too.

Using `new` with structures has two parts: creating a structure and returning pointers. To create a structure, you use the `new` operator to create an unnamed structure of the `inflatable` type and store its address in a pointer. You can use the following:

```
inflatable * ps = new inflatable;
```

This assigns to `ps` the address of a chunk of memory that is the size and structure of the `inflatable` type. Note that the syntax is the same as for the built-in types.

The tricky part is accessing members. When you create a structure, you use the dot membership operator with the structure name. All you have is its address. C++ provides the arrow membership operator (`->`). This operator uses the greater-than symbol, does for pointers to structures what the dot does for structure names. For example, if `ps` points to a type `inflatable`, `ps->price` is the `price` member of the pointed-to structure.

Sometimes new C++ users become confused as to how to use the arrow operator to specify a structure member. If the identifier is the name of a structure, use the dot operator. If the identifier is a structure, use the arrow operator.

A second, uglier approach to accessing structure members is to use a pointer to a structure, then `*ps` represents the structure. Because `*ps` is a structure, `(*ps).price` is the correct notation. The operator precedence rules require that you use parentheses.

Listing 4.21 uses `new` to create an unnamed structure. The notations for accessing structure members.

Listing 4.21 **newstrct.cpp**

```
// newstrct.cpp -- using new with a structure
#include <iostream>
struct inflatable // structure definition
{
    char name[20];
    float volume;
    double price;
};
int main()
{
    using namespace std;
```

An Example of Using new and delete

Let's look at an example that uses new and delete with the keyboard. Listing 4.22 defines a function getn that reads a string. This function reads the input into a large block of memory with an appropriate size to create a chunk of memory. After the function returns the pointer to the block. (This is a poor design for programs that read in a large number of strings, but it would be easier to use the string class, which handles this design.)

Suppose your program has to read 1,000 strings that are 1,000 characters long, but most of the strings are much shorter. If you try to hold the strings, you'd need 1,000 arrays of 1,000 characters. Much of that block of memory would wind up being unused. Instead, create an array of 1,000 pointers to char and then allocate memory for each string. That could save a lot of memory. Trying to use a large array for every string, you find that it's not feasible. You could also use new to find space to store only the strings that are needed. It's a little too ambitious for right now. Even using an array of pointers is ambitious for right now, but Listing 4.22 illustrates how delete works, the program uses it to free

```
char temp[80];           // temporary storage
cout << "Enter last name: ";
cin >> temp;
char * pn = new char[strlen(temp) + 1];
strcpy(pn, temp);        // copy string into new block

return pn;                // temp lost when function returns
}
```

Here is a sample run of the program in Listing 11.1.

```
Enter last name: Fredeldumpkin
Fredeldumpkin at 0x004326b8
Enter last name: Pook
Pook at 0x004301c8
```

Program Notes

Consider the function `getline()` in the program. The function reads the user's input word into the `temp` array. Next, it uses `strlen()` to determine the length of the word. Including the null character, the program then uses `new` to allocate a block to store the string, so that's the value given to `pn`. Finally, the program uses `strcpy()` to copy the word from `temp` to `pn`. The `getline()` uses the standard library function `getchar()` to read the characters from the new block. The function doesn't check to see if the user enters a null character.

the *free store* or *heap*. Data objects allocated in this way must be freed by the programmer, and we'll see how long they remain in existence. We'll take a look at a fourth form called *thread storage* that we'll discuss later.

Automatic Storage

Ordinary variables defined inside a function are called *automatic variables*. These terms mean that the variables are created when the function containing them is invoked, and they are destroyed when the function returns. For example, the `temp` array in Listing 4.22 exists only while the function `getname()` is executing. When program control returns to `main()`, the memory used by `temp` is freed. If `getname()` returned the address of `temp`, it would be pointing to a memory location that would soon be freed. To avoid this, we use `new` in `getname()`. Actually, automatic variables are stored on the stack. A *block* is a section of code enclosed between curly braces `{ }` in functions. But as you'll see in the next chapter, you can also define a variable inside one of those blocks. The scope of the variable is the block containing statements inside the block.

Automatic variables typically are stored on the stack. When execution enters a block of code, its variables are created and then are freed in reverse order when execution leaves the block. (This is called *last-in, first-out*, or *LIFO*, process.) So the stack grows

the life of a function. Using `new` and `delete` to manage memory is more complex than using `new` alone. Managing memory manually results in the part of memory in use always being the same. The difference between `new` and `delete` can leave holes in memory, making it difficult to allocate new memory requests more difficult.

Stacks, Heaps, and Memory Leaks

What happens if you *don't* call `delete` after `new`? The variable or construct created with the `new` operator? The variable or construct continues to persist if `delete` is not called, even if the pointer has been freed due to rules of scope. This is a way to access the construct on the free store even though it is gone. You have now created a *memory leak*. The memory is unusable through the life of the program; it's *leaked*. In extreme (though not uncommon) cases, memory leaks can use all the memory available to the application, causing it to crash. In addition, these leaks may negatively affect performance by running in the same memory space, causing the program to run slower. Even the best programmers and software companies make memory leaks. It's best to get into the habit of joining your `new` with `delete` for and entering the deletion of your construct. C++'s smart pointers (Chapter 15) help manage memory on the free store.

We can then access members using the member access operator:

```
s01.year = 1998;
```

We can create a pointer to such a structure:

```
antarctica_years_end * pa = &s02;
```

Provided the pointer has been set to a valid structure, we can use the membership operator to access members:

```
pa->year = 1999;
```

We can create arrays of structures:

```
antarctica_years_end trio[3]; // array of 3 structures
```

We then can use the membership operator to access members:

```
trio[0].year = 2003; // trio[0] is a structure
```

Here, `trio` is an array, but `trio[0]` is a structure. Because an array name is a pointer, we can use the membership operator:

```
(trio+1)->year = 2004; // same as trio[1].year
```

We can create an array of pointers:

```
const antarctica_years_end * arp[3] = {&s01, &s02, &s03};
```

operator.

```
std::cout << (*ppa)->year << std::endl;  
std::cout << (*(ppb+1))->year << std::endl;
```

Because `ppa` points to the first member of `arp`. So `(*ppa)->year` is the year member of the first element, `arp[0]`, which is correct. `ppb` points to the next element, `arp[1]`, which is also correct. For example, `*ppa->year` is the year of the first element, `arp[0]`, which is correct. `ppa->year`, which fails because the year member of `arp[0]` is not a pointer.

Is all this really true? Listing 4.23 incorporates a program.

Listing 4.23 **`mixtypes.cpp`**

```
// mixtypes.cpp -- some type combinations  
#include <iostream>  
  
struct antarctica_years_end  
{  
    int year;  
    /* some really interesting data, etc. */  
};  
  
int main()
```


The program compiles and works as promised.

Array Alternatives

Earlier this chapter mentioned the `vector` and `string` containers, as well as the built-in array. Let's take a brief look now at the advantages and benefits of using them.

The `vector` Template Class

The `vector` template class is similar to the `string` class. You can set the size of a `vector` object during run time, and you can add, delete, or insert new data in the middle. Basically, it's a dynamic array. Actually, the `vector` class does more than that, but it does so automatically.

At this time we won't venture very deeply into the details of the `vector` class. Instead, we'll look at a few basic practical matters. First, you need to include the `vector` header file. Second, the `vector` class is in the `std` namespace, so you can use a `using` directive, a `using` namespace directive, or a `std::` prefix.

The vector class has more capabilities than the built-in array at the cost of slightly less efficiency. If all you need is a simple array, it's better to use the built-in type. However, that has its drawbacks, such as lack of safety. C++11 responded to this situation by adding a new array type to the `std` namespace. Like the built-in array, it lives on the stack (or else static memory allocation) in order to avoid the inefficiency of built-in arrays. To this it adds convenience: to create an array object, you need to include the `array` header, just as you do for a vector:

```
#include <array>
...
using namespace std;
array<int, 5> ai;    // create array object
array<double, 4> ad = {1.2, 2.1, 3.43, 4.5};
```

More general, the following declaration creates an array of *typeName*:

```
array<typeName, n_elem> arr;
```

Unlike the case for vector, *n_elem* can't be a constant expression.

With C++11, you can use list-initialization for arrays, an option that was not an option with C++98 vector.

```

    a2[3] = 1.0/3.0;
// C++11 -- create and initialize array object
    array<double, 4> a3 = {3.14, 2.72, 1.62, 1.62};
    array<double, 4> a4;
    a4 = a3;      // valid for array object
// use array notation
    cout << "a1[2]: " << a1[2] << " at " << &a1[2] << endl;
    cout << "a2[2]: " << a2[2] << " at " << &a2[2] << endl;
    cout << "a3[2]: " << a3[2] << " at " << &a3[2] << endl;
    cout << "a4[2]: " << a4[2] << " at " << &a4[2] << endl;
// misdeed
    a1[-2] = 20.2;
    cout << "a1[-2]: " << a1[-2] << " at " << &a1[-2] << endl;
    cout << "a3[2]: " << a3[2] << " at " << &a3[2] << endl;
    cout << "a4[2]: " << a4[2] << " at " << &a4[2] << endl;
    return 0;
}

```

Here's some sample output:

```

a1[2]: 3.6 at 0x28cce8
a2[2]: 0.142857 at 0xca0328
a3[2]: 1.62 at 0x28ccc8
a4[2]: 1.62 at 0x28cca8

```

a4, and other compilers might make yet other unsafe behavior of built-in arrays.

Do the vector and array objects protect them. That is, you still can write unsafe code,

```
a2[-2] = .5;    // still allowed  
a3[200] = 1.4;
```

However, you have alternatives. One is using `getline()` member function with objects of the vector class. Another is using `at()` member function with objects of the vector class.

```
a2.at(1) = 2.3; // assign 2.3 to a2[1]
```

The difference between using bracket notation and `at()` is that if you use `at()`, an invalid index is caught due to range checking and the program aborts. This added checking does come at the cost of speed. C++ gives you the option of using either notation. The `at()` member function is a safer way of using objects that reduce the chances of errors. The `begin()` and `end()` member functions are used to iterate over the elements of a container. They are not accidentally exceeding the bounds. But we'll

string from one location to another. When using `cstring` or the `string.h` header file.

The C++ `string` class, supported by the `string` header file, provides more user-friendly means to deal with strings. It automatically resizes to accommodate stored strings, and it can copy a string.

The `&` operator lets you request memory address. The operator returns the address of the variable. The only means to access data object is a simple variable, you can use the `*` operator to access the value. If the data object is an array, you can use the `[]` operator to access the elements. If the data object is a structure, you can use the `->` operator to access structure members.

Pointers and arrays are closely connected. In C, `ar[i]` is interpreted as `*(ar + i)`, with the `ar` pointing to the first element of the array. Thus, the array name can be used as a pointer name with array notation to access array elements.

The `new` and `delete` operators let you explicitly allocate and deallocate memory. They are used to allocate memory dynamically and when they are returned to the memory pool. Variables declared within a function, and static variables declared with the keyword `static`, are less flexible.

Question 3 to the variable `even`.

5. Write a statement that displays the value of `ideas`.
6. Declare an array of `char` and initialize it.
7. Declare a `string` object and initialize it.
8. Devise a structure declaration that describes a person's name, kind, the weight in whole ounces, and the number of ideas.
9. Declare a variable of the type defined in 8.
10. Use `enum` to define a type called `Response` with values `Maybe`, `Yes`, and `No`. `Yes` should be 1, `No` should be 0.
11. Suppose `ted` is a `double` variable. Declare a pointer to `ted` and use the pointer to display `ted`'s value.
12. Suppose `treacle` is an array of 10 `float` values. Declare a pointer to an element of `treacle` and use the pointer to display the value.
13. Write a code fragment that asks the user for the number of elements in a dynamic array of that many `int` values and creates a `vector` object.

lowing example of output:

```
What is your first name? Betty Sue
What is your last name? Yewe
What letter grade do you deserve? B
What is your age? 22
Name: Yewe, Betty Sue
Grade: C
Age: 22
```

Note that the program should be able to handle more than one word. Also note that the program should handle one letter. Assume that the user requests a letter grade. Don't worry about the gap between a D and a C.

2. Rewrite Listing 4.4, using the C++ `string` class.
3. Write a program that asks the user to enter their first name, and that then constructs, stores, and outputs the user's last name followed by a comma, and a space. Use the functions from the `cstring` header file.

```
Enter your first name: Flip
```

```
Enter your last name: Fleming
```

```
Here's the information in a single s
```

7. William Wingate runs a pizza-analysis program that uses the following information:
- The name of the pizza company.
 - The diameter of the pizza
 - The weight of the pizza

Devise a structure that can hold this information. Write a program that declares a structure variable of that type. The program should prompt the user for the preceding items of information, and then store the information in the structure. Use `cin` (or its methods) and `cout` (or its methods).

8. Do Programming Exercise 7 but use `new` to allocate the structure variable. Also have the program prompt the user for the pizza company name.
9. Do Programming Exercise 6, but instead of using arrays and structures, use `new` to allocate the array.
10. Write a program that requests the user's name, the number of meters (40-meter, if you prefer) and then displays the data in a structure object to hold the data. (Use a built-in structure type.)

- Compound statements (blocks)
- The comma operator
- Relational operators: `>`, `>=`, `==`, `<=`, `<`, and `!=`
- The `while` loop
- The `typedef` facility
- The `do while` loop
- The `get()` character input method
- The end-of-file condition
- Nested loops and two-dimensional arrays

Computers do more than store data. They can calculate, simulate, model, predict, analyze, classify, and extrapolate, synthesize, and otherwise manipulate data. And while we may want to use computers to store, analyze, and trash data, but we'll try to steer clear of the manipulative miracles, programs need tools for making decisions. Of course, C++ provides such tools. `while` loops, `do while` loops, `if` statements, `switch` statements, and so on. As the book says, "C++ employs, so if you know C, you can zip through the material on 'Statements and Logical Operators.'" (But don't forget `cin` handles character input!) These various prepositional expressions and logical expressions to go with the

```
        return 0;  
    }
```

Here is the output from the program in Listing 1.1:

```
C++ knows loops.  
C++ knows loops.  
C++ knows loops.  
C++ knows loops.  
C++ knows loops.  
C++ knows when to stop.
```

This loop begins by setting the integer `i` to 0:

```
i = 0
```

This is the *loop initialization* part of the loop. Then the program checks whether `i` is less than 5:

```
i < 5
```

If it is, the program executes the following:

```
cout << "C++ knows loops.\n";
```

Then the program uses the *loop update* part:

```
i++
```

The C++ loop design positions these elements in parentheses. Each part is an expression, and each other. The statement following the control is executed as long as the test expression returns true.

```
for (initialization; test-expression; update-expression)  
    body
```

C++ syntax counts a complete for statement as a compound statement. If you want to incorporate one or more statements in the body, the statement requires using a compound statement, or a block.

The loop performs initialization just once, setting a variable to a starting value and then use the *test-expression* to determine whether the loop continues. The *test-expression* is a relational expression—that is, it compares the value of *i* to 5, checking whether the program executes the loop body. Actually, the loop body executes only if the expression evaluates to true/false comparisons. You can use any expression that evaluates to a boolean value. Thus, an expression with a value of 0 is false, and the loop terminates. If the expression evaluates to true, and the loop continues. Listing 5.2 demonstrates the test condition. (In the update section, *i* is incremented by 1 each time it's used.)

```
i = i - 1  
Done now that i = 0
```

Note that the loop terminates when `i` reaches 0.

How do relational expressions, such as `i < 0`, evaluate? Before the `bool` type was added to C++, the value of the expression `i < 0` was 1 if true and 0 if false. Thus, the value of the expression `5 < 5` was 0. Now that C++ has added the `bool` type, the expressions evaluate to the `bool` literals `true` and `false` instead of 1 and 0. This is a change to incompatibilities, however, because a C++ program can use 0 where integer values are expected, and it can use `true` and `false` where `bool` values are expected.

The `for` loop is an *entry-condition* loop. This means that the test condition is evaluated *before* each loop cycle. The loop never executes if the test condition is false. For example, suppose you rerun the program with the starting value 0. Because the test condition fails the very first time, the loop never gets executed:

```
Enter the starting countdown value: 0  
Done now that i = 0
```

This look-before-you-loop attitude can be contrasted with the *update-expression* which is evaluated at the end of each loop cycle. Typically, it's used to increase or decrease the loop counter.

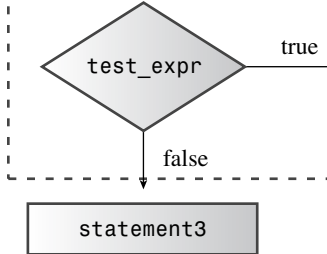


Figure 5.1 The c

A `for` statement looks something like a `function` by paired parentheses. However, `for`'s status as a statement prevents the compiler from thinking `for` is a function. It also prevents

Tip

Common C++ style is to place a space between `for` and the opening parenthesis. However, some programmers omit space between a function name and the opening parenthesis.

```
for (i = 6; i < 10; i++)  
    smart_function(i);
```

The expression `cooks = 4` has the value 4, just because C++ permits this behavior doesn't mean the same rule that makes this peculiar statement possible:

```
x = y = z = 0;
```

This is a fast way to set several variables to 0. Appendix D, “Operator Precedence”) reveals that first 0 is assigned to `z`, and then `z = 0` is assigned to `y`, and finally `y = 0` is assigned to `x`.

Finally, as mentioned previously, relational expressions have `bool` values `true` or `false`. The short program `express.cpp` prints the values of expressions. The `<<` operator has higher precedence than the `<` operator, so the code uses parentheses to ensure the correct order of evaluation.

Listing 5.3 **express.cpp**

```
// express.cpp -- values of expressions
#include <iostream>
int main()
{
    using namespace std;
    int x;

    x = 100;

    cout << "The expression x = 100 has the value " << x << endl;
```

Normally, `cout` converts `bool` values to 1 or 0. The `cout.setf(ios::boolalpha)` function call switches `cout` to use the words `true` and `false` instead of 1 and 0.

Note

A C++ expression is a value or a combination of values. Every expression has a value.

To evaluate the expression `x = 100`, C++ must perform two very different acts. The first act is to evaluate the expression `100`. Every act of evaluating an expression changes the state of the program. Evaluation has a *side effect*. Thus, evaluating an assignment expression has the side effect of changing the assignee's value. You might think of this from the standpoint of how C++ is constructed. The expression `100` has no side effect. Not all expressions have side effects. For example, the expression `x = 100` has a side effect. It creates a new value, but it doesn't change the value of `x`. The expression `x++` has a side effect because it involves incrementing `x`.

From expression to statement is a short step. The following is an expression:

```
age = 100
```

Whereas the following is a statement:

```
age = 100;
```


code such as the following invalid:

```
eggs = int toad * 1000;    // invalid, not  
cin >> int toad;          // can't combine
```

Similarly, you can't assign a for loop to a variable. A for loop is not an expression, so it has no value and cannot be assigned.

```
int fx = for (i = 0; i < 4; i++)  
    cout >> i;    // not possible
```

Bending the Rules

C++ adds a feature to C loops that requires some new syntax. This was the original syntax:

```
for (expression; expression; expression)  
    statement
```

In particular, the control section of a for statement was defined earlier in this chapter, separated by semicolons, like the following, however:

```
for (int i = 0; i < 5; i++)
```

That is, you can declare a variable in the initialization section, which is convenient, but it doesn't fit the original syntax. This once-outlaw behavior was originally accepted because it was so useful.

Another thing you should know is that so earlier rule and treat the preceding loop as if it available after the loop terminates.

Back to the for Loop

Let's be a bit more ambitious with loops. List the first 16 factorials. Factorials, which are handy in the following way. Zero factorial, written as $0!$, is defined as 1. Then, $1!$ is $1 * 0!$, or 1. Then, $2!$ is $2 * 1!$, or 2. Then, $3!$ is $3 * 2!$, or 6, and so on. In other words, $n!$ is the product of that integer with the preceding factorial. Victor Borge's best-known monologues featuring the exclamation mark is pronounced something like "factorial." However, in this case, "!" is pronounced "factorial." We'll calculate the values of successive factorials, storing them in an array, and display the results. Also the program introduces

Listing 5.4 `formore.cpp`

```
// formore.cpp -- more looping with for
#include <iostream>
const int ArSize = 16;          // example of
int main()
{
    long long factorials[ArSize];
```

```
13! = 6227020800  
14! = 87178291200  
15! = 1307674368000
```

Factorials get big fast!

Note

This listing uses the `long long` type. If your compiler can use `double`. However, the integer format numbers grow larger.

Program Notes

The program in Listing 5.4 creates an array to hold factorials. The first element of the array is 1!, and so on. Because the first two elements of the `factorials` array to hold the first two factorials (the first element of the array has an index value of 0.) After that, the program calculates the product of the index with the previous factorial value and stores the result in the array. The loop counter as a variable in the body of the loop.

The program in Listing 5.4 demonstrates how to use arrays by providing a convenient means to access array elements. `formore.cpp` uses `const` to create a symbolic constant for the array size. Then it uses `ArSize` wherever the array size is needed.

so we'll practice using them next.

Also this example reminds us that we can make selected standard names available.

Changing the Step Size

So far the loop examples in this chapter have one in each cycle. You can change that by changing the update expression. Listing 5.5, for example, increases the loop counter by more than one. Instead of using `i++` as the update expression, it uses `i += step`, where `step` is a user-selected step size.

Listing 5.5 `bigstep.cpp`

```
// bigstep.cpp -- count as directed
#include <iostream>
int main()
{
    int i = 0;
    int step = 1;

    using std::cout;    // a using declaration
    using std::cin;
    using std::endl;
    cout << "Enter an integer: ";
    int by;
```

value 100.

Finally, this example illustrates the use of us

Inside Strings with the `for` Loop

The `for` loop provides a direct way to access c. ple, Listing 5.6 enables you to enter a string and character, in reverse order. You could use either in this example because both allow you to use ters in a string; Listing 5.6 uses a `string` class yields the number of characters in the string; t expression to set `i` to the index of the last cha character. To count backward, the program use the array subscript by one in each loop. Also I relational operator (`>=`) to test whether the lo marize all the relational operators soon.

Listing 5.6 **forstr1.cpp**

```
// forstr1.cpp -- using for with a string
#include <iostream>
#include <string>
int main()
{
```

The Increment (++) and Decrement (--) Operators

C++ features several operators that are frequently used. In this section, we will examine them now. You've already seen two: the name C++, and the decrement operator --. Two of the most commonly used and increasingly common loop operations: increasing and decreasing a counter. However, there's more to their story than you've seen. There are two varieties. The *prefix* version comes before the operand, as in ++x. The *postfix* version comes after the operand, as in x++. The two varieties are used interchangeably, but they differ in terms of when they take place. The prefix version takes place in advance or afterward; both methods have to do with the timing of when the money gets added. Listing 5.7 shows the increment operator.

Listing 5.7 **plus_one.cpp**

```
// plus_one.cpp -- the increment operator
#include <iostream>

int main()
{
    using std::cout;
    int a = 20;
    int b = 20;
```

Using the increment and decrement operators is the common task of increasing or decreasing a variable.

The increment and decrement operators are postfix, meaning they come away from the variable and increment or decrement the same variable. The problem is that the use-then-change and change-then-use are ambiguous. That is, a statement such as the following is ambiguous in concurrent systems:

```
x = 2 * x++ * (3 - ++x);    // don't do it
```

C++ does not define correct behavior for this statement.

Side Effects and Sequence Points

Let's take a closer look at what C++ does and doesn't do when operators take effect. First, recall that a *side effect* is an operation that an expression modifies something, such as a variable or memory location. A *sequence point* in program execution at which all side effects must be complete before going on to the next step. In C++ the semicolon is a sequence point. That means all changes made by assignment or increment and decrement operators in a statement must take place before the statement. Some operators that we'll discuss in Chapter 10 have a side effect at the end of any full expression is a sequence point.

expression is the entire assignment statement. At this point, so all that C++ guarantees is that `x` will be updated before the program moves to the following statement. This is the only statement implemented after each subexpression is evaluated. This is the only statement evaluated, which is why you should avoid statements like

C++11 documentation has dropped the `volatile` keyword, which doesn't carry over well when discussing multiple threads. These threads are framed in terms of sequencing, with some events happening before other events. This descriptive approach is to provide language that can more clearly handle

Prefixing Versus Postfixing

Clearly, whether you use the prefix or postfix increment or decrement for some purpose, such as a function argument, is a matter of style. The value of an increment or decrement expression

```
x++;
```

and

```
++x;
```

are different from one another? Or are

```
for (n = lim; n > 0; --n)
```


The Increment/Decrement Operator

You can use increment operators with pointers. Adding an increment operator to a pointer increments the type it points to. The same rule holds for decrement.

```
double arr[5] = {21.1, 32.8, 23.4, 45.2, 30.9};
double *pt = arr; // pt points to arr[0],
++pt;             // pt points to arr[1],
```

You can also use these operators to change the value of a variable. When used in conjunction with the `*` operator, Applesoft asks questions of what gets dereferenced and what gets incremented or decremented. This is determined by the placement and precedence of the increment or decrement operator. The increment, decrement, and dereferencing operators all have the same precedence, which is higher than the prefix precedence, which is higher than the postfix precedence. The postfix increment and decrement operators are right to left.

The right-to-left association rule for prefix operators applies to the increment and decrement operators. The prefix increment operator applies `++` to `pt` (because the `++` is to the right of the dereferencing operator) to the value of `pt`:

```
double x = *++pt; // increment pointer, then dereference
```

the first member of an array, ++pt changes p

Combination Assignment Operator

Listing 5.5 uses the following expression to u

```
i = i + by
```

C++ has a combined addition and assignment operator to produce the same result more concisely:

```
i += by
```

The += operator adds the values of its two operands. The first operand is the left operand on the left. This implies that the left operand can assign a value, such as a variable, an array, or a pointer. The right operand identifies by dereferencing a pointer:

```
int k = 5;
k += 3;                                // ok, k set to 8
int *pa = new int[10];                // pa points to array of 10 ints
pa[4] = 12;
pa[4] += 6;                            // ok, pa[4] set to 18
*(pa + 4) += 7;                        // ok, pa[4] set to 25
pa += 2;                               // ok, pa points to 6th element
34 += 10;                             // quite wrong
```

to use paired braces to construct a *compound statement*. Braces enclose a block of statements, and the statements they enclose and, for each statement, the statements that follow it. For example, the program in Listing 5.8 groups several statements into a single block. This enables the program to read five numbers from the user's input, and do a calculation. The program calculates the sum of the five numbers entered, and this provides a natural occasion for

Listing 5.8 **block.cpp**

```
// block.cpp -- use a block statement
#include <iostream>
int main()
{
    using namespace std;
    cout << "The Amazing Accounto will sum\n";
    cout << "five numbers for you.\n";
    cout << "Please enter five values:\n";
    double number;
    double sum = 0.0;
    for (int i = 1; i <= 5; i++)
    {
        cout << "Value " << i << ": ";
        cin >> number;
        sum += number;
    }
```

```
        cin >> number;
        sum += number;
cout << "Five exquisite choices indeed! "
```

The compiler ignores indentation, so only
Thus, the loop would print the five prompts
pletes, the program moves to the following li

Compound statements have another inter
inside a block, the variable persists only as lon
within the block. When execution leaves the
the variable is known only within the block:

```
#include <iostream>
int main()
{
    using namespace std;
    int x = 20;
    {                                // block star
        int y = 100;
        cout << x << endl;  // ok
        cout << y << endl;  // ok
    }                                // block ends
    cout << x << endl;      // ok
    cout << y << endl;      // invalid, w
    return 0;
```

As you have seen, a block enables you to sneak two expressions into just one expression. The C++ syntax allows just one statement. The comma operator, enabling you to sneak two expressions into just one expression. For example, suppose you have a variable that increases one each cycle and a second variable decreases one each cycle. The update part of a `for` loop control section would normally contain just one expression there. The solution is to use the comma operator to sneak two expressions into one:

```
++j, --i    // two expressions count as one
```

The comma is not always a comma operator. In a list of names, the comma serves to separate adjacent names in a list. For example, in the following code, the comma is a separator here, not an operator.

```
int i, j;    // comma is a separator here, not an operator
```

Listing 5.9 uses the comma operator twice. The first time, it sneaks two expressions into one: the expression that increments the string class object. (You could also write the expression that increments the length of the word would be limited by your compiler. The second time, it sneaks two expressions into one: the expression that displays the contents of an array in reverse order. The program in Listing 5.9 sneaks two expressions into one.

Here is a sample run of the program in Listing 15-10:

```
Enter a word: stressed  
desserts  
Done
```

By the way, the `string` class offers more operations. I will leave those for Chapter 16, “The `string` Class.”

Program Notes

Look at the `for` control section of the program. The `++` operator to squeeze two initializations into one line. Then it uses the comma operator again. The `++` expression for the last part of the control section.

Next, look at the body. The program uses `strrev` to reverse the string in a single unit. In the body, the program reverses the array with the last element. Then it increments the pointer to refer to the next-to-the-first element and the program swaps those elements. Note that the program stops when it reaches the center of the array. If it were to continue, it would begin swapping the switched elements back to their original positions.

Another thing to note is the location for `word`. The code declares `i` and `j` before the loop because of the comma operator. That's because declarations also pose—separating items in a list. You can use a comma to create and initialize two variables, but it's a bit awkward:

```
int j = 0, i = word.size() - 1;
```

In this case the comma is just a list separator. The declaration declares and initializes both `j` and `i`. However,

Incidentally, you can declare `temp` inside the loop:

```
int temp = word[i];
```

This may result in `temp` being allocated and deallocated a bit slower than declaring `temp` once before the loop is finished, `temp` is discarded if it's declared inside the loop.

Comma Operator Tidbits

By far the most common use for the comma operator is to combine two expressions into a single `for` loop expression. But C++ does have other interesting properties. First, it guarantees that the first expression is evaluated before the second.

Computers are more than relentless number
pare values, and this capability is the foundation
relational operators embody this ability. C++
numbers. Because characters are represented
operators with characters, too. They don't wo
with `string` class objects. Each relational exp
comparison is true and to the `bool` value `false`
tors are well suited for use in a loop test expr
relational expressions to 1 and false relational
operators.

Table 5.2 **Relational Operators**

Operator
<code><<</code>
<code><=</code>
<code>==</code>
<code>></code>
<code>>=</code>
<code>!=</code>

Assignment, Comparison, and a M

Don't confuse testing the is-equal-to operator. The expression asks the musical question "Is music

```
musicians == 4    // comparison
```

The expression has the value `true` or `false`.
`musicians:`

```
musicians = 4     // assignment
```

The whole expression, in this case, has the v

The flexible design of the `for` loop creates a common error: programmers accidentally drop an equals sign (=) from the test part of the loop, instead of a relational expression for the test part of the loop. That's because you can use any valid C++ expression in the test part. Remember that nonzero values test as `true`, and the expression `assigns 4 to musicians` has the value 4 and is `true`. In a language, such as Pascal or BASIC, that uses `=` to assign, this is more prone to this slip.

Listing 5.10 shows a situation in which you attempt to examine an array of quiz scores and find out if there are not 20. It shows a loop that correctly uses comparison and assignment in the test condition. The program

```
    return 0;  
}
```

Because the program in Listing 5.10 has a bug, we need to add a check about it to actually running it. Here is some sample output:

Doing it right:

quiz 0 is a 20

quiz 1 is a 20

quiz 2 is a 20

quiz 3 is a 20

quiz 4 is a 20

Doing it dangerously wrong:

quiz 0 is a 20

quiz 1 is a 20

quiz 2 is a 20

quiz 3 is a 20

quiz 4 is a 20

quiz 5 is a 20

quiz 6 is a 20

quiz 7 is a 20

quiz 8 is a 20

quiz 9 is a 20

quiz 10 is a 20

Caution

Don't use `=` to compare for equality; use `==`.

Like C, C++ grants you more freedom than at the cost of requiring greater responsibility. Overplanning prevents a program from going beyond its intended scope. However, with C++ classes, you can design a program that prevents a lot of nonsense. Chapter 13, “Class Inheritance,” explains how to build the protection into your programs when you use classes. Listing 5.10 should include a test that keeps it from going even for the “good” loop. If all the scores were within the array bounds. In short, the loop needs to test the index. Chapter 6 shows how to use logical operators to simplify the condition.

Comparing C-Style Strings

Suppose you want to see if a string in a character array matches a string constant, the following test might not do what you want:

```
word == "mate"
```

Remember that the name of an array is a string constant. The address of the first element of a string constant is a synonym for its address. The

```
char little[6] = "Daffy";           // 5 letters plus
```

By the way, although you can't use relational operators to compare characters because characters are actually integers, the following is valid code, at least for the ASCII and Unicode characters of the alphabet:

```
for (ch = 'a'; ch <= 'z'; ch++)  
    cout << ch;
```

The program in Listing 5.11 uses `strcmp()` in the `test` program displays a word, changes its first letter, displays it again until `strcmp()` determines that word is the same as the original. The program includes the `cstring` file because it provides a function for comparing strings.

Listing 5.11 **compstr1.cpp**

```
// compstr1.cpp -- comparing strings using arrays  
#include <iostream>  
#include <cstring>           // prototype for strcmp()  
int main()  
{  
    using namespace std;  
    char word[5] = "?ate";
```

Kate

late

After loop ends, word is mate

Program Notes

The program in Listing 5.11 has some interesting behavior. It wants the loop to continue as long as `word` is not `mate`. It continues as long as `strcmp()` says the two strings are not equal. That is this:

```
strcmp(word, "mate") != 0    // strings are not equal
```

This statement has the value 1 (`true`) if the strings are not equal. But what about `strcmp(word, "mate")`? It has the value 0 (`false`) if the strings are equal and the value 1 (`true`) if the strings are unequal and the value -1 (`false`) if the function returns `true` if the strings are different. You can use just the function instead of the whole statement. It has the same behavior and involves less typing. Also it is the way that is traditionally used `strcmp()`.

Testing for Equality or Order

You can use `strcmp()` to test C-style strings for equality. It is true if `str1` and `str2` are identical:

```
strcmp(str1, str2) == 0
```

type `char` really is an integer type, so the operation is stored in the variable. Also note that using an individual character in a string:

```
word[0] = ch;
```

Comparing string Class Strings

Life is a bit simpler if you use `string` class strings. The `string` class design allows you to use relational operators with strings, possible because one can define class functions that overload the operators. In 12, “Classes and Dynamic Memory Allocation,” we looked into class designs, but from a practical standpoint, you can use the relational operators with `string` objects. You can use a `string` object instead of an array of characters.

Listing 5.12 `compstr2.cpp`

```
// compstr2.cpp -- comparing strings using string class
#include <iostream>
#include <string>      // string class
int main()
{
    using namespace std;
    string word = "?ate";
```

relational test expression, or as an aggregate object to extract individual characters.

As you can see, you can achieve the same results with string objects, but programming with string objects is more cumbersome.

Finally, unlike most of the `for` loops you have seen, `while` loops aren't counting loops. That is, they don't execute a fixed number of times. Instead, each of these loops watches a condition (e.g., `"mate"`) to signal that it's time to stop. More to come on this second kind of test, so let's examine that first.

The while Loop

The `while` loop is a `for` loop stripped of the counting part, leaving only the test condition and a body:

```
while (test-condition)
    body
```

First, a program evaluates the parenthesized test condition. If the condition evaluates to `true`, the program executes the body of the loop. If the condition evaluates to `false`, the program skips the body and finishes with the body, the program returns to the test condition. If the condition is nonzero, the program executes the body of the loop. The execution continues until the test condition evaluates to `false`.

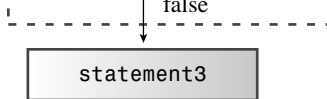


Figure 5.3 The structure of a while loop

Listing 5.13 puts a while loop to work. The program reads a string and displays the character and its ASCII value until it reaches the null character. This technique of stepping through a string until reaching the null character is a standard C++ technique. Because a string contains its own termination information, the program does not need information about how long a string is.

Listing 5.13 **while.cpp**

```
// while.cpp -- introducing the while loop
#include <iostream>
const int ArSize = 20;
int main()
{
    using namespace std;
    char name[ArSize];
```


they do add an endearing technoid tone to the

Program Notes

The `while` condition in Listing 5.13 looks like

```
while (name[i] != '\0')
```

It tests whether a particular character in the string is non-null. If the test fails, the loop eventually succeeds, the loop body needs to change `i` to the next character, and the loop increments `i` at the end of the loop body. Omitting the increment would cause the program to print the same array element, printing the character and the address of the character in memory. Getting such an infinite loop is one of the most common mistakes in programming. Often you can cause it when you forget to update a loop variable.

You can rewrite the `while` line this way:

```
while (name[i])
```

With this change, the program works just as well. When `name[i]` is an ordinary character, its value is the character's ASCII value. When `name[i]` is the null character, its character value is 0. This is more concise (and more commonly used) but it is not necessarily faster. Dumb compilers might produce faster code for the first version, but smart compilers produce the same code for both.

```
        update-expression;  
    }
```

Similarly, the `while` loop

```
while (test-expression)  
    body
```

could be rewritten this way:

```
for ( ; test-expression ; )  
    body
```

This `for` loop requires three expressions (by two expressions), but they can be empty and semicolons are mandatory. Incidentally, a missing `test-expression` is interpreted as `true`, so this loop runs forever:

```
for ( ; ; )  
    body
```

Because `for` loops and `while` loops are new, it's a matter of style. There are three differences. One difference is that the condition in a `for` loop is interpreted as `true` if it is missing. Another difference is that the first statement in a `for` loop to declare a variable is interpreted as `true` if it is missing. Finally, there is a slight difference in the `do` statement, which is discussed in Chapter 6. Typ-

```

i = 0;
while (name[i] != '\0')
    cout << name[i] << endl;
    i++;
cout << "Done\n";

```

The indentation tells you that the program a of the loop body. The absence of braces, ho consists solely of the first `cout` statement. Thu of the array indefinitely. The program never r side the loop.

The following example shows another potenti

```

i = 0;
while (name[i] != '\0');    // problem
{
    cout << name[i] << endl;
    i++;
}
cout << "Done\n";

```

This time the code gets the braces right, but a semicolon terminates a statement, so this words, the body of the loop is a *null statement* the material in braces now comes *after* the lo cycles, doing nothing forever. Beware the stra

unsigned long on others, and perhaps some

But the `ctime` header file (`time.h` on less to these problems. First, it defines a symbolic number of system time units per second. So seconds. Or you can multiply seconds by `CLOCKS_PER_SEC`. Second, `ctime` establishes `clock_t` as an alias “Type Aliases,” later in this chapter.) This means `clock_t`, and the compiler converts it to long type for your system.

Listing 5.14 shows how to use `clock()` and

Listing 5.14 **waiting.cpp**

```
// waiting.cpp -- using clock() in a time
#include <iostream>
#include <ctime> // describes clock() fun
int main()
{
    using namespace std;
    cout << "Enter the delay time, in sec
    float secs;
    cin >> secs;
    clock_t delay = secs * CLOCKS_PER_SEC
```

```
typedef char byte; // makes byte an alias for char
```

Here's the general form:

```
typedef typeName aliasName;
```

In other words, if you want `aliasName` to be an alias for `typeName` as if it were a variable of that type, you can use the `typedef` keyword. For example, to make `byte_pointer` an alias for `byte`, you could declare `byte_pointer` as a pointer-to-`byte`:

```
typedef char * byte_pointer; // pointer to byte
```

You could try something similar with `#define` and aliases for variables. For example, consider the following:

```
#define FLOAT_POINTER float *  
FLOAT_POINTER pa, pb;
```

Preprocessor substitution converts the declaration to:

```
float * pa, pb; // pa a pointer to float, pb a float
```

The `typedef` approach doesn't have that problem. The `typedef` approach makes using `typedef` a better choice than the `#define` approach as the only choice.

Notice that `typedef` doesn't create a new type. If you make `word` an alias for `int`, `cout` treats `word` as an `int`.

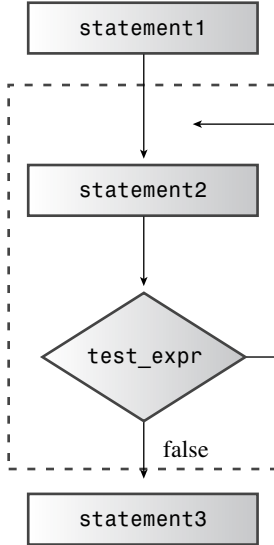


Figure 5.4 The struc

Usually, an entry-condition loop is a better choice than a do-while loop because the entry-condition loop checks before looping.

```
}
```

Here's a sample run of the program in Listing 10.1:

```
Enter numbers in the range 1-10 to find my favorite number: 
```

```
9
```

```
4
```

```
7
```

```
Yes, 7 is my favorite.
```

Strange for loops

It's not terribly common, but you may occasionally see a `for(;;)` loop, sometimes called a "forever" loop.

```
int I = 0;
for(;;) // sometimes called a "forever" loop
{
    I++;
    // do something ...
    if (30 >= I) break;    // if statement
}
```

Or here is another variation:

```
int I = 0;
for(;;I++)
{
    // do something ...
}
```

The Range-Based for Loop (C++11)

The C++11 adds a new form of loop called the *range-based for* loop—that of doing something with each element, generally, of one of the container classes, such as `vector`.

```
double prices[5] = {4.99, 10.99, 6.87, 7.99, 8.49};  
for (double x : prices)  
    cout << x << std::endl;
```

Here `x` initially represents the first member of the `prices` array. For the first element, the loop then cycles `x` to represent the remaining elements. In return, so this code would print all five members, one per line. Every value included in the range of the array.

To modify array values, you need a different syntax for the range-based for loop:

```
for (double &x : prices)  
    x = x * 0.80;           //20% off sale
```

The `&` symbol identifies `x` as a reference variable, a tool used in the book “Adventures in Functions.” The significance here is that the subsequent code to modify the array contents, whereas the previous code only read the values.

The range-based for loop also can be used with `initializer lists`:

```
for (int x : {3, 5, 2, 8, 6})  
    cout << x << " ";  
cout << '\n';
```


example, Listing 5.16 stops reading input when the user enters the # character. The test program counts the number of characters it reads and then displays the characters that have been read. (Pressing the # character stops the program; programs have to check for the # character. Typically, the operating system handles the input character. The test program echoes the input.) Listing 5.16 shows the total number of characters processed. Listing 5.16

Listing 5.16 **textin1.cpp**

```
// textin1.cpp -- reading chars with a while loop
#include <iostream>
int main()
{
    using namespace std;
    char ch;
    int count = 0;           // use basic input
    cout << "Enter characters; enter # to stop\n";
    cin >> ch;               // get a character
    while (ch != '#')        // test the character
    {
        cout << ch;          // echo the character
        ++count;            // count the character
        cin >> ch;           // get the next character
    }
}
```

plays the character, increments the count, and
vital. Without it, the loop repeatedly processes
last step, the program advances to the next character.
Note that the loop design follows the guideline that the loop
terminates the loop is if the last character read is not
a character before the loop starts. The condition is checked
the end of the loop.

This all sounds reasonable. So why does the program not work?
cin. When reading type char values, just as with scanf, it skips
spaces and newline characters. The spaces in the input are not
counted.

To further complicate things, the input to the program is not
you type don't get sent to the program until you press the Enter
type characters after the # when running the program. When you
Enter, the whole sequence of characters is sent to the program.
Processing the input after it reaches the # character.

cin.get(char) to the Rescue

Usually, programs that read input character-by-character use
including spaces, tabs, and newlines. The istream class, which
cin belongs, includes member functions that read characters.
The function cin.get(ch) reads the next character from the input.

```
}
```

Here is a sample run of the program in Listing 1.1:

```
Enter characters; enter # to quit:
```

```
Did you use a #2 pencil?
```

```
Did you use a
```

```
14 characters read
```

Now the program echoes and counts every character entered, so it is still possible to type more input.

If you are familiar with C, this program may look odd. The `cin.get(ch)` call places a value in the `ch` variable. In C you must pass the address of a variable, not the value of that variable. But the call to `cin.get(ch)` in C++ is fine. In C, code like this won't work. In C++ it can work because the argument is a *reference*. The reference type is supported by the `iostream` header file declares the argument to `get` as a reference. This way the function can alter the value of its argument. You can't do this in C, while, the C mavens among you can relax; or you can use `cin.get(&ch)` as it does in C. For `cin.get(ch)`, however, it

the same function with a single argument of `ch` because the language supports an OOP feature called function overloading that allows you to create different functions with the same name, as long as they have different argument lists. If, for example, you are using C++, the compiler finds the version of `cin.get()` that takes no arguments. But if you use `cin.get(ch)`, the compiler finds the version that takes a `char` argument. And if the code provides a version of `cin.get()` that takes no arguments.

Function overloading enables you to use the same function to perform the same basic task in different ways or to perform different tasks. We’re awaiting you in Chapter 8. Meanwhile, you can see more examples by using the `get()` examples that come with the C++ standard library. In the different function versions, we’ll include a comment. Thus, `cin.get()` means the version that takes no arguments, and `cin.get(ch)` means the version that takes one argument.

The End-of-File Condition

As Listing 5.17 shows, using a symbol such as `EOF` is not very satisfactory because such a symbol might be used for something other than the end-of-file condition. Other arbitrarily chosen symbols, such as `@` and `END`, can employ a much more powerful technique—`enum`.

GNU C++ for the PC recognize Ctrl+Z which requires a subsequent Enter. In short, many PCs treat Ctrl+Z as a simulated EOF, but the exact details vary (on a line, Enter key required or not required).

If your programming environment can test for EOF, as shown in Listing 5.17 with redirected files and you can simulate the EOF. That sounds useful, so let's see how.

When `cin` detects the EOF, it sets two bits in the `ios_base` member function named `eof()` to see whether the stream has reached EOF. It returns the `bool` value `true` if the EOF has been reached; otherwise, it returns `false`. The `fail()` member function returns `true` if the stream has reached EOF or if the last operation failed (e.g., if the last read operation returned 0 or -1 and `false` otherwise. Note that the `eof()` and `fail()` member functions report the most recent attempt to read; that is, they report the result of the last read operation. The `cin.eof()` or `cin.fail()` test should always be used to check for EOF. Listing 5.18 reflects this fact. It uses `fail()` instead of `eof()` because it appears to work with a broader range of implementations.

Note

Some systems do not support simulated EOF. If you have been using `cin.get()` to read characters, that won't work here because detecting the EOF is not supported.

Here is sample output from the program i

```
The green bird sings in the winter.<ENTER>
```

```
The green bird sings in the winter.
```

```
Yes, but the crow flies in the dawn.<ENTE
```

```
Yes, but the crow flies in the dawn.
```

```
<CTRL>+<Z><ENTER>
```

```
73 characters read
```

Because I ran the program on a Windows
simulate the EOF condition. Unix and Linux
in Unix and Unix-like systems, including Lin
of the program; the `fg` command lets executi

By using redirection, you can use the prog
report how many characters it has. This time,
characters from a two-line file on a Unix syst

```
$ textin3 < stuff
```

```
I am a Unix file. I am proud
```

```
to be a Unix file.
```

```
48 characters read
```

```
$
```

operator, which toggles true to false and vice versa. The test to look like this:

```
while (!cin.fail())    // while input has
```

The return value for the `cin.get(char)` member function of the `istream` class provides a function that can convert a `bool` value; this conversion function is called `boolalpha`. It is expected, such as in the test condition of a `while` loop, the conversion is `true` if the last attempted read operation means you can rewrite the `while` test to look like this:

```
while (cin)    // while input is successful
```

This is a bit more general than using `!cin.fail()`, as it covers other possible causes of failure, such as disk full or a bad file.

Finally, because the return value of `cin.get` can be formatted to this format:

```
while (cin.get(ch))    // while input is successful
{
    ...                // do stuff
}
```

Here, `cin.get(char)` is called once in the loop and once at the end of the loop.

```
cout.put(ch);
```

It works much like C's `putchar()`, except instead of type `int`.

Note

Originally, the `put()` member had the single `int` argument, which would then be type cast to `char`. However, some C++ implementations have `put(signed char)`, and `put(unsigned char)`. These implementations generate an error message when you pass an `int` to `put()`. An explicit type cast is a good choice for converting the `int`. An explicit type cast is required for `int` types.

To use `cin.get()` successfully, you need to know when the function reaches the EOF. When the function reaches the EOF, there are no more characters to read. `cin.get()` returns a special value, represented by the constant `EOF`, which is defined in the `iostream` header file. The `EOF` character value so that the program won't continue to read. `EOF` is defined as the value `-1` because no character can have that value. You can see the heart of Listing 5.18 looks like this:

You should realize that EOF does not represent a character, so that there are no more characters.

There's a subtle but important point about EOF so far. Because EOF represents a value outside the range of the char type, it might not be compatible with the char type. If char is signed, so a char variable could never have the value of EOF. If char is unsigned, you use cin.get() (with no argument) and then test for EOF to type int instead of to type char. Also if you use char, you might have to do a type cast to char when testing for EOF.

Listing 5.19 incorporates the cin.get() approach and also condenses the code by combining character

Listing 5.19 **textin4.cpp**

```
// textin4.cpp -- reading chars with cin.get()
#include <iostream>
int main(void)
{
    using namespace std;
    int ch;                                // store char
    int count = 0;

    while ((ch = cin.get()) != EOF) // test for EOF
```

Let's analyze the loop condition:

```
while ((ch = cin.get()) != EOF)
```

The parentheses that enclose the subexpression `cin.get()` evaluate that expression first. To do the evaluation, the `cin.get()` function. Next, it assigns the function's return value to `ch`. The assignment statement is the value of the left-hand side of the assignment to the value of `ch`. If this value is `EOF`, the loop condition fails. The loop test condition needs all the parentheses. Suppose

```
while (ch = cin.get() != EOF)
```

The `!=` operator has higher precedence than the assignment operator. The `cin.get()`'s return value to `EOF`. A comparison of `ch` to `EOF` is converted to 0 or 1, and that's the value of the expression.

Using `cin.get(ch)` (with an argument) for `ch` avoids any type problems. Remember that the `cin.get()` returns the value to `ch` at the `EOF`. In fact, it doesn't assign the value to `ch`. It's called on to hold a non-char value. Table 5.3 shows the difference between `cin.get(char)` and `cin.get()`.

`cout.put()` methods of `iostream`. You just globally replace `getchar()` and `putchar()` with the old code uses a type `int` variable for input. Your implementation has multiple prototypes

Nested Loops and Two-Dimensional Arrays

Earlier in this chapter you saw that the `for` loop can be nested. Now let's go a step further and look at how a two-dimensional array serves to handle two-dimensional arrays.

First, let's examine what a two-dimensional array is. Arrays that are termed *one-dimensional arrays* because they store a single row of data. You can visualize a two-dimensional array as a grid of rows and columns of data. You can use a two-dimensional array to store quarterly sales figures for six separate districts, or you can use a two-dimensional array to represent a computerized game board.

C++ doesn't provide a special two-dimensional array type for which each element is itself an array. For example, to store maximum temperature data for five cities over a 4-day period, you can use an array as follows:

```
int maxtemps[4][5];
```

The expression `maxtemps[0]` is the first element of `maxtemps`. The expression `maxtemps[0]` is itself an array of five `ints`. The expression `maxtemps[0][0]`, and this element is a single `int`. You can think of the second subscript as representing the column.

```
int maxtemps[4][5];
```

The `maxtemps` array viewed as a table:

		0	1
<code>maxtemps[0]</code>	0	<code>maxtemps[0][0]</code>	<code>maxtemps[0][1]</code>
<code>maxtemps[1]</code>	1	<code>maxtemps[1][0]</code>	<code>maxtemps[1][1]</code>
<code>maxtemps[2]</code>	2	<code>maxtemps[2][0]</code>	<code>maxtemps[2][1]</code>
<code>maxtemps[3]</code>	3	<code>maxtemps[3][0]</code>	<code>maxtemps[3][1]</code>

Figure 5.6 Accessing array elements

For a two-dimensional array, each element is accessed by using a form like that in the previous example. The initialization of `maxtemps` consists of a comma-separated series of one-dimensional arrays, each enclosed in a set of braces:

```
int maxtemps[4][5] =    // 2-D array
{
    {96, 100, 87, 101, 105},    // values for city 0
    {96, 98, 91, 107, 104},    // values for city 1
    {97, 101, 93, 108, 107},    // values for city 2
    {98, 103, 95, 109, 108}    // values for city 3
};
```

You can visualize `maxtemps` as four rows of data. The initialization `{96, 103, 101}` initializes the first row, representing the maximum temperature for city 0. By placing each row of data on its own line, if possible, you can make the code more readable.

Using a Two-Dimensional Array

Listing 5.20 incorporates an initialized two-dimensional array into a program. This time the program reverses the order of the loops (city index) on the outside and the row loop on the inside.

```

int maxtemps[Years][Cities] =    // 2-
{
    {96, 100, 87, 101, 105},    // val
    {96, 98, 91, 107, 104},    // val
    {97, 101, 93, 108, 107},    // val
    {98, 103, 95, 109, 108}    // val
};

cout << "Maximum temperatures for 200
for (int city = 0; city < Cities; ++c
{
    cout << cities[city] << ":\t";
    for (int year = 0; year < Years;
        cout << maxtemps[year][city]
    cout << endl;
}
    // cin.get();
return 0;
}

```

```
};
```

This approach limits each of the five strings of pointers stores the addresses of the five string copies each of the five string literals to the column array of pointers is much more economical in memory. To modify any of the strings, the two-dimensional array is enough, both choices use the same initialization to play the strings.

Also you could use an array of `string` class objects to store the string data. The declaration would look like this:

```
const string cities[Cities] = // array of strings
{
    "Gribble City",
    "Gribbletown",
    "New Gribble",
    "San Gribble",
    "Gribble Vista"
};
```

provides several ways to do this. If `ch` is a `type`,
reads the next input character into `ch`:

```
cin >> ch;
```

However, it skips over spaces, newlines, and
reads the next input character, regardless of it

```
cin.get(ch);
```

The member function call `cin.get()` returns
spaces, newlines, and tabs, so it can be used as

```
ch = cin.get();
```

The `cin.get(char)` member function calls
by returning a value with the `bool` conversion
function call reports the EOF by returning the
`istream` file.

A nested loop is a loop within a loop. Nested
two-dimensional arrays.

5. What would the following code fragment

```
int k = 8;  
do  
    cout <<" k = " << k << endl;  
while (k++ < 5);
```

6. Write a for loop that prints the values
a counting variable by a factor of two in
7. How do you make a loop body include
8. Is the following statement valid? If not,

```
int x = (1,024);
```

What about the following?

```
int y;  
y = 1,024;
```

9. How does `cin>>ch` differ from `cin.get`
input?

- Cleo earns 5% of \$100 the first year, gets \$105, or \$5.25, and so on. Write a program that calculates the value of Cleo's investment to exceed \$1000 and displays the value of both investments at that time.
5. You sell the book *C++ for Fools*. Write a program that tracks the number of monthly sales (in terms of number of copies sold). Use a loop to prompt you by month, using either integer or string objects, if you prefer) initialized to zero. Store the sales data in an array of `int`. Then, the program should calculate the contents and report the total sales for the year.
 6. Do Programming Exercise 5 but use a `double` array for years of monthly sales. Report the total sales for the combined years.
 7. Design a structure called `car` that holds information about a mobile: its make, as a string in a character array; when it was built, as an integer. Write a program that uses `new` to log. The program should then use `new` to create an array of structures. Next, it should prompt the user for the make (of more than one word) and year information. Note that this requires some care because it alternates

You should include the `cstring` header to make the comparison test.

9. Write a program that matches the description of Exercise 8, but use a `string` class object instead of a C string. Include the `string` header file and use a relational operator to compare the strings.
10. Write a program using nested loops that displays a pattern of asterisks. The number of rows to display. It should then display one asterisk in the first row, two in the second row, and so on. The number of asterisks are preceded by the number of spaces. The total number of characters displayed should be a total number of characters equal to the number of rows squared. The output would look like this:

Enter number of rows: 5

```
.....*  
....**  
...***  
..****  
.*****  
*****
```

- The ctype library of character functions
- The conditional operator: `?:`
- The `switch` statement
- The `continue` and `break` statements
- Number-reading loops
- Basic file input/output

One of the keys to designing intelligent programs is making decisions. Chapter 5, “Loops and Relational Operators,” discusses decision making—looping—in which a program decides whether to repeat an action. Chapter 6, “Decision Making,” investigates how C++ lets you use branch statements to make decisions. Which vampire-protection scheme (garlic, holy water, or a wooden stake) has the user selected? Did the user choose the `switch` statements to implement decisions, and Chapter 7, “The Conditional Operator,” shows how to use the conditional operator. Chapter 8, “Logical Operators,” also looks at the conditional operator, the `choice`, and the logical operators, which let you combine conditions. Chapter 9, “File Input/Output,” takes a first look at file input/output.

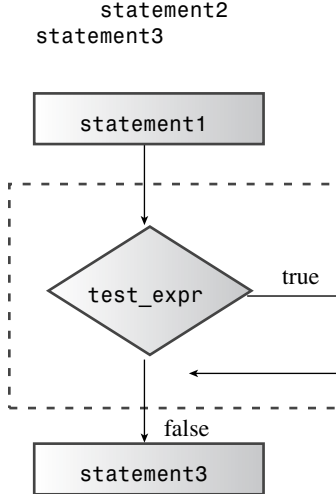


Figure 6.1 The struc

Most often, *test-condition* is a relational loops. Suppose, for example, that you want a p as well as the total number of characters. You

```
        cout << " characters total in sentence\n";  
        return 0;  
    }
```

Here's some sample output from the program:

```
The balloonist was an airhead  
with lofty goals.
```

6 spaces, 46 characters total in sentence

As the comments in Listing 6.1 indicate, the `if` statement is executed only when `ch` is a space. Because it is outside the loop, it is executed in every loop cycle. Note that the text "The balloonist was an airhead with lofty goals." is generated by pressing Enter.

The `if else` Statement

Whereas an `if` statement lets a program decide whether a statement is executed, an `if else` statement lets a program decide which of two statements is executed. It's an invaluable statement for C++ because the C++ `if else` statement is modeled after similar statements in other languages. For example, if you have a Cookie card, you get a Cookie Plus Plus, else you get a Cookie. The `if else` statement has this general form:

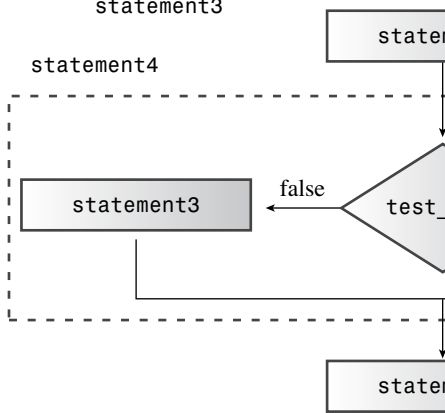


Figure 6.2 The structure

For example, suppose you want to alter inc keeping the newline character intact. In that output line of equal length. This means you w action for newline characters and a different c Listing 6.2 shows, `if else` makes this task eas fier, one of the alternatives to a using directive

}
Here's some sample output from the program:

Type, and I shall repeat.

An ineffable joy suffused me as I beheld

Bo!jofggbcmf!kpz!tvggvtfe!nf!bt!J!cfifme

the wonders of modern computing.

uif!xpofbst!pg!npefso!dpnqvujoh

Please excuse the slight confusion.

Note that one of the comments in Listing 10-10 produces an interesting effect. Can you deduce what it does? The program can explain what's happening. (Hint: Think about the `if` statement.)

Formatting `if else` Statements

Keep in mind that the two alternatives in an `if else` statement are single block statements. If you need more than one statement in a block, you must use braces to make the statements a block. The C compiler does not automatically consider everything between `if` and `else` to be a block. This can lead to a compiler error:

they enclose the statements. The preceding code is another:

```
if (ch == 'Z') {  
    zorro++;  
    cout << "Another Zorro candidate\n";  
}  
else {  
    dull++;  
    cout << "Not a Zorro candidate\n";  
}
```

The first form emphasizes the block structure. The second form more closely ties the blocks to the keywords. Both are consistent and should serve you well; however, the first is the employer with strong and specific views on the matter.

The if else if else Construct

Computer programs, like life, might present you with multiple situations. You can extend the C++ if else statement to handle the else should be followed by a single statement. If the else statement itself is a single statement, it can be written as follows:

Listing 6.3 `ifelseif.cpp`

```
// ifelseif.cpp -- using if else if else
#include <iostream>
const int Fave = 27;
int main()
{
    using namespace std;
    int n;

    cout << "Enter a number in the range
    cout << "my favorite number: ";
    do
    {
        cin >> n;
        if (n < Fave)
            cout << "Too low -- guess again\n";
        else if (n > Fave)
            cout << "Too high -- guess again\n";
        else
            cout << Fave << " is right!\n";
    } while (n != Fave);
    return 0;
}
```

```
if (myNumber = 3)
```

The compiler would simply assign the value 3 to `myNumber` and the program would run—a very common error, and a difficult one to debug. The compiler issues a warning, which you would be wise to heed. The `if` statement allows the compiler to find errors is much easier than debugging faulty results.

Logical Expressions

Often you must test for more than one condition. For example, if a variable is a lowercase letter, its value must be greater than 'a' and less than 'z'. Or, if you ask a user to respond with a **Y** or **N** as well as lowercase. To meet this kind of need, C++ provides operators to combine or modify existing expressions. The logical AND operator, written `&&`; and logical NOT, written `!`.

The Logical OR Operator: `||`

In English, the word *or* can indicate when one condition is sufficient for a statement. For example, you can go to the MegaMicro website or you can work for MegaMicro, Inc. The C++ equivalent is the logical OR operator. This operator combines two expressions into a single expression.

ance of C++11, the subexpression to the left of the `||` operator is evaluated first, then the subexpression to the right.) For example, con-

```
i++ < 6 || i == j
```

Suppose `i` originally has the value 10. By the time the subexpression `i++ < 6` has the value 11. Also C++ won't bother evaluating the subexpression on the right if the subexpression on the left is true, for it only takes a true or false value. (The semicolon and the `||` operator are short-circuit operators.)

Listing 6.4 uses the `||` operator in an `if` statement to check if a character is a lowercase version of a character. Also it uses the `std::string` type from Chapter 4, “Compound Types”) to spread a s-

Listing 6.4 `or.cpp`

```
// or.cpp -- using the logical OR operator
#include <iostream>
int main()
{
    using namespace std;
    cout << "This program may reformat your data.\n";
    cout << "and destroy all your data.\n";
    cout << "Do you wish to continue? <y/n> ";
    char ch;
```

The Logical AND Operator: &&

The logical AND operator, written `&&`, also combines two expressions. The resulting expression has the value `true` only if both expressions are `true`. Here are some examples:

```
5 == 5 && 4 == 4    // true because both expressions are true
5 == 3 && 4 == 4    // false because first expression is false
5 > 3 && 5 > 10      // false because second expression is false
5 > 8 && 5 < 10      // false because first expression is false
5 < 8 && 5 > 2       // true because both expressions are true
5 > 8 && 5 < 2       // false because both expressions are false
```

Because the `&&` has a lower precedence than most other operators, you should use parentheses in these expressions. Like the sequence point, so the left side is evaluated, and then the right side is evaluated. If the left side is false, then C++ doesn't bother evaluating the right side. This is how the `&&` operator works.

```

{
    using namespace std;
    float naaq[ArSize];
    cout << "Enter the NAAQs (New Age Awareness)
            << "of\nyour neighbors. Program
            << "when you make\n" << ArSize <<
            << "or enter a negative value.\n";

    int i = 0;
    float temp;
    cout << "First value: ";
    cin >> temp;
    while (i < ArSize && temp >= 0) // 2
    {
        naaq[i] = temp;
        ++i;
        if (i < ArSize) // row
        {
            cout << "Next value: ";
            cin >> temp; // so
        }
    }
    if (i == 0)
        cout << "No data--bye\n";
}

```

First value: 28

Next value: 72

Next value: 15

Next value: 6

Next value: 130

Next value: 145

Enter your NAAQ: 50

3 of your neighbors have greater awareness
the New Age than you do.

The second run terminates after a negative

Enter the NAAQs (New Age Awareness Quotient)
your neighbors. Program terminates when you
6 entries or enter a negative value.

First value: 123

Next value: 119

Next value: 4

Next value: 89

Next value: -1

Enter your NAAQ: 123.031

0 of your neighbors have greater awareness
the New Age than you do.

ber. Note that the loop reads another value in the array, only if there is still room left in the array.

After it gets data, the program uses an `if` statement to see if `data` was entered (that is, if the first entry was a negative value) and prints it if present.

Setting Up Ranges with `&&`

The `&&` operator also lets you set up a series of conditions, each a choice corresponding to a particular range of values. Listing 6.5 also shows a useful technique for handling a series of strings. A `string` variable can identify a single string by pointing to its first character. A `char*` can identify a series of strings. You simply use the address of an array element. Listing 6.6 uses the `qualify` array. For example, `qualify[1]` holds the address of the second element. The program can then use `qualify[1]` as it would any other string, either with `cout` or with `strlen()` or `strcmp()`. Using `qualify` protects from accidental alterations.


```
        else if (age >= 30 && age < 65)
            index = 2;
        else
            index = 3;

        cout << "You qualify for the " << qual
        return 0;
    }
```

Here is a sample run of the program in Listing 6.6:

```
Enter your age in years: 87
You qualify for the pie-throwing festival.
```

The entered age doesn't match any of the three ranges, so the program prints the default string and then prints the corresponding string.

Program Notes

In Listing 6.6, the expression `age > 17 && age < 35` selects values—that is, ages in the range 18–34. The `<` operator is used to exclude 35. To include 35 in its range, which is the case for the `>=` operator to include 35 in its range, which is the case for the `<=` operator. If `age < 50`, the value 35 would be missed by all three ranges. You should check that the ranges don't have holes.

The Logical NOT Operator: !

The `!` operator negates, or reverses the truth value of an expression. If `expression` is true, then `!expression` is false. If `expression` is false, or nonzero, then `!expression` is true. You can call the exclamation point *bang*, making `!x` “bang x”.

Usually you can more clearly express a relationship with `$x > 5$` than with `!(x <= 5)`.

```
if (!(x > 5))           // if (x <= 5)
```

But the `!` operator can be useful with functions that can be interpreted that way. For example, `strcmp(s1, s2)` returns 0 if the two C-style strings `s1` and `s2` are the same. This implies that `!strcmp(s1, s2)` is true if they are the same.

Listing 6.7 uses the technique of applying the `is_int()` function to screen numeric input for suitability to be assigned to an `int`. `is_int()`, which we’ll discuss further in a moment, returns true if the value is in the range of values that can be assigned to type `int`. The code `while(!is_int(num))` to reject values that do not.

Listing 6.7 `not.cpp`

```
// not.cpp -- using the not operator
#include <iostream>
#include <climits>
```

```
        return false;  
    }  
}
```

Here is a sample run of the program in Listing 6.7:

```
Yo, dude! Enter an integer value: 62341286  
Out of range -- please try again: -8000222  
Out of range -- please try again: 99999  
You've entered the integer 99999  
Bye
```

Program Notes

If you enter a too-large value to a program receiving integer input, the program simply truncates the value to fit, without warning. The program in Listing 6.7 avoids that by first reading the input as a `long`. The `long` type has more than enough precision to hold values much greater than `int`. Another choice for holding the input would be `long long`, knowing that it is wider than `int`.

The Boolean function `is_int()` uses the `std::numeric_limits::is_integer` (defined in the `climits` file) to determine whether its argument is within the range of an integer. If the value is `true`, it returns `true`; otherwise, it returns `false`.

The logical AND operator has a higher precedence than the OR operator, which gets you the result of this expression:

```
age > 30 && age < 45 || weight > 300
```

means the following:

```
(age > 30 && age < 45) || weight > 300
```

That is, one condition is that age be in the range 30 to 45, and that weight be greater than 300. The entire expression is true if either of these conditions are true.

You can, of course, use parentheses to tell the compiler what you mean. For example, suppose you want to use && to combine the conditions that age be 50 or weight be greater than 300 with the condition that donation be greater than 1,000. You have to enclose the OR part with parentheses:

```
(age > 50 || weight > 300) && donation > 1000
```

Otherwise, the compiler combines the weight condition with the age condition instead of with the donation condition.

Although the C++ operator precedence rules allow you to compound comparisons without using parentheses, it is a good idea to use parentheses to group the tests, whether or not the compiler requires it. Code that is easier to read, it doesn't force someone

Table 6.3 **Logical Operators: Alternative Representations**

Operator	Alternative Representation
<code>&&</code>	and
<code> </code>	or
<code>!</code>	not

The `cctype` Library of Character Classification Functions

C++ has inherited from C a handy package of character classification functions in the `cctype` header file (`cctype.h`, in the older C headers). The `isupper(ch)` function returns a nonzero value if `ch` is an uppercase letter or a macro `isalpha(ch)` returns a nonzero value if `ch` is an alphabetic character. Similarly, the `ispunct(ch)` function returns a nonzero value if `ch` is a punctuation character, such as a comma or period. (These functions return an `int` value, not a `bool`, but the usual `bool` conversions allow you to treat them as `bool`.)

Using these functions is more convenient than using the usual `bool` conversions. For example, here's how you might use AND and OR to test if `ch` is an alphabetic character:

```
if ((ch >= 'a' && ch <= 'z') || (ch >= 'A' && ch <= 'Z'))
```

```
char ch;
int whitespace = 0;
int digits = 0;
int chars = 0;
int punct = 0;
int others = 0;

cin.get(ch); // get first character
while (ch != '@') // test for end of input
{
    if(isalpha(ch)) // is it a letter?
        chars++;
    else if(isspace(ch)) // is it a space?
        whitespace++;
    else if(isdigit(ch)) // is it a digit?
        digits++;
    else if(ispunct(ch)) // is it a punctuation mark?
        punct++;
    else
        others++;
    cin.get(ch); // get next character
}
```

Function Name**Return Value**`isalnum()`

This function returns a letter or a digit).

`isalpha()`

This function returns

`isblank()`

This function returns tal tab.

`iscntrl()`

This function returns

`isdigit()`

This function returns

`isgraph()`

This function returns other than a space.

`islower()`

This function returns

`isprint()`

This function returns including a space.

`ispunct()`

This function returns character.

`isspace()`

This function returns space character (that return, horizontal tab

`isupper()`

This function returns

*expression*3. Here are two examples that show how the conditional operator works. In the first example, the expression `5 > 3` is true, so the expression evaluates to 10; otherwise, it evaluates to 12. In the second example, the expression `3 == 9` is false, so the expression evaluates to 25; otherwise, it evaluates to 18.

We can paraphrase the first example this way: if the expression `5 > 3` evaluates to true, it evaluates to 10; otherwise, it evaluates to 12. In real-world expressions would involve variables.

Listing 6.9 uses the conditional operator to find the larger of two integers.

Listing 6.9 **condit.cpp**

```
// condit.cpp -- using the conditional operator
#include <iostream>
int main()
{
    using namespace std;
    int a, b;
    cout << "Enter two integers: ";
    cin >> a >> b;
    cout << "The larger of " << a << " and " << b << " is ";
    int c = a > b ? a : b;    // c = a if a > b, else c = b
    cout << " is " << c << endl;
    return 0;
}
```

the following mild example shows:

```
const char x[2] [20] = {"Jason ", "at your  
const char * y = "Quillstone ";
```

```
for (int i = 0; i < 3; i++)  
    cout << ((i < 2)? !i ? x [i] : y : x[1
```

This is merely an obscure (but, by no means
strings in the following order:

Jason Quillstone at your service

In terms of readability, the conditional operator
and simple expression values:

```
x = (x > y) ? x : y;
```

If the code becomes more involved, it can
else statement.

The switch Statement

Suppose you create a screen menu that asks the user to
example, Cheap, Moderate, Expensive, Extravagant.
else if else sequence to handle five alternatives.

a very important way. Each C++ case label is
ary between choices. That is, after a program
sequentially executes all the statements follow
itly direct it otherwise. Execution does *not* au
execution stop at the end of a particular grou
statement. This causes execution to jump to t

Listing 6.10 shows how to use `switch` and
menu for executives. The program uses a sho
A `switch` statement then selects an action ba

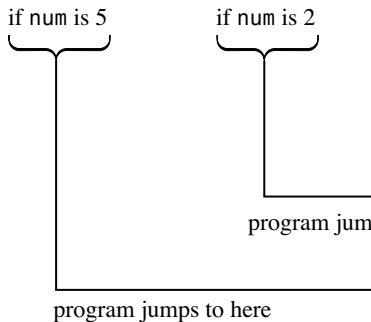


Figure 6.3 The structure

```

        case 2 :    report();
                  break;
        case 3 :    cout << "The boss
                  break;
        case 4 :    comfort();
                  break;
        default :    cout << "That's no
    }
    showmenu();
    cin >> choice;
}
cout << "Bye!\n";
return 0;
}

```

```

void showmenu()
{
    cout << "Please enter 1, 2, 3, 4, or 5
           "1) alarm           2) report\
           "3) alibi           4) comfort
           "5) quit\n";
}

void report()

```

```
3) alibi                4) comfort
5) quit
2
It's been an excellent week for business.
Sales are up 120%. Expenses are down 35%.
Please enter 1, 2, 3, 4, or 5:
1) alarm                2) report
3) alibi                4) comfort
5) quit
```

```
6
That's not a choice.
Please enter 1, 2, 3, 4, or 5:
1) alarm                2) report
3) alibi                4) comfort
5) quit
5
Bye!
```

The while loop terminates when the user enters a corresponding choice from the switch list, and the program terminates.

Note that input has to be an integer for the while loop to work. If you enter a letter, the input statement will fail, and the program will kill the program. To deal with those who don't follow the instructions, we can add a check for non-integer input.

```

                                break;
        case 'c':
        case 'C': comfort();
                                break;
        default : cout << "That's not a ch
    }
    showmenu();
    cin >> choice;
}

```

Because there is no `break` immediately following the `case 'C':` statement, the program continues on to the next line, which is the statement following the `default` case.

Using Enumerators as Labels

Listing 6.11 illustrates using `enum` to define a set of constants in a `switch` statement. In general, `enum` constants (because the compiler can't know how you will define them), so the `switch` statement compares the `int` value of the `enum` constant to `int`. Also the enumerators are just labels for the `switch` condition.

```
        case indigo : cout << "Her hair is indigo.\n";  
        }  
        cout << "Enter color code (0-6): ";  
        cin >> code;  
    }  
    cout << "Bye\n";  
    return 0;  
}
```

Here's sample output from the program in

```
Enter color code (0-6): 3  
Her nails were green.  
Enter color code (0-6): 5  
Her eyes were violet.  
Enter color code (0-6): 2  
Her shoes were yellow.  
Enter color code (0-6): 8  
Bye
```

switch and if else

Both the switch statement and the if else alternatives. The if else is the more versatile ranges, as in the following:

is to use `switch` if you have three or more al

The `break` and `continue`

The `break` and `continue` statements enable a program can use the `break` statement in a `switch` statement to pass to the next statement. The `continue` statement is used in loops and cause the loop and then start a new loop cycle (see

Listing 6.12 shows how the two statements work in text. The loop echoes each character and uses `break` if a period. This shows how you can use `break` when some condition becomes true. Next the program uses `continue` if the character is a space. The loop uses `continue` to skip over the character if it isn't a space.

Listing 6.12 `jump.cpp`

```
// jump.cpp -- using continue and break
#include <iostream>
const int ArSize = 80;
int main()
{
    using namespace std;
```

```
while (cin.get() != '\n')
{
    statement1;
    if (ch == '\n')
        continue;
    statement2;
}
statement3;
```

continue skips rest of loop body

```
while (cin.get() != '\n')
{
    statement1;
    if (ch == '\n')
        break;
    statement2;
}
statement3;
```

break skips rest of loop and exits loop

Figure 6.4 The structure of `continue` and `break`

However, the `continue` statement can make statements follow the `continue`. That way, you can use it as part of an `if` statement.

C++, like C, also has a `goto` statement. A statement location bearing the `goto` label:

```
goto paris;
```

That is, you can have code like this:

```
char ch;  
cin >> ch;  
if (ch == 'P')  
    goto paris;  
cout << ...  
...  
paris: cout << "You've just arrived at Paris";
```

In most circumstances (some would say in all), `goto` is a bad idea, and you should use structured controls, such as `if` and `while`, to control program flow.

look at a couple examples that illustrate these

Say you want to write a program that calculates the total weight of fish. There's a five-fish limit, so a five-element array is a reasonable choice. It's possible that you could catch fewer fish. Listing 6.13 shows a program that is full or if you enter non-numeric input.

Listing 6.13 **cinfish.cpp**

```
// cinfish.cpp -- non-numeric input terminates
#include <iostream>
const int Max = 5;
int main()
{
    using namespace std;
    // get data
    double fish[Max];
    cout << "Please enter the weights of " << Max << " fish.\n";
    cout << "You may enter up to " << Max << " fish <q to terminate>.\n";
    cout << "fish #1: ";
    int i = 0;
    while (i < Max && cin >> fish[i]) {
        if (++i < Max)
            cout << "fish #" << i+1 << " ";
```

```

        cin.get();    // read q
    }
    cin.get();        // read end of line a
    cin.get();        // wait for user to p

```

You also could use code similar to this in Listing 6.14 to read more input after exiting the loop.

Listing 6.14 further illustrates using the `cin` object.

The expression `cin >> fish[i]` in Listing 6.14 is a function call, and the function returns `cin`. If `cin` is part of a loop condition, the conversion value is `true` if input succeeded. If the expression terminates the loop. By the way, the variable `q` is a character.

```

Please enter the weights of your fish.
You may enter up to 5 fish <q to terminate>
fish #1: 30
fish #2: 35
fish #3: 25
fish #4: 40
fish #5: q
32.5 = average weight of 4 fish
Done.

```

Listing 6.14 **cingolf.cpp**

```
// cingolf.cpp -- non-numeric input skipping
#include <iostream>
const int Max = 5;
int main()
{
    using namespace std;
    // get data
    int golf[Max];
    cout << "Please enter your golf score\n";
    cout << "You must enter " << Max << " numbers\n";
    int i;
    for (i = 0; i < Max; i++)
    {
        cout << "round #" << i+1 << ": ";
        while (!(cin >> golf[i])) {
            cin.clear();      // reset input
            while (cin.get() != '\n')
                continue;    // get rid of bad input
            cout << "Please enter a number: ";
        }
    }
    // calculate average
```

The heart of the error-handling code in Listing 8.8 is

```
while (!(cin >> golf[i])) {
    cin.clear();          // reset input
    while (cin.get() != '\n')
        continue; // get rid of bad input
    cout << "Please enter a number: ";
}
```

If the user enters **88**, the `cin` expression is false. Furthermore, because `cin` is true, the expression `!(cin >> golf[i])` is true, and the inner loop terminates. But if the user enters **m**, the character `m` is placed into the array, the expression `!(cin >> golf[i])` is false, and the program enters the inner while loop. The first statement in the inner loop resets input. If you omit this statement, the program would not work. In the program, the program uses `cin.get()` in a while loop instead of `cin >>` at the end of the line. This gets rid of the bad input, along with the rest of the line. This approach is to read to the next whitespace, which is the same as the approach instead of one line at a time. Finally, the program

Let's see how this line of input is handled:

First, let's try type `char`:

```
char ch;  
cin >> ch;
```

The first character in the input line is assigned to `ch`, the digit 3, and the character code (in binary) for the digit 3. Since both the source (the input) and the destination are both characters, so no translation occurs. The character code for the digit 3 is the numeric value 3 that is stored; rather, it is the character code for the digit 3. In the next input statement, the digit character 8 is the next character examined by the next input statement.

Next, let's try the `int` type with the same input:

```
int n;  
cin >> n;
```

In this case, `cin` reads up to the first non-numeric character, the 8 digit, leaving the period as the next character. The program computes that these two characters correspond to the integer 38. The character for 38 is copied to `n`.

Next, let's try the `double` type:

```
double x;  
cin >> x;
```

On output, the opposite translations take place: sequences of digit characters, and floating-point digits and other characters (for example, 284.5) are translated to no translation.

The main point to this is that all the input relevant to console input is a text file—that is, a file in text code. Not all files are text files. For example, data files are in numeric forms—that is, in binary integer or floating-point forms. Processing files may contain text information, but they do not describe formatting, fonts, printers, and the like.

The file I/O discussed in this chapter parallels the I/O with text files. To create a text file for input, you use `ifstream` in Windows, or `vi` or `emacs` for Unix/Linux. You use `ofstream` to save the file in text format. The code editors `vi` and `emacs` indeed, the source code files are examples of text files. To look at files created with text output.

Writing to a Text File

For file output, C++ uses analogs to `cout`. So, in this section, we discuss basic facts about using `cout` for console output.

Note that although the `iostream` header is called `cout`, you have to declare your own object associating it with a file. Here's how you declare

```
ofstream outFile;           // outFile an ofstream  
ofstream fout;              // fout an ofstream
```

Here's how you can associate the objects with a file:

```
outFile.open("fish.txt");    // outFile use  
char filename[50];  
cin >> filename;            // user specifies  
fout.open(filename);        // fout used
```

Note that the `open()` method requires a C++ literal string or a string stored in an array.

Here's how you can use these objects:

```
double wt = 125.8;  
outFile << wt;               // write a number  
char line[81] = "Objects are closer than  
fout << line << endl;      // write a line of
```

The important point is that after you've declared an object with a file, you use it exactly as you would use


```
char automobile[50];
int year;
double a_price;
double d_price;

ofstream outFile;           // create
outFile.open("carinfo.txt"); // assume

cout << "Enter the make and model of a car: ";
cin.getline(automobile, 50);
cout << "Enter the model year: ";
cin >> year;
cout << "Enter the original asking price: ";
cin >> a_price;
d_price = 0.913 * a_price;

// display information on screen with cout

cout << fixed;
cout.precision(2);
cout.setf(ios_base::showpoint);
cout << "Make and model: " << automobile;
cout << "Year: " << year << endl;
```

```
Enter the original asking price: 13500
Make and model: Flitz Perky
Year: 2009
Was asking $13500.00
Now asking $12325.50
```

The screen output comes from using `cout`. If you `ch` contains the executable program, you should find a new may be in some other folder, depending on how the co the output generated by using `outFile`. If you open it the following contents:

```
Make and model: Flitz Perky
Year: 2009
Was asking $13500.00
Now asking $12325.50
```

As you can see, `outFile` sends precisely the same sequence of characters to the `carinfo.txt` file that `cout` sends to the display.

Program Notes

After the program in Listing 6.15 declares an `ofstream` object, it uses the `open` method to associate the object with a particular file:

```
ofstream outFile;           // create object f
outFile.open("carinfo.txt"); // associate with
```

tents. The contents are then replaced with the new contents, overriding this default behavior.

Caution

When you open an existing file for output, by default the file is truncated so the contents are lost.

It is possible that an attempt to open a file with a name having the requested name might already exist. A careful programmer would check to see if the file already exists before opening it. We will discuss this technique for this in the next example.

Reading from a Text File

Next, let's examine text file input. It's based on the `ifstream` class. So let's begin with a summary of those elements.

- You must include the `iostream` header file.
- The `iostream` header file defines an `ifstream` class.
- The `iostream` header file declares an `ifstream` class.
- You must account for the `std` namespace. You can use the `std::` prefix for elements such as `std::ifstream`.

- You can use an `ifstream` object with the `get()` method to read characters and with the `getline()` method to read lines.
- You can use an `ifstream` object with the `isatty()` method to test for the success of an input attempt.
- An `ifstream` object itself, when used as a Boolean value, returns the Boolean value `true` if the last read attempt was successful.

Note that although the `iostream` header defines two objects called `cin` and `cout`, you have to declare your own `ifstream` object and associate it with a file. Here's how you declare an `ifstream`:

```
ifstream inFile;           // inFile an ifstream object
ifstream fin;              // fin an ifstream object
```

Here's how you can associate them with particular files:

```
inFile.open("bowling.txt"); // inFile used to read from bowling.txt
char filename[50];
cin >> filename;           // user specifies filename
fin.open(filename);         // fin used to read from filename
```

Note that the `open()` method requires a C++-style literal string or a string stored in an array.

the `cstdlib` header file, which also defines `EXIT_SUCCESS` to communicate with the operating system. The

The `is_open()` method is relatively new to you can use the older `good()` method instead. `is_open()` checks quite as extensively as `is_open()` for problems.

The program in Listing 6.16 opens a file specified by the user, reports the number of values, their sum, and the average. The program is designed to design the input loop correctly, and the following code is in more detail. Notice that this program benefits from the

Listing 6.16 **sumafile.cpp**

```
// sumafile.cpp -- functions with an array
#include <iostream>
#include <fstream>           // file I/O support
#include <cstdlib>           // support for EXIT_SUCCESS
const int SIZE = 60;
int main()
{
    using namespace std;
    char filename[SIZE];
    ifstream inFile;         // object for
```

```

        cout << "Input terminated by data\n";
    else
        cout << "Input terminated for unknown reason\n";
    if (count == 0)
        cout << "No data processed.\n";
    else
    {
        cout << "Items read: " << count << endl;
        cout << "Sum: " << sum << endl;
        cout << "Average: " << sum / count << endl;
    }
    inFile.close();           // finished with file
    return 0;
}

```

To use the program in Listing 6.16, you first create a file containing the numbers. You can use a text editor, such as the one shown in Figure 6.16, to create this file. Let's assume that the file is named `data.txt` and its contents:

```

18 19 18.5 13.5 14
16 19.5 20 18 12 18.5
17.5

```

name in the character array filename. Then t
`inFile.open(filename);`

As discussed earlier in this chapter, it's vital succeeded. Here are a few of the things that n file might be located in another directory or f user might mistype the name or omit a file ex time trying to figure what's wrong with a file that the program didn't open the file. Testing misspent effort.

You need to pay close attention to the pro several things to test for when reading from a read past the EOF. The `eof()` method returns data ran into the EOF. Second, the program n instance, Listing 6.16 expects a file containing `true` if the most recent read attempt encount returns `true` if the EOF is encountered.) Fina for example, a corrupted file or a hardware fa most recent read attempt encountered such a tions individually, it's simpler to use the good when wrong:

immediately before the loop, just before the input statement at the end of the loop, just before

```
// standard file-reading loop design
inFile >> value;           // get first value
while (inFile.good())      // while input good
{
    // loop body goes here
    inFile >> value;        // get next value
}
```

You can condense this somewhat by using `inFile >> value` instead of `inFile.get()` and that `inFile`, when placed in the loop test, evaluates to `inFile.good()`—that is, the loop test becomes

```
inFile >> value
```

Thus, you can replace the two input statements with one, and the loop test with a single loop test. That is, you can replace the preceding code with

```
// abbreviated file-reading loop design
// omit pre-loop input
while (inFile >> value)    // read and test
{
    // loop body goes here
    // omit end-of-loop input
}
```


The conditional operator (`?:`) provides a compact way to write conditional expressions.

The `cctype` library of character functions provides a set of tools for analyzing character input.

Loops and selection statements are useful tools for writing console I/O. After you declare `ifstream` and `ofstream` files, you can use these objects in the same manner as `fstream`.

With C++'s loops and decision-making statements, you can write interesting, intelligent, and powerful programs that demonstrate the real powers of C++. Next, we'll look at functions.

Chapter Review

1. Consider the following two code fragments:

```
// Version 1
while (cin.get(ch))    // quit on eof
{
    if (ch == ' ')
        spaces++;
    if (ch == '\n')
        newlines++;
}

// Version 2
```

```

        cout << ch;
        ct1++;
        if (ch == '$')
            ct2++;
        cout << ch;
    }
    cout << "ct1 = " << ct1 << ", ct2 = " << ct2 << endl;
    return 0;
}

```

Suppose you provide the following input line:

Hi!

Send \$10 or \$20 now!

What is the output? (Recall that input is read in characters.)

4. Construct logical expressions to represent the following conditions.
 - a. weight is greater than or equal to 150.
 - b. ch is q or Q.
 - c. x is even but is not 26.
 - d. x is even but is not a multiple of 10.
 - e. donation is in the range 1,000–1,500.

```
f_grade++;
```

8. In Listing 6.10, what advantage would `t` and `c`, instead of numbers for the menu, have about what happens if the user types `q` or types `5` in either case.)
9. Consider the following code fragment:

```
int line = 0;
char ch;
while (cin.get(ch))
{
    if (ch == 'Q')
        break;
    if (ch != '\n')
        continue;
    line++;
}
```

Rewrite this code without using `break`

Please enter a c, p, t, or g. t
A maple is a tree.

4. When you join the Benevolent Order meetings by your real name, your job title, and your preference. Write a program that can list members by real name, job title, or preference. Base the program on the following code.

```
// Benevolent Order of Programmers  
struct bop {  
    char fullname[STRSIZE]; // real name  
    char title[STRSIZE];    // job title  
    char bopname[STRSIZE];  // secret name  
    int preference;         // 0 = none  
};
```

In the program, create a small array of members. Have the program run a loop that lets the user choose from the following alternatives:

- a. display by name b. display by job title
- c. display by bopname d. display by preference
- q. quit

5. The Kingdom of Neutronia, where the following income tax code:

First 5,000 tvarps: 0% tax

Next 10,000 tvarps: 10% tax

Next 20,000 tvarps: 15% tax

Tvarps after 35,000: 20% tax

For example, someone earning 38,000 tvarps would pay
 $0.10 \times 20,000 + 0.15 \times 3,000 + 0.20 \times 3,000$ tvarps in taxes.
Write a program that solicits incomes and to report taxes.
If the user enters a negative number or no

6. Put together a program that keeps track of donations for the Preservation of Rightful Influence. It should solicit the number of contributors and then solicit the contribution from each contributor. The information should be stored in a list of structures. Each structure should have a string object to store the name and a float object to store the contribution. After reading all the data, the program should report the amounts donated for all donors who contributed.

9. Do Programming Exercise 6 but modify it so that each item in the file should be the number of lines that consist of pairs of lines, with the first line being a name and the second line being a contribution.

4

Sam Stone

2000

Freida Flass

100500

Tammy Tubbs

5000

Rich Raptor

55000

- Designing functions to process arrays
- Using const pointer parameters
- Designing functions to process text strings
- Designing functions to process structures
- Designing functions to process objects
- Functions that call themselves (recursion)
- Pointers to functions

Fun is where you find it. Look closely, and you'll find it everywhere. C++ has a large library of useful functions (the standard library), but real programming pleasure comes from writing your own. On the other hand, real programming productivity comes from knowing what you can do with the STL and the BOOST C++ libraries. In this chapter, "Adventures in Functions," examine how to use functions, pass data to them, and retrieve information from them. Another chapter concentrates on how to use function templates and function structures. Finally, it touches on recursion and recursion. If you're a C++ expert, you'll find much of this chapter familiar. If you're a beginner, expertise. C++ has made several additions to the C language, and this chapter deals primarily with those. Meanwhile, let's a


```
// calling.cpp -- defining, prototyping, and using a function
#include <iostream>

void simple();    // function prototype

int main()
{
    using namespace std;
    cout << "main() will call the simple() function.\n";
    simple();      // function call
    cout << "main() is finished with the simple() function.\n";
    // cin.get();
    return 0;
}

// function definition
void simple()
{
    using namespace std;
    cout << "I'm but a simple function.\n";
}
```

```
}
```

Here *parameterList* specifies the types a to the function. This chapter more fully investigates the `return` statement marks the end of the function. Other than the opening brace. Type `void` functions correspond to old BASIC subprogram procedures. Functions perform some sort of action. For example, a function that prints `n` times could look like this:

```
void cheers(int n)           // no return value
{
    for (int i = 0; i < n; i++)
        std::cout << "Cheers! ";
    std::cout << std::endl;
}
```

The `int n` parameter list means that `cheers` takes `n` as an argument when you call this function.

A function with a return value produces a value when called it. In other words, if the function returns a value.

return value to a specified CPU register or memory location. Both the returning function and the caller agree on the type of data at that location. The caller knows what to expect, and the function definition tells the compiler (see Figure 7.1). Providing the same information in two places may seem like extra work, but it makes good sense. If you're at something from your desk at the office, you expect to find it right if you provide a description of what you're looking for at the office.

A function terminates after executing a return statement—for example, as alternative 1. A function terminates after it executes the first return statement. In the following example, the `else` isn't needed, but it stands the intent:

```
int bigger(int a, int b)
{
    if (a > b )
        return a; // if a > b, function returns a
    else
        return b; // otherwise, function returns b
}
```

(Usually, having multiple return statements is confusing, and some compilers might issue a warning enough to understand.)

Functions with return values are much like BASIC. They return a value to the calling program variable, display the value, or otherwise use it. The cube of a type double value:

```
double cube(double x)    // x times x times x
{
    return x * x * x;    // a type double value
}
```

For example, the function call `cube(1.2)` executes the return statement uses an expression. The function returns the value (1.728, in this case) and returns the value.

Prototyping and Calling a Function

By now you are familiar with making function prototypes with function prototyping because that's often the case. 7.2 shows the `cheers()` and `cube()` function prototypes.

```

        for (int i = 0; i < n; i++)
            cout << "Cheers! ";
        cout << endl;
    }

double cube(double x)
{
    return x * x * x;
}

```

The program in Listing 7.2 places a using statement at the top to make the members of the std namespace. Here's a sample run:

```

Cheers! Cheers! Cheers! Cheers! Cheers!
Give me a number: 5
A 5-foot cube has a volume of 125 cubic feet
Cheers! Cheers! Cheers! Cheers! Cheers! Cheers! Cheers! Cheers! Cheers! Cheers!

```

Note that main() calls the type void function cheer() with no arguments followed by a semicolon: cheer();. The cube() statement. But because cube() has a return value, it needs a return statement:

```
double volume = cube(side);
```

could only guess, and that is something compilers can't do.

Still, you might wonder, why does the compiler search further in the file and see how the functions are defined. The approach is that it is not very efficient. The compiler holds `main()` on hold while searching the rest of the file. The fact that the function might not even be in the file, or over several files, which you can compile individually. In this case, the compiler might not have access to the definition of `main()`. The same is true if the function is placed in a separate file. The function prototype is to place the function definition in a separate file, if possible. Also the C++ programming style is to place the function prototype in a header file, which provides the structure for the whole program.

Prototype Syntax

A function prototype is a statement, so it must be a valid statement. The way to get a prototype is to copy the function definition and add a semicolon. That's what the program in the next section does.

```
double cube(double x); // add ; to header
```

```
void say_bye(...); // C++ abdication
```

Normally this use of an ellipsis is needed only for a variable number of arguments, such as `printf()`.

What Prototypes Do for You

You've seen that prototypes help the compiler reduce the chances of program errors. In particular,

- The compiler correctly handles the function arguments.
- The compiler checks that you use the correct number of arguments.
- The compiler checks that you use the correct types, and converts the arguments to the correct type if necessary.

We've already discussed how to correctly handle function arguments. What happens when you use the wrong number of arguments? Let's make the following call:

```
double z = cube();
```

A compiler that doesn't use function prototypes, when it's called, it looks where the call to `cube()` should be, and if the value happens to be there. This is how C works.

verted correctly to a mere `int`. Some compilers are more strict, and `int` is an automatic conversion from a larger type.

Also prototyping results in type conversion. For example, convert an integer to a structure or a float.

Prototyping takes place during compile time. Without it, type checking, as you've just seen, catches many errors that would catch during runtime.

Function Arguments and Return Values

It's time to take a closer look at function arguments and return values. That means the numeric value of the argument is assigned to a new variable. For example, Listing 10.1 shows a function that calculates the volume of a cube.

```
double volume = cube(side);
```

Here `side` is a variable that, in the sample program, is the argument to `cube()`, recall, was this:

```
double cube(double x)
```

When this function is called, it creates a new variable and initializes it with the value 5. This insulates data in the function from `cube()` because `cube()` works with a copy of the data.

Variables, including parameters, declared within a function are called *local variables*. When a function is called, the computer allocates memory for these variables. When the function terminates, the computer deallocates the memory for these variables. (Some C++ literature refers to this as *creating and destroying variables*. That does make it sound like a fancy term, but it's just called *local variables* because they are localized to the function. This helps preserve data integrity. It also means that a variable called `x` in `main()` and another variable called `x` in some other function are two related variables, much as the Albany in California and the Albany in New York (see Figure 7.3). Such variables are also termed *local variables* and are created and deallocated automatically during program execution.

Multiple Arguments

A function can have more than one argument. You can pass multiple arguments with commas:

```
n_chars('R', 25);
```

This passes two arguments to the function.

```
    n    i    y
variables in main()
```

Figure 7.3

Similarly, when you define the function, you have to declare the variables in the function header:

```
void n_chars(char c, int n) // two arguments
```

This function header states that the function takes one type `char` argument and one type `int` argument. The parameters are declared to the function. If a function has two parameters, you declare the type of each parameter separately. You can't declare regular variables:

```
void fifi(float a, float b) // declare e
void fufu(float a, b)      // NOT accep
```

As with other functions, you just add a semicolon to the prototype:

```
void n_chars(char c, int n); // prototype
```

As with single arguments, you don't have to declare the type as in the definition, and you can omit the `char` argument:

```
void n_chars(char, int); // prototype
```

```

    {
        cout << "Enter an integer: ";
        cin >> times;
        n_chars(ch, times); // function wi
        cout << "\nEnter another character
            " q-key to quit: ";
            cin >> ch;
        }
        cout << "The value of times is " << ti
        cout << "Bye\n";
        return 0;
    }

void n_chars(char c, int n) // displays c
{
    while (n-- > 0)          // continue un
        cout << c;
}

```

The program in Listing 7.3 illustrates placing initializations rather than within the functions. Here

Enter a character: **W**

Enter an integer: **50**

```
while (n-- > 0)           // continue until  
    cout << c;
```

Notice that the program keeps count by decreasing `n`, a formal parameter from the argument list. This is a local variable in `main()`. The `while` loop then decrements `n`. Since `n` is a local variable, changing the value of `n` has no effect on the value of `n` of times in `main()`, the value of `n` in `main()` is still `n_chars()`.

Another Two-Argument Function

Let's create a more ambitious function—one that calculates the probability of winning. Also the function illustrates the use of local variables and arguments.

Many states in the United States now sponsor a state Lottery. The Lottery lets you pick a certain number of choices from a card having 51 numbers. You are required to pick six numbers from a card having 51 numbers at random. If your choice exactly matches the winning numbers, or so. Our function will calculate the probability of winning. The function that successfully predicts the winning numbers. C++, although powerful, has yet to implement a function that successfully predicts the winning numbers.

For example, compare

```
(10 * 9) / (2 * 1)
```

with

```
(10 / 2) * (9 / 1)
```

The first evaluates to $90 / 2$ and then to 45, then to 45. Both give the same answer, but the first gets a more accurate value (90) than does the second. The more accurate value gets. For large numbers, this strategy of alternating multiplication and division can prevent the calculation from overflowing the maximum value of the integer type.

Listing 7.4 incorporates this formula into a function that calculates the number of picks and the total number of combinations. The function uses the unsigned int type (unsigned, for short, is an integer type). Several integers can produce pretty large results, so the function uses unsigned integers for the function's return value. Also terms such as unsigned int are integer types.

Note

Some C++ implementations don't support type unsigned int. If you fall into that category, try ordinary double instead.

```

// the following function calculates the
// numbers correctly from numbers choices
long double probability(unsigned numbers,
{
    long double result = 1.0; // here co
    long double n;
    unsigned p;

    for (n = numbers, p = picks; p > 0; n
        result = result * n / p ;
    return result;
}

```

Here's a sample run of the program in Lis

Enter the total number of choices on the
the number of picks allowed:

49 6

You have one chance in 1.39838e+007 of wi

Next two numbers (q to quit): **51 6**

You have one chance in 1.80095e+007 of wi

Next two numbers (q to quit): **38 6**

You have one chance in 2.76068e+006 of wi

Next two numbers (q to quit): **q**

bye

corresponds to the number of cookies that people want. That's easy to find; you just use a loop to add them up. Finding the sum of an array of numbers is such a common task that it makes sense to have a function to do it. In fact, you won't have to write a new loop every time you want to find the sum of an array.

Let's consider what the function interface is. The function is called `sum_arr`, so the function name is `sum`, it should return the answer. If you keep your function simple, you can return the answer with a type `int` return value. So that the function can be used in a variety of contexts, you pass the array name as an argument. And to make the function more flexible, you restrict it to an array of a particular size, you pass the size as an argument. The only ingredient here is that you have to declare the function name. Let's see what that and the rest of the function looks like.

```
int sum_arr(int arr[], int n) // arr = array of n integers
```

This looks plausible. The brackets seem to indicate that the function takes an array of integers. The fact that the brackets are empty seems to indicate that the array can be of any size. But things are not always as they seem. The good news is that you can write the rest of the function. First, let's use an example to check that this approach works.

Listing 7.5 illustrates using a pointer as if it were an array. It initializes the array to some values and uses the `sum_arr` function to find the sum. The `sum_arr()` function uses `arr` as if it were an array of integers.

```
for (int i = 0; i < n; i++)  
    total = total + arr[i];  
return total;  
}
```

Here is the output of the program in Listing 7.5:

```
Total cookies eaten: 255
```

As you can see, the program works. Now let's look at how it works.

How Pointers Enable Array-Processing

The key to the program in Listing 7.5 is that it uses the name of an array as if it were a pointer. Recall that C++ interprets an array name as the address of the first element of the array.

```
cookies == &cookies[0] // array name is address of first element
```

(There are a few exceptions to this rule. First, you can use the `sizeof` operator to label the storage. Second, applying `sizeof` to an array name returns the address of the whole array, not the first element. Third, as mentioned in Chapter 6, the `&` operator applied to an array name returns the address of the whole array, not the first element. For example, `&cookies` returns the address of a 32-byte block of memory if `int` is 32 bits wide.)

of an array. Whether `arr` is a pointer or an array, the fourth element of the array. And it probably works for both of the following two identities:

```
arr[i] == *(ar + i)    // values in two notations
&arr[i] == ar + i      // addresses in two notations
```

Remember that adding one to a pointer, increments the value equal to the size, in bytes, of the type to which it points, and array subscription are two equivalent ways of accessing an element of an array.

The Implications of Using Arrays as Arguments

Let's look at the implications of Listing 7.5. The `sum_arr(Array, ArrSize)` passes the address of the first element of the array to the `sum_arr()` function. The `cookies` address to the pointer variable `arr` and `ArrSize` to the `n` variable. This means Listing 7.5 doesn't really pass the array to the function where the array is (the address), but rather how many elements it has (the `n` variable). (See Listing 7.5 for the function then uses the original array. If you pass a copy, it works with a copy. But if you pass an array, this difference doesn't violate C++'s pass-by-value rule.)

memory, but it has to spend time copying large arrays. Copying large arrays with the original data raises the possibility of a memory problem in classic C, but ANSI C and C++'s compilers soon see an example. But first, let's alter Listing 7.5 so that array functions operate. Listing 7.6 demonstrates how to use a value. It also shows how the pointer concept is more powerful than it may have appeared at first. To provide a more complete look, it looks like, the program uses the `std::` qualification for access to `cout` and `endl`.

Listing 7.6 **arrfun2.cpp**

```
// arrfun2.cpp -- functions with an array
#include <iostream>
const int ArSize = 8;
int sum_arr(int arr[], int n);
// use std:: instead of using directive
int main()
{
    int cookies[ArSize] = {1,2,4,8,16,32,64,128};
    // some systems require preceding int with std::
    // enable array initialization
```

```
        return total;  
    }  
}
```

Here's the output of the program in Listing 7.6:

```
003EF9FC = array address, 32 = sizeof cookies  
003EF9FC = arr, 4 = sizeof arr  
Total cookies eaten: 255  
003EF9FC = arr, 4 = sizeof arr  
First three eaters ate 7 cookies.  
003EFA0C = arr, 4 = sizeof arr  
Last four eaters ate 240 cookies.
```

Note that the address values and the array names are in hexadecimal. Also some implementations will display the values of `sizeof` in hexadecimal. Others will use hexadecimal for the addresses but decimal for the values.

Program Notes

Listing 7.6 illustrates some very interesting points. The expressions `sizeof cookies` and `sizeof arr` both evaluate to the same address. The value of `sizeof cookies` is 32, whereas `sizeof arr` is only 4. This is because `cookies` is the address of the whole array, whereas `arr` is the address of the first element. The execution takes place on a system that uses 4-byte integers.

```
void fillArray(int arr[], int size);
```

Don't try to pass the array size by using braces.

```
void fillArray(int arr[size]);
```

More Array Function Examples

When you choose to use an array to represent data, your design decisions should go beyond how data is stored and how data is used. Often you'll find it profitable to choose data operations. (The profits here include increased efficiency, and ease of debugging.) Also when you choose data operations when you think about a program, you're using an OOP mind-set; that, too, might prove profitable.

Let's examine a simple case. Suppose you want to store similar values of your real estate. (If necessary, suppose you decide what type to use. Certainly, `double` is less restrictive than `float` and provides enough significant digits to represent most real estate values.) Decide on the number of array elements. (We'll put off that decision, but let's keep things simple and assume five properties, so you can use an array of five elements.)

Now consider the possible operations you can perform on the array. Two very basic ones are reading values from the array and writing values to the array.

number of items to be read, and the function
example, if you use this function with an array
argument. If you then enter only three values,

You can use a loop to read successive values
the loop early? One way is to use a special value
property should have a negative value, you can
of input. Also the function should do something
ther input. Given these considerations, you can

```
int fill_array(double ar[], int limit)
{
    using namespace std;
    double temp;
    int i;
    for (i = 0; i < limit; i++)
    {
        cout << "Enter value #" << (i + 1) << " : ";
        cin >> temp;
        if (!cin)        // bad input
        {
            cin.clear();
            while (cin.get() != '\n')
                continue;
            cout << "Bad input; input processed\n";
        }
    }
}
```

arguments because C++ passes them by value. Functions that use an array work with the original array. The function is able to do its job. To keep a function from changing an array argument, you can use the keyword `const` (“constant Data”) when you declare the formal argument:

```
void show_array(const double ar[], int n)
```

The declaration states that the pointer `ar` can’t use `ar` to change the data. That is, you can’t change that value. Note that this doesn’t mean `show_array()` just means that you can’t use `ar` in the `show_array()` treats the array as read-only data. This restriction is by doing something like the following:

```
ar[0] += 10;
```

In this case, the compiler will put a stop to it. For example, gives an error message like this (edited):

```
Cannot modify a const object in function
show_array(const double *,int)
```

Other compilers may choose to express the

```
        ar[i] *= r;  
    }  
}
```

Because this function is supposed to alter `ar`, you declare `ar`.

Putting the Pieces Together

Now that you've defined a data type in terms of `int` (three functions), you can put together a program that uses what you've already built all the array-handling tools. The program does check to see if the user enters a value for the multiplication factor with a number. In this case, rather than a single operation, the program uses a loop to ask the user to do the operation. The main work consists of having `main()` call the functions and then showing the result. It places a `using` directive in the program to use the `std` facilities.

Listing 7.7 `arrfun3.cpp`

```
// arrfun3.cpp -- array functions and constants  
#include <iostream>  
const int Max = 5;
```

```
        revalue(factor, properties, size)
        show_array(properties, size);
    }
    cout << "Done.\n";
    cin.get();
    cin.get();
    return 0;
}
```

```
int fill_array(double ar[], int limit)
{
    using namespace std;
    double temp;
    int i;
    for (i = 0; i < limit; i++)
    {
        cout << "Enter value #" << (i + 1) << ": ";
        cin >> temp;
        if (!cin)    // bad input
        {
            cin.clear();
            while (cin.get() != '\n')
                continue;
            cout << "Bad input; input proc
```



```
{  
    for (int i = 0; i < n; i++)  
        ar[i] *= r;  
}
```

Here are two sample runs of the program i

```
Enter value #1: 100000  
Enter value #2: 80000  
Enter value #3: 222000  
Enter value #4: 240000  
Enter value #5: 118000  
Property #1: $100000  
Property #2: $80000  
Property #3: $222000  
Property #4: $240000  
Property #5: $118000  
Enter revaluation factor: 0.8  
Property #1: $80000  
Property #2: $64000  
Property #3: $177600  
Property #4: $192000  
Property #5: $94400  
Done.
```

into a program. This is sometimes called *bottom-up* design, because it moves from the component parts to the whole. The *top-down* approach, which concentrates on data representation and algorithm design, is more like programming, on the other hand, leans toward developing a modular grand design first and then filling in the details. Both methods are useful, and both lead to modular designs.

The Usual Array Function Idiom

Suppose you want a function to process an array. If the function is intended to modify the array, the prototype is

```
void f_modify(double ar[], int n);
```

If the function preserves values, the prototype is

```
void _f_no_change(const double ar[], int n);
```

Of course, you can omit the variable name `n` if the array size is known. It can be something other than `void`. The main point is that the function takes an array element of the passed array and that because the function returns `void`, either function can be used with any size array.

```
double rewards[1000];  
double faults[50];  
...  
f_modify(rewards, 1000);  
f_modify(faults, 50);
```

```
double elbuod[20];
```

Then the two pointers `elbuod` and `elbuod` the name of the array, points to the first element the last element (that is, `elbuod[19]`), so `elbuod` array. Passing a range to a function tells it which Listing 7.6 to use two pointers to specify a range.

Listing 7.8 **arrfun4.cpp**

```
// arrfun4.cpp -- functions with an array
#include <iostream>
const int ArSize = 8;
int sum_arr(const int * begin, const int * end)
int main()
{
    using namespace std;
    int cookies[ArSize] = {1,2,4,8,16,32,64,128};
    // some systems require preceding int with unsigned
    // enable array initialization

    int sum = sum_arr(cookies, cookies + ArSize);
    cout << "Total cookies eaten: " << sum << endl;
    sum = sum_arr(cookies, cookies + 3);
    cout << "First three eaters ate " << sum << endl;
```

Program Notes

In Listing 7.8, notice the for loop in the sum

```
for (pt = begin; pt != end; pt++)  
    total = total + *pt;
```

It sets `pt` to point to the first element to be summed and adds `*pt` (the value of the element) to `total`. Then it increments it, causing it to point to the next element. When `pt` finally equals `end`, it's pointing past the range, so the loop halts.

Second, notice how the different function

```
int sum = sum_arr(cookies, cookies + ArSize - 1);  
...  
sum = sum_arr(cookies, cookies + 3);  
...  
sum = sum_arr(cookies + 4, cookies + 8);
```

The pointer value `cookies + ArSize` points to the element after the last element in the array. (The array has `ArSize` elements, so the last element's address is `cookies + ArSize - 1`.) So the first call to `sum_arr` sums the entire array. Similarly, `cookies + 4`, `cookies + 8`, and so on.

```
cin >> *pt;           // INVALID for the same
```

Now for a subtle point. This declaration for `pt` points to is really a constant; it just means that `pt` is concerned. For example, `pt` points to `age`, and `age` is a constant. You can modify `age` directly by using the `age` variable, but you can't modify it through the pointer:

```
*pt = 20;              // INVALID because pt points to a const  
age = 20;              // VALID because age is not const
```

Previous examples have assigned the address of a regular variable to a pointer. This example assigns the address of a regular variable to a pointer. There are two other possibilities: assigning the address of a const variable to a pointer, and assigning the address of a const to a regular pointer. The first is valid, and the second isn't:

```
const float g_earth = 9.80;  
const float * pe = &g_earth;    // VALID
```

```
const float g_moon = 1.63;  
float * pm = &g_moon;           // INVALID
```

For the first case, you can use neither `g_earth` nor `pe` to modify `g_earth`. `g_earth` doesn't allow the second case for a simple reason: `g_earth` is a const, so you can't modify it. To modify `g_earth`, you can use `pm` to alter the value of `g_earth`.

allows `p1` to be used to alter `const` data. So the address or pointer to a `const` pointer works only for example, if the pointer points to a fundamental

Note

You can assign the address of either `const` or non-`const` data provided that the data type is not itself a pointer. You can assign `const` data only to a non-`const` pointer.

Suppose you have an array of `const` data:

```
const int months[12] = {31,28,31,30,31,30,
```

The prohibition against assigning the address of `const` data does not pass the array name as an argument to a function. The following argument:

```
int sum(int arr[], int n); // should have been const
...
int j = sum(months, 12);    // not allowed
```

This function call attempts to assign a `const` value to `arr` (the `arr` argument), and the compiler disallows the function call.

```
int sloth = 3;  
const int * ps = &sloth;      // a pointer  
int * const finger = &sloth;  // a const p
```

Note that the last declaration has repositioning constrained `finger` to point only to `sloth` and not alter the value of `sloth`. The middle declaration constrains the value of `sloth`, but it permits you to have `ps` and `*ps` and `*ps` are both `const`, and `*finger` and `ps`

If you like, you can declare a `const` pointer:

```
double trouble = 2.0E30;  
const double * const stick = &trouble;
```

Here `stick` can point only to `trouble`, and `*stick` is `const` of `trouble`. In short, both `stick` and `*stick`

Typically you use the pointer-to-`const` for function arguments. For example, recall the `show_array`

```
void show_array(const double ar[], int n);
```

Using `const` in this declaration means that any array that is passed to it. This technique works in the other direction. Here, for example, the array elements are `double` pointers or pointers-to-pointers, you wouldn't

Figure 7.5 Pointers-to-c

Functions and Two-Dimen

To write a function that has a two-dimensional array as an argument, remember that the name of an array is treated as its address. The argument is a pointer, just as for one-dimensional arrays. The function must be called correctly. Suppose, for example, that you start with

```
int data[3][4] = {{1,2,3,4}, {9,8,7,6}, {5,4,3,2}};
int total = sum(data, 3);
```

What should the prototype for `sum()` look like? It should take the number of rows (3) as an argument and not a pointer.

Well, `data` is the name of an array with three rows and four columns. It is an array of four `int` values. Thus, the type of `data` is `int (*)[4]`. The appropriate prototype would be this:

```
int sum(int (*ar2)[4], int size);
```



```
int total3 = sum(a, 10); // sum first 10 elements of a
int total4 = sum(a+10, 20); // sum next 20 elements of a
```

Given that the parameter `ar2` is a pointer to an array, how do we write the function definition? The simplest way is to use `ar2` as if it were an array. Here's a possible function definition:

```
int sum(int ar2[][4], int size)
{
    int total = 0;
    for (int r = 0; r < size; r++)
        for (int c = 0; c < 4; c++)
            total += ar2[r][c];
    return total;
}
```

Again, note that the number of rows is what is passed in, and the number of columns is fixed at four, both in the function definition and in the inner `for` loop.

Here's why you can use array notation. Because `ar2` is a pointer to the first element (element 0) of an array whose elements are array-objects, `ar2[r]` is a pointer to element number `r`. Therefore `ar2[r]` is element number `r` of an array-of-four-int, so `ar2[r]` is the name of the array object that contains the element. Accessing an element of an array by giving an array name gives an array element, so `ar2[r][c]` is an array element, hence is a single `int` value. The pointer `ar2` is a pointer to an array of four `int`s, so it is a pointer to an array of four `int`s.

tions, too. For example, passing a string as an argument to a function can use `const` to protect a string argument from being modified. We'll twist to strings that we'll unravel now.

Functions with C-Style String Arguments

Suppose you want to pass a string as an argument to a function representing a string:

- An array of `char`
- A quoted string constant (also called a string literal)
- A pointer-to-`char` set to the address of the first character in the string

All three choices, however, are type pointers. When you pass a string, you can use all three as arguments to string-processing functions.

```
char ghost[15] = "galloping";  
char * str = "galumphing";  
int n1 = strlen(ghost);           // ghost is an array of char  
int n2 = strlen(str);             // str is a pointer to char  
int n3 = strlen("gamboling");     // address of first char in string
```

Informally, you can say that you're passing a string to a function by passing the address of the first character in the string. The function type should use type `char *` as the type for the argument.

```

        unsigned int ms = c_in_str(mmm, 'm');
        unsigned int us = c_in_str(wail, 'u');
        cout << ms << " m characters in " << m;
        cout << us << " u characters in " << w;
        return 0;
    }

// this function counts the number of ch c
// in the string str
unsigned int c_in_str(const char * str, char ch)
{
    unsigned int count = 0;

    while (*str)           // quit when *str == '\0'
    {
        if (*str == ch)
            count++;
        str++;              // move pointer to next char
    }
    return count;
}

```

acter itself. For example, immediately after the first character in `minimum`. As long as the character is nonzero, so the loop continues. At the end of the loop, the pointer is incremented by 1 byte so that it points to the next character. When it points to the terminating null character, making the condition false, the loop ends. The code for the null character. That condition tests if the character is nonzero. Are these functions ruthless? Because they stop at nothing.

Functions That Return C-Style Strings

Now suppose you want to write a function that returns a C-style string. The function `buildstr` that we saw earlier. But it can return the address of a string. For example, defines a function called `buildstr` that takes two arguments: a character and a number. Using the character and the number, the function creates a string of length equals the number, and then it initializes the string with the character. Finally, it returns a pointer to the new string.

Listing 7.10 `strgback.cpp`

```
// strgback.cpp -- a function that returns a C-style string
#include <iostream>
char * buildstr(char c, int n);          // prototype
int main()
{
```


Let's move from arrays to structures. It's easier with arrays. Although structure variables resemble array items, structure variables behave like basic, simple variables. That is, unlike an array, a structure ties its name to a single entity that will be treated as a unit. Recall that you can pass arrays by reference; you can pass structures by value, just as you do with basic types. A function works with a copy of the original structure. There's no funny business like the name of an array changing in a function. The name of a structure is simply the name of the variable. The address, you have to use the & address operator to denote the address operator. C++ additionally has structure pointers (to be discussed in Chapter 8.)

The most direct way to program by using structures is to pass them as arguments of the basic types—that is, pass them as arguments of functions. However, there is one disadvantage to passing structures by value: the space and effort involved in making a copy of the structure. This can increase requirements and slow down the system. For this reason, C++ allows the passing of structures by value), many programmers prefer to use a pointer of a structure and then using a pointer to access the structure. A third alternative, called *passing by reference*, that is, passing a pointer to the structure, is another two choices now, beginning with passing

```
travel_time sum(travel_time t1, travel_time t2)
```

To add two times, you first add the minutes, then the number of hours to carry over, and the minutes left. Listing 7.11 incorporates this approach into the `show_time()` function to display the contents of a `travel_time` structure.

Listing 7.11 **travel.cpp**

```
// travel.cpp -- using structures with functions
#include <iostream>
using namespace std;

struct travel_time
{
    int hours;
    int mins;
};

const int Mins_per_hr = 60;

travel_time sum(travel_time t1, travel_time t2)
{
    return travel_time{t1.hours + t2.hours, t1.mins + t2.mins};
}

void show_time(travel_time t);

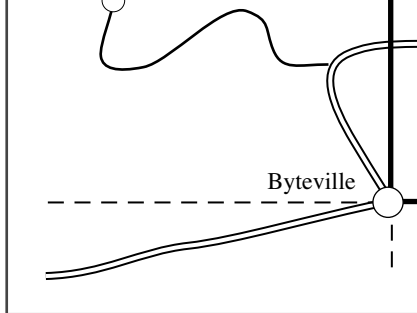
int main()
{
    using namespace std;

    travel_time t1{1, 45};
    travel_time t2{2, 30};
    travel_time t3 = sum(t1, t2);
    show_time(t3);
    return 0;
}
```

```
void show_time(travel_time t)
{
    using namespace std;
    cout << t.hours << " hours, "
         << t.mins << " minutes\n";
}
```

Here `travel_time` acts just like a standard C++ type: variables, function return types, and function arguments. If `t1` and `t2` are `travel_time` structures, you can add them. Note that because the `sum()` function returns a `travel_time` structure, you must pass it an argument for the `show_time()` function. If `sum(t1, t2)` returns by value, the `show_time(sum(t1, t2))` call in order to find the sum of `t1` and `t2` then passes `sum()`'s return value, not the function itself. Here is the output of the program in Listing 7.11:

```
Two-day total: 10 hours, 40 minutes
Three-day total: 15 hours, 12 minutes
```

rectangular coordinates of Mi

Figure 7.6 Recta

```
struct rect
{
    double x;           // horizontal di
    double y;           // vertical dist
};
```

A second way to describe the position of a origin and in what direction it is (for example mathematicians have measured the angle coun



polar coordinates of Micro

Figure 7.7 Po

Let's construct a function that displays the functions in the C++ library (borrowed from need to measure angles in that unit. But for c measure to degrees. This means multiplying b 57.29577951. Here's the function:

```
// show polar coordinates, converting ang
void show_polar (polar dapos)
{
    using namespace std;
    const double Rad_to_deg = 57.29577951

    cout << "distance = " << dapos.distan
    cout << ", angle = " << dapos.angle *
    cout << " degrees\n";
}
```

```

polar rect_to_polar(rect xypos)    // type
{
    polar answer;

    answer.distance =
        sqrt( xypos.x * xypos.x + xypos.y
    answer.angle = atan2(xypos.y, xypos.x)
    return answer;    // returns a polar
}

```

Now that the functions are ready, writing the main program. Listing 7.12 presents the result.

Listing 7.12 **strctfun.cpp**

```

// strctfun.cpp -- functions with a struct
#include <iostream>
#include <cmath>

// structure declarations
struct polar
{
    double distance;    // distance from origin
    double angle;       // direction from origin
};
struct rect

```

```

        return 0;
    }

    // convert rectangular to polar coordinates
    polar rect_to_polar(rect xypos)
    {
        using namespace std;
        polar answer;

        answer.distance =
            sqrt( xypos.x * xypos.x + xypos.y * xypos.y );
        answer.angle = atan2(xypos.y, xypos.x);
        return answer;        // returns a polar coordinate
    }

    // show polar coordinates, converting angles to degrees
    void show_polar (polar dapos)
    {
        using namespace std;
        const double Rad_to_deg = 57.29577951;

        cout << "distance = " << dapos.distance << endl;
        cout << "angle = " << dapos.angle * Rad_to_deg << endl;
        cout << " degrees\n";
    }

```

What really happens when you use `cin >> rplace.x` is that `cin` returns a type `istream` value. If you apply `rplace.x` to that `istream` value, you get the `istream` object (as in `cin >> rplace.x >> rplace.x`). The `istream` class. Thus, the entire while loop tests whether `cin >> rplace.x` returns a non-zero `bool` value of `true` or `false`, depending on whether `cin` expects the user to enter `q`, as shown in the sample output, `cin >> rplace.x` leaves the `q` in the input queue and returns a non-zero value, so the loop.

Compare that approach for reading numbers:

```
for (int i = 0; i < limit; i++)
{
    cout << "Enter value #" << (i + 1) << " : ";
    cin >> temp;
    if (temp < 0)
        break;
    ar[i] = temp;
}
```

To terminate this loop early, you enter a negative value. This restriction fits the needs of

After these changes are made, the function

```
// show polar coordinates, converting angles
void show_polar (const polar * pda)
{
    using namespace std;
    const double Rad_to_deg = 57.29577951;

    cout << "distance = " << pda->distance;
    cout << ", angle = " << pda->angle *
    cout << " degrees\n";
}
```

Next, let's alter `rect_to_polar`. This is more complicated. The `rect_to_polar` function returns a structure. If you should use a pointer instead of a return value. The first points to the structure that's to hold the conversion. The function *modifies* an existing structure in the first argument is `const` pointer, the second is not. The principles used to convert `show_polar()` to pointer are used in the reworked program.

```

    {
        using namespace std;
        rect rplace;
        polar pplace;

        cout << "Enter the x and y values: ";
        while (cin >> rplace.x >> rplace.y)
        {
            rect_to_polar(&rplace, &pplace);
            show_polar(&pplace);          // pas
            cout << "Next two numbers (q to qu
        }
        cout << "Done.\n";
        return 0;
    }

// show polar coordinates, converting angl
void show_polar (const polar * pda)
{
    using namespace std;
    const double Rad_to_deg = 57.29577951;

    cout << "distance = " << pda->distance
    cout << ", angle = " << pda->angle * R

```

Functions and string C

Although C-style strings and `string` class objects, a `string` class object is more closely related to a structure than a C-style string. You can assign a structure to another structure and an array of structures to a function, and you can pass a structure as a complete entity to a function, and you can pass a `string` object as an entity. If you need several strings, you can declare an array of `string` objects instead of a two-dimensional array of C-style strings.

Listing 7.14 provides a short example that passes an array of `string` objects to a function that displays the contents of the array.

Listing 7.14 `topfive.cpp`

```
// topfive.cpp -- handling an array of strings
#include <iostream>
#include <string>
using namespace std;
const int SIZE = 5;
void display(const string sa[], int n);
int main()
{
    string list[SIZE];           // an array of strings
    cout << "Enter your " << SIZE << " favorite songs:\n";
```



```
3: Saturn
4: Jupiter
5: Moon
Your list:
1: Orion Nebula
2: M13
3: Saturn
4: Jupiter
5: Moon
```

The main point to note in this example is that the program treats `string` just as it would treat an array. If you want an array of `string`, you just use the usual array syntax:

```
string list[SIZE];    // an array holding
```

Each element of the `list` array, then, is a `string` object. To get a line of input, you use `getline`:

```
getline(cin, list[i]);
```

Similarly, the formal argument `sa` is a pointer to a `string` object and can be used accordingly:

```
cout << i + 1 << ": " << sa[i] << endl;
```

How can we declare these two functions?
so that's what must appear in the prototypes:

```
void show(std::array<double, 4> da);    //  
void fill(std::array<double, 4> * pa); //
```

These considerations form the core of the more features. First, it replaces 4 with a symbol

```
const int Seasons = 4;
```

Second, it adds a `const` array object containing four seasons:

```
const std::array<std::string, Seasons> Seasons  
    {"Spring", "Summer", "Fall", "Winter"}
```

Note that the array template is not limited to class types too. Listing 7.15 presents the program

Listing 7.15 **arrobj.cpp**

```
//arrobj.cpp -- functions with array objects  
#include <iostream>  
#include <array>  
#include <string>  
// constant data
```

```

        cin >> (*pa)[i];
    }
}

void show(std::array<double, Seasons> da)
{
    using namespace std;
    double total = 0.0;
    cout << "\nEXPENSES\n";
    for (int i = 0; i < Seasons; i++)
    {
        cout << Snames[i] << ": $" << da[i] << "\n";
        total += da[i];
    }
    cout << "Total Expenses: $" << total << "\n";
}

```

Here's a sample run:

```

Enter Spring expenses: 212
Enter Summer expenses: 256
Enter Fall expenses: 208
Enter Winter expenses: 244

```

programming look more complicated.

```
fill(&expenses);    // don't forget the &  
...  
cin >> (*pa)[i];
```

In the last statement, `pa` is a pointer to an object, and `(*pa)[i]` is an element in the object. The logic is straightforward, but the operator precedence makes it easy to make errors.

Using references, as discussed in Chapter 8, avoids these notational problems.

Recursion

And now for something completely different. A characteristic of recursion is that it can call itself. (Unlike C, however, this ability is termed *recursion*. Recursion is an important concept in many fields, such as artificial intelligence, but we'll just take a look at how it works.

calls, first the *statements1* section is executed. Then the *statements2* section is executed, and then the *statements1* section is executed again, and so on. In other words, the order from the order in which the functions were called is reversed. After the base case is reached, recursion, the program then has to back out the recursive calls. This behavior is illustrated in Figure 7.16.

Listing 7.16 **recur.cpp**

```
// recur.cpp -- using recursion
#include <iostream>
using namespace std;

void countdown(int n);

int main()
{
    countdown(4);           // call the recursive function
    return 0;
}

void countdown(int n)
{
    using namespace std;
    cout << "Counting down ... " << n << endl;
    if (n > 0)
        countdown(n - 1);
}
```

```
cout << n << ": Kaboom!"; << " (n at 0012FE0C)
```

Doing so produces output like the following:

```
Counting down ... 4 (n at 0012FE0C)
Counting down ... 3 (n at 0012FD34)
Counting down ... 2 (n at 0012FC5C)
Counting down ... 1 (n at 0012FB84)
Counting down ... 0 (n at 0012FAAC)
0: Kaboom!           (n at 0012FAAC)
1: Kaboom!           (n at 0012FB84)
2: Kaboom!           (n at 0012FC5C)
3: Kaboom!           (n at 0012FD34)
4: Kaboom!           (n at 0012FE0C)
```

Note how the `n` having the value 4 is stored at memory address 0012FE0C in this example), the `n` having the value 3 is stored at memory address 0012FD34, and so on. Also note that the memory address for `n` during the “Counting down” stage is the same as the memory address for `n` during the “Kaboom!” stage.

Recursion with Multiple Recursive Calls

Recursion is particularly useful for situations where a task can be broken down into two smaller, similar tasks. For example, calculating the factorial of a number can be done recursively by multiplying the number by the factorial of the number minus one.

```

        ruler[Len - 1] = '\\0';
        int max = Len - 2;
        int min = 0;
        ruler[min] = ruler[max] = '|';
        std::cout << ruler << std::endl;
        for (i = 1; i <= Divs; i++)
        {
            subdivide(ruler,min,max, i);
            std::cout << ruler << std::endl;
            for (int j = 1; j < Len - 2; j++)
                ruler[j] = ' '; // reset to blank
        }

        return 0;
    }

void subdivide(char ar[], int low, int high)
{
    if (level == 0)
        return;
    int mid = (high + low) / 2;

```

meanly. That is, one call generates two, which generates four, and so on. That's why the level 6 call is able to fill the stack. The exponential doubling of the number of function calls (and the associated overhead) make this form of recursion a poor choice if you can avoid it. It is an elegant and simple choice if the necessary.

Pointers to Functions

No discussion of C or C++ functions would be complete without pointers to functions. We'll take a quick look at this topic in this section, and then return to more advanced texts.

Functions, like data items, have addresses. A pointer variable, which stores the stored machine language code for a function, is as important nor useful for you or the user to know as a data item. In a program, for example, it's possible to write a function as an argument. That enables the first function to call the second function. This approach is more awkward than storing the address of the second one directly, but it leaves open the possibility of calling a function different times. That means the first function

The `process()` call enables the `process()` from within `process()`. The `thought()` call passes the return value of `think()` to the `thought()`.

Declaring a Pointer to a Function

To declare pointers to a data type, the declaration the pointer points. Similarly, a pointer to a function the pointer points. This means the declaration type and the function's signature (its arguments) the same information about a function that a pointer to a function. Pam LeCoder has written a time-estimator.

```
double pam(int); // prototype
```

Here's what a declaration of an appropriate function looks like:

```
double (*pf)(int); // pf points to a function  
                  // one int argument and  
                  // returns type double
```

Tip

In general, to declare a pointer to a particular function, declare the pointer as a pointer to a regular function of the desired kind and use the pointer in the form `(*pf)`. In this case, `pf` is a pointer to a function.

such as the `pam()` function. It could have the

```
void estimate(int lines, double (*pf)(int
```

This declaration says the second argument
argument and a `double` return value. To have
pass `pam()`'s address to it:

```
estimate(50, pam); // function call telli
```

Clearly, the tricky part about using pointers
whereas passing the address is very simple.

Using a Pointer to Invoke a Function

Now we get to the final part of the technique
pointed-to function. The clue comes in the p
the same role as a function name. Thus, all yo
tion name:

```
double pam(int);  
double (*pf)(int);  
pf = pam;           // pf now points to  
double x = pam(4);  // call pam() using  
double y = (*pf)(5); // call pam() using
```

design facilitates future program development. For estimating time, he doesn't have to rewrite supply his own `ralph()` function, making sure type. Of course, rewriting `estimate()` isn't a to more complex code. Also the function pointer behavior of `estimate()`, even if he doesn't have `estimate()`.

Listing 7.18 **fun_ptr.cpp**

```
// fun_ptr.cpp -- pointers to functions
#include <iostream>
double betsy(int);
double pam(int);

// second argument is pointer to a type double
// takes a type int argument
void estimate(int lines, double (*pf)(int))

int main()
{
    using namespace std;
    int code;
```

```
    cout << (*pf)(lines) << " hour(s)\n";  
}
```

Here is a sample run of the program in Listing 10.1.

```
How many lines of code do you need? 30  
Here's Betsy's estimate:  
30 lines will take 1.5 hour(s)  
Here's Pam's estimate:  
30 lines will take 1.26 hour(s)
```

Here is a second sample run of the program in Listing 10.1.

```
How many lines of code do you need? 100  
Here's Betsy's estimate:  
100 lines will take 5 hour(s)  
Here's Pam's estimate:  
100 lines will take 7 hour(s)
```

Variations on the Theme of Function Pointers

With function pointers, the notation can get a bit cumbersome. This section illustrates some of the challenges of function pointers. To begin, here are prototypes for some functions that will be used in the examples.

Now consider the following statements.

```
cout << (*p1)(av,3) << ": " << *(*p1)(av,  
cout << p2(av,3) << ": " << *p2(av,3) << endl;
```

Both `(*p1)(av,3)` and `p2(av,3)`, recall, return `double` values (the return values of `f1()` and `f2()`, in this case) with `av` and `3` as arguments (the first is, address of `double` values). So the first part of the statement prints the address of a `double` value. To see the actual value, we use the `*` operator to these addresses, and that's what `*p2(av,3)` does.

With three functions to work with, it could be useful to have an array of pointers. Then one can use a `for` loop to call them. What would that look like? Clearly, it should look something like `pa[3]` is a pointer, but there should be a `[3]` somewhere. The question is where. And here's the answer (including the `main` function):

```
const double * (*pa[3])(const double *, int);
```

Why put the `[3]` there? Well, `pa` is an array of pointers. Declaring an array of three things is this: `pa[3]`. The type of thing to be placed in the array. Open the `main` function. `*pa[3]` says `pa` is an array of three pointers. The pointers each pointer points to: a function with a signature

because the result can be initialed with a single

```
auto pc = &pa; // C++11 automatic type deduction
```

What if you prefer to do it yourself? Clear the notation for `pa`, but because there is one more level stuck somewhere. In particular, if we call the `pa` pointer, not an array name. This suggests the following. The parentheses bind the `pd` identifier to the

```
*pd[3]    // an array of 3 pointers  
(*pd)[3]  // a pointer to an array of 3 elements
```

In other words, `pd` is a pointer, and it points to things are is described by the rest of the original code. The following:

```
const double *(*(*pd)[3])(const double *,
```

To call a function, realize that if `pd` points to `(*pd)[i]` is an array element, which is a pointer. The for the function call is `(*pd)[i](av, 3)`, and the returned pointer points to. Alternatively, you can call the function with a pointer and use `(*(*pd)[i])(av, 3)` for the pointed-to double value.

```

const double * f3(const double *, int);

int main()
{
    using namespace std;
    double av[3] = {1112.3, 1542.6, 2227.9};

    // pointer to a function
    const double *(*p1)(const double *, int);
    auto p2 = f2; // C++11 automatic type deduction
    // pre-C++11 can use the following code
    // const double *(*p2)(const double *, int);
    cout << "Using pointers to functions:\n";
    cout << " Address Value\n";
    cout << (*p1)(av,3) << ": " << *(*p1)(av,3)\n";
    cout << p2(av,3) << ": " << *p2(av,3)\n;

    // pa an array of pointers
    // auto doesn't work with list initialization
    const double *(*pa[3])(const double *, int);
    // but it does work for initializing the array
    // pb a pointer to first element of pa
    auto pb = pa;
    // pre-C++11 can use the following code

```

```
        // alternative notation
        cout << ((*pd)[2])(av,3) << ": " <<
        // cin.get();
        return 0;
    }
```

```
// some rather dull functions
```

```
const double * f1(const double * ar, int
{
    return ar;
}
const double * f2(const double ar[], int
{
    return ar+1;
}
const double * f3(const double ar[], int
{
    return ar+2;
}
```

The addresses shown are the locations of the

This example may seem esoteric, but point to memory is not unheard of. Indeed, the usual implementation of “Class Inheritance”) uses this technique. Fortunately,

Appreciating `auto`

One of the goals of C++11 is to make C++ easier to use by concentrating more on design and less on details. Listing 7.19 illustrates

```
auto pc = &pa;
const double *(*(pd)[3])(const double
```

The automatic type deduction feature reflects the changes in the compiler. In C++98, the compiler uses its knowledge of the type of the variable at least with this feature, it uses its knowledge of the type of the variable.

There is a potential drawback. Automatic type deduction is not always possible. The compiler is able to match the type of the initializer, but it is not always possible to deduce the wrong type of initializer:

```
auto pc = *pa;    // oops! used *pa instead of &pa
```

This declaration would make `pc` match the type of `pa`. This is a compile-time error when Listing 7.19 later uses `pc`, as

Functions are the C++ programming module. You have a function definition and a prototype, and you have to write the code that implements what the function does. You have to define the function interface: how many and what kind of arguments, what sort of return type, if any, to get from it. The code calls the function arguments to the function and to the function code.

By default, C++ functions pass arguments by value. The parameters in the function definition are new variables created by the function call. Thus, C++ functions pass arguments by working with copies.

C++ treats an array name argument as the address of the array. Technically, this is still passing by value because the function gets the address, but the function uses the pointer to access the array. When you declare formal parameters for a function, the following two declarations are equivalent:

```
typeName arr[];  
typeName * arr;
```

Both of these mean that `arr` is a pointer to an array. In the code, however, you can use `arr` as if it were an array: `arr[i]`. Even when passing pointers, you can pass arrays by declaring the formal argument to be a pointer to an array.

Chapter Review

1. What are the three steps in using a function?
2. Construct function prototypes that match the following functions.
 - a. `igor()` takes no arguments and has no return value.
 - b. `tofu()` takes an `int` argument and returns a `long` value.
 - c. `mpg()` takes two type `double` arguments and returns a `double` value.
 - d. `summation()` takes the name of a variable and returns a `long` value.
 - e. `doctor()` takes a `string` argument and returns a `double` value.
 - f. `ofcourse()` takes a `boss` structure and returns a `double` value.
 - g. `plot()` takes a pointer to a `map` structure and returns a `double` value.

10. C++ enables you to pass a structure by value or by reference. If `glitz` is a structure variable, how would you pass its address? What are the advantages?
11. The function `judge()` has a type `int (*)` as its return type. It takes the address of a function. The function which it calls passes to a `const char` as an argument and returns an `int`.
12. Suppose we have the following structure:
- ```
struct applicant {
 char name[30];
 int credit_ratings[3];
};
```
- a. Write a function that takes an `applicant` structure and displays its contents.
- b. Write a function that takes the address of an `applicant` structure as an argument and displays the contents of the structure.
13. Suppose the functions `f1()` and `f2()` have the following definitions:
- ```
void f1(applicant * a);  
const char * f2(const applicant * a);
```

```

{
    char maker[40];
    float height;
    float width;
    float length;
    float volume;
};

```

- a. Write a function that passes a box value of each member.
 - b. Write a function that passes the a volume member to the product of
 - c. Write a simple program that uses
4. Many state lotteries use a variation of these variations you choose several numbers. For example, you might select also pick a single number (called a megaball) and range, such as 1–27. To win the grand prize, you must pick the numbers correctly. The chance of winning is the probability of picking the field numbers times the probability of picking the megaball.

show the array.

7. Redo Listing 7.7, modifying the three pointer parameters to represent a range, returning the actual number of items read after the last location filled; the other first argument to identify the end of the data.
8. Redo Listing 7.15 without using the array class.
 - a. Use an ordinary array of `const` son names, and use an ordinary array for the expenses.
 - b. Use an ordinary array of `const` son names, and use a structure with a `double` for the expenses. (This does not use the array class.)
9. This exercise provides practice in writing structures. The following is a program skeleton. Write the functions:

```
#include <iostream>
using namespace std;
```

```
// display3() takes the address of t  
// of student structures and the num  
// arguments and displays the conten  
void display3(const student pa[], in
```

```
int main()  
{  
    cout << "Enter class size: ";  
    int class_size;  
    cin >> class_size;  
    while (cin.get() != '\n')  
        continue;  
  
    student * ptr_stu = new student[  
    int entered = getinfo(ptr_stu, c  
    for (int i = 0; i < entered; i++)  
    {  
        display1(ptr_stu[i]);  
        display2(&ptr_stu[i]);  
    }  
    display3(ptr_stu, entered);  
    delete [] ptr_stu;  
    cout << "Done\n";  
    return 0;  
}
```

by using these pointers. Hint: Here's how:

```
double (*pf[3])(double, double);
```

You can initialize such an array by using function names as addresses.

- Function template specializations

With Chapter 7, “Functions: C++’s Program Building Blocks,” you know a lot about C++ functions, but there’s more to learn. New function features that separate C++ from other languages include inline functions, by-reference variable passing (polymorphism), and template functions. This chapter, which you read so far, explores features found in C++ but not in C, and into plus-plussedness.

C++ Inline Functions

Inline functions are a C++ enhancement designed to blur the distinction between normal functions and inline functions. It’s all in how the C++ compiler incorporates them into the program. The distinction between inline functions and normal functions is less important than the program’s innards than we have so far. Let’s dig in.

The final product of the compilation process is a set of machine language instructions. When the program loads these instructions into the computer’s memory, it loads them at a regular memory address. The computer then goes through the instructions. Sometimes, as when you have a loop or a branch,

the non-inline call. On the other hand, you are doing a quick process, so the absolute time savings may not be called frequently.

To use this feature, you must take at least one of the following steps:

- Preface the function declaration with the `inline` keyword.
- Preface the function definition with the `inline` keyword.

A common practice is to omit the prototype (meaning the function header and all the functions that normally go).

The compiler does not have to honor your request. It may decide the function is too large or notice that the function is not indeed possible for inline functions), or the feature is not implemented for your particular compiler.

¹ It's a bit like having to leave off reading some text, finishing the footnote, returning to where you were reading.

```
        cout << "hubba!";  
        cout << "\n";  
    }
```

A regular function transfers program execution to a separate function.

Figure 8.1 Inline function

Listing 8.1 illustrates the inline technique squares its argument. Note that the entire definition of the function is placed in the header file, but if the definition doesn't fit on one or two lines (because of long identifiers), the function is probably a poor candidate for inlining.

Listing 8.1 **inline.cpp**

```
// inline.cpp -- using an inline function  
#include <iostream>  
  
// an inline function definition  
inline double square(double x) { return x*x; }  
  
int main()  
{
```

Even though the program doesn't provide features are still in play. That's because the entire program's first use, serves as a prototype. This means that the program must be compiled with a long argument, and the program must be double before passing it to the function.

Inline Versus Macros

The `inline` facility is an addition to C++. C++ provides *macros*, which are crude implementations of a macro for squaring a number:

```
#define SQUARE(X) X*X
```

This works not by passing arguments but through a symbolic label for the "argument":

```
a = SQUARE(5.0); is replaced by a = 5.0*5.0  
b = SQUARE(4.5 + 7.5); is replaced by b = 4.5 + 7.5  
d = SQUARE(c++); is replaced by d = c++
```

Only the first example here works properly. You need parentheses:

```
#define SQUARE(X) ((X)*(X))
```

Creating a Reference Variable

You might recall that C and C++ use the & symbol. C++ assigns an additional meaning to the & symbol: creating references. For example, to make rodents refer to rats, you could do the following:

```
int rats;  
int & rodents = rats;    // makes rodents refer to rats
```

In this context, & is not the address operator. Just as `char *` in a declaration means pointer to `char`, `int &` means pointer to `int`. The reference declaration allows you to make rodents refer to the same value and the same memory as rats. This is the claim.

Listing 8.2 **firstref.cpp**

```
// firstref.cpp -- defining and using a reference  
#include <iostream>  
int main()  
{  
    using namespace std;  
    int rats = 101;  
    int & rodents = rats;    // rodents is a reference to rats  
}
```

```
rats = 101, rodents = 101
rats = 102, rodents = 102
rats address = 0x0065fd48, rodents address
```

As you can see, both `rats` and `rodents` have the same value. (The address values and display format vary from platform to platform.) Changing the value of `rodents` by one affects both variables. More precisely, `rodents` represents a single variable for which there are two different ways to access it. This example shows you how a reference works. A reference is a reference, which is as a function parameter, parameter passing, and so on. We'll look into these uses pretty soon.)

References tend to be a bit confusing at first because they are tantalizingly reminiscent of pointers, but they are not. You can create both a reference and a pointer to refer to the same variable.

```
int rats = 101;
int &rodents = rats;    // rodents a reference
int *prats = &rats;    // prats a pointer
```

Then you could use the expressions `rodents` and `prats` to access the value of `rats`. From a C++ standpoint, a reference looks a lot like a pointer. The only difference is that the dereferencing operator is understood implicitly for a reference. But there are differences besides that.

Listing 8.3 **secref.cpp**

```
// secref.cpp -- defining and using a reference
#include <iostream>
int main()
{
    using namespace std;
    int rats = 101;
    int & rodents = rats;    // rodents is a reference to rats

    cout << "rats = " << rats;
    cout << ", rodents = " << rodents << endl;

    cout << "rats address = " << &rats;
    cout << ", rodents address = " << &rodents << endl;

    int bunnies = 50;
    rodents = bunnies;        // can we change the value of rats?
    cout << "bunnies = " << bunnies;
    cout << ", rats = " << rats;
    cout << ", rodents = " << rodents << endl;

    cout << "bunnies address = " << &bunnies;
    cout << ", rodents address = " << &rodents << endl;
    return 0;
}
```

```
int & rodents = *pt;  
int bunnies = 50;  
pt = &bunnies;
```

Initializing `rodents` to `*pt` makes `rodents` point to `bunnies` does not alter the fact that `pt` points to `bunnies`.

References as Function Parameters

Most often, references are used as function parameters to provide an alias for a variable in the calling program. This is called *passing by reference*. Passing by reference allows the function to modify the variable in the calling program. C++'s addition of the feature allows for passing by value. Passing by value, recall, results in the call to the function creating a copy of the variable from the calling program (see Figure 8.2). Of course, this has the value limitation by using pointers.

Let's compare using references and using pointers to swap the values of two variables. A swap function that takes two variables in the calling program. That means the function will work because the function will work with the original variables instead of the variables themselves. The function can work with the original data. Alternatively, the function can access the original data. Listing 8.4 shows all the code needed to work, so that you can compare them.

```

    }
    void grumpy(int &x)
    {
        ...
    }

```

→ makes x an alias for times

Figure 8.2 Passing by value and pas

Listing 8.4 **swaps.cpp**

```

// swaps.cpp -- swapping with references and with pointers
#include <iostream>

void swapr(int & a, int & b);    // a, b are aliases
void swapp(int * p, int * q);    // p, q are addresses
void swapv(int a, int b);        // a, b are new variables

int main()
{
    using namespace std;
    int wallet1 = 300;
    int wallet2 = 350;

```

```
        temp = a;           // use a, b for values
        a = b;
        b = temp;
    }
```

```
void swapp(int * p, int * q)    // use pointers
{
    int temp;

    temp = *p;                // use *p, *q for values
    *p = *q;
    *q = temp;
}
```

```
void swapv(int a, int b)       // try using values
{
    int temp;

    temp = a;                 // use a, b for values
    a = b;
    b = temp;
}
```

((swapp(&wallet1, &wallet2))). (Recall that a pointer to an int and therefore the arguments such as &wallet1.)

Next, compare the code for the functions (passing by value). The only outward difference parameters are declared:

```
void swapr(int & a, int & b)
void swapv(int a, int b)
```

The internal difference, of course, is that in swapv(), aliases for wallet1 and wallet2, so swapping wallet1 and wallet2, the variables a and b are new variables, so swapping a and b has no effect on wallet1 and wallet2, so swapping a and b has no effect on wallet1 and wallet2.

Finally, compare the functions swapr() (passing by pointer). The first difference is in how the function parameters are declared:

```
void swapr(int & a, int & b)
void swapp(int * p, int * q)
```

The second difference is that the pointer variable p is dereferenced throughout when the function uses p.

Earlier, I said you should initialize a reference variable. When a function call initializes its parameters with argument variables, the function arguments are initialized to the argument values.

```
        cout << cube(x) ;  
        cout << " = cube of " << x << endl;  
        cout << refcube(x) ;  
        cout << " = cube of " << x << endl;  
        return 0;  
    }
```

```
double cube(double a)  
{  
    a *= a * a;  
    return a;  
}
```

```
double refcube(double &ra)  
{  
    ra *= ra * ra;  
    return ra;  
}
```

Here is the output of the program in Listing 10.1:

```
27 = cube of 3  
27 = cube of 27
```

```
double yo[3] = { 2.2, 3.3, 4.4};  
z = cube (yo[2]);           // pass the
```

Suppose you try similar arguments for a function. It may seem that passing a reference should be more efficient. If you use a name for a variable, then the actual argument is a variable. The following doesn't appear to make sense because

```
double z = refcube(x + 3.0); // should not
```

For example, you can't assign a value to sum:

```
x + 3.0 = 5.0; // nonsensical
```

What happens if you try a function call like `refcube(x + 3.0)` in C++, that's an error, and most compilers will give you a warning along the following lines:

```
Warning: Temporary used for parameter 'ra
```

The reason for this milder response is that you can pass expressions to a reference variable. In some cases, `x + 3.0` is not a type `double` variable, the program initializes it to the value of the expression `x + 3.0` in a temporary variable. Let's take a closer look at how these temporaries are not created.

Now, to return to our example, suppose you have the following constant reference argument:

```
double refcube(const double &ra)
{
    return ra * ra * ra;
}
```

Next, consider the following code:

```
double side = 3.0;
double * pd = &side;
double & rd = side;
long edge = 5L;
double lens[4] = { 2.0, 5.0, 10.0, 12.0 };
double c1 = refcube(side);           // ra
double c2 = refcube(lens[2]);        // ra
double c3 = refcube(rd);             // ra
double c4 = refcube(*pd);            // ra
double c5 = refcube(edge);           // ra
double c6 = refcube(7.0);            // ra
double c7 = refcube(side + 10.0);    // ra
```

The arguments `side`, `lens[2]`, `rd`, and `*pd` are all lvalues, so it is possible to generate a reference for them, (Recall that an element of an array behaves like a variable.)

leaving `a` and `b` unaltered.

In short, if the intent of a function with references passed as arguments, situations that create temporary variables is a solution is to prohibit creating temporary variables. This is what the C++ Standard now does. (However, some compilers still generate some of error messages, so if you see a warning about temporary variables, it's a warning about temporary variables.)

Now think about the `refcube()` function. It's a function that takes a reference to a cube to modify them, so temporary variables cause problems. The general in the sorts of arguments it can handle. The difference is `const`, C++ generates temporary variables. If a function with a `const` reference formal argument, it mimics the traditional passing by value behavior. It's unaltered and using a temporary variable to hold the value.

Note

If a function call argument isn't an lvalue or a `const` reference parameter, C++ creates an unnamed temporary variable and stores the value of the function call argument to that variable. The function then refers to that variable.

semantics. The original reference type (the one with lvalue reference).

Using References with a Structure

References work wonderfully with structures. Indeed, references were introduced primarily for basic built-in types.

The method for using a reference to a structure is the same as the method for using a reference to a basic variable when declaring a structure parameter. For example, here is a declaration of a structure:

```
struct free_throws
{
    std::string name;
    int made;
    int attempts;
    float percent;
};
```

Then a function using a reference to this type can be declared as follows:

```
void set_pc(free_throws & ft);    // use a reference
```

If the intent is that the function doesn't alter the structure, then the function can be declared as follows:

```
void display(const free_throws & ft);    // use a const reference
```

```

// partial initializations - remaining me
    free_throws one = {"Ifelsa Branch", 1
    free_throws two = {"Andor Knott", 10,
    free_throws three = {"Minnie Max", 7,
    free_throws four = {"Whily Looper", 5
    free_throws five = {"Long Long", 6, f
    free_throws team = {"Throwgoods", 0,
// no initialization
    free_throws dup;

    set_pc(one);
    display(one);
    accumulate(team, one);
    display(team);
// use return value as argument
    display(accumulate(team, two));
    accumulate(accumulate(team, three), f
    display(team);
// use return value in assignment
    dup = accumulate(team, five);
    std::cout << "Displaying team:\n";
    display(team);
    std::cout << "Displaying dup after as
    display(dup);

```

```
free_throws & accumulate(free_throws & target)
{
    target.attempts += source.attempts;
    target.made += source.made;
    set_pc(target);
    return target;
}
```

Here is the program output:

```
Name: Ifelsa Branch
    Made: 13      Attempts: 14      Percent: 93
Name: Throwgoods
    Made: 13      Attempts: 14      Percent: 93
Name: Throwgoods
    Made: 23      Attempts: 30      Percent: 77
Name: Throwgoods
    Made: 35      Attempts: 48      Percent: 73
Displaying team:
Name: Throwgoods
    Made: 41      Attempts: 62      Percent: 66
Displaying dup after assignment:
Name: Throwgoods
```

```

        if (pt->attempts != 0)
            pt->percent = 100.0f *float(pt->m
        else
            pt->percent = 0;
    }

```

The next function call is this:

```
display(one);
```

Because `display()` displays the contents of `one`, the function uses a `const` reference parameter. In this case, the structure is passed by value, but using a reference is more efficient because it creates a copy of the original structure.

The next function call is this:

```
accumulate(team, one);
```

The `accumulate()` function takes two structures as parameters: `team` and `one`. `team` is a reference to the first structure and `one` is a reference to the second structure. Only the first structure is modified, whereas the second parameter is a `const` reference. The following code shows the `free_throws` and `accumulate` functions being called on the `free_throws` and `accumulate` functions.

```
display(accumulate(team, two));
```

is the same as that of the following:

```
accumulate(team, two);  
display(team);
```

The same logic applies to this statement:

```
accumulate(accumulate(team, three), four);
```

This has the same effect as the following:

```
accumulate(team, three);  
accumulate(team, four);
```

Next, the program uses an assignment statement:

```
dup = accumulate(team, five);
```

As you might expect, this copies the values.

Finally, the program uses `accumulate()` in

```
accumulate(dup, five) = four;
```

This statement—that is, assigning a value to a variable whose value is a reference. The code won't compile if the return value is a reference to `dup`, this code

If `accumulate()` returned a structure instead of a reference, it would involve copying the entire structure to a temporary variable, then returning to `dup`. But with a reference return value, `tea` is a much simpler approach.

Note

A function that returns a reference is actually

Being Careful About What a Return Reference

The single most important point to remember about returning a reference to a memory location is to be careful. What you want to avoid is code along the lines of

```
const free_throws & clone2(free_throws &ft)
{
    free_throws newguy; // first step to clone
    newguy = ft;         // copy info
    return newguy;       // return reference
}
```

This has the unfortunate effect of returning a reference to `newguy` that passes from existence as soon as the function returns. In “Chapter 10, Memory Models and Namespaces,” discusses the problem in more detail. Similarly, you should avoid returning pointers to

no longer needed. A call to `clone()` conceals to use `delete` later. The `auto_ptr` template in Chapter 16, “The `string` Class and the Stack,” handles the deletion process.

Why Use `const` with a Reference Return

Listing 8.6, as you’ll recall, had this statement:

```
accumulate(dup, five) = four;
```

It had the effect of first adding data from `five` to `dup` with the contents of `four`. Why does this work? Because `accumulate` returns a modifiable lvalue on the left. That is, the subexpression should identify a block of memory. The function returned a reference to `dup`, which does indeed make the statement valid.

Regular (non reference) return types, on the other hand, are accessed by address. Such expressions can appear on the right-hand side of an assignment but not the left. Other examples of rvalues are expressions such as `x + y`. Clearly, it doesn’t make sense to assign such as `10.0`, but why is a normal function return value, you’ll recall, resides in a temporary memory location even until the next statement.

And of course you still could use accumulated statement.

By omitting `const`, you can write shorter

Usually, you're better off avoiding the add obscure features often expand the opportunity type a `const` reference therefore protects you sionally, however, omitting `const` does make s in Chapter 11, "Working with Classes," is an

Using References with a Class O

The usual C++ practice for passing class objects instance, you would use reference parameters `ostream`, `istream`, `ofstream`, and `ifstream` c

Let's look at an example that uses the `str` choices, some of them bad. The general idea to each end of another string. Listing 8.7 pro this. However, one of the designs is so flawed even not compile.


```

        result = version2(input, "###");
        cout << "Your string enhanced: " << result;
        cout << "Your original string: " << input;

        cout << "Resetting original string.\n";
        input = copy;
        result = version3(input, "@@@");
        cout << "Your string enhanced: " << result;
        cout << "Your original string: " << input;

        return 0;
    }

string version1(const string & s1, const string & s2)
{
    string temp;

    temp = s2 + s1 + s2;
    return temp;
}

const string & version2(string & s1, const string & s2)
{

```

At this point the program crashed.

Program Notes

Version 1 of the function in Listing 8.7 is the

```
string version1(const string & s1, const  
{  
    string temp;  
  
    temp = s2 + s1 + s2;  
    return temp;  
}
```

It takes two `string` arguments and uses `string` that has the desired properties. Note that the `string` objects are brand-new. The function would produce the same

```
string version4(string s1, string s2) /
```

In this case, `s1` and `s2` would be brand-new `string` objects. This is more efficient because the function doesn't have to copy the old objects to the new. The use of the `string` objects is to use, but not modify, the original strings.

The `temp` object is a new object, local to the function. It exists only when the function terminates. Thus, returning

the common case of the return type being a quoted string literal, a null-terminated array of `char`. Hence the following works fine:

```
result = version1(input, "****");
```

The `version2()` function doesn't create a new string, it just returns a reference to the original string:

```
const string & version2(string & s1, const string & s2)
{
    s1 = s2 + s1 + s2;
    // safe to return reference passed to function
    return s1;
}
```

This function is allowed to alter `s1` because it's a reference.

Because `s1` is a reference to an object (input), it's safe to alter it. Because `s1` is a reference to `input`, the following code is safe:

```
result = version2(input, "###");
```

`input` essentially becomes equivalent to the following:

```
version2(input, "###");    // input altered
result = input;            // reference to original
```

The program attempts to refer to memory

Another Object Lesson: Objects,

The `ostream` and `ofstream` classes bring an
As you may recall from Chapter 6, “Branching
of the `ofstream` type can use `ostream` methods
same forms as console input/output. The language
features from one class to another is called *inheritance*.
discusses this feature in detail. In brief, `ostream`
`ofstream` class is based on it) and `ofstream` is
from `ostream`). A derived class inherits the base
`ofstream` object can use base class features such as
writing methods.

Another aspect of inheritance is that a base
object without requiring a type cast. The practical
function having a base class reference parameter
class objects and also with derived objects. For
parameter can accept an `ostream` object, such as
might declare, equally well.

Listing 8.8 demonstrates this point by using
and to display the data onscreen; only the function

```

    {
        cout << "Can't open " << fn << ".
        exit(EXIT_FAILURE);
    }
    double objective;
    cout << "Enter the focal length of your
        "telescope objective in mm: ";
    cin >> objective;
    double eps[LIMIT];
    cout << "Enter the focal lengths, in m
        << " eyepieces:\n";
    for (int i = 0; i < LIMIT; i++)
    {
        cout << "Eyepiece #" << i + 1 << "
        cin >> eps[i];
    }
    file_it(fout, objective, eps, LIMIT);
    file_it(cout, objective, eps, LIMIT);
    cout << "Done\n";
    return 0;
}

```

Here is a sample run of the program in L

```
Enter the focal length of your telescope
Enter the focal lengths, in mm, of 5 eyep
Eyepiece #1: 30
Eyepiece #2: 19
Eyepiece #3: 14
Eyepiece #4: 8.8
Eyepiece #5: 7.5
Focal length of objective: 1800 mm
f.l. eyepiece   magnification
           30.0             60
           19.0             95
           14.0            129
            8.8            205
            7.5            240
Done
```

The following line writes the eyepiece data

```
file_it(fout, objective, eps, LIMIT);
```

And this line writes the identical informat

```
file_it(cout, objective, eps, LIMIT);
```


- If the data object is a class object, use a

Of course, these are just guidelines, and there are other choices. For example, `cin` uses references for `&n` instead of `cin >> &n`.

Default Arguments

Let's look at another topic from C++'s bag of tricks. A *default argument* is a value that's used automatically if an argument is omitted from a function call. For example, if you set `n` to the value of 1, the function call `wow()` is the same as `wow(1)` if you use a function. Suppose you have a function `left` that takes characters of a string, with the string and `n` as arguments. It returns a pointer to a new string consisting of the first `n` characters. For example, the call `left("theory", 3)` returns a pointer to it. Now suppose you establish a default value for `n`. The call `left("theory", 3)` would work as before. But the call `left("theory")`, instead of taking an argument of 1 and return a pointer to the string, the program often needs to extract a one-character from longer strings.


```
beeps = harpo(2);           // same as harpo(2)
beeps = harpo(1,8);         // same as harpo(1,8)
beeps = harpo (8,7,6);      // no default arguments
```

The actual arguments are assigned to the correct variables; you can't skip over arguments. Thus, the following is invalid:

```
beeps = harpo(3, ,8);      // invalid, doesn't compile
```

Default arguments aren't a major programming feature, but they do provide some convenience. When you begin working with classes, you'll find a large number of constructors, methods, and member functions that use default arguments.

Listing 8.9 puts default arguments to use. Note that the function has a default argument. The function definition is the same as the function declaration.

Listing 8.9 **left.cpp**

```
// left.cpp -- string function with a default argument
#include <iostream>
const int ArSize = 80;
char * left(const char * str, int n = 1);
int main()
{
    using namespace std;
    char sample[ArSize];
    cout << "Enter a string:\n";
```

```
}
```

Here's a sample run of the program in Listing 8.9:

```
Enter a string:
```

```
forthcoming
```

```
fort
```

```
f
```

Program Notes

The program in Listing 8.9 uses `new` to create a string of characters. One awkward possibility is that an user enters a number of characters. In that case, the function `strlen` returns the null string. Another awkward possibility is that a user request more characters than the string contains. To avoid these, using a combined test:

```
i < n && str[i]
```

The `i < n` test stops the loop after `n` characters. In the test, the expression `str[i]`, is the code for the loop reaches the null character, the code is 0,

```
int m = 0;
while (m <= n && str[m] != '\0')
    m++;
char * p = new char[m+1]:
// use m instead of n in rest of code
```

Remember, the expression `str[m] != '\0'` is `false` when it is the null character and `true` when it is not. The expression `str[m] != '\0'` is converted to `true` in an `&&` expression and zero is converted to `false`. It can be written this way:

```
while (m<=n && str[m])
```

Function Overloading

Function polymorphism is a neat C++ addition. C++ overloading lets you call the same function by using different names. *Polymorphism*, also called *function overloading*, lets you have many names for the same function. The word *polymorphism* means having many forms. Similarly, the expression *overloading* means attaching more than one function to the same name. These expressions boil down to the same thing, but we'll use the word *overloading*—it sounds harder working. You can have many versions of functions that do essentially the same thing.

```
print(1999.0, 10);           // use #2  
print(1999, 12);             // use #4  
print(1999L, 15);            // use #3
```

For example, `print("Pancakes", 15)` uses prototype #1.

When you use overloaded functions, you must specify the types in the function call. For example, consider the following code:

```
unsigned int year = 3210;  
print(year, 6);              // ambiguous call
```

Which prototype does the `print()` call match? The lack of a matching prototype doesn't automatically cause an error because C++ will try to use standard type conversions. If the `print()` prototype were #2, the function call would convert the `year` value to type `double`. But in the earlier code, `year` is an integer, so the number is the first argument, providing three different prototypes. With this ambiguous situation, C++ rejects the code.

Some signatures that appear to be different are actually the same. For example, consider these two prototypes:

```
double cube(double x);  
double cube(double & x);
```

```
dabble(p2);           // dabble(char *);  
drivel(p1);           // drivel(const char *);  
drivel(p2);           // drivel(const char *);
```

The `dribble()` function has two prototypes with different pointer arguments—regular pointers—and the compiler selects one of them if the first actual argument is `const`. The `dabble()` function has one argument, but the `drivel()` function matches two different arguments. The reason for this difference in behavior is that it's valid to assign a non-`const` value to a `const` variable.

Keep in mind that the signature, not the function body, determines the return type. For example, the following two declarations are identical:

```
long gronk(int n, float m);      // same signature  
double gronk(int n, float m);   // hence same return type
```

Therefore, C++ doesn't permit you to overload a function with different return types, but only if the signatures are identical.

```
long gronk(int n, float m);      // different signature  
double gronk(float n, float m);  // hence different return type
```

After we discuss templates later in this chapter, we'll see how template matching works.

```
double x = 55.5;
const double y = 32.0;
stove(x);           // calls stove(d
stove(y);           // calls stove(c
stove(x+y);         // calls stove(d
```

If, say, you omit the `stove(double &&)` function and use the `stove(const double &)` function instead

An Overloading Example

In this chapter we've already developed a `left` function that returns the first `n` characters in a string. Let's add a second function that returns the first `n` digits in an integer. You can use it, for example, to extract the postal zip code stored as an integer, which is

The integer function is a bit more difficult because you don't have the benefit of each digit being a character. One approach is to first compute the number of digits in the integer, then loop 10 logs off one digit, so you can use division to extract the digit with a loop, like this:

```
unsigned digits = 1;
while (n /= 10)
    digits++;
```

from that of the old `left()`, you can use both

Listing 8.10 **leftover.cpp**

```
// leftover.cpp -- overloading the left()
#include <iostream>

unsigned long left(unsigned long num, unsigned
char * left(const char * str, int n = 1);

int main()
{
    using namespace std;
    char * trip = "Hawaii!!";    // test variable
    unsigned long n = 12345678; // test variable
    int i;
    char * temp;

    for (i = 1; i < 10; i++)
    {
        cout << left(n, i) << endl;
        temp = left(trip, i);
        cout << temp << endl;
        delete [] temp; // point to temporary
    }
    return 0;
```

```
// This function returns a pointer to a new string
// consisting of the first n characters of str
char * left(const char * str, int n)
{
    if(n < 0)
        n = 0;
    char * p = new char[n+1];
    int i;
    for (i = 0; i < n && str[i]; i++)
        p[i] = str[i]; // copy characters
    while (i <= n)
        p[i++] = '\0'; // set rest of string to 0
    return p;
}
```

Here's the output of the program in Listing 12.3:

```
1
H
12
Ha
123
Haw
1234
```


But using the single function with a default argument to write instead of two, and the program instead of two. If you decide to modify the function if you require different types of arguments, in that case, you should use function overloading.

What Is Name Decoration?

How does C++ keep track of which overloaded function to use for each of these functions. When you use the `g++` and compile programs, your C++ compiler performs *name decoration* or *name mangling*—throughout the program on the formal parameter types specified in the function's undecorated function prototype:

```
long MyFunctionFoo(int, float);
```

This format is fine for us humans; we know that `int` and `float`, and it returns a value of type `long`. The compiler implements this interface by transforming the name of the function into an unsightly appearance, perhaps something like

```
?MyFunctionFoo@@YAXH
```

The apparent gibberish decorating the original function name (the *mangled name*) encodes the number and types of parameters. The result in a different set of symbols being added to the program's symbol table. These conventions for their efforts at decorating.

```
double x; // intended change of  
short doubleerval; // unintended change
```

C++'s function template capability automating greater reliability.

Function templates enable you to define an example, you can set up a swapping template

```
template <typename AnyType>  
void Swap(AnyType &a, AnyType &b)  
{  
    AnyType temp;  
    temp = a;  
    a = b;  
    b = temp;  
}
```

The first line specifies that you are setting an arbitrary type `AnyType`. The keywords `template` and `<` you can use the keyword `class` instead of `typename` if you wish. The type name (`AnyType`, in this example) follows the usual C++ naming rules; many programmers agree that, to admit, is simple indeed. The rest of the code defines the function of type `AnyType`. The template does not provide the compiler with directions about how to define the function.

You should use templates if you need function types. If you aren't concerned with backward typing a longer word, you can use the keyword `decltype` to declare type parameters.

To let the compiler know that you need a function, you use a function called `Swap()` in your program. The compiler generates the corresponding code for you and then generates the corresponding object code. The program layout follows the usual pattern: the function prototype near the top of the file, followed by `main()`. The example follows the more modern style using `decltype` as the type parameter.

Listing 8.11 **funtemp.cpp**

```
// funtemp.cpp -- using a function template
#include <iostream>
// function template prototype
template <typename T> // or class T
void Swap(T &a, T &b);

int main()
{
    using namespace std;
```

```
    a = b;  
    b = temp;  
}
```

The first `Swap()` function in Listing 8.11 creates an `int` version of the function. That is, a definition that looks like this:

```
void Swap(int &a, int &b)  
{  
    int temp;  
    temp = a;  
    a = b;  
    b = temp;  
}
```

You don't see this code, but the compiler generates it. The second `Swap()` function has two `double` versions. That is, it replaces `T` with `double`.

```
void Swap(double &a, double &b)  
{  
    double temp;  
    temp = a;
```

You use templates when you need functions to work with different types, as in Listing 8.11. It might be, however, that you need a function to perform a specific algorithm. To handle this possibility, you can overload functions with different regular function definitions. As with ordinary functions, overloaded functions have distinct function signatures. For example, Listing 8.12 shows two versions of a function for swapping elements of two arrays. The original function has the signature `void Swap(T &a, T &b);`, whereas the new template has the signature `template <typename T> void Swap(T *a, T *b, int n);`. The first parameter in this case happens to be a specific type, but in general, all template arguments have to be template parameters.

When, in `twotemps.cpp`, the compiler encounters the first definition, it sees that it has two `int` arguments and matches `Swap` with the first definition. If, however, it has two `int` arrays and an `int` value for the third argument, it uses the second template.

Listing 8.12 **twotemps.cpp**

```
// twotemps.cpp -- using overloaded templates
#include <iostream>

template <typename T>          // original template
void Swap(T &a, T &b);

// new template
template <typename T>
void Swap(T *a, T *b, int n);
```

```
}
```

```
template <typename T>  
void Swap(T &a, T &b)  
{  
    T temp;  
    temp = a;  
    a = b;  
    b = temp;  
}
```

```
template <typename T>  
void Swap(T a[], T b[], int n)  
{  
    T temp;  
    for (int i = 0; i < n; i++)  
    {  
        temp = a[i];  
        a[i] = b[i];  
        b[i] = temp;  
    }  
}
```

```
void Show(int a[])  
{
```

Suppose you have a template function:

```
template <class T>          // or template <type T>
void f(T a, T b)
{...}
```

Often the code makes assumptions about `T`. For instance, the following statement assumes that `a > b` is true if type `T` is a built-in array type:

```
a = b;
```

Similarly, the following assumes `a > b` is defined for a structure:

```
if (a > b)
```

Also the `>` operator is defined for array names. `a > b` compares the addresses of the arrays, which makes no sense. The following assumes the multiplication operator `*` is defined if `T` is an array, a pointer, or a structure:

```
T c = a*b;
```

In short, it's easy to write a template function. On the other hand, sometimes a generalization makes no sense. It allows for it. For example, it could make sense to write `a + b` if `T` is a coordinate, even though the `+` operator isn't defined for `T`.

Because C++ allows you to assign one structure type to another, the function `swap()` for type `T` is a job structure. But suppose you only want to swap the members, keeping the name members unchanged. The arguments to `swap()` would be the same as for the original function, so you can't use template overloading to supply a specialized function.

However, you can supply a specialized function, called a *specialization*, with the required code. If the code exactly matches a function call, it uses that definition.

The specialization mechanism has changed in the current form as mandated by the C++ Standard.

Third-Generation Specialization (ISO/ANSI C++)

After some youthful experimentation with other approaches, the committee has settled on this approach:

- For a given function name, you can have a regular function, a template specialization, and an explicit specialization template. The compiler chooses one of all of these.
- The prototype and definition for an explicit specialization template should mention the specialization.
- A specialization overrides the regular function. A regular function overrides both.


```

int main()
{
    double u, v;
    ...
    Swap(u,v); // use template
    job a, b;
    ...
    Swap(a,b); // use void Swap<job>(job
}

```

The `<job>` in `Swap<job>` is optional because this is a specialization for `job`. Thus, the prototype is

```
template <> void Swap(job &, job &); //
```

In case you have to work with an older compiler, standard usage soon, but first, let's see how explicit specialization

An Example of Explicit Specialization

Listing 8.13 illustrates how explicit specialization

```

    int i = 10, j = 20;
    cout << "i, j = " << i << ", " << j << "\n";
    cout << "Using compiler-generated int
Swap(i,j);    // generates void Swap(
cout << "Now i, j = " << i << ", " << j << "\n";

    job sue = {"Susan Yaffee", 73000.60,
    job sidney = {"Sidney Taffee", 78060.00,
    cout << "Before job swapping:\n";
    Show(sue);
    Show(sidney);
    Swap(sue, sidney); // uses void Swap(
    cout << "After job swapping:\n";
    Show(sue);
    Show(sidney);
    // cin.get();
    return 0;
}

```

```

template <typename T>
void Swap(T &a, T &b)    // general version
{
    T temp;
    temp = a;
    a = b;
    b = temp;
}

```

```
}
```

Here's the output of the program in Listing

```
i, j = 10, 20.
```

```
Using compiler-generated int swapper:
```

```
Now i, j = 20, 10.
```

```
Before job swapping:
```

```
Susan Yaffee: $73000.60 on floor 7
```

```
Sidney Taffee: $78060.72 on floor 9
```

```
After job swapping:
```

```
Susan Yaffee: $78060.72 on floor 9
```

```
Sidney Taffee: $73000.60 on floor 7
```

Instantiations and Specializations

To extend your understanding of templates, let's look at *instantiation*. Keep in mind that including a function template does not generate a function definition. It's merely a placeholder. When the compiler uses the template to generate code for a specific type, the result is termed an *instantiation* of the template. For example, the function call `Swap(i, j)` causes the compiler to instantiate the `Swap` template with `int` as the type. The template *is not* a function.

explicitly defined for the `int` type. These provide function definitions. The explicit specialization plate, whereas the explicit instantiation omits

Caution

It is an error to try to use both an explicit instantiation and an explicit specialization for the same type(s) in the same file, or, more gener-

Explicit instantiations also can be created. For an example of an explicit instantiation, consider the following:

```
template <class T>
T Add(T a, T b)      // pass by value
{
    return a + b;
}
...
int m = 6;
double x = 10.2;
cout << Add<double>(x, m) << endl;  // explicit instantiation
```

The template would fail to match the function call because the compiler expects both function arguments to be of the type `double`, but the first forces the type `double` instantiation, and the second forces the type `int` match the second parameter of the `Add<double>`.

```

{
    template void Swap<char>(char &, char &)
    short a, b;
    ...
    Swap(a,b);    // implicit template instantiation
    job n, m;
    ...
    Swap(n, m);    // use explicit specialization
    char g, h;
    ...
    Swap(g, h);    // use explicit template instantiation
    ...
}

```

When the compiler reaches the explicit instantiation to generate a char version of `Swap()`. The compiler matches a template to the actual arguments when the compiler reaches the function call `Swap(n, m)` because the two arguments are type `short`. When the compiler reaches `Swap(g, h)`, it uses the separate definition (the `char` version) generated when it processed the explicit instantiation.

First, the compiler rounds up the suspects, that have the name `may()`. Then, it finds those that match the following pass muster because they have one argument:

```
void may(int);  
float may(float, float = 3);  
void may(char);  
char * may(const char *);  
char may(const char &);  
template<class T> void may(const T &);  
template<class T> void may(T *);
```

Note that just the signatures and not the return types of candidates (#4 and #7), however, are not viable. The first is implicitly `void` (that is, without an explicit type can only be `void`) is viable because it can be used to generate a `void` expression. That leaves five viable functions, each of which is declared.

Next, the compiler has to determine which of the five is the best. The conversion required to make the function argument match the argument. In general, the ranking from best to

that match to an int or formal parameter, so these rules include converting char & to char* entry means that a function name as an actual formal parameter, as long as both have the same number function pointers from Chapter 7. Also recall as an argument to a function that expects a pointer volatile keyword later in Chapter 9.

Table 8.1 Trivial Conversions Allowed for an

From an Actual Argument	To a Formal Parameter
Type	Type &
Type &	Type
Type []	* Type
Type (argument-list)	Type (*)
Type	const Type
Type	volatile Type
Type *	const Type *
Type *	volatile Type *

Another case in which one exact match is better than a non-template function and the other isn't. It's better than a template, including explicit specialization.

If you wind up with two exact matches the compiler chooses the template function that is the more specialized. For example, that an explicit specialization is chosen over the template pattern:

```
struct blot {int a; char b[10];};  
template <class Type> void recycle (Type t);  
template <> void recycle<blot> (blot & t);  
...  
blot ink = {25, "spots"};  
...  
recycle(ink); // use specialization
```

The term *most specialized* doesn't necessarily mean the most specialized. Generally, it indicates that fewer conversions take place to use. For example, consider the following two functions:

```
template <class Type> void recycle (Type t);  
template <class Type> void recycle (Type t);
```


Let's examine a complete program that uses the `ShowArray` template definition to use. Listing 8.14 shows the contents of an array. The first definition (Template A) is passed as an argument contains the data to be shown. Template B) assumes that the array elements are pointers.

Listing 8.14 **tempover.cpp**

```
// tempover.cpp -- template overloading
#include <iostream>

template <typename T>                // template A
void ShowArray(T arr[], int n);

template <typename T>                // template B
void ShowArray(T * arr[], int n);

struct debts
{
    char name[50];
    double amount;
};

int main()
```

```
template <typename T>
void ShowArray(T arr[], int n)
{
    using namespace std;
    cout << "template A\n";
    for (int i = 0; i < n; i++)
        cout << arr[i] << ' ';
    cout << endl;
}
```

```
template <typename T>
void ShowArray(T * arr[], int n)
{
    using namespace std;
    cout << "template B\n";
    for (int i = 0; i < n; i++)
        cout << *arr[i] << ' ';
    cout << endl;
}
```

```
}

int main()
{
    using namespace std;
    int m = 20;
    int n = -30;
    double x = 15.5;
    double y = 25.9;

    cout << lesser(m, n) << endl;      /
    cout << lesser(x, y) << endl;      /
    cout << lesser<>(m, n) << endl;    /
    cout << lesser<int>(x, y) << endl; /

    return 0;
}
```

(The final function call converts double to int things about that.)

should choose a template function rather than noting that the actual arguments are type `int`.

Finally, consider this statement:

```
cout << lesser<int>(x, y) << endl; // use
```

Here we have a request for an explicit instantiation that gets used. The values of `x` and `y` are `int`, so the `lesser` function returns an `int` value, which is why the program

Functions with Multiple Type Arguments

Where matters really get involved is when a function is matched to prototypes with multiple type arguments for all the arguments. If it can find a function that is a better match than any other functions, it is chosen. For one function to be better than another, it must be at least as good a match for all arguments and a better match for at least one argument.

This book does not intend to challenge the rules. The rules are there so that there is a well-defined way to choose between prototypes and templates.

Template Function Evolution

In the early days of C++, most people didn't use templates. Functions and template classes would prove to

The proper type might be T1 or T2 or some other type. T1 could be double and T2 could be int, in which case the resultant type is double. T1 could be short and T2 could be int, in which case the resultant type is int. T1 could be short and T2 is char. Then addition invokes the resultant type is int. Also the + operator can complicate the options further. Therefore, in C++11, we have the solution of xpy.

The decltype Keyword (C++11)

The C++11 solution is a new keyword: `decltype`.

```
int x;  
decltype(x) y;    // make y the same type as x
```

The argument to `decltype` can be an expression. Here is this code:

```
decltype(x + y) xpy;    // make xpy the same type as x + y  
xpy = x + y;
```

Alternatively, we could combine these two lines of code:

```
decltype(x + y) xpy = x + y;
```

```
decltype(rx) u = y; // u is type double &
decltype(pd) v;      // v is type const double
```

Stage 2: If *expression* is a function call, return type:

```
long indeed(int);
decltype (indeed(3)) m; // m is type int
```

Note

The call *expression* isn't evaluated. In this case, we get the return type; there's no need to actually

Stage 3: If *expression* is an lvalue, then *expression* has the same type as the variable it refers to. This might seem to imply that earlier examples would be wrong, but they're not. For example, `double w = 4.4;` has type `double`, given that `w` is an lvalue. However, keep in mind that `w` is a variable, not a function call. Stage 1. For this stage to apply, *expression* can't be a function call. What can it be? One obvious possibility is a pointer to a variable.

```
double xx = 4.4;
decltype ((xx)) r2 = xx; // r2 is double
decltype(xx) w = xx;    // w is double
```

```

        xytype xpy = x + y;
        xytype arr[10];
        xytype & rxy = arr[2];    // rxy a ref
        ...
    }

```

Alternative Function Syntax (C++11 Tr

The decltype mechanism by itself leaves and incomplete template function:

```

template<class T1, class T2>
?type? gt(T1 x, T2 y)
{
    ...
    return x + y;
}

```

Again, we don't know in advance what type that we could use `decltype(x + y)` for the return type. At the point where the code, the parameters `x` and `y` have not yet been declared (and are thus unusable to the compiler). The `decltype` mechanism is only usable after the variables are declared. To make this possible, C++11 allows `decltype` to be used in function declarations. Here's how it works using built-in

Summary

C++ has expanded C function capabilities. By placing a function definition ahead of the first function call, you suggest to the C++ compiler that it make the function call. The program then jumps to a separate section of code to execute the function call with the corresponding code. This is useful only when the function is short.

A reference variable is a kind of disguised pointer (it has a second name) for a variable. Reference variables are used in functions that process structures and class objects. A reference to a particular type can refer only to data of that type. A reference to a derived type can refer only to data of that type or derived from another, such as `ofstream` from `ostream`. References also refer to the derived type.

C++ prototypes enable you to define default values for arguments. If you omit the corresponding argument, the program uses the default. If you include an argument value, the program uses that value. Default arguments can be provided only from right to left. To provide a default value for a particular argument, you must provide arguments to the right of that argument.

3. Write overloaded versions of `quote()` enclosed in double quotation marks. Write one for a double argument, and one for

4. The following is a structure template:

```
struct box
{
    char maker[40];
    float height;
    float width;
    float length;
    float volume;
};
```

- Write a function that has a reference parameter and displays the value of each member.
 - Write a function that has a reference parameter and sets the volume member.
5. What changes would need be made to `show()` use reference parameters?

8. Given the template of Chapter Review Question 4, provide a template that takes two `int` arguments and returns the one with the larger value.
9. What types are assigned to `v1`, `v2`, `v3`, `v4`, and `v5` in the code below (assuming the code is part of a complete program)?
- ```
int g(int x);
...
float m = 5.5f;
float & rm = m;
decltype(m) v1 = m;
decltype(rm) v2 = m;
decltype((m)) v3 = m;
decltype (g(100)) v4;
decltype (2.0 * m) v5;
```

## Programming Exercises

1. Write a function that normally takes one string argument and prints that string once. However, if a second argument is nonzero, the function should print the string the number of times that function has been called. The number of times the string is printed is not equal to the number of times the function is called.

```
Next string (q to quit): q
Bye.
```

4. The following is a program skeleton:

```
#include <iostream>
using namespace std;
#include <cstring> // for strlen
struct stringy {
 char * str; // points to
 int ct; // length of
};

// prototypes for set(), show(), and
int main()
{
 stringy beany;
 char testing[] = "Reality isn't

 set(beany, testing); // first
 // allocates space
 // sets str member
 // new block, copies
 // and sets ct member
```

6. Write a template function `maxn()` that takes a reference to an array of type `T` and an integer representing the number of elements. The function returns the largest item in the array. Test the template with an array of six `int` values and a string array. The program should also include a specialization for strings that takes an argument and the number of pointers to strings. It returns the address of the longest string. If multiple strings have the same length, the function should return the address of the first. Test the specialization with an array of five strings.
7. Modify Listing 8.14 so that it uses two template functions. One should return the sum of the array contents in `double`. The other program now should report the total number of elements.

C++ offers many choices for storing data. The duration of data remains in memory (storage duration) and the visibility of data (access to data (scope and linkage)). You can also use the `extern` keyword and placement `new` offers a variation on that. The `volatile` keyword provides additional control over access. Larger programs are often split into multiple source code files that may share some data in common. The linker performs compilation of the program files, and this chapter

## Separate Compilation

C++, like C, allows and even encourages you to write a program in separate files. As Chapter 1, “Getting Started,” explains, you can compile the files separately and then link them together. The C++ compiler typically compiles programs as a single unit. If you modify just one file, you can recompile just that file and then link the compiled versions of the other files. This facilitates incremental compilation. Furthermore, most C++ environments provide a build system for file management. Unix and Linux systems, for example, use `make` to track which files a program depends on and when they were last compiled, and it detects that you’ve changed one or more files. `make` then make remembers the proper steps needed to recompile the program. Many development environments (IDEs), including

that use those structures

- A source code file that contains the code for those functions
- A source code file that contains the code for those functions

This is a useful strategy for organizing a project. If you have a program that uses those same functions, you can include the functions file to the project or make list. Also, this is a common approach. One file, the header file, contains the function prototypes. The second file contains the function code for those functions. When they form a package you can use for a variety of programs.

You shouldn't put function definitions or variable definitions in a header file. It might work for a simple setup but usually it leads to errors. If you put a function definition in a header file and then include that header file in multiple programs that are part of a single program, you'd wind up with multiple definitions in a single program, which is an error, unless the program is specifically designed to handle this. This is commonly found in header files:

- Function prototypes
- Symbolic constants defined using `#define`
- Structure declarations
- Class declarations

compiler (bcc32.exe) also behave that way. And Borland and Microsoft Visual C++ go through essentially the same process. In Chapter 1, you initiate the process differently and associate source code files with it. Note that you add header files, to projects. That's because the #include directive can be used to include multiple declarations.

### Caution

In IDEs, don't add header files to the project. Instead, add source code files in other source code files.

#### Listing 9.1 **coordin.h**

```
// coordin.h -- structure templates and functions
// structure templates
#ifndef COORDIN_H_
#define COORDIN_H_

struct polar
{
 double distance; // distance from origin
 double angle; // direction from positive x-axis
};
```



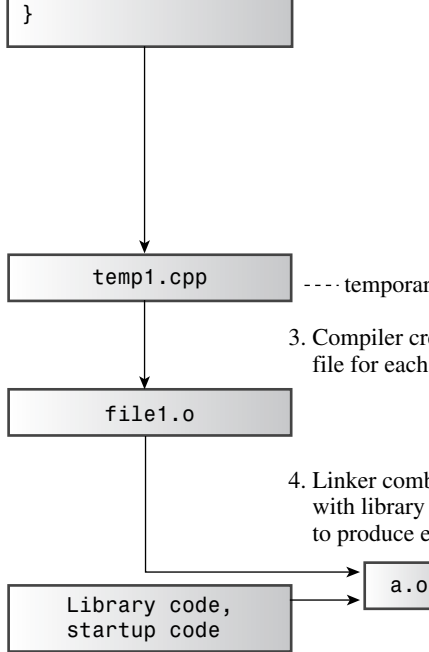


Figure 9.1 Compiling a multifile

```
#define COORDIN_H_
// place include file contents here
#endif
```

The first time the compiler encounters the file `coordin.h` (I chose a name based on the `include` files in the directory to create a name that is unlikely to be defined elsewhere). The compiler looks at the material between the `#ifn` and `#endif`. In the process of looking at the material, the compiler defines `COORDIN_H_`. If it then encounters a second inclusion of `coordin.h`, it sees that `COORDIN_H_` is defined and skips to the end of the file. This method doesn't keep the compiler from including the file, but it ignores the contents of all but the first inclusion. Most C++ files use this guarding scheme. Otherwise you could include the same file more than once, and that will produce a compile error.

## Listing 9.2 **file1.cpp**

```
// file1.cpp -- example of a three-file program
#include <iostream>
#include "coordin.h" // structure template
using namespace std;
int main()
```

```
// convert rectangular to polar coordinates
polar rect_to_polar(rect xypos)
{
 using namespace std;
 polar answer;

 answer.distance =
 sqrt(xypos.x * xypos.x + xypos.y
 * xypos.y);
 answer.angle = atan2(xypos.y, xypos.x);
 return answer; // returns a polar
}

// show polar coordinates, converting angles
void show_polar (polar dapos)
{
 using namespace std;
 const double Rad_to_deg = 57.29577951;

 cout << "distance = " << dapos.distance;
 cout << ", angle = " << dapos.angle *
 Rad_to_deg << " degrees\n";
}
```

---

If you are provided with the source code, you compile the source with your compiler.

## Storage Duration, Scope,

Now that you've seen a multifile program, it's time to look at memory schemes in Chapter 4, "Compound Types and Information." Information can be shared across files. It might be useful to review Chapter 4, so let's review what it says about memory (four under C++11) for storing data, and then look at data in memory:

- **Automatic storage duration**—Variables with automatic storage duration, including function parameters—have a storage duration that ends when program execution enters the function's return statement, and the memory used for them is freed. C++ has two kinds of automatic storage duration.
- **Static storage duration**—Variables declared with the keyword `static` have static storage duration, which lasts for the entire time a program is running. C++ has two kinds of static storage duration.
- **Thread storage duration (C++11)**—Variables with thread storage duration are local to a thread. These are CPUs that can handle simultaneous execution. This allows a program to split computations

A C++ variable can have one of several scopes. A variable with *block scope* is known only within the block in which it is defined. A block is a series of statements enclosed in braces, and it is possible to have blocks nested within blocks, but you can have other blocks nested within the same block. A variable with *file scope* (also termed *file scope*) is known throughout the entire source file in which it is defined. Automatic variables have local scope, but their scope is limited to the block in which they are defined, depending on how it is defined. Names used in function arguments are known only within the parentheses enclosing the argument list (regardless of what they are or if they are even present.) (See Chapter 10, “Objects and Classes”). Variables with *namespace scope* (Chapter 10, “Objects and Classes”). Variables with *namespace scope* (Chapter 10, “Objects and Classes”). Variables with *namespace scope* (Chapter 10, “Objects and Classes”). (Now that namespaces have been added to the language, *namespace scope* becomes a special case of namespace scope.)

C++ functions can have class scope or namespace scope, but they can't have local scope. (Because a function with local scope would have to be called by another function. Such a function could not be called.)

The various C++ storage choices are characterized by their storage duration, their scope, and their linkage. Let's look at C++'s storage choices. We'll begin by examining the situation before namespaces were introduced, and then how namespaces modify the picture.

```

 cout << "hello\n";
 int websight = -2; // websight
 cout << websight << ' ' << teledeli;
 } // websight
 cout << teledeli << endl;
 ...
} // teledeli expires

```

But what if you name the variable in the inner block the same as the variable in the outer block? In this case, the program interprets the inner block variable while the program executes statements in the inner block. The new definition *hides* the prior definition. The new definition is temporarily out of scope. When the program comes back into scope (see Figure 9.2).

Listing 9.4 illustrates how automatic variables are used in blocks that contain them.

**Listing 9.4    autoscp.cpp**

```
// autoscp.cpp -- illustrating scope of au
#include <iostream>
void oil(int x);
int main()
{
 using namespace std;

 int texas = 31;
 int year = 2011;
 cout << "In main(), texas = " << texas
 cout << &texas << endl;
 cout << "In main(), year = " << year <
 cout << &year << endl;
 oil(texas);
 cout << "In main(), texas = " << texas
 cout << &texas << endl;
 cout << "In main(), year = " << year <
 cout << &year << endl;
 return 0;
}
```

```
In main(), year = 2011, &year = 0012FEC8
In oil(), texas = 5, &texas = 0012FDE4
In oil(), x = 31, &x = 0012FDF4
In block, texas = 113, &texas = 0012FDD8
In block, x = 31, &x = 0012FDF4
Post-block texas = 5, &texas = 0012FDE4
In main(), texas = 31, &texas = 0012FED4
In main(), year = 2011, &year = 0012FEC8
```

Notice that each of the three `texas` variables has a different address, and that the program uses only the particular `texas` variable containing the value 113 to the `texas` in the inner block. The other `texas` variables of the same name. (As usual, the address changes as the program runs.)

Let's summarize the sequence of events. When the program starts, it allocates space for `texas` and `year`, and these variables are in memory. When `oil()` is called, these variables remain in memory but new variables for `texas` and `x` are allocated and come into scope. When the block in `oil()` ends, the new `texas` passes out of scope and an even newer definition. The variable `x`, however, remains in scope. When the block defines a new `x`, the previous `x` is freed, and `texas #2` comes back into scope. When the program ends, `texas` and `x` expire, and the original `texas` and `year` remain.





and compilers developed in sophistication, the `extern` keyword was heavily used and perhaps the compiler didn't even require it. With C++11, even that hint is being dropped. You can explicitly identify a variable as being automatic, but you have to do so with variables that would be automatic anyway. It's a way to indicate that you really do want to use an automatic variable, not as an external variable. This is the same purpose as the `register` keyword, an important reason for retaining `register`, however, is that it uses that keyword.

## Static Duration Variables

C++, like C, provides static storage duration variables with three linkage (accessible across files), internal linkage (accessible within a file), and no linkage (accessible to just one function). All three last for the duration of the program; they are not destroyed. Because the number of static variables a program doesn't need a special device such as a symbol table. It just allocates a fixed block of memory to hold all static variables present as long as the program executes. Also, for a static variable, the compiler sets it to 0. Static arrays and static member variables are set to 0 by default.

```
{
 ...
}
```

4. Stack after function terminates

Figure 9.3 Passing arguments

### Note

Classic K&R C does not allow you to initialize static arrays and structures. But some older C++ translators use C++ static storage classes for initializing arrays and structures.

```
}
```

As stated previously, all the static duration (this example) persist from the time the program starts until the program ends. The variable `count`, which is declared inside `func1()`, means it can be used only inside the `func1()` function. But unlike `llama`, `count` remains in memory even when `func1()` is not being executed. Both `global` and `one_file` variables are visible from the point of declaration until the end of the program. In `main()`, `func1()`, and `func2()`. Because `one_file` is only in the file containing this code. Because `global` is used in other files that are part of the program.

All static duration variables share the following characteristics. A static variable has all its bits set to 0. Such a variable is initialized to 0.

Table 9.1 summarizes the storage class features. In the next section, we'll examine the static duration varieties in more detail.

Note that the keyword `static` has somewhat different meanings, as shown in Table 9.1. When used with a local variable, `static` has no linkage, `static` indicates the kind of storage class. When used outside of a block, `static` indicates internal linkage. This is a bit of a *keyword overloading* situation. One might term this *keyword overloading* in the context of storage classes.

appropriate type. For example, the null pointer may have a nonzero internal representation, so that value. Structure members are zero-initialized.

Zero-initialization and constant-expression *initialization*. This means the variable is initialized (at translation unit). Dynamic initialization means

So what determines which form of initialization variables are zero-initialized, whether or not any variable is initialized using a constant expression from the file contents (including included header files) initialization. The compiler is prepared to do simple enough information at this time, the variable is initialized following:

```
#include <cmath>
int x; // zero-initialized
int y = 5; // constant expression
long z = 13 * 13; // constant expression
const double pi = 4.0 * atan(1.0); // dynamic initialization
```

First, `x`, `y`, `z`, and `pi` are zero-initialized. Then `y` and `z` are initialized using constant expressions and initializes `y` and `z` to 5 and 169.

On the other hand, C++ has the one definition rule that there can be only one definition of a variable. There are two kinds of variable declarations. One is the one that causes storage for the variable to be allocated. The other, simply, a *declaration*. It does not cause storage for the variable.

A referencing declaration uses the keyword `extern`. Otherwise, a declaration is a definition and causes storage to be allocated.

```
double up; // definition, up defined elsewhere
extern int blem; // blem defined elsewhere
extern char gr = 'z'; // definition because of initialization
```

If you use an external variable in several files, you must declare that variable (per the one definition rule). But you must declare that variable using the keyword `extern` in all but one file.

```
// file01.cpp
extern int cats = 20; // definition because of initialization
int dogs = 22; // also a definition
int fleas; // also a definition
...
// file02.cpp
// use cats and dogs from file01.cpp
extern int cats; // not definitions
```

```
void promise ();
int main()
{
 ...
}

void promise ()
{
 ...
}
```

This file defines the variable `process_status`, causing the compiler to allocate space for it.

Figure 9.4 Defining declarati

Note that the one definition rule doesn't m with a given name. For example, automatic va in different functions are separate variables, in its own address. Also as later examples show, a the same name. However, although a program name, each version can have only one definiti

```
 cout << "Global warming is " << warming;
 update(0.1); // call function
 cout << "Global warming is " << warming;
 local(); // call function
 cout << "Global warming is " << warming;
 return 0;
 }
```

---

### Listing 9.6    **support.cpp**

---

```
// support.cpp -- use external variable
// compile with external.cpp
#include <iostream>
extern double warming; // use warming from main

// function prototypes
void update(double dt);
void local();

using std::cout;
void update(double dt) // modifies global variable
{
 extern double warming; // optional redeclaration
```



Global warming is 0.4 degrees.

## Program Notes

The output of the program in Listings 9.5 and 9.6 shows that `update()` can access the external variable `warming`. The change made to `warming` shows up in subsequent uses of `warming`.

The definition for `warming` is in Listing 9.5.

```
double warming = 0.3; // warming definition
```

Listing 9.6 uses `extern` to make the `warming` variable available in that file:

```
extern double warming; // use warming from other file
```

As the comment indicates, this declaration is only a declaration, not a definition, and it is used to tell the compiler that `warming` is defined externally elsewhere.”

In addition, the `update()` function re-declares `warming` with the keyword `extern`. This keyword means “Use the variable `warming` defined externally.” Because that is what `update()` wants, the `extern` declaration, this declaration is optional. It serves to tell the compiler to use the external variable.

The `local()` function demonstrates that when a function has the same name as a global variable, the local version of the function

constant data because you can use the keyword `const` to

```
const char * const months[12] =
{
 "January", "February", "March", "April",
 "June", "July", "August", "September",
 "November", "December"
};
```

In this example, the first `const` protects the data, and the second makes sure that each pointer in the array remains pointed initially.

## Static Duration, Internal Linkage

Applying the `static` modifier to a file-scope variable creates a difference between internal linkage and external linkage. In that context, a variable with internal linkage is only visible within the file. But a regular external variable has external linkage and is visible in other files, as the previous example showed.

What if you want to use the same name to declare a variable in multiple files? Can you just omit the `extern`?

```

void flush()
{
 cout << errors; // uses errors def
 ...
}

```

This doesn't violate the one definition rule that the identifier `errors` has internal linkage and a single definition.

## Note

In a multifile program, you can define an external variable in one file. Files using that variable have to declare that variable.

You can use an external variable to share data between files in a program. You can use a static variable with internal linkage, which is found in just one file. (Namespaces offer an alternative to a file-scope variable. If you use `static`, you needn't worry about variables found in other files.)

Listings 9.7 and 9.8 show how C++ handles external linkage. Listing 9.7 (`twofile1.cpp`) defines the external variable `harry`. The `main()` function uses `harry` and then calls the `remote_access()` function. Listing 9.8 (`twofile2.cpp`) shows that file. In this file, `tom` uses the `extern` keyword to share `tom` with

```
 cout << &tom << " = &tom, " << &dick
 cout << &harry << " = &harry\n";
 remote_access();
 return 0;
 }
```

---

### Listing 9.8    **twofile2.cpp**

---

```
// twofile2.cpp -- variables with internal linkage
#include <iostream>

extern int tom; // tom defined elsewhere
static int dick = 10; // overrides external definition
int harry = 200; // external variable
 // no conflict with static dick

void remote_access()
{
 using namespace std;
 cout << "remote_access() reports the
 cout << &tom << " = &tom, " << &dick
 cout << &harry << " = &harry\n";
}
```

---

tion—you could use them to pass secret account information.) Also if you initialize a static local variable once, when the program starts up. Subsequent accesses to the variable are done the way they do for automatic variables.

### Listing 9.9    **static.cpp**

---

```
// static.cpp -- using a static local variable
#include <iostream>
// constants
const int ArSize = 10;

// function prototype
void strcount(const char * str);

int main()
{
 using namespace std;
 char input[ArSize];
 char next;

 cout << "Enter a line:\n";
 cin.get(input, ArSize);
 while (cin)
 {
```

```
}
```

Incidentally, the program in Listing 9.9 should not exceed the size of the destination array. Recall that the `getline` method reads up to the end of the line or up to the specified size, whichever is first. It leaves the newline character in the input buffer. To read the character that follows the line input, you must make a preceding call to `cin.get(input, ArSize)`. If the input buffer contains a newline character, there are more characters left in the buffer. You can choose to reject the rest of the line, but you can modify the program to skip the next input cycle. The program also uses the `getline` method. The `get(char *, int)` causes `cin` to test as if the input buffer is empty.

Here is the output of the program in Listing 9.9:

```
Enter a line:
```

```
nice pants
```

```
"nice pants" contains 9 characters
```

```
9 characters total
```

```
Enter next line (empty line to quit):
```

```
thanks
```

```
"thanks" contains 6 characters
```

```
15 characters total
```

```
static
extern
thread_local (added by C++11)
mutable
```

You've already seen most of these, and you've seen the `register` keyword, except that `thread_local` can't be used prior to C++11, the keyword `auto` could be used to indicate that a variable is an automatic variable. (In C++11, `auto` is used to deduce the type of a variable.) The keyword `register` is used in a declaration to indicate that a variable is a register variable. In C++11, simply is an explicit way of saying `static`, when used with a file-scope declaration. When used with a local declaration, it indicates static storage duration. `extern` indicates a reference declaration—that is, the variable is defined elsewhere. The keyword `thread_local` indicates that the variable's lifetime is the duration of the containing thread. A thread-local variable is a regular static variable is to the whole program. The keyword `const` is used to declare a constant variable, so let's look at the cv-qualifiers first.

## Cv-Qualifiers

Here are the cv-qualifiers:

```
const
volatile
```

```

{
 char name[30];
 mutable int accesses;
 ...
};
const data veep = {"Claybourne Clodde", 0};
strcpy(veep.name, "Joye Joux"); // not
veep.accesses++; // allo

```

The `const` qualifier to `veep` prevents a program from adding the `mutable` specifier to the `accesses` member.

This book doesn't use `volatile` or `mutable`.

## More About `const`

In C++ (but not C), the `const` modifier alters the linkage of a global variable that has external linkage by default. That is, C++ treats a global constant variable as if the `static` specifier had been used.

```

const int fingers = 10; // same as static
int main(void)
{
 ...
}

```



ing them. Each definition is private to the file it is in. If you put constant definitions in a header file. Then include that header file in two source code files, they receive two definitions.

If, for some reason, you want to make a constant global, you can use the `extern` keyword to override the default internal linkage.

```
extern const int states = 50; // definition
```

You then must use the `extern` keyword to declare the constant. This differs from regular external variable linkage keyword `extern` when you define a variable, because you must initialize the variable. Keep in mind, however, now that a single source file can use initialization.

When you declare a `const` within a function, the constant is usable only when the program is in that function. This means that you can create constants within a function. Be aware, however, about the names conflicting with constants declared elsewhere.

## Functions and Linkage

Like variables, functions have linkage properties. The difference is greater than for variables. C++, like C, does not allow functions to have static storage class, so all functions automatically have static storage class. This means that, as long as the program is running. By default, functions can be shared across files. You can, in fact, use

each file that uses the function should have a definition in a header file. Thus, each file that includes the header file must have an inline function definition. However, C++ does not require that all inline function definitions for a particular function be identical.

## Where C++ Finds Functions

Suppose you call a function in a particular file. Where does the compiler look for the function definition? If the function prototype is in a header file, the compiler looks only in that file for the function definition. The linker (and the linker, too) looks in all the program files. If you get an error message because you can have only one definition, it means the linker fails to find any definition in the files, the function is not defined. If you define a function that has the same name as a function in a standard library, that if you define a function that has the same name as a function in a standard library, it uses your version rather than the library version. If you define a function that has the same name as a standard library function, so you shouldn't need to give explicit instructions to identify which libraries

## Language Linking

Another form of linking, called *language linking*, is used to link a program with a library. A linker needs a different symbolic name for each function in a library. For example, to implement because there can be only one C function for each purpose, a C compiler might translate a C function

C and C++ language linkage are the only  
But implementations have the option of provid

## Storage Schemes and Dynamic Al

You've seen the five schemes, excluding thread  
for variables (including arrays and structures).  
using the C++ `new` operator (or by using the  
kind of memory *dynamic memory*. As you saw i  
by the `new` and `delete` operators, not by scop  
can be allocated from one function and freed  
memory, dynamic memory is not LIFO; the c  
when and how `new` and `delete` are used. Typi  
ory chunks: one for static variables (this chunk  
variables, and one for dynamic storage.

Although the storage scheme concepts don  
to automatic and static pointer variables used  
example, suppose you have the following state

```
float * p_fees = new float [20];
```

The 80 bytes (assuming that a `float` is 4 b  
memory until the `delete` operator frees it. Bu  
when program execution exits the block cont

enclosed in parentheses:

```
int *pi = new int (6); // *pi set to 6
double * pd = new double (99.99); // *pd
```

The parentheses syntax also can be used w  
we haven't got that far yet.

To initialize an ordinary structure or an ar  
initialization using braces. The new standard :

```
struct where {double x; double y; double
where * one = new where {2.5, 5.3, 7.2};
int * ar = new int [4] {2,4,6,7};
```

With C++11, you also can use the brace i

```
int *pin = new int {}); // *pi set to 6
double * pdo = new double {99.99}; // *p
```

## When new Fails

It may be that new can't find the requested ar  
handled that eventuality by having new return  
throws a `std::bad_alloc` exception. Chapte  
vides some short examples showing how each

```
int * pa = new(40 * sizeof(int));
```

As you've seen, a statement with a `new` operator, so, in general, using the `new` operator may do

Similarly,

```
delete pi;
```

invokes the following function call:

```
delete (pi);
```

Interestingly, C++ terms these functions *replacement new*. If, by expertise and desire, you can supply replacement new functions to meet your specific requirements. One can define replacement functions with class scope so that they are particular to a particular class. Your code would use the `new` operator, but it would call upon the replacement `new()` function.

## The Placement new Operator

Normally, the `new` operator has the responsibility of allocating memory that is large enough to handle the amount of memory requested. It has a variation, called *placement new*, that allows a programmer to specify a memory location. A programmer might use this feature to set up

```
// now, the two forms of placement new
 p2 = new (buffer1) chaff; // place
 p4 = new (buffer2) int[20]; // place
...

```

For simplicity, this example uses two static memory buffers. So this code allocates space for a `chaff` array in `buffer1` and an array of 20 ints in `buffer2`.

Now that you've made your acquaintance with placement new, let's write a program. Listing 9.10 uses both `new` and placement `new` to create arrays. This program illustrates some important features of placement `new` that we'll discuss after seeing the output.

#### Listing 9.10 **newplace.cpp**

---

```
// newplace.cpp -- using placement new
#include <iostream>
#include <new> // for placement new
const int BUF = 512;
const int N = 5;
char buffer[BUF]; // chunk of memory
int main()
{
 using namespace std;

```

```

 for (i = 0; i < N; i++)
 {
 cout << pd3[i] << " at " << &pd3[i] << "\n";
 cout << pd4[i] << " at " << &pd4[i] << "\n";
 }

 cout << "\nCalling new and placement new\n";
 delete [] pd1;
 pd1= new double[N];
 pd2 = new (buffer + N * sizeof(double)) double[N];
 for (i = 0; i < N; i++)
 pd2[i] = pd1[i] = 1000 + 60.0 * i;
 cout << "Memory contents:\n";
 for (i = 0; i < N; i++)
 {
 cout << pd1[i] << " at " << &pd1[i] << "\n";
 cout << pd2[i] << " at " << &pd2[i] << "\n";
 }
 delete [] pd1;
 delete [] pd3;
 return 0;
}

```

---

1060 at 006E4AB8; 1060 at 00FD9168  
1120 at 006E4AC0; 1120 at 00FD9170  
1180 at 006E4AC8; 1180 at 00FD9178  
1240 at 006E4AD0; 1240 at 00FD9180

## Program Notes

The first thing to note about Listing 9.10 is the `pd2` array in the `buffer` array; both `pd2` and `buffer` are, however, of different types; `pd1` is pointer-to-double, in other words, the way, that's why the program uses a `(void*)` cast to try to display a string.) Meanwhile, regular new memory, at location 006E4AB0, which is part of the

The second point to note is that the second new block of memory—one beginning at 00FD9138—new results in the same block of memory being freed, beginning at 00FD9138. The important fact here is that `new` is passed to it; it doesn't keep track of where the memory is and it doesn't search the block for unused memory. It's up to the programmer to manage memory. For example, the program provides an offset into the `buffer` array so that

```
pd2 = new (buffer + N * sizeof(double)) double[N];
```



## Other Forms of Placement new

Just as regular new invokes a new function with no arguments, placement new invokes a new function with two arguments:

```
int * p1 = new int; // invokes new(int)
int * p2 = new(buffer) int; // invokes new(buffer, int)
int * p3 = new(buffer) int[40]; // invokes new(buffer, int, 40)
```

The placement new function is not replaced by an overloaded new function. It takes at least two parameters, the first of which always specifies the memory location, and the second, the number of bytes requested. Any such overload must have the same signature as new, or if the additional parameters don't specify a location, it must be new(int).

## Namespaces

Names in C++ can refer to variables, functions, and structure members. When programming projects grow, the number of name conflicts increases. When you use class libraries, the number of name conflicts increases. For example, two libraries might both define a class named `Node`, but in incompatible ways. You might use `Node` from one library, and `Tree` from the other, and each might expect the other to use `Node`. These problems are termed *namespace problems*.

the declarative region is the file) usually see the variable is termed the *scope*, which along. Figures 9.5 and 9.6 illustrate the terms

C++’s rules about global and local variables. A declarative region can declare names that are visible in other declarative regions. A local variable declared in one function is not a variable declared in a second function.

## New Namespace Features

C++ now adds the ability to create named namespaces. In a namespace, one whose main purpose is to prevent names in one namespace don’t conflict with names in other namespaces, and there are mechanisms for letting one namespace use names from another namespace. The following code, for example, defines two namespaces, Jack and Jill:

```
namespace Jack {
 double pail; //
 void fetch(); //
 int pal; //
 struct Well { ... }; //
}
namespace Jill {
```

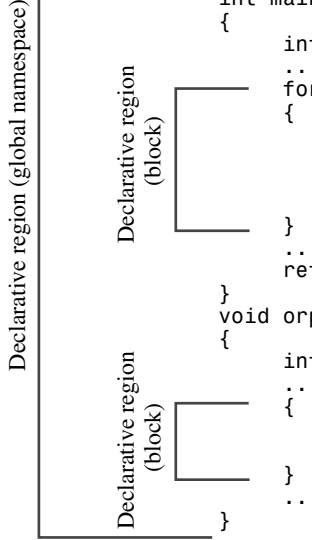


Figure 9.5 Declarative regions

Namespaces are *open*, meaning that you can add to them. For example, the following statement adds the namespace `Jill`:

```
namespace Jill {
 char * goose(const char *);
}
```

Similarly, the original `Jack` namespace provides the function `fetch()`. You can provide the code for the function later, in a separate `Jack` namespace again:

```
namespace Jack {
 void fetch()
 {
 ...
 }
}
```

```

namespace Jill {
 double bucket(double n) { ... }
 double fetch;
 struct Hill { ... };
}
char fetch;
int main()
{
 using Jill::fetch; // put fetch into
 double fetch; // Error! Already
 cin >> fetch; // read a value i
 cin >> ::fetch; // read a value i
 ...
}

```

Because a using declaration adds the name, it precludes creating another local variable by the same name. If we had a local variable, fetch, it would override a global variable.

Placing a using declaration at the external

```

void other();
namespace Jill {
 double bucket(double n) { ... }
 double fetch;
 struct Hill { ... };
}

```

```
#include <iostream> // places names in
using namespace std; // make names avail
```

Placing a using directive in a particular function. Here's an example:

```
int main()
{
 using namespace jack; // make names a
 ...
}
```

You've seen this form often in this book with

One thing to keep in mind about using directives is that they increase the possibility of name conflicts. That is, if you use namespace jill available, and you use the scope

```
jack::pal = 3;
jill::pal = 10;
```

The variables jack::pal and jill::pal are in different memory locations. However, if you employ using directives, you can

```
using jack::pal;
using jill::pal;
pal = 4; // which one? now have
```

```

int main()
{
 using namespace Jill; // import a
 Hill Thrill; // create a
 double water = bucket(2); // use Jill
 double fetch; // not an e
 cin >> fetch; // read a v
 cin >> ::fetch; // read a v
 cin >> Jill::fetch; // read a v
 ...
}

int foom()
{
 Hill top; // ERROR
 Jill::Hill crest; // valid
}

```

Here, in `main()`, the name `Jill::fetch` is have local scope, so it doesn't override the glo hides both `Jill::fetch` and the global `fetch` ables are available if you use the scope-resolut this example to the preceding one, which uses

This is the approach used for most of this

```
#include <iostream>
int main()
{
 using namespace std;
```

First, the `iostream` header file puts every  
directive makes the names available within ma

```
#include <iostream>
using namespace std;
int main()
{
```

This exports everything from the `std` na  
main rationale for this approach is expedien  
have namespaces, you can replace the first tw  
inal form:

```
#include <iostream.h>
```

However, namespace proponents hope th  
the scope-resolution operator or the `using` o  
following:

```
using namespace std; // avoid as too inc
```



```

 {
 int flame;

 ...
 }
 float water;
}

```

In this case, you refer to the flame variable. You can make the inner names available with this:

```
using namespace elements::fire;
```

Also you can use using directives and using declarations:

```

namespace myth
{
 using Jill::fetch;
 using namespace elements;
 using std::cout;
 using std::cin;
}

```

Suppose you want to access `Jill::fetch`. You can use the `myth` namespace, where it can be called `fetch`.  
`std::cin >> myth::fetch;`

You can make `mvft` an alias for `my_very_` statement:

```
namespace mvft = my_very_favorite_things;
```

You can use this technique to simplify using

```
namespace MEF = myth::elements::fire;
using MEF::flame;
```

## Unnamed Namespaces

You can create an *unnamed namespace* by omitting

```
namespace // unnamed namespace
{
 int ice;
 int bandycoot;
}
```

This code behaves as if it were followed by `using namespace` declared in this namespace are in potential scope of the namespace that contains the unnamed namespace. In this case, the variables are like global variables. However, if a namespace is created using `using directive` or `using declaration` to make

```

{
...
}

int other()
{
...
}

```

## A Namespace Example

Let's take a look at a multifile example that demonstrates namespaces. The first file in this example (see Listing 10-1) contains items normally found in header files—constants, typedefs, and function prototypes. In this case, the items are placed in two namespaces. The first namespace, `person`, contains a definition of a `Person` structure, plus a function that displays a person's name and a function that displays a person's debts. The second namespace, `debts`, defines a structure for storing money owed to that person. This structure uses the `person` namespace. The `debts` namespace has a `using` directive to make the names in the `person` namespace available in the `debts` namespace. The `debts` namespace also

```
};
void getDebt (Debt &);
void showDebt (const Debt &);
double sumDebts (const Debt ar[], int
}
```

---

The second file in this example (see Listing 9.12) source code file provide definitions for functions, which are declared in a namespace, has to be in the same namespace as the declarations. This comes in handy. The original namespace (refer to Listing 9.11). The file then adds the namespace as shown in Listing 9.12. Also the namespace is added to the std namespace with the using declaration and

#### Listing 9.12    **namesp.cpp**

---

```
// namesp.cpp -- namespaces
#include <iostream>
#include "namesp.h"

namespace pers
{
```

```

 void showDebt(const Debt & rd)
 {
 showPerson(rd.name);
 std::cout <<" : $" << rd.amount <<
 }
 double sumDebts(const Debt ar[], int n)
 {
 double total = 0;
 for (int i = 0; i < n; i++)
 total += ar[i].amount;
 return total;
 }
}

```

---

Finally, the third file of this program (see Listing 9.13) defines the structures and functions declared and defined in the first two files. It shows several methods of making the namespace identifier

### Listing 9.13    **usenmsp.cpp**

---

```

// usenmsp.cpp -- using namespaces
#include <iostream>
#include "namesp.h"

```

```

 int i;
 for (i = 0; i < 3; i++)
 getDebt(zippy[i]);

 for (i = 0; i < 3; i++)
 showDebt(zippy[i]);
 cout << "Total debt: $" << sumDebts(z);
 return;
 }

void another(void)
{
 using pers::Person;
 Person collector = { "Milo", "Rightsh
 pers::showPerson(collector);
 std::cout << std::endl;
}

```

---

In Listing 9.13, `main()` begins by using two

```

using debts::Debt; // makes the Debt
using debts::showDebt; // makes the sho

```

```
Enter last name: Binx
Enter debt: 100
Enter first name: Cleve
Enter last name: Delaproux
Enter debt: 120
Enter first name: Eddie
Enter last name: Fiotox
Enter debt: 200
Binx, Arabella: $100
Delaproux, Cleve: $120
Fiotox, Eddie: $200
Total debt: $420
Rightshift, Milo
```

## Namespaces and the Future

As programmers become more familiar with namespaces, new guidelines will emerge. Here are some current guidelines:

- Use variables in a named namespace inside a function.
- Use variables in an unnamed namespace inside a function.
- If you develop a library of functions or classes, C++ currently already calls for placing them in a namespace.

# Summary

C++ encourages the use of multiple files in a program. A common organizational strategy is to use a header file to define the user-defined type and for functions to manipulate the user types. You put the function definitions. Together, the header file and the source file implement the user-defined type and how it can be used. Functions using those functions can go into a third file.

C++'s storage schemes determine how long a variable exists (its duration) and what parts of a program have access to it. Automatic variables are variables that are defined within a function. They exist and are known only within the body. They exist and are known only within the block that contains the definition. Automatic variables are declared with the storage class specifier `register` or with no specifier. The `register` specifier is automatically automatic. The `register` specifier is used to indicate that a variable is heavily used, but that use is deprecated.

Static variables exist for the duration of the entire program. A static variable is known to all functions in the file. A static variable can be made available to other files in the program (extern). If you use such a variable, that file must declare it by using `extern`. A variable shared across files should have a defining declaration in one file, but it can be used if combined with initialization in other files.



- b. The secret variable is to be shared with the public.
  - c. The topsecret variable is to be shared with the public, but hidden from other files.
  - d. beencalled keeps track of how many times beencalled has been called.
2. Describe the differences between a user and a process.
3. Rewrite the following so that it doesn't have any errors.

```
#include <iostream>
using namespace std;
int main()
{
 double x;
 cout << "Enter value: ";
 while (! (cin >> x))
 {
 cout << "Bad input. Please enter a number: ";
 cin.clear();
 while (cin.get() != '\n')
 continue;
 }
}
```

- }
5. Suppose you want the average (3, 6) function to take two `int` arguments when it is called in one program. How could you set this up?
6. What will the following two-file program do?

```
// file1.cpp
#include <iostream>
using namespace std;
void other();
void another();
int x = 10;
int y;

int main()
{
 cout << x << endl;
 {
 int x = 4;
 cout << x << endl;
 cout << y << endl;
 }
}
```

```
 cout << "another(): " << x << ",
 }
```

7. What will the following program display?

```
#include <iostream>
using namespace std;
void other();
namespace n1
{
 int x = 1;
}
```

```
namespace n2
{
 int x = 2;
}
```

```
int main()
{
 using namespace n1;
 cout << x << endl;
 {
```

# Programming Exercises

1. Here is a header file:

```
// golf.h -- for pe9-1.cpp
```

```
const int Len = 40;
struct golf
{
 char fullname[Len];
 int handicap;
};
```

```
// non-interactive version:
```

```
// function sets golf structure to
// using values passed as argument.
void setgolf(golf & g, const char *
```

```
// interactive version:
```

```
// function solicits name and hand
// and sets the members of g to th
// returns 1 if name is entered, 0
int setgolf(golf & g);
```

prototyped functions. For example, a loop that iterates over the array of structures and terminate when the array reaches the end of the array for the golfer's name. The `main()` function calls the functions to access the golf structures.

2. Redo Listing 9.9, replacing the character array with a pointer. It should no longer have to check whether the input string is empty. Use the input string to check for an empty string.
3. Begin with the following structure declaration:

```
struct chaff
{
 char dross[20];
 int slag;
};
```

Write a program that uses placement new to allocate a buffer. Then assign values to the structure members (use `strcpy()` for the char array) and use a pointer to access the structure. Use a static array, like that in Listing 9.1, to store the pointers to the new buffers.

}

The first file should be a header file that should be a source code file that extends the three prototyped functions. The third should use the interactive version of `showSales` and the non-interactive version of `showSales` and structure. It should display the contents of `showSales()`.



- Class methods (also called class functions)
- Creating and using class objects
- Class constructors and destructors
- `const` member functions
- The `this` pointer
- Creating arrays of objects
- Class scope
- Abstract data types

Object-oriented programming (OOP) is a programming paradigm that uses objects and classes to structure applications. C++ is a general-purpose programming language that supports OOP, and C++ has enhanced C with features that support the OOP approach. The following are the most important features of OOP:

- Abstraction
- Encapsulation and data hiding
- Polymorphism
- Inheritance
- Reusability of code



Let's see, I want to enter the name, times  
those who don't follow baseball or softball  
divided by the player's official number of times  
player gets on base or makes an out, but only  
count as official times at bat), and all those  
Wait, the computer is supposed to make line  
out some of that stuff, such as the batting  
the results. How should I organize this? I g  
tions. Yeah, I'll make `main()` call a function  
make the calculations, and then call a third  
what happens when I get data from the net  
again. Okay, I can add a function to update  
in `main()` to select between entering, calc  
Hmmm...how am I going to represent the c  
the players' names, another array to hold t  
to hold the hits, and so on. No, that's dumb  
information for a single player and then use  
the whole team.

In short, with a procedural approach, you first  
follow and then think about how to represent  
the program running the whole season, you print  
a file and read data from a file.)

# Abstraction and Classes

Life is full of complexities, and one way we cope with them is through abstractions. You are a collection of more than just physical attributes. Your mind would say that your mind is a collection of thoughts, feelings, and experiences. It's much simpler to think of yourself as a single entity. This is a special step of representing information in terms of a single concept. To abstract the essential operational features of a system into a set of terms. In the softball statistics example, the program updates, and displays the data. From abstraction to implementation, which in C++ is a class design that implements the abstraction.

## What Is a Type?

Let's think a little more about what constitutes a type. If you subscribe to the popular stereotype, you might think of a nerd as someone with black-rimmed glasses, pocket protector full of pens, and a love for obscure facts. You might conclude that a nerd is better defined by their behavior than by their appearance. If he or she responds to an awkward social situation with a awkward silence, you don't mind stretched analogies, with a procedure to handle the situation. You think of a data type in terms of its appearance and behavior.

information yourself. In exchange for this extension, you have to custom fit new data types to match real-world

## Classes in C++

A *class* is a C++ vehicle for translating an abstract data representation and methods for manipulating it. Let's look at a class that represents stocks.

First, you have to think a bit about how to represent a stock as the basic unit and define a class to represent it. That you would need 100 objects to represent 100 different stocks you can represent a person's current holdings. The number of shares owned would be part of the class. The company would have to maintain records of such things as the purchase for tax purposes. Also it would have to keep track of the share price. It seems a bit ambitious for a first effort at defining a class, but it is a simplified view of matters. In particular, it is a simplified form to the following:

- Acquire stock in a company.
- Buy more shares of the same stock.
- Sell stock.
- Update the per-share value of a stock.
- Display information about the holdings.

definitions supply the details.

## What Is an Interface?

An *interface* is a shared framework for interaction between a computer and a printer or between the user and a program. For example, if you are a user, the user might be you and the program might be a word processor; you don't transfer words directly from the keyboard to the processor. Instead, you interact with the interface provided by the word processor. The computer shows you a character on the screen, you move the cursor on the screen. You click the mouse, and the cursor moves to the paragraph you were typing. The program interprets your intentions to specific information stored in memory.

For classes, we speak of the public interface of a class. For the class, the interacting system consists of the methods provided by whoever wrote the class. For the user, to write code that interacts with class objects, the user must know the public interface of the class objects. For example, to find the number of elements in a vector, you don't open up the object to what is inside; you use the `size()` method provided by the class creators. It turns out that the class definition of `vector` is private, but the public interface between the user and a `vector` object is public. The `size()` method is part of the `vector` class public interface.

```
// stock00.h -- Stock class interface
// version 00
#ifndef STOCK00_H_
#define STOCK00_H_

#include <string>

class Stock // class declaration
{
private:
 std::string company;
 long shares;
 double share_val;
 double total_val;
 void set_tot() { total_val = shares *
public:
 void acquire(const std::string & co, l
 void buy(long num, double price);
 void sell(long num, double price);
 void update(double price);
 void show();
}; // note semicolon at the end

#endif
```

---

tions come later in the implementation file, but they define the public and protected function interfaces. The binding of data and member functions is a defining feature of the class. Because of this design, creating an object defines the rules governing how that object can be used.

You've already seen how the `istream` and `ostream` classes use member functions such as `get()` and `getline()`. The function prototypes in the header file demonstrate how member functions are established. For example, `istream` has a `getline()` prototype in the `istream` header file.

## Access Control

Also new are the keywords `private` and `public` to control access to class members. Any program that uses an object can access the public portions directly. A program can access the private portions only through public member functions (or, as you'll see in Chapter 11, through a friend function). For example, the only way to access the `stock` class is to use one of the `stock` member functions. The `stock` class acts as go-betweens between a program and an object. The `stock` class is an interface between object and program. This interface is called *data hiding*. (C++ provides a technique called *encapsulation* that we'll discuss when we cover class inheritance in Chapter 11. See Figure 10.1.) Whereas data hiding may be an optional feature in a prospectus, it's a good practice in computing.

that constitute the public interface for the class (abstraction)

Figure 10.1 T

A class design attempts to separate the public interface from the implementation. The public interface represents the operations that are available to the user, separating the implementation details together and hiding them from the user. This is called *encapsulation*. *Data hiding* (putting data in a private section of the class) is an instance of encapsulation, and so is hiding function details from the user. This is the usual practice of placing class function declarations in a header file and their implementation in a source file.

## OOP and C++

OOP is a programming style that you can use to design a program. In C++, you can incorporate many OOP ideas into ordinary C++ code. Listing 9.1 provides an example (see Listings 9.1, 9.2, 9.3). In C++, you can declare a prototype along with the prototypes for functions. In C++, a struct or class function simply defines variables of that structure or class. In C++, `main()` does not directly call any function. In that example defines an abstract type that places data and functions in types in a header file, hiding the actual data and functions.

## Member Access Control: Public or Private

You can declare class members, whether they are data or functions, in the public or the private section of a class. One of the goals of OOP is to hide the data, data items normally called attributes, and the functions that constitute the class interface. Other programs can call those functions from a program. As the standard practice is to put the member functions in the private section. You can use them from a program, but the public methods can use them. This is a way to handle implementation details that don't fit the public interface.

You don't have to use the keyword `private` to indicate the default access control for class objects:

```
class World
{
 float mass; // private by default
 char name[20]; // private by default
public:
 void tellall(void);
 ...
};
```

However, this book explicitly uses the `private` keyword for data hiding.



(::) to indicate to which class the function belongs. The `update()` member function looks like this:

```
void Stock::update(double price)
```

This notation means you are defining the `update()` member function for the `Stock` class. Not only does this identify `update()` as a member function, it also uses the same name for a member function for another class. A member function for a `Buffoon` class would have this form:

```
void Buffoon::update()
```

Thus, the scope-resolution operator resolves the ambiguity of which method definition applies. We say that the identifier `update()` for member functions of the `Stock` class can, if necessary, be qualified with the scope-resolution operator. That's because `update()` is in scope. Using `update()` outside the class, however, requires special measures, which we'll discuss later.

One way of looking at method names is that they are qualified names that include the class name. `Stock::update()` is qualified, while `update()` is a simple `update()`. On the other hand, is an abbreviation for a qualified name—one that can be used just in class scope.

The second special characteristic of method names is that they are members of a class. For example, the `show()` method is a member of the `Stock` class.

```

 {
 company = co;
 if (n < 0)
 {
 std::cout << "Number of shares ca
 << company << " shares
 shares = 0;
 }
 else
 shares = n;
 share_val = pr;
 set_tot();
 }

```

```

void Stock::buy(long num, double price)
{
 if (num < 0)
 {
 std::cout << "Number of shares pu
 << "Transaction is aborted.\
 }
 else
 {
 shares += num;
 }
}

```

```

 }
}

void Stock::update(double price)
{
 share_val = price;
 set_tot();
}

void Stock::show()
{
 std::cout << "Company: " << company
 << " Shares: " << shares << "\n"
 << " Share Price: $" << share_price << "\n"
 << " Total Worth: $" << total_worth << "\n";
}

```

---

## Member Function Notes

The `acquire()` function manages the first acquisition of shares, whereas `buy()` and `sell()` manage adding to or subtracting from the shares. The `buy()` and `sell()` methods make sure that the number of shares is not a negative number. Also if the user attempts to sell more shares than are currently held, the program will terminate with an error message.

the class implementation section.

```
class Stock
{
private:
 ...
 void set_tot(); // definition kept s
public:
 ...
};

inline void Stock::set_tot() // use inli
{
 total_val = shares * share_val;
}
```

The special rules for inline functions require that the functions be defined where they are used. The easiest way to make sure that in a multifile program is to include the inline definitions in the source file in which the corresponding class is defined. (Some developers use preprocessor directives that allow the inline definitions to go into a separate file.)

Incidentally, according to the *rewrite rule*, defining a function as inline is equivalent to replacing the method definition with the function definition as an inline function immediately after the class definition.

## Note

When you call a member function, it uses the `self` parameter to invoke the member function.

Similarly, the function call `kate.sell()` invokes `sell()` on the `kate` object, causing that function to get `kate` as the `self` parameter.

Each new object you create contains storage for its own members. But all objects of the same class share one copy of each method. Suppose, for example, that `kate.shares` occupies one chunk of memory and `kate.show()` occupies another chunk of memory. But `kate.show()` and `john.show()` are the same method—that is, both execute the same block of code on their respective data. Calling a member function is what sends a message to an object. Thus, sending the same message to two different objects applies it to two different objects (see Figure 1).

## Using Classes

In this chapter you've seen how to define a class and how to produce a program that creates and uses objects. In this section, we'll use classes as similar as possible to using the basic, built-in types. To create a class object by declaring a class variable,

Figure 10.2 Objects, da

class type. You can pass objects as arguments, and you can assign one object to another. C++ provides functions `cin` and `cout` to recognize objects, and even provides functions to create objects of similar classes. It will be a while before you can use objects now with the simpler properties. Indeed, you can create objects and call a member function. Listing 10.3 provides a simple implementation file. It creates a `Stock` object, and it tests the features built in to the `Stock` class. The techniques for multifile programs described in this book, “C++,” and in Chapter 9. In particular, compilation and linking are present in the same directory or folder.

### Listing 10.3 `usestock0.cpp`

---

```
// usestock0.cpp -- the client program
// compile with stock00.cpp
#include <iostream>
#include "stock00.h"
```

```
Company: NanoSmart Shares: 35
 Share Price: $18.125 Total Worth: $634.375
Company: NanoSmart Shares: 300035
 Share Price: $40.125 Total Worth: $12037.875
Company: NanoSmart Shares: 35
 Share Price: $0.125 Total Worth: $4.375
```

Note that `main()` is just a vehicle for testing. If the `Stock` class works as you want it to, you can use it in other programs. The critical point in using the new type is that it does what you want it to do; you shouldn't have to think about the details. This is the point of the sidebar, "The Client/Server Model."

## The Client/Server Model

OOP programmers often discuss program design in terms of conceptualization, the *client* is a program that uses the class methods, constitute the *server*, and the programs that need it. The client uses the server's interface. This means that the client's only responsibility is to know that interface. The server's designer's responsibility, is to see that the server performs according to that interface. Any change in the design should be to details of implementation. Programmers to improve the client and the server independently, without the server having unforeseen repercussions on the client.

Listing 8.8, using return values for the setting

```
std::streamsize prec =
 std::cout.precision(3); // save prec
...
std::cout.precision(prec); // r

// store original flags
std::ios_base::fmtflags orig = std::cout.
...
// reset to stored values
std::cout.setf(orig, std::ios_base::float
```

As you may recall, `fmtflags` is a type defined in the `std` namespace, hence the rather long name. The `setf` method sets the flags, and the reset statement uses that info to reset the `floatfield` section, which includes flags for formatting. Third, let's not worry too much about the details of the changes; they are confined to the implementation of the `cout` aspects of the program using the class.



After this replacement and leaving the header file, recompile the program. Now the output would be:

```
Company: NanoSmart Shares: 20
 Share Price: $12.500 Total Worth: $250.00
Company: NanoSmart Shares: 35
 Share Price: $18.125 Total Worth: $634.375
You can't sell more than you have! Transaction failed!
Company: NanoSmart Shares: 35
 Share Price: $18.125 Total Worth: $634.375
Company: NanoSmart Shares: 300035
 Share Price: $40.125 Total Worth: $12037.875
Company: NanoSmart Shares: 35
 Share Price: $0.125 Total Worth: $4.375
```

## Reviewing Our Story to Date

The first step in specifying a class design is to define the class. The class definition is modeled after a structure declaration and lists the members. The declaration has a private section for members that can be accessed only through the member functions. The members declared there can be accessed only through the member functions.

```
char * Bozo::Retort()
```

In other words, `Retort()` is not just a type that belongs to the `Bozo` class. The full, qualified name, `Bozo::Retort()`. The name `Retort()`, on the other hand, is a qualified name, and it can be used only in certain class methods.

Another way of describing this situation is that the scope-resolution operator is needed to refer to a class declaration and a class method.

To create an object, which is a particular class, you use the name it were a type name:

```
Bozo bozetta;
```

This works because a class *is* a user-defined type.

You invoke a class member function, or method, by using the dot membership operator:

```
cout << Bozetta.Retort();
```

This invokes the `Retort()` member function. If you refer to a particular data member, the function name, `bozetta`, refers to the `bozetta` object.

ceed in initializing an object. (You could initialize the data members public instead of private, but that's not one of the main justifications for using classes: data hiding.)

In general, it's best that all objects be initialized. To illustrate, consider the following code:

```
Stock gift;
gift.buy(10, 24.75);
```

With the current implementation of the `Stock` class, this code creates the company member. The class design assumes that there will be any other member functions, but there is no way to ensure that. Around this difficulty is to have objects initialized. To accomplish this, C++ provides for special member functions, called constructors, especially for constructing new objects and assigning values to them. Precisely, C++ provides a name for these member functions: `constructor`. If you provide the method definition. The name of the constructor for the `Stock` class is a member function. The constructor prototype and header have an interesting property: no return value, it's not declared type `void`. In

```

 if (n < 0)
 {
 std::cerr << "Number of shares ca
 << company << " shares
 shares = 0;
 }
 else
 shares = n;
 share_val = pr;
 set_tot();
 }

```

This is the same code that the `acquire()` reference is that in this case, a program automatically declares an object.

## Member Names and Parameter Names

Often those new to constructors try to use the same names in the constructor, as in this example:

```

// NO!
Stock::Stock(const string & company, 1
{

```

```
long shares_;
...
```

With either convention, you then can use `company` and `shares` in the public interface.

## Using Constructors

C++ provides two ways to initialize an object: one way is to use the constructor explicitly:

```
Stock food = Stock("World Cabbage", 250, 1.5);
```

This sets the `company` member of the `food` object to "World Cabbage", the `shares` member to 250, and so on.

The second way is to call the constructor implicitly:

```
Stock garment("Furry Mason", 50, 2.5);
```

This more compact form is equivalent to the first form:

```
Stock garment = Stock("Furry Mason", 50, 2.5);
```

C++ uses a class constructor whenever you use `new` for dynamic memory allocation.

```
Stock *pstock = new Stock("Electroshock Games");
```

structor would look like this:

```
Stock::Stock() { }
```

The net result is that the `fluffy_the_cat` is initialized, just as the following creates `x` without p

```
int x;
```

The fact that the default constructor has not appear in the declaration.

A curious fact about default constructors is that they don't define any constructors. After you define a class, it is your responsibility for providing a default constructor for that class. If you don't provide a nondefault constructor, such as `Stock(int, const string & pr)`, and don't provide your own version of a default constructor, this becomes an error:

```
Stock stock1; // not possible with current compiler
```

The reason for this behavior is that you must initialize all nonstatic data members of uninitialized objects. If, however, you wish to use a class without defining it, you must define your own default constructor for that class. If you don't define a default constructor, the compiler will generate one for you. You must define a default constructor to use the class. You can define a default constructor to take all the arguments to the existing constructor:

```
Stock(const string & co = "Error", int n
```

After you've used either method (no argument or argument), to create the default constructor, you can declare it explicitly:

```
Stock first; // calls default constructor
Stock first = Stock(); // calls it explicitly
Stock *prelief = new Stock; // calls it implicitly
```

However, you shouldn't be misled by the implicit declaration.

```
Stock first("Concrete Conglomerate");
Stock second();
Stock third;
```

The first declaration here calls the nondefault constructor with one argument. The second declaration states that there is a default constructor for the object. When you implicitly call the default constructor,

## Destructors

When you use a constructor to create an object, you are responsible for keeping track of that object until it expires. At that time, you call a special member function bearing the formidable name destructor. It picks up any debris, so it actually serves a useful purpose. When you are new to allocate memory, the destructor should be declared. The constructor doesn't do anything fancy like using

When should a destructor be called? The code shouldn't explicitly call a destructor. (See `new` in Chapter 12, "Classes and Dynamic Memory".) If you create a static storage class object, its destructor is called when the program terminates. If you create an automatic storage class object while doing something, its destructor is called automatically when the scope in which the object is defined terminates. If the object is created in the free store, or the heap, its destructor is called when you free the memory. Finally, a program can create objects and perform operations; in that case, the program automatically calls the destructor when it has finished using it.

Because a destructor is called automatically, you don't need to write one. If you don't provide one, the compiler will generate one for you. If it detects code that leads to the destruction of an object, it will call the destructor. If you don't provide a definition for the destructor, the linker will generate an error.

## Improving the Stock Class

At this point we need to incorporate the concrete member function definitions. Given the significance of the changes, we move them from `stock00.h` to `stock10.h`. The class member functions are defined in `stock10.cpp`. Finally, we place the program using these resources in `stock10main.cpp`.



```
// two constructors
 Stock(); // default constructor
 Stock(const std::string & co, long n = 100);
 ~Stock(); // noisy destructor
 void buy(long num, double price);
 void sell(long num, double price);
 void update(double price);
 void show();
};

#endif
```

---

## The Implementation File

Listing 10.5 provides the method definitions for the `stock10.h` file in order to provide the class definition in double quotation marks instead of in single quotation marks. This file is located at the same location where your source files are. It includes the `iostream` header file to provide I/O support. It also includes the `string` header file and qualified names (such as `std::string`) to use the `string` header files. This file adds the constructor and destructor methods. To help you see when these methods

```

 if (n < 0)
 {
 std::cout << "Number of shares ca
 << company << " shares
 shares = 0;
 }
 else
 shares = n;
 share_val = pr;
 set_tot();
 }
 // class destructor
 Stock::~Stock() // verbose class d
 {
 std::cout << "Bye, " << company << "!"
 }

 // other methods
 void Stock::buy(long num, double price)
 {
 if (num < 0)
 {
 std::cout << "Number of shares pu

```

```

 else
 {
 shares -= num;
 share_val = price;
 set_tot();
 }
 }

void Stock::update(double price)
{
 share_val = price;
 set_tot();
}

void Stock::show()
{
 using std::cout;
 using std::ios_base;
 // set format to #.###
 ios_base::fmtflags orig =
 cout.setf(ios_base::fixed, ios_base::floatfield);
 std::streamsize prec = cout.precision();

 cout << "Company: " << company

```

```
// compile with stock10.cpp
#include <iostream>
#include "stock10.h"

int main()
{
 {
 using std::cout;
 cout << "Using constructors to create\n";
 Stock stock1("NanoSmart", 12, 20.0);
 stock1.show();
 Stock stock2 = Stock ("Boffo Objects");
 stock2.show();

 cout << "Assigning stock1 to stock2:\n";
 stock2 = stock1;
 cout << "Listing stock1 and stock2:\n";
 stock1.show();
 stock2.show();

 cout << "Using a constructor to reset\n";
 stock1 = Stock("Nifty Foods", 10, 50.0);
 cout << "Revised stock1:\n";
 }
}
```

```
Constructor using Nifty Foods called
Bye, Nifty Foods!
Revised stock1:
Company: Nifty Foods Shares: 10
 Share Price: $50.00 Total Worth: $500.00
Done
Bye, NanoSmart!
Bye, Nifty Foods!
```

Some compilers may produce a program with one additional line:

```
Using constructors to create new objects
Constructor using NanoSmart called
Company: NanoSmart Shares: 12
 Share Price: $20.00 Total Worth: $240.00
Constructor using Boffo Objects called
Bye, Boffo Objects!
Company: Boffo Objects Shares: 2
 Share Price: $2.00 Total Worth: $4.00
...
```

The following “Program Notes” section explains this output.

The C++ Standard gives a compiler a couple of options for how to handle this. The first option is to make it behave exactly like the first syntax. The second option is to allow the call to the constructor using Boffo Objects called Company: Boffo Objects Shares: 2

The second way is to allow the call to the destructor to be delayed until the end of the scope. In this option, the destructor is called for the temporary object at the end of the scope. The output of the program is then as follows:  
Constructor using Boffo Objects called  
Bye, Boffo Objects!  
Company: Boffo Objects Shares: 2

The compiler that produced this output delayed the destructor call until the end of the scope, but it's possible that a compiler might wait until the end of the program to call the destructor. In this case, the output would be displayed later.

The following statement illustrates that you can assign objects of the same type:

```
stock2 = stock1; // object assignment
```

As with structure assignment, class object assignment copies one object to the other. In this case, the original object is not destroyed.

Finally, at the end, the program displays this:

```
Done
Bye, NanoSmart!
Bye, Nifty Foods!
```

When the `main()` function terminates, its memory is freed from your plane of existence. Because such an object created is the first deleted, and the first "NanoSmart" was originally in `stock1` but was reset to "Nifty Foods".)

The output points out that there is a fundamental difference between the two statements:

```
Stock stock2 = Stock ("Boffo Objects", 2, 50.0);
stock1 = Stock("Nifty Foods", 10, 50.0);
```

The first of these statements invokes initialization, which sets the object's value, and it may or may not create a temporary object. The second invokes assignment. Using a constructor in an assignment statement causes the creation of a temporary object before the assignment.

## Tip

If you can set object values either through initialization or assignment, use initialization. It is usually more efficient.

```
const Stock land = Stock("Kludgenhorn Prop
land.show();
```

With current C++, the compiler should code for `show()` fails to guarantee that it won't be altered. We can't declare a function's argument to be a `const` because of a syntax problem: The `show()` method doesn't take an argument. Instead, the object it uses is provided implicitly. C++ has a new syntax, one that says a function promise not to alter the object. A C++ solution is to place the `const` keyword before the function name. The `show()` declaration should look like this:

```
void show() const; // promises not to alter
```

Similarly, the beginning of the function definition is

```
void Stock::show() const // promises not to alter
```

Class functions declared and defined this way promise not to alter the object as you should use `const` references and pointers. When it's appropriate, you should make class methods `const`. The object invoking the method is the invoking object. We'll follow this rule from here on.



```
Bozo *pc = new Bozo{"Popo", "Le Peu"};
```

If a constructor has just one argument, that object to a value that has the same type as the pose you have this constructor prototype:

```
Bozo(int age);
```

Then you can use any of the following for

```
Bozo dribble = bozo(44); // primary form
Bozo roon(66); // secondary fo
Bozo tubby = 32; // special form
```

Actually, the third example is a new point, time to tell you about it. Chapter 11 mention lead to unpleasant surprises.

## Caution

A constructor that you can use with a single a to initialize an object to a value:

```
Classname object = value;
```

This feature can cause problems, but it can b

You can do still more with the `Stock` class. So far, we've seen how to use it with but a single object: the object that invoked `main()`. But we may need to deal with two objects, and doing so may require some changes to `this`. Let's look at how the need for this can arise.

Although the `Stock` class declaration displays the `total_val` member, for example, by looking at the `show()` output, you can see that it's not the greatest value, but the program can't tell because it doesn't know the most direct way of letting a program know about the return values. Typically, you use inline code for this:

```
class Stock
{
private:
 ...
 double total_val;
 ...
public:
 double total() const { return total_val; }
 ...
};
```

This definition, in effect, makes `total_val` a public member of the class. If program access is concerned. That is, you can use `total_val` directly, but the class doesn't provide a method to access it.

reference to one of those two objects. The compiler won't modify the explicitly accessed object, and states that the function won't modify the implicit object. The function returns a reference to one of the two constant objects, and the compiler returns that reference.

Suppose, then, that you want to compare the two objects and assign the one with the greater total value to `stock1`. The following statements do so:

```
stock1 = stock1.topval(stock2);
stock2 = stock2.topval(stock1);
```

The first form accesses `stock1` implicitly and `stock2` explicitly. The second form accesses `stock1` explicitly and `stock2` implicitly. The function compares the two objects and returns a reference to the one with the greater total value.

Actually, this notation is a bit confusing. It uses the relational operator `>` to compare the two objects, but the `topval` function is a member function, not a relational operator. Chapter 11 discusses operator overloading, which Chapter 11 discusses.

Meanwhile, there's still the implementation problem. Here's a partial implementation that

```
const Stock & Stock::topval(const Stock & s)
{
 if (s.total_val > total_val)
```

Here `s.total_val` is the total value for the object `s`. If `total_val` is the total value for the object to which the function returns, and if `total_val` is greater than `total_val`, the function returns a reference to the object used to evoke the method (the object which the `topval` message is sent.) Here's the code. If you make the call `stock1.topval(stock2)`, `stock2` is an alias for `stock2`, but there is no alias for `stock1`.

The C++ solution to this problem is to use a pointer. (The pointer points to the object used to invoke the method, as a hidden argument to the method.) Thus, the function sets `this` to the address of the `stock1` object. The `topval()` method. Similarly, the function calls `stock2` the address of the `stock2` object. In general, all calls to the method use the address of the object that invokes the method. (The shorthand notation for `this->total_val`. (Remember that you use the `->` operator to access structure members for class members.) (See Figure 10.4.)

```
so s is joe, this points to kate,
and *this is kate
```

Figure 10.4 this pointer

## Note

Each member function, including constructors, has a special property of the `this` pointer is that it points to the object that invoked the function. (Because it refers to the invoking object as a whole, it can't be used as a pointer to a member.) If you use the `const` qualifier after the function argument parentheses, `const this`, then you can't use `this` to change the object's state; in that case, you can't use `this` to change the object's state.

What you want to return, however, is not a pointer to the object. You want to return the object itself, and you can do that by applying the dereferencing operator `*` to a pointer to the object (points.) Now you can complete the method `topval` of the `Stock` class, returning the invoking object:

```
const Stock & Stock::topval(const Stock & s) const
{
 if (s.total_val > total_val)
 return s; // argument object
 else
 return *this; // invoking object
}
```

```

 void buy(long num, double price);
 void sell(long num, double price);
 void update(double price);
 void show()const;
 const Stock & topval(const Stock & s);
};

#endif

```

---

Listing 10.8 presents the revised class method. Also now that you've seen how the 10.8 replaces them with silent versions.

#### Listing 10.8    **stock20.cpp**

---

```

// stock20.cpp -- augmented version
#include <iostream>
#include "stock20.h"

// constructors
Stock::Stock() // default constructor
{
 company = "no name";
 shares = 0;
}

```

```

 }

 // other methods
 void Stock::buy(long num, double price)
 {
 if (num < 0)
 {
 std::cout << "Number of shares purchased is negative.\n";
 << "Transaction is aborted.\n";
 }
 else
 {
 shares += num;
 share_val = price;
 set_tot();
 }
 }
}

```

```

void Stock::sell(long num, double price)
{
 using std::cout;
 if (num < 0)
 {

```

```

{
 using std::cout;
 using std::ios_base;
 // set format to #.###
 ios_base::fmtflags orig =
 cout.setf(ios_base::fixed, ios_base::floatfield);
 std::streamsize prec = cout.precision();

 cout << "Company: " << company
 << " Shares: " << shares << "\n";
 cout << " Share Price: $" << share_val;
 // set format to #.##
 cout.precision(2);
 cout << " Total Worth: $" << total_val;

 // restore original format
 cout.setf(orig, ios_base::floatfield);
 cout.precision(prec);
}

const Stock & Stock::topval(const Stock & s)
{
 if (s.total_val > total_val)

```



ment—`mystuff[0]`, `mystuff[1]`, and so on—with the `Stock` methods:

```
mystuff[0].update(); // apply update
mystuff[3].show(); // apply show()
const Stock * tops = mystuff[2].topval(mystuff)
 // compare 3rd and 2nd elements and
 // to point at the one with a higher
```

You can use a constructor to initialize the `Stock` objects. You can use the constructor for each individual element:

```
const int STKS = 4;
Stock stocks[STKS] = {
 Stock("NanoSmart", 12.5, 20),
 Stock("Boffo Objects", 200, 2.0),
 Stock("Monolithic Obelisks", 130, 3.25),
 Stock("Fleep Enterprises", 60, 6.5)
};
```

Here the code uses the standard form for initialization of an array of values enclosed in braces. In this case, a call

```
// usestok2.cpp -- using the Stock class
// compile with stock20.cpp
#include <iostream>
#include "stock20.h"

const int STKS = 4;
int main()
{
 // create an array of initialized objects
 Stock stocks[STKS] = {
 Stock("NanoSmart", 12, 20.0),
 Stock("Boffo Objects", 200, 2.0),
 Stock("Monolithic Obelisks", 130, 3.25),
 Stock("Fleep Enterprises", 60, 6.75)
 };

 std::cout << "Stock holdings:\n";
 int st;
 for (st = 0; st < STKS; st++)
 stocks[st].show();
 // set pointer to first element
 const Stock * top = &stocks[0];
 for (st = 1; st < STKS; st++)
 top = &top->topval(stocks[st]);
}
```

class. When that's done, writing the program is under the skin. For example, the original Unix `cfront` that converted C++ programs to C programs had to do is convert a C++ method definition

```
void Stock::show() const
{
 cout << "Company: " << company
 << " Shares: " << shares << "\n"
 << " Share Price: $" << share_price
 << " Total Worth: $" << total_worth
}
```

to the following C-style definition:

```
void show(const Stock * this)
{
 cout << "Company: " << this->company
 << " Shares: " << this->shares
 << " Share Price: $" << this->share_price
 << " Total Worth: $" << this->total_worth
}
```

without conflict. For example, the `shares` member of a `JobRide` class. Also class members of a class from the outside world. This is, to invoke a public member function, you

```
Stock sleeper("Exclusive Ore", 100, 0.25)
sleeper.show(); // use object to invoke
show(); // invalid -- can't
```

Similarly, you have to use the scope-resolution functions:

```
void Stock::update(double price)
{
 ...
}
```

In short, within a class declaration or a member unadorned member name (the unqualified member function. A constructor name is recorded the same as the class name. Otherwise, you must use the indirect membership operator (`->`), or the `scope::` operator, on the context, when you use a class member. This illustrates how identifiers with class scope can be

## Class Scope Constants

Sometimes it would be nice to have symbolic constants in a class declaration might use the literal 30 to specify the number of days in the month. If the same value is used for the same for all objects, it would be nice to have a constant. You might think the following would be a solution:

```
class Bakery
{
private:
 const int Months = 12; // declare a constant
 double costs[Months];
 ...
};
```

But this won't work because declaring a constant in a class doesn't create an object. Hence, until you create an object, the value doesn't exist. (Actually, C++11 provides for member static constants to make the preceding array declaration work; C++11 also provides, however, a couple ways to achieve essentially the same thing.)

First, you can declare an enumeration within a class. An enumeration declaration has class scope, so you can use enumeration names for integer constants. That is, you can use



But you can do an explicit type conversion:

```
int Frodo = int(t_shirt::Small); // Frodo
```

Enumerations are represented by some underlying type. The choice was implementation-dependent. Thus, enumerations might be of different sizes on different systems. C++ has unscoped enumerations. By default, the underlying type is `int`. Furthermore, there's a syntax for indicating the underlying type:

```
// underlying type for pizza is short
enum class : short pizza {Small, Medium, Large};
```

The `: short` specifies the underlying type of the enumeration is an integer type. Under C++11, you also can specify the underlying type for an unscoped enumeration, but if you don't, the compiler makes it implementation-dependent.

## Abstract Data Types

The `Stack` class is pretty specific. Often, however, we want more general concepts. For example, using class-like structures, computer scientists describe as *abstract data types* (ADTs). An ADT describes a data type in a general fashion without implementation details. Consider, for example, the stack. It's a

tion. Instead, it should be expressed in general terms, and so on. Listing 10.10 shows one approach implemented. If it hasn't been implemented rather than `bool`, `false`, and `true`.

#### Listing 10.10    **stack.h**

---

```
// stack.h -- class definition for the stack
#ifndef STACK_H_
#define STACK_H_

typedef unsigned long Item;

class Stack
{
private:
 enum {MAX = 10}; // constant specifying size
 Item items[MAX]; // holds stack items
 int top; // index for top of stack
public:
 Stack();
 bool isempty() const;
 bool isfull() const;
 // push() returns false if stack already full
 bool push(const Item & item); // add item to stack
```



method for isolating from the class design the

Next, you need to implement the class me

#### Listing 10.11 **stack.cpp**

---

```
// stack.cpp -- Stack member functions
#include "stack.h"
Stack::Stack() // create an empty stack
{
 top = 0;
}

bool Stack::isempty() const
{
 return top == 0;
}

bool Stack::isfull() const
{
 return top == MAX;
}

bool Stack::push(const Item & item)
{

```

like this are one of the things that make OOP a good idea. You need a separate array to represent the stack and an index to keep track of the top. In that case, it is your responsibility to make sure you have a new stack. Without the protection that private methods provide, you are making some program blunder that alters data.

Let's test this stack. Listing 10.12 models a customer who takes orders from the top of his in-basket, using the stack.

#### Listing 10.12    **stacker.cpp**

---

```
// stacker.cpp -- testing the Stack class
#include <iostream>
#include <cctype> // or ctype.h
#include "stack.h"
int main()
{
 using namespace std;
 Stack st; // create an empty stack
 char ch;
 unsigned long po;
 cout << "Please enter A to add a purchase,\n";
 cout << "P to process a PO, or Q to quit\n";
 while (cin >> ch && toupper(ch) != 'Q')
 {
```

```

 }
 break;
 }
 cout << "Please enter A to add a purchase order, P to process a PO, or Q to quit." << endl;
}
cout << "Bye\n";
return 0;
}

```

---

The little while loop in Listing 10.12 that we added is absolutely necessary at this point, but it will come in handy again in Chapter 14. Here's a sample run:

```

Please enter A to add a purchase order,
P to process a PO, or Q to quit.

```

**A**

```

Enter a PO number to add: 17885

```

```

Please enter A to add a purchase order,
P to process a PO, or Q to quit.

```

**P**

```

PO #17885 popped

```

```

Please enter A to add a purchase order,

```

## Summary

OOP emphasizes how a program represents a programming problem by using the OOP approach. It defines an interface with the program, specifying how to use the class that implements the interface. Typically, private data is hidden, whereas public member functions, also called methods, are exposed. The class combines data and methods in a single unit and accomplishes data hiding.

Usually, you separate a class declaration from its implementation. The class declaration proper goes into a header file, and the implementation goes into a source file. The source code that defines the methods goes into a separate file. This approach separates the design from the implementation. In principle, you need to know only the class declaration to use the class. Of course, you can look at the implementation (but only in compiled form only), but your program should not depend on it. For example, you can know that a particular value is stored in a class communicate only through methods of the class. You can use either part separately without worrying about the other.

the class methods provide an implementation

## Chapter Review

1. What is a class?
2. How does a class accomplish abstraction?
3. What is the relationship between an object and a class?
4. In what way, aside from being functions, are classes different from class data members?
5. Define a class to represent a bank account. The class should have methods for a depositor's name, the account number (ID), and the balance. The operations should allow the following:
  - Creating an object and initializing it
  - Displaying the depositor's name, account number, and balance
  - Depositing an amount of money
  - Withdrawing an amount of money

Just show the class declaration, not the implementation. Exercise 1 provides you with an opportunity to implement the class.

```

public:
 Person() {lname = ""; fname[0] = '\0';}
 Person(const string & ln, const string & fn) {
 // the following methods display lname and fname
 void Show() const; // first form
 void FormalShow() const; // last form
 };

```

(It uses both a string object and a character array, so both of the two forms are used.) Write a program that uses the class providing code for the undefined methods. Your program should also use the three possible constructors (one with no arguments and two arguments) and the two display methods. Here are the constructors and methods:

```

Person one; // first form
Person two("Smythecraft"); // second form
Person three("Dimwiddy", "Sam"); // third form
one.Show();
cout << endl;
one.FormalShow();
// etc. for two and three

```

3. Do Programming Exercise 1 from Chapter 10, and provide an appropriate golf class declaration. Remember to use the

typedef declaration so that Item is typ

## 6. Here's a class declaration:

```
class Move
{
private:
 double x;
 double y;
public:
 Move(double a = 0, double b = 0)
 showmove() const;
 Move add(const Move & m) const;
 // this function adds x of m to x of
 // adds y of m to y of invoking obje
 // move object initialized to new x,
 reset(double a = 0, double b = 0
};
```

Create member function definitions and a pro

- You can add items to the list.
- You can determine whether the
- You can determine whether the
- You can visit each item in the list

As you can see, this list really is simple; it does

Design a `List` class to represent this abstract data type. Write a header file with the class declaration and the implementations. You should also create

The main reason for keeping the list simple is to make it a good programming exercise. You can implement the data type, as a linked list. But the point is, it's a choice. That is, the public interface should be simple and so on. It should be expressed in the





- Class conversion functions

C++ classes are feature-rich, complex, and powerful. In “Classes,” you began a journey toward object-oriented programming and use a simple class. You saw how a class describes data to be used to represent an object and by the operations that can be performed with the object. You saw member functions, the constructor and the destructor. Creating objects made to a class specification. This chapter is an exploration of class properties, concentrating on the general principles. You’ll probably find some concepts simple and some a bit more subtle. To best understand them, study the examples and experiment with them: What happens if you pass instead of a reference argument for this function? What happens out of a destructor? Don’t be afraid to make mistakes. It’s better unraveling an error than by doing something that might assume that a maelstrom of mistakes inevitably leads to a crash. You’ll be rewarded with a fuller understanding of what C++ can do for you.

This chapter starts with operator overloading, such as `=` and `+` with class objects. Then it discusses letting nonmember functions access private class members.

# Operator Overloading

Let's look at a technique for giving object operators an example of C++ polymorphism. In Chapter 11, we saw how C++ enables you to define several functions that have different signatures (argument lists). This is called *functional polymorphism*. Its purpose is to let you perform a basic operation, even though you apply the operation in a way that is how awkward English would be if you had to use the same type of object—for example, lift\_left your left leg. Operator overloading extends the overloading concept to operators. In C++, it applies to C++ operators. Actually, many C++ operators are overloaded. For example, the \* operator, when applied to an address, yields a pointer. But applying \* to two numbers yields a product. The compiler uses the number and type of operands to decide which definition to use.

C++ lets you extend operator overloading to use the + symbol to add two objects. Again, the compiler uses the operands to determine which definition of + to use. This often makes code look more natural. For example, adding two arrays. Usually, this winds up looking like

```
for (int i = 0; i < 20; i++)
 evening[i] = sam[i] + janet[i]; //
```

The compiler, recognizing the operands and the operator, replaces the operator with the corresponding function call:

```
district2 = sid.operator+(sara);
```

The function then uses the `sid` object implicitly and the `sara` object explicitly (because it's passed explicitly). The function then returns. Of course, the nice part is that it's more concise than the clunky function notation.

C++ imposes some restrictions on operator overloading. We'll stand after you've seen how overloading works. We'll review the process and then discuss the limitations.

## Time on Our Hands: Developer's Example

If you worked on the Priggs account for 2 hours and 40 minutes in the afternoon, how long did you work? The example where the concept of addition for time (a mixture of hours and minutes) don't match the standard C++'s "Programming Modules," handles a similar situation. It defines a `Time` class, using a method to handle addition. Let's

```
#endif
```

---

The `Time` class provides methods for adjusting values, and for adding two times. Listing 11.2 shows the `AddMin()` and `Sum()` methods use integer arithmetic to adjust the minutes and hours values when the total minutes exceeds 60 because the only `iostream` feature used is `cout`. It is economical to use `std::cout` rather than use `cout`.

#### Listing 11.2 **mytime0.cpp**

---

```
// mytime0.cpp -- implementing Time methods
#include <iostream>
#include "mytime0.h"

Time::Time()
{
 hours = minutes = 0;
}

Time::Time(int h, int m)
{
 hours = h;
```

```

 Time sum;
 sum.minutes = minutes + t.minutes;
 sum.hours = hours + t.hours + sum.minutes / 60;
 sum.minutes %= 60;
 return sum;
}

void Time::Show() const
{
 std::cout << hours << " hours, " << minutes << " minutes\n";
}

```

---

Consider the code for the `Sum()` function that the return type is not a reference. The result is less efficient. The code would produce the same value, but it's usually faster and more memory efficient.

However, the return value cannot be a reference. It creates a new `Time` object (`sum`) that represents the sum of the two `Time` objects. Returning the object, as this code does, creates a copy of the object. If the return type were `Time &`, the function could return a reference to the `sum` object. But the `sum` object is a local variable and its lifetime ends when the function returns, so the reference would be a reference to a dangling reference.

```
Time planning;
Time coding(2, 40);
Time fixing(5, 55);
Time total;

cout << "planning time = ";
planning.Show();
cout << endl;

cout << "coding time = ";
coding.Show();
cout << endl;

cout << "fixing time = ";
fixing.Show();
cout << endl;

total = coding.Sum(fixing);
cout << "coding.Sum(fixing) = ";
total.Show();
cout << endl;

return 0;
}
```

---

```
{
private:
 int hours;
 int minutes;
public:
 Time();
 Time(int h, int m = 0);
 void AddMin(int m);
 void AddHr(int h);
 void Reset(int h = 0, int m = 0);
 Time operator+(const Time & t) const;
 void Show() const;
};
#endif
```

---

### Listing 11.5 **mytime1.cpp**

---

```
// mytime1.cpp -- implementing Time meth
#include <iostream>
#include "mytime1.h"

Time::Time()
{
```



```

 hours = h;
 minutes = m;
 }

Time Time::operator+(const Time & t) const
{
 Time sum;
 sum.minutes = minutes + t.minutes;
 sum.hours = hours + t.hours + sum.minutes / 60;
 sum.minutes %= 60;
 return sum;
}

void Time::Show() const
{
 std::cout << hours << " hours, " << minutes << " minutes\n";
}

```

---

Like `Sum()`, `operator+()` is invoked by a `Time` object as the left-hand argument, and returns a `Time` object. Thus, you can add two `Time` objects using the same syntax that `Sum()` uses:

```
total = coding.operator+(fixing); // full name
```

```
cout << "planning time = ";
planning.Show();
cout << endl;
```

```
cout << "coding time = ";
coding.Show();
cout << endl;
```

```
cout << "fixing time = ";
fixing.Show();
cout << endl;
```

```
total = coding + fixing;
// operator notation
cout << "coding + fixing = ";
total.Show();
cout << endl;
```

```
Time morefixing(3, 28);
cout << "more fixing time = ";
morefixing.Show();
cout << endl;
total = morefixing.operator+(total);
```

```
C = A + B; // use addition as defined
```

Can you add more than two objects? For `e` objects, can you do the following?

```
t4 = t1 + t2 + t3; //
```

The way to answer this is to consider how `+` calls. Because addition is a left-to-right operator,

```
t4 = t1.operator+(t2 + t3); //
```

Then the function argument is itself translated:

```
t4 = t1.operator+(t2.operator+(t3)); //
```

Is this valid? Yes, it is. The function call `t2.operator+(t3)` represents the sum of `t2` and `t3`. This object then becomes the argument to the `t1.operator+( )` function call, and that call returns the object that represents the sum of `t2` and `t3`. In short, `t4` is the sum of `t1` and `t3`, just as desired.

addition.

- You can't create new operator symbols
- `operator**()` function to denote exponentiation
- You cannot overload the following operators

| Operator                      | Description                           |
|-------------------------------|---------------------------------------|
| <code>sizeof</code>           | The size of a type                    |
| <code>.</code>                | The member access operator            |
| <code>.*</code>               | The pointer-to-member access operator |
| <code>::</code>               | The scope resolution operator         |
| <code>?:</code>               | The conditional operator              |
| <code>typeid</code>           | An RTTI operator                      |
| <code>const_cast</code>       | A type cast operator                  |
| <code>dynamic_cast</code>     | A type cast operator                  |
| <code>reinterpret_cast</code> | A type cast operator                  |
| <code>static_cast</code>      | A type cast operator                  |

This still leaves all the operators

|     |    |     |        |
|-----|----|-----|--------|
| >   | += | -=  | *=     |
| ^=  | &= | =   | <<     |
| <<= | == | !=  | <=     |
|     | ++ | --  | ,      |
| ()  | [] | new | delete |

---

In addition to these formal restrictions, you should avoid overloading operators. For example, you shouldn't overload the `+` operator to add the data members of two `Time` objects. Nothing is wrong with the `+` operator did, so it would be better to define a function such as `Swap()`.

## More Overloaded Operators

Some other operations make sense for the `Time` class: you can subtract one time from another or multiply a time by a scalar. To overload the subtraction and multiplication operators. To overload the subtraction operator: you create `operator-()` and `operator*` and adding prototypes to the class declaration:

```
Time operator-(const Time & t) const;
Time operator*(double n) const;
```

```
};
#endif
```

---

Then you add definitions for the new methods.  
Listing 11.8.

#### Listing 11.8    **mytime2.cpp**

---

```
// mytime2.cpp -- implementing Time methods
#include <iostream>
#include "mytime2.h"

Time::Time()
{
 hours = minutes = 0;
}

Time::Time(int h, int m)
{
 hours = h;
 minutes = m;
}

void Time::AddMin(int m)
{
```

```
Time Time::operator-(const Time & t) const
{
 Time diff;
 int tot1, tot2;
 tot1 = t.minutes + 60 * t.hours;
 tot2 = minutes + 60 * hours;
 diff.minutes = (tot2 - tot1) % 60;
 diff.hours = (tot2 - tot1) / 60;
 return diff;
}
```

```
Time Time::operator*(double mult) const
{
 Time result;
 long totalminutes = hours * mult * 60;
 result.hours = totalminutes / 60;
 result.minutes = totalminutes % 60;
 return result;
}
```

```
void Time::Show() const
{
 std::cout << hours << " hours, " << mi
}

```

---

```
cout << "waxing time = ";
waxing.Show();
cout << endl;
```

```
cout << "total work time = ";
total = weeding + waxing; // use op
total.Show();
cout << endl;
```

```
diff = weeding - waxing; // use op
cout << "weeding time - waxing time =
diff.Show();
cout << endl;
```

```
adjusted = total * 1.5; // use o
cout << "adjusted work time = ";
adjusted.Show();
cout << endl;
```

```
return 0;
```

```
}
```

---



Before seeing how to make friends, let's look at overloading a binary operator (that is, an operator that operates on two operands). Multiplying a `Time` object by a scalar value is a natural operation, so let's study that case.

In the previous `Time` class example, the overloaded `+` operator was defined in terms of the `operator+` member function. From the other two overloaded operators in the example, the `+` operator and the `-` operator, the addition and subtraction operators each combine two `Time` values. The multiplication operator combines a `Time` value with a double value, so it can be used. Remember, the left operand is the object of the class.

```
A = B * 2.75;
```

This statement translates to the following member function call:

```
A = B.operator*(2.75);
```

But what about the following statement?

```
A = 2.75 * B; // cannot correspond to a member function
```

Conceptually, `2.75 * B` should be the same as `B * 2.75`. However, it does not correspond to a member function because the left operand is the invoking object, but `2.75` is not a `Time` object. To replace the expression with a member function call, we need to define a friend function.

can't already access private data in a class. We lack access. But there is a special category of access private members of a class.

## Creating Friends

The first step toward creating a friend function and prefix the declaration with the keyword

```
friend Time operator*(double m, const Time
```

This prototype has two implications:

- Although the `operator*()` function is a member function. So it isn't invoked by
- Although the `operator*()` function is a member function. So it has the same access rights as a member function.

The second step is to write the function definition. When writing the function definition, you don't use the `Time::` qualifier. Also, you don't use the `friend` keyword in the definition. The definition should look like this:

```
Time operator*(double m, const Time & t)
{
 Time result;
 long totalminutes = t.hours * mult * m;
 result.hours = totalminutes / 60;
```

keep in mind that only a class declaration can be a friend. The class declaration still controls which functions are public and friends are simply two different mechanisms.

Actually, you can write this particular friend definition so that it switches which value comes first.

```
Time operator*(double m, const Time & t)
{
 return t * m; // use t.operator*(m)
}
```

The original version accessed `t.minutes` and was a friend. This version only uses the `Time` object and doesn't handle the private values, so this version doesn't need to be a friend. There are two reasons to make this version a friend, too. First, it's part of the official class interface. Second, if you want to access private data directly, you only have to change the prototype.

## Tip

If you want to overload an operator for a class, you can use a class term as the first operand, you can use a class term as the first operand, you can use a class term as the first operand.

## The First Version of Overloading <<

To teach the Time class to use cout, you can implement like the following uses two objects, with

```
cout << trip;
```

If you use a Time member function to overload << as it did when you overloaded the \* operator, you would have to use the << operator this way:

```
trip << cout; // if operator<<() were a member
```

This would be confusing. But by using a free function for this way:

```
void operator<<(ostream & os, const Time & t)
{
 os << t.hours << " hours, " << t.minutes << " minutes";
}
```

This lets you use

```
cout << trip;
```

to print data in the following format:

```
4 hours, 23 minutes
```

output to a file. Through the magic of inheritance, `ofstream` objects can use `ostream` methods. This allows `ofstream` to write `Time` data to files as well as to `cout`. The `ofstream` object instead of `cout` as the first

The call `cout << trip` should use the `cout` object. It passes the object as a reference instead of by value. This causes `os` to be an alias for `cout`, and the expression for `cerr`. The `Time` object can be passed by value. This makes the object values available to the function. It takes more memory and time than passing by value.

## The Second Version of Overloading <<

The implementation just presented has a problem. The code `cout << trip;`

But the implementation doesn't allow you to use the `cout` object normally uses:

```
cout << "Trip time: " << trip << " (Tuesday"
```

To understand why this doesn't work and why you need to know a bit more about how `cout` operates.

```
}
```

Note that the return type is `ostream &`. `Return` returns a reference to an `ostream` object. Because the function to begin with, the net effect is that the `cout` object passed to it. That is, the statement

```
cout << trip;
```

becomes the following function call:

```
operator<<(cout, trip);
```

And that call returns the `cout` object. So now the original statement becomes

```
cout << "Trip time: " << trip << " (Tuesday)\n";
```

Let's break this into separate steps to see how the particular `ostream` definition of `<<` that displays a `string` works.

```
cout << "Trip time: "
```

So the expression `cout << "Trip time: "` evaluates to its return value, `cout`. This reduces the original statement to

```
cout << trip << " (Tuesday)\n";
```

friend function with a definition in this form:

```
ostream & operator<<(ostream & os, const
{
 os << ... ; // display object contents
 return os;
}
```

Listing 11.10 shows the class definition as `mytime3.h`. It implements `operator*()` and `operator<<()`. It implements `operator<<()` because the code is so short. (When the definition is short, you can use the `friend` prefix.)

## Caution

You use the `friend` keyword only in the prototype. You do not use it in the function definition unless the definition is in the header file.

### Listing 11.10 `mytime3.h`

```
// mytime3.h -- Time class with friends
#ifndef MYTIME3_H_
#define MYTIME3_H_
#include <iostream>

class Time
{
```

mytime3.h in mytime3.cpp provides support

### Listing 11.11 **mytime3.cpp**

---

```
// mytime3.cpp -- implementing Time meth
#include "mytime3.h"
```

```
Time::Time()
{
 hours = minutes = 0;
}
```

```
Time::Time(int h, int m)
{
 hours = h;
 minutes = m;
}
```

```
void Time::AddMin(int m)
{
 minutes += m;
 hours += minutes / 60;
 minutes %= 60;
}
```



```
 int tot1, tot2;
 tot1 = t.minutes + 60 * t.hours;
 tot2 = minutes + 60 * hours;
 diff.minutes = (tot2 - tot1) % 60;
 diff.hours = (tot2 - tot1) / 60;
 return diff;
 }
```

```
Time Time::operator*(double mult) const
{
 Time result;
 long totalminutes = hours * mult * 60;
 result.hours = totalminutes / 60;
 result.minutes = totalminutes % 60;
 return result;
}
```

```
std::ostream & operator<<(std::ostream & o
{
 o << t.hours << " hours, " << t.minutes;
 return o;
}
```

---

```
 temp = aida + toska; // operator+(
 cout << "Aida + Tosca: " << temp << endl;
 temp = aida* 1.17; // member operator
 cout << "Aida * 1.17: " << temp << endl;
 cout << "10.0 * Tosca: " << 10.0 * toska << endl;

 return 0;
}
```

---

Here is the output of the program in Listing 13.1:

Aida and Tosca:

3 hours, 35 minutes; 2 hours, 48 minutes

Aida + Tosca: 6 hours, 23 minutes

Aida \* 1.17: 4 hours, 11 minutes

10.0 \* Tosca: 28 hours, 0 minutes

## Overloaded Operators: Member and Nonmember Functions

For many operators, you have a choice between member functions and nonmember functions to implement operator overloading.

```
T1 = T2.operator+(T3); // member function
T1 = operator+(T2, T3); // nonmember function
```

Keep in mind that you must choose one or the other form of the `operator+` operator, but not both. Because both forms may be used, this forms is an ambiguity error, leading to a compilation error.

Which form, then, is it best to use? For some applications, the member function is the only valid choice. Otherwise, the nonmember function. Sometimes, depending on the class design, the nonmember function is the better choice, particularly if you have defined type conversions and Friends,” near the end of this chapter.

## More Overloading: A Vector

Let’s look at another class design that uses operator overloading for representing vectors. This class also illustrates further applications of operator overloading. In describing two different ways of describing the same thing, we will see how, for vectors, you can use many of the new techniques. A vector, as the term is used in engineering and physics, is a quantity that has a magnitude (size) and a direction. For example, if you know the magnitude and direction of a push, how hard you push (the magnitude) and in what direction can save a tottering vase, whereas a push without direction can lead to doom. To fully describe the motion of your

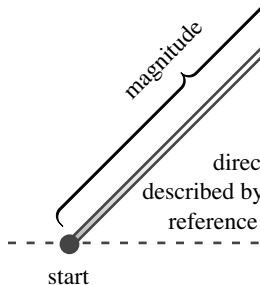


Figure 11.1  
displacement

Now say you're Lhanappa, the great mammoth hunter, who has traveled 10 kilometers to the northwest. But because of the wind, you decide to travel 4.1 kilometers from the southeast. So you go 10 kilometers northwest and then 4.1 kilometers southeast, ending up finding the herd from the south. You know these two locations are the same location as the single 14.1-kilometer vector. The mammoth hunter, also knows how to add vectors.

Vectors are a natural choice for operator overloading. A vector with a single number, so it makes sense to think of vectors as having analogs to ordinary arithmetic operations. This parallel suggests overloading the corresponding operators for vectors.

To keep things simple, in this section we'll use a vector to represent a screen displacement, instead of a three-dimensional movement of a helicopter or a gymnast. You need a two-dimensional vector, but you have a choice of how to represent it.

- You can describe a vector by its magnitude and direction.
- You can represent a vector by its x and y components.

The components are a horizontal vector (the x component) and a vertical vector (the y component), which add up to the final vector. For example, moving a point 30 units to the right and 40 units up puts the point at the same spot as moving 50 units diagonally. Therefore, a vector with a magnitude of 50 and a direction of 53.1° is the same as a vector having a horizontal component of 30 and a vertical component of 40. In physics, a displacement vector is where you start and where you end up.

presents a class declaration. To refresh your memory, we repeat the class declaration inside the `VECTOR` namespace, with a couple constants (`RECT` and `POL`) for identifying the technique in Chapter 10, so we may as well use it.

### Listing 11.13    **vect.h**

---

```
// vect.h -- Vector class with <<, mode s
#ifndef VECTOR_H_
#define VECTOR_H_
#include <iostream>
namespace VECTOR
{
 class Vector
 {
 public:
 enum Mode {RECT, POL};
 // RECT for rectangular, POL for Polar
 private:
 double x; // horizontal
 double y; // vertical va
 double mag; // length of v
 double ang; // direction c
 Mode mode; // RECT or POL
```

```

 friend std::ostream &
 operator<<(std::ostream & o,
 const vector<point> & v) {
 };

 } // end namespace VECTOR
#endif

```

---

Notice that the four functions in Listing 11.14 are all defined in the class declaration. This automation is possible because that these functions are so short makes them convenient to place in the header file. Since all of these functions should alter object data, so they are declared `static`. As you recall from Chapter 10, this is the syntax for declaring a static member function of a class. The `vector` object it implicitly accesses.

Listing 11.14 shows all the methods and functions for the `vector` class. The listing uses the open nature of namespaces to place the `vector` namespace. Note how the constructor `vector()` set both the rectangular and the polar representations. The `atan2` function is available immediately without further comment. As mentioned in Chapter 4, “Compound Types,” the `atan2` function uses angles in radians, so the functions `atan2` and `atan` methods. The `vector` class implementation handles the conversion from polar coordinates to rectangular coordinates or con-

```
void Vector::set_mag()
{
 mag = sqrt(x * x + y * y);
}
```

```
void Vector::set_ang()
{
 if (x == 0.0 && y == 0.0)
 ang = 0.0;
 else
 ang = atan2(y, x);
}
```

```
// set x from polar coordinate
void Vector::set_x()
{
 x = mag * cos(ang);
}
```

```
// set y from polar coordinate
void Vector::set_y()
{
 y = mag * sin(ang);
}
```



```

 set_x();
 set_y();
 }
 else
 {
 cout << "Incorrect 3rd argument\n";
 cout << "vector set to 0\n";
 x = y = mag = ang = 0.0;
 mode = RECT;
 }
}

```

```

// reset vector from rectangular coordinates
// RECT (the default) or else from polar coordinates
// form is POL

```

```

void Vector::reset(double n1, double n2, int form)
{
 mode = form;
 if (form == RECT)
 {
 x = n1;
 y = n2;
 set_mag();
 }
}

```

```
 {
 mode = POL;
 }
```

```
void Vector::rect_mode() // set t
{
 mode = RECT;
}
```

```
// operator overloading
// add two Vectors
Vector Vector::operator+(const Vector
{
 return Vector(x + b.x, y + b.y);
}
```

```
// subtract Vector b from a
Vector Vector::operator-(const Vector
{
 return Vector(x - b.x, y - b.y);
}
```

```
// reverse sign of Vector
```

```

 os << "(x,y) = (" << v.x << "
 else if (v.mode == Vector::POL)
 {
 os << "(m,a) = (" << v.mag << "
 << v.ang * Rad_to_deg << "
 }
 else
 os << "Vector object mode is
 return os;
}

} // end namespace VECTOR

```

---

You could design the class differently. For example, polar coordinates and not the polar coordinates could be moved to the `magval()` method. Since which conversions are seldom used, this could be a good idea. The `reset()` method isn't really needed. Suppose the following code:

```
shove.reset(100,300);
```

means, look at the code for the constructor:

```
Vector::Vector(double n1, double n2, Mode
{
 mode = form;
 if (form == RECT)
 {
 x = n1;
 y = n2;
 set_mag();
 set_ang();
 }
 else if (form == POL)
 {
 mag = n1;
 ang = n2 / Rad_to_deg;
 set_x();
 set_y();
 }
 else
 {
 cout << "Incorrect 3rd argument
 cout << "vector set to 0\n";
```

enum type:

```
Vector rector(20.0, 30.0, 2); // type mis
```

Still, the resourceful and curious user could wonder what happens:

```
Vector rector(20.0, 30.0, VECTOR::Vector::
```

In this case, he gets admonished.

Next, the operator<<() function uses the

```
// display rectangular coordinates if mode == RECT
// else display polar coordinates if mode == POL
std::ostream & operator<<(std::ostream & os, const Vector & v)
{
 if (v.mode == Vector::RECT)
 os << "(x,y) = (" << v.x << ", " << v.y << ")";
 else if (v.mode == Vector::POL)
 {
 os << "(m,a) = (" << v.mag << ", " << v.ang * Rad_to_deg << ")";
 }
 else
 os << "Vector object mode is invalid";
 return os;
}
```

## Overloading Arithmetic Operators

Adding two vectors is very simple when you have the x and y components to get the x component of the answer. Similarly, you can get the y component of the answer. From this code:

```
Vector Vector::operator+(const Vector & b)
{
 Vector sum;
 sum.x = x + b.x;
 sum.y = y + b.y;
 return sum; // incomplete vector
}
```

And this would be fine if the object stored the angle. Unfortunately, this version of the code fails to set the angle. Adding more code:

```
Vector Vector::operator+(const Vector & b)
{
 Vector sum;
 sum.x = x + b.x;
 sum.y = y + b.y;
 sum.set_ang(sum.x, sum.y);
}
```

In visual terms, multiplying a vector by a number scales it by that factor. So multiplying a vector by 3 produces a vector still pointed in the same direction. It's easy to think of a vector class represents a vector. In polar terms, you multiply the magnitude alone. In rectangular terms, you multiply a vector's x and y components separately by the number. That is, multiplying by 3 makes the components 15 and 9. Here's the multiplication operator does:

```
Vector Vector::operator*(double n) const
{
 return Vector(n * x, n * y);
}
```

As with overloaded addition, the code lets you create a new Vector object from the new x and y components. The `n` is a `double` value. Just as in the `Time` example, you can also overload `double times Vector`:

```
Vector operator*(double n, const Vector &a)
{
 return a * n; // convert double time to Vector
}
```

implicit vector argument, so you should use `x`.

Next, consider the unary minus operator, `-`. Applying the `-` operator to a regular number, as in `-x`, changes the sign of `x`. Applying the `-` operator to a vector reverses the sign of each component. The `-` operator should return a new vector that is the reverse of `x`. The `-` operator leaves the magnitude unchanged but reverses the direction. People with no mathematical training nonetheless have an intuitive understanding of `-`. It is the prototype and definition for overloaded unary operators.

```
Vector operator-() const;
Vector Vector::operator-() const
{
 return Vector (-x, -y);
}
```

Note that now there are two separate definitions of the `-` operator. The two definitions have different signatures. You can have two definitions of the `-` operator because C++ provides two different versions of the `-` operator to begin with. An operator that has a non-void return type can only be overloaded as a binary operator.



look-up is faster. If an application often needs the implementation used in this example would only infrequently, the other implementation would be a better choice. You can use the first implementation in one program and the second in another, but they must have the same user interface for both.

## **Taking the Vector Class on a Random Walk**

Listing 11.15 provides a short program that uses the famous Drunkard's Walk problem. Actually, not a serious health problem rather than as a source of random walk problem. The idea is that you place a person on a grid and let them walk in a straight line, but the direction of each step varies randomly. One way of phrasing the problem is to let a random walker to travel, say, 50 feet away from the starting point by adding a bunch of randomly oriented vectors.

Listing 11.15 lets you select the target distance and the number of steps the wanderer's step. It maintains a running total of the distance traveled (represented as a vector) and reports the number of steps taken, the distance, along with the walker's location (in both x and y coordinates). The progress is quite inefficient. A journey of 1,000 steps might leave the walker only 50 feet from the starting point. The distance traveled (50 feet, in this case) by the number of steps taken.

```
double dstep;
cout << "Enter target distance (q to
while (cin >> target)
{
 cout << "Enter step length: ";
 if (!(cin >> dstep))
 break;

 while (result.magval() < target)
 {
 direction = rand() % 360;
 step.reset(dstep, direction,
 result = result + step;
 steps++;
 }
 cout << "After " << steps << " st
 "has the following location:\n";
 cout << result << endl;
 result.polar_mode();
 cout << " or\n" << result << endl;
 cout << "Average outward distance
 << result.magval()/steps << e
 steps = 0;
 result.reset(0.0, 0.0);
```

(x,y) = (-21.9577, 45.3019)

or

(m,a) = (50.3429, 115.8593)

Average outward distance per step = 0.0529

Enter target distance (q to quit): **50**

Enter step length: **1**

After 1716 steps, the subject has the foll

(x,y) = (40.0164, 31.1244)

or

(m,a) = (50.6956, 37.8755)

Average outward distance per step = 0.0295

Enter target distance (q to quit): **q**

Bye!

The random nature of the process produces different results even if the initial conditions are the same. On average, it quadruples the number of steps needed to cover a distance. It suggests that, on average, the number of steps ( $N$ ) needed to cover a distance of  $D$  is given by the following equation:

$$N = (D/s)^2$$

This is just an average, but there will be considerable variation. For example, 1,000 trials of attempting to travel 50 steps resulted in 636 steps (close to the theoretical value of 625).

the default seed value and initiate a different sequence of random numbers. `time()` uses the return value of `time(0)` to set the seed, which is the current calendar time, often implemented as the number of seconds since 1970. (More generally, `time()` takes the address of a variable, sets the seed to that variable and also returns it. Using 0 for the address of the variable, otherwise unneeded `time_t` variable.) Thus, the seed changes each time you run the program, making the random numbers more unpredictable.

```
srand(time(0));
```

The `cstdlib` header file (formerly `stdlib.h`) contains the `srand()` and `rand()` functions, whereas `ctime` (formerly `time.h`) contains the `time()` function. `cstdlib` provides more extensive random number support than `stdlib.h` (the latter header file.)

The program uses the `result` vector to keep track of the current state of the cycle of the inner loop, the program sets the `result` vector to the current `result` vector. When the magnitude of the `result` vector exceeds the loop terminates.

By setting the vector mode, the program computes the magnitude of the vector and in polar terms.

Incidentally, the following statement has the effect of setting the vector mode, regardless of the initial modes of `result`.

```
result = result + step;
```

# Automatic Conversions and Classes

The next topic on the class menu is type conversions to and from user-defined types. The compiler handles conversions for its built-in types. When you convert a value of one standard type to a variable of another standard type, the value is converted to the same type as the receiving variable, if compatible. For example, the following statements

```
long count = 8; // int value 8 converted to long
double time = 11; // int value 11 converted to double
int side = 3.33; // double value 3.33 converted to int
```

These assignments work because C++ recognizes that `count`, `time`, and `side` represent the same basic thing—a number—and it knows how to do it for making the conversions. Recall from Chapter 2 that when you can lose some precision in these conversions, the variable `side` results in `side` getting the value 3.

The C++ language does not automatically convert between user-defined types. For example, the following statement fails because the right side is a number:

```
int * p = 10; // type clash
```

```

class Stonewt
{
private:
 enum {Lbs_per_stn = 14}; // pounds per stone
 int stone; // whole stones
 double pds_left; // fractional part
 double pounds; // entire pounds
public:
 Stonewt(double lbs); // constructor from lbs
 Stonewt(int stn, double lbs); // constructor from stn and lbs
 Stonewt(); // default constructor
 ~Stonewt(); // destructor
 void show_lbs() const; // show in pounds
 void show_stn() const; // show in stones
};
#endif

```

---

As mentioned in Chapter 10, enum provides constants, provided that they are integers. Or you can use a static const int:

```
static const int Lbs_per_stn = 14;
```

```
// construct Stonewt object from double lbs
Stonewt::Stonewt(double lbs)
{
 stone = int (lbs) / Lbs_per_stn; //
 pds_left = int (lbs) % Lbs_per_stn + 1;
 pounds = lbs;
}

// construct Stonewt object from stone, do
Stonewt::Stonewt(int stn, double lbs)
{
 stone = stn;
 pds_left = lbs;
 pounds = stn * Lbs_per_stn + lbs;
}

Stonewt::Stonewt() // default constructor
{
 stone = pounds = pds_left = 0;
}

Stonewt::~~Stonewt() // destructor
{
}
```

The program uses the `Stonewt(double)` constructor to create a `Stonewt` object, using `19.6` as the initialization value. The compiler then copies the contents of the temporary object into `stn`. This is a *conversion* because it happens automatically, without an explicit cast.

Only a constructor that can be used with just one argument is allowed. The following constructor has two arguments:

```
Stonewt(int stn, double lbs); // not a constructor
```

However, it would act as a guide to the compiler to use the `int` as the second parameter:

```
Stonewt(int stn, double lbs = 0); // int constructor
```

Having a constructor work as an automatic conversion is a useful feature. As programmers acquired more experience with C++, they found that the automatic aspect isn't always desirable. Sometimes conversions are not what you want. So C++ added a new keyword, `explicit`. That is, you can declare the constructor this way:

```
explicit Stonewt(double lbs); // no implicit conversions
```

This turns off implicit conversions such as `Stonewt stn;` and forces explicit conversions—that is, conversions using a cast:



Let's look at the last point in more detail. The function prototyping lets the Stonewt (double) numerical types. That is, both of the following double and then using the Stonewt (double)

```
Stonewt Jumbo(7000); // uses Stonewt(double)
Jumbo = 7300; // uses Stonewt(double)
```

However, this two-step conversion process is a poor choice. That is, if the class also defined a Stonewt operator, it would reject these statements, probably pointing out that 7000 is too long or a double, so the call is ambiguous.

Listing 11.18 uses the class constructors to handle type conversions. Be sure to compile Listing

#### Listing 11.18 **stone.cpp**

---

```
// stone.cpp -- user-defined conversions
// compile with stonewt.cpp
#include <iostream>
using std::cout;
#include "stonewt.h"
void display(const Stonewt & st, int n);
int main()
{
```

```

void display(const Stonewt & st, int n)
{
 for (int i = 0; i < n; i++)
 {
 cout << "Wow! ";
 st.show_stn();
 }
}

```

---

Here is the output of the program in Listing 12.1:

```

The celebrity weighed 19 stone, 9 pounds
The detective weighed 20 stone, 5.7 pounds
The President weighed 302 pounds
After dinner, the celebrity weighed 19 stone, 9 pounds
After dinner, the President weighed 325 pounds
Wow! 23 stone, 3 pounds
Wow! 23 stone, 3 pounds
The wrestler weighed even more.
Wow! 30 stone, 2 pounds
Wow! 30 stone, 2 pounds
No stone left unearned

```

double and then uses `Stonewt(double)` to se

Finally, note the following function call:

```
display(422, 2); // convert 422 to doub
```

The prototype for `display()` indicates that the first argument is a `Stonewt` object. (Either a `Stonewt` or a `Stonewt &` form is acceptable.) Confronted with an `int` argument, the compiler looks for a constructor to convert the `int` to the desired `Stonewt` object. The compiler finds the `Stonewt(int)` constructor, which converts the `int` to a `double` and then uses `Stonewt(double)` to construct the `Stonewt` object.

## Conversion Functions

Listing 11.18 converts a number to a `Stonewt` object. Listing 11.19 converts a `Stonewt` object to a `double` value.

```
Stonewt wolfe(285.7);
double host = wolfe; // ?? possible ??
```

The answer is that you can do this—but not in the way shown. The C++ standard provides a special form of a C++ operator function called a *conversion function* for converting another type *to* the class type.

- The conversion function must not specify a return type.
- The conversion function must have no parameters.

For example, a function to convert to type `double`:

```
operator double();
```

The *typeName* part (in this case *typeName* is `double`) specifies the type to which to convert, so no return type is needed. The `operator` keyword means it has to be invoked by a particular class. The `double` means a `double` value to convert. Thus, the function doesn't need any arguments.

To add functions that convert `stone_wt` or `stone_lb` to `double` requires adding the following prototypes to the header file:

```
operator int();
operator double();
```

Listing 11.19 shows the modified class declaration.

#### Listing 11.19    **stonewt1.h**

---

```
// stonewt1.h -- revised definition for the stone class
#ifndef STONEWT1_H_
#define STONEWT1_H_
class Stonewt
{
```

0.5 is 114.9, and `int (114.9)` is 114. But if  
and `int (115.1)` is 115.

### Listing 11.20    **stonewt1.cpp**

---

```
// stonewt1.cpp -- Stonewt class methods +
#include <iostream>
using std::cout;
#include "stonewt1.h"

// construct Stonewt object from double value
Stonewt::Stonewt(double lbs)
{
 stone = int (lbs) / Lbs_per_stn; //
 pds_left = int (lbs) % Lbs_per_stn + 1;
 pounds = lbs;
}

// construct Stonewt object from stone, do
Stonewt::Stonewt(int stn, double lbs)
{
 stone = stn;
 pds_left = lbs;
 pounds = stn * Lbs_per_stn +lbs;
}
```

```
{

 return int (pounds + 0.5);

}

Stonewt::operator double() const
{
 return pounds;
}
```

---

Listing 11.21 tests the new conversion function. The program uses an implicit conversion, whereas the previous program used an explicit cast. Be sure to compile Listing 11.20 along with Listing 11.21.

#### Listing 11.21    **stone1.cpp**

---

```
// stone1.cpp -- user-defined conversion
// compile with stonewt1.cpp
#include <iostream>
#include "stonewt1.h"

int main()
{
```

```
double p_wt = poppins;
```

The answer is no. In the `p_wt` example, the `poppins` is converted to type `double`. But in the `cout` example, the conversion should be to `int` or to `double`. Facing this ambiguity, the compiler would complain that you were using an ambiguous conversion. It indicates what type to use.

Interestingly, if the class defined only the `double` conversion, the compiler would accept the statement. That's because with only one conversion, there is no ambiguity.

You can have a similar situation with assignment. If a class defines two conversions, the compiler rejects the following statement as ambiguous:

```
long gone = poppins; // ambiguous
```

In C++, you can assign both `int` and `double` values to a `long` variable. You can legitimately use either conversion function. The compiler has the responsibility of choosing which. But if you eliminate the ambiguity, the compiler accepts the statement. For example, if you define only the `int` conversion, then the compiler will use the `int` conversion function. Then it converts the `int` value to type `long` via the `long` constructor.

When the class defines two or more conversion functions, you must explicitly indicate which conversion function to use. You can do this by casting the value to the

implicit conversions. In C++98, the keyword `explicit` was used to prevent implicit conversions, but C++11 removes that limitation and introduces the `explicit` operator as `explicit`:

```
class Stonewt
{
 ...
 // conversion functions
 explicit operator int() const;
 explicit operator double() const;
};
```

With these declarations in place, you would write:

Another approach is to replace a conversion function with a static member function that does the same task—but only if called explicitly. For example, you could replace `Stonewt::operator int()` with

with

```
int Stonewt::Stone_to_Int() { return int (pou...
```

This disallows the following:

```
int plb = poppins;
```



## Conversions and Friends

Let's bring addition to the `Stonewt` class. As in the previous class, you can use either a member function or a friend function to simplify matters, assume that no conversion functions are defined.) You can implement addition with a member function:

```
Stonewt Stonewt::operator+(const Stonewt & st) const
{
 double pds = pounds + st.pounds;
 Stonewt sum(pds);
 return sum;
}
```

Or you can implement addition as a friend function:

```
Stonewt operator+(const Stonewt & st1, const Stonewt & st2)
{
 double pds = st1.pounds + st2.pounds;
 Stonewt sum(pds);
 return sum;
}
```

Remember, you can provide the method `operator+` as a member or as a friend. Both forms let you do the following:

or else

```
total = operator+(jennySt, bennySt); //
```

In either case, the actual argument types matter. The member function is invoked, as required, by a `std::string`.

Next,

```
total = jennySt + kennyD;
```

becomes

```
total = jennySt.operator+(kennyD); // member
```

or else

```
total = operator+(jennySt, kennyD); // free
```

Again, the member function is invoked, as required. In each case, one argument (`kennyD`) is type `double`. The `std::string` constructor to convert the argument to a `std::string`.

By the way, having an `operator double (const std::string&)` causes confusion at this point because that would convert `kennyD`. Instead of converting `kennyD` to `double` and then performing the addition, it could convert `jennySt` to `double` and perform the addition. This version functions creates ambiguities.

choices. The first, as you just saw, is to define the Stonewt(double) constructor handle conversion Stonewt arguments:

```
operator+(const Stonewt &, const Stonewt &)
```

The second choice is to further overload the operator+ to explicitly use one type double argument:

```
Stonewt operator+(double x); // member function
friend Stonewt operator+(double x, Stonewt y);
```

That way, the following statement exactly matches the member function:

```
total = jennySt + kennyD; // Stonewt + double
```

And the following statement exactly matches the friend function:  
s) friend function:

```
total = pennyD + jennySt; // double + Stonewt
```

Earlier, we did something similar for Vector.

Each choice has advantages. The first choice is a shorter program because you define fewer functions and fewer chances to mess up. The disadvantage is the extra code needed to invoke the conversion constructor.

operator *op* has this form:

```
operatorop(argument-list)
```

*argument-list* represents operands for the member function, then the first operand is the *argument-list*. For example, this chapter overloads the `operator+()` member function for the `vector` class. For vectors, you can use either of the following statements:

```
result = up.operator+(right);
result = up + right;
```

For the second version, the fact that the `operator+` is a member function of `vector` allows C++ to use the `vector` definition of addition.

When an operator function is a member function, the first operand is the invoking object. In the preceding statement, `up` is the invoking object. If you want to define an operator function for a class object, you must use a friend function. The friend function definition in whichever order you want.

One of the most common tasks for operator overloading is to allow an `ostream` object to be the first operand in an output operation.

```
bean = String("pinto"); // converts type
```

To convert from a class to another type, you provide instruction about how to make the conversion as a member function. If it is to convert to type `double`, the prototype:

```
operator typeName();
```

Note that it must have no declared return type (despite having no declared return type) return type. The function is used to convert type `Vector` to type `double` value.

```
Vector::operator double()
{
 ...
 return a_double_value;
}
```

Experience has shown that often it is better to use member functions.

As you might have noticed, classes require more than do simple C-style structures. In return, they do

into a file. Label each position with the  
the initial conditions (target distance and  
the file. The file contents might look like this:

```

Target Distance: 100, Step Size: 20
0: (x,y) = (0, 0)
1: (x,y) = (-11.4715, 16.383)
2: (x,y) = (-8.68807, -3.42232)
...
26: (x,y) = (42.2919, -78.2594)
27: (x,y) = (58.6749, -89.7309)
After 27 steps, the subject has the
(x,y) = (58.6749, -89.7309)
or
(m,a) = (107.212, -56.8194)
Average outward distance per step =

```

2. Modify the `Vector` class header and implementation so that the magnitude and angle are not calculated until they should be calculated on demand when the methods are called. You should leave the public methods (the methods with the same arguments) but alter the private method and the method implementations.

or equal to 11 stone. (The simplest approach is to initialize to 11 stone and to compare the other stones to it.)

7. A complex number has two parts: a real part and an imaginary part. To write an imaginary number is this:  $(3.0i)$ , where  $i$  is the imaginary part. Suppose  $a = (A, Bi)$ . What are the operations:

- Addition:  $a + c = (A + C, (B + D)i)$
- Subtraction:  $a - c = (A - C, (B - D)i)$
- Multiplication:  $a * c = (A * C - B * D, (A * D + B * C)i)$
- Multiplication: (x a real number):  $r * a = (r * A, r * B * i)$
- Conjugation:  $\sim a = (A, -Bi)$

Define a complex class so that the following results:

```
#include <iostream>
using namespace std;
#include "complex0.h" // to avoid conflict
int main()
{
 complex a(3.0, 4.0); // initialize
```

```
Enter a complex number (q to quit):
real: 10
imaginary: 12
c is (10,12i)
complex conjugate is (10,-12i)
a is (3,4i)
a + c is (13,16i)
a - c is (-7,-8i)
a * c is (-18,76i)
2 * c is (20,24i)
Enter a complex number (q to quit):
real: q
Done!
```

Note that `cin >> c`, through overload parts.





- Using static class members
- Using placement new with objects
- Using pointers to objects
- Implementing a queue abstract data type

This chapter looks at how to use `new` and some of the subtle problems that using `dynamic_cast` presents. It has a short list of topics, but these topics affect compiler and linker behavior, so they are important for understanding the compiler and linker overloading.

Let's look at a specific example of how C++ handles memory. The simplest, most primitive way is to use a character array. This has some drawbacks. You might use a 14-character array to hold the name "Smeadsbury-Crafthovingham". Or to be safer, you might use a 2,000-character array. Or you might create an array of 2,000 such objects, each of which is a character array that is only partly filled. (At that point, you are using a lot of memory.) There is an alternative.

Often it is much better to decide many things when a program runs rather than when it's compiled. One thing a name in an object is to use the `new` operator to allocate a certain amount of memory while the program is running.



under development. It's the first stage of development, allocation, and it does the obvious things correctly in the constructors and destructor. It omits doing some additional good things that the problems the class has should help you understand. Changes you will make later, when you convert

You should note two points about this definition. First, instead of a `char` array to represent a name. The class allocates storage space for the string itself. Instead of allocating space for the string. This arrangement avoids a predefined limit to the string size.

Second, the definition declares the `num_strings` static storage class. A *static class member* has a special status of a static class variable, regardless of the number. The number is shared among all objects of that class, not among all members of a family. If, say, you create 10 `str` members and 10 `len` members, but just one `num_strings` (Figure 12.1). This is convenient for data that should have the same value for all class objects. The `num_strings` is intended to keep track of the number of objects.



Figure 12.1 A string class

By the way, Listing 12.1 uses the `num_strings` static data member, illustrating static data members and as a device to count the number of strings. In general, a string class doesn't need such a member.

Take a look at the implementation of the class methods in Listing 12.2, which handles using a pointer and a static member.

#### Listing 12.2 **strngbad.cpp**

```
// strngbad.cpp -- StringBad class methods
#include <cstring> // s
#include "strngbad.h"
using std::cout;

// initializing static class member
int StringBad::num_strings = 0;

// class methods
```



with a regular C string:

```
StringBad::StringBad(const char * s)
{
 len = std::strlen(s); // set
 str = new char[len + 1]; // allo
 std::strcpy(str, s); // init
 num_strings++; // set
 cout << num_strings << ": \"" << str
 << "\" object created\n"; // For
}
```

Recall that the class `str` member is just a pointer to memory for holding a string. You can pass a string to initialize an object:

```
String boston("Boston");
```

The constructor must then allocate enough memory and must copy the string to that location. Let's go

First, the function initializes the `len` member to be the length of the string. Next, it uses `new` to allocate the new memory, then it assigns the address of the new memory to `str`. Finally, it returns the length of a string, not counting the null terminator. The constructor adds one to `len` to allow space for the null terminator.

}

The destructor begins by announcing when the memory is no longer needed. This is informative but not essential. However, the destructor must free the memory. If the destructor member points to memory allocated with `new`, the pointer expires. But the memory `str` pointed to by `str` is not freed. Deleting an object frees the memory pointed to by `str`, but it does not automatically free memory pointed to by `str`. That is, you must use the destructor. By placing the destructor in the same class as the constructor, you ensure that the memory that a constructor allocated is freed when the object expires.

### Warning

Whenever you use `new` in a constructor to allocate memory, you must provide a corresponding destructor to free that memory. If you use `new` in a function, you should use `delete []` (with brackets).

Listing 12.3, which is taken from a program that demonstrates the `StringBad` class, illustrates when and how the `StringBad` constructor places the object declarations within an inner block. When the inner block's execution exits the block in which an object is declared, the destructor would be called after execution exits the block. In some environments from seeing the destructor call, the program may not close. Be sure to compile Listing 12.2 along



```

 cout << "headline1: " << headline1;
 callme2(headline2);
 cout << "headline2: " << headline2;
 cout << "Initialize one object to\n";
 StringBad sailor = sports;
 cout << "sailor: " << sailor << endl;
 cout << "Assign one object to another\n";
 StringBad knot;
 knot = headline1;
 cout << "knot: " << knot << endl;
 cout << "Exiting the block.\n";
 }
 cout << "End of main()\n";

 return 0;
}

void callme1(StringBad & rsb)
{
 cout << "String passed by reference:\n";
 cout << " \"" << rsb << "\"\n";
}

void callme2(StringBad sb)

```

```
headline1: Celery Stalks at Midnight
String passed by value:
 "Lettuce Prey"
"Lettuce Prey" object deleted, 2 left
headline2: Dû°
Initialize one object to another:
sailor: Spinach Leaves Bowl for Dollars
Assign one object to another:
3: "C++" default object created
knot: Celery Stalks at Midnight
Exiting the block.
"Celery Stalks at Midnight" object deleted
"Spinach Leaves Bowl for Dollars" object
"Spinach Leaves Bowl for Doll8" object de
"@g" object deleted, -1 left
"-|" object deleted, -2 left
End of main()
```

The various nonstandard characters that appear in the output are not valid in a C system; they are one of the signs that String objects are being created and destroyed. the negative object count. Newer compiler/OS combinations will catch this error in the program just before displaying the line above and will cause the program to report a General Protection Fault (GPF). A C program should not access a memory location forbidden to it; this is another

This section of code seems to work fine, to

But then the program executes the followi

```
callme2(headline2);
cout << "headline2: " << headline2 << endl
```

Here, `callme2()` passes `headline2` by value  
cates a serious problem:

String passed by value:

```
"Lettuce Prey"
```

```
"Lettuce Prey" object deleted, 2 left
headline2: D  
```

First, passing `headline2` as a function argu  
called. Second, although passing by value is su  
from change, the function messes up the origi  
nonstandard characters get displayed. (The exa  
pens to sitting in memory.)

Even worse, look at the end of the output  
cally for each of the objects created earlier:

Exiting the block.

```
"Celery Stalks at Midnight" object deleted
```

```
"Spinach Leaves Bowl for Dollars" object d
```

```
"Spinach Leaves Bowl for Doll18" object del
```

tax for the following:

```
StringBad sailor = StringBad(sports); //c
```

Because `sports` is type `StringBad`, a mat  
`StringBad(const StringBad &);`

And it turns out that the compiler autom  
*copy constructor* because it makes a copy of an  
another. The automatic version would not k  
variable, so it would mess up the counting s  
by this example stem from member function  
so let's look at that topic now.

## Special Member Functions

The problems with the `StringBad` class stem  
member functions that are defined automatic  
of these member functions is inappropriate to  
C++ automatically provides the following m

- A default constructor if you define no
- A default destructor if you don't define
- A copy constructor if you don't define

compiler supplies the following default:

```
Klunk::Klunk() { } // implicit default constructor
```

That is, it supplies a constructor (the *default constructor*) that does nothing. It's needed because the compiler needs a constructor to create objects. Without a constructor, the compiler would have to create objects with default values, which is not what you want. The default constructor is the only constructor that does nothing.

```
Klunk klunk; // invokes default constructor
```

The default constructor makes `klunk` like an array. The value at initialization is unknown.

After you define any constructor, C++ does not provide a default constructor. If you want to create objects that aren't initialized, you must define a default constructor explicitly. It's a constructor that sets particular values:

```
Klunk::Klunk() // explicit default constructor
{
 klunk_ct = 0;
 ...
}
```

A constructor with arguments still can be a default constructor. For example, the `Klunk` class could have a constructor that takes an integer argument and sets `klunk_ct` to that value.

```
Klunk(int n = 0) { klunk_ct = n; }
```

You must know two things about a copy

## When a Copy Constructor Is Used

A copy constructor is invoked whenever a new object is created from an existing object of the same kind. This happens in the following situation: when you explicitly initialize a new object with an existing object. For example, given that `motto` is a `StringBad` object, the following code calls the copy constructor:

```
StringBad ditto(motto); // calls StringBad(StringBad)
StringBad metoo = motto; // calls StringBad(StringBad)
StringBad also = StringBad(motto); // calls StringBad(StringBad)
StringBad * pStringBad = new StringBad(motto); // calls StringBad(StringBad)
```

Depending on the implementation, the copy constructor may call the constructor directly to create `metoo` and `also`, or it may create temporary objects whose contents are then assigned to `metoo` and `also`. In the third case, `StringBad` initializes an anonymous object to `motto` and then passes a pointer to that object to the `pStringBad` pointer.

Less obviously, a compiler uses a copy constructor to create copies of an object. In particular, it's used when a function calls `callme2()` does in Listing 12.3) or when a function

```
private members is not allowed);

StringBad sailor;
sailor.str = sports.str;
sailor.len = sports.len;
```

If a member is itself a class object, the copy constructor copies one member object to another. Static member variables are not copied because they belong to the class as a whole instead of to an object. This illustrates the action of an implicit copy constructor.

## Back to Stringbad: Where the Copy Constructor Lives

You are now in a position to understand the output of the program. I will assume that the output is the one shown after the program has run. The program output indicates two more objects were created than that the program does create two additional objects. The copy constructor is used to initialize the objects. When the copy constructor function is called, and it is used to initialize the objects. The copy constructor doesn't vocalize its activities, so it doesn't increment the `num_strings` counter. However, when the copy constructor is invoked upon the demise of all objects, regardless of how many objects are created, the weirdness is a problem because it means the program doesn't count. The solution is to provide an explicit constructor.

3.

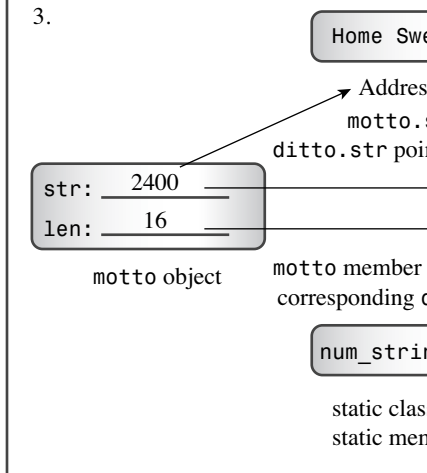


Figure 12.2 An inside look

### Tip

If your class has a static data member whose address you need, you should provide an explicit copy constructor.



the destructor for `sailor`, and this results in undefined behavior. In the case of Listing 12.3, the program produces memory mismanagement.

Another disturbing symptom is that attempts to delete the object may cause the program to abort. Microsoft Visual C++ plays an error message window saying “Debug Assertion Failed!” and reports “double free or corruption” and aborts the program. Sometimes it shows messages or even no message, but the same event occurs.

## Fixing the Problem by Defining an Explicit Copy Constructor

The cure for the problems in the class design is to define an explicit copy constructor. Instead of just copying the address of the string, the copy constructor must allocate memory and assign the address of the duplicate to the `str` member of the new object. This is a string rather than referring to another object's `str` member. This is a different string rather than making duplicate references to the same string. This is how you can code the `String` copy constructor.

```
StringBad::StringBad(const StringBad & st)
{
 num_strings++; // handle static member
 len = st.len; // same length
 str = new char [len + 1]; // allot space
 std::strcpy(str, st.str); // copy string
```

2. `String ditto(motto); // deep`

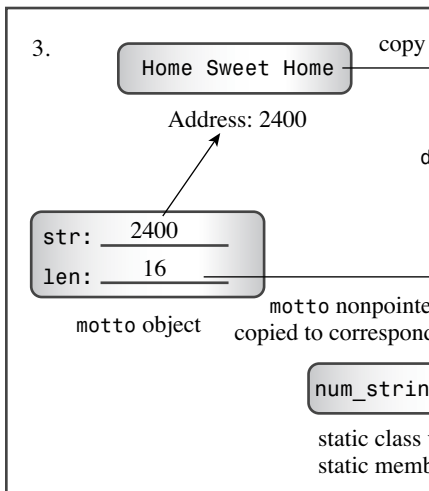


Figure 12.3 An inside

## When an Assignment Operator Is Used

An overloaded assignment operator is used when creating a new object from an existing object:

```
StringBad headline1("Celery Stalks at Midnight");
...
StringBad knot;
knot = headline1; // assignment operator
```

An assignment operator is not necessarily used.

```
StringBad metoo = knot; // use copy constructor
```

Here `metoo` is a newly created object being assigned. The copy constructor is used. However, as mentioned before, C++ handles this statement in two steps: using the copy constructor to create the object and then using assignment to copy the object's data. Every assignment always invokes a copy constructor, and for this reason, the assignment operator is used.

Like a copy constructor, an implicit implementation of the assignment operator forms a member-to-member copy. If a member function uses the assignment operator defined for a particular member, static data members are unaffected.

## Fixing Assignment

The solution for the problems created by an `operator` to provide your own assignment operator definition. The implementation is similar to that of the copy

- Because the target object may already exist, the function should use `delete []` to free former objects.
- The function should protect against assigning the same memory described previously. Objects that are reassigned.
- The function returns a reference to the object.

By returning an object, the function can be used with built-in types can be chained. That is, if `s0`, `s1`, and `s2` are `String` objects, write the following:

```
s0 = s1 = s2;
```

In function notation, this becomes the following:

```
s0.operator=(s1.operator=(s2));
```

don't first apply the delete operator, the previous version of the program no longer has a pointer to the old object.

Next, the program proceeds like a copy constructor: it creates a new string and then copying the string from the old object.

When it is finished, the program returns `*t` to the caller.

Assignment does not create a new object, so the program does not need static data member `num_strings`.

Adding the copy constructor and the assignment operator to the `StringBad` class clears up all the problems. Here is the output after these changes have been made:

```
End of main()
"Celery Stalks at Midnight" object deleted
"Spinach Leaves Bowl for Dollars" object deleted
"Spinach Leaves Bowl for Dollars" object deleted
"Lettuce Prey" object deleted, 1 left
"Celery Stalks at Midnight" object deleted
```

The object counting is correct now, and no more memory leaks.

return the normally (/),

The first new method returns the length of a string. The other two methods are comparison functions that allow you to compare strings. The `len` function returns the number of characters in a string. The static class method `len` returns the number of elements in a slice. The member `num_strings`. Let's look at some details.

## The Revised Default Constructor

The new default constructor merits notice. It is defined as follows:

```
String::String()
{
 len = 0;
 str = new char[1];
 str[0] = '\\0'; // default character
}
```

You might wonder why the code uses

```
str = new char[1];
```

and not this:

```
str = new char;
```

## C++11 Null Pointer

In C++98, the literal `0` has two meanings—it can be a null pointer—thus making it difficult for the programmer to distinguish between the two. Sometimes programmers use `(void *)0` to represent the null pointer, but the pointer itself may have a nonzero internal representation. C++11 provides a better solution by introducing `nullptr`, a new keyword for the null pointer. C++11 provides a better solution by introducing `nullptr`, a new keyword for the null pointer. You still can use `0` as before, but code would be invalidated—but henceforth the

```
str = nullptr; // C++11 null pointer
```

## Comparison Members

Three of the methods in the `String` class perform string comparison. The `compare` method returns `true` if the first string comes before the second (lexicographically, precisely, in the machine collating sequence). The `compare` method returns `false` if the first string is greater than the second. The `compare` method returns `0` if the two strings are equal. The `compare` method is used in the standard string comparison functions is to use the standard string comparison functions. The `compare` method returns a negative value if its first argument precedes the second, a zero value if the first argument is equal to the second, and a positive value if the first follows the second. The `compare` method is used in the standard string comparison functions like this:

```
bool operator<(const String &st1, const String &st2)
{
 if (std::strcmp(st1.str, st2.str) < 0)
```

good choice for an inline function.

Making the comparison functions friends of string objects and regular C strings. For example, suppose you have the following code:

```
if ("love" == answer)
```

This gets translated to the following:

```
if (operator==("love", answer))
```

The compiler then uses one of the constructors:

```
if (operator==(String("love"), answer))
```

And this matches the prototype.

## Accessing Characters by Using Bracket Operator

With a standard C-style string, you can use bracket notation to access individual characters:

```
char city[40] = "Amsterdam";
cout << city[0] << endl; // display the first character
```

In C++ the two bracket symbols constitute the subscript operator. In addition, you can overload this operator by using a member function. The C++ operator (one with two operands) puts the value of the first operand to the power of the second operand. For example,  $2 + 5$ . But the bracket operator places one o



```
cout << opera[4];
```

becomes this:

```
cout << opera.operator[4];
```

The return value is `opera.str[4]`, or the access to private data.

Declaring the return type as type `char &` argument. For example, you can use the following

```
String means("might");
means[0] = 'r';
```

The second statement is converted to an `operator[]` operation:

```
means.operator[] [0] = 'r';
```

This assigns 'r' to the method's return value `means.str[0]`, making the code equivalent to:

```
means.str[0] = 'r';
```

This last line of code violates private access. In a public method, it is allowed to alter the array content. The string becomes "right".

```
cout << text[1]; // ok, uses non-const
cout << answer[1]; // ok, uses const ver
cin >> text[1]; // ok, uses non-const
cin >> answer[1]; // compile-time error
```

## Static Class Member Functions

It's possible to declare a member function as appear in the function declaration but not in rate.) This has two important consequences.

First, a static member function doesn't have doesn't even get a `this` pointer to play with. the public section, it can be invoked using the tor. For instance, you can give the `String` class with the following prototype/definition in the

```
static int HowMany() { return num_strings;
```

It could be invoked like this:

```
int count = String::HowMany(); // invoking
```

The second consequence is that because a with a particular object, the only data member

serves as a conversion function.

2. In Listing 12.6, later in this chapter, the `operator=(const String &)` function to copy the contents of the `String` object to the name object.
3. The program calls the `~String()` destructor to destroy the `String` object.

The simplest way to make the process more efficient is to overload the `operator+` so that it works directly with ordinary `char*` strings, thus avoiding creating and destroying a temporary object. Here is the code:

```
String & String::operator=(const char * s)
{
 delete [] str;
 len = std::strlen(s);
 str = new char[len + 1];
 std::strcpy(str, s);
 return *this;
}
```

As usual, you must deallocate memory for the old string and allocate memory for the new string.

Listing 12.4 shows the revised class declaration. As mentioned, it defines the constant `CINLIM`, which is the maximum length of a string.

```

 int length () const { return len; }

// overloaded operator methods
 String & operator=(const String &);
 String & operator=(const char *);
 char & operator[] (int i);
 const char & operator[] (int i) const;

// overloaded operator friends
 friend bool operator<(const String &s
 friend bool operator>(const String &s
 friend bool operator==(const String &s
 friend ostream & operator<<(ostream &
 friend istream & operator>>(istream &

// static function
 static int HowMany();
};
#endif

```

---

Listing 12.5 presents the revised method definitions.

#### Listing 12.5    **string1.cpp**

---

```

// string1.cpp -- String class methods
#include <cstring> // str
#include "string1.h" // inc

```

```

{
 len = 4;
 str = new char[1];
 str[0] = '\\0'; // default
 num_strings++;
}

```

```

String::String(const String & st)
{
 num_strings++; // handle s
 len = st.len; // same len
 str = new char [len + 1]; // allot sp
 std::strcpy(str, st.str); // copy str
}

```

```

String::~~String() // r
{
 --num_strings; // r
 delete [] str; // r
}

```

// overloaded operator methods

// assign a String to a String

```

 {
 return str[i];
 }

 // read-only char access for const String
 const char & String::operator[](int i) const
 {
 return str[i];
 }

 // overloaded operator friends

 bool operator<(const String &st1, const String &st2)
 {
 return (std::strcmp(st1.str, st2.str) < 0);
 }

 bool operator>(const String &st1, const String &st2)
 {
 return st2 < st1;
 }

 bool operator==(const String &st1, const String &st2)
 {

```

The overloaded `>>` operator provides a similar interface to `getline`, returning a `String` object. It assumes an input line of size `ArSize` and discards any characters beyond that limit. Keep in mind that the expression `!input` in an `if` condition evaluates to `false` if input is not empty, and to `true` in an end-of-file condition, or in the case of getting an error.

Listing 12.6 exercises the `String` class with arrays of `String` objects. The program has the user enter sayings, stores them in an array of `String` objects, plays them, and reports which string is the shortest.

### Listing 12.6    **sayings1.cpp**

---

```
// sayings1.cpp -- using expanded String class
// compile with string1.cpp
#include <iostream>
#include "string1.h"
const int ArSize = 10;
const int MaxLen = 81;
int main()
{
 using std::cout;
 using std::cin;
 using std::endl;
 String name;
```

```
 cout << "Here are your sayings:\n";
 for (i = 0; i < total; i++)
 cout << sayings[i][0] << ": ";

 int shortest = 0;
 int first = 0;
 for (i = 1; i < total; i++)
 {
 if (sayings[i].length() < sayings[first].length())
 shortest = i;
 if (sayings[i] < sayings[first])
 first = i;
 }
 cout << "Shortest saying:\n" << sayings[shortest];
 cout << "First alphabetically:\n" << sayings[first];
 cout << "This program used " << St << " String objects. Bye.\n";
 }
 else
 cout << "No input! Bye.\n";
 return 0;
}
```

---



1: a fool and his money are soon parted  
2: penny wise, pound foolish  
3: the love of money is the root of much e  
4: out of sight, out of mind  
5: absence makes the heart grow fonder  
6: absinthe makes the hart grow fonder  
7:

Here are your sayings:

a: a fool and his money are soon parted  
p: penny wise, pound foolish  
t: the love of money is the root of much e  
o: out of sight, out of mind  
a: absence makes the heart grow fonder  
a: absinthe makes the hart grow fonder

Shortest saying:

penny wise, pound foolish

First alphabetically:

a fool and his money are soon parted

This program used 11 String objects. Bye.

character as a visual reminder that this value favored a simple 0 instead of the equivalent the `nullptr` keyword as a better alternative

- You should define a copy constructor to do deep copying. Typically, the constructor

```
String::String(const String & st)
{
 num_strings++; // ha
 len = st.len; // sa
 str = new char [len + 1]; // al
 std::strcpy(str, st.str); // co
}
```

In particular, the copy constructor should do deep copying; it should copy the data, not just the address. It should not copy static class members whose value would be shared.

- You should define an assignment operator to do deep copying. Typically, the class example:

```

 }

String::String(const char * s)
{
 len = std::strlen(s);
 str = new char; // oops, no
 std::strcpy(str, s); // oops, no
}

```

```

String::String(const String & st)
{
 len = st.len;
 str = new char[len + 1]; // go
 std::strcpy(str, st.str); // go
}

```

The first constructor fails to use `new` to initialize the default object, applies `delete` to `str`. The result of initializing by `new` is undefined, but it is probably

```

String::String()
{
 len = 0;
 str = new char[1]; // uses new with 1
 str[0] = '\0';
}

```

```
{
 delete str; // oops, should be d
}
```

The destructor uses `delete` incorrectly. Because `str` is a pointer to a dynamically allocated array of characters, the destructor should delete an array.

## Memberwise Copying for Classes

Suppose you use the `String` class, or, for that matter, `std::string`, for class members:

```
class Magazine
{
private:
 String title;
 string publisher;
 ...
};
```

`String` and `string` both use dynamic memory management. To write a copy constructor and assignment operator for `Magazine`, you must call `String` and `string` methods, not in itself. The default memberwise copying of `String` and `string` does the right thing, but you must use the `String` and `string` methods. If you copy or assign one `Magazine` object to another, the copy constructors and assignment operators must call the `String` and `string` methods.

in Chapter 11. The function would be used in

```
Vector force1(50,60);
Vector force2(10,70);
Vector max;
max = Max(force1, force2);
```

Either of the following two implementations

```
// version 1
Vector Max(const Vector & v1, const Vector & v2)
{
 if (v1.magval() > v2.magval())
 return v1;
 else
 return v2;
}
```

```
// version 2
const Vector & Max(const Vector & v1, const Vector & v2)
{
 if (v1.magval() > v2.magval())
 return v1;
 else
 return v2;
}
```

operator=() method returns a reference to s

The return value of operator<<() is used

```
String s1("Good stuff");
cout << s1 << "is coming!";
```

Here, the return value of operator<<(cout) is the string "is coming!". Here, the return type is ostream. Using an ostream return type would be a good idea, and, as it turns out, the ostream class does exactly that. Fortunately, returning a reference to cout poses no problem in the calling function.

## Returning an Object

If the object being returned is local to the calling function, it is returned by reference because the local object has its own memory address. Thus, when control returns to the calling function, the reference can refer. In these circumstances, returning by reference. Typically, overloaded arithmetic operators are returned by reference, which uses the Vector class again:

```
Vector force1(50,60);
Vector force2(10,70);
Vector net;
net = force1 + force2;
```

```
net = force1 + force2;
```

However, the definition also allows you to

```
force1 + force2 = net;
```

```
cout << (force1 + force2 = net).magval() <
```

Three questions immediately arise. Why would they be possible? What do they do?

First, there is no sensible reason for writing these sensible reasons. People, even programmers, might think that the operator==() were defined for the vector class.

```
if (force1 + force2 = net)
```

instead of this:

```
if (force1 + force2 == net)
```

Also programmers tend to be ingenious, and sometimes make mistakes.

Second, this code is possible because the compiler creates a temporary object to represent the return value. So in the first statement, force2 stands for that temporary object. In the second statement, force2 stands for net. In Statements 2 and 3, net is assigned

cally. Another approach is to use pointers to p  
gories. Listing 12.7 implements this approach  
shortest pointer points to the first object in  
object with a shorter string, it resets shortes  
pointer tracks the alphabetically earliest string  
new objects; they merely point to existing ob  
allocate additional memory.

For variety, the program in Listing 12.7 us  
object:

```
String * favorite = new String(sayings[ch
```

Here the pointer `favorite` provides the o  
new. This particular syntax means to initialize  
`sayings[choice]`. That invokes the copy co  
copy constructor (`const String &`) matches  
The program uses `srand()`, `rand()`, and `time`

#### Listing 12.7    **sayings2.cpp**

---

```
// sayings2.cpp -- using pointers to obje
// compile with string1.cpp
#include <iostream>
#include <cstdlib> // (or stdlib.h)
#include <ctime> // (or time.h) fo
```



```

 else
 sayings[i] = temp; // overw
 }
 int total = i; // tota

 if (total > 0)
 {
 cout << "Here are your sayings:\n";
 for (i = 0; i < total; i++)
 cout << sayings[i] << "\n";

 // use pointers to keep track of shortest
 String * shortest = &sayings[0]; //
 String * first = &sayings[0];
 for (i = 1; i < total; i++)
 {
 if (sayings[i].length() < shortest->length())
 shortest = &sayings[i];
 if (sayings[i] < *first) //
 first = &sayings[i]; //
 }
 cout << "Shortest saying:\n" << *shortest;
 cout << "First alphabetically:\n" << *first;
 srand(time(0));
 }

```

Also the usual conversions invoked by `prototype` takes place as long as there is no ambiguity in the default constructor:

```
Class_name * ptr = new Class_name;
```

Here's a sample run of the program in Listing 10.1:

```
Hi, what's your name?
```

```
>> Kirt Rood
```

```
Kirt Rood, please enter up to 10 short sayings:
```

```
1: a friend in need is a friend indeed
```

```
2: neither a borrower nor a lender be
```

```
3: a stitch in time saves nine
```

```
4: a niche in time saves stine
```

```
5: it takes a crook to catch a crook
```

```
6: cold hands, warm heart
```

```
7:
```

```
Here are your sayings:
```

```
a friend in need is a friend indeed
```

```
neither a borrower nor a lender be
```

```
a stitch in time saves nine
```

```
a niche in time saves stine
```

```
it takes a crook to catch a crook
```

```
cold hands, warm heart
```



- You declare a pointer to an object by using

```
String * glamour;
```

- You can initialize a pointer to point to

```
String * first = &sayings[0];
```

- You can initialize a pointer by using new

```
String * favorite = new String(sayings[0]);
```

Also see Figure 12.6 for a more detailed example of using new.

- Using new with a class invokes the appropriate constructor to create a newly created object:

```
// invokes default constructor
String * gleep = new String;
```

```
// invokes the String(const char *) constructor
String * glop = new String("my my my");
```

initializing a pointer using new  
and the String(const char\*)  
class constructor:

String \* p = new String("g");

Initializing a pointer using new  
and the String(const String &)  
class constructor:

String \* p = new String(\*f);

Using the -> operator  
to access a class  
method via a pointer:

if (saying  
    {  
        objec

Using the \* deferencing  
operator to obtain an  
object from a pointer:

if (saying  
    {  
        objec

Figure 12.5 Pointers

## Looking Again at Placement new

Recall that placement new allows you to specify memory. Chapter 9, “Memory Models and N context of built-in types. Using placement new 12.8 uses placement new along with regular new a class with a chatty constructor and destructor objects.

### Listing 12.8    `placnew1.cpp`

---

```
// placenew1.cpp -- new, placement new,
#include <iostream>
#include <string>
#include <new>
using namespace std;
const int BUF = 512;
```

```
 cout << pc1 << ": ";
 pc1->Show();
 cout << pc2 << ": ";
 pc2->Show();

 JustTesting *pc3, *pc4;
 pc3 = new (buffer) JustTesting("Bad Id");
 pc4 = new JustTesting("Heap2", 10);

 cout << "Memory contents:\n";
 cout << pc3 << ": ";
 pc3->Show();
 cout << pc4 << ": ";
 pc4->Show();

 delete pc2;
 delete pc4;
 delete [] buffer;
 cout << "Done\n";
 return 0;
 }
```

---

There are a couple problems with placement new. First, creating a second object, placement new simply creates a new object, not a new object. Not only is the first object with a new object. Not only is the first object called for the first object. This, of course, would require dynamic memory allocation for its members.

Second, using delete with pc2 and pc4 are two objects that pc2 and pc4 point to. But using delete the destructors for the objects created with placement new.

One lesson to be learned here is the same as before. You have to manage the memory locations in a buffer. If you have different locations, you provide two different buffers. If the locations don't overlap. You can, for example,

```
pc1 = new (buffer) JustTesting;
pc3 = new (buffer + sizeof (JustTesting))
```

Here the pointer pc3 is offset from pc1 by sizeof (JustTesting).

The second lesson to be learned here is that if you have two objects on the heap, you need to arrange for their destructors to be called. If you have objects on the heap, you can use this:

```
delete pc2; // delete object pointed to by pc2
```



of the rare cases that require an explicit call. And finally, by deferring the destruction of the object to be destroyed. Because there are three pointers to the object, these pointers:

```
pc3->~JustTesting(); // destroy object pc3
pc1->~JustTesting(); // destroy object pc1
```

Listing 12.9 fixes Listing 12.8 by managing the object's lifetime and by adding appropriate uses of `delete` and `delete[]`. The problem in fact is the proper order of deletion. The objects must be destroyed in order opposite that in which they were created. In principle, a later object might have dependence on an earlier one to hold the objects should be freed only after the later object has been freed.

#### Listing 12.9 **placnew2.cpp**

---

```
// placnew2.cpp -- new, placement new, new[]
#include <iostream>
#include <string>
#include <new>
using namespace std;
const int BUF = 512;

class JustTesting
{
```

```
 cout << pc2 << ": ";
 pc2->Show();

 JustTesting *pc3, *pc4;
 // fix placement new location
 pc3 = new (buffer + sizeof (JustTesting))
 JustTesting("Better Idea");
 pc4 = new JustTesting("Heap2", 10);

 cout << "Memory contents:\n";
 cout << pc3 << ": ";
 pc3->Show();
 cout << pc4 << ": ";
 pc4->Show();

 delete pc2; // free Heap1
 delete pc4; // free Heap2
 // explicitly destroy placement new object
 pc3->~JustTesting(); // destroy object
 pc1->~JustTesting(); // destroy object
 delete [] buffer; // free buffer
 cout << "Done\n";
 return 0;
 }
```

---

# Reviewing Techniques

By now, you've encountered several programming class-related problems, and you may be having trouble. The following sections summarize several techniques.

## Overloading the << Operator

To redefine the << operator so that you use it with your class, you define a friend operator function that has the following signature:

```
ostream & operator<<(ostream & os, const C_Name & c)
{
 os << ... ; // display object contents
 return os;
}
```

Here *c\_name* represents the name of the class. To return the required contents, you can use those methods in the class. Remember to declare the friend function with the friend status.

does not know them and thus won't catch you

- Any class member that points to memory should be null after the delete operator applied to it in the class destructor.
- If a destructor frees memory by applying delete, every constructor for that class should make sure by setting the pointer to the null pointer.
- Constructors should settle on using either new or delete. Both. The destructor should use delete. If the constructors use new, the destructor should use delete if the constructors use new.
- You should define a copy constructor to create a new object by copying an existing object. This is useful for creating a pointer to existing memory. This is useful for creating a new object to another. The constructor should be defined as follows:

```
className(const className &)
```

- You should define a class member function that has a function definition with the following signature: `void c_name::function_name(args)`. This is a member of the `c_name` class and has the following example assumes that the constructor is defined using new []:

A rather natural way of representing the problem is as a queue. A queue is an abstract data type (ADT) that holds items. Items are added to the rear of the queue, and items are removed from the front. It's a bit like a stack, except that a stack has additional operations. A stack is a LIFO (last in, first out) structure, while a queue is a FIFO (first in, first out) structure. Conceptually, a queue is like a line at a bank. The queue is well suited to the task. So one part of the project is to implement a queue. In the next chapter you'll read about the Standard Template Library, which provides a queue. It's better to develop your own than by just reading about it.

The items in the queue will be customers. The queue is a line at a bank. That is, that, on average, a third of the customers will take one minute, a third will take two minutes, and a third will take three minutes. The random intervals, but the average number of customers in the queue is 1.5. The next more parts of your project will be to design a program that simulates the interaction of the queue. (See Figure 12.7).

## A Queue Class

The first order of business is to design a Queue class. The first thing you need to decide is the kind of queue you'll need:

- A queue holds an ordered sequence of items.
- A queue has a limit on the number of items it can hold.
- You should be able to create an empty queue.
- You should be able to check whether a queue is empty.
- You should be able to check whether a queue is full.
- You should be able to add an item to the queue.
- You should be able to remove an item from the queue.
- You should be able to determine the number of items in the queue.

As usual when designing a class, you need to decide on the implementation.

### The Queue Class Interface

The queue attributes listed in the preceding section are the ones you need for a queue class:

## The Queue Class Implementation

After you determine the interface, you can implement the queue. One approach is to use an array to represent the queue data. One approach is to use a linked list to represent the queue data. One approach is to use a queue of the required number of elements. However, an array-based queue has limitations. For example, removing an item from the queue requires shifting every remaining element one unit closer to the front of the queue. Something more elaborate, such as treating the array as a circular queue, is a reasonable fit to the requirements of a queue. A linked list is a reasonable fit to the requirements of a queue. Each *node* contains the information to be stored in the queue. Each node in the list. For the queue in this example, you can use a structure to represent a node:

```
struct Node
{
 Item item; // data stored in node
 struct Node * next; // pointer to next node
};
```

Figure 12.8 illustrates a linked list.

The example shown in Figure 12.8 is called a singly linked list. Each node contains a pointer to the next node in the list. If the pointer to the next node in the list is set to NULL (or, equivalently, 0), the list ends. With C++11, you should use the new

a queue always adds a new item to the end of the queue. The `rear` member points to the last node, too (see Figure 10.10). The `items` member keeps track of the maximum number of items in the queue. The current number of items is stored in the `items` member. Thus, the private part of the `Queue` class is:

```
class Queue
{
private:
 // class scope definitions
 // Node is a nested structure definition
 struct Node { Item item; struct Node *next; };
 enum {Q_SIZE = 10};
 // private class members
 Node * front; // pointer to front of queue
 Node * rear; // pointer to rear of queue
 int items; // current number of items in queue
 const int qsize; // maximum number of items in queue
 ...
public:
 //...
};
```

The declaration uses the C++ ability to nest structures. By placing the `Node` declaration inside the `Queue` class, `Node` is a type that you can use to declare class members.



type is restricted to the class. That way, you don't have a `Node` conflicting with some global declaration of `Node` in the class. Some obsolescent compilers do not support nested structures. If you don't, then you have to define a `Node` structure globally.

## Nested Structures and Classes

A structure, a class, or an enumeration declaration can be *nested* in the class. It has class scope. Such a declaration is *private*. Rather, it specifies a type that can be used in the class. If the declaration is made in the private section of the class, then it is private. If the declaration is made in the public section of the class, then it is public. If the declaration is used out of the class, through use of the scope resolution operator, then it is public. If the declaration were declared in the public section of the `Queue` class, then you could use `Queue::Node` outside the `Queue` class.

After you settle on a data representation, then you can implement the class.

## The Class Methods

A class constructor should provide values for the data members. If the class, for example, begins in an empty state, you should provide a constructor that initializes the data members.

```

Queue::Queue(int qs) : qsize(qs) // ini
{
 front = rear = NULL;
 items = 0;
}

```

In general, the initial value can involve complex calculation on the argument list. The technique is not limited to the Queue constructor like this:

```

Queue::Queue(int qs) : qsize(qs), front(N
{
}

```

Only constructors can use this initializer-list syntax for const class members. You also can declare class members as references:

```

class Agency {...};
class Agent
{
private:
 Agency & belong; // must use initialization
 ...
};
Agent::Agent(Agency & a) : belong(a) {...}

```

Data members are initialized in the order in which they are declared in the order in which initializers are listed.

## Caution

You can't use the member initializer list syntax

The parenthesized form used in the member initializations, too. That is, if you like, you can re

```
int games = 162;
double talk = 2.71828;
```

with

```
int games(162);
double talk(2.71828);
```

This allows initializing built-in types to local

## C++11 Member In-Class Initialization

C++11 allows you to do what would seem to

```
class Classy
{
 int mem1 = 10; // in-class in
 const int mem2 = 20; // in-class in
```

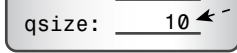
```

 if (isfull())
 return false;
 Node * add = new Node; // create node
// on failure, new throws std::bad_alloc
 add->item = item; // set node p
 add->next = NULL; // or nullptr
 items++;
 if (front == NULL) // if queue i
 front = add; // place item
 else
 rear->next = add; // else place
 rear = add; // have rear
 return true;
 }

```

In brief, the method goes through the fol

1. Terminate if the queue is already full. (is selected by the user via the construct
2. Create a new node. If new can't do so, i up in Chapter 15, "Friends, Exceptions unless one provides additional program terminates.

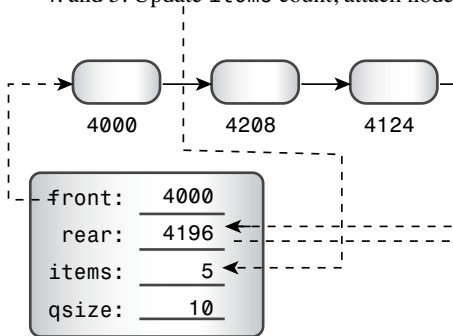


Queue object

2. Create n

3. Copy values to node and set next pointer

4. and 5. Update items count, attach node



Queue object

Figure 12.10 E

3. Decrease the item count (`items`) by one.
4. Save the location of the front node for later.
5. Take the node off the queue. This is accomplished by changing the pointer `front` to point to the next node in the queue.
6. To conserve memory, delete the former front node.
7. If the list is now empty, set `rear` to `NULL`. (In this case, after setting `front->next` to `NULL`, in C++11, you can use `nullptr`.

Step 4 is necessary because step 5 erases the node that `front` points to. The next node is.

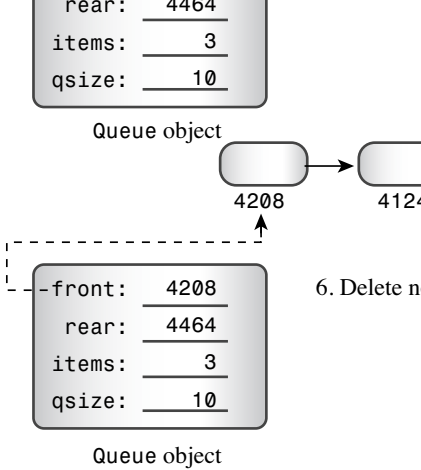


Figure 12.11 D

answer is, “Does the default memberwise copy of a Queue object work?” Memberwise copying of a Queue object would copy the front and rear of the same linked list as the original. If the original Queue object changes the shared linked list, then the copy’s rear pointer gets updated, essentially pointing to the original object. Clearly, then, cloning or copying a Queue constructor and an assignment constructor that copy the original object.

Of course, that raises the question of why you would want to save snapshots of a queue. Perhaps you would like to provide identical input to multiple simulations. Or you would like to be useful to have operations that split a queue into two, like opening an additional checkout stand. Similarly, you might want to truncate a queue.

But suppose you don’t want to do any of these things. You might simply ignore those concerns and use the method as is. However, at some time in the future, you might need to do something. And you might forget that you failed to provide a copy constructor. Your programs will compile and run, but they will be buggy. So it would seem that it’s best to provide a copy constructor even though you don’t need them now.



Chapter 18 returns to this topic.

Are there any other effects to note? Yes. Reference objects are passed (or returned) by value. However, it is a preferred practice of passing objects as references to other temporary objects. But the `Queue` definition uses objects, such as overloading the addition operator.

## The Customer Class

At this point, we need to design a customer class with properties, such as a name, account numbers, and other properties you need for the simulation are what is required for the customer's transaction. When a customer arrives, the program should create a new customer object with an arrival and a randomly generated value for the processing time. When the customer reaches the front of the queue, the program should calculate the customer's joining time to get the customer's waiting time. We will implement the `Customer` class:

```
class Customer
{
private:
 long arrive; // arrival time for customer
 int processtime; // processing time for customer
```

```

// This queue will contain customer items
class Customer
{
private:
 long arrive; // arrival time f
 int processtime; // processing tim
public:
 Customer() { arrive = processtime = 0; }
 void set(long when);
 long when() const { return arrive; }
 int ptime() const { return processtim
};

```

```

typedef Customer Item;

```

```

class Queue
{
private:
// class scope definitions
 // Node is a nested structure definit
 struct Node { Item item; struct Node
 enum {Q_SIZE = 10};
// private class members
 Node * front; // pointer to fro

```

```

// Queue methods
Queue::Queue(int qs) : qsize(qs)
{
 front = rear = NULL; // or nullptr
 items = 0;
}

Queue::~~Queue()
{
 Node * temp;
 while (front != NULL) // while queue
 {
 temp = front; // save address
 front = front->next; // reset pointer
 delete temp; // delete from
 }
}

bool Queue::isempty() const
{
 return items == 0;
}

```

```

 rear->next = add; // else place
rear = add; // have rear
return true;
}

```

```

// Place front item into item variable and
bool Queue::dequeue(Item & item)

```

```

{
 if (front == NULL)
 return false;
 item = front->item; // set item to
items--;
 Node * temp = front; // save location
 front = front->next; // reset front
 delete temp; // delete front
 if (items == 0)
 rear = NULL;
 return true;
}

```

```

// customer method

```

```

// when is the time at which the customer
// the arrival time is set to when and then

```

3. If a customer is being processed, decrement the number of customers in the queue.
4. Track various quantities, such as the number of customers turned away, cumulative time, and average queue length.

When the simulation cycle is finished, the program prints out the findings.

An interesting matter is how the program computes the time a customer arrived. Suppose that on average, 10 customers arrive at a counter every 6 minutes. The program computes `min_per_cust`. However, having a customer service time of `min_per_cust` is not realistic. What you really want (at least most of the time) is a customer to arrive to a customer every 6 minutes. The program uses a random number to show a customer shows up during a cycle:

```
bool newcustomer(double x)
{
 return (std::rand() * x / RAND_MAX < 1);
}
```

Here's how it works. The value `RAND_MAX` is defined in `stdlib.h` and represents the largest value the random number generator can return (the maximum value). Suppose that `x`, the average time between arrivals, is 6. Then `x / RAND_MAX` will be somewhere between

```
int main()
{
 using std::cin;
 using std::cout;
 using std::endl;
 using std::ios_base;
 // setting things up
 std::srand(std::time(0)); // rand

 cout << "Case Study: Bank of Heather\n";
 cout << "Enter maximum size of queue:\n";
 int qs;
 cin >> qs;
 Queue line(qs); // line queue

 cout << "Enter the number of simulations:\n";
 int hours; // hours of simulation
 cin >> hours;
 // simulation will run 1 cycle per minute
 long cyclelimit = MIN_PER_HR * hours;

 cout << "Enter the average number of customers per hour:\n";
 double perhour; // average # of customers per hour
```

```

 temp.set(cycle); // cycle
 line.enqueue(temp); // add to line
 }
}
if (wait_time <= 0 && !line.isEmpty())
{
 line.dequeue(temp); // a customer served
 wait_time = temp.ptime(); // find next time
 line_wait += cycle - temp.when_served++;
}
if (wait_time > 0)
 wait_time--;
sum_line += line.queuecount();
}

// reporting results
if (customers > 0)
{
 cout << "customers accepted: " << customers;
 cout << " customers served: " << customers_served;
 cout << " turnaways: " << turnaways;
 cout << "average queue size: ";
 cout.precision(2);
 cout.setf(ios_base::fixed, ios_base::floatfield);
 cout << sum_line / customers << endl;
}

```

RAND\_MAX yourself.

Here are a few sample runs of the program

```
Case Study: Bank of Heather Automatic Tel
Enter maximum size of queue: 10
Enter the number of simulation hours: 100
Enter the average number of customers per
customers accepted: 1485
 customers served: 1485
 turnaways: 0
average queue size: 0.15
 average wait time: 0.63 minutes
Done!
```

```
Case Study: Bank of Heather Automatic Tel
Enter maximum size of queue: 10
Enter the number of simulation hours: 100
Enter the average number of customers per
customers accepted: 2896
 customers served: 2888
 turnaways: 101
average queue size: 4.64
 average wait time: 9.63 minutes
Done!
```



```
customers accepted: 114
 customers served: 110
 turnaways: 0
average queue size: 2.15
 average wait time: 4.52 minutes
Done!
```

```
Case Study: Bank of Heather Automatic Tell
Enter maximum size of queue: 10
Enter the number of simulation hours: 4
Enter the average number of customers per
customers accepted: 121
 customers served: 116
 turnaways: 5
average queue size: 5.28
 average wait time: 10.72 minutes
Done!
```

```
Case Study: Bank of Heather Automatic Tell
Enter maximum size of queue: 10
Enter the number of simulation hours: 4
Enter the average number of customers per
customers accepted: 112
 customers served: 109
```

demise of an object automatically triggers the

Objects that have members pointing to memory with initializing one object to another or assignment. C++ uses memberwise initialization and assignment. The assigned-to object winds up with exact copies of the original member points to a block of data, though. When the program eventually deletes the two objects, it delete the same block of memory twice, which is a special copy constructor that redefines initialization. In each case, the new definition should create a new object point to the copies. That way you have separate but identical data, with no overlap. The assignment operator. In each case, the goal is to have real data and not just pointers to the data.

When an object has automatic storage or static storage, the object is called automatically when the object is created. You can create an object by using `new` and assign its address to a pointer. The object is called automatically when you apply `delete` to the pointer. You can manage for class objects by using placement `new`. The responsibility of calling the destructor for that object is the responsibility of calling the destructor for that object. The method with a pointer to the object. C++ allows static member definitions inside a class. Such nested

```
 int items = 0;
 const int qsize = Q_SIZE;
 ...
};
```

This is equivalent to using a member initialization list. If a membership initialization list will override the default value, the default value is ignored.

As you might have noticed, classes require more code than do simple C-style structures. In return, they do

## Chapter Review

1. Suppose a String class has the following

```
class String
{
private:
 char * str; // points to string
 int len; // holds length of string
 //...
};
```

4. Identify and correct the errors in the following code:

```
class nifty
{
// data
 char personality[];
 int talents;
// methods
 nifty();
 nifty(char * s);
 ostream & operator<<(ostream & o) const;
}

nifty::nifty()
{
 personality = NULL;
 talents = 0;
}

nifty::nifty(char * s)
{
 personality = new char [strlen(s)];
 personality = s;
 talents = 0;
}
```

```
Golfer lulu("Little Lulu");
Golfer roy("Roy Hobbs", 12);
Golfer * par = new Golfer;
Golfer next = lulu;
Golfer hazzard = "Weed Thwacker";
*par = nancy;
nancy = "Nancy Putter";
```

- b. Clearly, the class requires several methods. If a particular method does it require to pass

## Programming Exercises

1. Consider the following class declaration

```
class Cow {
 char name[20];
 char * hobby;
 double weight;
public:
 Cow();
 Cow(const char * nm, const char * h);
 Cow(const Cow c&);
```

```

// pe12_2.cpp
#include <iostream>
using namespace std;
#include "string2.h"
int main()
{
 String s1(" and I am a C++ student");
 String s2 = "Please enter your name: ";
 String s3;
 cout << s2;
 cin >> s3;
 s2 = "My name is " + s3;
 cout << s2 << ".\n";
 s2 = s2 + s1;
 s2.stringup();
 cout << "The string\n" << s2 << "\n"
 << " 'A' characters in it.\n";
 s1 = "red";
 String s4(s1);
 // String(const String&)
 // then String s4 = s1;
 String rgb[3] = { String(s1), String(s2), String(s4) };
 cout << "Enter the name of a primary color: ";
 String ans;
 bool success = false;
 while (cin >> ans)

```

My name is Fretta Farbo.

The string

MY NAME IS FRETТА FARBO AND I AM A C

contains 6 'A' characters in it.

Enter the name of a primary color fo

Try again!

**BLUE**

That's right!

Bye

3. Rewrite the `Stock` class, as described in 10.9, so that it uses dynamically allocated memory objects to hold the stock names. Also re-implement the overloaded operator `<<()` definition. Test your code.

4. Consider the following variation of the

```
// stack.h -- class declaration for a stack of items
typedef unsigned long Item;
```

```
class Stack
{
```

5. The Bank of Heather has performed a simulation and found that a customer will wait more than one minute in line. Use the simulation to find the value for number of customers per hour that will result in a maximum wait time of one minute. (Use at least a 100-hour trial period.)
6. The Bank of Heather would like to know the probability that a customer will join the first queue if it is longer than the second queue and that the customer will join the second queue if it is longer than the first queue. Use the simulation to find the value for number of customers per hour that will result in a maximum wait time of one minute. (Note: This is a nonlinear problem. Doubling the number of customers doesn't double the number of customers per hour that will result in a maximum wait time of one minute.)





- Virtual member functions
- Early (static) binding and late (dynamic)
- Abstract base classes
- Pure virtual functions
- When and how to use public inheritance

One of the main goals of object-oriented programming is reuse. When you develop a new project, particularly a large one, you reuse proven code rather than to reinvent it. Code that has already been used and tested, can help you develop your program. Also the less you have to concern yourself with details, the more you concentrate on overall program strategy.

Traditional C function libraries provide reusable functions, such as `strlen()` and `rand()`, that programmers use. Vendors furnish specialized C libraries that provide specialized library. For example, you can purchase libraries for screen control functions. However, function libraries do not supply the source code for its library functions. You cannot modify the functions to meet your particular needs or your program to meet the workings of the library. Even if you do, you run the risk of unintentionally modifying the relationships among library functions as you

Of course, you could accomplish the same thing by creating, testing, and modifying it, but the inheritance mechanism makes it easier to add the new features. You don't even need access to the source code. If you purchase a class library that provides only a set of abstract class methods, you can still derive new classes from it. If you can distribute your own classes to others, keep them private, or even still giving your clients the option of adding features to the existing ones.

Inheritance is a splendid concept, and its proper use is a key to managing inheritance so that it works properly. This chapter looks at both the simple and the complex.

## Beginning with a Simple

When one class inherits from another, the original class is called a *base class*. The inheriting class is called a *derived class*. So to illustrate inheritance, let's consider a simple class. The Webtown Social Club has decided to play table tennis. As head programmer for the club, you decide to create a `TableTennisPlayer` class defined in Listings 13.1 and 13.2.

### Listing 13.1 `tabtenn0.h`

---

```
// tabtenn0.h -- a table-tennis base class
#ifndef TABTENNO_H_
#define TABTENNO_H_
#include <string>
```

```

TableTennisPlayer::TableTennisPlayer (const
 const string & ln, bool ht) : firstna
 lastname(ln), hasTable(ht) {}

void TableTennisPlayer::Name() const
{
 std::cout << lastname << ", " << first
}

```

---

All the `TableTennisPlayer` class does is let them have tables. There are a couple of points to the `string` class to hold the names. This is more of a character array. And it is rather more professional. “Classes and Dynamic Memory Allocation.” The initializer list syntax introduced in Chapter 1.

```

TableTennisPlayer::TableTennisPlayer (const
 const string & ln, bool ht) : firstna
{
 firstname = fn;
 lastname = ln;
 hasTable = ht;
}

```

```

 if (player2.HasTable())
 cout << ": has a table";
 else
 cout << ": hasn't a table.\n";

 return 0;
 }

```

---

And here's the output of the program in L

Blizzard, Chuck: has a table.

Boomdea, Tara: hasn't a table.

Note that the program uses constructors w

```
TableTennisPlayer player1("Chuck", "Blizza
```

```
TableTennisPlayer player2("Tara", "Boomdea
```

But the formal parameters for the construc

This is a type mismatch, but the `string` class,

has a constructor with a `const char *` param

cally to create a `string` object initialized by t

either a `string` object or a C-style string as a

constructor. The first invokes a `string` constr

and the second invokes a `string` constructor

- An object of the derived type has store of the derived type. (The derived class inherits the base-class interface.)
- An object of the derived type can use the methods of the base class. (The derived class inherits the base-class interface.)

Thus, a `RatedPlayer` object can store the player's rating and whether the player has a table. Also a `RatedPlayer` object has `HasTable()`, and `ResetTable()` methods from `TableTennisPlayer` (see Figure 13.1 for another example).

What needs to be added to these inherited methods?

- A derived class needs its own constructor.
- A derived class can add additional data members.

In this particular case, the class needs one additional data member: the rating. It should also have a method for retrieving the rating. So the class declaration could look like this:

```
// simple derived class
class RatedPlayer : public TableTennisPlayer
{
private:
 unsigned int rating; // add a data member
```

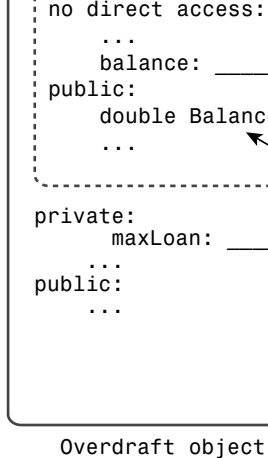


Figure 13.1 Base-class a

The constructors have to provide data for  
ited members. The first `RatedPlayer` constru  
each member, and the second `RatedPlayer` c  
parameter, which bundles three items (`first`  
gle unit.

example, a program has the following declaration:

```
RatedPlayer rplayer1(1140, "Mallory", "Du
```

The `RatedPlayer` constructor assigns the value `true` to the formal parameters `fn`, `ln`, and `ht` and passes the arguments to the `TableTennisPlayer` constructor. The `RatedPlayer` constructor creates an embedded `TableTennisPlayer` object and stores a pointer to it in `tblt`. Then the program enters the body of the constructor, constructs the `RatedPlayer` object, and assigns the value `1140` to the `rating` member (see Figure 13.2).

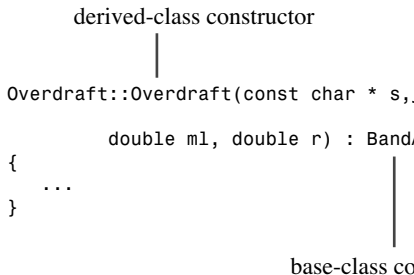


Figure 13.2 Passing arguments to a base-class constructor



```
 rating = r;
 }
```

Again, the `TableTennisPlayer` information constructor:

```
TableTennisPlayer(tp)
```

Because `tp` is type `const TableTennisPlayer`, the constructor. The base class didn't define a copy constructor, so the compiler automatically generates a copy constructor. If you haven't defined one already. In this case, the implicitly generated copy constructor, is fine because the class does not have any non-static data members. (The `string` members do use dynamic memory, but the compiler will use the `string` class copy constructor.)

You may, if you like, also use member initialization in the constructor. In this case, you use the member name in the constructor. The second constructor can also be written in this way:

```
// alternative version
RatedPlayer::RatedPlayer(unsigned int r, const TableTennisPlayer &tp)
 : TableTennisPlayer(tp), rating(r)
{
 }
}
```

when an object of a derived class expires, the destructor and then calls the base-class destructor.

## Member Initializer Lists

A constructor for a derived class can use the `base` constructor to call a base-class constructor. Consider this example:

```
derived::derived(type1 x, type2 y) : base(x)
{
 ...
}
```

Here, `derived` is the derived class, `base` is the base class, and `x` is passed to the base-class constructor. If, say, the derived class has two data members, `10` and `12`, this mechanism then passes `10` and `12` to the base-class constructor. (see Chapter 14, “Reusing Code in C++”), and the base class. However, that class can use the `base` constructor to call its immediate base class, and so on. If you use a member initializer list, the program uses the `base` constructor. The member initializer list can be used *only* in constructors.

```

 void Name() const;
 bool HasTable() const { return hasTable; }
 void ResetTable(bool v) { hasTable = v; }
};

// simple derived class
class RatedPlayer : public TableTennisPlayer
{
private:
 unsigned int rating;
public:
 RatedPlayer(unsigned int r = 0, const TableTennisPlayer & ln = "none") : TableTennisPlayer(ln)
 { rating = r; }
 RatedPlayer(unsigned int r, const TableTennisPlayer & ln) : TableTennisPlayer(ln)
 { rating = r; }
 unsigned int Rating() const { return rating; }
 void ResetRating(unsigned int r) { rating = r; }
};

#endif

```

---

Listing 13.5 provides the method definitions for the `RatedPlayer` class. While it is possible to separate files, but it's simpler to keep the definitions in the same file as the class declaration.

```
{
}
```

---

Listing 13.6 creates objects of both the `TableTennisPlayer` and `RatedPlayer` classes. Notice that objects of both classes can use the `HasTable()` methods.

### Listing 13.6 `usett1.cpp`

---

```
// usett1.cpp -- using base class and derived class
#include <iostream>
#include "tabtenn1.h"

int main (void)
{
 using std::cout;
 using std::endl;
 TableTennisPlayer player1("Tara", "Borwick");
 RatedPlayer rplayer1(1140, "Mallory", "Mallory");
 rplayer1.Name(); // derived class
 if (rplayer1.HasTable())
 cout << "Mallory has a table.\n";
 else
```

## Special Relationships Between Derived and Base Classes

A derived class has some special relationships with its base class. One, as we have just seen, is that a derived-class object can use base-class methods. Another, which we have not seen, is that base-class methods can use derived-class objects. These relationships are not private:

```
RatedPlayer rplayer1(1140, "Mallory", "Duck");
rplayer1.Name(); // derived object uses base-class method
```

Two other important relationships are that a base-class object can use a derived-class object without an explicit type cast and that a derived-class object can use a base-class object without an explicit type cast.

```
RatedPlayer rplayer1(1140, "Mallory", "Duck");
TableTennisPlayer & rt = rplayer1;
TableTennisPlayer * pt = &rplayer1;
rt.Name(); // invoke Name() with reference
pt->Name(); // invoke Name() with pointer
```

However, a base-class pointer or reference couldn't use `rt` or `pt` to invoke, say, the derived-class method `Play`:

```

{
 using std::cout;
 cout << "Name: ";
 rt.Name();
 cout << "\nTable: ";
 if (rt.HasTable())
 cout << "yes\n";
 else
 cout << "no\n";
}

```

The formal parameter `rt` is a reference to an object or to a derived-class object. Thus, you can pass a `TableTennisPlayer` argument or a `RatedPlayer` argument:

```

TableTennisPlayer player1("Tara", "Boomde", "Duchess");
RatedPlayer rplayer1(1140, "Mallory", "Duchess");
Show(player1); // works with TableTennisPlayer
Show(rplayer1); // works with RatedPlayer

```

A similar relationship would hold for a function parameter; it could be used with either the actual argument or a derived-class object as an actual argument:

TableTennisPlayer object embedded in the

Similarly, you can assign a derived-class obj

```
RatedPlayer olaf1(1840, "Olaf", "Loaf", tr
TableTennisPlayer winner;
winner = olaf1; // assign derived to base
```

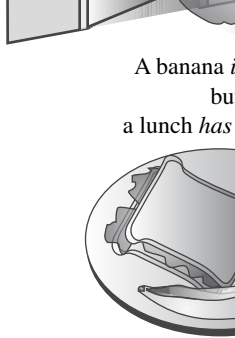
In this case, the program uses the implicit c

```
TableTennisPlayer & operator=(const TableT
```

Again, a base-class reference refers to a der  
tion of `olaf1` is copied to `winner`.

## Inheritance: An *Is-a* Relat

The special relationship between a derived cla  
model for C++ inheritance. Actually, C++ ha  
tected, and private. Public inheritance is the m  
relationship. This is shorthand for saying that a  
object of the base class. Anything you do with  
with a derived-class object. Suppose, for exam  
store, say, the weight and caloric content of a f  
fruit, you could derive a `Banana` class from the  
the data members of the original class, so a `Ba`  
ing the weight and caloric content of a banan



A banana *is* a  
but  
a lunch *has*

Figure 13.3 Is-a and

Public inheritance doesn't model an *is-like* relationship. It's often pointed out that lawyers are like sharks. For example, sharks can live underwater. A class derived from a `Shark` class. Inheritance can add properties from a base class. In some cases, sharing a class encompassing those characteristics is a *has-a* relationship, to define the related classes.



you can have multiple behaviors for a method  
key mechanisms for implementing polymorphism

- Redefining base-class methods in a derived class
- Using virtual methods

It's time for another example. You have been promoted at Social Club to become head programmer for the bank. The bank asks you to do is develop two classes: the first class is the Savings Account, the plan, the Brass Account, and the second class is the Overdraft Account, which adds an overdraft protection feature to the Savings Account (but not too much larger) than his or her balance. The bank also wants the user for the excess payment and adding a new feature to the accounts in terms of data to be stored and operations to be performed.

First, here is the information for a Brass Account:

- Client name
- Account number
- Current balance

And here are the operations to be represented by the Brass Account:

- Creating an account
- Depositing money into the account

BrassPlus class meet the *is-a* test? Sure. Even though `isA` is true for a BrassPlus object. Both store a client's name. With both, you can make deposits and withdrawals. But, since that the *is-a* relationship is not, in general, symmetric, similarly, a Brass object won't have all the capabilities of

## Developing the Brass and BrassPlus classes

The Brass Account class information is pretty straightforward. You get enough details about how the overdraft facility works. For further information, the friendly Pontoon bank website is following:

- A Brass Plus Account limits how much a customer can withdraw in drafts. The default value is \$500, but some customers may have a higher limit.
- The bank may change a customer's overdraft limit.
- A Brass Plus Account charges interest on overdrafts. Some customers may start with a different interest rate.
- The bank may change a customer's interest rate.
- The account keeps track of how much the customer owes (principal plus interest). The user cannot pay off the account without closing it.

```

 private:
 std::string fullName;
 long acctNum;
 double balance;
 public:
 Brass(const std::string & s = "Nullbod
 double bal = 0.0);
 void Deposit(double amt);
 virtual void Withdraw(double amt);
 double Balance() const;
 virtual void ViewAcct() const;
 virtual ~Brass() {}
};

```

```

//Brass Plus Account Class
class BrassPlus : public Brass
{
 private:
 double maxLoan;
 double rate;
 double owesBank;
 public:
 BrassPlus(const std::string & s = "Nul
 double bal = 0.0, double ml =

```

- The Brass class also declares a virtual `ViewAcct()` method, but it does nothing.

The first point in the list is nothing new. The `Brass` class was already there when it added a new data member and two new methods.

The second point in the list is how the `ViewAcct()` method is used differently for the derived class. The two `ViewAcct()` methods must be two separate method definitions. The qualified name `Brass::ViewAcct()` and the qualified name `BrassPlus::ViewAcct()`. A program will use the `ViewAcct()` method definition to use:

```
Brass dom("Dominic Banker", 11224, 4183.41);
BrassPlus dot("Dorothy Banker", 12118, 2500.00);
dom.ViewAcct(); // use Brass::ViewAcct()
dot.ViewAcct(); // use BrassPlus::ViewAcct()
```

Similarly, there will be two versions of `Withdraw()` and one that's used by `BrassPlus` objects. Methods such as `Deposit()` and `Balance()`, are declared in the `Brass` class.

The third point (the use of `virtual`) is more important. It determines which method is used if the method is called by a pointer instead of by an object. If you don't use the `virtual` keyword, the method is called based on the reference type or pointer type.

In this case, both references are type `Brass`. `BrassPlus::ViewAcct()` is used for it. Using `ViewAcct()` results in similar behavior.

It turns out, as you'll see in a bit, that this behavior is not unique to `ViewAcct()`. Therefore, it is the common practice to declare `virtual` functions that might be redefined in a derived class. When a function is declared `virtual` in a base class, it is automatically `virtual` in the derived class, and you can declare which functions are `virtual` by using the keyword `virtual` in the declarations, too.

The fourth point is that the base class declares `virtual` functions so that the correct sequence of destructors is called. We will discuss this point in more detail later in this chapter.

## Note

If you redefine a base-class method in a derived class, you must declare the base-class method as `virtual`. This makes the object type instead of the type of a reference. You must also declare a `virtual` destructor for the base class.

```
Brass::Brass(const string & s, long an, d
{
 fullName = s;
 acctNum = an;
 balance = bal;
}
```

```
void Brass::Deposit(double amt)
{
 if (amt < 0)
 cout << "Negative deposit not all
 << "deposit is cancelled.\n"
 else
 balance += amt;
}
```

```
void Brass::Withdraw(double amt)
{
 // set up ###.## format
 format initialState = setFormat();
 precis prec = cout.precision(2);

 if (amt < 0)
 cout << "Withdrawal amount must b
```

```

 }

// BrassPlus Methods
BrassPlus::BrassPlus(const string & s, long
 double ml, double r) : Brass(s,
{
 maxLoan = ml;
 owesBank = 0.0;
 rate = r;
}

BrassPlus::BrassPlus(const Brass & ba, double
 : Brass(ba) // uses implicit
{
 maxLoan = ml;
 owesBank = 0.0;
 rate = r;
}

// redefine how ViewAcct() works
void BrassPlus::ViewAcct() const
{
 // set up ###.## format
 format initialState = setFormat();

```

```

 owesBank += advance * (1.0 + rate);
 cout << "Bank advance: $" << advance;
 cout << "Finance charge: $" << advance * rate;
 Deposit(advance);
 Brass::Withdraw(amt);
 }
 else
 cout << "Credit limit exceeded. Try again." << endl;
 restore(initialState, prec);
}

format setFormat()
{
 // set up ###.## format
 return cout.setf(std::ios_base::fixed)
 .setf(std::ios_base::floatfield);
}

void restore(format f, precis p)
{
 cout.setf(f, std::ios_base::floatfield);
 cout.precision(p);
}

```

---



```

 owesBank = 0.15;
 rate = r;
 }

```

Each of these constructors uses the member information to a base-class constructor and then adds new data items added by the `BrassPlus` class.

Non-constructors can't use the member information. A non-constructive method can call a public base-class method. For example, the core of the `BrassPlus` version of `ViewAcct` is

```

// redefine how ViewAcct() works
void BrassPlus::ViewAcct() const
{
 ...
 Brass::ViewAcct(); // display base p
 cout << "Maximum loan: $" << maxLoan << endl;
 cout << "Owed to bank: $" << owesBank << endl;
 cout << "Loan Rate: " << 100 * rate << endl;
 ...
}

```

In other words, `BrassPlus::ViewAcct()` calls on the base-class method `Brass::ViewAcct()` to display the base-class members. Using the scope-resolution operator to call a base-class method is a standard technique.

```

{
...
 double bal = Balance();
 if (amt <= bal)
 Brass::Withdraw(amt);
 else if (amt <= bal + maxLoan - owesBank)
 {
 double advance = amt - bal;
 owesBank += advance * (1.0 + rate);
 cout << "Bank advance: $" << advance << endl;
 cout << "Finance charge: $" << advance * rate << endl;
 Deposit(advance);
 Brass::Withdraw(amt);
 }
 else
 cout << "Credit limit exceeded. Transaction not allowed." << endl;
...
}

```

Note that the method uses the base-class method `Balance()` to get the current balance. The code doesn't have to use the `GetBalance()` method because this method has not been redefined in the derived class.

The `ViewAcct()` and the `Withdraw()` methods use the `cout` and `endl` formatting methods to set the output mode for the derived class.

```
 cout.setf(f, std::ios_base::floatfield);
 cout.precision(p);
 }
```

You can find more details about formatting

## Using the Brass and BrassPlus Class

Listing 13.9 shows the class definitions with a

### Listing 13.9 `usebrass1.cpp`

---

```
// usebrass1.cpp -- testing bank account c
// compile with brass.cpp
#include <iostream>
#include "brass.h"

int main()
{
 using std::cout;
 using std::endl;

 Brass Piggy("Porcelot Pigg", 381299, 4
 BrassPlus Hoggy("Horatio Hogg", 382288
 Piggy.ViewAcct();
 cout << endl;
```

Account Number: 382288  
Balance: \$3000.00  
Maximum loan: \$500.00  
Owed to bank: \$0.00  
Loan Rate: 11.125%

Depositing \$1000 into the Hogg Account:  
New balance: \$4000  
Withdrawing \$4200 from the Pigg Account:  
Withdrawal amount of \$4200.00 exceeds your  
Withdrawal canceled.  
Pigg account balance: \$4000  
Withdrawing \$4200 from the Hogg Account:  
Bank advance: \$200.00  
Finance charge: \$22.25  
Client: Horatio Hogg  
Account Number: 382288  
Balance: \$0.00  
Maximum loan: \$500.00  
Owed to bank: \$222.25  
Loan Rate: 11.125%

```
int main()
{
 using std::cin;
 using std::cout;
 using std::endl;

 Brass * p_clients[CLIENTS];
 std::string temp;
 long tempnum;
 double tempbal;
 char kind;

 for (int i = 0; i < CLIENTS; i++)
 {
 cout << "Enter client's name: ";
 getline(cin,temp);
 cout << "Enter client's account num";
 cin >> tempnum;
 cout << "Enter opening balance: $";
 cin >> tempbal;
 cout << "Enter 1 for Brass Account";
 cout << "2 for BrassPlus Account: ";
 while (cin >> kind && (kind != '1'
 && kind != '2'))
 cout << "Enter either 1 or 2: ";
 }
}
```

```
 for (int i = 0; i < CLIENTS; i++)
 {
 delete p_clients[i]; // free memo
 }
 cout << "Done.\n";
 return 0;
 }
```

---

The program in Listing 13.10 lets user input new clients added and then uses `new` to create and initialize. Recall that `getline(cin, temp)` reads a line of input into object `temp`.

Here is a sample run of the program in Listing 13.10:

```
Enter client's name: Harry Fishsong
Enter client's account number: 112233
Enter opening balance: $1500
Enter 1 for Brass Account or 2 for BrassP
Enter client's name: Dinah Otternoe
Enter client's account number: 121213
Enter opening balance: $1800
Enter 1 for Brass Account or 2 for BrassP
Enter the overdraft limit: $350
Enter the interest rate as a decimal frac
```

Account Number: 212118  
Balance: \$5200.00  
Maximum loan: \$800.00  
Owed to bank: \$0.00  
Loan Rate: 10.00%

Client: Tim Turtletop  
Account Number: 233255  
Balance: \$688.00

Done.

The polymorphic aspect is provided by the

```
for (i = 0; i < CLIENTS; i++)
{
 p_clients[i]->ViewAcct();
 cout << endl;
}
```

If the array member points to a Brass object, the  
array member points to a BrassPlus object, Brass::ViewAcct() were been declared as virtual  
in all cases.

with C++, the task is more complex because to look at the function arguments as well as the function to use. Nonetheless, this kind of binding is determined during the compiling process; binding that occurs at compile time is called *static binding* (or *early binding*). However, virtual methods are resolved at run time because the compiler doesn't know which method to choose to make. Therefore, the compiler has no way to determine which virtual method to be selected as the program runs (this is called *dynamic binding*).

Now that you've seen virtual methods at a little more depth, beginning with how C++ handles pointers,

## Pointer and Reference Type Comp

Dynamic binding in C++ is associated with polymorphism, and this is governed, in part, by the inheritance mechanism. The *is-a* relationship is in how it handles polymorphism. C++ does not allow you to assign an address to a reference. Nor does it let a reference to one type refer to another.

```
double x = 2.5;
int * pi = &x; // invalid assignment, mismatched types
long & rl = x; // invalid assignment, mismatched types
```



these data members wouldn't apply to the base Singer class from an Employee class, adding a range and a member function, called range(). It wouldn't make sense to apply the range() method to the address of an Employee object and use implicit downcasting were allowed, you could (see Figure 13.4).

Upcasting also takes place for function call parameters. Consider the following code fragment the virtual method ViewAcct():

```
void fr(Brass & rb); // uses rb.ViewAcct()
void fp(Brass * pb); // uses pb->ViewAcct()
void fv(Brass b); // uses b.ViewAcct()

int main()
{
 Brass b("Billy Bee", 123432, 10000.0);
 BrassPlus bp("Betty Beep", 232313, 123
 fr(b); // uses Brass::ViewAcct()
 fr(bp); // uses BrassPlus::ViewAcct()
 fp(b); // uses Brass::ViewAcct()
 fp(bp); // uses BrassPlus::ViewAcct()
```

>ViewAcct () goes by the pointer type (Brass) because the pointer type is known at compile time, so the compiler calls Brass::ViewAcct () at compile time. In short, static binding for virtual methods.

But if ViewAcct () is declared as virtual in Brass, and the object type (BrassPlus) and invokes BrassPlus::ViewAcct () to see that the object type is BrassPlus, but, in this case, the type might only be determined when the program runs. The compiler generates code that binds ViewAcct () to BrassPlus::ViewAcct () depending on the object type, while the program runs. This is dynamic binding for virtual methods.

In most cases, dynamic binding is a good thing. It allows the method designed for a particular type. Given the following:

- Why have two kinds of binding?
- If dynamic binding is so good, why isn't it used everywhere?
- How does it work?

We'll look at answers to these questions next.

ods you expect to be redefined.

### Tip

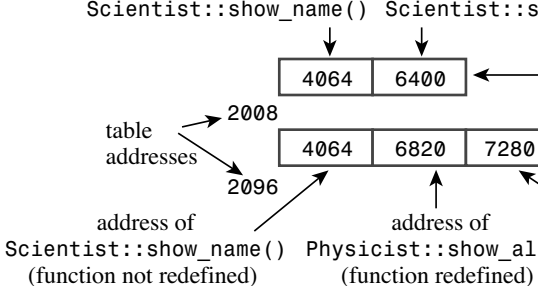
If a method in a base class will be redefined, the method should not be redefined, you should

Of course, when you design a class, it's not a method falls. Like many aspects of real life, cla

## How Virtual Functions Work

C++ specifies how virtual functions should be the compiler writer. You don't need to know virtual functions, but seeing how it is done may help take a look.

The usual way compilers handle virtual functions object. The hidden member holds a pointer to an array is usually termed a *virtual function table* (*virtual functions declared for objects of that class* contains a pointer to a table of addresses of all the a derived class contains a pointer to a separate provides a new definition of a virtual function, the



|     |             |             |
|-----|-------------|-------------|
| ... | Sophie Fant | <b>2008</b> |
|-----|-------------|-------------|

...                      name                      vptr

|     |              |             |    |
|-----|--------------|-------------|----|
| ... | Adam Crusher | <b>2096</b> | nu |
|-----|--------------|-------------|----|

...                      name                      vptr

```
Physicist adam("Adam Crusher", "nuclear
Scientist * psc = &adam;
psc->show_all());
```

1. Find va
2. Go to t
3. Find ad
4. Go to th

Figure 13.5    A virtual

- If a virtual method is invoked by using a pointer to an object, the program uses the method defined for the reference *late, binding*. This behavior is important for a pointer or reference to refer to an object.
- If you're defining a class that will be used as a base class, declare as virtual functions the class methods that you want derived classes.

There are several other things you may need to know about virtual functions which have been mentioned in passing already.

## Constructors

Constructors can't be virtual. Creating a derived-class constructor, not a base-class constructor. The derived-class constructor, but this sequence is distinct from the base-class constructor, so it doesn't inherit the base-class constructors, so it can't be virtual, anyway.

## Destructors

Destructors should be virtual unless a class isn't a base class. Suppose `Employee` is a base class and `Singer` is a derived class.

Normally, you should provide a base class with a destructor. If you don't, you need a destructor.

## Friends

Friends can't be virtual functions because friends can be virtual functions. If this poses a problem, you can sidestep it by having the friend function use a virtual function.

## No Redefinition

If a derived class fails to redefine a function (or a virtual version of the function). If a derived class is partially defined, the most recently defined version of the function is hidden, as described next.

## Redefinition Hides Methods

Suppose you create something like the following:

```
class Dwelling
{
public:
 virtual void showperks(int a) const;
 ...
};
class Hovel : public Dwelling
```

new exception to this rule is that a return type class can be replaced by a reference or pointer *covariance of return type* because the return type is a class type:

```
class Dwelling
{
public:
 // a base method
 virtual Dwelling & build(int n);
 ...
};
class Hovel : public Dwelling
{
public:
 // a derived method with a covariant return type
 virtual Hovel & build(int n); // same as build in Dwelling
 ...
};
```

Note that this exception applies only to references and pointers.  
Second, if the base class declaration is overridden by a different  
class versions in the derived class:

## Access Control: protected

So far the class examples in this book have used `private` to control access to class members. There is one more access keyword, `protected`. The `protected` keyword is used to restrict access to class members in a protected section of a class. The difference between `private` and `protected` is that `protected` members can be accessed from the base class. Members of a derived class can access `protected` members of the base class directly, but they cannot directly access `private` members. `protected` members in the protected category behave like `private` members in the `private` category but behave like public members in the `public` category.

For example, suppose the `Brass` class declares

```
class Brass
{
protected:
 double balance;
 ...
};
```

In this case, the `BrassPlus` class could access `balance` directly. For example, the core of `BrassPlus` could be



you could write code like this:

```
void BrassPlus::Reset(double amt)
{
 balance = amt;
}
```

The `Brass` class was designed so that the `Reset` method provides the only means for altering `balance`. But `balance` is a public variable as far as `BrassPlus` is concerned. For example, the safeguards found in `Withdraw()`.

### Caution

You should prefer private to protected access control. You should use base-class methods to provide de-

However, protected access control can be useful to restrict derived classes access to internal functions that

## Abstract Base Classes

So far you've seen simple inheritance and the `virtual` keyword. The next step in increasing sophistication is the abstract base class. In some programming situations that provide the

```

public:
 ...
 void Move(int nx, ny) { x = nx; y = ny; }
 virtual double Area() const { return 3.14 * x * y; }
 virtual void Rotate(double nang) { ang = nang; }
 virtual void Scale(double sa, double sb) { x = x * sa; y = y * sb; }
 ...
};

```

Now suppose you derive a `Circle` class from `Ellipse`:

```

class Circle : public Ellipse
{
 ...
};

```

Although a circle is an ellipse, this derivation is not ideal because a circle has only a single value, its radius, to describe its size, not two values for its major axis (a) and semiminor axis (b). The `Circle` class could be defined to set the same value to both the `a` and `b` members, thus eliminating the redundancy of the same information. The `angle` parameter and `Rotate()` method make sense for a circle, and the `Scale()` method makes sense for a non-circle by scaling the two axes differently. Finally, we can put a redefined `Rotate()` method in the `Circle` class:

Circle and Ellipse classes from the ABC. The Ellipse class uses base-class pointers to manage a mixture of Ellipse and Circle objects; it can use a polymorphic approach. In this case, what would be the Move() method? The coordinates of the center of the shape; a Move() method, which works differently for an Ellipse and a Circle. The Move() method can't even be implemented for the ABC class. The ABC class has no data members. C++ has a way to provide an *abstract function*. A pure virtual function has = 0 as a return type. The Move() method:

```
class BaseEllipse // abstract base class
{
private:
 double x; // x-coordinate of center
 double y; // y-coordinate of center
 ...
public:
 BaseEllipse(double x0 = 0, double y0 = 0) {}
 virtual ~BaseEllipse() {}
 void Move(int nx, ny) { x = nx; y = ny; }
 virtual double Area() const = 0; // a pure virtual function
 ...
}
```

A program using these classes would be able to create `BaseCircle` objects but no `BaseEllipse` objects. Because `BaseCircle` is the base class, a collection of such objects can be represented by pointers. Classes such as `Circle` and `Ellipse` inherit from `BaseCircle` indicate that you can create objects of those classes.

In short, an ABC describes an interface that all classes derived from an ABC use regular virtual functions. In terms of the properties of the particular derived class, the

## Applying the ABC Concept

You'd probably like to see a complete example of an ABC representing the `Brass` and `BrassPlus` accounts. This class should contain all methods and data for both `Brass` and the `BrassPlus` classes. The methods for `BrassPlus` class than they do for the `Brass` class. At least one virtual function should be a pure virtual function. The `AcctABC` class abstract.

Listing 13.11 is a header file that declares the `Brass` and `BrassPlus` classes (both concrete classes) and the `AcctABC` class. The `AcctABC` provides some protected methods that derived-class methods can call but that are not available to derived-class objects. `AcctABC` also provides a

```

 };

 const std::string & FullName() const;
 long AcctNum() const {return acctNum;}
 Formatting SetFormat() const;
 void Restore(Formatting & f) const;
public:
 AcctABC(const std::string & s = "Null",
 double bal = 0.0);
 void Deposit(double amt) ;
 virtual void Withdraw(double amt) =
 double Balance() const {return bal;}
 virtual void ViewAcct() const = 0;
 virtual ~AcctABC() {}
};

```

```

// Brass Account Class
class Brass :public AcctABC
{
public:
 Brass(const std::string & s = "Null",
 double bal = 0.0) : AcctABC(s,
 virtual void Withdraw(double amt);
 virtual void ViewAcct() const;
 virtual ~Brass() {}
};

```

### Listing 13.12    **acctabc.cpp**

---

```
// acctabc.cpp -- bank account class meth
#include <iostream>
#include "acctabc.h"
using std::cout;
using std::ios_base;
using std::endl;
using std::string;

// Abstract Base Class
AcctABC::AcctABC(const string & s, long a
{
 fullName = s;
 acctNum = an;
 balance = bal;
}

void AcctABC::Deposit(double amt)
{
 if (amt < 0)
 cout << "Negative deposit not all
 << "deposit is cancelled.\n"
```

```

 cout.precision(f.pr);
 }

// Brass methods
void Brass::Withdraw(double amt)
{
 if (amt < 0)
 cout << "Withdrawal amount must be positive\n"
 << "withdrawal canceled.\n";
 else if (amt <= Balance())
 AcctABC::Withdraw(amt);
 else
 cout << "Withdrawal amount of $" << amt
 << " exceeds your balance.\n"
 << "Withdrawal canceled.\n";
}

void Brass::ViewAcct() const
{
 Formatting f = SetFormat();
 cout << "Brass Client: " << FullName() << "\n";
 cout << "Account Number: " << AcctNum() << "\n";
}

```

```

 Formatting f = SetFormat();

 cout << "BrassPlus Client: " << FullN
 cout << "Account Number: " << AcctNum
 cout << "Balance: $" << Balance() <<
 cout << "Maximum loan: $" << maxLoan
 cout << "Owed to bank: $" << owesBank
 cout.precision(3);
 cout << "Loan Rate: " << 100 * rate <
 Restore(f);
 }

```

```

void BrassPlus::Withdraw(double amt)
{
 Formatting f = SetFormat();

 double bal = Balance();
 if (amt <= bal)
 AcctABC::Withdraw(amt);
 else if (amt <= bal + maxLoan - owesBank
 {
 double advance = amt - bal;
 owesBank += advance * (1.0 + rate

```



structure to set and restore formats with just t

```
struct Formatting
{
 std::ios_base::fmtflags flag;
 std::streamsize pr;
};
...
Formatting f = SetFormat();
...
Restore(f);
```

It's a neater look.

A problem with the older version was that were standalone functions, so those function n functions of the same name. There are several declare both functions `static`, making them processor, `acctabc.cpp`. A second is to place bot definition into a namespace. But one of the to this example places the structure definition an class definition. This makes them available to t hiding them from the outside world.

```
const int CLIENTS = 4;
```

```
int main()
```

```
{
```

```
 using std::cin;
```

```
 using std::cout;
```

```
 using std::endl;
```

```
 AcctABC * p_clients[CLIENTS];
```

```
 std::string temp;
```

```
 long tempnum;
```

```
 double tempbal;
```

```
 char kind;
```

```
 for (int i = 0; i < CLIENTS; i++)
```

```
 {
```

```
 cout << "Enter client's name: "
```

```
 getline(cin,temp);
```

```
 cout << "Enter client's account
```

```
 cin >> tempnum;
```

```
 cout << "Enter opening balance:
```

```
 cin >> tempbal;
```

```
 cout << "Enter 1 for Brass Acco
```

```
 << "2 for BrassPlus Account
```

```
 }

 for (int i = 0; i < CLIENTS; i++)
 {
 delete p_clients[i]; // free memory
 }
 cout << "Done.\n";

 return 0;
}
```

---

The program itself behaves the same as the same input as for Listing 13.10, the output was

## **ABC Philosophy**

The ABC methodology is a much more systematic than the more ad hoc, spur-of-the-moment example. Before designing an ABC, you first needed to represent a programming problem. This school of thought holds that if you design an concrete classes should be those that never seen produce cleaner designs with fewer complications.

```

// Base Class Using DMA
class baseDMA
{
private:
 char * label;
 int rating;

public:
 baseDMA(const char * l = "null", int
 baseDMA(const baseDMA & rs);
 virtual ~baseDMA();
 baseDMA & operator=(const baseDMA & r
 ...
};

```

The declaration contains the special methods new: a destructor, a copy constructor, and an

Now suppose you derive a lackDMA class new or have other unusual design features tha

```

// derived class without DMA
class lacksDMA :public baseDMA
{
private:
 char color[40];

```

Essentially the same situation holds for assignment. A class automatically uses the base-class assignment operator. So it, too, is fine.

These properties of inherited objects also hold for pointer-to-object types. They manage themselves objects. For example, Chapter 10, “Object Management,” shows how to manage objects by using a string object to represent the company name. Our String example, which uses dynamic memory allocation, can create problems. The default Stock copy constructor to copy the company member of an object would use the string assignment operator on the object, and the Stock destructor (default or overloaded) would use the string destructor.

## Case 2: Derived Class Does Use new

Suppose that the derived class uses new:

```
// derived class with DMA
class hasDMA :public baseDMA
{
private:
 char * style; // use new in constructor
public:
 ...
};
```

```

 {
 label = new char[std::strlen(rs.label);
 std::strcpy(label, rs.label);
 rating = rs.rating;
 }

```

The `hasDMA` copy constructor only has access to the `baseDMA` copy constructor to handle the base class portion.

```

hasDMA::hasDMA(const hasDMA & hs)
 : baseDMA(hs)
{
 style = new char[std::strlen(hs.style);
 std::strcpy(style, hs.style);
}

```

The point to note is that the member initialization is done in the `baseDMA` constructor. There is no `baseDMA` copy constructor. There is no `baseDMA` copy constructor parameter, but none is needed. That's because the `baseDMA` reference parameter, and a base class pointer. The `baseDMA` copy constructor uses the `baseDMA` copy constructor to construct the `baseDMA` portion of the new object.

Next, consider assignment operators. The usual pattern:

```

 style = new char[std::strlen(hs.style)];
 std::strcpy(style, hs.style);
 return *this;
 }

```

The following statement might look a little

```
baseDMA::operator=(hs); // copy base port
```

But using function notation instead of operator notation. In effect, the statement means that

```
*this = hs; // use baseDMA::operator=()
```

But, of course, the compiler ignores comments. The compiler would use `baseDMA::operator=()` if the function notation gets the correct assignment.

In summary, when both the base class and the derived class have a constructor, a destructor, a copy constructor, the derived-class destructor, copy constructor, and assignment operator, the requirement is accomplished three different ways. For a constructor, it is accomplished by the member initialization list, or else the default constructor. For a destructor, it is accomplished by the destructor definition. For the assignment operator, it is accomplished by the explicit call of the base-class assignment operator.

```

baseDMA(const char * l = "null", int
baseDMA(const baseDMA & rs);
virtual ~baseDMA();
baseDMA & operator=(const baseDMA & r
friend std::ostream & operator<<(std:
cons
};

// derived class without DMA
// no destructor needed
// uses implicit copy constructor
// uses implicit assignment operator
class lacksDMA :public baseDMA
{
private:
 enum { COL_LEN = 40};
 char color[COL_LEN];
public:
 lacksDMA(const char * c = "blank", co
 int r = 0);
 lacksDMA(const char * c, const baseDM
 friend std::ostream & operator<<(std:
 cons
};

```



```
// dma.cpp --dma class methods
```

```
#include "dma.h"
```

```
#include <cstring>
```

```
// baseDMA methods
```

```
baseDMA::baseDMA(const char * l, int r)
```

```
{
```

```
 label = new char[std::strlen(l) + 1];
```

```
 std::strcpy(label, l);
```

```
 rating = r;
```

```
}
```

```
baseDMA::baseDMA(const baseDMA & rs)
```

```
{
```

```
 label = new char[std::strlen(rs.label) + 1];
```

```
 std::strcpy(label, rs.label);
```

```
 rating = rs.rating;
```

```
}
```

```
baseDMA::~baseDMA()
```

```
{
```

```
 delete [] label;
```

```
}
```

```

 color[39] = '\\0';
 }

 lacksDMA::lacksDMA(const char * c, const
 : baseDMA(rs)
 {
 std::strncpy(color, c, COL_LEN - 1);
 color[COL_LEN - 1] = '\\0';
 }

 std::ostream & operator<<(std::ostream &
 {
 os << (const baseDMA &) ls;
 os << "Color: " << ls.color << std::e
 return os;
 }

 // hasDMA methods
 hasDMA::hasDMA(const char * s, const char
 : baseDMA(l, r)
 {
 style = new char[std::strlen(s) + 1];
 std::strcpy(style, s);
 }

```

```

 baseDMA::operator=(hs); // copy base
 delete [] style; // prepare for
 style = new char[std::strlen(hs.style)+1];
 std::strcpy(style, hs.style);
 return *this;
 }

std::ostream & operator<<(std::ostream & os, const baseDMA & hs)
{
 os << (const baseDMA &) hs;
 os << "Style: " << hs.style << std::endl;
 return os;
}

```

---

The new feature to note in Listings 13.14 is the use of a friend to a base class. Consider, for example,

```

friend std::ostream & operator<<(std::ostream & os, const hasDMA & hs)
{
 os << "hasDMA: " << hs.label << " rating: " << hs.rating << std::endl;
 return os;
}

```

Being a friend to the `hasDMA` class gives the `operator<<` function access to the `label` and `rating` members. There's a problem: This function is not a friend of the `baseDMA` class. The solution is to make the `operator<<` function a friend of the `baseDMA` class as well. The solution is to add the following line to the `baseDMA` class definition:

```
 using std::cout;
 using std::endl;

 baseDMA shirt("Portabelly", 8);
 lacksDMA balloon("red", "Blimpo", 4);
 hasDMA map("Mercator", "Buffalo Keys");
 cout << "Displaying baseDMA object:\n";
 cout << shirt << endl;
 cout << "Displaying lacksDMA object:\n";
 cout << balloon << endl;
 cout << "Displaying hasDMA object:\n";
 cout << map << endl;
 lacksDMA balloon2(balloon);
 cout << "Result of lacksDMA copy:\n";
 cout << balloon2 << endl;
 hasDMA map2;
 map2 = map;
 cout << "Result of hasDMA assignment:\n";
 cout << map2 << endl;
 return 0;
 }
}
```

---

## Class Design Review

C++ can be applied to a wide variety of program class design to some paint-by-numbers routine often apply, and this is as good a time as any to earlier discussions.

## Member Functions That the Compiler

As first discussed in Chapter 12, the compiler member functions, termed *special member functions* these special member functions are particularly them now.

## Default Constructors

A default constructor is one that has no arguments have default arguments. If you don't default constructor for you. Its existence allow pose `Star` is a class. You need a default constructor

A class copy constructor is used in the following cases:

- When a new object is initialized to an existing object
- When an object is passed to a function by value
- When a function returns an object by value
- When the compiler generates a temporary object

If a program doesn't use a copy constructor, the compiler provides a prototype but not a function definition. A copy constructor that performs memberwise initialization of an object is initialized to the value of the corresponding member. If a member is itself a class object, then memberwise initialization is defined for that particular class.

In some cases, memberwise initialization is not sufficient. Objects initialized with `new` generally require that you delete them. Or a class example. Or a class may have a static variable. In these cases, you need to define your own copy constructor.

## **Assignment Operators**

A default assignment operator handles assignment of objects of a class. Don't confuse assignment with initialization. Assignment uses initialization, and if a statement alters the value of an object,

conversion function approach can lead to confusion.

Chapter 18, “Visiting with the New C++”

adds added by C++11: the move constructor and

## **Other Class Method Considerations**

There are several other points to keep in mind. This section lists some of them.

### **Constructor Considerations**

Constructors are different from other class methods. Whereas other methods are invoked by existing objects, constructors aren't inherited. Inheritance means a derived class inherits the methods of its base class. In the case of constructors, the object doesn't exist until the constructor is called.

### **Destructor Considerations**

You need to remember to define an explicit destructor by new in the class constructors and takes care of destroying a class object requires. If the class is a base class, you should provide a virtual destructor even if the class doesn't have any virtual methods.

```

...
public:
 explicit Star(const char *);
...
};
...
Star north;
north = "polaris"; // not allowed
north = Star("polaris"); // allowed

```

To convert from a class object to some other type (see Chapter 11, “Working with Classes”). A function with no arguments or declared return type is converted to. Despite having no declared return type, it has a conversion value. Here are some examples:

```

Star::Star double() {...} // conversion to double
Star::Star const char * () {...} // conversion to const char*

```

You should be judicious with such functions. In some sense. Also with some class designs, having conversion operators can lead to ambiguous code. For example, suppose you have a `Vector` type of Chapter 11, and suppose you have

```

Vector ius(6.0, 0.0);
Vector lux = ius + 20.2; // ambiguous

```



Some class methods return objects. You've probably seen methods that return objects directly whereas others return references to objects. If you return an object, but if it isn't necessary, you should use references as closely.

First, the only coding difference between returning an object and a reference is in the function prototype and header.

```
Star nova1(const Star &); // returns a reference to a Star object
Star & nova2(const Star &); // returns a reference to a Star object
```

Next, the reason you should return a reference to an object involves generating a temporary copy of the object made available to the calling program. Thus, returning an object involves calling a copy constructor to generate the copy, then calling a destructor to get rid of the copy. Returning a reference to an object directly is analogous to passing an object by reference. Temporary copies. Similarly, returning a reference to an object is analogous to passing an object by reference. Both the calling and the called function operate on the same object.

However, it's not always possible to return a reference to a temporary object created in the function. The reference becomes invalid when the function terminates and the object is destroyed. You should return an object in order to generate a copy of the object for the calling program.

method doesn't modify an argument:

```
Star::Star(const char * s) {...} // won't
```

You can use `const` to guarantee that a me

```
void Star::show() const {...} // won't ch
```

Here `const` means `const Star *` this, v

Normally, a function that returns a referen  
statement, which really means you can assign  
can use `const` to ensure that a reference or p  
data in an object:

```
const Stock & Stock::topval(const Stock &
{
 if (s.total_val > total_val)
 return s; // argument ob
 else
 return *this; // invoking ob
}
```

Here the method returns a reference either  
both declared `const`, the function is not allow  
returned reference also must be declared `con`

explicit type cast. Depending on the class declaration, the cast may or may not make sense. (You might

## **What's Not Inherited**

Constructors are not inherited. That is, creating a derived-class constructor. However, derived-class constructors use the member-initializer list syntax to call on base-class constructors. If the derived-class constructor is not defined, the compiler generates one for you. If the derived-class constructor is defined, the compiler generates one for you by using the member-initializer list syntax to call on the base-class constructor. In an inheritance chain, each class calls the constructor of its immediate base class. C++ does not support the inheritance of constructors. However, the default behavior is to inherit.

Destructors are not inherited either. However, when a program first calls the derived destructor and then the base class destructor, the compiler generates a destructor for the base class. In other words, if a class serves as a base class, its destructor is inherited.

Assignment operators are not inherited. The compiler generates the same function signature in a derived class as in the base class. If an assignment operator has a function signature that is not the same as the base class, the compiler generates a formal parameter that is the class type. Assignment operators are not inherited, which we'll look at next.

## Assignment Operator Considerations

If the compiler detects that a program assigns an object of a class, the compiler automatically supplies that class with an assignment operator. If a class uses memberwise assignment of this operator uses memberwise assignment, then the object being assigned the value of the corresponding object. If the object belongs to a derived class, the compiler uses the assignment operator to handle assignment for the base-class portion of the object. If a class explicitly provided an assignment operator for the base class, then the compiler will use that operator. Similarly, if a class contains a member that is an object of a class, the assignment operator for that class is used for that member.

As you've seen several times, you need to define an assignment operator for a class. Class constructors use `new` to initialize pointers. You need to define an assignment operator for the base part of derived objects. You need to define an assignment operator for a derived class *unless* it inherits from a base class that defines an assignment operator. For example, the `baseDMA` class defines an assignment operator. The `derivedDMA` class uses the implicit assignment operator generated by the compiler.

lated into a method that is invoked by the left

```
blips.operator=(snips);
```

Here the left-hand object is a Brass object. The function `Brass::operator=(const Brass &)` function can only refer to a derived-class object, such as `blips`, because it deals with base-class members, so the `maxLoan` member of `snips` is ignored in the assignment. In short, `snips` is a derived-class object, and only the base-class members are in scope.

What about the reverse? Can you assign a Brass object to a BrassPlus object? Look at this example:

```
Brass gp("Griff Hexbait", 21234, 1200);
BrassPlus temp;
temp = gp; // possible?
```

Here the assignment statement would be true if the `BrassPlus` class had a conversion constructor. Otherwise, the statement `temp.operator=(gp);` would be required.

The left-hand object is a `BrassPlus` object. The function `BrassPlus::operator=(const BrassPlus &)` function cannot automatically refer to a base-class object. However, `BrassPlus` is *also* a conversion constructor:

```
BrassPlus(const Brass &);
```

## Private Versus Protected Members

Remember that protected members act like public members when concerned, but they act like private members when it comes to access protected members of a base class directly. You can access via base-class member functions. Thus, making them public for security, whereas making them protected simply makes them a trap, in his book *The Design and Evolution of C++*, Stroustrup, data members than protected data members.

## Virtual Method Considerations

When you design a base class, you have to decide whether or not. If you want a derived class to be able to redefine a virtual method in the base class. This enables late, or dynamic, binding. To be redefined, you don't make it virtual. The method, but it should be interpreted as non-virtual.

Note that inappropriate code can circumvent the following two functions:

```
void show(const Brass & rba)
{
 rba.ViewAcct();
 cout << endl;
}
```

a derived object via a base-class pointer or reference. The derived-class destructor followed by the base-class destructor.

## Friend Considerations

Because a friend function is not actually a class member, a derived class might still want a friend to a derived class to be able to access a friend. To accomplish this is to type cast a derived-class pointer to a base-class pointer and to then use the type cast reference operator to access the friend.

```
ostream & operator<<(ostream & os, const HS & hs)
{
 // type cast to match operator<<(ostream & os, const baseDMA &) hs;
 os << (const baseDMA &) hs;
 os << "Style: " << hs.style << endl;
 return os;
}
```

You can also use the `dynamic_cast<>` operator to cast a derived-class pointer to a base-class pointer. See “Exceptions, and More,” for the type cast:

```
os << dynamic_cast<const baseDMA &> (hs);
```

For reasons discussed in Chapter 15, this works only if the base class is a polymorphic class.

*ence Manual*, summarizes these properties. In operators of the form +=, \*=, and so on. Note no different from those of the “other operators” is to point out that these operators don

Table 13.1    **Member Function Properties**

| Function    | Inherited | Member or Friend |
|-------------|-----------|------------------|
| Constructor | No        | Member           |
| Destructor  | No        | Member           |
| =           | No        | Member           |
| &           | Yes       | Either           |
| Conversion  | Yes       | Member           |
| ()          | Yes       | Member           |
| []          | Yes       | Member           |
| ->          | Yes       | Member           |
| op=         | Yes       | Either           |



public and protected base-class methods. You can use the methods to the class, and you can use the derived class methods. Each derived class requires its own constructor. When a program deletes an object, it first calls the derived class destructor.

If a class is meant to be a base class, you may want to use private members so that derived classes can use them. Using private members, in general, reduces the number of methods that a derived class can redefine a base-class method by declaring it with the keyword `virtual`. Virtual methods or references to be handled on the basis of the reference type or pointer type. In particular, the method must be virtual.

You might want to define an ABC that defines the implementation matters. For example, you could define a class which particular shape classes, such as `Circle`, which include at least one pure virtual method. You can use `0` before the closing semicolon of the declaration to indicate a pure virtual method.

```
virtual double area() const = 0;
```

8. Can you assign the address of an object to a pointer of a base class? Can you assign the address of an object of a derived class to a pointer of a base class?
9. Can you assign an object of a derived class to a pointer of a base class? Can you assign an object of a base class to a pointer of a derived class?
10. Suppose you define a function that takes a pointer to a base class as an argument. Why can this function also use a pointer to a derived class?
11. Suppose you define a function that takes a pointer to a base class as an argument. Why can the function pass a base-class object but not a derived-class object as an argument?
12. Why is it usually better to pass objects by reference than by value?
13. Suppose `Corporation` is a base class and `Employee` is a derived class. Suppose that each class defines a `head()` method. What does `ph.head()` mean if `ph` is of the `Corporation` type, and that `ph` is a pointer to an `Employee` object. How is `ph->head()` interpreted?
  - a. Regular nonvirtual method
  - b. Virtual method

```

// Base class
class Cd { // represents a CD disk
private:
 char performers[50];
 char label[20];
 int selections; // number of s
 double playtime; // playing tim
public:
 Cd(char * s1, char * s2, int n,
 Cd(const Cd & d);
 Cd();
 ~Cd();
 void Report() const; // reports
 Cd & operator=(const Cd & d);
};

```

Derive a classic class that adds an array identifying the primary work on the CD. If the methods be virtual, modify the base-class constructor. If the method is not needed, remove it from the program. The following program:

```

#include <iostream>
using namespace std;
#include "classic.h" // which wi

```

```

 copy = c2;
 copy.Report();

 return 0;
}

void Bravo(const Cd & disk)
{
 disk.Report();
}

```

2. Do Programming Exercise 1 but use dynamic size arrays for the various strings tracked.
3. Revise the `baseDMA-lacksDMA-hasDMA` derived from an `ABC`. Test the result with 13.10. That is, it should feature an array to make runtime decisions as to what to use. Add `View()` methods to the class definition.
4. The Benevolent Order of Programmers: describe it, the BOP Portmaster has de

Brand: Gallo  
Kind: tawny  
Bottles: 20

The operator<<() function presents in  
newline character at the end):

Gallo, tawny, 20

The Portmaster completed the method  
derived the VintagePort class as follow  
accidentally routing a bottle of '45 Coc  
tal barbecue sauce:

```
class VintagePort : public Port // s
{
private:
 char * nickname; // i.
 int year; // vi
public:
 VintagePort();
 VintagePort(const char * br, int
 VintagePort(const VintagePort &
 ~VintagePort() { delete [] nickn
 VintagePort & operator=(const Vi
```





- Virtual base classes
- Creating class templates
- Using class templates
- Template specializations

One of the main goals of C++ is to facilitate one mechanism for achieving this goal, but it offers other choices. One technique is to use class members of one class to hold objects of another class. This is referred to as *containment*. Another technique is to use private or protected inheritance. Containment and inheritance are typically used to implement *composition*, which the new class has an object of another class as a member. A `BluRayPlayer` might have a `BluRayPlayer` object. Multiple inheritance, in which a class inherits from two or more base classes, combines containment and inheritance.

Chapter 10, “Objects and Classes,” introduces composition. Chapter 11 looks at class templates, which provide another way to achieve composition. You define a class in generic terms. Then you create specific classes defined for specific types. For example, you define a `Stack` class and then use the template to create one class for `int` and another class that represents a stack of `double` values. A `Stack` represents a stack of stacks.



your project.)

Representing the quiz scores presents similar problems, which places a size limitation. You could use a `vector` or a large body of supporting code. You could use `std::string` and memory allocation to represent an array. You could use `std::string` that is capable of representing the data.

Developing your own class is not out of the question, but it is that difficult because an array of `double` shares a common memory. You could base the design of an array-of-`double` on `std::vector`, but in fact, that is what earlier editions of this book did.

But, of course, it is even easier if the library already does: the `valarray` class.

## **The `valarray` Class: A Quick Look**

The `valarray` class is supported by the `valarray` header. The `valarray` class is targeted to deal with numeric values (integers, floats, etc.). It supports operations such as summing the contents of an array of values in an array. So that it can handle different data types, the `valarray` class. Later, this chapter goes into how to use the `valarray` class. What you know now is how to use one.

the values from an ordinary array. With C++:

```
valarray<int> v5 = {20, 32, 17, 9}; // C++
```

Next, here are a few of the methods:

- The `operator[]()` method provides a
- The `size()` method returns the number
- The `sum()` method returns the sum of
- The `max()` method returns the largest
- The `min()` method returns the smallest

There are many more methods, some of which you've already seen more than enough to pro-

## The Student Class Design

At this point, the design plan for the `Student` class is to have a `name` and a `valarray<double>` object to store grades. What methods should be done? You might be tempted to publicly declare methods like `getGrade()`. That would be an example of multiple public methods for a single attribute, which would be inappropriate here. The reason is that the `Student` class and the `valarray` classes doesn't fit the *is-a* model. A student is not a



## The Student Class Example

At this point you need to provide the `Student` class with constructors and at least a few functions. Listing 14.1 does this, defining all the class members and friends for input and output.

### Listing 14.1 `studentc.h`

---

```
// studentc.h -- defining a Student class
#ifndef STUDENTC_H_
#define STUDENTC_H_

#include <iostream>
#include <string>
#include <valarray>

class Student
{
private:
 typedef std::valarray<double> ArrayDb;
 std::string name; // contained
 ArrayDb scores; // contained
 // private method for scores output
```

```

 // output
 friend std::ostream & operator<<(std::ostream& os, const Student& s) {
 os << s.name() << " " << s.scores() << "\n";
 };

#endif

```

---

In order to simplify notation, the `Student` class uses a `typedef` to define `ArrayDb` as a `std::valarray<double>`:

This enables the remaining code to use the `ArrayDb` type instead of `std::valarray<double>`. Thus, methods are defined as `void SetScores(ArrayDb& scores)`. Placing this `typedef` in the private portion of the `Student` class definition is not ideal. Placing it internally in the `Student` implementation but outside the `Student` class definition is better.

Note the use of the keyword `explicit`:

```

explicit Student(const std::string & s)
 : name(s), scores() {}
explicit Student(int n) : name("Nully"), scores(n) {}

```

Recall that a constructor that can be called with a single argument is called a `conversion function` from the argument type to the class type. The first constructor is an example of this idea. In the second constructor, for instance, the argument is the number of elements in an array rather than a value for the array.

```
Queue::Queue(int qs) : qsize(qs) {...} /
```

This code uses the name of the data member objects as the names of the constructors from previous examples, such as `hasDMA`, to initialize the base-class portion of a derived class object.

```
hasDMA::hasDMA(const hasDMA & hs) : baseDMA(hs) {}
```

For *inherited* objects, constructors use the `baseDMA` name to invoke a specific base-class constructor. For *non-inherited* objects, the constructor name. For example, look at the last constructor in the `Student` class.

```
Student(const char * str, const double * scores, int n) : name(str), scores(pd, n) {}
```

Because it initializes member objects, not just member names, not the class names, in the initialization list invokes the matching constructor. That is, `name(str)` invokes the `name(char *)` constructor, and `scores(pd, n)` invokes the `scores(double *, int)` constructor, which, because of the typedef, matches the `scores(double *, int)` constructor.

What happens if you don't use the initialization list? In C++, requires that all member objects be constructed. So if you omit the initialization list, you are in trouble for the member objects' classes.

```

 else
 return 0;
 }

```

This defines a method that can be invoked on a `valarray` using the `size()` and `sum()` methods. That's because `valarray` methods invoke the member functions of the `valarray` class, `Student` methods invoke `valarray` methods, and the `Student` method `sum()` invokes `valarray` methods.

Similarly, you can define a friend function that overloads the `<<` operator:

```

// use string version of operator<<()
ostream & operator<<(ostream & os, const Student & stu)
{
 os << "Scores for " << stu.name << ":\n";
 ...
}

```

Because `stu.name` is a `string` object, it invokes the `operator<<(string &)` function, which is provided as part of the `string` library. The `operator<<(ostream & os, const Student & stu)` function is a friend function of the `Student` class so that it can access the `name` member variable. It uses the public `Name()` method instead of the private `name` member variable.

}  
Using a helper like this gathers the messy coding of the friend function neater:

```
// use string version of operator<<()
ostream & operator<<(ostream & os, const
{
 os << "Scores for " << stu.name << ":
 stu.arr_out(os); // use private meth
 return os;
}
```

The helper function could also act as a builder, should you choose to provide them.

Listing 14.2 shows the class methods file for `studentc` to allow you to use the `[]` operator to access individual elements.

#### Listing 14.2    `studentc.cpp`

```
// studentc.cpp -- Student class using containers
#include "studentc.h"
using std::ostream;
using std::endl;
using std::istream;
using std::string;
```



```
// private method
ostream & Student::arr_out(ostream & os) c
{
 int i;
 int lim = scores.size();
 if (lim > 0)
 {
 for (i = 0; i < lim; i++)
 {
 os << scores[i] << " ";
 if (i % 5 == 4)
 os << endl;
 }
 if (i % 5 != 0)
 os << endl;
 }
 else
 os << " empty array ";
 return os;
}

// friends
```

Aside from the private helper method, Listing 14.3 shows how to use the `Student` class. Using containment allows you to take advantage of the code already written.

## Using the New Student Class

Let's put together a small program to test the `Student` class. The program should use an array of just three `Student` objects. The program should use an unsophisticated input cycle that reads until the user cuts the input process short. Listing 14.3 presents the program along with `studentc.cpp`.

### Listing 14.3 `use_stuc.cpp`

---

```
// use_stuc.cpp -- using a composite class
// compile with studentc.cpp
#include <iostream>
#include "studentc.h"
using std::cin;
using std::cout;
using std::endl;

void set(Student & sa, int n);
```

```
void set(Student & sa, int n)
{
 cout << "Please enter the student's name: ";
 getline(cin, sa);
 cout << "Please enter " << n << " quiz scores: ";
 for (int i = 0; i < n; i++)
 cin >> sa[i];
 while (cin.get() != '\n')
 continue;
}
```

---

Here is a sample run of the program in Listing 11.10:

```
Please enter the student's name: Gil Bayts
Please enter 5 quiz scores:
92 94 96 93 95
Please enter the student's name: Pat Roone
Please enter 5 quiz scores:
83 89 72 78 95
Please enter the student's name: Fleur O'D
Please enter 5 quiz scores:
92 89 96 74 64
```

*private inheritance*, public and protected members of the derived class. This means the methods of the public interface of the derived object. They are the functions of the derived class.

Let's look at the interface topic more closely. The methods of the base class become public methods of the derived class. The derived class inherits the base-class interface. This is part of the public interface, the public methods of the base class become the public methods of the derived class. In short, the derived class does not inherit the base-class interface. For objects, this lack of inheritance is part of the public interface.

With private inheritance, a class does inherit the base class's interface. For example, if you base a `Student` class on a `string` class, the `Student` class inherits the `string` class component that can be used to store a string. The `Student` class's methods can use the `string` methods internally.

Containment adds an object to a class as a member. Private inheritance adds an object to a class as an unnamed member. The term *subobject* to denote an object added by inheritance.

Private inheritance, then, provides the same interface as the base class. The derived class's implementation, don't acquire the interface. The derived class has a *has-a* relationship. In fact, you can produce a class that has the same public interface as the contained class. The differences between the two approaches affect the implementation. You can use private inheritance to redesign the

## Initializing Base-Class Components

Having implicitly inherited components instead of explicitly inheriting them is not a good idea in this example because you can no longer use `name`. Instead, you have to go back to the technique used in the previous example. To avoid this problem, consider constructors. Containment uses the following code:

```
Student(const char * str, const double * p, const int n)
: name(str), scores(pd, n) {} // u
```

The new version should use the member initializer list, which uses the *class* name instead of a *member* name:

```
Student(const char * str, const double * p, const int n)
: std::string(str), ArrayDb(pd, n) {}
```

Here, as in the preceding example, `ArrayDb` is a base class. Be sure to note that the member initializer list uses `std::string` instead of `name(str)`. This is the second main change.

Listing 14.4 shows the new class declaration and the use of explicit object names and the use of class names in constructors.

```

 : std::string(s), ArrayDb(a)
Student(const char * str, const double a)
 : std::string(str), ArrayDb(p)
~Student() {}
double Average() const;
double & operator[] (int i);
double operator[] (int i) const;
const std::string & Name() const;
// friends
// input
friend std::istream & operator>>(std::istream & is, Student & s)
friend std::istream & getline(std::istream & is, Student & s)

// output
friend std::ostream & operator<<(std::ostream & os, const Student & s)

};

#endif

```

---

`string::size()` ←

`valarray<double> object`

`valarray<double>::sum()` ←

```
class Student:private str
private valarray<double>
{
 ...
};
```

Figure 14.2 Objects within

Here, however, inheritance lets you use the  
tor to invoke base-class methods:

```
double Student::Average() const
{
 if (ArrayDb::size() > 0)
 return ArrayDb::sum()/ArrayDb::siz
```

This code returns a reference to the inherited `Student` object.

## Accessing Base-Class Friends

The technique of explicitly qualifying a function with `Base::` is used for friend functions because a friend function is not a member function and therefore cannot use an explicit type cast to the base class to invoke the same technique used to access a base-class member function. If you have a name for the `Student` object, so that you can call the friend function, for example, consider the following friend function:

```
ostream & operator<<(ostream & os, const Student & stu)
{
 os << "Scores for " << (const String & stu.name) << " are: ";
 ...
}
```

If `plato` is a `Student` object, then the following code prints the scores for `plato`, `stu` being a reference to `plato` and `os` being `cout`:

```
cout << plato;
```



```
using std::ostream;
using std::endl;
using std::istream;
using std::string;

// public methods
double Student::Average() const
{
 if (ArrayDb::size() > 0)
 return ArrayDb::sum()/ArrayDb::size();
 else
 return 0;
}

const string & Student::Name() const
{
 return (const string &) *this;
}

double & Student::operator[] (int i)
{
 return ArrayDb::operator[] (i);
}
```

```
 return os;
 }

 // friends
 // use String version of operator>>()
 istream & operator>>(istream & is, Student & stu)
 {
 is >> (string &stu);
 return is;
 }

 // use string friend getline(ostream & os, string & str, istream & is)
 istream & getline(istream & is, Student & stu)
 {
 getline(is, (string &stu));
 return is;
 }

 // use string version of operator<<()
 ostream & operator<<(ostream & os, const Student & stu)
 {
 os << "Scores for " << (const string &stu.name) << endl;
 stu.arr_out(os); // use private method
 return os;
 }
}
```

---

```

const int pupils = 3;
const int quizzes = 5;

int main()
{
 Student ada[pupils] =
 {Student(quizzes), Student(quizzes)

 int i;
 for (i = 0; i < pupils; i++)
 set(ada[i], quizzes);
 cout << "\nStudent List:\n";
 for (i = 0; i < pupils; ++i)
 cout << ada[i].Name() << endl;
 cout << "\nResults:";
 for (i = 0; i < pupils; i++)
 {
 cout << endl << ada[i];
 cout << "average: " << ada[i].Aver
 }
 cout << "Done.\n";
 return 0;
}

```

Student List:

Gil Bayts

Pat Roone

Fleur O'Day

Results:

Scores for Gil Bayts:

92 94 96 93 95

average: 94

Scores for Pat Roone:

83 89 72 78 95

average: 83.4

Scores for Fleur O'Day:

92 89 96 74 64

average: 83

Done.

The same input as before leads to the same results.  
produces.

Another situation that calls for using private virtual functions. Again, this is a privilege accorded to the base class. With private inheritance, the redefined function is private to the derived class, not publicly.

### Tip

In general, you should use containment to model composition. Use private inheritance if the new class needs to access the base class's private members. Use protected inheritance if it needs to redefine virtual functions.

## Protected Inheritance

Protected inheritance is a variation on private inheritance. It is used when listing a base class:

```
class Student : protected std::string,
 protected std::valarray<double>
{...};
```

With protected inheritance, public and protected members of the derived class. As with private inheritance, the base class is available to the derived class but not to the outside world. The difference between private and protected inheritance occurs when the derived class redefines a virtual function. With private inheritance, this function is private to the derived class. With protected inheritance, this function is protected to the derived class. That's because the derived class is using the base-class interface. That's because the

---

## Redefining Access with using

Public members of a base class become protected in a private derivation. Suppose you want to make a method publicly accessible in the derived class. One option is to make the method a static base-class method. For example, suppose you have a `std::valarray sum()` method. You can declare a static method, then define the method this way:

```
double Student::sum() const // public static method
{
 return std::valarray<double>::sum();
}
```

Then a `Student` object can invoke `Student::sum()` or `std::valarray<double>::sum()` method to the effect of `sum()`. If the `std::valarray` typedef is in scope, you can use `ArrayDb::sum()` instead.

There is an alternative to wrapping one function in a namespace (such as those used with namespaces) to make it publicly accessible. A member function can be used by the derived class, even though it is protected. Suppose you want to be able to use the `valarray` member function.

remove the existing prototypes and definitions. This declaration approach works only for inheritance.

There is an older way to redeclare base-class methods: place the method name in the public section of the derived class. To do that:

```
class Student : private std::string, private std::vector<double>
{
public:
 std::valarray<double>::operator[]; // ...
};
```

This looks like a using declaration without the *deprecated*, meaning that the intention is to phase it out. Using a using declaration, you can use it to make a member of the derived class.

## Multiple Inheritance

MI describes a class that has more than one inheritance. In public MI should express an *is-a* relationship. If you have a *Singer* class and a *Singer* class, you could derive a *Singer*

which is the circumstance that causes the mo  
tions for the Worker, Waiter, and Singer cla

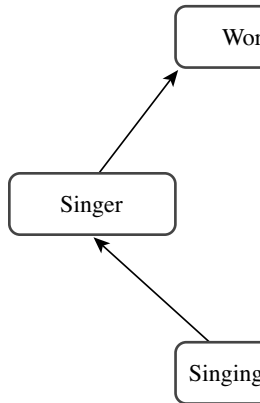


Figure 14.3 MI w



```

 {
private:
 int panache;
public:
 Waiter() : Worker(), panache(0) {}
 Waiter(const std::string & s, long n,
 : Worker(s, n), panache(p) {}
 Waiter(const Worker & wk, int p = 0)
 : Worker(wk), panache(p) {}
 void Set();
 void Show() const;
};

class Singer : public Worker
{
protected:
 enum {other, alto, contralto, soprano,
 bass, baritone, tenor}
 enum {Vtypes = 7};
private:
 static char *pv[Vtypes]; // string
 int voice;
public:
 Singer() : Worker(), voice(other) {}

```

```
using std::endl;
// Worker methods

// must implement virtual destructor, even if not needed
Worker::~Worker() {}

void Worker::Set()
{
 cout << "Enter worker's name: ";
 getline(cin, fullname);
 cout << "Enter worker's ID: ";
 cin >> id;
 while (cin.get() != '\n')
 continue;
}

void Worker::Show() const
{
 cout << "Name: " << fullname << "\n";
 cout << "Employee ID: " << id << "\n";
}

// Waiter methods
void Waiter::Set()
```

```

 cout << "Enter number for singer's voice: ";
 int i;
 for (i = 0; i < Vtypes; i++)
 {
 cout << i << ": " << pv[i] << " ";
 if (i % 4 == 3)
 cout << endl;
 }
 if (i % 4 != 0)
 cout << endl;
 while (cin >> voice && (voice < 0 || voice >= Vtypes))
 cout << "Please enter a value >= 0 and <= " << Vtypes << endl;

 while (cin.get() != '\n')
 continue;
 }

void Singer::Show() const
{
 cout << "Category: singer\n";
 Worker::Show();
 cout << "Vocal range: " << pv[voice] << endl;
}

```

---

```
std::cout << std::endl;
}

return 0;
}
```

---

Here is the output of the program in Listing 10.1:

```
Enter waiter's name: Waldo Dropmaster
Enter worker's ID: 442
Enter waiter's panache rating: 3
Enter singer's name: Sylvie Sireenne
Enter worker's ID: 555
Enter number for singer's vocal range:
0: other 1: alto 2: contralto 3: soprano
4: bass 5: baritone 6: tenor
3
Category: waiter
Name: Bob Apple
Employee ID: 314
Panache rating: 5

Category: singer
Name: Beverly Hills
```

## How Many Workers?

Suppose you begin by publicly deriving `Singer`

```
class SingingWaiter: public Singer, public
```

Because both `Singer` and `Waiter` inherit a `Worker` component (see Figure 10.1), `SingingWaiter` ends up with two `Worker` components (see Figure 10.2).

As you might expect, this raises problems. For example, how do you pass the address of a derived-class object to a base-class function?

```
SingingWaiter ed;
Worker * pw = &ed; // ambiguous
```

Normally, such an assignment sets a base-class pointer to the address of the `Worker` object within the derived object. But `ed` contains two `Worker` objects, with two addresses from which to choose. You could specify

```
Worker * pw1 = (Waiter *) &ed; // the Worker object in Waiter
Worker * pw2 = (Singer *) &ed; // the Worker object in Singer
```

This certainly complicates the technique of passing pointers to objects. We will refer to a variety of objects (polymorphism).

Having two copies of a `Worker` object causes another problem. The issue is why should you have two copies of a `Worker` object? If any other worker, should have just one name. To solve this problem, in the face of this of tricks, it added a virtual base class to make

## Virtual Base Classes

Virtual base classes allow an object derived from a common base to inherit just one object of the base. If you would make `Worker` a virtual base class to `Singer` and `Waiter`, virtual in the class declarations (virtual and public),

```
class Singer : virtual public Worker {...}
class Waiter : public virtual Worker {...}
```

Then you would define `SingingWaiter` as follows:

```
class SingingWaiter: public Singer, public Waiter {...}
```

Now a `SingingWaiter` object will contain one `Worker` subobject in essence, the inherited `Singer` and `Waiter` objects, each bringing in its own copy (see Figure 14.4). If you have one `Worker` subobject, you can use polymorphism.

Let's look at some questions you might have:

- Why the term *virtual*?
- Why don't we dispense with declaring the norm for MI?
- Are there any catches?

First, why the term *virtual*? After all, there's a difference between the concepts of virtual functions and virtual inheritance. It's from the C++ community to resist the introduction of a new keyword corresponding to a new facility—a bit of keyword overloading.

Next, why don't we dispense with declaring the norm for MI? First, there are cases where a base class is virtual. Second, making a base class virtual requires accounting, and you shouldn't have to pay for it. Third, there are the disadvantages presented in the next paragraph.

Finally, are there catches? Yes. Making virtual functions C++ rules, and you have to code some things differently. This may involve changing existing code. For example,

```
class C : public B
{
 int c;
public:
 C(int q = 0, int m = 0, int n = 0) :
 ...
};
```



return; Suppose you do want a new version of `Show()` for this object to invoke an inherited `Show()` method:

```
SingingWaiter newhire("Elise Hawks", 2005, 1000);
newhire.Show(); // ambiguous
```

With single inheritance, failing to redefine `Show()` in the derived class causes the compiler to use the `Show()` from the ancestral definition. In this case, each direct ancestor has its own `Show()`, so this call is ambiguous.

## Caution

Multiple Inheritance can result in ambiguous calls. For example, `Draw()` methods in `ChessPlayer` and `PokerPlayer` could inherit two quite different `Draw()` methods from the `Player` class.

You can use the scope-resolution operator to specify which `Show()` to use:

```
SingingWaiter newhire("Elise Hawks", 2005, 1000);
newhire.Singer::Show(); // use Singer version
```

However, a better approach is to redefine `Show()` in the derived class to specify which `Show()` to use. For example, if you want to use the `Singer` version, you could use this:

```

 waiter::Show();
 }
 cout << "Presence rating: " << presence;
}

```

This incremental approach fails for the `show` method because it ignores the `waiter` component.

```

void SingingWaiter::Show()
{
 Singer::Show();
}

```

You can remedy that by calling the `waiter` component's `show` method.

```

void SingingWaiter::Show()
{
 Singer::Show();
 Waiter::Show();
}

```

However, this displays a person's name and address twice, because both `waiter::Show()` and `worker::Show()` call `worker::Show()`.

How can you fix this? One way is to use a *delegation* approach. That is, you can provide a method `show` that delegates the call to `waiter::Show()` (and another method that displays only `waiter` components), and another that displays only

```

}

void SingingWaiter::Show() const
{
 cout << "Category: singing waiter\n";
 Worker::Data();
 Data();
}

```

Similarly, the other `Show()` methods would be implemented in the `Worker` components.

With this approach, objects would still use `protected` methods, on the other hand, should be internally used to facilitate the public interface. `protected` methods used to facilitate the public interface would prevent, say, `Waiter` code from using `protected` methods of situation for which the `protected` access class was intended, they can be used internally by all the objects derived from the outside world.

Another approach would be to make all the methods `protected`, but using `protected` methods instead of `protected` data. This would allowable access to the data.

The `Set()` methods, which solicit data for the objects, would be `protected`. For example, `SingingWaiter::Set()` should be `protected`.

```

 virtual void Get();
public:
 Worker() : fullname("no one"), id(0L)
 Worker(const std::string & s, long n)
 : fullname(s), id(n) {}
 virtual ~Worker() = 0; // pure virtual
 virtual void Set() = 0;
 virtual void Show() const = 0;
};

```

```

class Waiter : virtual public Worker
{
private:
 int panache;
protected:
 void Data() const;
 void Get();
public:
 Waiter() : Worker(), panache(0) {}
 Waiter(const std::string & s, long n,
 : Worker(s, n), panache(p) {}
 Waiter(const Worker & wk, int p = 0)
 : Worker(wk), panache(p) {}
}

```

```

};

// multiple inheritance
class SingingWaiter : public Singer, public Worker
{
protected:
 void Data() const;
 void Get();
public:
 SingingWaiter() {}
 SingingWaiter(const std::string & s, int v = 0,
 int v = other)
 : Worker(s,n), Waiter(s, n, p) {}
 SingingWaiter(const Worker & wk, int p)
 : Worker(wk), Waiter(wk,p), Singer(wk) {}
 SingingWaiter(const Waiter & wt, int v)
 : Worker(wt),Waiter(wt), Singer(wt) {}
 SingingWaiter(const Singer & wt, int p)
 : Worker(wt),Waiter(wt,p), Singer(wt) {}
 void Set();
 void Show() const;
};

#endif

```

---

```
 while (cin.get() != '\n')
 continue;
 }

 // Waiter methods
 void Waiter::Set()
 {
 cout << "Enter waiter's name: ";
 Worker::Get();
 Get();
 }

 void Waiter::Show() const
 {
 cout << "Category: waiter\n";
 Worker::Data();
 Data();
 }

 // protected methods
 void Waiter::Data() const
 {
 cout << "Panache rating: " << panache
 }
}
```

```

 Worker::Data();
 Data();
 }

 // protected methods
 void Singer::Data() const
 {
 cout << "Vocal range: " << pv[voice] << endl;
 }

 void Singer::Get()
 {
 cout << "Enter number for singer's voice: ";
 int i;
 for (i = 0; i < Vtypes; i++)
 {
 cout << i << ": " << pv[i] << " ";
 if (i % 4 == 3)
 cout << endl;
 }
 if (i % 4 != 0)
 cout << '\n';
 cin >> voice;
 }

```

```

void SingingWaiter::Show() const
{
 cout << "Category: singing waiter\n";
 Worker::Data();
 Data();
}

```

---

Of course, curiosity demands that you test the code to do so. Note that the program makes addresses of various kinds of classes to base-class. The C-style string library function `strchr()` in `string.h` is used in the following code snippet. The `while` loop continues until `strchr("wstq", choice) == NULL`.

This function returns the address of the first occurrence of the character `choice` in the string `"wstq"`; the function returns `NULL` if the character is not found. This test is simpler to write than an `if` statement that checks for each character individually.

Be sure to compile Listing 14.12 along with the other programs in the chapter.



```

 << "w: waiter s: singer "
 << "t: singing waiter q: quit\n";
 cin >> choice;
 while (strchr("wstq", choice) == NULL)
 {
 cout << "Please enter a w, s, t or q\n";
 cin >> choice;
 }
 if (choice == 'q')
 break;
 switch(choice)
 {
 case 'w': lolas[ct] = new Waiter();
 break;
 case 's': lolas[ct] = new Singer();
 break;
 case 't': lolas[ct] = new SingingWaiter();
 break;
 }
 cin.get();
 lolas[ct]->Set();
}

cout << "\nHere is your staff:\n";

```

Enter singer's name: **Sinclair Parma**

Enter worker's ID: **1044**

Enter number for singer's vocal range:

0: other    1: alto    2: contralto    3: soprano

4: bass    5: baritone    6: tenor

**5**

Enter the employee category:

w: waiter    s: singer    t: singing waiter

**t**

Enter singing waiter's name: **Natasha Garg**

Enter worker's ID: **1021**

Enter waiter's panache rating: **6**

Enter number for singer's vocal range:

0: other    1: alto    2: contralto    3: soprano

4: bass    5: baritone    6: tenor

**3**

Enter the employee category:

w: waiter    s: singer    t: singing waiter

**w**

Here is your staff:

Category: waiter

Name: Wally Slipshod

x, and y. In this case, class M contains one class ancestor (that is, classes C and D) and a separate ancestor (that is, classes x and y). So, all told, it's one. When a class inherits a particular base class through multiple virtual paths, the class has one base-class subobject. It has a separate base-class subobject to represent each

## Virtual Base Classes and Dominance

Using virtual base classes alters how C++ resolves names in classes, the rules are simple. If a class inherits two classes with the same name from different classes, using the name of the class name is ambiguous. If virtual base classes are used, the name may not be ambiguous. In this case, if one name dominates unambiguously without a qualifier.

So how does one member name dominate another? It dominates the same name in any ancestor class, when considering the following definitions:

```
class B
{
public:
 short q();
 ...
}
```

```
{
 ...
};
```

Here the definition of `q()` in class `C` dominates the definition of `q()` derived from `B`. Thus, methods in `F` can use `q()` and neither definition of `omg()` dominates the other. Therefore, an attempt by `F` to use `omg()` is ambiguous.

The virtual ambiguity rules pay no attention to access. `E::omg()` is private and hence not directly accessible. Similarly, even if `C::q()` were private, it could call `B::q()` in class `F`, but an unqualified reference to `C::q()` is ambiguous.

## Multiple Inheritance Synopsis

First, let's review MI without virtual base classes. However, if a class inherits two members with the same name, you need to use class qualifiers in the derived class. That is, methods in the `BadDude` class, derived from `Gunslinger` and `PokerFace`, would use `Gunslinger::draw()` and `PokerFace::draw()` methods inherited from the two classes to avoid ambiguous usage.

# Class Templates

Inheritance (public, private, or protected) and when you want to reuse code. Consider, for example, and the `Queue` class (see Chapter 12). These are classes designed to hold other objects or data. For example, `unsigned long` values. You could be storing double values or string objects. The code for an object stored. However, rather than write new code, you could define a stack in a generic (that is, type-independent) way, with a specific type as a parameter to the class. Then you could produce stacks of different kinds of values. In C++, `typedef` as a first pass at dealing with this design. Drawbacks. First, you have to edit the header file for each type you can use the technique to generate just one header file. You can't have a `typedef` represent two different types. You can't use a method to define a stack of `ints` and a stack of `double`s.

C++'s class templates provide a better way to do this. C++ originally did not support templates, and since then it has continued to evolve, so it is possible that your compiler does not support the templates presented here.) Templates provide *parameterization* of a class. You present a type name as an argument to a recipe for building a class.

```

 stack();
 bool isempty() const;
 bool isfull() const;
 // push() returns false if stack already full
 bool push(const Item & item); // add item
 // pop() returns false if stack already empty
 bool pop(Item & item); // pop item
};

```

The template approach will replace the `Stack` class and the `Stack` member functions with template functions, you preface a template class with `template`

```
template <class Type>
```

The keyword `template` informs the compiler that the part in angle brackets is analogous to an angle bracketed name, the keyword `class` as serving as a type name, and of `Type` as representing a name for this variable.

Using `class` here doesn't mean that `Type` serves as a generic type specifier for which a template is used. Newer C++ implementations allow `typename` instead of `class` in this context:

```
template <typename Type> // newer choice
```

```
...
}
```

becomes the following:

```
template <class Type>
bool Stack<Type>::push(const Type & item)
{
 ...
}
```

If you define a method within the class definition, you must include the template preface and the class qualifier.

Listing 14.13 shows the combined class and function definitions. (It is important to realize that these templates are not class and function definitions; they are instructions to the C++ compiler about how to generate definitions. A particular actualization of a template, such as the creation of string objects, is called an *instantiation* or a *specialization*. Functions and functions in a separate implementation file would normally provide the keyword `export` to allow such a specialization. Some vendors implemented it. C++11 discontinues its use. (The `export` keyword for possible future use.) Because the templates are compiled separately, templates have to be used in the same way as instantiations of templates. The simplest way to

```
template <class Type>
```

```
Stack<Type>::Stack()
```

```
{
 top = 0;
}
```

```
template <class Type>
```

```
bool Stack<Type>::isempty()
```

```
{
 return top == 0;
}
```

```
template <class Type>
```

```
bool Stack<Type>::isfull()
```

```
{
 return top == MAX;
}
```

```
template <class Type>
```

```
bool Stack<Type>::push(const Type & item)
```

```
{
 if (top < MAX)
 {
```



## Using a Template Class

Merely including a template in a program does not ask for an instantiation. To do this, you declare using the generic type name with the particular type you would create two stacks, one for stacking

```
Stack<int> kernels; // create a stack of integers
Stack<string> colonels; // create a stack of strings
```

Seeing these two declarations, the compiler will generate two separate class declarations and two separate class definitions. The `Stack<int>` class declaration will replace `Type` in the `Stack` definition. The `Stack<string>` class declaration will replace `Type` in the `Stack` definition. The algorithms you use have to be consistent with the type. For example, assumes that you can assign one item to another. For types, structures, and classes (unless you make them copyable), this is fine. For arrays, this is not possible.

Generic type identifiers such as `Type` in the `Stack` definition act something like variables, but instead of a variable, you assign a type to it. So in the kernel stack, the type is `int`.

Notice that you have to provide the desired function templates, for which the compiler will figure out what kind of function to generate.

```

 cout << "Please enter A to add a purchase

 << "P to process a PO, or Q to quit\n";

while (cin >> ch && std::toupper(ch) != 'Q')

{

 while (cin.get() != '\n')

 continue;

 if (!std::isalpha(ch))

 {

 cout << "\a";

 continue;

 }

 switch(ch)

 {

 case 'A':

 case 'a': cout << "Enter a PO number: ";

 cin >> po;

 if (st.isfull())

 cout << "stack is full\n";

 else

 st.push(po);

 break;

 case 'P':

 case 'p': if (st.isempty())

 cout << "stack is empty\n";

 else

 st.pop();

 break;

 }

}

```

Please enter A to add a purchase order,  
P to process a PO, or Q to quit.

**A**

Enter a PO number to add: **silver747boeing**

Please enter A to add a purchase order,  
P to process a PO, or Q to quit.

**P**

PO #silver747boeing popped

Please enter A to add a purchase order,  
P to process a PO, or Q to quit.

**P**

PO #blueR8audi popped

Please enter A to add a purchase order,  
P to process a PO, or Q to quit.

**P**

PO #red911porsche popped

Please enter A to add a purchase order,  
P to process a PO, or Q to quit.

**P**

stack already empty

Please enter A to add a purchase order,  
P to process a PO, or Q to quit.

**Q**

Bye

```
char * po;
```

The idea is to use a char pointer instead of a char array for input. This approach fails immediately because there is no space to hold the input strings. (The program crashes after cin tried to store input in some invalid memory.)

Version 2 replaces

```
string po;
```

with

```
char po[40];
```

This allocates space for an input string. Unfortunately, it is placed on the stack. But an array is fundamental to the pop() method:

```
template <class Type>
bool Stack<Type>::pop(Type & item)
{
 if (top > 0)
 {
 item = items[--top];
 return true;
 }
}
```

## Using a Stack of Pointers Correctly

One way to use a stack of pointers is to have a set of pointers, with each pointer pointing to a different object. This then makes sense because each pointer will represent the responsibility of the calling program, not the stack's job is to manage the pointers, not create objects.

For example, suppose you have to simulate a cart of folders to Plodson. If Plodson's cart is empty, he takes a file from the cart and places it in his in-basket. If the cart is not empty nor full, Plodson may process the top file from the cart and put it into his in-basket. If the cart is full, cap self-expression, he flips a coin to decide whether to investigate the effects of his method on the other files.

You can model this with an array of pointers. Each string will contain the name of the file. The stack to represent the in-basket, and you can use the out-basket. Adding a file to the in-basket is adding an input array onto the stack, and processing a file is popping the stack and adding it to the out-basket.

```

 Stack(const Stack & st);
 ~Stack() { delete [] items; }
 bool isempty() { return top == 0; }
 bool isfull() { return top == stacksize; }
 bool push(const Type & item); // add
 bool pop(Type & item); // pop
 Stack & operator=(const Stack & st);
};

```

```

template <class Type>
Stack<Type>::Stack(int ss) : stacksize(ss)
{
 items = new Type [stacksize];
}

```

```

template <class Type>
Stack<Type>::Stack(const Stack & st)
{
 stacksize = st.stacksize;
 top = st.top;
 items = new Type [stacksize];
 for (int i = 0; i < top; i++)
 items[i] = st.items[i];
}

```

```

template <class Type>
Stack<Type> & Stack<Type>::operator=(const Stack<Type> & st)
{
 if (this == &st)
 return *this;
 delete [] items;
 stacksize = st.stacksize;
 top = st.top;
 items = new Type [stacksize];
 for (int i = 0; i < top; i++)
 items[i] = st.items[i];
 return *this;
}

#endif

```

---

Notice that the prototype declares the return type to be a reference to `Stack`, and the actual type as `Stack<Type>`. The former is an abbreviation within the class scope. That is, you can use `st` inside the template function definitions, but not

```

// in basket
 const char * in[Num] = {
 " 1: Hank Gilgamesh", " 2: Ki
 " 3: Betty Rocker", " 4: Ian
 " 5: Wolfgang Kibble", " 6: P
 " 7: Joy Almondo", " 8: Xaver
 " 9: Juan Moore", "10: Misha
 };

// out basket
 const char * out[Num];

 int processed = 0;
 int nextin = 0;
 while (processed < Num)
 {
 if (st.isempty())
 st.push(in[nextin++]);
 else if (st.isfull())
 st.pop(out[processed++]);
 else if (std::rand() % 2 && nextin
 st.push(in[nextin++]);
 else
 st.pop(out[processed++]);
 }

```



Bye

Please enter stack size: 5

3: Betty Rocker  
5: Wolfgang Kibble  
6: Portia Koop  
4: Ian Flagranti  
8: Xaverie Paprika  
9: Juan Moore  
10: Misha Mache  
7: Joy Almondo  
2: Kiki Ishtar  
1: Hank Gilgamesh

Bye

## Program Notes

The strings in Listing 14.16 never move. Pushing a new pointer to an existing string. That is, it creates a new string. And popping a string off the stack removes it from the array.

The program uses `const char *` as a type for a set of string constants.

```

//arrayTP.h
#ifndef ARRAYTP_H_
#define ARRAYTP_H_

#include <iostream>
#include <cstdlib>

template <class T, int n>
class ArrayTP
{
private:
 T ar[n];
public:
 ArrayTP() {};
 explicit ArrayTP(const T & v);
 virtual T & operator[] (int i);
 virtual T operator[] (int i) const;
};

template <class T, int n>
ArrayTP<T,n>::ArrayTP(const T & v)
{
 for (int i = 0; i < n; i++)
 ar[i] = v;
}

```

```
#endif
```

---

Note the template heading in Listing 14.1.

```
template <class T, int n>
```

The keyword `class` (or, equivalently in this case, `typename`) is the first template parameter, or type argument. `int` identifies `n` as the second template parameter, one that specifies a particular type. This second parameter is called a *non-type*, or *expression, argument*. Suppose we have

```
ArrayTP<double, 12> eggweights;
```

This causes the compiler to define a class `ArrayTP` of type `double` and an `eggweights` object of that class. When defining `ArrayTP`, the compiler uses `double` and `n` with `12`.

Expression arguments have some restrictions. They cannot be a pointer type, an enumeration type, a reference, or a function type. However, `double &rm` and `double * pm` are allowed. `rm` is the name of the argument or take its address. Thus, in the example above, `&n` or `&n` would not be allowed. Also when you use an expression argument should be a constant expression.

## Template Versatility

You can apply the same techniques to templates. Template classes can serve as base classes, and they can themselves be type arguments to other templates. You can use a template by using an array template. Or you can use a template to construct an array whose elements are stacks. Here is some code along the following lines:

```
template <typename T> // or <class T>
class Array
{
private:
 T entry;
 ...
};

template <typename Type>
class GrowArray : public Array<Type> {...}

template <typename Tp>
class Stack
{
 Array<Tp> ar; // use an Array
 ...
}
```

cout.width(2) causes the next item to be displayed with a width of 2 characters, unless a larger width is needed to show the value.

### Listing 14.18    **twod.cpp**

---

```
// twod.cpp -- making a 2-d array
#include <iostream>
#include "arraytp.h"
int main(void)
{
 using std::cout;
 using std::endl;
 ArrayTP<int, 10> sums;
 ArrayTP<double, 10> aves;
 ArrayTP< ArrayTP<int,5>, 10> twodee;

 int i, j;

 for (i = 0; i < 10; i++)
 {
 sums[i] = 0;
 for (j = 0; j < 5; j++)
 {
 twodee[i][j] = (i + 1) * (j + 1);
 }
 }
}
```

of `twodee`, each of which is a five-element array. The average of an element of `twodee`:

```
1 2 3 4 5 : sum = 15, average = 1.5
2 4 6 8 10 : sum = 30, average = 3
3 6 9 12 15 : sum = 45, average = 4.5
4 8 12 16 20 : sum = 60, average = 6
5 10 15 20 25 : sum = 75, average = 7.5
6 12 18 24 30 : sum = 90, average = 9
7 14 21 28 35 : sum = 105, average = 10.5
8 16 24 32 40 : sum = 120, average = 12
9 18 27 36 45 : sum = 135, average = 13.5
10 20 30 40 50 : sum = 150, average = 15
Done.
```

## Using More Than One Type Parameter

You can have a template with more than one type parameter. You might want a class that holds two kinds of values. You could use a `pair` for holding two disparate values. (Incidentally, `pair` is a standard library pair.) The short program in Listing 14.19 shows how. The `first()` and `second()` `const` methods report the stored values. The `get1()` and `get2()` methods, by virtue of returning references to the stored values by using assignment.

```

 return a;
 }
 template<class T1, class T2>
 T2 & Pair<T1,T2>::second()
 {
 return b;
 }

int main()
{
 using std::cout;
 using std::endl;
 using std::string;
 Pair<string, int> ratings[4] =
 {
 Pair<string, int>("The Purpled Duck", 4),
 Pair<string, int>("Jaquie's Frisco", 3),
 Pair<string, int>("Cafe Souffle", 2),
 Pair<string, int>("Bertie's Eats", 1),
 };

 int joints = sizeof(ratings) / sizeof(Pair<string, int>);
 cout << "Rating:\t Eatery\n";
 for (int i = 0; i < joints; i++)

```

## Default Type Template Parameters

Another new class template feature is that you can specify default type parameters:

```
template <class T1, class T2 = int> class
```

This causes the compiler to use `int` for the second parameter.

```
Topo<double, double> m1; // T1 is double, T2 is double
Topo<double> m2; // T1 is double, T2 is int
```

The STL (discussed in Chapter 16) often uses this feature to specify a class.

Although you can provide default values for class template parameters, you can't do so for function template parameters. However, you can provide default type parameters for both class and function templates.

## Template Specializations

Class templates are like function templates in that they can be used with explicit instantiations, and explicit specializations.



```
template class ArrayTP<string, 100>; // ge
```

In this case, the compiler generates the class even though no object of the class has yet been created. In other words, even without explicit instantiation, the general template is used to create a specialization.

## Explicit Specializations

An *explicit specialization* is a definition for a particular specialization instead of the general template. Sometimes you want a class to behave differently when instantiated for a particular type. You can do this with an explicit specialization. Suppose, for example, that `SortedArray` represents a sorted array for which items are sorted in ascending order.

```
template <typename T>
class SortedArray
{
 ...// details omitted
};
```

Also suppose the template uses the `>` operator to compare numbers. It will work if `T` represents a class type that has a `T::operator>()` method. But it won't work if `T` is a primitive type like `char *`. Actually, the template will work, but it

```
SortedArray<int> scores; // use
SortedArray<const char *> dates; // use
```

## Partial Specializations

C++ allows for *partial specializations*, which p  
For example, a partial specialization can prov  
parameters:

```
// general template
 template <class T1, class T2> class Pair
// specialization with T2 set to int
 template <class T1> class Pair<T1, int>
```

The <> following the keyword template c  
unspecialized. So the second declaration speci  
specifying all the types leads to an empty brace

```
// specialization with T1 and T2 set to i
 template <> class Pair<int, int> {...
```

The compiler uses the most specialized te  
would happen given the preceding three tem

```

 template <class T1, class T2, class T3> class Trio<T1, T2, T3>
// specialization with T3 set to T2
 template <class T1, class T2> class Trio<T1, T2, T2>
// specialization with T3 and T2 set to T1
 template <class T1> class Trio<T1, T1, T1>

```

Given these declarations, the compiler would

```

Trio<int, short, char *> t1; // use general Trio
Trio<int, short> t2; // use Trio<T1, T2, T2>
Trio<char, char *, char *> t3; use Trio<T1, T1, T1>

```

## Member Templates

A template can be a member of a structure, class, or union, and can have the same feature to fully implement its design. Listing 14.20 shows a class with a nested template class and a template member function.

Listing 14.20 **tempmemb.cpp**

---

```

// tempmemb.cpp -- template members
#include <iostream>
using std::cout;
using std::endl;

```

```

int main()
{
 beta<double> guy(3.5, 3);
 cout << "T was set to double\n";
 guy.Show();
 cout << "V was set to T, which is double\n";
 cout << guy.blab(10, 2.3) << endl;
 cout << "U was set to int\n";
 cout << guy.blab(10.0, 2.3) << endl;
 cout << "U was set to double\n";
 cout << "Done\n";
 return 0;
}

```

---

The hold template is declared in the private section of the beta class, and is used only within the beta class scope. The beta class has two data members:

```

hold<T> q; // template object
hold<int> n; // template object

```

28.2809

U was set to double

Done

Note that replacing 10 with 10.0 in the call to `blab` would make `blab` return a `double`, making the return type `double`, which is not the case.

As mentioned previously, the type of the `beta` object is not set by the declaration of the `guy` object. Unlike the first example, the type of the `beta` parameter is not set by the function call. For instance, you can implement `blab()` as `blab(int, double)`, and the type of the `beta` parameter by the usual function prototype rules:

```
cout << guy.blab(10, 3) << endl;
```

You can declare the `hold` class and `blab` method outside the `beta` template. However, you cannot declare template members at all, and others that accept the `beta` type as a parameter. However, if you want to define the template methods outside the `beta` template, you can use the following code:

```
template <typename T>
class beta
{
private:
 template <typename V> // declaration
```

```
// member definition
template <typename T>
 template <typename U>
 U beta<T>::blab(U u, T t)
 {
 return (n.Value() + q.Value()) * u;
 }
 }
```

The definitions have to identify `T`, `v`, and `u`. If the templates are nested, you have to use the

```
template <typename T>
 template <typename V>
```

syntax instead of this syntax:

```
template<typename T, typename V>
```

The definitions also must indicate that `blab` is a member of `beta` class, and they use the scope-resolution operator `::`.

## Templates As Parameters

You've seen that a template can have type parameters. It can also have non-type parameters, such as `int n`. A template class `template<int n> class`

declaration:

```
Crab<Stack> nebula;
```

Hence, in this case, `Thing<int>` is instantiated as `Stack<double>`. In short, the template type is used as a template argument.

The `Crab` class declaration makes three functions represented by `Thing`. The class should have a `pop()` method, and these methods should have any template class that matches the `Thing` site `push()` and `pop()` methods. This chapter template defined in `stacktp.h`, so the example

#### Listing 14.21 **tempparm.cpp**

---

```
// tempparm.cpp - templates as parameters
#include <iostream>
#include "stacktp.h"

template <template <typename T> class Thing>
class Crab
{
private:
 Thing<int> s1;
```

```

 while (nebula.pop(ni, nb))
 cout << ni << ", " << nb << endl;
 cout << "Done.\n";

 return 0;
 }

```

---

Here is a sample run of the program in Listing 11.1:

```

Enter int double pairs, such as 4 3.5 (0 to stop):
50 22.48
25 33.87
60 19.12
0 0
60, 19.12
25, 33.87
50, 22.48
Done.

```

You can mix template parameters with regular parameters. The function declaration could start out like this:

```

template <template <typename T> class This, typename... Args>
class Crab

```



- Unbound template friends, meaning that

to each specialization of the class

Let's look at examples of each.

## Non-Template Friend Functions to Template

Let's declare an ordinary function in a template

```
template <class T>
class HasFriend
{
public:
 friend void counts(); // friend to
 ...
};
```

This declaration makes the `counts()` function a friend of the template. For example, it would be a friend of the `HasFriend<string>` class.

The `counts()` function is not invoked by a template specialization, and it has no object parameters, so how is it invoked? There are several possibilities. It could access a global variable; it could use a global pointer; it could create its own static members of a template class, which exist separately

```

 friend void report (HasFriend<int> &);
 ...
};

```

That is, `report()` with a `HasFriend<int>` `HasFriend<int>` class. Similarly, `report()` will be an overloaded version of `report()` that is

Note that `report()` is not itself a template. This means that you have to define a plan to use:

```

void report (HasFriend<short> &) {...}; //
void report (HasFriend<int> &) {...}; //

```

Listing 14.22 illustrates these points. The program Note that this means that each particular specialization is a friend to one particular `HasFriend` specialization. The `counts()` method, which is a friend to `HasFriend<double>`, reports the value of `ct` for two particular specializations of `HasFriend<double>`. The program also provides a `friend` declaration for `HasFriend<double>` to be a friend to one particular `HasFriend` specialization.

```

// non-template friend to all HasFriend<T>
void counts()
{
 cout << "int count: " << HasFriend<int>::count;
 cout << "double count: " << HasFriend<double>::count;
}

// non-template friend to the HasFriend<int>
void reports(HasFriend<int> & hf)
{
 cout << "HasFriend<int>: " << hf.item << endl;
}

// non-template friend to the HasFriend<double>
void reports(HasFriend<double> & hf)
{
 cout << "HasFriend<double>: " << hf.item << endl;
}

int main()
{
 cout << "No objects declared: ";
 counts();
}

```

```
HasFriend<int>: 10
HasFriend<int>: 20
HasFriend<double>: 10.5
```

## Bound Template Friend Functions to Templates

You can modify the preceding example by making the functions themselves. In particular, you can set things up so that for each specialization of a class gets a matching specialization of the functions. This is more complex than for non-template friends and is more complex than for non-template functions.

For the first step, you declare each template function as follows:

```
template <typename T> void counts();
template <typename T> void report(T &);
```

Next, you declare the templates again as functions. You can then declare specializations based on the class template specialization.

```
template <typename TT>
class HasFriendT
{
...
 friend void counts<TT>();
 friend void report<>(HasFriendT<TT> &);
};
```

```
};
```

One specialization is based on `TT`, which becomes `HasFriendT<TT>`, which becomes `HasFriendT<TT>::count<int>()` and `report<HasFriendT<int>>()` and `HasFriendT<int>` class.

The third requirement the program must make is that the friends. Listing 14.23 illustrates these three `count()` function that is a friend to all `HasFriendT` `count()` functions, one of which is a friend to the `count()` function calls have no function parameter. To deduce the desired specialization, these calls use forms to indicate the specialization. For the case, use the argument type to deduce the specialization, the same effect:

```
report<HasFriendT<int>> >(hfi2); // same as
```

#### Listing 14.23 `tmp2tmp.cpp`

```
// tmp2tmp.cpp -- template friends to a template class
#include <iostream>
using std::cout;
using std::endl;
```

```

template <typename T>
void counts()
{
 cout << "template size: " << sizeof(H
 cout << "template counts(): " << HasF
}

template <typename T>
void report(T & hf)
{
 cout << hf.item << endl;
}

int main()
{
 counts<int>();
 HasFriendT<int> hfi1(10);
 HasFriendT<int> hfi2(20);
 HasFriendT<double> hfdb(10.5);
 report(hfi1); // generate report(Has
 report(hfi2); // generate report(Has
 report(hfdb); // generate report(Has
 cout << "counts<int>() output:\n";
 counts<int>();
}

```

of a template declared outside a class. An int specialization, and so on. By declaring a template friend functions for which every function specialization. For unbound friends, the friend template class type parameters:

```
template <typename T>
class ManyFriend
{
...
 template <typename C, typename D> friend
};
```

Listing 14.24 shows an example that uses a `show2(hfi1, hfi2)` gets matched to the following:

```
void show2<ManyFriend<int> &, ManyFriend<int> &>
 (ManyFriend<int> & c, ManyFriend<int> & d)
```

Because it is a friend to all specializations of `ManyFriend`, the item members of all specializations. But it is not a friend to the objects.

Similarly, `show2(hfd, hfi2)` gets matched to the following:

```
void show2<ManyFriend<double> &, ManyFriend<double> &>
 (ManyFriend<double> & c, ManyFriend<double> & d)
```

```
 cout << c.item << ", " << d.item << endl;
 }

int main()
{
 ManyFriend<int> hfi1(10);
 ManyFriend<int> hfi2(20);
 ManyFriend<double> hfdb(10.5);
 cout << "hfi1, hfi2: ";
 show2(hfi1, hfi2);
 cout << "hfdb, hfi2: ";
 show2(hfdb, hfi2);

 return 0;
}
```

---

Here's the output of the program in Listing 10.1:

```
hfi1, hfi2: 10, 20
hfdb, hfi2: 10.5, 20
```



```
arrtype<int> days; // days is t
arrtype<std::string> months; // months is
```

In short, `arrtype<T>` means type `std::array<T>`.

C++11 extends the `using =` syntax to non-`typedef` types, making it equivalent to an ordinary `typedef`:

```
typedef const char * pc1; // typedef
using pc2 = const char *; // using =
typedef const int *(*pa1)[10]; // typedef
using pa2 = const int *(*)[10]; // using =
```

As you get used to it, you may find the new syntax for defining the type name from type information more convenient.

Another C++11 addition to templates is the ability to define a template class or function that can take a type argument. 18, “Visiting with the New C++ Standard,” later in this book.

## Summary

C++ provides several means for reusing code. Chapter 13, “Class Inheritance,” enables you to model inheritance, which is able to reuse the code of base classes. Private inheritance reuses base-class code, this time modeling *has-a* relationships.

bases. You can use class qualifiers to resolve names and avoid multi-inherited bases. However, using virtual inheritance and writing initialization lists for constructors and destructors.

Class templates let you create a generic class. A class type, is represented by a type parameter. A type parameter

```
template <class T>
class Ic
{
 T v;
 ...
public:
 Ic(const T & val) : v(val) { }
 ...
};
```

Here T is the type parameter, and it acts as a placeholder until later time. (This parameter can have any valid type choices.) You can also use `typename` instead of `class`.

```
template <typename T> // same as template <class T>
class Rev { ... } ;
```

Class definitions (instantiations) are generated for each type that specifies a particular type. For example, the following

```

public:
 Ic(const char * s) : str(s) { }
 ...
};

```

Then a declaration of the following form rather than using the general template:

```
class Ic<char *> chic;
```

A class template can specify more than one parameters:

```

template <class T, class TT, int n>
class Pals {...};

```

The following declaration would generate string for TT, and 6 for n:

```
Pals<double, string, 6> mix;
```

A class template can also have parameters t

```

template < template <typename T> class CL,
class Trophy {...};

```

Here z stands for an int value, U stands for template declared using template <typename

```
class Person
```

```
class Person
```

```
class Person, class Automobile
```

2. Suppose you have the following defini

```
class Frabjous {
```

```
private:
```

```
 char fab[20];
```

```
public:
```

```
 Frabjous(const char * s = "C+
```

```
 virtual void tell() { cout <<
```

```
};
```

```
class Gloam {
```

```
private:
```

```
 int glip;
```

```
 Frabjous fb;
```

```
public:
```

```
 Gloam(int g = 0, const char *
```

```
 Gloam(int g, const Frabjous &
```

```
 void tell();
```

```
};
```

provide definitions for the three `Gloam`

4. Suppose you have the following definition 14.13 and the `Worker` class of Listing 14.13

```
Stack<Worker *> sw;
```

Write out the class declaration that will not use inline class methods.

5. Use the template definitions in this chapter to write the definitions for the following:
  - An array of `string` objects
  - A stack of arrays of `double`
  - An array of stacks of pointers to `Worker`

How many template class definitions are there?

6. Describe the differences between virtual and non-virtual functions.

```
// ...
Wine(const char * l, int y);
```

The wine class should have a method `Count()` that returns the number of years, prompts the user to enter the number of bottles, and returns the bottle counts. A method `Label()` should return the label. A method `sum()` should return the total. A `valarray<int>` object in the `Pair` object.

The program should prompt the user to enter the number of elements of the array, and the year and bottle counts. The program should use this data to create a `Pair` object and store the information stored in the object. For example:

```
// pe14-1.cpp -- using Wine class
#include <iostream>
#include "winec.h"

int main (void)
{
 using std::cin;
 using std::cout;
 using std::endl;

 cout << "Enter name of wine: ";
 char lab[50];
```

Enter name of wine: **Gully Wash**  
Enter number of years: **4**  
Enter Gully Wash data for 4 year(s):  
Enter year: **1988**  
Enter bottles for that year: **42**  
Enter year: **1994**  
Enter bottles for that year: **58**  
Enter year: **1998**  
Enter bottles for that year: **122**  
Enter year: **2001**  
Enter bottles for that year: **144**  
Wine: Gully Wash

| Year | Bottles |
|------|---------|
| 1988 | 42      |
| 1994 | 58      |
| 1998 | 122     |
| 2001 | 144     |

Wine: Gushing Grape Red

| Year | Bottles |
|------|---------|
| 1993 | 48      |
| 1995 | 60      |
| 1998 | 72      |

Total bottles for Gushing Grape Red:  
Bye

value. (Optionally, you could define a `Draw()` and use a `Card` return value for `Draw()` `show()` function. The `BadDude` class defines two `PokerPlayer` classes. It has a `Gdraw()` and a `Cdraw()` member that returns the `Show()` function. Define all these classes and necessary methods (such as methods for setting up the program similar to that in Listing 14.12).

5. Here are some class declarations:

```
// emp.h -- header file for abstr_emp

#include <iostream>
#include <string>

class abstr_emp
{
private:
 std::string fname; // abstr_emp
 std::string lname; // abstr_emp
 std::string job;
```



```

{
private:
 int inchargeof; // number
protected:
 int InChargeOf() const { return
 int & InChargeOf(){ return incha
public:
 manager();
 manager(const std::string & fn,
 const std::string & j, i
 manager(const abstr_emp & e, int
 manager(const manager & m);
 virtual void ShowAll() const;
 virtual void SetAll();
};

```

```

class fink: virtual public abstr_emp
{
private:
 std::string reportsto; //
protected:
 const std::string ReportsTo() co
 std::string & ReportsTo(){ retur

```

```
};
```

Note that the class hierarchy uses MI with special rules for constructor initialization of some protected-access methods. This applies to `highfink` methods. (Note, for example, that `manager::ShowAll()` calls `fink::ShowAll()` and `manager::abstr_emp::ShowAll()` twice.) Provide the classes in a program. Here is a mini

```
// pe14-5.cpp
// useempl.cpp -- using the abstr_emp
```

```
#include <iostream>
using namespace std;
#include "emp.h"
```

```
int main(void)
{
 employee em("Trip", "Harris", " ")
 cout << em << endl;
 em.ShowAll();
}
```

Why is no assignment operator defined?

Why are `ShowAll()` and `SetAll()` virtual?

Why is `abstr_emp` a virtual base class?

Why does the `highfink` class have no constructor?

Why is only one version of `operator<` defined?

What would happen if the end of the program was

```
abstr_emp tri[4] = {em, fi, hf, hf2};
for (int i = 0; i < 4; i++)
 tri[i].ShowAll();
```

- `dynamic_cast` and `typeid`
- `static_cast`, `const_cast`, and `reinterpret_cast`

This chapter ties up some loose ends and the C++ language. The loose ends include friend nested classes, which are classes declared within other classes, here are exceptions, runtime type identification, C++ exception handling provides a mechanism that otherwise would bring a program to a halt for object types. The new type cast operators and other facilities are fairly new to C++, and older code

## Friends

Several examples in this book so far use friend functions for a class. Such functions are not the only kind of function that can be a friend. In that case, any method of the class can be a friend. In that case, any method of the class can be a friend. In that case, any method of the class can be a friend. Also you can declare particular member functions of a class to be friends. Functions, member functions, or classes are friends if they are outside. Thus, although friends do grant outside

- Cable or antenna tuning mode
- TV tuner or A/V input

The tuning mode reflects the fact that, in the past, the tuning for channels 14 and up is different from the tuning for channels 2 to 13 for over-the-air broadcast reception. The input selection chooses between broadcast TV, and a DVD. Some sets may offer other inputs, but this list is enough for the purpose of this class.

Also a television has some parameters that can be set. The number of channels they can store is a common one. The channel track that value.

Next, you must provide the class with methods to control the television. Most television sets these days hide their controls behind a remote control. The remote has buttons to change channels, and so on, without the need to go up or down one channel at a time but can jump to any channel. There's usually a button for increasing the volume.

A remote control should duplicate the controls of the television. The methods can be implemented by using the `TV` methods. The `TV` class provides random access channel selection from channel 2 to channel 20 without going through all the channels. The `TV` can work in two or more modes, for example, as a channel controller.

```
enum {Antenna, Cable};
enum {TV, DVD};
```

```
Tv(int s = Off, int mc = 125) : state
 maxchannel(mc), channel(2), mode(
void onoff() {state = (state == On)?
bool ison() const {return state == On
bool volup();
bool voldown();
void chanup();
void chandown();
void set_mode() {mode = (mode == Ante
void set_input() {input = (input == T
void settings() const; // display all
private:
 int state; // on or off
 int volume; // assumed to
 int maxchannel; // maximum num
 int channel; // current cha
 int mode; // broadcast c
 int input; // TV or DVD
};
```

use wraparound, with the lowest channel setting and the highest channel setting, `maxchannel`.

Many of the methods use the conditional settings:

```
void onoff() {state = (state == On)? Off
```

Provided that the two state values are true, this can be done more compactly by using the complement operator (`^=`), as discussed in Appendix E.

```
void onoff() {state ^= 1;}
```

In fact, you could store up to eight bivalent variables and toggle them individually, but that's beyond the bitwise operators discussed in Appendix E.

## Listing 15.2    **tv.cpp**

---

```
// tv.cpp -- methods for the Tv class (Revised)
#include <iostream>
#include "tv.h"

bool Tv::volup()
{
```

```

 else
 channel = 1;
 }

void Tv::chardown()
{
 if (channel > 1)
 channel--;
 else
 channel = maxchannel;
}

void Tv::settings() const
{
 using std::cout;
 using std::endl;
 cout << "TV is " << (state == Off? "Off" : "On") << endl;
 if (state == On)
 {
 cout << "Volume setting = " << volume << endl;
 cout << "Channel setting = " << channel << endl;
 cout << "Mode = " << (mode == Antenna? "antenna" : "cable") << endl;
 }
}

```



```

 s42.settings();

 Remote grey;

 grey.set_chan(s42, 10);
 grey.volup(s42);
 grey.volup(s42);
 cout << "\n42\" settings after using r
 s42.settings();

 Tv s58(Tv::On);
 s58.set_mode();
 grey.set_chan(s58,28);
 cout << "\n58\" settings:\n";
 s58.settings();
 return 0;
}

```

---

Here is the output of the program in Listing 10.1:

```

Initial settings for 42" TV:
TV is Off

```

make the private parts of the `Tv` class public and `Remote` a friend of `Tv` that encompasses both a television and a remote control. This is a good idea to reflect the fact that a single remote control can control both a television and a remote control.

## Friend Member Functions

Looking at the code for the last example, you can see that the `set_chan()` methods are implemented by using the public interface of `Tv`. In other words, the methods don't really need friend status. Indeed, the only private `Tv` member directly is `Remote::set_chan()`. You can make `Remote` to be a friend. You do have the option of making `Remote` a friend of `Tv` rather than making the entire `Remote` class a friend. You need to be careful about the order in which you define the friend definitions. Let's look at why.

The way to make `Remote::set_chan()` a friend in the `Tv` class declaration:

```
class Tv
{
 friend void Remote::set_chan(Tv & t,
 ...
};
```

method in particular.

Another difficulty remains. In Listing 15.1 such as the following:

```
void onoff(Tv & t) { t.onoff(); }
```

Because this calls a `Tv` method, the compiler has to know about the `Tv` class definition at this point so that it knows what method `onoff` is. The compiler's job is not necessarily follows the `Remote` declaration. The compiler has to place the `Remote` to method *declarations* and to place the `Tv` to method *definitions*. This leads to the following ordering:

```
class Tv; // forward declaration
class Remote { ... }; // Tv-using method declarations
class Tv { ... }; // Tv method declarations
// put Remote method definitions here
```

The `Remote` prototypes look like this:

```
void onoff(Tv & t);
```

All the compiler needs to know when inspecting the forward declaration supplies that information. When it reads the actual method definitions, it has already read the information needed to compile those methods.

```

 bool voldown(Tv & t);
 void onoff(Tv & t) ;
 void chanup(Tv & t) ;
 void chandown(Tv & t) ;
 void set_mode(Tv & t) ;
 void set_input(Tv & t);
 void set_chan(Tv & t, int c);

};

class Tv
{
public:
 friend void Remote::set_chan(Tv & t,
 enum State{Off, On};
 enum {MinVal,MaxVal = 20};
 enum {Antenna, Cable};
 enum {TV, DVD};

 Tv(int s = Off, int mc = 125) : state
 maxchannel(mc), channel(2), mode(0)
 void onoff() {state = (state == On)?
 bool ison() const {return state == On
 bool volup();

```

```
inline void Remote::set_chan(Tv & t, int c) {
 #endif
```

---

If you include `tvfm.h` instead of `tv.h` in the program, the program behaves the same as the original. The difference is that `Remote` instead of all the `Remote` methods—is a friend of `TV`. This is a friendly difference.

Recall that inline functions have internal linkage. Their definition must be in the file that uses the function. Here, the definition of `set_chan` is in the file, so including the header file in the file that uses the function is the right place. You could place the definition of `set_chan` in the header file, provided that you remove the `inline` keyword, to avoid multiple definition.

By the way, making the entire `Remote` class a friend of `TV` is possible, because the friend statement itself identifies `Remote` as a class. Here is how:

```
friend class Remote;
```

## Other Friendly Relationships

Other combinations of friends and classes besides the ones shown here are possible. Let's take a brief look at some of them.

```
t
 friend void Remote::set_chan
 ...
};
// Remote methods here
```

```
private:

public:
void set_chan(Tv & t, int c)
```

Remote object

Just `Remote::set_chan()` can

Figure 15.1 Class friends

Suppose the advance of technology brings an interactive remote control unit might let you watch a television program, and the television might respond to your response is wrong. Ignoring the possibility of a program the viewers, let's just look at the C++ program that benefit from mutual friendship, with some Remote

```
...
}
```

Because the `Remote` declaration follows the definition, it is not defined in the class declaration. However, the compiler will consider the `Tv` declaration so that the definition of `buzz()` is not inlined. If you don't want `buzz()` to be inline, you need to declare it as `inline` in the class declaration.

## Shared Friends

Another use for friends is when a function needs to access the private members of two classes. Logically, such a function should be a member of one of the classes, but it is impossible. It could be a member of one class, but it is not a member of the other. It is more reasonable to make the function a friend of both classes. For example, you might have a `Probe` class that represents some sort of probe, and an `Analyzer` class that represents some sort of probe analyzer. The `Analyzer` class has an internal clock, and you would like to be able to access it. You might use something along the following lines:

```
class Analyzer; // forward declaration
class Probe
{
 friend void sync(Analyzer & a, const Probe & p);
 friend void sync(Probe & p, const Analyzer & a);
 ...
}
```

In C++, you can place a class declaration inside another class. A class declared inside another is called a *nested class*, and it helps avoid namespace pollution. Member functions of the class contain objects of the nested class. The outside world can use the nested class only if the declaration is in the public section and if you use the `using` keyword. The standard of C++ either don't allow nested classes or it is not recommended.

Nesting classes is not the same as containing a class object as a member of another class. Nesting creates a new class, not a class member. Instead, it defines a type that contains the nested class declaration.

The usual reasons for nesting a class are to avoid namespace pollution and to avoid name conflicts. The `Queue` class in the following example illustrates nesting a structure definition:

```
class Queue
{
private:
 // class scope definitions
 // Node is a nested structure definition
 struct Node {Item item; struct Node *next;
 ...
};
```



```

class Queue
{
// class scope definitions
 // Node is a nested class definition
 class Node
 {
 public:
 Item item;
 Node * next;
 Node(const Item & i) : item(i), next(
 };
 ...
};

```

This constructor initializes the node's item to 0, which is one way of writing the null pointer (including a header file that defines `NULL`. Use `nullptr`.) Because all nodes created by the `Queue` use this pointer, this is the only constructor the class needs.

Next, you need to rewrite `enqueue()` by using

```

bool Queue::enqueue(const Item & item)
{
 if (isfull())
 return false;

```

## Scope

If a nested class is declared in a private section of the first class, it is visible only to the second class. This applies, for example, to the preceding example. Hence, `Queue` member `Node` objects, but other parts of a program do not have access to them. If you were to derive a class from `Queue`, `Node` objects in the derived class can't directly access the private members of `Queue`.

If the nested class is declared in a protected section of the first class but invisible to the outside world. However, the first class has information about the nested class and could directly create objects of it.

If a nested class is declared in a public section of the first class, it is visible to the second class, to classes derived from the second class, and to the outside world. However, because the nested class is not a friend of the first class, the first class must use the `friend` qualifier in the outside world. For example, suppose we have the following code:

```
class Team
{
public:
 class Coach { ... };
 ...
};
```

## Access Control

After a class is in scope, access control comes into play. In the `Queue` example, the nested class that governs access to a regular class is the `Node` class. The `Queue` declaration does not grant the `Queue` class any special access to the `Node` class, nor does it grant the `Node` class any special access to the `Queue` class. A `Queue` class object can access only the public members of the `Node` class. For this reason, the `Queue` example makes all the member functions of the `Node` class private. This is the usual practice of making data members private and member functions public. The implementation feature of the `Queue` class and is not a member of the `Queue` class. Since the `Node` class is declared in the private section of the `Queue` class, its methods can access `Node` members directly, a class can access its own private members.

In short, the location of a class declaration determines its access. Given that a particular class is in scope, the user of the class (public, private, friend) determine the access a program has to the class.

## Nesting in a Template

You've seen that templates are a good choice for the `Queue` class. You may be wondering whether it is possible to convert the `Queue` class definition to a template. The answer is yes, and the result is a template class.

```

 Node * rear; // pointer to rear
 int items; // current number of items
 const int qsize; // maximum number of items
 QueueTP(const QueueTP & q) : qsize(0), items(0) {}
 QueueTP & operator=(const QueueTP & q) {
public:
 QueueTP(int qs = Q_SIZE);
 ~QueueTP();
 bool isempty() const
 {
 return items == 0;
 }
 bool isfull() const
 {
 return items == qsize;
 }
 int queuecount() const
 {
 return items;
 }
 bool enqueue(const Item &item); // add item to queue
 bool dequeue(Item &item); // remove item from queue
 };

```

```

{
 if (isfull())
 return false;
 Node * add = new Node(item); // create new node
 // on failure, new throws std::bad_alloc exception
 items++;
 if (front == 0) // if queue is empty
 front = add; // place item at front
 else
 rear->next = add; // else place item at rear
 rear = add; // have rear point to new node
 return true;
}

```

// Place front item into item variable and return it

```

template <class Item>
bool QueueTP<Item>::dequeue(Item & item)
{
 if (front == 0)
 return false;
 item = front->item; // set item to front item
 items--;
 Node * temp = front; // save location of front
 front = front->next; // move front to next node
 delete temp; // delete old front
 return true;
}

```

## Listing 15.6    **nested.cpp**

---

```
// nested.cpp -- using a queue that has a
#include <iostream>

#include <string>
#include "queuetp.h"

int main()
{
 using std::string;
 using std::cin;
 using std::cout;

 QueueTP<string> cs(5);
 string temp;

 while(!cs.isfull())
 {
 cout << "Please enter your name.
 "served in the order of a
 "name: ";
 getline(cin, temp);
 cs.enqueue(temp);
 }
}
```

The queue is full. Processing begins!  
Now processing Kinsey Millhone...  
Now processing Adam Dalgliesh...  
Now processing Andrew Dalziel...  
Now processing Kay Scarpetta...  
Now processing Richard Jury...

## Exceptions

Programs sometimes encounter runtime problems that are not handled normally. For example, a program may try to allocate more memory than is available, or it may encounter an error. Programmers try to anticipate such calamities. C++ provides a powerful tool for dealing with these situations. Exceptions are a feature of C++, so some older compilers haven't implemented them. This feature is turned off by default, so you may have to use a compiler that has it on.

Before examining exceptions, let's look at a simple example. The harmonic mean is available to programmers. As a test case, let's look at the harmonic mean of two numbers. The *harmonic mean* is the reciprocal of the average of the inverses. This can be reduced to the following formula:

$$2.0 \quad x \quad y \quad / \quad (x + y)$$

```
//error1.cpp -- using the abort() function
#include <iostream>
#include <cstdlib>
double hmean(double a, double b);

int main()
{
 double x, y, z;

 std::cout << "Enter two numbers: ";
 while (std::cin >> x >> y)
 {
 z = hmean(x,y);
 std::cout << "Harmonic mean of "
 << " is " << z << std::endl;
 std::cout << "Enter next set of numbers: ";
 }
 std::cout << "Bye!\n";
 return 0;
}

double hmean(double a, double b)
{
 if (a == 0 || b == 0)
 abort();
 return 2*a*b/(a+b);
}
```



(You may, perhaps, enjoy the conceit that a program could avoid aborting by checking the `hmean()` function. However, it's not safe to rely enough to perform such a check.

## Returning an Error Code

A more flexible approach than aborting is to return an error code. For example, the `get(void)` member function of `istream` returns the ASCII code for the next input character, but returns 0 if it reaches the end-of-file. This approach doesn't work for `hmean()` because there's no special value for a valid return value, so there's no special value for an error. In this situation, you can use a pointer argument to return an error code to the calling program and use the function return value to indicate the success or failure, yet still perform actions other than aborting. Listing 15.8 shows `hmean()` as a `bool` function whose return value is a pointer to the argument for obtaining the answer.

```

 std::cout << "Bye!\n";
 return 0;
 }

 bool hmean(double a, double b, double * ans)
 {
 if (a == -b)
 {
 *ans = DBL_MAX;
 return false;
 }
 else
 {
 *ans = 2.0 * a * b / (a + b);
 return true;
 }
 }
}

```

---

Here's a sample run of the program in Lis

```

Enter two numbers: 3 6
Harmonic mean of 3 and 6 is 4
Enter next set of numbers <q to quit>: 10
One value should not be the negative of t

```

## The Exception Mechanism

Now let's see how you can handle problems by using exceptions. An *exception* is a response to an exceptional circumstance, such as an attempt to divide by zero. Exceptions are used to transfer control from one part of a program to another. Handling exceptions involves the following:

- Throwing an exception
- Catching an exception with a handler
- Using a `try` block

A program throws an exception when a problem occurs. For example, we modify `hmean()` in Listing 15.7 to throw an exception. A `throw` statement, in essence, is a jump; it transfers control to a statement at another location. The `throw` keyword is followed by a value, such as a character string, that identifies the exception.

A program catches an exception with an `except` clause, where you want to handle the problem. The `except` clause catches the exception. A handler begins with the keyword `except` (followed by parentheses) that indicates the type of exception being caught, followed by a brace-enclosed block of code that handles the exception. The keyword, along with the exception type, serves

```

 }
 z = hmean(x,y);
} // end of
catch (const char * s) // start
{
 std::cout << s << std::endl;
 std::cout << "Enter a new pair
 continue;
} // end of
std::cout << "Harmonic mean of "
 << " is " << z << std::endl;
std::cout << "Enter next set of n
}
std::cout << "Bye!\n";
return 0;
}

```

```

double hmean(double a, double b)
{
 if (a == -b)
 throw "bad hmean() arguments: a =
 return 2.0 * a * b / (a + b);
}

```

---

Throwing an exception looks like this:

```
if (a == -b)
 throw "bad hmean() arguments: a = -b"
```

In this case, the thrown exception is the string "bad hmean() arguments: a = -b". The exception type can be a `std::string` or a custom class type is the usual choice, as later examples show.

Executing the `throw` is a bit like executing a `return` statement. It ends the current function execution. However, instead of returning a value, it causes a program to back up through the sequence of function calls to the function that contains the `try` block. In C++, this is the calling function. Soon you'll see an example of exception handling. Meanwhile, in this case, the `throw` passes control back to the program looks for an exception handler (following the `try` block). If no handler is found, the program terminates. If one is found, the handler of exception thrown.

The handler, or `catch` block, looks like this:

```
catch (char * s) // start of exception handler
{
 std::cout << s << std::endl;
 std::cout << "Enter a new pair of numbers: ";
 continue;
} // end of handler
```

loop illustrates the fact that handler statements on the same line acts like a label directing program flow (

You might wonder what happens if a function does not have a block or else no matching handler. By default, the function will throw an exception, but you can modify that behavior. V

## Using Objects as Exceptions


Typically, functions that throw exceptions throw a specific type. This is that you can use different exception types to represent different situations that produce exceptions. Also an object can use this information to help identify the exception that was thrown. Also in principle a catch block could take a specific course of action to pursue. Here, for example, an exception could be thrown by the `hmean()` function:

```
class bad_hmean
{
private:
 double v1;
 double v2;
public:
 bad_hmean(int a = 0, int b = 0) : v1(a), v2(b) {}
};
```

```

 cout << "Enter next set of
 }
 ...
 double hmean(double a, double b)
 {
 if (a == -b)
 throw "bad hmean() argu
 return 2.0 * a * b / (a + b)
 }

```



1. The program calls `hmean()` within a try
2. `hmean()` throws an exception, transferring  
assigning the exception string to `s`.
3. The catch block transfers execution back

Figure 15.2 Program

A `bad_hmean` object can be initialized to the  
method can be used to report the problem, in  
can use code like this:

```

if (a == -b)
 throw bad_hmean(a,b);

```

`gmean()` throws a `bad_gmean` exception, the  
and gets caught by the second.

#### Listing 15.10    **exc\_mean.h**

---

```
// exc_mean.h -- exception classes for h
#include <iostream>

class bad_hmean
{
private:
 double v1;
 double v2;
public:
 bad_hmean(double a = 0, double b = 0)
 void mesg();
};

inline void bad_hmean::mesg()
{
 std::cout << "hmean(" << v1 << ", " << v2 << "
 << "invalid arguments: a = " << v1 << ", b = " << v2 << "\n";
}
```



```

double gmean(double a, double b);

int main()
{
 using std::cout;
 using std::cin;
 using std::endl;

 double x, y, z;

 cout << "Enter two numbers: ";
 while (cin >> x >> y)
 {
 try {
 // start of try block
 z = hmean(x,y);
 cout << "Harmonic mean of " << x << " and " << y << "
 << " is " << z << endl;
 cout << "Geometric mean of " << x << " and " << y << "
 << " is " << gmean(x,y) << endl;
 cout << "Enter next set of numbers: ";
 } // end of try block
 catch (bad_hmean & bg) // start of catch block
 {
 bg.mesg();
 cout << "Try again.\n";
 }
 }
}

```

```
{
 if (a < 0 || b < 0)
 throw bad_gmean(a,b);
 return std::sqrt(a * b);
}
```

---

Here's a sample run of the program, one to test the `gmean()` function:

```
Enter two numbers: 4 12
Harmonic mean of 4 and 12 is 6
Geometric mean of 4 and 12 is 6.9282
Enter next set of numbers <q to quit>: 5
hmean(5, -5): invalid arguments: a = -b
Try again.
5 -2
Harmonic mean of 5 and -2 is -6.66667
gmean() arguments should be >= 0
Values used: 5, -2
Sorry, you don't get to play any more.
Bye!
```

One point to notice is that the `bad_hmean` handler uses a `break` statement, while the `bad_gmean` handler uses a `throw` statement.

try block. However, that can be accomplished for another reason was to allow the compiler to add code to the exception specification was violated. This can happen without throwing an exception, but it might call a function that throws an exception. And maybe that function didn't throw an exception, but after a library update, it now does. Anyway, the C++ programming community, particularly among those who write exception-safe code, was that this feature is better than exception specifications with the blessings of the C++ standards committee.

However, C++11 does allow for one specific use of `noexcept`. The `noexcept` specifier can be used to indicate a function that does not throw an exception.

```
double marm() noexcept; // marm() doesn't throw
```

There is some debate about the necessity of this feature. Some people feel it's better to avoid using it (at least in C++11), but enough about the need to introduce a new keyword. The fact that a function shouldn't throw an exception can be thought of as a promise made by the programmer. It should be thought of as a promise made by the programmer.

There also is a `noexcept()` operator (see [A.10](#)). It can be used to indicate its operand could throw an exception.

return address that resides in a `try` block (see exception handlers at the end of the block for a function call. This process is called *unwinding*. The throw mechanism is that, just as with functions, any automatic class objects on the stack. However, they are put on the stack by that function, whereas they are removed from the stack by the entire sequence of function calls. Without the unwinding-the-stack feature, a program could not place automatic class objects placed on the stack by a function.

Listing 15.12 provides an example of unwinding. The `main()` function, which in turn calls `hmean()` and `gmean()`. The `main()` and `means()` create objects of the `MyClass` type when its constructor and destructor are used. The `try` block catches `bad_hmean` and `bad_gmean` exceptions, and the `try` block catches `bad_gmean` exception. This `catch` block has the following

```
catch (bad_hmean & bg) // start of catch block
{
 bg.mesg();
 std::cout << "Caught in means()\n";
 throw; // rethrows the exception
}
```

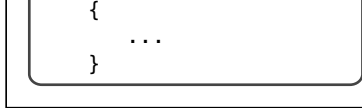


Figure 15.3 thro

After the code responds by displaying mess means, in this case, sending the exception on exception rises to the next try-catch combin exception. If no handler is found, the program same header file (`exc_mean.h` in listing 15.10)

#### Listing 15.12 **error5.cpp**

```
//error5.cpp -- unwinding the stack
#include <iostream>
#include <cmath> // or math.h, unix users
#include <string>
#include "exc_mean.h"

class demo
{
private:
 std::string word;
public:
 demo (const std::string & str)
```

```

using std::cin;
using std::endl;

double x, y, z;
{
 demo d1("found in block in main()");
 cout << "Enter two numbers: ";
 while (cin >> x >> y)
 {
 try {
 //
 z = means(x,y);
 cout << "The mean mean
 << " is " << z
 cout << "Enter next pair: ";
 } // end of try block
 catch (bad_hmean & bg)
 {
 bg.mesg();
 cout << "Try again.\n";
 continue;
 }
 catch (bad_gmean & hg)
 {
 cout << hg.mesg();

```

```

 if (a < 0 || b < 0)
 throw bad_gmean(a,b);
 return std::sqrt(a * b);
 }

double means(double a, double b)
{
 double am, hm, gm;
 demo d2("found in means()");
 am = (a + b) / 2.0; // arithmetic m
 try
 {
 hm = hmean(a,b);
 gm = gmean(a,b);
 }
 catch (bad_hmean & bg) // start of catch
 {
 bg.mesg();
 std::cout << "Caught in means()\n";
 throw; // rethrows the
 }
 d2.show();
 return (am + hm + gm) / 3.0;
}

```

---

Bye!

## Program Notes

Let's trace through the course of the sample : the demo constructor announces, an object is means () is called, and another demo object is values 6 and 12 on to hmean () and gmean (), means (), which calculates a result and return invokes d2.show(). After returning the result d2 is called automatically:

```
demo found in means() lives!
demo found in means() destroyed
```

The next input cycle sends the values 6 and a new demo object and relays the values to hmean bad\_hmean exception, which is caught by the following output:

```
hmean(6, -6): invalid arguments: a = -b
Caught in means()
```



demo found in means() destroyed

Finally, the `bad_gmean` handler in `main()` calls `gmean()` arguments should be `>= 0`  
Values used: 6, -8  
Sorry, you don't get to play any more.

Then the program terminates normally, displaying the message, calling the destructor for `d1`. If the `catch` block had not occurred, `break`, the program would terminate immediately.  
demo found in main() lives!  
Bye!

However, you would still see this message:  
demo found in main() destroyed

Again, the exception mechanism would not return to the caller on the stack.

## More Exception Features

Although the `throw-catch` mechanism is similar to the `return` mechanism, there are a few differences. One is that a `return` statement in a function `fun()` transfers control back to the caller.

```

 catch(problem & p)
 {
 // statements
 }

```

Here, `p` would refer to a copy of `oops` rather than `oops` itself. `oops` no longer exists after `super()` termination. The `throw` instruction with the `throw`:

```

throw problem(); // construct and throw

```

At this point you might wonder why the `throw` creates a copy. After all, the usual reason for using a reference is not having to make a copy. The answer is that `throw` is not a property: A base-class reference can also refer to a collection of exception types that are related to the property. The specification need only list a reference to the base class of the derived objects thrown.

Suppose you have a class hierarchy of exception types separately. A base-class reference can catch any object. An object can only catch that object and objects derived from it. An object is caught by the first `catch` block that matches the object. `catch` blocks in inverse order of derivation:

```
{ // statements }
```

If the `bad_1` & handler were first, it would handle the exceptions. With the inverse order, a `bad_3` exception would not be caught by the handler.

### Tip

If you have an inheritance hierarchy of exceptions, arrange the `catch` blocks so that the exception of the most specific type (the one furthest down the class hierarchy sequence) is caught last.

Arranging `catch` blocks in the proper sequence ensures that each type of exception is handled. But sometimes you may not know the exception to expect. For instance, say you write a function that calls another function that you don't know whether that other function throws an exception even if you don't know the type. Then you can use an ellipsis for the exception type:

```
catch (...) { // statements } // catches any exception
```

The main intent for C++ exceptions is to provide a way to write fault-tolerant programs. That is, exceptions must be integrated into a program design so you don't have to take fault handling as an afterthought. The flexibility and robustness encourage programmers to integrate fault handling when it's appropriate. In short, exceptions are the kind of natural approach to programming.

Newer C++ compilers are incorporating support for the exception header file (formerly `except.h`). The `exception` class that C++ uses as a base class for other exception classes. Your code, too, can throw an exception object. One virtual member function is named `what()`, which is implementation dependent. However, you can redefine it in a class derived from `exception`.

```
#include <exception>

class bad_hmean : public std::exception
{
public:
 const char * what() { return "bad argument"; }
 ...
};

class bad_gmean : public std::exception
{
```

```

class logic_error : public exception {
public:
 explicit logic_error(const string& what_arg) : exception(what_arg) {}
 ...
};

class domain_error : public logic_error {
public:
 explicit domain_error(const string& what_arg) : logic_error(what_arg) {}
 ...
};

```

Note that the constructors take a string of the character data returned as a C-style string.

These two new classes serve, in turn, as bases for the `logic_error` family describes, as you might expect. Sound programming could avoid such errors, but the C++ standard library does not. The name of each class indicates the sort of error it represents.

```

 domain_error
 invalid_argument
 length_error
 out_of_bounds

```

out\_of\_bounds exception if the index used is

Next, the `runtime_error` family describes errors that could not easily be predicted and provide a sort of error it is intended to report:

```
range_error
overflow_error
underflow_error
```

Each class has a constructor like that of `runtime_error` and a `what()` method that returns a string to be returned by the `what()` method.

An underflow error can occur in floating-point arithmetic. The smallest nonzero magnitude that a floating-point type can represent would produce a smaller value would cause a underflow. This can occur with either integer or floating-point types. If a calculation would exceed the largest representable value, the result can lie outside the valid range of a function. This is a runtime error, and you can use the `range_error` exception.

In general, an exception of the `logic_error` family is amenable to a programming fix, whereas a `runtime_error` is unavoidable trouble. All these error classes have a `what()` method. The distinction is that the different class names all

throw a `bad_alloc` exception. The new header class, which is publicly derived from the exception, returned a null pointer when it couldn't allocate.

Listing 15.13 demonstrates the current approach. The program displays the implementation-dependent error message and terminates early.

#### Listing 15.13    **newexcp.cpp**

---

```
// newexcp.cpp -- the bad_alloc exception
#include <iostream>
#include <new>
#include <cstdlib> // for exit(), EXIT_FAILURE
using namespace std;

struct Big
{
 double stuff[20000];
};

int main()
{
 Big * pb;
 try {
```

In this case, the `what()` method returns the amount of memory requested.

## The Null Pointer and `new`

Much code was written when `new` (the old `new`) was the only way to allocate memory. Some compilers handled the transition to the C++ standard by providing a switch to choose which behavior she wanted. The C++ standard provides a native form of `new` that still returns a null pointer if it fails.

```
int * pi = new (std::nothrow) int;
int * pa = new (std::nothrow) int[500];
```

Using this form, you could rewrite the code as follows:

```
Big * pb;

pb = new (std::nothrow) Big[10000]; // 1,000,000 bytes
if (pb == 0)
{
 cout << "Could not allocate memory. Bye!\n";
 exit(EXIT_FAILURE);
}
```



```

private:
 int bi; // bad index value
public:
 explicit bad_index(int ix,
 const std::string & s = "Index")
 : bi_val(ix) {}
 int bi_val() const {return bi;}
 virtual ~bad_index() throw() {}
};

explicit Sales(int yy = 0);
Sales(int yy, const double * gr, int n) : year(yy), gross(gr) {}
virtual ~Sales() { }
int Year() const { return year; }
virtual double operator[](int i) const { return gross[i]; }
virtual double & operator[](int i);

private:
 double gross[MONTHS];
 int year;
};

class LabeledSales : public Sales
{
public:
 class nbad_index : public Sales::bad_index
 {

```

available as a type to client catch blocks. Note that `bad_index` can be identified as `Sales::bad_index`. This class is a friend of the `Sales` class. The `bad_index` class has the ability to store and report an array index.

The `nbad_index` class is nested in the public section of the `Sales` class, available to client code as `LabeledSales::nbad_index`. This class has the ability to store and report the label of a `Label` object. In addition, from `logic_error`, `nbad_index` also ultimately inherits the ability to throw.

Both classes have overloaded `operator[]` that return a reference to an individual array element stored in an object. The `bad_index` class also has a set of bounds.

Both the `bad_index` and `nbad_index` classes inherit from `logic_error`. The reason is that both eventually inherit from `std::exception`. The `bad_index` destructor uses the `throw()` exception specifier. In C++11, the exception destructor doesn't have a destructor.

Listing 15.15 shows the implementation of the `bad_index` class, nested inline in Listing 15.14. Note that nested classes can be used more than once. Also note that the `operator[]` can be used to access an array index is out of bounds.

```

 for (i = 0; i < 11m; ++i)
 gross[i] = gr[i];
 // for i > n and i < MONTHS
 for (; i < MONTHS; ++i)
 gross[i] = 0;
 }

double Sales::operator[] (int i) const
{
 if(i < 0 || i >= MONTHS)
 throw bad_index(i);
 return gross[i];
}

double & Sales::operator[] (int i)
{
 if(i < 0 || i >= MONTHS)
 throw bad_index(i);
 return gross[i];
}

LabeledSales::nbad_index::nbad_index(const
 const string & s) : Sales::bad
{

```

```
{
 if(i < 0 || i >= MONTHS)
 throw nbad_index(Label(), i);
 return Sales::operator[] (i);
}
```

---

Listing 15.16 uses the classes in a program array in the LabeledSales object sales2 and Sales object sales1. These attempts are made each kind of exception.

#### Listing 15.16    **use\_sales.cpp**

---

```
// use_sales.cpp -- nested exceptions
#include <iostream>
#include "sales.h"

int main()
{
 using std::cout;
 using std::cin;
 using std::endl;
```

```

 cout << sales1[i] << ' ';
 if (i % 6 == 5)
 cout << endl;
 }
 cout << "Year = " << sales2.Year()
 cout << "Label = " << sales2.Label()
 for (i = 0; i <= 12; ++i)
 {

 cout << sales2[i] << ' ';
 if (i % 6 == 5)
 cout << endl;
 }
 cout << "End of try block 1.\n";
}
catch(LabeledSales::nbad_index & bad)
{
 cout << bad.what();
 cout << "Company: " << bad.label_val()
 cout << "bad index: " << bad.bi_val()
}
catch(Sales::bad_index & bad)
{
 cout << bad.what();
}

```

```
}
```

---

Here is the output of the program in Listing 10.1:

```
First try block:
```

```
Year = 2011
```

```
1220 1100 1122 2212 1232 2334
```

```
2884 2393 3302 2922 3002 3544
```

```
Year = 2012
```

```
Label = Blogstar
```

```
12 11 22 21 32 34
```

```
28 29 33 29 32 35
```

```
Index error in LabeledSales object
```

```
Company: Blogstar
```

```
bad index: 12
```

```
Next try block:
```

```
Index error in Sales object
```

```
bad index: 20
```

```
done
```

```
terminate_handler set_terminate(terminate_
void terminate(); // C++98
void terminate() noexcept; // C++11
```

Here the typedef makes `terminate_handler` a type that has no arguments and no return value. The first argument, its argument, the name of a function (that is, `void`), is the `void` return type. It returns the address of the function. The second argument, the `set_terminate()` function more than once, is the most recent call to `set_terminate()`.

Let's look at an example. Suppose you want to print a message to that effect and then call `std::terminate()` with a value of 5. First, you include the exception handling header, which is available with a `using` directive or appropriate `std::` qualifier:

```
#include <exception>
using namespace std;
```

Next, you design a function that does the work. Here's a prototype:

```
void myQuit()
{
 cout << "Terminating due to uncaught exception";
}
```

It's good to know which exceptions to catch by default, aborts the program.

However, there's a bit more to the story. In addition to include exceptions thrown by functions called `Argh()` calls a `Duh()` function that can throw an exception. `Argh()` should appear in the `Argh()` exception specification. Unless you write all the functions in the library that this will get done correctly. You might look more closely at what happens if a function throws an exception specification. (It also suggests that the exception specification might be unwieldy, which is part of the reason for the new exception specification.)

The behavior is much like that for uncaught exceptions. If an exception is thrown and not caught, the program calls the unexpected handler. If the handler is not set, the program calls `unexpected()` function? No one expects the handler to be set. In turn, calls `terminate()`, which, by default, calls `set_terminate()` function that modifies the handler. `set_unexpected()` function that modifies the handler. Both functions are also declared in the exception handling header.

```
typedef void (*unexpected_handler)();
unexpected_handler set_unexpected(unexpected_handler);
unexpected_handler set_unexpected(unexpected_handler);
```



exception type and is declared in the e

- If the newly thrown exception does not  
and if the original exception specification  
type, the unmatched exception is replaced  
std::bad\_exception type.

In short, if you'd like to catch all exceptions  
thing like the following. First, you make sure the  
available:

```
#include <exception>
using namespace std;
```

Next, you design a replacement function that  
bad\_exception type and that has the proper p

```
void myUnexpected()
{
 throw std::bad_exception(); //or just
}
```

Just using throw without an exception cause  
However, the exception will be replaced with  
specification includes that type.

From the preceding discussion of using exceptions correctly) that exception handling should be designed. Doing this has some disadvantages, though. First, it adds and subtracts from the speed of a program. Even with templates because template functions might be specialized on the particular specialization used. Exceptions and templates don't always work that well together.

Let's look a little further at dynamic memory management in the following function:

```
void test1(int n)
{
 string msg("I'm trapped in an endless loop");
 ...
 if (oh_no)
 throw exception();
 ...
 return;
}
```

The `string` class uses dynamic memory allocation. When the function `test1` for `msg` would be called when the function returns, during stack unwinding, the `throw` statement, even though it fails, still allows the destructor to be called.

```

 try {
 if (oh_no)
 throw exception();
 }
 catch(exception & ex)
 {
 delete [] ar;
 throw;
 }
 ...
 delete [] ar;
 return;
 }

```

However, this clearly enhances the opportunity for error checking. Another solution is to use one of the smart pointers. “The string Class and the Standard Template Library”

In short, although exception handling is expensive, it does have costs in terms of programming effort, program size. On the other hand, the cost of no error checking can be much higher.

already provided ways to do so for their own support in C++, each vendor's mechanism is different from the others. Creating a language standard for RTTI compatible with each other.

## **What Is RTTI For?**

Suppose you have a hierarchy of classes descended from a base class. You have a base-class pointer to point to an object of a derived class. You want to call a function that, after processing some information, returns an object of that type, and returns its address. How can you tell what kind of object it points to?

Before answering this question, you need to know the type. Perhaps you want to invoke the correct method. In that case, you don't really need to know the object's type. The function possessed by all members of the class is the same. The object has an uninherited method. In that case, you need to know the object's type. Or maybe, for debugging purposes, you would like to know the objects were generated. For these last two cases,

particular type. Let's look at what that means.

```
class Grand { // has virtual methods};
class Superb : public Grand { ... };
class Magnificent : public Superb { ... };
```

Next, suppose you have the following pointers:

```
Grand * pg = new Grand;
Grand * ps = new Superb;
Grand * pm = new Magnificent;
```

Finally, consider the following type casts:

```
Magnificent * p1 = (Magnificent *) pm;
Magnificent * p2 = (Magnificent *) pg;
Superb * p3 = (Magnificent *) pm;
```

Which of these type casts are safe? Depending on the compiler, any of them could be safe, but the only ones guaranteed to be safe are those that cast to the same type as the object or else a direct or indirect base class. For example, Type Cast #1 is safe because it sets a type `Magnificent` pointer to a `Magnificent` object. Type Cast #2 is not safe because it sets a `Magnificent` pointer to a `Grand` object. Type Cast #3 is not safe because it sets a `Superb` pointer to a `Magnificent` object. A `Grand` object (base-class object) to a derived-class pointer (e.g., `Magnificent` pointer) is not safe because the base-class object to have derived-class properties. A `Magnificent` object, for example, might have

pointed-to object (\*pt) is of type Type or else

```
dynamic_cast<Type *>(pt)
```

Otherwise, the expression evaluates to 0, the

Listing 15.17 illustrates the process. First it defines the `Grand`, `Superb`, and `Magnificent` classes, each of which defines a `Speak()` function, which `Magnificent` redefines (see Figure 15.17). The `main()` function that randomly creates and initializes five objects and returns the address as a type `Grand *` pointer. (The user is the active user making decisions.) A loop assigns the pointer to `pg` and then uses `pg` to invoke the `Speak()` function. The code correctly invokes the `Speak()` version of the

```
for (int i = 0; i < 5; i++)
{
 pg = GetOne();
 pg->Speak();
 ...
}
```

You can't use this exact approach (using a pointer to `Superb` for the condition; it's not defined for the `Grand` class. However, you can use `dynamic_cast` to see if `pg` can be type cast to a pointer to `Superb` or `Magnificent`. In either case, you can then use the `if` statement to see if `ps` is a pointer to `Superb`.

```
if (ps = dynamic_cast<Superb *>(pg))
 ps->Say();
```

Recall that the value of an assignment expression is the value of the right-hand side. Thus, the value of the `if` condition is `ps`. If the type cast fails, `ps` is `0`. If the type cast fails, which it will if `pg` points to a `Grand` object, `ps` is `0`. Listing 15.17 shows the full code. (By the way, some programmers use the `==` operator in an `if` statement condition, but this is not a good idea.)

### Listing 15.17 `rtti1.cpp`

---

```
// rtti1.cpp -- using the RTTI dynamic_cast
#include <iostream>
#include <cstdlib>
#include <ctime>

using std::cout;
```

```

public:
 Magnificent(int h = 0, char c = 'A')
 void Speak() const {cout << "I am a m
 void Say() const {cout << "I hold the
 " and the integer " << Va
};

```

```

Grand * GetOne();

```

```

int main()
{
 std::srand(std::time(0));
 Grand * pg;
 Superb * ps;
 for (int i = 0; i < 5; i++)
 {
 pg = GetOne();
 pg->Speak();
 if (ps = dynamic_cast<Superb *>(p
 ps->Say();
 }
 return 0;
}

```



```
functions when possible and RTTI only when
I am a superb class!!
I hold the superb value of 68!
I am a magnificent class!!!
I hold the character R and the integer 68!
I am a magnificent class!!!
I hold the character D and the integer 12!
I am a magnificent class!!!
I hold the character V and the integer 59!
I am a grand class!
```

As you can see, the `Say()` methods were in the classes. (The output will vary from run to run depending on the object type.)

You can use `dynamic_cast` with reference types. If there is no reference value corresponding to the null pointer, `dynamic_cast` returns a null pointer value that can be used to indicate failure. Instead, `dynamic_cast` throws a `type_bad_cast` exception if the object is not of the target class and defined in the `typeinfo` header file. Here is an example where `rg` is a reference to a `Grand` object:

```
#include <typeinfo> // for bad_cast
...
try {
```

exception type is derived from the exception header file.

The implementation of the `typeid` class `name()` member that returns an implementation (not necessarily) the name of the class. For example, the string defined for the class of the object to which it is applied.

```
cout << "Now processing type " << typeid(
```

Listing 15.18 modifies Listing 15.17 so that it uses the `typeid` member function. Note that they are used for debugging purposes only. The `typeid` test is used for debugging functions don't handle. The `typeid` test is used for debugging method, so it can't be invoked by a class pointer. Note how the method can be used in debugging. Note the header file.

#### Listing 15.18 `rtti2.cpp`

---

```
// rtti2.cpp -- using dynamic_cast, typeid
#include <iostream>
#include <cstdlib>
#include <ctime>
#include <typeinfo>
using namespace std;
```

```

public:
 Magnificent(int h = 0, char cv = 'A')
 void Speak() const {cout << "I am a ma
 void Say() const {cout << "I hold the
 " and the integer " << Val
};

Grand * GetOne();

int main()
{
 srand(time(0));
 Grand * pg;
 Superb * ps;
 for (int i = 0; i < 5; i++)
 {
 pg = GetOne();
 cout << "Now processing type " <<
 pg->Speak();
 if(ps = dynamic_cast<Superb *>(pg)
 ps->Say();
 if (typeid(Magnificent) == typeid(
 cout << "Yes, you're really ma
 }
 return 0;

```

```
Yes, you're really magnificent.
Now processing type Superb.
I am a superb class!!
I hold the superb value of 37!
Now processing type Grand.
I am a grand class!
Now processing type Superb.
I am a superb class!!
I hold the superb value of 18!
Now processing type Grand.
I am a grand class!
```

As with the preceding example, the exact program uses `rand()` to select types. Also soon when `name()` is called, for example, `5Grand` is

## Misusing RTTI

RTTI has many vocal critics within the C++ community, a potential source of program inefficiency and bad programming practices. Without delving into the details of programming that you should avoid.

```

 pm->Speak();
 }
 else if (typeid(Superb) == typeid(*pg))
 {
 ps = (Superb *) pg;
 ps->Speak();
 ps->Say();
 }
 else
 pg->Speak();
}

```

Not only is this uglier and longer than the other version, it also requires each class explicitly. Suppose, for example, that we have a class `Insufferable` derived from the `Magnificent` class. We would have to modify the `for` loop code, add a `break` statement, and so on. The version, however, requires no changes at all. The code is derived from `Grand`:

```
pg->Speak();
```

And this statement works for all classes derived from `Grand`:

```

if (ps = dynamic_cast<Superb *>(pg))
 ps->Say();

```

implausible, none of them make much sense. allowed? In C, all of them are. Stroustrup's re is allowable for a general type cast and to add discipline for the casting process:

```
dynamic_cast
const_cast
static_cast
reinterpret_cast
```

Instead of using a general type cast, you can use a cast of a particular purpose. This documents the intended result and gives you a chance to check that you did what you thought you did.

You've already seen the `dynamic_cast` operator. Suppose `ph` and `p1` are two classes, that `ph` is type `High *`, and that `p1` is type `Low *`. The statement assigns a `Low *` pointer to `p1` only if `ph` is (or indirectly) to `High`:

```
p1 = dynamic_cast<Low *> ph;
```

Otherwise, the statement assigns the null pointer to `p1`. This is this syntax:

```
dynamic_cast < type-name > (expression)
```

declare the value as const and use const\_cast could be done using the general type cast, but this would not allow us to simultaneously change the type:

```
High bar;
const High * pbar = &bar;
...
High * pb = (High *) (pbar); // valid
Low * pl = (Low *) (pbar); // also valid
```

Because the simultaneous change of type and value is a programming slip, using the const\_cast operator is not recommended.

The const\_cast is not all powerful. It cannot change the effect of attempting to change a quantity that is declared const. We can clarify this statement with the short example

#### Listing 15.19 `constcast.cpp`

---

```
// constcast.cpp -- using const_cast<>
#include <iostream>
using std::cout;
using std::endl;
```

The `const_cast` operator can remove the restriction that prevents the compiler to accept the following statement:

```
*pc += n;
```

However, because `pop2` is declared as `const`, you cannot change it, as is shown by the following sample:

```
pop1, pop2: 38383, 2000
pop1, pop2: 38280, 2000
```

As you can see, the calls to `change()` alter `pop1`, but `pop2` remains unchanged. This is because `change()` is declared as `const int *`, so it can only be pointed-to `int`. The pointer `pc` has the `const` attribute, so it can be pointed-to `value`, but only if that value wasn't `const`. You can alter `pop1` but not `pop2`.

The `static_cast` operator has the same syntax:

```
static_cast < type-name > (expression)
```

It's valid only if `type_name` can be converted to the type of `expression` has, or vice versa. Otherwise, the conversion is invalid. For example, the conversion from `Low` to `High` is valid, but a conversion from `High` to `Low` is not, and `Low` to `High` are valid, but a conversion from `High` to `Low` is not.

```
High bar;
Low blow;
...
```



```
dat * pd = reinterpret_cast< dat *> (&value);
cout << hex << pd->a; // display first 2
```

Typically, such type casts would be used for programming and would not be portable. For example, a multibyte value in a different order than does

The `reinterpret_cast` operator doesn't allow you to cast a pointer type to an integer type or an integer representation, but you can't cast a pointer to a pointer type. Another restriction is that you can't cast a pointer or vice versa.

The plain type cast in C++ is also restricted. It can do the type casts that C can do, plus some combinations, such as `static_cast<int*>(&1)` followed by a `const_cast`, but it can't do anything that is allowed in C but, typically, not in C++ because the `char` type is too small to hold a pointer implementation.

```
char ch = char (&d); // type cast
```

These restrictions make sense, but if you find them restrictive, you still have C available.

The RTTI features allow a program to determine the dynamic type of an object. The `dynamic_cast` operator is used to cast a derived object to its base class. The main purpose is to ensure that it's okay to invoke a virtual function. The `dynamic_cast` operator returns a `type_info` object. Two type codes are used to determine whether an object is of a specific type. The `dynamic_cast` can be used to obtain information about an object.

The `dynamic_cast`, `static_cast`, `const_cast`, and `reinterpret_cast` provide safer, better-documented type casts than the C-style casts.

## Chapter Review

1. What's wrong with the following attempt?

```
a. class snap {
 friend clasp;
 ...
};
class clasp { ... };
```

b. class cuff {  
 public:  
 void snip(muff &) { ...  
 ...  
};

```

 {
private:
 class Sauce
 {
 int soy;
 int sugar;
 public:
 Sauce(int s1, int s2) : soy(s1), sugar(s2) {}
 };
 ...
};

```

4. How does throw differ from return?
5. Suppose you have a hierarchy of exception classes. In what order should you catch them?
6. Consider the Grand, Superb, and Magnificent classes. Suppose pg is a type Grand \* pointer that is a pointer to one of these three classes and that ps is a type Superb \* pointer. How would you know the following two code samples both work?
 

```

if (ps = dynamic_cast<Superb *>(pg))
 ps->say(); // sample #1

```

objects need not hold the bad values; t

3. This exercise is the same as Programming Exercise 15.15. The function `label_val()` should be derived from a base class (its arguments are two argument values, the exceptions should be thrown as well as the function name, and a single exception should be used for both exceptions). The loop to terminate.
4. Listing 15.16 uses two catch blocks after the `try` block. An exception leads to the `label_val()` method. That it uses a single catch block after the `try` block. Invoking `label_val()` only when app



- Container classes
- Iterators
- Function objects (functors)
- STL algorithms
- The `initializer_list` template

By now you are familiar with the C++ `iostream` and `fstream` classes, when you can reuse code written by others. There are many commercially available C++ class libraries, some of which are available as part of the C++ package. For example, you can use the `ostream` class as part of the `ostream` header file. This chapter will show you how to use your programming pleasure.

You've already encountered the `string` class in Chapter 1. This chapter looks at `string` classes. Then the chapter looks at "smart pointers" and how to use them to manage dynamic memory a bit easier. Next, the chapter looks at the Standard Template Library (STL), a collection of useful templates for handling containers. The STL exemplifies the programming paradigm of templates. This chapter looks at the `initializer_list` template and how to use it using initializer-list syntax with STL objects.

constructors used in Listing 16.1, in that order. The constructor representations are simplified; really is a typedef for a template specialization, an optional argument relating to memory management, chapter and in Appendix F, “The `string` Template Implementation—dependent integral type definition defines `string::npos` as the maximum possible value equal the maximum value of an unsigned integer. The function `NBTS` for null-byte-terminated string—terminated with a null character.

Table 16.1 **`string` Class Constructors**

| Constructor                                 |        |
|---------------------------------------------|--------|
| <code>string(const char * s)</code>         | l<br>t |
| <code>string(size_type n, char c)</code>    | C<br>i |
| <code>string(const string &amp; str)</code> | l<br>s |

**Listing 16.1    str1.cpp**

```
// str1.cpp -- introducing the string class
#include <iostream>
#include <string>
// using string constructors

int main()
{
 using namespace std;
 string one("Lottery Winner!"); //
 cout << one << endl; //
 string two(20, '$'); //
 cout << two << endl;
 string three(one); //
 cout << three << endl;
 one += " Oops!"; //
 cout << one << endl;
 two = "Sorry! That was ";
 three[0] = 'P';
 string four; //
```



```
Lottery Winner!
Lottery Winner! Oops!
Sorry! That was Pottery Winner!
All's well that ends!
well, well...
That was Pottery in motion!
```

## Program Notes

The start of the program in Listing 16.1 illustrates how to convert a C++ string object to a regular C-style string and display it by using the `c_str()` member function.

```
string one("Lottery Winner!"); // ctor
cout << one << endl; // overload
```

The next constructor initializes the string object with a specific number of characters:

```
string two(20, '$'); // ctor
```

The copy constructor initializes the string object with a copy of another string object:

```
string three(one); // ctor
```

The overloaded `+=` operator appends the contents of one string object to the end of another string object:

```
one += " Oops!"; // overload
```

might expect, the `+` operator concatenates its operands. The operator is multiply overloaded, so the same expression can create a C-style string or a `char` value.

The fifth constructor takes a C-style string and an integer indicating how many characters to copy:

```
char alls[] = "All's well that ends well";
string five(alls,20); // ctor 5
```

Here, as the output shows, just the first 20 characters of `alls` are used to initialize the `five` object. As Table 16-1 shows, the length of the C-style string, the requested number of characters, and the length of the resulting string are all 20. Using 20 with 40 in the preceding example would have resulted in a string of 40 characters at the end of `five`. (That is, the constructor would have copied the following the string "All's well that ends well" twice.)

The sixth constructor has a template argument and two iterator arguments:

```
template<class Iter> string(Iter begin, Iter end);
```

The intent is that *begin* and *end* act like pointers to the beginning and end of a range of memory. (In general, *begin* and *end* can be iterators, which are used in the STL.) The constructor then uses the iterators pointed to by *begin* and *end* to initialize the `string` object. The word *range*, borrowed from mathematics, means the range of values between *begin* and *end*.

Now suppose you want to use this constructor with a string object—say, the object `five`. The following code is correct:

```
string seven(five + 6, five + 10);
```

The reason is that the name of an object, `five`, is not the address of an object, hence `five` is not a pointer. However, `five[6]` is a `char` value, so `&five[6]` is a pointer to the constructor:

```
string seven(&five[6], &five[10]); // ctor
```

The seventh constructor copies a portion of an object:

```
string eight(four, 7, 16); // ctor
```

This statement copies 16 characters from `four` (starting at the eighth character) in `four`.

## C++11 Constructors

The `string(string && str) noexcept` constructor guarantees that the new `string` is a copy of `str`. However, the compiler does not guarantee that `str` will be treated as `const`. This is a *const constructor*. The compiler can use it in some situations.

```
string stuff;
cin >> stuff; // read a word
getline(cin, stuff); // read a line,
```

Both versions of `getline()` allow for an character to use to delimit input:

```
cin.getline(info, 100, ':'); // read up
getline(stuff, ':'); // read up
```

The main operational difference is that the `string` object to hold the input character

```
char fname[10];
string lname;
cin >> fname; // could be a problem
cin >> lname; // can read a very, ve
cin.getline(fname, 10); // may truncate
getline(cin, fname); // no truncation
```

The automatic sizing feature allows the `string` object to use the numeric parameter that limits the number of characters read.

A design difference is that the C-style `string` versions use the `istream` class, whereas the `string` versions use the `string` class. The `istream` class is the invoking object for C-style string input and output. This applies to the `>>` form, too, which is evident in the following examples.

the stream. In this system, setting `eofbit` registers failbit registers detecting an input error; setting `failbit` registers detecting a hardware failure; and setting `badbit` registers detecting a software failure. Chapter 17, “Input, Output, and Files,” discusses this further.

The operator `>>()` function for the `string` object is used for reading to and discarding a delimiting character and leaves that character in the input queue. A space or tab character or more generally, any character.

So far in this book, you’ve seen several examples of input functions for `string` objects work with files. You can also use them for file input. Listing 16.2 shows an example of the file. It assumes that the file contains strings separated by the `getline()` method of specifying a delimiter. It reads one string to an output line.

#### Listing 16.2 **strfile.cpp**

---

```
// strfile.cpp -- read strings from a file
#include <iostream>
#include <fstream>
#include <string>
#include <cstdlib>
int main()
```

Here is a sample `tobuy.txt` file:

```
sardines:chocolate ice cream:pop corn:leeks
cottage cheese:olive oil:butter:tofu:
```

Typically, for the program to find the text file as the executable program or sometimes you can provide the full path name. On a Windows string the escape sequence `\\` represents a single backslash. The following code opens the file `C:\\CPP\\Progs\\tobuy.txt`:

```
fin.open("C:\\CPP\\Progs\\tobuy.txt"); //
```

Here is the output of the program in Listing 10.1:

```
1: sardines
2: chocolate ice cream
3: pop corn
4: leeks
5:
cottage cheese
6: olive oil
7: butter
8: tofu
9:
```

Done

```

if (snake1 < snake 2) // operator
...
if (snake1 == snake3) // operator
...
if (snake3 != snake2) // operator
...

```

You can determine the size of a string. Both `length()` and `size()` return the number of characters in a string.

```

if (snake1.length() == snake2.size())
 cout << "Both strings have the same length."

```

Why two functions that do the same thing? There are two versions of the `string` class, and `size()` was added to the C++ standard library in C++11.

You can search a string for a given substring. `string::npos` is the maximum possible number of characters in a string, which is the largest unsigned `int` or unsigned `long` value on the system.

---

The `string` library also provides the related `find_last_of()`, `find_first_not_of()`, and a set of overloaded function signatures as the first, the last occurrence of a substring or character in the first occurrence in the invoking string of any example, the following statement would return the index 3) because that's the first occurrence

```
int where = snake1.find_first_of("hark");
```

The `find_last_of()` method works the same way. Thus, the following statement would return

```
int where = snake1.last_first_of("hark");
```

The `find_first_not_of()` method finds the first character that is not a character in the argument. So the following would return the index 0 because `c` is not found in `hark`:

```
int where = snake1.find_first_not_of("hark");
```

(You'll learn about `find_last_not_of()` in the next chapter.)



```
fluid, quire, remove, scold,
"valid", "whence", "xenon", "yearn",
```

```
int main()
{
 using std::cout;
 using std::cin;
 using std::tolower;
 using std::endl;
 std::srand(std::time(0));
 char play;
 cout << "Will you play a word game? <y
 cin >> play;
 play = tolower(play);
 while (play == 'y')
 {
 string target = wordlist[std::rand
 int length = target.length();
 string attempt(length, '-');
 string badchars;
 int guesses = 6;
 cout << "Guess my secret word. It
 << " letters, and you guess\n"
 << "one letter at a time. You
```

```

 attempt[loc]=letter;
 // check if letter appears
 loc = target.find(letter,
 while (loc != string::npos)
 {
 attempt[loc]=letter;
 loc = target.find(let
 }
 }
 cout << "Your word: " << atte
 if (attempt != target)
 {
 if (badchars.length() > 0)
 cout << "Bad choices:
 cout << guesses << " bad
 }
}
if (guesses > 0)
 cout << "That's right!\n";
else
 cout << "Sorry, the word is "

```

Your word: a--a--

Bad choices: e

5 bad guesses left

Guess a letter: **t**

Oh, bad guess!

Your word: a--a--

Bad choices: et

4 bad guesses left

Guess a letter: **r**

Good guess!

Your word: a--ar-

Bad choices: et

4 bad guesses left

Guess a letter: **y**

Good guess!

Your word: a--ary

Bad choices: et

4 bad guesses left

Guess a letter: **i**

Good guess!

Your word: a-iary

Bad choices: et

4 bad guesses left

Guess a letter: **p**

failure to find a character or a string.

The program makes use of the fact that operator lets you append individual characters to

```
badchars += letter; // append a char to
```

The heart of the program begins by checking every word:

```
int loc = target.find(letter);
```

If `loc` is a valid value, the letter can be placed in the answer string:

```
attempt[loc]=letter;
```

However, a given letter might occur more than once. The program has to keep checking. The program uses `find()` which specifies a starting place in the string for the search. If the letter was found at location `loc`, the next search starts at `loc + 1`. It keeps the search going until no more occurrences are found. `find()` indicates failure if `loc` is after the end of the string.

```
// check if letter appears again
```

```
loc = target.find(letter, loc + 1);
```

```
while (loc != string::npos)
```

```
{
```

exceeds that size, the program allocates a new block, grows without continuous resizing. The capacity() method returns the current block size, and the reserve() method allows you to reserve a new block. Listing 16.4 shows an example that uses capacity().

#### Listing 16.4    **str2.cpp**

---

```
// str2.cpp -- capacity() and reserve()
#include <iostream>
#include <string>
int main()
{
 using namespace std;
 string empty;
 string small = "bit";
 string larger = "Elephants are a girl's best friend";
 cout << "Sizes:\n";
 cout << "\tempty: " << empty.size() << endl;
 cout << "\tsmall: " << small.size() << endl;
 cout << "\tlarger: " << larger.size() << endl;
 cout << "Capacities:\n";
 cout << "\tempty: " << empty.capacity() << endl;
 cout << "\tsmall: " << small.capacity() << endl;
 cout << "\tlarger: " << larger.capacity() << endl;
 empty.reserve(50);
}
```

```
string filename;
cout << "Enter file name: ";
cin >> filename;
ofstream fout;
```

The bad news is that the `open()` method news is that the `c_str()` method returns a pointer to the contents as the invoking string object. So you can write

```
fout.open(filename.c_str());
```

## String Varieties

This section treats the `string` class as if it were mentioned earlier, the string library really is based on the `basic_string` template

```
template<class charT, class traits = char_traits<charT>,
 class Allocator = allocator<charT>>
basic_string {...};
```

The `basic_string` template comes with a typedef name:

```
typedef basic_string<char> string;
typedef basic_string<wchar_t> wstring;
typedef basic_string<char16_t> u16string;
typedef basic_string<char32_t> u32string;
```

```
 str = ps;
 return;
}
```

You probably see its flaw. Each time the function allocates memory on the heap but never frees the memory, thus creating a memory leak. A simple solution—just remember to free the allocated memory just before the return statement:

```
delete ps;
```

However, a solution involving the phrase “return” is not a solution. Sometimes you won’t remember. Or you’ll forget to comment out the code. And even if you do remember, you’ll consider the following variation:

```
void remodel(std::string & str)
{
 std::string * ps = new std::string(str);
 ...
 if (weird_thing())
 throw exception();
 str = *ps;
 delete ps;
 return;
}
```





#1: Creates storage for ap and a

ap     10080  
6000

#2: Copies value into dynamic

ap     10080  
6000

#3: Discards ap, and ap's destr

Figure 16.2    A regular p

Here `new double` is a pointer returned by  
It is the argument to the `auto_ptr<double>`  
corresponding to the formal parameter `p` in th  
an actual argument for a constructor. The oth

```
unique_ptr<double> pdu(new double); // p
shared_ptr<string> pss(new string); // ps
```

C++11 `shared_ptr` and `weak_ptr` classes  
the pointer expires when execution leaves the scope  
ods to report when an object is created or deleted

### Listing 16.5 `smrtptrs.cpp`

```
// smrtptrs.cpp -- using three kinds of smart pointers
// requires support of C++11 shared_ptr and weak_ptr
#include <iostream>
#include <string>
#include <memory>

class Report
{
private:
 std::string str;
public:
 Report(const std::string s) : str(s)
 { std::cout << "Object created\n"; }
 ~Report() { std::cout << "Object deleted\n"; }
 void comment() const { std::cout << str << "\n"; }
};

int main()
{
```

```
Object created!
using unique_ptr
Object deleted!
```

Each of these classes has an explicit constructor. Thus, there is no automatic type cast from a pointer

```
shared_ptr<double> pd;
double *p_reg = new double;
pd = p_reg; // not allowed
pd = shared_ptr<double>(p_reg); // allowed
shared_ptr<double> pshared = p_reg; // not allowed
shared_ptr<double> pshared(p_reg); // allowed
```

The smart pointer template classes are defined so that the object acts like a regular pointer. For example, you can dereference it (`*ps`), use it to access structures, or to a regular pointer that points to the same type. You can also point an object to another of the same type, but that raises a question

But first, here's something you should avoid:

```
string vacation("I wandered lonely as a cloud");
shared_ptr<string> pvac(&vacation); // Not allowed
```

When `pvac` expires, the program would attempt to delete the memory, which is wrong.

the object. Then have assignment transfer ownership to `auto_ptr` and for `unique_ptr`, although

- Create an even smarter pointer that keeps track of the number of pointers pointing to a particular object. This is called *reference counting*. It increases the count by one, and the expected lifetime by one. Only when the final pointer expires does the object get destroyed. This is the `shared_ptr` strategy.

The same strategies we've discussed for `auto_ptr` and `weak_ptr` also apply to copy constructors.

Each approach has its uses. Listing 16.6 shows a strategy that is poorly suited.

#### Listing 16.6 `fowl.cpp`

---

```
// fowl.cpp -- auto_ptr a poor choice
#include <iostream>
#include <string>
#include <memory>

int main()
{
 using namespace std;
 auto_ptr<string> films[5] =
```

Segmentation fault (core dumped)

The “core dumped” message should help find the problem. (The behavior for this sort of pointer encounter different behavior, depending on your compiler. The following statement transfers ownership from

```
pwin = films[2]; // films[2] loses ownership
```

That causes `films[2]` to no longer refer to the object with ownership of an object, it no longer provides a valid pointer. It goes to print the string pointed to by `films[2]`, which is entirely an unpleasant surprise.

Suppose you go back to Listing 16.6 but use a compiler that need a compiler that supports the C++11 standard and gives this output:

```
The nominees for best avian baseball film are
Fowl Balls
Duck Walks
Chicken Runs
Turkey Errors
Goose Eggs
The winner is Chicken Runs!
```

```
auto_ptr<string> p2;
p2 = p1;
```

When, in statement #3, `p2` takes over ownership of ownership. This, recall, is good because it prevents `p1` from trying to delete the same object. But it is bad because `p1` can no longer use `p1` because `p1` no longer points to valid data.

Now consider the `unique_ptr` equivalent:

```
unique_ptr<string> p3(new string("auto"));
unique_ptr<string> p4;
p4 = p3;
```

In this case, the compiler does not allow `p4` to be assigned `p3` because `p4` would not be pointing to valid data. Hence, `unique_ptr` is safer (less potential program crash) than `auto_ptr`.

But there are times when assigning one smart pointer to another can leave a dangerous orphan behind. Suppose we have this:

```
unique_ptr<string> demo(const char * s)
{
 unique_ptr<string> temp(new string(s));
 return temp;
}
```

source of mischief. Assignment #2 leaves no unique\_ptr constructor, which constructs a temporary. Ownership is transferred to pu3. This selective move is superior to auto\_ptr, which would allow both. That auto\_ptr is banned (by recommendation) for container objects, whereas unique\_ptr is allowed, is something along the lines of assignment #1 to #2. If you use auto\_ptr, you get a compile-time error. If the algorithm is assignment #2, that's okay, and the program proceeds. With unique\_ptr, you could lead to undefined behavior and mysterious crashes.

Of course, it could be possible that you read the assignment. The assignment is unsafe only if you use the auto\_ptr manner, such as dereferencing it. But you could assign a new value to it. C++ has a standard library function to assign one unique\_ptr to another. Here is an example of a demo() function, which returns a unique\_ptr.

```
using namespace std;
unique_ptr<string> ps1, ps2;
ps1 = demo("Uniquely special");
ps2 = move(ps1); // enable move semantics
ps1 = demo(" and more");
cout << *ps2 << *ps1 << endl;
```

you might have an STL container of pointers or assignment operations that will work with (you'll get a compiler warning) or `auto_ptr`. If the compiler doesn't offer `shared_ptr`, you can get a

If the program doesn't need multiple pointers, `unique_ptr` is a good choice for the return type for a function created by `new`. That way, ownership is transferred to the caller, and that smart pointer takes on the responsibility of deleting the object. `unique_ptr` objects in an STL container provide algorithms, such as `sort()`, that copy or assign objects without assuming the proper includes and using static memory. `unique_ptr` could be used in a program:

```
unique_ptr<int> make_int(int n)
{
 return unique_ptr<int>(new int(n));
}

void show(unique_ptr<int> & pi)
{
 cout << *pi << ' ';
}

int main()
{
 ...
}
```



ously owned by the `unique_ptr`.  
You would use `auto_ptr` in the same situation. If your compiler doesn't provide unique library `scoped_ptr`, which is similar.

## The Standard Template Library

The STL provides a collection of templates for containers, objects, and algorithms. A container is a unit, like a list, that holds other objects. STL containers are homogeneous; that is, they hold objects of the same type. STL algorithms are recipes for accomplishing particular tasks, such as finding a particular value in a list. Iterators are objects that let you move through a container. Pointers let you move through an array; they are a type of iterator. Functors are objects that act like functions; they are used in STL algorithms (including function names because a function is an object). The STL provides a way to construct a variety of containers, including arrays, lists, and maps. It also provides a variety of operations, including searching, sorting, and inserting.

Alex Stepanov and Meng Lee developed STL. They first implemented it in 1994. The ISO/ANSI C++ Standard adopted STL as a part of the C++ Standard. The STL is not a library; it is a standard. Naming. Instead, it represents a different programming paradigm. This makes STL interesting both in terms of

vector object, assign one vector object to another vector elements. To make the class generic, you can do what STL does, defining a vector template in the header file.

To create a vector template object, you use a data type to be used. Also the vector template constructor use an initialization argument to indicate how many elements to create.

```
#include <vector>
using namespace std;
vector<int> ratings(5); // a vector of 5 integers
int n;
cin >> n;
vector<double> scores(n); // a vector of n doubles
```

After you create a vector object, you can use the usual array notation for accessing individual elements.

```
ratings[0] = 9;
for (int i = 0; i < n; i++)
 cout << scores[i] << endl;
```

## Allocators Again

Like the string class, the various STL containers have a template argument that specifies what allocator object to use. The template begins like this:

```
using std::cout;
using std::endl;

vector<int> ratings(NUM);
vector<string> titles(NUM);
cout << "You will do exactly as told.
 << NUM << " book titles and your
int i;
for (i = 0; i < NUM; i++)
{
 cout << "Enter title #" << i + 1 << endl;
 getline(cin,titles[i]);
 cout << "Enter your rating (0-10): ";
 cin >> ratings[i];
 cin.get();
}
cout << "Thank you. You entered the following
 << "Rating\tBook\n";
for (i = 0; i < NUM; i++)
{
 cout << ratings[i] << "\t" << titles[i] << endl;
}

return 0;
}
```

---

dynamically allocated array. The next section discusses methods.

## Things to Do to Vectors

Besides allocating storage, what else can the `vector` containers provide? Certain basic methods, including inserting elements in a container, `swap()`, which exchanges the contents of two containers, `begin()`, which returns an iterator that refers to the first element, and `end()`, which returns an iterator that represents the end of the container.

What's an iterator? It's a generalization of a pointer. It can be an object for which pointer-like operations are defined (the `operator*()` and incrementing (for example, `++it`; you'll see later, generalizing pointers to iterators). The `vector` interface for a variety of container classes, including `vector`, wouldn't work. Each container class defines a `typedef` called `iterator`. For example, for the type `double` specialization of `vector`, you would write

```
vector<double>::iterator pd; // pd an iterator
```

Suppose `scores` is a `vector<double>` object. Here's how to declare and construct it:

```
vector<double> scores;
```

```
for (pd = scores.begin(); pd != scores.end();
 cout << *pd << endl;;
```

All containers have the methods just discussed. `push_back()` is a method that only some STL containers have. `push_back()` adds an element to the end of a vector. While `push_back()` adds an element so that the vector size increases to accommodate the new element, you can write code like the following:

```
vector<double> scores; // create an empty vector
double temp;
while (cin >> temp && temp >= 0)
 scores.push_back(temp);
cout << "You entered " << scores.size() << " elements\n";
```

Each loop cycle adds one more element to the vector. You can see the number of elements in the vector when you write the `cout` statement. The program runs as long as the program has access to sufficient memory. The program terminates when memory is necessary.

The `erase()` method removes a given range of elements from a container. The first argument defines the range to be removed. It's important to note that the STL defines ranges using two iterators. The first iterator defines the beginning of the range, and the second iterator is one beyond the end of the range.

**Note**

A range `[it1, it2)` is specified by two iterators, `it1` and `it2`, where `it1` is the first element, and `it2` is the element not including, `it2`.

An `insert()` method complements `erase()`. The first parameter gives the position ahead of which new elements are to be inserted. The second and third iterator parameters define the range to be inserted into another container object. For example, the following code inserts the elements of the `new_v` vector ahead of the first element of the `old_v` vector.

```
vector<int> old_v;
vector<int> new_v;
...
old_v.insert(old_v.begin(), new_v.begin(), new_v.end());
```

Incidentally, this is a case where having a `begin()` method makes it simple to append something to the end of a container. `old_v.insert(old_v.end(), new_v.begin(), new_v.end())` inserts `new_v` ahead of `old_v.end()`, meaning it's appended to the end of `old_v`.

```

vector<Review> books;
Review temp;
while (FillReview(temp))
 books.push_back(temp);
int num = books.size();
if (num > 0)
{
 cout << "Thank you. You entered the following books:\n";
 << "Rating\tBook\n";
 for (int i = 0; i < num; i++)
 ShowReview(books[i]);
 cout << "Reprising:\n";
 << "Rating\tBook\n";
 vector<Review>::iterator pr;
 for (pr = books.begin(); pr != books.end(); pr++)
 ShowReview(*pr);
 vector <Review> oldlist(books);
 if (num > 3)
 {
 // remove 2 items
 books.erase(books.begin() + 1, books.begin() + 3);
 cout << "After erasure:\n";
 for (pr = books.begin(); pr != books.end(); pr++)
 ShowReview(*pr);
 }
}

```

```

 std::cout << "Enter book rating: ";
 std::cin >> rr.rating;
 if (!std::cin)
 return false;
 // get rid of rest of input line
 while (std::cin.get() != '\n')
 continue;
 return true;
 }

void ShowReview(const Review & rr)
{
 std::cout << rr.rating << "\t" << rr.
}

```

---

Here is a sample run of the program in Li

```

Enter book title (quit to quit): The Cat
Enter book rating: 5
Enter book title (quit to quit): Candid C
Enter book rating: 7
Enter book title (quit to quit): Warriors
Enter book rating: 4
Enter book title (quit to quit): Quantum

```



|   |                  |
|---|------------------|
| 7 | Candid Canines   |
| 4 | Warriors of Wonk |
| 8 | Quantum Manners  |

## More Things to Do to Vectors

There are many things programmers commonly do to vectors, such as sort them, randomize the order, and so on. Does the STL support these common operations? No! The STL takes a different approach. For each container class, it defines a single `find` function for all container classes. This design philosophy is based on the fact that suppose you had 8 container classes and 10 operations. If you had a member function, you'd need 8 \* 10, or 80, separate member functions. In the STL approach, you'd need just 10 separate nonmember functions. If you defined a new container class, provided that it inherited from an existing one, you could use the existing 10 nonmember functions.

On the other hand, the STL sometimes defines a nonmember function for the same task as a member function. For example, there is a class-specific algorithm that is more efficient than the `swap` member function. Therefore, the vector `swap()` will be more efficient than the `swap` nonmember function.

random\_shuffle(begin(), end(), books.begin(), books.end());

Unlike `for_each`, which works with any container class allow random access, which the

The `sort()` function, too, requires that the container support random access in two versions. The first version takes two iterators to define the range by using the `<` operator defined for the type. For example, the following sorts the contents of a `vector`. The built-in `<` operator to compare values:

```
vector<int> coolstuff;
...
sort(coolstuff.begin(), coolstuff.end());
```

If the container elements are user-defined, you need to provide an `operator<()` function defined that works with the type. For example, you could sort a vector containing `Review` objects. Review member function or a nonmember function. If `Review` is a structure, its members are public, and a nonmember function is defined.

```
bool operator<(const Review & r1, const Review & r2)
{
 if (r1.title < r2.title)
 return true;
```

```

 else
 return false;
 }

```

With this function in place, you can use the `sort` function to sort a vector of `Review` objects in order of increasing rating. The following code sorts `books` in order of increasing rating:

```
sort(books.begin(), books.end(), WorseThan());
```

Note that the `WorseThan()` function does not change the ordering of `Review` objects. If two objects have the same rating, the function sorts by using the `rating` member. But if two objects have the same rating, `WorseThan()` treats them as equivalent. This is called *weak ordering*, and the second kind is called *strict weak ordering*. If `a < b` and `b < a` are false, then `a` and `b` must be identical. They might be identical, or they might just have the same rating. In the `WorseThan()` example, if two objects have the same rating member, the best you can say for strict weak ordering is that they are identical.

Listing 16.9 illustrates the use of these STL functions.

#### Listing 16.9 **vect3.cpp**

---

```

// vect3.cpp -- using STL functions
#include <iostream>
#include <string>

```

```

 << books.size() << " ratings\n";
 << "Rating\tBook\n";
 for_each(books.begin(), books.end(),
 [](const Review & r) {
 cout << r.rating << "\t" << r.title << "\n";
 });

 sort(books.begin(), books.end());
 cout << "Sorted by title:\n";
 for_each(books.begin(), books.end(),
 [](const Review & r) {
 cout << r.rating << "\t" << r.title << "\n";
 });

 sort(books.begin(), books.end(),
 [](const Review & r1, const Review & r2) {
 return r1.rating < r2.rating;
 });
 cout << "Sorted by rating:\n";
 for_each(books.begin(), books.end(),
 [](const Review & r) {
 cout << r.rating << "\t" << r.title << "\n";
 });

 random_shuffle(books.begin(), books.end());
 cout << "After shuffling:\n";
 for_each(books.begin(), books.end(),
 [](const Review & r) {
 cout << r.rating << "\t" << r.title << "\n";
 });
 }
 else
 {
 cout << "No entries. ";
 cout << "Bye.\n";
 return 0;
 }
}

bool operator<(const Review & r1, const Review & r2)
{

```

```

 std::cin >> rr.rating;
 if (!std::cin)
 return false;
 // get rid of rest of input line
 while (std::cin.get() != '\n')
 continue;
 return true;
 }

void ShowReview(const Review & rr)
{
 std::cout << rr.rating << "\t" << rr.t
}

```

---

Here's a sample run of the program in Listing 11.1:

```

Enter book title (quit to quit): The Cat W
Enter book rating: 8
Enter book title (quit to quit): The Dogs
Enter book rating: 6
Enter book title (quit to quit): The Wimps
Enter book rating: 3
Enter book title (quit to quit): Farewell
Enter book rating: 7

```

6           The Dogs of Dharma  
8           The Cat Who Can Teach You Weight  
Bye.

## The Range-Based for Loop (C++)

The range-based for loop, mentioned in Chapter 16, is designed to work with the STL. To review,

```
double prices[5] = {4.99, 10.99, 6.87, 7.99, 5.99};
for (double x : prices)
 cout << x << std::endl;
```

The contents of the parentheses for the for loop are the beginning and end iterators in a container and then the name of the container. The range-based for loop uses a named variable to access each container element. The range-based for loop is shown in Listing 16.9:

```
for_each(books.begin(), books.end(), ShowReview);
```

It can be replaced with the following range-based for loop:

```
for (auto x : books) ShowReview(x);
```

The compiler will use the type of books, which is vector, and the Review type Review, and the loop will pass each Review object to ShowReview.

## Why Iterators?

Understanding iterators is perhaps the key to make algorithms independent of the type of container used. The STL's generic approach.

To see why iterators are needed, let's look at a solution for two different data representations and a generic approach. First, let's consider a function that searches for a particular value. You could write the function as follows:

```
double * find_ar(double * ar, int n, const double &val)
{
 for (int i = 0; i < n; i++)
 if (ar[i] == val)
 return &ar[i];
 return 0; // or, in C++11, return nullptr;
}
```

If the function finds the value in the array, it returns a pointer to the value; otherwise, it returns the null pointer. This approach works through the array. You could use a template to make it more generic. Nonetheless, this algorithm is still tied to the array.

the linked list.

If you consider details of implementation, algorithms: One uses array indexing to move through start to start->p\_next. But broadly, the two move value with each value in the container in sequence.

The goal of generic programming in this context is that would work with arrays or linked lists or other. Should the function be independent of the data type, independent of the data structure of the container, independent of the data type stored in a container, independent of the process of moving through the values, independent of the generalized representation.

What properties should an iterator have in a short list:

- You should be able to dereference an iterator it refers. That is, if p is an iterator, \*p should be defined.
- You should be able to assign one iterator to another. The expression p = q should be defined.
- You should be able to compare one iterator to another. If p and q are iterators, the expressions p == q and p != q should be defined.



```

typedef double * iterator;
iterator find_ar(iterator begin, iterator
{
 iterator ar;
 for (ar = begin; ar != end; ar++)
 if (*ar == val)
 return ar;
 return end; // indicates val not found
}

```

For the `find_ll()` function, you can define operators:

```

struct Node
{
 double item;
 Node * p_next;
};

class iterator
{
 Node * pt;
public:
 iterator() : pt(0) {}
 iterator (Node * pn) : pt(pn) {}
}

```

```

 {
 iterator start;
 for (start = head; start != 0; ++start)
 if (*start == val)
 return start;
 return 0;
 }

```

This is very nearly the same as `find_ar()`. Both functions determine whether they've reached the end. The `find_ar()` function uses an iterator to one-past-the-end value stored in the final node. Remove that condition and the two functions are identical. For example, you could require that the container have an element after the last official element. That is, you could require that the list have a past-the-end element, and you could return that element as the past-the-end position. Then `find_ar()` and `find()` would be identical. Detecting the end of data and becoming identical to `find()` is a common end element move from making requirements to the container class.

The STL follows the approach just outlined. `deque`, and so on) defines an iterator type appropriate to the container. An iterator might be a pointer; for another, it might be a reference. The iterator will provide the needed operations.

```
list<double>::iterator pr;
for (pr = scores.begin(); pr != scores.end(); ++pr)
 cout << *pr << endl;
```

The only change is in the type declared for `pr`. The STL approach is to use appropriate iterators and designing the classes in a way that can reuse the same code for containers that have quite dissimilar internal structures.

With C++ automatic type deduction, you can write the same code with either the vector or the list:

```
for (auto pr = scores.begin(); pr != scores.end(); ++pr)
 cout << *pr << endl;
```

Actually, as a matter of style, it's better to avoid using iterators if possible. You should use an STL function, such as `for_each`, if available for you. Alternatively, use the C++11 range-based `for` loop:

```
for (auto x : scores) cout << x << endl;
```

So to summarize the STL approach, you start with an algorithm and a container. You express it in as general terms as possible, independent of container type. To make the general algorithm work, you design iterators that meet the needs of the algorithm. Finally, you place the design. That is, basic iterator properties and container properties are placed on the algorithm.

them) and can be compared for equality (using `==`). The `!=` operator, possibly overloading, should produce the same result as `iter1 != iter2` if and only if `iter1 == iter2` is false.

```
iter1 == iter2
```

is true, then the following is also true:

```
*iter1 == *iter2
```

Of course, these properties hold true for built-in iterators. The above requirements are guides for what you must do when you write an iterator class. Now let's look at other iterator classes.

## Input Iterators

The term *input* is used from the viewpoint of the container. The way the container provides the container to the program is considered input. So an *input* iterator is an iterator that provides values from a container. In particular, dereferencing an input iterator is used to read a value from a container, but it needn't be used to write to a container. Algorithms that require an input iterator are algorithms that require a container.

An input iterator has to allow you to access the first element in a container and increment the iterator to access the next element.

In short, you can use an input iterator for single-pass, read-only algorithms and an output operator for single-pass, write-only algorithms.

## Forward Iterators

Like input and output iterators, forward iterators traverse through a container. So a forward iterator can only dereference the element at a time. However, unlike input and output iterators, a sequence of values in the same order each time. For a forward iterator, you can still dereference the pointer to the same value. These properties make multiple passes over the data.

A forward iterator can allow you to both read and write to read it:

```
int * pirw; // read-write iterator
const int * pir; // read-only iterator
```

## Bidirectional Iterators

Suppose you have an algorithm that needs to traverse a container in both directions. For example, a reverse function could start with a pointer to the first element, decrement the pointer, and traverse the process. A bidirectional iterator has all the features of a forward iterator for the two decrement operators (prefix and postfix).

|            |                     |
|------------|---------------------|
| $a < b$    | True if $b - a > 0$ |
| $a > b$    | True if $b < a$     |
| $a \geq b$ | True if $!(a < b)$  |
| $a \leq b$ | True if $!(a > b)$  |

---

Expressions such as  $a + n$  are valid only if  $n$  is a non-negative integer, and  $a - n$  is valid only if  $n$  is a non-negative integer and  $a$  is not the container (including past-the-end).

## Iterator Hierarchy

You have probably noticed that the iterator `const_iterator` has all the capabilities of an input iterator and of a bidirectional iterator, and the iterator `bidirectional_iterator` has all the capabilities of a bidirectional iterator and of a random access iterator. And a random access iterator has all the capabilities of a random access iterator. Table 16.4 summarizes the main capabilities. Table 16.5 summarizes the main capabilities of the `const_iterator` and `reverse_iterator`.  $n$  is an integer.

Table 16.4 Iterator Capabilities

| Iterator Capability | Input | Output |
|---------------------|-------|--------|
| Dereferencing read  | Yes   | No     |
| Dereferencing write | No    | Yes    |

random access iterator, can be used just with random access characterizations. As mentioned earlier, each container has a name called `iterator`. So the `vector<int>` container has `vector<int>::iterator`. But the documentation says that all iterators are random access iterators. That, in turn, means any iterator type because a random access iterator is a `list<int>` class has iterators of type `list<int>::iterator`. doubly linked list, so it uses a bidirectional iterator. It has random access iterators, but it can use algorithms that require bidirectional iterators.

## Concepts, Refinements, and Models

The STL has several features, such as kinds of iterators, and the C++ language. That is, although you can design, say, a bidirectional iterator, you can't have the compiler restrict an iterator to be a random access iterator. The reason is that the forward iterator is a set of requirements that must be satisfied by an iterator class you've designed. A pointer is a primary pointer. An STL algorithm works with a pointer. A bidirectional iterator is a requirement. STL literature uses the word *concept*. There is an input iterator concept, a forward iterator concept, a bidirectional iterator concept, a random access iterator concept. If you do need iterators for, say, a container class you've designed, which include iterator templates for the standard STL algorithms.

The STL `sort()` function, recall, takes as element in a container and an iterator pointing just `Receipts`) is the address of the first element `+ SIZE`) is the address of the element following function call sorts the array:

```
sort(Receipts, Receipts + SIZE);
```

C++ guarantees that the expression `Receipts` in the array or one past-the-end. Thus, C++ pointers into an array, and this makes it possible arrays. Thus, the fact that pointers are iterators it possible to apply STL algorithms to ordinary algorithms to data forms of your own design, pointers (which may be pointers or objects) and past-

### **`copy()`, `ostream_iterator`, and `istream_iterator`**

The STL provides some predefined iterators. There is an algorithm called `copy()` for copying an algorithm is expressed in terms of iterators, so another or even from or to an array, because iterators. For example, the following copies an array



The `out_iter` iterator now becomes an input iterator, providing the information about the output stream being sent to the output stream. The second template argument specifies the character type used by the output stream (e.g., `wchar_t`.) The first constructor argument (the stream) is being used. It could also be a stream used for output. The second argument is a separator to be displayed after each iteration.

You could use the iterator like this:

```
*out_iter++ = 15; // works like cout << 15;
```

For a regular pointer, this would mean assigning the value 15 to the pointer and then incrementing the pointer. For the iterator, it means send 15 and then a string consisting of 15 spaces to the output stream by `cout`. Then it should get ready for the next iteration. This is done with `copy()` as follows:

```
copy(dice.begin(), dice.end(), out_iter);
```

This would mean to copy the entire range of the iterator to the output stream—that is, to display the contents of the stream.

Or you could skip creating a named iterator and use the `ostream_iterator` instead. That is, you could use the adapter like this:

```
copy(dice.begin(), dice.end(), ostream_iterator<int, char>(cout, " "));
```

contents to the output stream:

```
ostream_iterator<int, char> out_iter(cout);
copy(dice.begin(), dice.end(), out_iter);
```

Now suppose you want to print the contents of the dice array in reverse order (for forming time-reversal studies.) There are several ways to do this, but rather than wallow in them, let's go to one that does it in a simple way. There is a function called `rbegin()` that returns a reverse iterator. There is also a function `rend()` that returns a reverse iterator pointing to the end of the array. If you have a reverse iterator and you decrement it, a reverse iterator makes it decrement, you can print the contents backward:

```
copy(dice.rbegin(), dice.rend(), out_iter);
```

You don't even have to declare a reverse iterator.

## Note

Both `rbegin()` and `end()` return the same value (a reverse iterator). Since `reverse_iterator` is a specialization of `iterator`, you can use the same value (an iterator to the first element),

Reverse pointers have to make a special case for the first element. If you initialize `*dice` to `dice.rbegin()`. What should `*dice` point to? If you dereference that address, you shouldn't try to dereference that address. If you dereference the address of the first element, `copy()` stops one location before the end of the array.

```
 // copy from vector to output
 copy(dice.begin(), dice.end(), out_iter);
 cout << endl;
 cout <<"Implicit use of reverse iterator\n";
 copy(dice.rbegin(), dice.rend(), out_iter);
 cout << endl;
 cout <<"Explicit use of reverse iterator\n";
 vector<int>::reverse_iterator ri;
 for (ri = dice.rbegin(); ri != dice.rend(); ++ri)
 cout << *ri << ' ';
 cout << endl;

 return 0;
 }
}
```

---

Here is the output of the program in Listing 10.1:

```
Let the dice be cast!
6 7 2 9 4 11 8 7 10 5
Implicit use of reverse iterator.
5 10 7 8 11 4 9 2 7 6
Explicit use of reverse iterator.
5 10 7 8 11 4 9 2 7 6
```

`front_insert_iterator` inserts items at the beginning of the container, inserting new items in front of the location specified as an argument. All three of these iterators are models of the `Insertable` concept.

There are restrictions. A `back_insert_iterator` is only defined for container types that allow rapid insertion at the end. (Remember that the `vector` class is the only container type that allows rapid insertion at the end. See section “Container Concepts,” later in this chapter for more details.) The `vector` class qualifies. A `front_insert_iterator` is only defined for container types that allow constant time insertion at the beginning. The `list` class doesn’t qualify, but the `queue` class does. The `deque` class qualifies for both. Thus, you can use it to insert material at the beginning of a `deque`. The `front_insert_iterator` does so faster for the `deque` than the `insert` method.

**Tip**  
You can use an `insert_iterator` to convert a `vector` to a `list`. It inserts data.

These iterators take the container type as a template argument and the container identifier as a constructor argument. To insert data into a `vector<int>` container called `dice`, you use `dice.insert(dice.begin(), 1, 6)`. To use a `back_insert_iterator`, you use `back_insert_iterator<vector<int> > back_i(dice); back_i.insert(1, 6)`. To use a `front_insert_iterator`, you use `front_insert_iterator<vector<int> > f_i(dice); f_i.insert(1, 6)`.

The reason you have to declare the container type is that the `insert` method uses the appropriate container method. The `vector` class uses `push_back` to insert at the end, the `list` class uses `push_front` to insert at the beginning, and the `deque` class uses `push_front` to insert at the beginning and `push_back` to insert at the end.

```
{
 using namespace std;
 string s1[4] = {"fine", "fish", "fash", "fash"};
 string s2[2] = {"busy", "bats"};
 string s3[2] = {"silly", "singers"};
 vector<string> words(4);
 copy(s1, s1 + 4, words.begin());
 for_each(words.begin(), words.end(), cout << endl;

 // construct anonymous back_insert_iterator
 copy(s2, s2 + 2, back_insert_iterator<vector<string>>(words));
 for_each(words.begin(), words.end(), cout << endl;

 // construct anonymous insert_iterator object
 copy(s3, s3 + 2, insert_iterator<vector<string>>(words, words.begin()));

 for_each(words.begin(), words.end(), cout << endl;
 return 0;
}
```

---

predefined iterators multiply the capabilities

## Kinds of Containers

The STL has both container concepts and categories with names such as container, sequence container, and container adaptor. The 11 container types are deque, list, queue, priority\_queue, stack, vector, unordered\_map, unordered\_multimap, unordered\_multiset, unordered\_set, and bitset. (This category is not a container for dealing with data at the bit level.) The unordered\_\* containers are unordered\_map, unordered\_multimap, unordered\_multiset, and unordered\_set. The bitset container is bitset. The unordered\_\* containers move bitset from the container category into the container adaptor category. The concepts categorize the types, let's start with

## Container Concepts

No type corresponds to the basic container concept. It's sort of a virtual concept common to all the container classes. It's sort of virtual because the container classes don't actually implement it. In another way, the container concept lays down the requirements that all container classes must satisfy.

A *container* is an object that stores other objects. The objects stored may be objects in the OOP sense. Data stored in a container is *owned* by the container.

|                                 |                        |                                                                                       |
|---------------------------------|------------------------|---------------------------------------------------------------------------------------|
| <code>X::value_type</code>      | <code>T</code>         | The type of the elements stored in the container                                      |
| <code>X u;</code>               |                        | Creates a container <code>u</code> of type <code>X</code>                             |
| <code>X();</code>               |                        | Creates a container of type <code>X</code>                                            |
| <code>X u(a);</code>            |                        | Copy constructs a container <code>u</code> from <code>a</code><br><code>u == a</code> |
| <code>X u = a;</code>           |                        | Same effect as <code>X u(a);</code>                                                   |
| <code>r = a;</code>             | <code>X&amp;</code>    | Copy assigns <code>a</code> to <code>r</code><br><code>r == a</code>                  |
| <code>(&amp;a)-&gt;~X();</code> | <code>void</code>      | Applies destructor to <code>a</code>                                                  |
| <code>a.begin();</code>         | iterator               | Returns an iterator to the first element of <code>a</code>                            |
| <code>a.end();</code>           | iterator               | Returns an iterator to the end value of <code>a</code>                                |
| <code>a.size();</code>          | unsigned integral type | Returns the number of elements in <code>a</code><br><code>a.end()</code>              |
| <code>a.swap(b);</code>         | <code>void</code>      | Swaps contents of <code>a</code> and <code>b</code>                                   |

$a == b$  has linear complexity because the  $==$  element of the container. Actually, that is a w  
ferent sizes, no individual comparisons need

## Constant-Time and Linear-Time Complexity

Imagine a long, narrow box filled with large packages. The box is open at just one end. Suppose your task is to fetch the first package. This is a constant time task. Whether there are 10 packages or one at the end makes no difference.

Now suppose your task is to fetch the package at the other end. This is a linear time task. If there are 10 packages altogether, it takes 10 times longer. If there are 100 packages, it takes 100 times longer. Assuming that you are a tireless worker, the task will take 10 times longer than the first task.

Now suppose your task is to fetch an arbitrary package. The package you are supposed to get is the first one at the other end. The number of packages you have to move is still proportional to the distance, so the task still has linear-time complexity.

Replacing the long, narrow box with a similar one that is open at both ends to constant-time complexity because then you can go straight to the package and remove it without moving the others.



| Expression              | Return Type                 | Description                                                                                                                                  |
|-------------------------|-----------------------------|----------------------------------------------------------------------------------------------------------------------------------------------|
| <code>X u(rv);</code>   |                             | Move construct <code>u</code> from <code>rv</code> .<br>The value of <code>rv</code> is destroyed after the construction of <code>u</code> . |
| <code>X u = rv;</code>  |                             | Same as <code>X u(rv);</code>                                                                                                                |
| <code>a = rv;</code>    | <code>X&amp;</code>         | Move assign <code>a</code> from <code>rv</code> .<br>The value of <code>rv</code> is destroyed after the assignment of <code>a</code> .      |
| <code>a.cbegin()</code> | <code>const_iterator</code> | Returns a constant iterator to the first element of <code>a</code> .                                                                         |
| <code>a.cend()</code>   | <code>const_iterator</code> | Returns a constant iterator to the end of <code>a</code> .                                                                                   |

---

The difference between copy construction, move construction and move assignment on the one hand and copy assignment on the other is that a copy operation leaves the original unchanged, whereas a move operation transfers ownership without doing any copying. We will see how these operations can provide more efficient code than copy operations in the next section. We will discuss move semantics further.

Table 16.7 Sequence Requirements

| Expression                   | Return Type | Description         |
|------------------------------|-------------|---------------------|
| <code>x(n,t);</code>         |             | Declare             |
| <code>x(n, t)</code>         |             | Creates             |
| <code>x(a(i, j)</code>       |             | Declare<br>range [i |
| <code>x(i, j)</code>         |             | Creates<br>content  |
| <code>a.insert(p,t)</code>   | iterator    | Inserts             |
| <code>a.insert(p,n,t)</code> | void        | Inserts             |
| <code>a.insert(p,i,j)</code> | void        | Inserts             |
| <code>a.erase(p)</code>      | iterator    | Erases              |
| <code>a.erase(p,q)</code>    | iterator    | Erases              |
| <code>a.clear()</code>       | void        | Is the s            |

Because the deque, list, queue, priority queue, and stack are all models of the sequence concept, they



The iterator returned by the two methods. Recall that incrementing such an iterator causes it to move in reverse order.

The `vector` template class is the simplest type that should be used by default unless the particular virtues of the other types.

## **deque**

The `deque` template class (declared in the `deque` header, a type often called a *deque* (pronounced “deck”), in the STL, it’s a lot like a `vector` container, supporting constant-time insertion and removal at both ends, that inserting and removing items from the beginning and end are constant-time operations instead of being linear-time. Most operations take place at the beginning and end using a `deque` data structure.

The goal of constant-time insertion and removal makes the design of a `deque` object more complex than a `vector`. To offer random access to elements and linear-time insertion and removal of a sequence, the `vector` container should a

complete list of STL methods and functions, so the parameter is one you normally don't have to worry about.

Table 16.9    **Some `list` Member Functions**

| Function                                                       | Description                                                                                                                                |
|----------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------|
| <code>void merge(list&lt;T, Alloc&gt;&amp; x)</code>           | Merge <code>x</code> into the container. The merge can be sorted. Complexity: linear in the size of the container and <code>x</code> .     |
| <code>void remove(const T &amp; val)</code>                    | Remove all elements equal to <code>val</code> . Complexity: linear in the size of the container.                                           |
| <code>void sort()</code>                                       | Sort the container. Complexity: $O(n \log n)$ .                                                                                            |
| <code>void splice(iterator pos, list&lt;T, Alloc&gt; x)</code> | Insert the elements of <code>x</code> into the container at position <code>pos</code> . Complexity: linear in the size of <code>x</code> . |
| <code>void unique()</code>                                     | Remove duplicate elements. Complexity: linear in the size of the container.                                                                |

```
cout << "List one: ";
for_each(one.begin(),one.end(), outi
cout << endl << "List two: ";
for_each(two.begin(), two.end(), outi
cout << endl << "List three: ";
for_each(three.begin(), three.end(),
three.remove(2);
cout << endl << "List three minus 2s:
for_each(three.begin(), three.end(),
three.splice(three.begin(), one);
cout << endl << "List three after spl
for_each(three.begin(), three.end(),
cout << endl << "List one: ";
for_each(one.begin(), one.end(), outi
three.unique();
cout << endl << "List three after uni
for_each(three.begin(), three.end(),
three.sort();
three.unique();
cout << endl << "List three after sor
for_each(three.begin(), three.end(),
two.sort();
three.merge(two);
```

```
for (auto x : three) cout << x << " ";
```

The main difference between `insert()` and `splice()` is that `insert()` inserts elements from the original range into the destination, while `splice()` moves elements from the source range into the destination. Thus, after the contents of `three` are moved to `one`, `three` is empty. (The `splice()` method has additional prototypes for moving multiple ranges of elements.) The `splice()` method leaves its first argument (the source range) as an iterator to point to an element in `one`, that iterator is then moved to `three`, and `splice()` relocates it in `three`.

Notice that `unique()` only reduces adjacent duplicates. If the program executes `three.unique()`, `three` still contains three elements, but they weren't adjacent. But applying `sort()` and then `unique()` results in a single appearance.

There is a nonmember `sort()` function (Library `<algorithm>`) that works with iterators. Because the trade-off for rapid insertion is that it requires a lot of memory, use the nonmember `sort()` function with a `list` instead of a `vector`. A `list` version that works within the restrictions of the STL.

## The `list` Toolbox

The `list` methods form a handy toolbox. Suppose you have a `list` of `int`s. Adding lists to organize. You could sort each list, merge two lists, or remove multiple entries.

add an element to the rear of a queue, remove the values of the front and rear elements, check if a queue is empty. Table 16.10 lists these operations.

Note that `pop()` is a data removal method. To use a value from a queue, you first use `front()` to remove it from the queue.

Table 16.10 **queue Operations**

| Method                                 | Description                                     |
|----------------------------------------|-------------------------------------------------|
| <code>bool empty() const</code>        | Returns true if the queue is empty              |
| <code>size_type size() const</code>    | Returns the number of elements in the queue     |
| <code>T&amp; front()</code>            | Returns a reference to the front element        |
| <code>T&amp; back()</code>             | Returns a reference to the back element         |
| <code>void push(const T&amp; x)</code> | Inserts <code>x</code> at the rear of the queue |
| <code>void pop()</code>                | Removes the front element                       |

## **priority\_queue**

The `priority_queue` template class (declared in `<queue>`) is a container adaptor that supports the same operations as `queue`. It is a priority queue, that with `priority_queue`, the largest item gets removed first.



|                                        |                               |
|----------------------------------------|-------------------------------|
| <code>bool empty() const</code>        | Returns true if               |
| <code>size_type size() const</code>    | Returns the nu                |
| <code>T&amp; top()</code>              | Returns a refer               |
| <code>void push(const T&amp; x)</code> | Inserts <code>x</code> at the |
| <code>void pop()</code>                | Removes the e                 |

---

Much as with `queue`, if you want to use a `stack` to push a value, you use `push()` to push the value, and then you use `pop()` to

## **array (C++11)**

The `array` template class, introduced in Chapter 11, is not an STL container because it has a fixed size. Unlike a container, such as `push_back()` and `insert()`, `array` has functions that do make sense, such as `operator[]`. `array` uses many standard STL algorithms, such as `sort`.

## **Associative Containers**

An *associative container* is another refinement of a container. An associative container associates a value with a key and uses the key to find the value. The values could be structures representing employees, with fields for office number, home and work phones, health

more than one value with the same key. For a `multiset` object could hold, say 1, 2, 2, 2, 3, 5.

For the `map` type, the value type is different with only one value per key. The `multimap` type is associated with multiple values.

There's too much information about these containers (G does list the methods), so let's just look at an example that uses `multimap`.

## A set Example

The STL `set` models several concepts. It is an unordered container and the keys are unique, so it can hold no more than one of each. And `list`, `set` uses a template parameter to provide a comparison function:

```
set<string> A; // a set of string objects
```

An optional second template argument can be used to provide an object to be used to order the key. By default, `less` is used. Older C++ implementations may not provide an explicit template parameter:

```
set<string, less<string> > A; // older i
```

automatically satisfy the precondition for using the container be sorted. The `set_union()` function takes two iterators to define a range in one set, the second two define a range in the other. The third argument is an output iterator that identifies a location to store the result. For example, to display the union of sets A and B, you could write:

```
set_union(A.begin(), A.end(), B.begin(), B.end(),
 ostream_iterator<string, char>(cout, " "));
```

Suppose you want to place the result into a container. You would want the last argument to be an iterator to the container, like `C.begin()`, but that doesn't work for two reasons. First, `set` treats keys as constant values, so the iterator returned by `C.begin()` and can't be used as an output iterator. The second reason is that `set_union()`, like `copy()`, overwrites existing information. The container has to have sufficient space to hold the result of the union to satisfy that requirement. But the `insert_iterator` class was designed to solve these problems. Earlier you saw that it converts a `copy` operation into an `insert` operation. It's an iterator concept, so you can use it to write to a container. You can use the `insert_iterator` to copy information from one container to another. You pass the `insert_iterator` to copy information to as the last argument of the container and an iterator as arguments:

```
set_union(A.begin(), A.end(), B.begin(), B.end(),
 insert_iterator<set<string>>(C, C.begin()));
```

```
#include <string>
#include <set>
#include <algorithm>
#include <iterator>

int main()
{
 using namespace std;
 const int N = 6;
 string s1[N] = {"buffoon", "thinkers"}
 string s2[N] = {"metal", "any", "food"}

 set<string> A(s1, s1 + N);
 set<string> B(s2, s2 + N);

 ostream_iterator<string, char> out(co
 cout << "Set A: ";
 copy(A.begin(), A.end(), out);
 cout << endl;
 cout << "Set B: ";
 copy(B.begin(), B.end(), out);
 cout << endl;
```

```

 cout << endl;

 cout << "Showing a range:\n";
 copy(C.lower_bound("ghost"), C.upper_bound("ghost"),
 cout << endl;

 return 0;
}

```

---

Here is the output of the program in Listing 10.1:

```

Set A: buffoon can for heavy thinkers
Set B: any deliver elegant food for metal
Union of A and B:
any buffoon can deliver elegant food for heavy metal
Intersection of A and B:
for
Difference of A and B:
buffoon can heavy thinkers
Set C:
any buffoon can deliver elegant food for heavy metal
Set C after insertion:
any buffoon can deliver elegant food for ghost
Showing a range:
grungy heavy metal

```

```
make_pair<int, string> codes;
```

An optional third template argument can be used to specify an object to be used to order the key. By default, the container is used with the key type as its parameter. Older versions of the STL used the template parameter explicitly.

To keep information together, the actual value is stored as a single data type into a single pair. To do this, the STL provides a `pair` class for storing two kinds of values in a single container. The first datatype is the type of the stored data, the value type, `datatype>`. For example, the value type for the `map` is `pair<const int, string>`.

Suppose that you want to store city names with their corresponding codes. To fit the `codes` declaration, which uses an `int` for the key, the approach is to create a `pair` and then insert it into the `map`.  

```
pair<const int, string> item(213, "Los Angeles");
codes.insert(item);
```

Or you can create an anonymous `pair` object and insert it directly.  

```
codes.insert(pair<const int, string> (213, "Los Angeles"));
```

Because items are sorted by key, there's no need to specify the

tion feature, which allows you to simplify the

```
auto range = codes.equal_range(718);
cout << "Cities with area code 718:\n";
for (auto it = range.first; it != range.second; ++it)
 cout << (*it).second << endl;
```

Listing 16.14 demonstrates most of these techniques, and shows some of the code writing.

#### Listing 16.14 **multimap.cpp**

---

```
// multimap.cpp -- use a multimap
#include <iostream>
#include <string>
#include <map>
#include <algorithm>

typedef int KeyType;
typedef std::pair<const KeyType, std::string> Value;
typedef std::multimap<KeyType, std::string> Multimap;

int main()
{
```

```

 = codes.equal_range(718);
 cout << "Cities with area code 718:\n";
 for (it = range.first; it != range.second; ++it)
 cout << (*it).second << endl;

 return 0;
}

```

---

Here is the output of the program in Listing 12.1:

```

Number of cities with area code 415: 2
Number of cities with area code 718: 2
Number of cities with area code 510: 2
Area Code City
 415 San Francisco
 415 San Rafael
 510 Oakland
 510 Berkeley
 718 Brooklyn
 718 Staten Island
Cities with area code 718:
Brooklyn
Staten Island

```



```

{
private:
 double slope;
 double y0;
public:
 Linear(double sl_ = 1, double y_ = 0)
 : slope(sl_), y0(y_) {}
 double operator()(double x) {return y0
};

```

The overloaded `()` operator then allows you

```

Linear f1;
Linear f2(2.5, 10.0);
double y1 = f1(12.5); // right-hand side
double y2 = f2(0.4);

```

Here `y1` is calculated using the expression `f1(12.5)`, which is equivalent to the expression `10.0 + 2.5 * 0.4`. In the expression `f1(12.5)`, the `12.5` and the `slope` come from the constructor for the object `f1`, and the `10.0` comes from the member variable `y0`. The `operator()` method is called on the object `f1` to calculate the value of `y1`.

Remember the `for_each` function? It applies a function to each element in a range:

```

for_each(books.begin(), books.end(), ShowF

```

just as the STL defines concepts for containers.

- A *generator* is a functor that can be called repeatedly to produce a sequence of values.
- A *unary function* is a functor that can be applied to a single argument.
- A *binary function* is a functor that can be applied to two arguments.

For example, the functor supplied to `for_each` is applied to one container element at a time.

Of course, these concepts come with refinements.

- A unary function that returns a `bool` value (a *predicate*).
- A binary function that returns a `bool` value (a *comparison function*).

Several STL functions require predicate or comparison functions. Listing 16.9 uses a version of `sort()` that takes a comparison function.

```
bool WorseThan(const Review & r1, const Review & r2) {
 ...
 return r1.rating() < r2.rating();
}

sort(books.begin(), books.end(), WorseThan);
```

The `list` template has a `remove_if()` member function. It applies the predicate to each member in the list and removes for which the predicate returns `true`. For example, to remove elements greater than 100 from the list `three`:

objects to different cut-off values to be used in the algorithm. This illustrates the technique.

### Listing 16.15    **functor.cpp**

---

```
// functor.cpp -- using a functor
#include <iostream>
#include <list>
#include <iterator>
#include <algorithm>

template<class T> // functor class definition
class TooBig
{
private:
 T cutoff;
public:
 TooBig(const T & t) : cutoff(t) {}
 bool operator()(const T & v) { return
};

void outint(int n) {std::cout << n << " ";
```

```
 for_each(etccetera.begin(), etccetera.e
 cout << endl;
 return 0;
 }
```

---

One functor (`f100`) is a declared object, and an anonymous object created by a constructor call.

Listing 16.15:

Original lists:

```
50 100 90 180 60 210 415 88 188 201
```

```
50 100 90 180 60 210 415 88 188 201
```

Trimmed lists:

```
50 100 90 60 88
```

```
50 100 90 180 60 88 188
```

Suppose that you already have a template

```
template <class T>
```

```
bool tooBig(const T & val, const T & lim)
```

```
{
```

```
 return val > lim;
```

```
}
```

As noted in the listing, C++11's `initializer_list` can replace

```
int vals[10] = {50, 100, 90, 180, 60, 210, 10, 10, 10, 10};
list<int> yadayada(vals, vals + 10); // range [vals, vals + 10)
list<int> etcetera(vals, vals + 10);
```

with this:

```
list<int> yadayada = {50, 100, 90, 180, 60, 210, 10, 10, 10, 10};
list<int> etcetera {50, 100, 90, 180, 60, 210, 10, 10, 10, 10};
```

## Predefined Functors

The STL defines several elementary functors. `accumulate` takes two values and comparing two values for equality. `adjacent_find` finds functions that take functions as arguments. `all_of` tests if all elements of a function. It has two versions. The first version takes two arguments are iterators that specify a range in a container (using that approach.) The third is an iterator that specifies a range and a functor that is applied to each element in the range. The result. For example, consider the following:

```
const int LIM = 5;
double arr1[LIM] = {36, 39, 42, 45, 48};
```

But then you'd have to define a separate `plus` function to define a template, except that you don't have a `functional` (formerly `function.h`) header to include, including one called `plus<>()`.

Using the `plus<>` class for ordinary addition:

```
#include <functional>
...
plus<double> add; // create a plus<double> object
double y = add(2.2, 3.4); // using plus<double>
```

But it makes it easy to provide a function object:

```
transform(gr8.begin(), gr8.end(), m8.begin(),
```

Here, rather than create a named object, `transform` is used to construct a functor to do the adding. (The `transform` function; what's passed to `transform()` is the `plus<double>` functor.)

The STL provides functor equivalents for all the arithmetic operators. Table 16.12 shows the names for the functors used with the C++ built-in types or with any user-defined type corresponding operator.

---

## Caution

Older C++ implementations use the functor n

## Adaptable Functors and Function

The predefined functors in Table 16.12 are all concepts: adaptable generators, adaptable unar adaptable predicates, and adaptable binary pre

What makes a functor adaptable is that it c argument types and return type. The member `first_argument_type`, and `second_argumen` like. For example, the return type of a `plus<i plus<int>::result_type`, and this would be

The significance of a functor being adaptable is that it allows the use of function adapter objects, which assume the existence of a function with an argument that is an adaptable functor. To declare a variable that matches the function

simply using the `binder1st` class. You give a pointer to construct a `binder1st` object, and it returns an object that converts the binary function `multiplies()` to a unary function. 2.5. Just do this:

```
binder1st(multiplies<double>(), 2.5)
```

Thus, the solution to multiplying every element by 2.5 is this:

```
transform(gr8.begin(), gr8.end(), out,
 binder1st(multiplies<double>(), 2.5))
```

The `binder2nd` class is similar, except that it binds the second argument instead of to the first. It has a helper function `to_binder1st`.



```
 for_each(gr8.begin(), gr8.end(), Show);
 cout << endl;
 cout << "m8: \t";
 for_each(m8.begin(), m8.end(), Show);
 cout << endl;
```

```
 vector<double> sum(LIM);
 transform(gr8.begin(), gr8.end(), m8.begin(),
 plus<double>());
 cout << "sum:\t";
 for_each(sum.begin(), sum.end(), Show);
 cout << endl;
```

```
 vector<double> prod(LIM);
 transform(gr8.begin(), gr8.end(), prod.begin(),
 bind1st(multiplies<double>(), 1));
 cout << "prod:\t";
 for_each(prod.begin(), prod.end(), Show);
 cout << endl;
```

```
 return 0;
```

```
}
```

ranges to be processed and to identify where the object argument to be used as part of the data.

There are two main generic components that use templates to provide generic types. Second, a representation for accessing data in a container, a container that holds type `double` values in a linked list, or with a container that has a structure, such as is used by `set`. Because pointers and functions such as `copy()` can be used with ordinary arrays.

The uniform container design allows many different kinds. For example, you can use `copy` to copy a vector object, from a vector object to a `list` object. You can use `==` to compare different kinds of vector. This is possible because the overloaded `==` compares contents, so a deque object and a vector with the same content in the same order.

## Algorithm Groups

The STL divides the algorithm library into four groups:

- Nonmodifying sequence operations
- Mutating sequence operations

## General Properties of Algorithms

As you've seen again and again in this chapter, iterators range. The function prototype indicates that. For example, the `copy()` function has the

```
template<class InputIterator, class OutputIterator>
OutputIterator copy(InputIterator first, InputIterator last,
 OutputIterator result);
```

Because the identifiers `InputIterator` and `OutputIterator` are used, they just as easily could have been `T` and `U`. However, the template parameter names to indicate the container types. The declaration tells you that the range parameters must be iterators, and the iterator indicating where the result goes must be an `OutputIterator`.

One way of classifying algorithms is on the basis of where the result is placed. Some algorithms do their work in place, meaning that when the `sort()` function is finished, the result is in the original data did. So `sort()` is an *in-place algorithm*. Other algorithms move the result of its work to another location, so it is an *out-of-place algorithm*. `copy()` can do both. Like `copy()`, it uses an output iterator. Unlike `copy()`, `transform()` allows the output iterator to be the same as the input range, so it can copy the transformed values back into the original range.

action conditionally, depending on the result of the predicate. These versions typically append `_if` to the function name. `replace_if()` replaces an old value with a new value if the predicate returns the value `true`. Here's the prototype:

```
template<class ForwardIterator, class Predicate, class T>
void replace_if(ForwardIterator first, ForwardIterator last,
 Predicate pred, const T& new_value);
```

(Recall that a predicate is a unary function object that returns a `bool` value.) A version called `replace_copy_if()`. You can find the prototype in the STL header. Its prototype is like:

As with `InputIterator`, `Predicate` is a function object that can easily be called `T` or `U`. However, the STL chooses a model where `bool` is the return type, so that the actual argument should be a model of the `UnaryPredicate` concept. The STL uses terms such as `Generator` and `BinaryPredicate` to describe these concepts. Keep in mind that the STL model other function object concepts. Keep in mind that the STL can remind you what the iterator or functor concepts are. The STL can remind you what the STL thing the compiler can check. If you use the STL, you can get a long list of error messages as the compiler tries to compile the code.

```

 string letters;
 cout << "Enter the letter grouping (quit to quit): ";
 while (cin >> letters && letters != "quit")
 {
 cout << "Permutations of " << letters << endl;
 sort(letters.begin(), letters.end());
 cout << letters << endl;
 while (next_permutation(letters.begin(), letters.end()))
 cout << letters << endl;
 cout << "Enter next sequence (quit to quit): ";
 }
 cout << "Done.\n";
 return 0;
}

```

---

Here's a sample run of the program in Listing 13.1:

```

Enter the letter grouping (quit to quit): awl
Permutations of awl
awl
awl
law
lwa
wal

```

```
la.remove(4); // remove all 4s from the
```

After this method call, all elements with the value 4 in the list are removed, and the list is automatically resized.

There is also an STL algorithm called `remove` that can be invoked by an object, it takes range arguments, and the function could look like this:

```
remove(lb.begin(), lb.end(), 4);
```

However, because this `remove()` is not a member function, it doesn't know the list's size. Instead, it makes sure all the nonremoved items are shifted one position to the left, and then it returns an iterator to the new past-the-end value. This means the list's size is not updated to the new list size. For example, you can use the `list::erase` method, which describes the part of the list that is no longer needed, and the process works.

#### Listing 16.18 **listrmv.cpp**

---

```
// listrmv.cpp -- applying the STL to a list
#include <iostream>
#include <list>
#include <algorithm>
```

```

 cout << "After using the erase() method\n";
 cout << "lb:\t";
 for_each(lb.begin(), lb.end(), Show);
 cout << endl;
 return 0;
 }

void Show(int v)
{
 std::cout << v << ' ';
}

```

---

Here's the output of the program in Listing 11.1:

Original list contents:

4 5 4 2 2 3 4 8 1 4

After using the remove() method:

la: 5 2 2 3 8 1

After using the remove() function:

lb: 5 2 2 3 8 1 4 8 1 4

After using the erase() method:

lb: 5 2 2 3 8 1

```
vector<string> words;
string input;
while (cin >> input && input != "quit")
 words.push_back(input);
```

What about getting the alphabetic word list? `unique()`, but that approach overwrites the original vector. There is an easier way that avoids creating a new object and copy (using an insert iterator) that automatically sorts its contents, which means you only need one copy of a key, so that takes the place of `sort()` called for ignoring the case differences. One can use `copy()` instead of `copy()` to copy data from the vector to a new vector, you can use one that converts a string to

```
set<string> wordset;
transform(words.begin(), words.end(),
 insert_iterator<set<string> > (wordset,
```

The `ToLower()` function is easy to write. We pass the `tolower()` function to each element in the source array, and store the result in the destination. Remember, `string` objects, too, are passed by reference, so returning the string as a reference means the original string is modified, without having to make copies. Here's the code for `ToLower()`:



```
wordmap.insert(pair<string, int>(*si,
words.end(), *si)));
```

The map class has an interesting feature: You can use indexes to access the stored values. For example, you can use the value associated with the key "the", which is the value of the string "the". Because the wordset contains the string "the", you can use the following code as an alternative approach:

```
for (si = wordset.begin(); si != wordset.end(); ++si)
 wordmap[*si] = count(words.begin(), words.end(), *si);
```

Because `si` points to a string in the wordset, you can use `*si` as a key for `wordmap`. This code places both keys and values in the map.

Similarly, you can use the array notation to access the values:

```
for (si = wordset.begin(); si != wordset.end(); ++si)
 cout << *si << ": " << wordmap[*si] << endl;
```

If a key is invalid, the corresponding value is 0.

Listing 16.19 puts these ideas together and shows how to use the three containers (a vector with the input, a set with the unique words, and a map with the word count).

```
words.push_back(input);
```

```
cout << "You entered the following words:\n";
for_each(words.begin(), words.end(),
cout << endl;
```

```
// place words in set, converting to lowercase
set<string> wordset;
transform(words.begin(), words.end(),
 insert_iterator<set<string> > (wordset.begin(),
 ToLower);
cout << "\nAlphabetic list of words:\n";
for_each(wordset.begin(), wordset.end(),
cout << endl;
```

```
// place word and frequency in map
map<string, int> wordmap;
set<string>::iterator si;
for (si = wordset.begin(); si != wordset.end(); ++si)
 wordmap[*si] = count(words.begin(), words.end(), *si);
```

```
// display map contents
cout << "\nWord frequency:\n";
```

You entered the following words:

The dog saw the cat and thought the cat fa

Alphabetic list of words:

and cat dog fat perfect saw the thought

Word frequency:

and: 1

cat: 4

dog: 1

fat: 1

perfect: 1

saw: 1

the: 5

thought: 2

The moral here is that your attitude when writing code is important. Write as much code as possible. STL's generic and flexible design is a result of this attitude. Also the STL designers are algorithm people. Efficiency is their priority. So the algorithms are well chosen and efficient.

class template, on the other hand, is oriented part of the STL. It doesn't have `push_back()` does provide a simple, intuitive interface for is designed as a substitute for the built-in array efficiency of that type with a better, safer interface `push_back()` and `insert()`, but it does offer `begin()`, `end()`, `rbegin()`, and `rend()`, making objects.

Suppose, for example, that you have these

```
vector<double> ved1(10), ved2(10), ved3(10);
array<double, 10> vod1, vod2, vod3;
valarray<double> vad1(10), vad2(10), vad3(10);
```

Furthermore, assume that `ved1`, `ved2`, `ved3` contain values. Suppose you want to assign the sum of each element of a third array, and so on. With the

```
transform(ved1.begin(), ved1.end(), ved2.begin(),
 plus<double>());
```

You can do the same with the array class

```
transform(vod1.begin(), vod1.end(), vod2.begin(),
 plus<double>());
```

```
log);
```

The `valarray` class overloads the usual mathematical operators and to return a `valarray` object, so you can use

```
vad3 = log(vad1); // log() overloaded
```

Or you could use the `apply()` method, which returns

```
vad3 = vad1.apply(log);
```

The `apply()` method doesn't alter the invoker, but returns a `valarray` that contains the resulting values.

The simplicity of the `valarray` interface is demonstrated in the next step calculation:

```
vad3 = 10.0 * ((vad1 + vad2) / 2.0 + vad1 * vad2);
```

The vector-STL version is left as an exercise for the reader.

The `valarray` class also provides a `sum()` method that returns a `valarray` object, a `size()` method that returns the number of elements, a `max()` method that returns the largest value in an object, and a `min()` method that returns the smallest value.

As you can see, `valarray` has a clear notation for mathematical operations, but it is also much less versatile than

vide `vad` and `vad + 10`, as the following code  
`sort(vad, vad + 10); // NO, vad an objec`

You can use the address operator:

```
sort(&vad[0], &vad[10]); // maybe?
```

But the behavior of using a subscript one  
doesn't necessarily mean using `&vad[10]` wor  
pilers used to test this code.) But it does mea  
fail, you probably would need a very unlikely  
against the end of the block of memory set a  
sion depended on your code, you might not

C++11 remedies the situation by providi  
that take a `valarray` object as an argument.  
`vad.begin()`. These functions return values  
requirements:

```
sort(begin(vad), end(vad)); // C++11 fix!
```

Listing 16.20 illustrates some of the relativ  
classes. It uses `push_back()` and the automa  
Then after sorting the numbers, the program  
`valarray` object of the same size and does a

```

 sq_rts = sqrt(numbers);
 valarray<double> results(size);
 results = numbers + 2.0 * sq_rts;
 cout.setf(ios_base::fixed);
 cout.precision(4);
 for (i = 0; i < size; i++)
 {
 cout.width(8);
 cout << numbers[i] << ": ";
 cout.width(8);
 cout << results[i] << endl;
 }
 cout << "done\n";
 return 0;
}

```

---

Here is a sample run of the program in Lis

Enter numbers (<=0 to quit):

**3.3 1.8 5.2 10 14.4 21.6 26.9 0**

1.8000: 4.4833

3.3000: 6.9332

5.2000: 9.7607

10.0000: 16.3246

This special subscripting facility allows you to represent two-dimensional data. For example, with 4 rows and 3 columns. You can store the object. Then a `slice(0,3,1)` object used as 0 and 2—that is, the first row. Similarly, a `slice(0,3,6, and 9—that is, the first column. Listing`

#### Listing 16.21 **vslice.cpp**

---

```
// vslice.cpp -- using valarray slices
#include <iostream>
#include <valarray>
#include <cstdlib>

const int SIZE = 12;
typedef std::valarray<int> vint; // s
void show(const vint & v, int cols);
int main()
{
 using std::slice; // f
 using std::cout;
 vint valint(SIZE); // t

 int i;
```



```

 using std::cout;
 using std::endl;

 int lim = v.size();
 for (int i = 0; i < lim; ++i)
 {
 cout.width(3);
 cout << v[i];
 if (i % cols == cols - 1)
 cout << endl;
 else
 cout << ' ';
 }
 if (lim % cols != 0)
 cout << endl;
 }

```

---

The `+` operator is defined for `valarray` of single int element, such as `valint[1]`. But `+` operator isn't defined for `slice`-subscripted `valarray`. Therefore, the program constructs full objects

```

vint (valint[slice(1,4,3)]) // calls a s

```

```
12 2 10
19 9 10
```

Because values are set using `rand()`, different values.

There's more, including the `gslice` class that should be enough to give you a sense of what

## The `initializer_list` Template

The `initializer_list` template is another way to use the initializer-list syntax to initialize an `std::vector`.

```
std::vector<double> payments {45.99, 39.21, 41.50, 38.95};
```

This would create a container for four elements with the values in the list. What makes this possible is that `std::vector` has a constructor taking an `initializer_list<T>` argument. `std::vector` has a constructor that accepts an `initializer_list` argument. The previous declaration is the same as this:

```
std::vector<double> payments({45.99, 39.21, 41.50, 38.95});
```

The usual list restrictions on narrowing apply. For example, the following is not allowed:

```
std::vector<int> values = {10, 8, 5.5};
```

Here, the element type is `int`, and the implicit conversion from `double` to `int` is not allowed.

It doesn't make sense to provide an initializer list constructor meant to handle lists of varying sizes. For instance, the constructor for a class taking a fixed number of arguments does not provide an `initializer_list` constructor:

```
class Position
{
private:
 int x;
 int y;
 int z;
public:
 Position(int xx = 0, int yy = 0, int zz = 0)
 : x(xx), y(yy), z(zz) {}
 // no initializer_list constructor
 ...
};
```

```
using std::cout;
```

```
cout << "List 1: sum = " << sum({2,3,4})
 << ", ave = " << average({2,3,4})
std::initializer_list<double> dl = {1,2,3,4,5,6,7,8,9,10}
cout << "List 2: sum = " << sum(dl)
 << ", ave = " << average(dl) << "
dl = {16.0, 25.0, 36.0, 40.0, 64.0};
cout << "List 3: sum = " << sum(dl)
 << ", ave = " << average(dl) << "
return 0;
```

```
}
```

```
double sum(std::initializer_list<double> il)
{
 double tot = 0;
 for (auto p = il.begin(); p != il.end(); ++p)
 tot += *p;
 return tot;
}
```

```
}
```

```
double average(const std::initializer_list<double> il)
{
 double tot = 0;
 int n = il.size();
 if (n > 0)
 tot = sum(il) / n;
 return tot;
}
```

choice is not a major performance issue. (The

The function argument can be a list literal like `dl`.

The iterator types for `initializer_list` in a list:

```
*dl.begin() = 2011.6; // no
```

But, as Listing 16.22 shows, you can attach

```
dl = {16.0, 25.0, 36.0, 40.0, 64.0}; // a
```

However, the intended use of the `initializer` constructor or some other function.

## Summary

C++ includes a powerful set of libraries that programming problems and the tools to simplify provides a convenient means to handle strings as management and a host of methods and functions these methods and functions allow you to compare another, reverse a string, search a string for characters and output operations.

underlying container class to give it the character of a stack. The `stack` class template name. Thus, `stack`, although built on top of a container, allows insertion and removal only at the top of the stack. C++ provides `unordered_multiset`, `unordered_map`, and `unordered_set`.

Some algorithms are expressed as container member functions, while others are as general, nonmember functions. This is made possible by the close relationship between containers and algorithms. One advantage is that there can be just one `for_each()` or `copy()` function, instead of one for each container. A second advantage is that STL algorithms work with all STL containers, such as ordinary arrays, `string` objects, and `vector`. This design is consistent with the STL iterator and container design.

Both containers and algorithms are characterized by their requirements or need. You should check that a container fulfills the requirements of an algorithm's needs. For example, the `for_each` algorithm's requirements are met by all the STL containers. The `random_access_iterator` requirements are met by all the STL containers that offer random access iterators, which not all containers do. Some containers offer a specialized method as an option if it does not meet the requirements of an algorithm. For example, the `list` class has a `splice` method, so it can use that method instead of the `splice` algorithm.

The STL also provides function objects, or functors. A functor is a class whose `operator()` is overloaded—that is, for which the

```

 RQ1(const char * s)
 {st = new char [strlen(s) + 1];
 RQ1(const RQ1 & rq)
 {st = new char [strlen(rq.st) +
 ~RQ1() {delete [] st};
 RQ & operator=(const RQ & rq);
 // more stuff
};

```

Convert this to a declaration that uses a longer name and no longer need explicit definitions?

2. Name at least two advantages string objects have over char arrays of ease-of-use.
3. Write a function that takes a reference to a string object and converts the string object to all upper case.
4. Which of the following are not examples (or are bad examples) of `auto_ptr`? (Assume that the following are all legal.)

```

auto_ptr<int> pia(new int[20]);
auto_ptr<string> (new string);
int rigue = 7;
auto_ptr<int>pr(&rigue);
auto_ptr dbl (new double);

```

1. A *palindrome* is a string that is the same forwards and backwards. “otto” and “otto” are rather short palindromes. Write a function `is_palindrome(string)` that takes a `string` object and that passes to a `bool` function. The function should return `true` if the string is a palindrome, and `false` otherwise. Don’t worry about complications such as capitalization. That is, this simple version should reject “Otto”. To scan the list of string methods in Appendix A.
2. Do the same problem as given in Program 16.2, but with complications such as capitalization, spaces, and punctuation. “Adam” should test as a palindrome. For example, “madamimadam” and “taco cat” should test as palindromes. Forget the useful `cctype` library. You may use `isalpha` and `isspace` although not necessary.
3. Redo Listing 16.3 so that it gets it working with a `vector<string>` object instead of an array. Use `push_back()` to copy how ever many words into the `vector<string>` object and use the `size()` method to get the size of the word list. Because the program should read words separated by spaces, tabs, or newlines, you should use the `>>` operator rather than `get()`.



ments. The first should be the number of spots selected. The second should be the number of spots selected. The third should be a `vector<int>` object that contains, in sorted order, the numbers that should be selected. For example, you could use the function `vector<int> winners;`  
`winners = Lotto(51,6);`

This would assign to `winners` a vector of 6 numbers from the range 1 through 51. Note that this is not a good job because it may produce duplicate values. A better job would be to create a vector that contains all the possible values and then select the beginning of the shuffled vector to obtain the winners. This is the function that lets you test the function.

8. Mat and Pat want to invite their friends to a party. Write a program that does the following:
  - Allows Mat to enter a list of his friends' names and then displays them in sorted order.
  - Allows Pat to enter a list of her friends' names and then displays them in sorted order.
  - Creates a third container that merges the two lists and displays the contents of this container.

```
clock_t end = clock();
cout << (double)(end - start)/CLOCKS
```

This is by no means a definitive test because of many factors, including available memory, which limits the size of the array or list. (One would expect the time for an array over the list to increase with the size of the array; you have a choice between a default constructor and a constructor for the measurement. With today's speeds, you can use as large an array as possible to get meaningful results: 100,000 elements, 1,000,000 elements, and so on.)

10. Modify Listing 16.9 (`vect3.cpp`) as follows:
  - a. Add a `price` member to the `Review` class.
  - b. Instead of using a vector of `Review` objects, use a vector of `shared_ptr<Review>` objects, each initialized with a pointer returned by `make_shared`.
  - c. Follow the input stage with a loop that sorts the books for displaying books: in original order, in order of increasing ratings, in order of decreasing ratings, in order of decreasing price, and so on.



- `istream` class methods
- Stream states
- File I/O
- Using the `ifstream` class for input from files
- Using the `ofstream` class for output to files
- Using the `fstream` class file input and output
- Command-line processing
- Binary files
- Random file access
- Incore formatting

Discussing C++ input and output (I/O) is a topic that every program uses input and output, and learning the tasks facing someone learning a computer language is one of its more advanced language features to implement. Thus, to really understand C++ I/O, you started, the early chapters of this book outline the object `cin` and the `ostream` class object `cout` and using `ifstream` and `ofstream` objects for file

ments of the target computer. In practice, most C++ I/O library functions originally developed for the standard C++ recognition of this I/O package, called the Standard C++ Library, a mandatory component of the standard C++ library. If you're familiar with the family of C functions, you'll find them in C++ programs. (Newer implementations of these functions.)

However, C++ relies on a C++ solution for I/O. This solution is a set of classes defined in the `iostream` (formerly `fstream.h`) header files. This class library defines the rules for how to do things, such as create classes and objects by following those rules. But just as C++ implements a family of functions, C++ comes with a standard library. The standard library was an informal standard consisting of the `iostream` and `fstream` header files. The ANSI/ISO C++ standard adopted the standard library as a standard class library and to add a new class, discussed in Chapter 16, "The `string` Class and `string` Literals." Chapter 16 discusses standard C++ I/O. But first, let's look at C++ I/O.

a device, such as a keyboard.) Similarly, managing the flow of information from a device to a program and associating some plumbing with bytes instead of water (see Fig. 1.10).

Usually, input and output can be handled by a block of memory used as an intermediate, transferring information from a device to a program or from a program to a device. For example, disk drives transfer information in blocks, but programs often process information 1 byte at a time. This is a waste of rates of information transfer. For example, assume a program counts the number of dollar signs in a hard-disk file. The program reads the file, processes it, reads the next character from the file, and so on. At a time from a disk requires a lot of hardware. It is much faster to read a large chunk from the disk, store it in memory, and then process one character at a time. Because it is much quicker to read from memory than from a hard disk, this approach is used in many operating systems. Of course, after the program reaches the end of the first chunk, it reads another chunk of data from the disk. This is like a reservoir that collects megagallons of runoff water from a mountain and releases it at your home at a more civilized rate of flow (see Fig. 1.11). In a program, the program can first fill the buffer and then transfer information from the buffer to the program, saving the buffer for the next batch of output. This is a good idea. You can come up with your own plumbing-based

Keyboard input provides one character at a time, so we need a buffer to help match different data transfer rates. It also allows the user to back up and correct input before the program normally flushes the input buffer when it reaches the end of a line. In this book we don't begin processing input until we have a full line. A C++ program normally flushes the output buffer when it reaches the end of a line. Depending on the implementation, a program can also flush too, such as at impending input. That is, when the program flushes any output currently in the output buffer. A C++ program consistent with ANSI C should behave in that manner.

## **Streams, Buffers, and the `iostream` Library**

The business of managing streams and buffers is handled by the `iostream` (formerly `iostream.h`) file brings in the `istream` and `ostream` classes to manage streams and buffers for you. The C++ standard defines templates in order to support both `char` and `wchar_t` character types. `char32_t` specializations. By using the typedefs and specializations of these templates mimic the traditional C++ `istream` and `ostream` classes. Here are some of those classes (see Figure 17.3):

refill stream buffer with n

Figure 17.2 A s

- The `streambuf` class provides memory for filling the buffer, accessing buffer contents, and managing buffer memory.
- The `ios_base` class represents general stream operations, open for reading and whether it's a binary or text stream.
- The `ios` class is based on `ios_base`, and provides a `streambuf` object.
- The `ostream` class derives from the `ios` class and provides output methods.
- The `istream` class derives from the `ios` class and provides input methods.
- The `iostream` class is based on the `ios` class and provides both input and output methods.

To use these facilities, you use objects of the `ostream` class, such as `cout` to handle output. The `cout` object automatically creates a buffer, and associates it with the `streambuf` member functions available to you.



## Redefining I/O

The ISO/ANSI C++98 Standard revised I/O and moved `ostream.h` to `ostream`, with `ostream` placing the I/O classes have been rewritten. To be able to handle international character sets that require the language added the `wchar_t` (or “wide”) character type. C++11 adds `char16_t` and `char32_t` facilities. Rather than develop two (or, now, four) sets, the committee developed a template set of I/O classes. `basic_ostream<charT>` and `basic_istream<charT>` template, in turn, is a template for any character type, such as how to compare for equality. It provides `char` and `wchar_t` specializations of `basic_ostream` and `basic_istream`. `ostream` and `istream` are typedefs for `char` specializations of `basic_ostream` and `basic_istream`. For example, there are `cout` and `cin` streams. The `ostream` header file contains the following:

Certain type-independent information that used to be in `ios.h` has been moved to the new `ios_base` class. This includes `ios::fixed`, which is now `ios_base::fixed`. The `ios` class aren’t available in the old `ios`.

ing information relating to output, such as the number of places after the decimal to use, the number of digits, and the address of a `streambuf` object that manages the output flow. A statement such as the following inserts the string "Bjarne free" into the buffer managed by `cout`:

```
cout << "Bjarne free";
```

The `ostream` class defines the operator `<<`. The `ostream` class also supports the `cout` data member, which is an instance of `ostream`, such as the ones this chapter discusses later. For each `ostream` object, the flow of bytes from the buffer is directed to the standard output of the operating system. In short, one end of a stream is connected to the standard output, and the other end is connected to a `streambuf` object, manages the flow of bytes.

## Redirection

The standard input and output streams normally write to the standard output. But many operating systems, including Unix, have a facility that lets you change the associations for the standard streams. Suppose, for example, that you have an executable program called `counter.exe` that counts the number of bytes written to the standard output. (From most versions of Windows you can right-click the Command Prompt icon to open a command prompt window like this:

this redirection syntax. (All of these other than the first one are for redirecting the standard error stream.) The characters between the redirection operators are the file names.

The standard output stream, represented by `cout`, is the default output. The standard error streams (represented by `cerr` and `clog`) are used for the program's error messages. By default, all three of these streams are connected to the standard output stream. But redirecting the standard output doesn't redirect the standard error streams. If you want to use these objects to print an error message, a program should redirect the standard error stream to a file or screen even if the regular `cout` output is redirected. Here's a code fragment:

```
if (success)
 std::cout << "Here come the goodies!\n";
else
{
 std::cerr << "Something horrible has happened!\n";
 exit(1);
}
```

If redirection is not in effect, whichever message stream is used, however, the program output has been redirected to a file. If, for example, the first redirection would go to the file but the second message, if any, would go to the standard error stream. In some operating systems, some operating systems permit redirecting the standard error stream. For example, the `2>` operator redirects the standard error stream to a file.

Most often, this book has used `cout` with the `<<` operator:

```
int clients = 22;
cout << clients;
```

In C++, as in C, by default the `<<` operator (see Appendix E, “Other Operators”). An expression `x << y` represents the binary representation of `x` and shift all the bits three positions to the left. This is a lot to do with output. But the `ostream` class overloads the `<<` operator for loading to output for the `ostream` class. In this book, the right-shift operator instead of the left-shift operator. The role through its visual aspect, which suggests that the right-shift operator is overloaded to recognize all the

- `unsigned char`
- `signed char`
- `char`
- `short`
- `unsigned short`
- `int`
- `unsigned int`

also indicates that the function returns a reference, which makes it possible to concatenate output, as in

```
cout << "I'm feeling sedimental over " <<
```

If you're a C programmer who has suffered from the problems and the problems that arise when you mismatch almost sinfully easy. (And C++ input, of course.)

## Output and Pointers

The `ostream` class defines insertion operator for

- `const signed char *`
- `const unsigned char *`
- `const char *`
- `void *`

C++ represents a string, don't forget, by using a pointer. The pointer can take the form of the name of a variable, a `to-char` or of a quoted string. Thus, all the following

```
char name[20] = "Dudly Diddlemore";
char * pn = "Violet D'Amore";
cout << "Hello!";
```

tions say that the reference is to the object us-  
operator function's return value is the same o-  
cout << "potluck" returns the cout object  
output by using insertion. For example, consi-  
cout << "We have " << count << " unhatched

The expression `cout << "We have "` disp-  
reducing the statement to the following:

```
cout << count << " unhatched chickens.\n"
```

Then the expression `cout << count` displ-  
returns `cout`, which can then handle the fina-  
17.4). This design technique really is a nice fe-  
loading the `<<` operator in the previous chap-

## **The Other ostream Methods**

Besides the various `operator<<()` functions,  
method for displaying characters and the `wri-`

Originally, the `put()` method had the following prototype:

```
ostream & put(char);
```

The current standard is equivalent, except that you can now invoke it by using the usual class method notation:

```
cout.put('W'); // display the W character
```

Here `cout` is the invoking object and `put()` is one of the overloaded operator functions, this function returns a reference to `cout` to concatenate output with it:

```
cout.put('I').put('t'); // displaying It with a space
```

The function call `cout.put('I')` returns a reference to `cout` for the `put('t')` call.

Given the proper prototype, you can use `put()` with any type other than `char`, such as `int`, and let function prototypes convert the value to the correct type `char` value. For example, you can write

```
cout.put(65); // display the A character
cout.put(66.3); // display the B character
```

the second argument indicates how many characters to write. The `write()` invokes the `char` specialization, so this example shows how the `write()` method works.

### Listing 17.1    **write.cpp**

---

```
// write.cpp -- using cout.write()
#include <iostream>
#include <cstring> // or else string.h

int main()
{
 using std::cout;
 using std::endl;
 const char * state1 = "Florida";
 const char * state2 = "Kansas";
 const char * state3 = "Euphoria";
 int len = std::strlen(state2);
 cout << "Increasing loop index:\n";
 int i;
 for (i = 1; i <= len; i++)
 {
 cout.write(state2,i);
 cout << endl;
 }
}
```



```
Kansas
Decreasing loop index:
Kansas
Kansa
Kans
Kan
Ka
K
Exceeding string length:
Kansas Euph
```

Note that the `cout.write()` call returns the `cout` object. The `write()` method returns a reference to the object that invokes it. This makes it possible to concatenate the output of `write()` by its return value, `cout`:

```
cout.write(state2,i) << endl;
```

Also, note that the `write()` method doesn't stop writing when it reaches the null character. It simply prints all the characters if that goes beyond the bounds of a particular string. For example, the string "Kansas" with two other strings so that the total length exceeds the buffer size. Compilers differ in the order in which they align memory. For example, "Kansas" occupies

Then the program *flushes* the buffer, sending new data. Typically, a buffer is 512 bytes or an time-saver when the standard output is connected to a device that doesn't want a program to access the hard disk directly. It's more effective to collect 512 bytes in a buffer and write them in one operation.

For screen output, however, filling the buffer is inconvenient if you had to reword the message until it had consumed the prerequisite 512 bytes to fill a buffer. In this case, the program doesn't necessarily wait until the buffer is full. The buffer, for example, normally flushes the buffer when the program implements flush the buffer when input is received. The following code:

```
cout << "Enter a number: ";
float num;
cin >> num;
```

The fact that the program expects input to be flushed (the "Enter a number: " message) immediately after a newline character. Without this feature, the program would be prompting the user with the cout message.

If your implementation doesn't flush output automatically, you can flush by using one of two manipulators. The first

to hold the number and, if present, a minus sign.

- Strings are displayed in a field equal in width to the string.

The default behavior for floating-point types differs between ancient and current C++ implementations.

- **New style**—Floating-point types are displayed in decimal notation. Trailing zeros aren't displayed. (Note that the number of trailing zeros has no connection with the precision to which the number is displayed.) The number is displayed in fixed-point notation or else in E notation, depending on the value of the number. The exponent is 6 or larger or -5 or smaller. Again, the number is displayed in decimal notation and, if present, a minus sign. The standard C library function `fprintf()` uses the new style.
- **Old style**—Floating-point types are displayed in decimal notation. Trailing zeros aren't displayed. The number of trailing zeros displayed has no connection with the precision to which the number is displayed. The number is displayed in fixed-point notation or else in E notation, depending on the value of the number. The number is displayed in decimal notation and, if present, a minus sign.

Because each value is displayed in a width of at least one character, there is a space between values explicitly; otherwise, consecutive values are displayed without spaces.

```
char ch = 'K';
int t = 273;
cout << ch << ":\n";
cout << t << ":\n";
cout << -t << ":\n";

double f1 = 1.200;
cout << f1 << ":\n";
cout << (f1 + 1.0 / 9.0) << ":\n";

double f2 = 1.67E2;
cout << f2 << ":\n";
f2 += 1.0 / 9.0;
cout << f2 << ":\n";
cout << (f2 * 1.0e4) << ":\n";

double f3 = 2.3e-4;
cout << f3 << ":\n";
cout << f3 / 10 << ":\n";

return 0;
}
```

---

to display integers. By using `ios_base` members and the number of places displayed to the right, the `ostream` class is an indirect base class for `ostream`, you (or descendants), such as `cout`.

## Note

The members and methods found in the `ios_base` class. Now `ios_base` is a base class to `ios` with `char` and `wchar_t` specializations, and

Let's look at how to set the number base to whether integers are displayed in base 10, base 8, or base 16 using the `hex`, `dec`, and `oct` manipulators. For example, the following sets the format state for the `cout` object to hexadecimal:

```
hex(cout);
```

After you do this, a program will print integers in hexadecimal. You can change the format state to another choice. Note that these manipulators are global functions, hence they don't have to be invoked by an object.

Although the manipulators really are functions, you can use them like this:

```
cout << hex;
```

```
// set to octal mode
 cout << oct << n << " " << n * n << endl;

// alternative way to call a manipulator
 dec(cout);
 cout << n << " " << n * n << " (d" << endl;

 return 0;
}
```

---

Here is some sample output from the program:

```
Enter an integer: 13
n n*n
13 169 (decimal)
d a9 (hexadecimal)
15 251 (octal)
13 169 (decimal)
```

right-justification. After that, the field width resets to 10 characters and the 24 are printed in fields equal to the

## Caution

The `width()` method affects only the next item printed. The default value afterward.

C++ never truncates data, so if you attempt to print a width of two, C++ expands the field to fit the data with asterisks if the data doesn't fit. The C/C++ standard is more important than keeping the columns neat. Listing 17.4 shows how the `width()` member function

### Listing 17.4 **width.cpp**

```
// width.cpp -- using the width method
#include <iostream>

int main()
{
 using std::cout;
 int w = cout.width(30);
 cout << "default field width = " << w;
```

The output displays values right-justified in 30 spaces. That is, `cout` achieves the full field width by inserting the spaces to the left of the value. This is the *fill character*. Right-justification is the default.

Note that the program in Listing 17.4 applies right-justification only to the first `cout` statement but not to the second. The `right` method affects only the next single item displayed. This is because `cout.width(30)` returns the previous width, because `cout.width(30)` returns the previous width just set. The fact that `w` is 0 means that zero is the default. To align column headings and data by using a width of eight characters for the second column and a width of eight characters for the second

## Fill Characters

By default, `cout` fills unused parts of a field with spaces. You can use the `fill` function to change that. For example, the following program uses an asterisk:

```
cout.fill('*');
```

That can be handy for, say, printing check marks. Listing 17.5 illustrates using this member function.



Waldo Whipsnade: \$\*\*\*\*900

Wilmarie Wooper: \$\*\*\*1350

Note that, unlike the field width, the new change it.

## Setting Floating-Point Display Precision

The meaning of floating-point *precision* depends on what it means the total number of digits displayed. (As discussed soon, *precision* means the number of digits after the decimal place. The precision default for C++, as you've seen, is 6. (Extra zeros are dropped.) The `precision()` member function of `cout` changes the precision. For example, the following statement causes `cout` to display two digits after the decimal point:

```
cout.precision(2);
```

Unlike the case with `width()`, but like the case with `setprecision()`, `precision()` stays in effect until it is reset. Listing 17.6 demonstrates this.

### Listing 17.6 `precise.cpp`

---

```
// precise.cpp -- setting the precision
#include <iostream>

int main()
```

fourth line displays a total of two digits.

## Printing Trailing Zeros and Decimal Points

Certain forms of output, such as prices or numbers, are retained. For example, the output to Listing 20.4. The `iostream` family of classes doesn't accomplish that. However, the `ios_base` class controls several formatting features. The class has arguments to this function. For example, to play trailing decimal points:

```
cout.setf(ios::showpoint);
```

In the default floating-point format, it also instead of displaying 2.00 as 2, `cout` will display 2.000000. This is in effect. Listing 17.7 adds this statement to

In case you're wondering about the notation, `ios::data_members` is a class-scope static constant that is defined in `ios.cpp`. This means that you have to use the scope-resolution operator `::` if you use the name outside a member function. For example, to use a constant defined in the `ios` base class,

Here is the output of the program in Listing 10-1:

```
"Furry Friends" is $20.4000!
"Fiery Fiends" is $2.78889!
"Furry Friends" is $20.!
"Fiery Fiends" is $2.8!
```

This output shows the trailing zeros for the first point but no trailing zeros because the precision of the second has been displayed.

## More About `setf()`

The `setf()` method controls several other formatting options. The point is displayed, so let's take a closer look at the `setf()` member in which individual bits (called *flags*) control various aspects, such as the number base and whether the output is called *setting the flag* (or bit) and means setting a bit, meaning equivalent to setting DIP switches to on or off. The `setf()` function provides another means of controlling the output.

The `setf()` function has two prototypes. The first prototype is `fmtflags setf(fmtflags);`

**Note** A *bitmask type* is a type that is used to store a type, an enum, or an STL `bitset` container. The `ios_base::showpos` is accessible and has its own meaning. The `ios_base::showpos` state information.

Because these formatting constants are defined in the `std` namespace, you can use the scope-resolution operator with them. That is, you can use `std::ios_base::showpos` instead of just uppercase. If you don't use a `using` directive, you must use the scope-resolution operator to indicate that the constant is from the `std` namespace. If you don't use the scope-resolution operator, you can use `std::ios_base::showpos`, and the `std::ios_base::showpos` constants are overridden. Listing 17.8 illustrates using `std::ios_base::showpos`.

#### Listing 17.8 **setf.cpp**

```
// setf.cpp -- using setf() to control formatting
#include <iostream>

int main()
{
 using std::cout;
 using std::endl;
 using std::ios_base;

 ios_base::showpos;

 int temperature = 63;
```

0X3F

How 0X1! oops -- How true!

Note that the plus sign is used only with the decimal and octal values as unsigned; therefore no sign is displayed. (Some implementations may still display a plus sign.)

The second `setf()` prototype takes two arguments. The first argument is the `fmtflags` object, and the second is the `fmtflags` value to be set. The following code sets the `fmtflags` object to the `fmtflags` value `0x1`:

This overloaded form of the function is used to set a bit to 1. The first argument, as before, is a `fmtflags` object. The second argument is a value that first specifies the bit to be set. Suppose setting bit 3 to 1 means base 10, setting bit 4 to 1 means base 16. Suppose output is in base 10. If you only do you have to set bit 5 to 1, you also have to set bit 4 to 1. The clever `hex` manipulator does both tasks. The `basefield` manipulator requires a bit more work because you use the `basefield` manipulator to clear and then use the first argument to indicate the base. This is indicated as it sounds because the `ios_base` class has a `basefield` member for this purpose. In particular, you should use the `ios_base::basefield` member and `ios_base::hex` as the first argument. The following function call has the same effect as the following:

```
cout.setf(ios_base::hex, ios_base::basefield);
```

ment. Left-justification means starting a value at the left end of the field. Right-justification means ending a value at the right end of the field. You can also place any signs or base prefixes at the left or right of the field. (Unfortunately, C++ does not allow this.)

Fixed-point notation means using the `123.456` format, where 123 is the integer part of the number, and 456 is the size of the number, and scientific notation means using the `1.23456e7` format, where 1.23456 is the mantissa, and 7 is the size of the number. If you are familiar with C, you know that the default C++ mode corresponds to the `%f` specifier, and `scientific` corresponds to the `%e` specifier.

Under the C++ Standard, both fixed and scientific notation have the following properties:

- *Precision* means the number of digits to the right of the decimal point. The number of digits to the left of the decimal point is determined by the number of digits.
- Trailing zeros are displayed.

The `setf()` function is a member function of the `ostream` class. For the `ostream` class, you can invoke the `setf()` function. For example, to request left-justification, you use the `ios_base::left` flag.

```
ios_base::fmtflags old = cout.setf(ios::left, ios::fmtflags);
```

To restore the previous setting, you use the `unsetf()` function. For example, to restore the previous setting, you use the `unsetf()` function.

```
cout.setf(old, ios::adjustfield);
```

```
cout.setf(ios_base::showpos);
cout.setf(ios_base::showpoint);
cout.precision(3);
// use e-notation and save old format
ios_base::fmtflags old = cout.setf(ios_base::floatfield);
cout << "Left Justification:\n";
long n;
for (n = 1; n <= 41; n+= 10)
{
 cout.width(4);
 cout << n << "|";
 cout.width(12);
 cout << sqrt(double(n)) << "|\n";
}

// change to internal justification
cout.setf(ios_base::internal, ios_base::floatfield);
// restore default floating-point display
cout.setf(old, ios_base::floatfield);

cout << "Internal Justification:\n";
for (n = 1; n <= 41; n+= 10)
{
```

Left Justification:

|     |            |  |
|-----|------------|--|
| +1  | +1.000e+00 |  |
| +11 | +3.317e+00 |  |
| +21 | +4.583e+00 |  |
| +31 | +5.568e+00 |  |
| +41 | +6.403e+00 |  |

Internal Justification:

|      |   |      |  |
|------|---|------|--|
| + 1  | + | 1.00 |  |
| + 11 | + | 3.32 |  |
| + 21 | + | 4.58 |  |
| + 31 | + | 5.57 |  |
| + 41 | + | 6.40 |  |

Right Justification:

|     |  |        |  |
|-----|--|--------|--|
| +1  |  | +1.000 |  |
| +11 |  | +3.317 |  |
| +21 |  | +4.583 |  |
| +31 |  | +5.568 |  |
| +41 |  | +6.403 |  |

Note how a precision of 3 causes the default justification in this program) to display a total of three modes display three digits to the right of the decimal point. The exponent for e-notation depends on the



If you knew for certain that `cout` were in the `ios_base::fixed` state, you could call `unsetf(ios_base::fixed)` as an argument to `unsetf()` regardless of the current state of `cout`, so it's a

## Standard Manipulators

Using `setf()` is not the most user-friendly approach. There are standard manipulators to invoke `setf()` for you, automatically. You've already seen `dec`, `hex`, and `oct`. These manipulators, like those in older C++ implementations, work like `hex` and `oct` on left-justification and the fixed decimal point.

```
cout << left << fixed;
```

Table 17.3 lists these along with several other manipulators.

**Table 17.3 Some Standard Manipulators**

| Manipulator              |
|--------------------------|
| <code>boolalpha</code>   |
| <code>noboolalpha</code> |
| <code>showbase</code>    |
| <code>noshowbase</code>  |

oct

fixed

scientific

---

### Tip

If your system supports these manipulators, have the option of using `setf()`.

## The `iomanip` Header File

Setting some format values, such as the field width, is a common task for many programs. To make life easier, C++ supplies additional manipulators in the `iomanip` header file. They provide the same services already discussed in a more convenient manner. The three most commonly used manipulators are `setw()` for setting the field width, `setfill()` for setting the fill character, and `setprecision()`. Unlike the manipulators discussed previously, `setw()` takes an integer argument that specifies the field width, `setfill()` takes a `char` argument that indicates the fill character, and `setprecision()` takes an integer argument that specifies the number of digits to display after the decimal point.

```

 cout << fixed << right;

// use iomanip manipulators
cout << setw(6) << "N" << setw(14) <<
 << setw(15) << "fourth root\n";

double root;
for (int n = 10; n <=100; n += 10)
{
 root = sqrt(double(n));
 cout << setw(6) << setfill('.') <<
 << setw(12) << setprecision(3) <<
 << setw(14) << setprecision(3) <<
 << endl;
}

return 0;
}

```

---

Here is the output of the program in Listing 10.1:

| N      | square root | fourth root |
|--------|-------------|-------------|
| ....10 | 3.162       | 1.7783      |
| ....20 | 4.472       | 2.1147      |

Typically, you use `cin` as follows:

```
cin >> value_holder;
```

Here `value_holder` identifies the memory location to be the name of a variable, a reference, a dereferenced pointer, or of a class. How `cin` interprets the input depends on the `istream` class, defined in the `iostream` header file. The `istream` class recognizes the following basic types:

- signed char &
- unsigned char &
- char &
- short &
- unsigned short &
- int &
- unsigned int &
- long &
- unsigned long &
- long long & (C++11)
- unsigned long long & (C++11)

int. In this case, the compiler matches

```
cin >> staff_size;
```

to the following prototype:

```
istream & operator>>(int &);
```

The function corresponding to that prototype is sent to the program—say, the characters 2, 3, 1, 8, 4. The function then converts these characters to the integer 23184. If, on the other hand, `staff_size` is a `double`, `operator>>(double &)` to convert the same input to the representation of the value 23184.0.

Incidentally, you can use the `hex`, `oct`, and `dec` manipulators. If the integer input is to be interpreted as hexadecimal, the following statement causes an input of 12 to be read as the decimal 18, and it causes `ff` or `FF` to be read as the decimal 255.

```
cin >> hex;
```

The `istream` class also overloads the `>>` operator for the following types:

- `signed char *`
- `char *`
- `unsigned char *`

The various versions of the extraction operator skip over white space (blank, tab, and newline) in the input stream. They skip over white space (blank, tab, and newline) until they find a non-white-space character. This is true even if the argument is type `char`, unsigned `char`, or `wchar_t`. C's character input functions (see Figure 17.5) are no exception. The `getchar` operator reads that character and assigns it to the variable. The `get` operator reads in one unit of the indicated type. The `getline` operator reads the initial non-white-space character up to the first white-space character of the given type.

For example, consider the following code:

```
int elevation;
cin >> elevation;
```

Suppose you type the following characters:

```
-123Z
```

The operator will read the `-`, `1`, `2`, and `3` characters and convert them to an integer. But the `z` character isn't valid, so the `z` character remains in the input stream, and the next call to `get` will read it. Meanwhile, the operator converts the characters `-123` to an integer and assigns it to `elevation`.

whether input meets the program requirements.

### Listing 17.11 **check\_it.cpp**

---

```
// check_it.cpp -- checking for valid input
#include <iostream>

int main()
{
 using namespace std;
 cout << "Enter numbers: ";

 int sum = 0;
 int input;
 while (cin >> input)
 {
 sum += input;
 }

 cout << "Last value entered = " << input << endl;
 cout << "Sum = " << sum << endl;
 return 0;
}
```

---

non-accessible file or trying to write to a write-only stream. (Implementations don't necessarily agree on which set `badbit`.) When all three of these stream error bits are set, programs can check the stream state and use that information to recover from the stream state. Table 17.4 lists these bits, along with some information about the stream state.

Table 17.4 Stream States

| Member               | Description                                                                                                                                       |
|----------------------|---------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>eofbit</code>  | Is set to 1 if the stream has reached the end of the file.                                                                                        |
| <code>badbit</code>  | Is set to 1 if the stream is in a bad state, meaning that there could be no further operations on the stream.                                     |
| <code>failbit</code> | Is set to 1 if the stream is in a fail state, meaning that characters could not be read or written, or that the expected character was not found. |
| <code>goodbit</code> | Just another name for <code>!badbit</code> .                                                                                                      |
| <code>good()</code>  | Returns <code>true</code> if the stream is in a good state (i.e., not cleared).                                                                   |
| <code>eof()</code>   | Returns <code>true</code> if the stream has reached the end of the file.                                                                          |



state, but they do so in a different fashion. The `clear()` method resets the stream state to the initial state. Thus, the following call uses the default state bits (`eofbit`, `badbit`, and `failbit`):

```
clear();
```

Similarly, the following call makes the state `failbit` set, and the other two state bits are cleared:

```
clear(eofbit);
```

The `setstate()` method, however, affects only one state bit. Thus, the following call sets `eofbit` without affecting the other state bits.

```
setstate(eofbit);
```

So if `failbit` was already set, it stays set.

Why would you reset the stream state? For example, you might want to use `clear()` with no argument to reopen the stream for input or end-of-file; whether doing so makes sense depends on what you're trying to accomplish. You'll see some examples shortly. The `clear()` method to provide a means for input and output functions to check the state. If it is an `int`, the following call can result in `eofbit`, `failbit`, or `badbit`:

```
cin >> num; // read an int
```

```
cin.exceptions(badbit | eofbit);
```

Listing 17.12 modifies Listing 17.11 so that an exception is thrown if failbit is set.

#### Listing 17.12 **cinexcp.cpp**

---

```
// cinexcp.cpp -- having cin throw an exception on failbit
#include <iostream>
#include <exception>

int main()
{
 using namespace std;
 // have failbit cause an exception to be thrown
 cin.exceptions(ios_base::failbit);
 cout << "Enter numbers: ";
 int sum = 0;
 int input;
 try {
 while (cin >> input)
 {
 sum += input;
 }
 } catch(ios_base::failure & bf)
```

tion for `badbit` because that circumstance would be designed to read numbers from a data file up to and including an exception for `failbit` because that would be

## Stream State Effects

An `if` or `while` test such as the following test (bits cleared):

```
while (cin >> input)
```

If a test fails, you can use the member functions to determine possible causes. For example, you could modify the test like this:

```
while (cin >> input)
{
 sum += input;
}
if (cin.eof())
 cout << "Loop terminated because EOF encountered"
```

Setting a stream state bit has a very important effect on further input or output until the bit is cleared. For example,

```
while (cin >> input)
{
```

way is to keep reading characters until reaching the end of the line. In Chapter 6, “Branching Statements and Logic,” we saw that `isspace()` returns true if its argument is a white-space character. We can use this to skip a line instead of just the next word:

```
while (cin.get() != '\n')
 continue; // get rid rest of line
```

This example assumes that the loop terminates because of reaching the end of the input, instead, that the loop terminated because of a failure. Then the new code disposing of bad input uses the `fail()` method to test whether the loop terminated because of a failure. For technical reasons, `fail()` returns true if either `failbit` or `badbit` are set. To exclude the latter case. The following code shows how to do this.

```
while (cin >> input)
{
 sum += input;
}
cout << "Last value entered = " << input << endl;
cout << "Sum = " << sum << endl;
if (cin.fail() && !cin.eof()) // failed
{
```

as it is, without skipping over white space and

Let's look at these two groups of `istream`

## Single-Character Input

When used with a `char` argument or no argument, `get` returns the next input character, even if it is a space, tab, or newline. The `istream` version assigns the input character to its argument. The `get` function returns the input character, converted to an integer type (e.g., `int`).

### The `get(char &)` Member Function

Let's try `get(char &)` first. Suppose you have

```
int ct = 0;
char ch;
cin.get(ch);
while (ch != '\n')
{
 cout << ch;
 ct++;
 cin.get(ch);
}
cout << ct << endl;
```

responding output to this:

IC++clearly.

Worse, the loop would never terminate! Between lines, the code would never assign the newlines, and the loop would never terminate the loop.

The `get(char &)` member function returns a reference to the character it reads. You can then invoke it. This means you can concatenate output like this:

```
char c1, c2, c3;
cin.get(c1).get(c2) >> c3;
```

First, `cin.get(c1)` assigns the first input character to `c1`, which is `cin`. This reduces the code to `cin.get(c2) >> c3`. The function call reads the second input character to `c2`. The function call returns a reference to `c2`. This, in turn, assigns the next non-white-space character to `c3`. The function call then winds up being assigned white space, but `c3` is already assigned.

If `cin.get(char &)` encounters the end of the file (Ctrl+Z for DOS and Windows command prompt, Ctrl+D for Unix), it does not assign a value to the variable. If the program has reached the end of the file, the method calls `setstate(failbit)`, which sets the `failbit` in the `ios_base::iostate` enum.

```
cout << c1 << endl;
```

The `get(void)` member function returns `int` (depending on the character set and locale). The

```
char c1, c2, c3;
cin.get().get() >> c3; // not valid
```

Here `cin.get()` returns a type `int` value. If `cin` is a non-`istream` object, you can't apply the membership operator `>>`. However, you can use `get()` at the end of an extra

```
char c1;
cin.get(c1).get(); // valid
```

The fact that `get(void)` returns type `int` is not a problem for the `>>` operator. But because `cin.get(c1)` returns `cin`, the following particular code would read the first input character, store it in `c1`, and input character and discard it.

Upon reaching the end-of-file, real or simulated, the constant `EOF`, which is a symbolic constant provided by the `iostream` header, allows the following construction for reading

```
int ch;
while ((ch = cin.get()) != EOF)
{
```

space is convenient for offering menu choices

```
cout << "a. annoy client b. bill
 << "c. calm client d. dece
 << "q.\n";
cout << "Enter a, b, c, d, or q: ";
char ch;
cin >> ch;
while (ch != 'q')
{
 switch(ch)
 {
 ...
 }
 cout << "Enter a, b, c, d, or q: ";
 cin >> ch;
}
```

To enter, say, a **b** response, you type **b** and p  
response **b\n**. If you used either form of `get (`  
`\n` character on each loop cycle, but the extra  
programmed in C, you've probably encountered  
to the program as an invalid response. It's an e



input as a string.) The third argument specifies the maximum number of characters to read. The versions with just two arguments use the default value of 1000. The `getline()` function reads up to the maximum characters or until the delimiter character is encountered, whichever comes first.

For example, the following code reads characters from `cin` into a string:

```
char line[50];
cin.getline(line, 50);
```

The `cin.getline()` function quits reading input characters or, by default, after encountering a newline character. The chief difference between `get()` and `getline()` is that `get()` leaves the delimiter character in the input stream, making it the first character of the next line. `getline()` extracts and discards the newline character.

Chapter 4 illustrated using the two-argument versions of `get()` and `getline()`. Now let's look at the three-argument versions. Encountering the delimiter character causes `get()` to stop reading, even if the number of characters hasn't been reached. So, if they reach the end of a line before reading the maximum number of characters, `get()` leaves the delimiter character in the input stream. In the default case, `get()` leaves the delimiter character in the input stream; `getline()` does not.

Listing 17.13 demonstrates how `getline()` works with the `ignore()` member function. `ignore()` takes the maximum number of characters to read and a character to ignore as arguments. It reads and discards

```
int main()
{
 using std::cout;
 using std::cin;
 using std::endl;

 char input[Limit];

 cout << "Enter a string for getline()
 cin.getline(input, Limit, '#');
 cout << "Here is your input:\n";
 cout << input << "\nDone with phase 1

 char ch;
 cin.get(ch);
 cout << "The next input character is

 if (ch != '\n')
 cin.ignore(Limit, '\n'); // di

 cout << "Enter a string for get() pro
 cin.get(input, Limit, '#');
 cout << "Here is your input:\n";
```

```
want my
Done with phase 2
The next input character is #
```

Note that the `getline()` function discards and the `get()` function does not.

## Unexpected String Input

Some forms of input for `get(char *, int)` and the other input functions, encountering end-of-the stream, such as device failure, sets `badbit`. This input that meets or exceeds the maximum number of characters. Let's look at those cases now.

If either method fails to extract any characters from the input string and uses `setstate()` to set `failbit`, how can we extract any characters? One possibility is if an input is off-file. For `get(char *, int)`, another possibility is if the input is off-file.

```
char temp[80];
while (cin.get(temp,80)) // terminates on EOF
 ...
```

too many input characters caused the method to read the entire line. If it's not a newline, then the end. This technique doesn't necessarily work. It reads and discards the newline, so looking at the next character. But if you use `get()`, you have the option of reading a single character. The next section includes an example of how to use `get()` to read a single character. This section summarizes these behaviors.

Table 17.6 **Input Behavior**

| Method                            | Behavior                                                                                                                                                                                                          |
|-----------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>getline(char *, int)</code> | Sets <code>failbit</code> if the number of characters read is greater than <code>int</code> . Sets <code>failbit</code> if the number of characters read is greater than <code>int</code> and more are available. |
| <code>get(char *, int)</code>     | Sets <code>failbit</code> if the number of characters read is greater than <code>int</code> .                                                                                                                     |

## Other `istream` Methods

Other `istream` methods besides the ones discussed above are `gcount()`, and `putback()`. The `read()` function reads characters from the specified location. For example, `read()` reads characters from the standard input and place them in the buffer.

The call to `cin.peek()` peeks at the next input character. Then the `while` loop test condition checks the condition. If this is the case, the loop reads the character into `ch`. When the loop terminates, the period or newline character is positioned to be the first character read by the next iteration. It appends a null character to the array, making it a C string.

The `gcount()` method returns the number of characters read by the extraction method. That means characters read by the `read()` method but not by the extraction operator. For example, suppose you've just read a line of text into the `myarray` array and you want to know how many characters it contains. You could use the `strlen()` function to count the characters, but it's quicker to use `cin.gcount()` to report how many characters were read from the input stream.

The `putback()` function inserts a character back into the input stream. The character then becomes the first character read by the next iteration. The `putback()` method takes one `char` argument, which is the character to be inserted. The argument type is `istream&`, which allows the call to be `cin.putback(ch)`. Using `peek()` is like using `get()` to read a character from the input stream. However, `putback()` puts the character back in the input stream. However, you can't put back a character that is different from the one that was read.

```
 else
 {
 cin.putback(ch); // reinse
 break;
 }
 }

 if (!cin.eof())
 {
 cin.get(ch);
 cout << endl << ch << " is next i
 }
 else
 {
 cout << "End of file reached.\n";
 std::exit(0);
 }

 while(cin.peek() != '#') // look a
 {
 cin.get(ch);
 cout << ch;
 }
 if (!cin.eof())
```

```

while(cin.get(ch)) // terminate
{
 if (ch != '#')
 cout << ch;
 else
 {
 cin.putback(ch); // reinsert char
 break;
 }
}

```

The expression `cin.get(ch)` returns false simulating end-of-file from the keyboard term first, the program puts the character back in th terminate the loop.

The second approach is simpler in appearance

```

while(cin.peek() != '#') // look ahead
{
 cin.get(ch);
 cout << ch;
}

```

```

 if (cin.peek() != '\n')
 cout << "Sorry, we only have enough room for " << name << endl;
 eatline();
 cout << "Dear " << name << ", enter your title: ";
 cin.get(title, SLEN);
 if (cin.peek() != '\n')
 cout << "We were forced to truncate your title: " << title << endl;
 eatline();
 cout << " Name: " << name << endl;
 cout << " Title: " << title << endl;

 return 0;
}

```

---

Here is a sample run of the program in Listing 10.1:

```

Enter your name: Ella Fishsniffer
Sorry, we only have enough room for Ella
Dear Ella Fish, enter your title:
Executive Adjunct
We were forced to truncate your title.

```



and so on. Unless you're programming on the  
have to worry about those things. What you do  
file, a way to have a program read the content  
ate and write to files. Redirection (as discussed  
file support, but it is more limited than explicit  
rection comes from the operating system, not  
tems. This book has already touched on file I/O  
thoroughly.

The C++ I/O class package handles file in-  
input and output. To write to a file, you create  
methods, such as the `<<` insertion operator or  
ifstream object and use the `istream` method  
`get()`. Files require more management than t  
example, you have to associate a newly opened  
read-only mode, write-only mode, or read-and-  
might want to create a new file, replace an old  
want to move back and forth through a file. Th  
eral new classes in the `fstream` (formerly `fst`  
class for file input and the `ofstream` class for t  
class for simultaneous file I/O. These classes ar  
header file, so objects of these new classes are  
learned.

```
ofstream fout("jar.txt"); // create fout
```

When you've gotten this far, you use `fout` in the same manner as `cout`. For example, if you want to write to the file, you can do the following:

```
fout << "Dull Data";
```

Indeed, because `ostream` is a base class for `ofstream`, `ofstream` inherits all `ostream` methods, including the various insertion operators, `<<` methods and manipulators. The `ofstream` class also allocates space for an output buffer when it creates an object. When you create two `ofstream` objects, the program creates two buffers. Each `ofstream` object such as `fout` collects output until the buffer is filled, it transfers the buffer contents to the file. Disk drives are designed to transfer data in large blocks. This approach greatly speeds up the transfer rate of data to and from the disk.

Opening a file for output this way creates a new file. If a file by that name exists prior to opening it for output, the file is truncated so that output starts with a clean file. Later in this chapter, you will learn how to open a file and retain its contents.

```

char ch;
fin >> ch; // read a character
char buf[80];
fin >> buf; // read a word from file
fin.getline(buf, 80); // read a line from file
string line;
getline(fin, line); // read from a file

```

Input, like output, is buffered, so creating an input buffer, which the `fin` object manages. Access is faster than byte-by-byte transfer.

The connections with a file are closed automatically—objects expire—for example, when the program terminates with a file explicitly by using the `close()` function:

```

fout.close(); // close output connection
fin.close(); // close input connection

```

Closing such a connection does not eliminate the file. However, the stream management apparatus object still exists, along with the input buffer. You can reconnect the stream to the same file or to another file.

Let's look at a short example. The program `writefile.cpp` creates a file that has that name, writes some information to the file, flushes the buffer, guaranteeing that the file

```

 fout << "For your eyes only!\n";
 cout << "Enter your secret number: ";
 float secret;
 cin >> secret;
 fout << "Your secret number is " << secret;
 fout.close(); // close file

// create input stream object for new file
ifstream fin(filename.c_str());
cout << "Here are the contents of " << filename << "\n";
char ch;
while (fin.get(ch)) // read character
 cout << ch; // write it to cout
cout << "Done\n";
fin.close();

return 0;
}

```

---

Here is a sample run of the program in Linux:

```

Enter name for new file: pythag
Enter your secret number: 3.14159

```

```
if (fin.fail()) // open attempt failed
{
 ...
}
```

Or because an `ifstream` object, like an `is`, where a `bool` type is expected, you could use

```
fin.open(argv[file]);
if (!fin) // open attempt failed
{
 ...
}
```

However, newer C++ implementations have been opened—the `is_open()` method:

```
if (!fin.is_open()) // open attempt failed
{
 ...
}
```

The reason this is better is that it tests for success or failure, as discussed in the following Caution.

want to count how many times a name appears. You can open a single stream and associate it with each file. This is more resources more effectively than opening a separate stream for each approach, you declare an `ifstream` object with a constructor. A `close()` method to associate the stream with a file. For example, to read two files in succession:

```
ifstream fin; // create stream object
fin.open("fat.txt"); // associate stream with file
... // do stuff
fin.close(); // terminate association
fin.clear(); // reset fin (may be necessary)
fin.open("rat.txt"); // associate stream with file
...
fin.close();
```

We'll look at an example shortly, but first, let's look at how to pass files to a program in a manner that allows the user to specify the files.

## Command-Line Processing

File-processing programs often use command-line arguments. *line arguments* are arguments that appear on the command line.

In this case, `argc` would be 1, `argv[0]` would be the command name, and so on. The following loop would print each command-line argument:

```
for (int i = 1; i < argc; i++)
 cout << argv[i] << endl;
```

Starting with `i = 1` just prints the command-line arguments. If `i = 0` would print the command name as well.

Command-line arguments, of course, go hand in hand with command-line systems such as the Windows command prompt. Even though Windows still allow you to use command-line arguments, there are some caveats:

- Many Windows IDEs (Integrated Development Environments) provide providing command-line arguments. Typically, you click on a set of menu choices that lead to a box into which you type arguments. The exact set of steps varies by IDE and Windows upgrade, so check your documentation.
- Many Windows IDEs can produce executable files that can be run in command prompt mode.

Listing 17.17 combines the command-line arguments with a loop that counts characters in files listed on the command line.

```

 if (!fin.is_open())
 {
 cerr << "Could not open " <<
 fin.clear();
 continue;
 }
 count = 0;
 while (fin.get(ch))
 count++;
 cout << count << " characters in
 total += count;
 fin.clear(); // needed
 fin.close(); // disconn
 }
 cout << total << " characters in all

 return 0;
}

```

---

## Note

Some C++ implementations require using `fin.clear()` if file descriptors do not. It depends on whether associating a file descriptor with a stream actually resets the stream state. In does no harm



filename or by using the `open()` method, you specifies the file mode:

```
ifstream fin("banjo", mode1); // construct
ofstream fout();
fout.open("harp", mode2); // open() w
```

The `ios_base` class defines an `openmode` type and `iostate` types, it is a bitmask type. (In the from several constants defined in the `ios_base` the constants and their meanings. C++ file I/O compatible with ANSI C file I/O.

**Table 17.7 File Mode Constants**

| Constant                      | Meaning                     |
|-------------------------------|-----------------------------|
| <code>ios_base::in</code>     | Open file for reading.      |
| <code>ios_base::out</code>    | Open file for writing.      |
| <code>ios_base::ate</code>    | Seek to end-of-file upon    |
| <code>ios_base::app</code>    | Append to end-of-file.      |
| <code>ios_base::trunc</code>  | Truncate file if it exists. |
| <code>ios_base::binary</code> | Binary file.                |

You can expect to find some differences as an example, some allow you to omit the `ios_base::trunc` don't. If you aren't using the default mode, then you have to specify the mode explicitly. Some compilers don't support all the modes, so you may offer choices beyond those in the table. Of course, you may have to make some alterations in the code to make it work in your system. The good news is that the development of C++ is moving towards more uniformity.

Standard C++ defines parts of file I/O in the `<fstream>` header. A C++ statement like

```
ifstream fin(filename, cplusplusmode);
```

is implemented as if it uses the C `fopen()` function:

```
fopen(filename, cmode);
```

Here `cplusplusmode` is a type `openmode` value, such as `ios_base::in`, corresponding to the C-mode string, such as `"r"`. Table 1 shows the mapping of C++ modes and C modes. Note that `ios_base::trunc` doesn't cause truncation when combined with `ios_base::in`, such as `ios_base::in | ios_base::trunc`, because the `is_open()` method detects this failure.

---

| C++ mode                                                                     | C mode               |
|------------------------------------------------------------------------------|----------------------|
| <code>ios_base::in</code>                                                    | <code>"r"</code>     |
| <code>ios_base::out</code>                                                   | <code>"w"</code>     |
| <code>ios_base::out  </code><br><code>ios_base::trunc</code>                 | <code>"w"</code>     |
| <code>ios_base::out  </code><br><code>ios_base::app</code>                   | <code>"a"</code>     |
| <code>ios_base::in  </code><br><code>ios_base::out</code>                    | <code>"r+"</code>    |
| <code>ios_base::in   ios_base</code><br><code>::out   ios_base::trunc</code> | <code>"w+"</code>    |
| <code>c++mode   ios_base::binary</code>                                      | <code>"modeb"</code> |
| <br>                                                                         |                      |
| <code>c++mode   ios_base::ate</code>                                         | <code>"mode"</code>  |

---

older compilers don't quite conform to the C++ standard, modes such as `nocreate` that are not part of the standard. Others require the `fin.clear()` call before opening.

### Listing 17.18    **append.cpp**

```
// append.cpp -- appending information to a file
#include <iostream>
#include <fstream>
#include <string>
#include <cstdlib> // (for exit())

const char * file = "guests.txt";
int main()
{
 using namespace std;
 char ch;
```

```
 while (getline(cin,name) && name.size() > 0)
 {
 fout << name << endl;
 }
 fout.close();

// show revised file
 fin.clear(); // not necessary for s
 fin.open(file);
 if (fin.is_open())
 {
 cout << "Here are the new contents\n";
 << file << " file:\n";
 while (fin.get(ch))
 cout << ch;
 fin.close();
 }
 cout << "Done.\n";
 return 0;

 }

```

---

Greta Greppo  
LaDonna Mobile  
Fannie Mae

Here are the new contents of the guests.txt  
Ghengis Kant  
Hank Attila  
Charles Bigg  
Greta Greppo  
LaDonna Mobile  
Fannie Mae  
Done.

You should be able to read the contents of  
the editor you use to write your source code

## Binary Files

When you store data in a file, you can store it in  
text form means you store everything as text, even  
2.324216e+07 in text form means storing the number  
as text. That requires converting the computer's internal  
representation of the number to character form, and that's exactly what  
fprintf does. Writing in binary form, on the other hand, means storing the number

Each format has advantages. The text format is easy to read and can be edited with an ordinary editor or word processor to read and modify. It is portable from one computer system to another. The binary format is better for numbers because it stores the exact internal representation, without truncation errors or round-off errors. Saving data in binary format is faster because of no format conversion and because you may be able to save more data in the same format. The binary format usually takes less space, depending on the data. Portability of the binary format system can be a problem, however, if the new system has a different representation for values. Even different compilers on the same system can have different representations for structure layouts. In these cases, you need a program to translate one data format to another.

Let's look at a more concrete example. Consider the following C declaration:

```
const int LIM = 20;
struct planet
{
 char name[LIM]; // name of planet
 double population; // its population
 double g; // its acceleration
};
planet pl;
```





**Tip**

The `read()` and `write()` member functions recover data that has been written to a file with

Listing 17.19 uses these methods to create a file. It is similar to Listing 17.18, but it uses `write()` to write data to the file and the `get()` method. It also uses manipulators

**Note**

Although the binary file concept is part of the C++ standard, not all compilers provide support for the binary file mode. Some operating systems have only one file type in the first place, so they implement `read()` and `write()` with the standard file operations. This rejects `ios_base::binary` as a valid constant. If your implementation doesn't support the binary file mode, use `cout.setf(ios_base::fixed, ios_base::left);` and `cout.setf(ios_base::right, ios_base::left);` to substitute `ios` for `ios_base`. Other compilers, particularly older ones, have idiosyncrasies.

```

// show initial contents
 ifstream fin;
 fin.open(file, ios_base::in | ios_base::app);
 //NOTE: some systems don't accept the ios_base::app flag
 if (fin.is_open())
 {
 cout << "Here are the current contents of the file\n";
 while (fin.read((char *) &pl, sizeof Planet))
 {
 cout << setw(20) << pl.name << " : " << pl.mass << endl;
 cout << setprecision(0) << setw(20) << pl.radius << endl;
 cout << setprecision(2) << setw(20) << pl.density << endl;
 }
 fin.close();
 }

// add new data
 ofstream fout(file,
 ios_base::out | ios_base::app);
 //NOTE: some systems don't accept the ios_base::app flag
 if (!fout.is_open())
 {

```

```

 fin.open(file, ios_base::in | ios_base::out);
 if (fin.is_open())
 {
 cout << "Here are the new contents\n";
 << file << " file:\n";
 while (fin.read((char *) &pl, sizeof(Planet)))
 {
 cout << setw(20) << pl.name << " " << pl.population << " " << pl.acceleration << "\n";
 << setprecision(0) << setprecision(2) << setprecision(2) << setprecision(2) << "\n";
 }
 fin.close();
 }
 cout << "Done.\n";
 return 0;
 }
}

```

---

Here is a sample initial run of the program

```

Enter planet name (enter a blank line to quit):
Earth
Enter planetary population: 6928198253
Enter planet's acceleration of gravity: 9.8

```

```
while (std::cin.get() != '\n') continue;
```

This reads and discards input up through the first newline character. The `cin.get()` input statement in the loop:

```
cin.get(pl.name, 20);
```

If the newline were left in place, this statement would read the rest of the line, terminating the loop.

You might wonder if this program could use a character array for the name member of the planet structure instead of a string. It would require some out major changes in design. The problem is that a string is not a character array; the string within itself; instead, it contains a pointer to the string. So if you copy the structure to a new one, you just copy the address of where the string was stored. That address is meaningless.

## Random Access

For our last file example, let's look at random access. Instead of moving to any location in the file instead of moving to the beginning, this approach is often used with database files. A pointer is given the location of data in the main data file. Then you can read the data there, and perhaps modify it. This is the most efficient way to access data in a file.



class was replaced with the `basic_istream` were replaced with the template-based types `streampos` and `streamoff` continue to exist `off_type`. Similarly, you can use the `wstream` `seekg()` with a `wistream` object.

Let's take a look at the arguments to the `seekg` `streamoff` type are used to measure offsets, `ios_base::off_type`. The `streamoff` argument represents the file position relative to one of three locations. (The type may be defined as `streamoff` or `streampos`.) The `seek_dir` argument is another integer type that defines the base, in the `ios_base` class. The constant `ios_base::beg` is the beginning of the file. The constant `ios_base::cur` is the current position. The constant `ios_base::end` is the end of the file. Here are some sample calls, assuming

```
fin.seekg(30, ios_base::beg); // 30 bytes from beginning
fin.seekg(-1, ios_base::cur); // back one byte
fin.seekg(0, ios_base::end); // go to end of file
```

Now let's look at the second prototype. `va_arg` is a macro that expands to the type of the argument in a file. It can be a class, but, if so, the class must have a `streampos` member and a constructor with an integer argument. The `streampos` argument is a `streampos` value and a constructor with an integer argument to `streampos` values. A `streampos` value represents the position in the file.

```

 {
 finout.seekg(0); // go to beginning
 cout << "Here are the current contents"
 << file << " file:\n";
 while (finout.read((char *) &pl, sizeof pl))
 {
 cout << ct++ << ": " << setw(LIM)
 << setprecision(0) << setw(12) << pl
 << setprecision(2) << setw(6) << pl << "
 }
 if (finout.eof())
 finout.clear(); // clear eof flag
 else
 {
 cerr << "Error in reading " << file << "\n";
 exit(EXIT_FAILURE);
 }
 }
else
{
 cerr << file << " could not be opened\n";
 exit(EXIT_FAILURE);
}

```

program asks the user to enter a record number. If the user enters a value that is not within the range of bytes in a record yields the byte number for the record number, the desired byte number

```
cout << "Enter the record number you wish
long rec;
cin >> rec;
eatline(); // get rid of new
if (rec < 0 || rec >= ct)
{
 cerr << "Invalid record number -- bye
 exit(EXIT_FAILURE);
}
streampos place = rec * sizeof pl; // co
finout.seekg(place); // random access
```

The variable `ct` represents the number of records beyond the limits of the file.

Next, the program displays the current record

```
finout.read((char *) &pl, sizeof pl);
cout << "Your selection:\n";
cout << rec << ": " << setw(LIM) << pl.name
<< setprecision(0) << setw(12) << pl.population
<< setprecision(2) << setw(6) << pl.gpa <<
if (finout.eof())
 finout.clear(); // clear eof flag
```



## Note

The older the implementation, the more likely systems don't recognize the binary flag, the f

### Listing 17.20 **random.cpp**

```
// random.cpp -- random access to a binary
#include <iostream> // not required by
#include <fstream>
#include <iomanip>
#include <cstdlib> // for exit()
const int LIM = 20;
struct planet
{
 char name[LIM]; // name of planet
 double population; // its population
 double g; // its acceleration
};

const char * file = "planets.dat"; // ASS
inline void eatline() { while (std::cin.get() != '\n') continue; }

int main()
{
```

```

 else
 {
 cerr << "Error in reading " <
 exit(EXIT_FAILURE);
 }
 }
else
{
 cerr << file << " could not be op
 exit(EXIT_FAILURE);
}

// change a record
cout << "Enter the record number you
long rec;
cin >> rec;
eatline(); // get rid of
if (rec < 0 || rec >= ct)
{
 cerr << "Invalid record number --
 exit(EXIT_FAILURE);
}
streampos place = rec * sizeof pl; /
finout.seekg(place); // random acc

```

```
 if (finout.fail())
 {
 cerr << "Error on attempted write\n";
 exit(EXIT_FAILURE);
 }

// show revised file
 ct = 0;
 finout.seekg(0); // go to beginning
 cout << "Here are the new contents of\n";
 << " file:\n";
 while (finout.read((char *) &pl, sizeof(pl)))
 {
 cout << ct++ << ": " << setw(LIM)
 << setprecision(0) << setw(12)
 << setprecision(2) << setw(6)
 << "\n";
 }
 finout.close();
 cout << "Done.\n";
 return 0;
}
```

---

```
5: Taanagoot: 361000004 10.
6: Marin: 252409 9.
Done.
```

By using the techniques in this program, you can add new material and delete records. If you were given the idea to reorganize it by using classes and functions, you could change the `planet` structure to a class definition; then you could use `cout << p1` to display the class data member values, so that `cout << p1` displays the class data member values. This example doesn't bother to verify input, so you can add input verification where appropriate.

## Working with Temporary Files

Developing applications often requires the use of temporary files. The user is transient and must be controlled by the program. How do you manage this in C++? It's really quite easy to create a temporary file, and delete the file. First of all, you need to know how to create a temporary file(s). But wait...how can you ensure that the file is deleted? The `tmpnam()` standard function declared in `<conio.h>` creates a temporary file name.

```
char* tmpnam(char* pszName);
```

The `tmpnam()` function creates a temporary file name. The file name is pointed to by `pszName`. The constants `L_tmpnam` and `MAX_PATH` are defined in `<conio.h>`.

more generally, by using `tmpnam()`, you can name files up to `L_tmpnam` characters per name. The name of the file you can run this program to see what names you

## Incore Formatting

The `iostream` family supports I/O between a program and a `string` object. That is, you can use `cout` to write formatted information into a `string` object, and methods such as `getline()` to read information from a `string` object. Reading formatted information from a `string` object is termed *incore* formatting. (The `sstream` family of `string` support superclasses also support *incore* support.)

The `sstream` header file defines an `ostrstream` class. (There is also a `wostream` class for wide character sets.) If you create an `ostrstream` object, you can use `cout` to write formatted information into it. You can use the same methods with an `ostrstream` object as you can with `cout`. That is, you can do something like the following:

```
ostrstream ostr;
double price = 380.0;
```

```
 string hdisk;
 cout << "What's the name of your hard disk? ";
 getline(cin, hdisk);
 int cap;
 cout << "What's its capacity in GB? ";
 cin >> cap;
 // write formatted information to ostream
 ostr << "The hard disk " << hdisk << " has a capacity of "
 << cap << " gigabytes.\n";
 string result = ostr.str(); // save the formatted string
 cout << result; // show the formatted string

 return 0;
}
```

---

Here's a sample run of the program in Listing 13.1:

```
What's the name of your hard disk? Datarapture
What's its capacity in GB? 2000
The hard disk Datarapture has a capacity of 2000 gigabytes.
```

The `istringstream` class lets you use the formatted string from an `istringstream` object, which can be used to write to a file or to a stream.

```
 string lit = "It was a dark and stormy
 " the full moon glowed br
 istringstream instr(lit); // use buf
 string word;
 while (instr >> word) // read a
 cout << word << endl;
 return 0;

 }
```

---

Here is the output of the program in Listing 10.1:

```
It
was
a
dark
and
stormy
day,
and
the
full
moon
glowed
brilliantly.
```

The I/O class library provides a variety of versions of the extraction operator (>>) that can convert character input to those types. The `getline` method provide further support for single-character input. The `ostream` class defines versions of the insertion operator (<<) for the basic C++ types and that convert them to strings. The `wostream` provides further support for single-character output. The `wostream` provide similar support for wide characters.

You can control how a program formats output by using manipulators (functions that can be used with the `ostream` and `wostream` files. These methods allow you to set the base, the field width, the number of decimal places for floating-point values, and other elements.

The `fstream` file provides class definitions for file I/O. The `ifstream` class derives from the `istream` class. If you have an object with a file, you can use all the `istream` methods. Associating an `ofstream` object with a file lets you use the `ostream` methods. And associating an `fstream` object with a file lets you use the `fstream` methods with the file.

To associate a file with a stream, you can pass the file name to the stream object or you can first create a file stream object and then use `open` to associate the stream with a file. The `close` method



1. What role does the `iostream` file play in C++?
2. Why does typing a number such as 121 result in an error? What version?
3. What's the difference between the standard output and error streams?
4. Why is `cout` able to display various C++ types without explicit instructions for each type?
5. What feature of the output method definition allows it to handle various types?
6. Write a program that requests an integer, a floating-point number, and a hexadecimal form. Display each form on a new line, 10 characters wide, and use the C++ number base constants.
7. Write a program that requests the following information and displays it as shown:

Enter your name: **Billy Gruff**

Enter your hourly wages: **12**

Enter number of hours worked: **7.5**

First format:

Billy Gruff: \$

Second format:

Billy Gruff : \$12.

```

 cin.get(ch);
 }
 cout << "ctl = " << ctl << "; " << endl;

 return 0;
}

```

What does it print, given the following

I see a q<Enter>

I see a q<Enter>

Here <Enter> signifies pressing the Enter key.

- Both of the following statements read a line of input until the end of a line. In what way does the behavior differ?

```

while (cin.get() != '\n')
 continue;
cin.ignore(80, '\n');

```

The resulting file would have these contents:

```
eggs kites donuts zero lassitude
balloons hammers finance drama
stones
```

5. Mat and Pat want to invite their friends to a picnic, following Exercise 8 in Chapter 16, except now they have two lists of friends. They have asked you to write a program that:
  - Reads a list of Mat's friends' names from `mat.dat`, one friend per line. The names are already in sorted order.
  - Reads a list of Pat's friends' names from `pat.dat`, one friend per line. The names are already in sorted order.
  - Merges the two lists, eliminating duplicates, and writes the result to `matnpat.dat`, one friend per line.

ate data from the user:

```
pc[i]->setall(); // invokes function
```

To save the data to a file, devise a virtual

```
for (i = 0; i < index; i++)
 pc[i]->writeall(fout); // fout of
```

## Note

Use text I/O, not binary I/O, for Programming. Do not include pointers to tables of pointers to virtual functions to a file. An object filled by using `read()` and pointers, which really messes up the behavior of the file to separate each data field from the next; then you could still use binary I/O, but not write out the class methods that apply the `write()` and `read()` usually rather than to the object as a whole. The intended data to a file.

```
int main()
{
 using namespace std;
 vector<string> vostr;
 string temp;

 // acquire strings
 cout << "Enter strings (empty line to stop): ";
 while (getline(cin,temp) && temp != "")
 vostr.push_back(temp);
 cout << "Here is your input.\n";
 for_each(vostr.begin(), vostr.end(), [](const string& s) {
 cout << s << " ";
 });

 // store in a file
 ofstream fout("strings.dat", ios_base::out);
 for_each(vostr.begin(), vostr.end(), [&fout](const string& s) {
 fout << s << "\n";
 });
 fout.close();

 // recover file contents
 vector<string> vistr;
 ifstream fin("strings.dat", ios_base::in);
 if (!fin.is_open())
 {
 cerr << "Could not open file\n";
 }
 while (fin.getline(temp, 1024))
 vistr.push_back(temp);
 for_each(vistr.begin(), vistr.end(), [](const string& s) {
 cout << s << " ";
 });
}
```

The `data()` member returns a pointer to the string. It's similar to the `c_str()` member, which returns a pointer to a null-terminated character array.

- Write a `GetString()` function that uses `read()` to obtain the size of a string, then copies that many characters from the file, appends a null character, and returns a pointer to the string. Because a string's data is stored in a character array, you get data into the string rather than



- Variadic templates

This chapter concentrates on the new C++11 features. The book already has covered several C++11 features. Then we'll look at some additional features in the C++11 additions that are beyond the scope of this draft. Since the C++11 draft is over 80% longer than C++98, we won't examine the BOOST library.

## C++11 Features Revisited

By now you may have lost track of the many features added. This section reviews them briefly.

### New Types

C++11 adds the `long long` and `unsigned long long` (or wider) and the `char16_t` and `char32_t` types for wide character representations, respectively. It also adds the `std::string_view` type. This section discusses these additions.



```

Stump s1(5,15.6); // Old style
Stump s2{5, 43.4}; // C++11
Stump s3 = {4, 32.1}; // C++11

```

However, if a class has a constructor whose template, then only that constructor can use the list-initialization were discussed in Chapters 3

## Narrowing

The initialization-list syntax provides protection against assigning a numeric value to a numeric type narrower than the value. This initialization allows you to do things that may

```

char c1 = 1.57e27; // double-to-char, undefined
char c2 = 459585821; // int-to-char, undefined

```

If you use initialization-list syntax, however, you can avoid these errors. Attempts that attempt to store values in a type “narrower” than the value are

```

char c1 {1.57e27}; // double-to-char, compile error
char c2 = {459585821}; // int-to-char, out of range error

```

However, conversions to wider types are allowed. Conversions to a type wider than the value is allowed if the value is within the range allowed

```

...
}
double sum(std::initializer_list<double>
{
 double tot = 0;
 for (auto p = il.begin(); p !=il.end()
 tot += *p;
 return tot;
}

```

## Declarations

C++11 implements several features that simplify programming arising when templates are used.

### **auto**

C++11 strips the keyword `auto` of its former meaning (Chapter 9, “Memory Models and Namespaces”) and introduces automatic type deduction, provided that the programmer sets the type of the variable to the type of the expression.

```

auto maton = 112; // maton is type int
auto pt = &maton; // pt is type int *
double fm(double, int);
auto pf = fm; // pf is type double (double (*)(double, int))

```

```

template<typename T, typename U>
void ef(T t, U u)
{
 decltype(T*U) tu;
 ...
}

```

Here, `tu` is of whatever type results from the expression `t*U`. This is well-defined. For example, if `T` is `char` and `U` is `short`, the result is `short` because of automatic integer promotions that take place.

The workings of `decltype` are more complex than those of `static_cast`. Types can be references and can be `const-qualified`. Here are some more examples:

```

int j = 3;
int &k = j;
const int &n = j;
decltype(n) i1; // i1 type const int
decltype(j) i2; // i2 type int
decltype((j)) i3; // i3 type int &
decltype(k + 1) i4; // i4 type int

```

See Chapter 8, “Adventures in Functions,”

had typedef for that purpose:

```
typedef std::vector<std::string>::iterator
```

C++11 provides a second syntax (discussed later) for creating aliases:

```
using itType = std::vector<std::string>::
```

The difference is that the new syntax also works for namespaces, whereas typedef can't:

```
template<typename T>
using arr12 = std::array<T,12>; // tem
```

This statement specializes the array<T, 12> template. For instance, consider these declarations:

```
std::array<double, 12> a1;
std::array<std::string, 12> a2;
```

They can be replaced with the following:

```
arr12<double> a1;
arr12<std::string> a2;
```

unique\_ptr, shared\_ptr, and weak\_ptr. Chapter 14, “Smart Pointers,” covers this in detail. All the new smart pointers have been designed with move semantics.

## Exception Specification Changes

C++ provides a syntax for specifying what exceptions a function can throw. This is covered in detail in Chapter 15, “Friends, Exceptions, and Move Semantics.”

```
void f501(int) throw(bad_dog); // can throw bad_dog
void f733(long long) throw(); // doesn't throw anything
```

As with auto\_ptr, the collective experience of the C++ community was that, in practice, exception specifications didn't work well. The C++11 standard deprecates exception specifications, and the community felt that there is some value in documenting this change. C++11 felt that there is some value in documenting this change, and it added the keyword noexcept for this purpose.

```
void f875(short, short) noexcept; // doesn't throw anything
```

## Scoped Enumerations

Traditional C++ enumerations provide a way to define a set of named constants. They come with a rather low level of type checking. C++11 introduced scoped enumerations, which are defined within the scope that encloses the enumeration definition.

conversions for classes could be established. Coming experience accumulated was that autolems in the form of unexpected conversions. problem by introducing the keyword `explicit` invoked by one-argument constructors:

```
class Plebe
{
 Plebe(int); // automatic int-to-plebe
 explicit Plebe(double); // requires
 ...
};
...
Plebe a, b;
a = 5; // implicit conversion,
b = 0.5; // not allowed
b = Plebe(0.5); // explicit conversion
```

C++11 extends the use of `explicit` (discussed in “C++11: New Features and Classes”) so that conversion functions can be

```
class Plebe
{
 ...
 // conversion functions
```

```

 Session(int n, double d, short s) : mem1(n), mem2(d), mem3(s) {}
 ...
};

```

You can use the equal sign or the brace for the smaller version. The result is the same as if you put the member initialization list entries for `mem1` and `mem2` in the first constructor.

```

Session() : mem1(10), mem2(1966.54) {}
Session(short s) : mem1(10), mem2(1966.54), mem3(s) {}

```

Note how using in-class initialization avoids the need for static variables, thus reducing work and the number of objects created by the programmer.

These default values are overridden by a constructor that takes a member initialization list, so the third constructor overrides the default values for `mem1` and `mem2`.

## Template and STL Changes

C++11 makes several changes extending the Standard Template Library in particular. Some of these changes are new features, some are changes of use. In this chapter we've already mentioned smart pointers.

```
x = std::rand();
```

## New STL Containers

C++11 adds `forward_list`, `unordered_map`, `unordered_multiset` to its collection of STL containers. The `forward_list` container is a singly linked list. It's simpler and more economical of space than the other four containers support implementing.

C++11 also adds the `array` template (discussed in Chapter 16), for which one specifies an element type and a size:

```
std::array<int,360> ar; // array of 360 ints
```

This template class does not satisfy all the requirements of a container because the size is fixed, you can't use any methods to change the size of a container. But `array` does have the `begin()` and `end()` methods you to use many of the range-based STL algorithms.

## New STL Methods

C++11 adds `cbegin()` and `cend()` to the list of container methods. The new methods return iterators to the first and one past the last element of the container, thus specifying a range encompassing all the elements.



To avoid confusion with the `>>` operator, C++11 removed nested template declarations:

```
std::vector<std::list<int> > vl; // >> not ok
```

C++11 removes that requirement:

```
std::vector<std::list<int>> vl; // >> ok
```

## The rvalue Reference

The traditional C++ reference, now called an lvalue. An lvalue is an expression, such as a variable, that represents data for which the program can obtain a memory address that could appear on the left side of an assignment. The `const` modifier is allowed for constructs that cannot be modified.

```
int n;
int * pt = new int;
const int b = 101; // can't assign to b,
int & rn = n; // n identifies datum
int & rt = *pt; // *pt identifies datum
const int & rb = b; // b identifies constant datum
```

C++11 adds the rvalue reference (discussed in Chapter 12), which can bind to rvalues—that is, values that can appear only on the right side of an assignment.

```
#include <iostream>
```

```
inline double f(double tf) {return 5.0*(t
int main()
{
 using namespace std;
 double tc = 21.5;
 double && rd1 = 7.07;
 double && rd2 = 1.8 * tc + 32.0;
 double && rd3 = f(rd2);
 cout << " tc value and address: " <<
 cout << "rd1 value and address: " <<
 cout << "rd2 value and address: " <<
 cout << "rd3 value and address: " <<
 cin.get();
 return 0;
}
```

---

Here is a sample output:

```
tc value and address: 21.5, 002FF744
rd1 value and address: 7.07, 002FF728
rd2 value and address: 70.7, 002FF70C
rd3 value and address: 21.5, 002FF6F0
```

constructor, which will use new to allocate memory. If we have a vector of 20,000 strings, each 1,000 characters long, then 20,000,000 characters will be copied from the old memory to the new memory controlled by `vstr_copy1`. That's a lot of copying to be done.

But does it “got to be done?” There are times when we can avoid this. Suppose we have a function that returns a vector of strings:

```
vector<string> allcaps(const vector<string>& vstr)
{
 vector<string> temp;
 // code that stores an all-uppercase version of each string in temp
 return temp;
}
```

Next, suppose we use it this way:

```
vector<string> vstr;
// build up a vector of 20,000 strings, each 1,000 characters long
vector<string> vstr_copy1(vstr);
vector<string> vstr_copy2(allcaps(vstr));
```

Superficially, statements #1 and #2 are similar: they both create a new `vector<string>` object. If we take this literally, then statement #1 has `temp` object managing 20,000,000 characters, and statement #2 goes through the effort of creating a 20,000,000 character

## A Move Example

Let's look at an example to see how move semantics work. Listing 18.2 defines and uses the `Useless` class, which has a regular copy constructor, and a move constructor, and a rvalue reference. In order to illustrate the problem, the destructor are unusually verbose, and the class manages a number of objects. Also some important methods are omitted. (Despite these omissions, the `Useless` class is an eco-friendly `Use_Less` class.)

### Listing 18.2 `useless.cpp`

---

```
// useless.cpp -- an otherwise useless class
#include <iostream>
using namespace std;

// interface
class Useless
{
private:
 int n; // number of elements
 char * pc; // pointer to data
 static int ct; // number of objects
```

```
}
```

```
Useless::Useless(int k) : n(k)
{
 ++ct;
 cout << "int constructor called; number is " << ct << endl;
 pc = new char[n];
 ShowObject();
}
```

```
Useless::Useless(int k, char ch) : n(k)
{
 ++ct;
 cout << "int, char constructor called; number is " << ct << endl;
 pc = new char[n];
 for (int i = 0; i < n; i++)
 pc[i] = ch;
 ShowObject();
}
```

```
Useless::Useless(const Useless & f): n(f.n)
{
 ++ct;
 cout << "copy constructor called; number is " << ct << endl;
 pc = new char[n];
 for (int i = 0; i < n; i++)
 pc[i] = f.pc[i];
 ShowObject();
}
```

```
 delete [] pc;
 }
```

```
Useless Useless::operator+(const Useless
{
 cout << "Entering operator+()\n";
 Useless temp = Useless(n + f.n);
 for (int i = 0; i < n; i++)
 temp.pc[i] = pc[i];
 for (int i = n; i < temp.n; i++)
 temp.pc[i] = f.pc[i - n];
 cout << "temp object:\n";
 cout << "Leaving operator+()\n";
 return temp;
}
```

```
void Useless::ShowObject() const
{
 cout << "Number of elements: " << n;
 cout << " Data address: " << (void *)
}
}
```

```

 three.ShowData();
 cout << "object four: ";
 four.ShowData();
 }
}

```

---

The crucial definitions are those of the two constructors. In addition to the output statements, here is the copy constructor:

```

Useless::Useless(const Useless & f): n(f.n)
{
 ++ct;
 pc = new char[n];
 for (int i = 0; i < n; i++)
 pc[i] = f.pc[i];
}

```

It does the usual deep copy, and it is the copy constructor statement:

```

Useless two = one; // calls copy constructor

```

The reference `f` refers to the lvalue object





```
Number of elements: 10 Data address: 0xa50
copy const called; number of objects: 2
Number of elements: 10 Data address: 0xa50
int, char constructor called; number of ob
Number of elements: 20 Data address: 0xa50
Entering operator+()
int constructor called; number of objects
Number of elements: 30 Data address: 0xa50
temp object:
Leaving operator+()
object one: xxxxxxxxxxxx
object two: xxxxxxxxxxxx
object three: oooooooooooooooooooooo
object four: xxxxxxxxxxxxoooooooooooooooo
destructor called; objects left: 3
deleted object:
Number of elements: 30 Data address: 0xa50
destructor called; objects left: 2
deleted object:
Number of elements: 20 Data address: 0xa50
destructor called; objects left: 1
deleted object:
Number of elements: 10 Data address: 0xa50
```

directs initialization for object four to the move semantics is coding the move constructor.

In short, the presence of one constructor or another with an rvalue reference sorts possibilities. Objects initialized with an lvalue object use the copy constructor. Objects initialized with an rvalue object use the move constructor. The two constructors have different behaviors.

This raises the question of what happens in C++ language. If there is no move constructor and no need for the copy constructor, what should happen? The compiler would invoke the copy constructor:

```
Useless four (one + three);
```

But an lvalue reference doesn't bind to an rvalue. From Chapter 8, a `const` reference parameter binds to an rvalue if the actual argument is an rvalue:

```
int twice(const & rx) {return 2 * rx;}
...
int main()
{
 int m = 6;
 // below, rx refers to m
 int n = twice(m);
```

Number of elements: 30 Data address: 01C337C0  
...

First, within the `Useless::operator+()` method, we allocate temporary storage for 30 elements at location `01C337C0`. Then, we create a temporary copy to which `f` will refer, copying the contents of `Useless`. Next, `temp`, which uses location `01C337C4`, gets constructed, reusing the recently freed memory address. Then, `temp` object, which used location `01C337E8`, gets destroyed. Finally, `Useless` is constructed, and two of them were destroyed. These techniques are meant to eliminate.

As the `g++` example shows, an optimizing compiler can do this on its own, but using an rvalue reference lets the programmer do it more appropriate.

## Assignment

The same considerations that make move semantics appropriate for assignment. Here, for example, we use the assignment and the move assignment operators.

```
Useless & Useless::operator=(const Useless& other)
{
 if (this == &f)
```

object. It's important that only one pointer points to the object. The move constructor copies the pointer in the source object to the null pointer in the destination object.

As with the move constructor, the move assignment operator is a lvalue reference because the method alters the source object.

## Forcing a Move

Move constructors and move assignment operators are not used with lvalues. How do we force the compiler to use them with lvalues? For instance, a program that has a set of candidate objects, select one object for further processing. It is convenient if you could use a move constructor to move the selected object. However, suppose you try to write

```
Useless choices[10];
Useless best;
int pick;
... // select one object, set pick to index of best
best = choices[pick];
```

The `choices[pick]` object is an lvalue, so the assignment operator, not the move assignment operator, is used. If `choices[pick]` look like an rvalue, then the move constructor or move assignment operator can be used. This can be done by using the `static_cast` to convert the lvalue to an rvalue: `Useless && choices[pick]`. C++11 provides a simpler way

```

 Useless(Useless && f); // move co
 ~Useless();
 Useless operator+(const Useless & f) co
 Useless & operator=(const Useless & f)
 Useless & operator=(Useless && f);
 void ShowData() const;
};

```

```

// implementation
int Useless::ct = 0;

```

```

Useless::Useless()
{
 ++ct;
 n = 0;
 pc = nullptr;
}

```

```

Useless::Useless(int k) : n(k)
{
 ++ct;
 pc = new char[n];
}

```

```
Useless::~~Useless()
{
 delete [] pc;
}
```

```
Useless & Useless::operator=(const Useless & f)
{
 std::cout << "copy assignment operator
 if (this == &f)
 return *this;
 delete [] pc;
 n = f.n;
 pc = new char[n];
 for (int i = 0; i < n; i++)
 pc[i] = f.pc[i];
 return *this;
}
```

```
Useless & Useless::operator=(Useless && f)
{
 std::cout << "move assignment operator
 if (this == &f)
```

```
}
```

```
void Useless::ShowData() const
{
 if (n == 0)
 std::cout << "(object empty)";
 else
 for (int i = 0; i < n; i++)
 std::cout << pc[i];
 std::cout << std::endl;
}

// application
int main()
{
 using std::cout;
 {
 Useless one(10, 'x');
 Useless two = one + one; // calls
 cout << "object one: ";
 one.ShowData();
 cout << "object two: ";
 two.ShowData();
 }
}
```

```
object two: xxxxxxxxxxxxxxxxxxxxxxxx
three = one
copy assignment operator called:
now object three = xxxxxxxxxx
and object one = xxxxxxxxxx
four = one + two
move assignment operator called:
now object four = xxxxxxxxxxxxxxxxxxxxxxxx
four = move(one)
move assignment operator called:
now object four = xxxxxxxxxx
and object one = (object empty)
```

As you can see, assigning one to three invokes `move` assignment operator. Assigning `move(one)` to four invokes `move` assignment operator.

You should realize that the `std::move()` operation. Suppose, for instance, that Chunk is defined as follows. The following code:

```
Chunk one;
...
Chunk two;
two = std::move(one); // move semantics?
```



The default constructor, recall, is a constructor the compiler provides one if you fail to define any. If a class has a default constructor, the absence of a default constructor is termed the *defaulted* constructor. The defaulted constructor leaves members of the built-in type uninitialized. The compiler provides default constructors for members that are class types.

Also the compiler provides a defaulted copy constructor if your code requires its use, and it now provides a defaulted move constructor if you don't provide one and if your code requires it. If a class has two defaulted constructors have the following form:

```
Someclass::Someclass(const Someclass &);
Someclass::Someclass(Someclass &&);
```

In similar circumstances, the compiler provides a defaulted move assignment operator with the following form:

```
Someclass & Someclass::operator=(const Someclass &);
Someclass & Someclass::operator=(Someclass &&);
```

Finally, the compiler provides a destructor.

There are various exceptions to this description. If a class has a copy constructor or a copy assignment operator, it may not have a move constructor or a move assignment operator. If a class has a move constructor or a move assignment operator, the compiler will not provide a copy constructor or a copy assignment operator.

```
Someclass & operator +(const Someclass
```

```
...
};
```

The compiler provides the same constructor if you had you not provided the move constructor.

The `delete` keyword, on the other hand, allows you to disable a particular method. For example, to prevent the compiler from generating a copy constructor, you can disable the copy constructor and copy assignment operator by using the `delete` keyword:

```
class Someclass
{
public:
 Someclass() = default; // use compiler-generated default constructor
 // disable copy constructor and copy assignment operator
 Someclass(const Someclass &) = delete;
 Someclass & operator=(const Someclass &) = delete;
 // use compiler-generated move constructor
 Someclass(Someclass &&) = default;
 Someclass & operator=(Someclass &&) = default;
 Someclass & operator+(const Someclass &) = default;
 ...
};
```

Then suppose we have the following code

```
Someclass sc;
sc.redo(5);
```

The `int` value 5 will be promoted to 5.0,

Now suppose the `Someclass` definition is

```
class Someclass
{
public:
...
 void redo(double);
 void redo(int) = delete;
...
};
```

In this case, the method call `sc.redo(5)` in the main program. The compiler will detect that fact and also detect that `redo(int)` is deleted. It will treat the call as a compile-time error. This illustrates that `delete` and `override` do exist as far as function look-up is concerned.

## Delegating Constructors

If you provide a class with several constructors, you may find yourself repeating code over and over. That is, some of the constructors may be identical except for the

the data members and to also do whatever the object finishes up doing whatever its own body requires.

## Inheriting Constructors

In another move to simplify coding, C++11 allows a class to inherit constructors from the base class. C++11 also allows a namespace to have a namespace available:

```
namespace Box
{
 int fn(int) { ... }
 int fn(double) { ... }
 int fn(const char *) { ... }
}
using Box::fn;
```

This makes all the overloaded `fn` functions available, making nonspecial member functions of a base class available. For example, consider the following code:

```
class C1
{
 ...
public:
 ...
}
```

```

class BS
{
 int q;
 double w;
public:
 BS() : q(0), w(0) {}
 BS(int k) : q(k), w(100) {}
 BS(double x) : q(-1), w(x) {}
 B0(int k, double x) : q(k), w(x) {}
 void Show() const {std::cout << q << "
};

```

```

class DR : public BS
{
 short j;
public:
 using BS::BS;
 DR() : j(-100) {} // DR needs its
 DR(double x) : BS(2*x), j(int(x)) {}
 DR(int i) : j(-2), BS(i, 0.5* i) {}
 void Show() const {std::cout << j << "
};

int main()

```

signature, you hide rather than override the c

```
class Action
{
 int a;
public:
 Action(int i = 0) : a(i) {}
 int val() const {return a;};
 virtual void f(char ch) const { std::cout << ch << endl; };
};

class Bingo : public Action
{
public:
 Bingo(int i = 0) : Action(i) {}
 virtual void f(char * ch) const { std::cout << *ch << endl; };
};
```

Because class Bingo uses `f(char * ch)` in its definition, it is not possible to call `f` on a Bingo object. This prevents a program from

```
Bingo b(10);
b.f('@'); // works for Action object, fails for Bingo object
```

With C++11, you can use the virtual specifier to override a virtual function. Place it after the function definition to match a base method, the compiler objects. This would generate a compile-time error message

grammer. You will have your suspicions seem-  
functions actually look—here’s an example:

```
[&count](int x){count += (x % 13 == 0);}
```

But they aren’t as arcane as they may look, especially with STL algorithms using function pointers.

## The How of Function Pointers, Functors, and Lambdas

Let’s look at an example using three approaches: function pointers, functors, and lambdas. We’ll treat these three forms as *function objects* so that we won’t have to say “functor or lambda.”) Suppose you wish to generate a list of 1000 random numbers and determine how many of them are divisible by 3 and how many are not. You might imagine that this is a quest you find absolutely

Generating the list is pretty straightforward. We’ll use a vector to hold the numbers and use the STL `generate` algorithm to generate random numbers:

```
#include <vector>
#include <algorithm>
#include <cmath>
...
std::vector<int> numbers(1000);
std::generate(numbers.begin(), numbers.end(),
```

recall from Chapter 16, is a class object than thanks to the class defining operator() () as a member function. One advantage of this operator in our example is that you can use the same operator for any integer value. Here is one possible definition:

```
class f_mod
{
private:
 int dv;
public:
 f_mod(int d = 1) : dv(d) {}
 bool operator()(int x) {return x % dv == 0;}
};
```

Recall how this works. You can use the constructor to set the divisor to a particular integer value:

```
f_mod obj(3); // f_mod.dv set to 3
```

This object can use the operator() method to test if a value is divisible by 3. For example:

```
bool is_div_by_3 = obj(7); // same as obj(7) % 3 == 0
```



lambda doesn't have a return statement, the type you would use this lambda as follows:

```
count3 = std::count_if(numbers.begin(), numbers.end(),
 [](int x){return x % 3 == 0;});
```

That is, you use the entire lambda expression as the return statement of the lambda constructor.

The automatic type deduction for lambdas works for the return statement. Otherwise, you need to use the return type explicitly:

```
[](double x)->double{int y = x; return x - y;}
```

Listing 18.4 illustrates the points just discussed.

#### Listing 18.4    **lambda0.cpp**

---

```
// lambda0.cpp -- using lambda expressions
#include <iostream>
#include <vector>
#include <algorithm>
#include <cmath>
#include <ctime>

const long Size1 = 39L;
const long Size2 = 100*Size1;
const long Size3 = 100*Size2;
```

```

// using a functor
class f_mod
{
private:
 int dv;
public:
 f_mod(int d = 1) : dv(d) {}
 bool operator()(int x) {return x
};

count3 = std::count_if(numbers.begin()
cout << "Count of numbers divisible b
count13 = std::count_if(numbers.begin()
cout << "Count of numbers divisible b

// increase number of numbers again
numbers.resize(Size3);
std::generate(numbers.begin(), number
cout << "Sample size = " << Size3 <<

// using lambdas
count3 = std::count_if(numbers.begin()
 [](int x){return x % 3 == 0;
cout << "Count of numbers divisible b

```

## The Why of Lambdas

You may be wondering what need, other than convenience, lambda serves. Let's examine this question in terms of efficiency, and capability.

Many programmers feel that it is useful to have a lambda used. That way, you don't have to scan through the third argument to a `count_if()` function call. If the code, all the components are close at hand. If the code is elsewhere, again all the components are at hand. If the code is elsewhere, again all the components are at hand. Because the definition is at the point of usage, a lambda cannot be defined inside other functions, so the code is far from the point of usage. Functors can be defined inside a functor class, can be defined inside a function, and can be defined at the point of use.

In terms of brevity, the functor code is more concise than lambda code. Functions and lambdas are appropriate for different exceptions would be if you had to use a lambda

```
count1 = std::count_if(n1.begin(), n1.end(),
 [](int x){return x % 3 == 0;});
count2 = std::count_if(n2.begin(), n2.end(),
 [](int x){return x % 3 == 0;});
```

ence. Using `[&]` provides access to all the automatic variables by reference. For example, in the following code, `count13` is an automatic variable. Using `[&]` provides access to `count13` by reference. Using `[=]` provides access to `count13` by value. Using `[=, &]` provides access to `count13` by reference, and `[=, &]` would provide access by value. Using `[=, &]` provides access by reference to all the automatic variables. In Listing 18.4, you can see how to use `[&]` to access automatic variables. In Listing 18.4, you can see how to use `[&]` to access automatic variables.

```
int count13;
...
count13 = std::count_if(numbers.begin(), numbers.end(),
 [](int x){return x % 13 == 0;})
```

with this:

```
int count13 = 0;
std::for_each(numbers.begin(), numbers.end(),
 [&count13](int x){count13 += x % 13 == 0;})
```

The `[&count13]` allows the lambda to use `count13` by reference. If `count13` is captured by reference, any changes to `count13` are reflected in the lambda. The expression `x % 13 == 0` evaluates to 1 when `x` is divisible by 13, and 0 otherwise. The lambda adds the result of `x % 13 == 0` to `count13`. The expression `x % 13 == 0` converts to 1 when added to `count13`. Similarly, the lambda adds the result of `x % 13 == 0` to `count13`. The `for_each()` applies the lambda expression to each element in the range `[numbers.begin(), numbers.end())`. The lambda expression adds the result of `x % 13 == 0` to `count13`. The number of elements divisible by 13 is stored in `count13`.

```

 using std::cout;
 std::vector<int> numbers(Size);

 std::srand(std::time(0));
 std::generate(numbers.begin(), numbers.end(), rand);
 cout << "Sample size = " << Size << " \n";
// using lambdas
 int count3 = std::count_if(numbers.begin(), numbers.end(),
 [](int x){return x % 3 == 0;});
 cout << "Count of numbers divisible by 3 = " << count3 << " \n";
 int count13 = 0;
 std::for_each(numbers.begin(), numbers.end(),
 [&count13](int x){count13 += x % 13 == 0;});
 cout << "Count of numbers divisible by 13 = " << count13 << " \n";
// using a single lambda
 count3 = count13 = 0;
 std::for_each(numbers.begin(), numbers.end(),
 [&](int x){count3 += x % 3 == 0; count13 += x % 13 == 0;});
 cout << "Count of numbers divisible by 3 = " << count3 << " \n";
 cout << "Count of numbers divisible by 13 = " << count13 << " \n";
 return 0;
 }

```

---

plied as an argument. C++11 provides additional templates, which provides a more flexible alternative. The `reference_wrapper` template, which allows a member function to be passed as an argument, and the `reference_wrapper` template allows you to pass a reference to a function, which can be copied, and the function wrapper template allows you to pass a function-like forms uniformly.

Let's look more closely at one example of the problem it addresses.

## The function Wrapper and Template

Consider the following line of code:

```
answer = ef(q);
```

What is `ef`? It could be the name of a function, it could be a function object. It could be a name of a variable, or it could be a name of a class. These are examples of *callable types*. The abundance of these types is a source of inefficiencies. To see this, let's examine a simple case.

First, let's define some templates in a header

```

class Fq
{
private:
 double z_;
public:
 Fq(double z = 1.0) : z_(z) {}
 double operator()(double q) { return z; }
};

```

---

The `use_f()` template uses the parameter `f` to return `f(v)`;

Next the program in Listing 18.7 calls the

#### Listing 18.7    **callable.cpp**

---

```

// callable.cpp -- callable types and template
#include "somedefs.h"
#include <iostream>

double dub(double x) {return 2.0*x;}
double square(double x) {return x*x;}

```

parameter F? Each time the actual argument is passed to the function object, the increment and returns a type double value, so it runs all six calls to `use_f()` and that the template following sample output shows, that belief is

Function pointer dub:

```
use_f count = 1, &count = 0x402028
2.42
```

Function pointer square:

```
use_f count = 2, &count = 0x402028
1.1
```

Function object Fp:

```
use_f count = 1, &count = 0x402020
6.05
```

Function object Fq:

```
use_f count = 1, &count = 0x402024
6.21
```

Lambda expression 1:

```
use_f count = 1, &count = 0x405020
1.4641
```

Lambda expression 2:

```
use_f count = 1, &count = 0x40501c
1.815
```



The function wrapper lets you rewrite the previous `use_f()` instead of five. Note that the function expressions in Listing 18.7 share a common body and each returns a type `double` value. We call this *signature*, which is described by the return type and parameter types enclosed in a pair of parentheses. We use `double(double)` as the call signature.

The function template, declared in the previous section, is in terms of a call signature, and it can be used to create a lambda expression having the same call signature. This section creates a function object `fdci` that takes the call signature `double(double)` and returns a type `double`:

```
std::function<double(char, int)> fdci;
```

You can then assign to `fdci` any function expression that takes type `char` and `int` arguments and returns a type `double`.

The various callable arguments in Listing 18.7 have the call signature `double(double)`. So to fix Listing 18.7 and reuse `function<double(double)>` to create six function objects and lambdas. Then all six calls to `use_f()` can be replaced with `(function<double(double)>)` for `F`, resulting in the result.

```
 cout << "Function pointer dub:\n";
 cout << " " << use_f(y, ef1) << endl;
 cout << "Function pointer square:\n";
 cout << " " << use_f(y, ef2) << endl;
 cout << "Function object Fp:\n";
 cout << " " << use_f(y, ef3) << endl;
 cout << "Function object Fq:\n";
 cout << " " << use_f(y, ef4) << endl;
 cout << "Lambda expression 1:\n";
 cout << " " << use_f(y, ef5) << endl;
 cout << "Lambda expression 2:\n";
 cout << " " << use_f(y, ef6) << endl;
 return 0;
}
```

---

Here is a sample output:

Function pointer dub:

```
use_f count = 1, &count = 0x404020
2.42
```

Function pointer sqrt:

```
use_f count = 2, &count = 0x404020
1.1
```

function:

```
typedef function<double(double)> fdd; // s
cout << use_f(y, fdd(dub)) << endl; // d
cout << use_f(y, fdd(square)) << endl;
...
```

Second, Listing 18.8 adapts the second argument to the function `f`. Another approach is to adapt the function to take the original arguments. This can be done by using the second parameter for the `use_f()` template definition.

```
#include <functional>
template <typename T>
T use_f(T v, std::function<T(T)> f) // f
{
 static int count = 0;
 count++;
 std::cout << " use_f count = " << count << endl;
 << " ", &count = " << &count << endl;
 return f(v);
}
```

Then the function calls can look like this:

```
cout << " " << use_f<double>(y, dub) << endl;
...
```

The goal is to be able to define `show_list` in the header file and lead to this output:

```
14, 2.71828
7.38905, !, 7, Mr. String objects!
```

There are a few key points to understand

- Template parameter packs
- Function parameter packs
- Unpacking a pack
- Recursion

## Template and Function Parameter Packs

As a starting point to see how parameter packs work in a function, one that displays a list consisting of

```
template<typename T>
void show_list0(T value)
{
 std::cout << value << ", ";
}
```

Next, much as

```
void show_list0(T value)
```

states that `value` is of type `T`, the line

```
void show_list1(Args... args) // args is
```

states that `args` is of type `Args`. More precisely, `args` contains a list of values that matches the list of types and in number. In this case, `args` contains

In this manner, the `show_list1()` variadic function calls:

```
show_list1();
```

```
show_list1(99);
```

```
show_list1(88.5, "cat");
```

```
show_list1(2,4,6,8, "who do we", std::string{"dog"});
```

In the last case, the `Args` template parameter is `int, int, const char *, and std::string`, and `args` contains the matching values `2, 4, 6, 8, "who do we", and std::string{"dog"}`.

## Unpacking the Packs

But how can the function access the contents of the argument list? That is, you can't use something like `Arg`

## Using Recursion in Variadic Templates

Although recursion dooms `show_list1()` as used recursion provides a solution to accessing the function parameter pack, process the first on to a recursive call, and so on, until the list is empty. It is important to make sure that there is a call that terminates the recursion. This involves changing the template heading to the following:

```
template<typename T, typename... Args>
void show_list3(T value, Args... args)
```

With this definition, the first argument to `show_list3` is assigned to `value`. The remaining arguments are passed to `args...`, which allows the function to do something with `value` and then make a recursive call with `args...`. The recursive call then prints a value and passes on the remaining arguments.

Listing 18.9 presents an implementation of this technique.

### Listing 18.9    **variadic1.cpp**

```
//variadic1.cpp -- using recursion to unpack a variadic template
#include <iostream>
#include <string>
```

Consider this function call:

```
show_list3(x*x, '!', 7, mr);
```

The first argument matches `T` to double and `char`, `int`, and `std::string` are placed in the `args` pack. `('!', 7, and mr)` are placed in the `args` pack.

Next, the `show_list3()` function uses `cout` and the string `" , "`. That takes care of displaying

Next comes this call:

```
show_list3(args...);
```

This, given the expansion of `args...`, is the

```
show_list3('!', 7, mr);
```

As promised, the list is shortened by one item and `!`, and the remaining two types and values. The next recursive call processes these recursively. The next recursive call processes these recursively. The version of `show_list3()` with no arguments

Here is the output for the two function calls

```
14, 2.71828, 7.38905, !, 7, Mr. String object
```

you can use this:

```
show_list3(const Args&... args);
```

That will cause each function parameter to be of type `std::string` instead of `std::string mr`, the final paring `std::string& mr`.

Listing 18.10 incorporates these two changes.

#### Listing 18.10    **variadic2.cpp**

---

```
// variadic2.cpp
#include <iostream>
#include <string>

// definition for 0 parameters
void show_list() {}

// definition for 1 parameter
template<typename T>
void show_list(const T& value)
{
 std::cout << value << '\n';
}
```



# More C++11 Features

C++11 adds many more features than this book covers, but many of them are not widely implemented at the time of writing, worth taking a quick look at the nature of some of them.

## Concurrent Programming

These days it's easier to improve computer performance by increasing processor speed. So computers with multiple cores and with multiple multicore processors are the norm. You can run threads of execution simultaneously. One processor handles your word document, another processor handles your spreadsheet.

Some activities can benefit from multiple threads. For example, searching for something in a singly linked list. You have to follow the links, in order, to the end of the list. If you have multiple processors, you could do to help. Now consider an unsorted array of data. If you have multiple arrays, you could start one thread from the beginning and another from the middle, thus halving the search time.

Multiple threads do raise many problems. What if you have two threads try to access the same data simultaneously?

tor can be represented by the widest integer for these numbers.

One of the most interesting additions is a regex header file. A regular expression specifies contents in a text string. For example, a bracket matches the brackets. Thus, `[cCkK]` matches a single `c`, `C`, `k`, or `K`. Other patterns include `cat`, `Cat`, `kat`, and `Kat`. Other patterns include `and` and many, many others. The fact that a backslash character in an escape sequence requires a pair of backslashes (to be written as the string literal `"\\d\\t\\w\\t"`) is one reason the raw string was introduced (see `same pattern as R"d\t\\w\\t"`).

Unix utilities such as `ed`, `grep`, and `awk` use regular expressions. The language Perl extended their capabilities. There are many to choose from several flavors of regular expressions.

## Low-Level Programming

The “low level” in low-level programming refers to the quality of the programming. Low level means close to hardware and machine language. Low-level programming is for increasing the efficiency of programs to those who do low-level programming.

to allow const variables to be stored in read-only memory, which is useful in embedded programming. (Variables, constants, and arrays at runtime, are stored in random-access memory.)

## Miscellaneous

C++11 follows the lead of C99 in allowing for extended integer types. Such types, for example, could be used to represent Extended types are supported in the C header `<stdint.h>`.

C++11 provides a mechanism, the *literal operator*, for defining integer literals. Using this mechanism, for instance, one can define a literal operator for which the corresponding literal operator will be used.

C++ has a debugging tool called `assert`. The `assert` macro checks if an assertion is true and which displays a message if it is not. The assertion would typically be about something that should be true at that point in the program. C++11 adds the `assert` macro used to test assertions during compile time. The `assert` macro is used in debug templates for which instantiation takes place at compile time.

C++11 provides more support for metaprogramming, which allows one to create or modify other programs or even their compilation at compile time using templates.

additional compiler support. And if they are not, they are distributed as text files in the form of header files.

One example of this sort of change is the Boost library, created by Stepanov, and made freely available by Hewlett-Packard. The programming community made it a candidate for the standard, which influenced other aspects of the emerging standard.

## **The Boost Project**

More recently, the Boost library has become a standard. It has had a significant influence on C++11. The Boost Project, led by Dawes, the then-chairman of the C++ library working group, and other members of the group and developed a standard within the confines of the standards committee. The basic idea was to create an open forum for people to post free C++ code, discuss licensing and programming practices, and it now has a large number of members. The result is a group of highly praised and highly used libraries in an environment in which the programming community can share ideas and provide feedback.

Boost has over 100 libraries at the time of writing. You can find a set from [www.boost.org](http://www.boost.org), as can documentation, including the appropriate header files.

types. The syntax is modeled after `dynamic_cast` and `static_cast` and takes a template parameter. Listing 18.11 shows a simple example.

#### Listing 18.11 `lexcast.cpp`

---

```
// lexcast.cpp -- simple cast from float to string
#include <iostream>
#include <string>
#include "boost/lexical_cast.hpp"
int main()
{
 using namespace std;
 cout << "Enter your weight: ";
 float weight;
 cin >> weight;
 string gain = "A 10% increase raises ";
 string wt = boost::lexical_cast<string>(weight);
 gain = gain + wt + " to "; // string
 weight = 1.1 * weight;
 gain = gain + boost::lexical_cast<string>(weight);
 cout << gain << endl;
 return 0;
}
```

---

# What Now?

If you have worked your way through this book, you know the basic rules of C++. However, that's just the beginning. The next stage is learning to use the language effectively in a real situation. The best situation is to be in a work or learning environment where you can see good C++ code and programmers. Also now you can focus on topics that concentrate on more advanced topics and techniques. See Appendix H, "Selected Readings and Internet Resources."

One promise of OOP is to facilitate the development of large projects. One of the essential activities of OOP is to model that represent the situation (called the *problem domain*). Problems are often complex, finding a suitable model for a complex system from scratch usually doesn't work. An evolutionary approach. Toward this end, practical techniques and strategies. In particular, it's important to see the evolution in the analysis and design stages as well as the actual code.

Two common techniques are *use-case analysis* and *use-case modeling*. A development team lists the common ways, or use cases, that the system to be used, identifying elements, actions,



```

int k(99);
Z200 zip(200,'Z',0.675);
std::vector<int> ai(5);
int ar[5] = {3, 9, 4, 7, 1};
for (auto pt = ai.begin(), int i =
 *pt = ai[i];

```

2. For the following short program, which of the following are valid calls, what does the reference argu

```

#include <iostream>
using namespace std;

double up(double x) { return 2.0* x; }
void r1(const double &rx) {cout << rx << endl;}
void r2(double &rx) {cout << rx << endl;}
void r3(double &&rx) {cout << rx << endl;}

int main()
{
 double w = 10.0;
 r1(w);
 r1(w+1);
 r1(up(w));
}

```



```

 r1(up(w));
 return 0;
}

```

- b. What does the following short program do?

```

#include <iostream>
using namespace std;

double up(double x) { return 2*x; }
void r1(double &rx) {cout << "r1: "; rx*=2; cout << " ";}
void r1(double &&rx) {cout << "r2: "; rx*=2; cout << " ";}

int main()
{
 double w = 10.0;
 r1(w);
 r1(w+1);
 r1(up(w));
 return 0;
}

```

- c. What does the following short program do?

```

#include <iostream>
using namespace std;

```

Why would this class not be a good candidate for a move function?  
What change in approach to storing the data would make it a good candidate for a move function?

6. Revise the following short program so that it uses a move function. Don't change `show2()`.

```
#include <iostream>
template<typename T>
 void show2(double x, T& fp) {std::cout << x << " " << fp << "\n";}
double f1(double x) { return 1.8*x;}
int main()
{
 show2(18.0, f1);
 return 0;
}
```

7. Revise the following short and ugly program so that it uses a move function instead of the `Adder` functor. Don't change the `main()` function.

```
#include <iostream>
#include <array>
const int Size = 5;
template<typename T>
```

```
 {
 fp(*pt);
 }
}
```

## Programming Exercises

1. Here is part of a short program:

```
int main()
{
 using namespace std;
 // list of double deduced from list
 auto q = average_list({15.4, 10.4, 12.5});
 cout << q << endl;
 // list of int deduced from list containing ints
 cout << average_list({20, 30, 19});
 // forced list of double
 auto ad = average_list<double>({15.4, 10.4, 12.5});
 cout << ad << endl;
 return 0;
}
```

```

 Cpmv & operator=(const Cpmv & mv);
 Cpmv operator+(const Cpmv & obj);
 void Display() const;
};

```

The `operator+()` function should create a new `Cpmv` object and the members concatenate the corresponding members. The `operator+()` implements move semantics for the `Cpmv` object. Write a program that uses all the methods of `Cpmv` and prints various methods verbose so that you can see the results.

3. Write and test a variadic template function that takes an arbitrarily long list of arguments with numeric types and returns the sum as a long double value.
4. Redo Listing 16.5 using lambdas. In particular, use a named lambda and replace the two `lambda` expressions.



(binary).

## Decimal Numbers (Base 10)

The method we use for writing numbers is based on powers of 10. For example, the number 2,468. The 2 represents 2 thousands, the 4 represents 4 hundreds, the 6 represents 6 tens, and the 8 represents 8 ones:

$$2,468 = 2 \times 1,000 + 4 \times 100 + 6 \times 10 + 8 \times 1$$

One thousand is  $10 \times 10 \times 10$ , which can be written as  $10^3$ . Using this notation, you can write the preceding equation as follows:

$$2,468 = 2 \times 10^3 + 4 \times 10^2 + 6 \times 10^1 + 8 \times 10^0$$

Because this number notation is based on powers of 10, it is called decimal notation. You can also use another number system called octal, which uses base 8 (octal) and base 16 (hexadecimal) notation. (Note that  $10^0$  is 1, as is any nonzero number raised to the power of 0.)

## Octal Integers (Base 8)

Octal numbers are based on powers of 8, so the octal number system uses powers of 8. C++ uses a 0 prefix to indicate octal numbers. You can use powers of 8 to find the equivalent base 10 value of an octal number.

| Hexadecimal Digit | Decimal Value |
|-------------------|---------------|
| a or A            | 10            |
| b or B            | 11            |
| c or C            | 12            |
| d or D            | 13            |
| e or E            | 14            |
| f or F            | 15            |

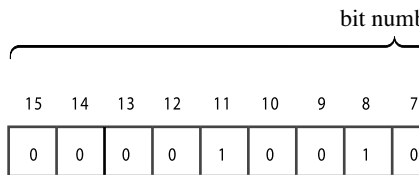
---

C++ uses 0x or 0X notation to indicate hexadecimal values. To find the decimal equivalent of a hexadecimal value, use the following formula:

| Hexadecimal | Decimal                                                                                           |
|-------------|---------------------------------------------------------------------------------------------------|
| 0x2B3       | $= 2 \cdot 16^2 + 11 \cdot 16^1 + 3 \cdot 16^0$ $= 2 \cdot 256 + 11 \cdot 16 + 3 \cdot 1$ $= 691$ |

Hardware documentation often uses hexadecimal values to represent memory locations and port numbers.

represents a 2-byte integer.



$$\begin{aligned}\text{value} &= 1 \times 2^{11} + \\ &= 2048 + 256 \\ &= 2338\end{aligned}$$

Figure A.1 A 2-

## Binary and Hex

Hex notation is often used to provide a more compact representation of memory addresses or integers holding bit-flags. Each hexadecimal digit corresponds to a 4-bit unit.



To convert a hex value to binary, you just find the binary equivalent. For example, the hex value 0100. Similarly, you can easily convert binary to hex. A 4-bit unit into the equivalent hex digit. For example, 0100 becomes 0x95.

## Big Endian and Little Endian

Oddly, two computing platforms that both use the same architecture represent the same number identically. Intel processors use the Little Endian architecture, whereas Motorola processors, and ARM processors employ the Big Endian architecture. Both systems can be configured to use either scheme.

The terms *Big Endian* and *Little Endian* can be confusing. “Little End In”—a reference to the order of bytes. On an Intel computer (Little Endian), the low-order byte of a value such as 0xABCD would be stored in memory first. A Big Endian machine would store the same value in reverse order: 0xAB 0xCD.

Jonathan Swift’s book *Gulliver’s Travels* is the first to satirize the irrationality of many political disputes by i





keywords. Keywords shown in boldface are additions to C++11. Keywords shown in italics are C++11 additions.

Table B.1 C++ Keywords

|                 |                         |                      |
|-----------------|-------------------------|----------------------|
| <i>alignas</i>  | <i>alignof</i>          | <i>asm</i>           |
| <b>break</b>    | <b>case</b>             | <i>catch</i>         |
| <i>char32_t</i> | <i>class</i>            | <b>constexpr</b>     |
| <b>continue</b> | <i>decltype</i>         | <b>default</b>       |
| <b>double</b>   | <i>dynamic_cast</i>     | <b>else</b>          |
| <i>export</i>   | <b>extern</b>           | <i>false</i>         |
| <i>friend</i>   | <b>goto</b>             | <b>if</b>            |
| <b>long</b>     | <i>mutable</i>          | <i>namespace</i>     |
| <i>nullptr</i>  | <i>operator</i>         | <i>priority</i>      |
| <b>register</b> | <i>reinterpret_cast</i> | <b>return</b>        |
| <b>sizeof</b>   | <b>static</b>           | <i>static_assert</i> |
| <b>switch</b>   | <i>template</i>         | <i>thread</i>        |
| <i>true</i>     | <i>try</i>              | <b>typedef</b>       |
| <b>union</b>    | <b>unsigned</b>         | <i>using</i>         |
| <b>volatile</b> | <i>wchar_t</i>          | <b>while</b>         |

## C++ Library Reserved Names

The compiler won't let you use keywords and a class of forbidden names for which the protection is provided. If you use a name which is a name reserved for use by the C++ library, the effect is undefined. That is, it might generate a warning, it might cause a program to run incorrectly, or it might cause a program to crash.

The C++ language reserves macro names beginning with `_`. If a header file includes a particular header file, then you should not use that header (or in headers included by that header) a name that is reserved. For example, if you include the header file `<climits>`, you should not use `CHAR_BIT` as an identifier because that name is reserved.

The C++ language reserves names beginning with `__` (double underscore) followed by an uppercase letter for any use except as a single underscore for use as a global variable. Some implementations reserve `__Lynx` in any case and names such as `__lynx`.

keyword and provided more than one use for. C++11 has implemented another way to avoid identifiers with special meanings. These identifiers they are used to implement language features. I will examine whether they are used as ordinary identifiers.

```
class F
{
 int final; //
public:
 ...
 virtual void unfold() {...} = final;
};
```

Here the `final` on line #1 is used as an ordinary identifier. It is used to invoke a language feature. The two uses are not conflicting.

Also C++ has many identifiers that are commonly reserved. These include header file names, library names, the required function with which execution begins, and so on. If you conflict, you can use these identifiers for other purposes. That is, nothing except a lack of common sense.

```
// allowable but silly
#include <iostream>
int iostream(int a);
```



prefix in the table, the ^ character denotes us

Table C.1 **ASCII Character Set Representati**

| <b>Decimal</b> | <b>Octal</b> | <b>Hex</b> | <b>Binary</b> |
|----------------|--------------|------------|---------------|
| 0              | 0            | 0          | 00000         |
| 1              | 01           | 0x1        | 00000         |
| 2              | 02           | 0x2        | 00000         |
| 3              | 03           | 0x3        | 00000         |
| 4              | 04           | 0x4        | 00000         |
| 5              | 05           | 0x5        | 00000         |
| 6              | 06           | 0x6        | 00000         |
| 7              | 07           | 0x7        | 00000         |
| 8              | 010          | 0x8        | 00001         |
| 9              | 011          | 0x9        | 00001         |
| 10             | 012          | 0xa        | 00001         |
| 11             | 013          | 0xb        | 00001         |
| 12             | 014          | 0xc        | 00001         |
| 13             | 015          | 0xd        | 00001         |



|    |     |      |        |
|----|-----|------|--------|
| 30 | 036 | 0x1e | 000111 |
| 31 | 037 | 0x1f | 000111 |
| 32 | 040 | 0x20 | 001000 |
| 33 | 041 | 0x21 | 001000 |
| 34 | 042 | 0x22 | 001000 |
| 35 | 043 | 0x23 | 001000 |
| 36 | 044 | 0x24 | 001000 |
| 37 | 045 | 0x25 | 001000 |
| 38 | 046 | 0x26 | 001000 |
| 39 | 047 | 0x27 | 001000 |
| 40 | 050 | 0x28 | 001010 |
| 41 | 051 | 0x29 | 001010 |
| 42 | 052 | 0x2a | 001010 |
| 43 | 053 | 0x2b | 001010 |
| 44 | 054 | 0x2c | 001011 |
| 45 | 055 | 0x2d | 001011 |
| 46 | 056 | 0x2e | 001011 |

|    |      |      |        |
|----|------|------|--------|
| 63 | 077  | 0x3f | 001111 |
| 64 | 0100 | 0x40 | 010000 |
| 65 | 0101 | 0x41 | 010000 |
| 66 | 0102 | 0x42 | 010000 |
| 67 | 0103 | 0x43 | 010000 |
| 68 | 0104 | 0x44 | 010000 |
| 69 | 0105 | 0x45 | 010000 |
| 70 | 0106 | 0x46 | 010000 |
| 71 | 0107 | 0x47 | 010000 |
| 72 | 0110 | 0x48 | 010011 |
| 73 | 0111 | 0x49 | 010011 |
| 74 | 0112 | 0x4a | 010011 |
| 75 | 0113 | 0x4b | 010011 |
| 76 | 0114 | 0x4c | 010011 |
| 77 | 0115 | 0x4d | 010011 |
| 78 | 0116 | 0x4e | 010011 |
| 79 | 0117 | 0x4f | 010011 |

|     |      |      |        |
|-----|------|------|--------|
| 96  | 0140 | 0x60 | 011000 |
| 97  | 0141 | 0x61 | 011000 |
| 98  | 0142 | 0x62 | 011000 |
| 99  | 0143 | 0x63 | 011000 |
| 100 | 0144 | 0x64 | 011000 |
| 101 | 0145 | 0x65 | 011000 |
| 102 | 0146 | 0x66 | 011000 |
| 103 | 0147 | 0x67 | 011000 |
| 104 | 0150 | 0x68 | 011010 |
| 105 | 0151 | 0x69 | 011010 |
| 106 | 0152 | 0x6a | 011010 |
| 107 | 0153 | 0x6b | 011010 |
| 108 | 0154 | 0x6c | 011010 |
| 109 | 0155 | 0x6d | 011010 |
| 110 | 0156 | 0x6e | 011010 |
| 111 | 0157 | 0x6f | 011010 |
| 112 | 0160 | 0x70 | 011100 |





the table). Left-to-right associativity means to apply the  
right-to-left associativity means to apply the

Table D.1 C++ Operator Precedence and As

| Operator                  | Assoc. | M  |
|---------------------------|--------|----|
| <i>Precedence Group 1</i> |        |    |
| ::                        |        | S  |
| <i>Precedence Group 2</i> |        |    |
| (expression)              |        | G  |
| ()                        | L-R    | F  |
| ()                        |        | V  |
| []                        |        | A  |
| .                         |        | D  |
| ->                        |        | In |
| ++                        |        | In |
| --                        |        | D  |
| const_cast                |        | S  |
| dynamic_cast              |        | S  |
| reinterpret_cast          |        | S  |
| static_cast               |        | S  |
| typeid                    |        | Ty |

noexcept

Fa

*Precedence Group 4*

. \* L-R

Me

->\*

In

*Precedence Group 5 (All Binary)*

\* L-R

Me

/

Di

%

Me

*Precedence Group 6 (All Binary)*

+ L-R

Ac

-

Su

*Precedence Group 7*

<< L-R

Le

>>

Ri

*Precedence Group 8*

< L-R

Le

<=

Le

>=

Gr

>

Gr

|                            |     |    |
|----------------------------|-----|----|
| =                          | R-L | S  |
| *=                         |     | M  |
| /=                         |     | D  |
| %=                         |     | Ta |
| +=                         |     | A  |
| -=                         |     | S  |
| &=                         |     | B  |
| ^=                         |     | B  |
| =                          |     | B  |
| <<=                        |     | L  |
| >>=                        |     | R  |
| <i>Precedence Group 17</i> |     |    |
| throw                      | L-R | T  |
| <i>Precedence Group 18</i> |     |    |
| ,                          | L-R | C  |

---



```
char ch = *str++;
```

The postfix `++` operator has higher precedence than the increment operator. So, `++str` increments the pointer, making it point to the next character pointed to. However, because `++` is the postfix operator, the original value of `*str` is assigned to `ch`. Thereafter, `str` is incremented to point to the next character. Here's a similar example:

```
char * str = "Whoa";
char ch = *++str;
```

The prefix `++` operator and the unary `*` operator associate right-to-left. So, again, `str` and not `*str` is incremented. Since the `++` operator is in prefix form, first `str` is incremented, and then `*str` is dereferenced. Thus, `str` moves to point to the next character.

Note that Table D.1 uses *binary* or *unary* to distinguish between two operators that use the same symbol. For example, the binary bitwise AND operator and the unary address-of operator both use the `&` symbol.

Appendix B, “C++ Reserved words,” lists the reserved keywords and operators.

## Bitwise Operators

The bitwise operators operate on the bits of an integer. The left-shift operator moves bits to the left, and the right-shift operator moves bits to the right. Each shift moves each 0 to a 1. Altogether, C++ has six such operators.

### The Shift Operators

The left-shift operator has the following syntax:

```
value << shift
```

Here *value* is the integer value to be shifted. The *shift* is the number of positions to shift. For example, the following shifts all the bits in the integer 13 to the left by 3 positions:

```
13 << 3
```

The vacated places are filled with zeros, as shown in Figure E.1).

Because each bit position represents a value of 2<sup>n</sup>, where *n* is the bit position (see Appendix A, “Number Bases”), shifting one bit position to the left is equivalent to multiplying the value by 2. Similarly, shifting two bit positions to the left is equivalent to multiplying the value by 4. Shifting *n* positions is equivalent to multiplying the value by 2<sup>*n*</sup>, or 104.

produce a new value, much as  $x + 3$  produces

If you want to use the left-shift operator to also use assignment. You can use regular assignment with shifting with assignment:

```
x = x << 4; // regular assignment
y <<= 2; // shift and assign
```

The right-shift operator ( $>>$ ), as you might expect, has the following syntax:

```
value >> shift
```

Here *value* is the integer value to be shifted. For example, the following shifts all the bits in

```
17 >> 2
```

For unsigned integers, the vacated places at the end are discarded. For signed integers, vacated places are filled with the value of the original leftmost bit. The choice of which is used is called the shift mode (Figure E.2 shows an example that illustrates this).

The right-shift operator. Shifting one place to the right is equivalent to division by 2. In general, shifting *n* places to the right is equivalent to

## The Logical Bitwise Operators

The logical bitwise operators are analogous to the arithmetic operators; they apply to a value on a bit-by-bit basis rather than to the entire value. The regular negation operator (!) and the bitwise NOT operator (~) convert a true (or nonzero) value to false (or zero). The bitwise NOT operator converts each individual bit to its opposite. For example, consider the unsigned char value of 3:

```
unsigned char x = 3;
```

The expression !x has the value 0. To see this, consider the binary value 00000011. Then you convert each 0 to 1 and each 1 to 0, resulting in 11111100, which in base 10 is the value 252. This new value is termed the *complement* of the original value.

The bitwise OR operator (|) combines two values. Each bit in the new value is set to 1 if either of the corresponding bits in the original values is set to 1. If both bits are set to 0 (see Figure E.4).

Table E.1 summarizes how the `|` operator

Table E.1 The Value of `b1 | b2`

| Bit Values          | <code>b1 = 0</code> | <code>b1</code> |
|---------------------|---------------------|-----------------|
| <code>b2 = 0</code> | 0                   | 1               |
| <code>b2 = 1</code> | 1                   | 1               |

The `|=` operator combines the bitwise OR operator with the assignment operator. For example, the code

```
a |= b; // set a to a | b
```

The bitwise XOR operator (`^`) combines two values. Each bit in the new value is set to 1 if one of the corresponding bits in the original values is set to 1. If both are 1, the final bit is set to 0 (see Figure E.5).

Table E.2 summarizes how the `^` operator

The  $\wedge$  operator combines the bitwise XOR

```
a ^= b; // set a to a ^ b
```

The bitwise AND operator ( $\&$ ) combines value. Each bit in the new value is set to 1 only if both original values are set to 1. If either or both are 0 (see Figure E.6).

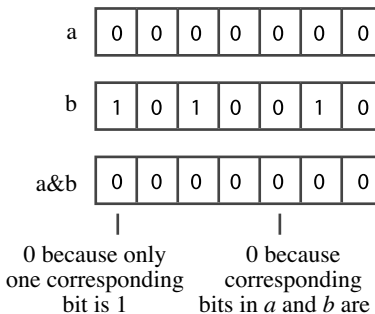


Figure E.6 The b

|                     |                     |
|---------------------|---------------------|
| <code>&amp;</code>  | <code>bitand</code> |
| <code>&amp;=</code> | <code>and_eq</code> |
| <code> </code>      | <code>bitor</code>  |
| <code> =</code>     | <code>or_eq</code>  |
| <code>~</code>      | <code>compl</code>  |
| <code>^</code>      | <code>xor</code>    |
| <code>^=</code>     | <code>xor_eq</code> |

---

These alternative forms let you write state

```
b = compl a bitand b; // same as b = ~a & b;
c = a xor b; // same as c = a ^ b;
```

## A Few Common Bitwise Operator

Often controlling hardware involves turning p status. The bitwise operators provide the mean the methods quickly.

In the following examples, `lottabits` repr the value corresponding to a particular bit. Bi ning with bit 0, so the value corresponding to

XORing 1 with 0 produces 1, turning an on bit on, turning an on bit off. All other bits in `lottabits` are unchanged. XORing 0 with 0 produces 0, and XORing

## Turning a Bit Off

The following operation turns off the bit in `lottabits` represented by `bit`:

```
lottabits = lottabits & ~bit;
```

These statements turn the bit off, regardless of its current value. XORing 1 with 1 produces an integer with all its bits set to 1 except for the bit that becomes 0. ANDing a 0 with any bit results in 0. All other bits in `lottabits` are unchanged. That's because `~bit` has a 0 in the position of `bit` and a 1 in all other positions. The value that bit had before.

Here's a briefer way of doing the same thing:

```
lottabits &= ~bit;
```

## Testing a Bit Value

Suppose you want to determine whether the bit in `lottabits` is on or off. The following test does not need

```
if (lottabits && bit) { // no good
```



```

 int feet;
 int inches;
public:
 Example();
 Example(int ft);
 ~Example();
 void show_in() const;
 void show_ft() const;
 void use_ptr() const;
};

```

Consider the `inches` member. Without a static class defines `inches` as a member identifier, but we have memory allocated:

```
Example ob; // now ob.inches exists
```

Thus, you specify an actual memory location in conjunction with a specific object. (In a member object, but then the object is understood to be a static member.)

C++ lets you define a member pointer to a static member:

```
int Example::*pt = &Example::inches;
```

This pointer is a bit different from a regular pointer to a specific memory location. But the `pt` pointer does not point to a specific memory location.

```
pt = &Example::feet; // reset pt
cout << ob1.*pt << endl; // display feet
```

In essence, the combination `*pt` takes the  
identify different member names (of the same

You can also use member pointers to identify  
relatively involved. Recall that declaring a pointer  
no arguments looks like this:

```
void (*pf)(); // pf points to a function
```

Declaring a pointer to a member function of a  
particular class. Here, for instance, is how to declare

```
void (Example::*pf)() const; // pf points to a
```

This indicates that `pf` can be used the same way  
Note that the term `Example::*pf` has to be qualified by  
particular member function to this pointer:

```
pf = &Example::show_inches;
```

Note that unlike in the case of ordinary functions,  
must use the address operator. Having made the pointer  
to invoke the member function:

```
Example ob3(20);
(ob3.*pf)(); // invoke show_inches()
```

```
 Example(int ft);
 ~Example();
 void show_in() const;
 void show_ft() const;
 void use_ptr() const;
};
```

```
Example::Example()
{
 feet = 0;
 inches = 0;
}
```

```
Example::Example(int ft)
{
 feet = ft;
 inches = 12 * feet;
}
```

```
Example::~~Example()
{
}
```

```
 pf = &Example::show_in;
 cout << "Set pf to &Example::show_in: ";
 cout << "Using (this->*pf)(): ";
 (this->*pf)();
 cout << "Using (yard.*pf)(): ";
 (yard.*pf)();
 }

int main()
{
 Example car(15);
 Example van(20);
 Example garage;

 cout << "car.use_ptr() output:\n";
 car.use_ptr();
 cout << "\nvan.use_ptr() output:\n";
 van.use_ptr();

 return 0;
}
```

---

This example assigned pointer values at compile time. If you can use member pointers to data members, the data associated with the pointer is determined at runtime.

## **alignof (C++11)**

Computer systems can have restrictions on how data is stored. For example, one system might require that a double value be stored at a memory location, whereas another might require the size of the value to be a multiple of eight. The `alignof` operator takes a type and returns a value indicating the required alignment type. Alignment is a way to determine how information is arranged within a structure.

### Listing E.2    **align.cpp**

---

```
// align.cpp -- checking alignment
#include <iostream>
using namespace std;
struct things1
{
 char ch;
 int a;
 double x;
```

Here is the output on one system:

```
char alignment: 1
int alignment: 4
double alignment: 8
things1 alignment: 8
things2 alignment: 8
things1 size: 16
things2 size: 24
```

Both structures have an alignment requirement. The structure size should be a multiple of eight so that each element is adjacent to the next and also so that the size is a multiple of eight. The individual members of `things1` are of 13 bits, but the requirement of using a multiple of eight needs some padding in it. Next, within each `things2` is on a multiple of eight. The different arrangement of `things2` results in `things2` needing more in the end, so it is out right.



class that defines necessary properties that a `basic_string` should have. For example, it should have a `length` member function that returns the length of the string, represented as an array of type `charT`. The `charT(0)`, the generalization of the null character, is cast of 0 to type `charT`. It could be just 0, as `0`, or be an object created by a `charT` constructor, depending on the implementation, and so on. The `Allocator` parameter is the allocator used for the string. The default allocator is `allocator<charT>`, which uses standard ways.

There are four predefined specializations:

```
typedef basic_string<char> string;
typedef basic_string<char16_t> u16string;
typedef basic_string<char32_t> u32string;
typedef basic_string<wchar_t> wstring;
```

These specializations, in turn, use the following traits and allocators:

```
char_traits<char>
allocator<char>
char_traits<char16_t>
allocator<char16_t>
char_traits<char32_t>
allocator<char32_t>
char_traits<wchar_t>
allocator<wchar_t>
```



```
typedef typename traits::char_type
```

The keyword `typename` is used to tell the compiler that `traits::char_type` is a type. For the `string` specialization, `char_type` is `char`.

`size_type` is used like `size_of`, except that it measures the size of the stored type. For the `string` specialization, `size_type` is the same as `size_of`. It is a constant expression.

`difference_type` is used to measure the difference between two iterators, again in units corresponding to the size of a stored element. It is a signed version of the type underlying `size_type`.

For the `char` specialization, `pointer` is `char*`. For the `string` specialization, `pointer` is `string*`. However, if you create a specialization for a type that is not a pointer (pointer and reference) could refer to class or fundamental types, basic pointers and references.

To allow Standard Template Library (STL) containers to be used, the `string` plate defines some iterator types:

```
typedef (models random access iterator)
typedef (models random access iterator)
typedef std::reverse_iterator<iterator>
typedef std::reverse_iterator<const_iterator>
```

|                         |                                                                                                                                                                                         |
|-------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>cbegin()</code>   | A <code>const_iterator</code> .                                                                                                                                                         |
| <code>end()</code>      | An iterator that points to the end of the container.                                                                                                                                    |
| <code>cend()</code>     | A <code>const_iterator</code> pointing to the end of the container.                                                                                                                     |
| <code>rbegin()</code>   | A reverse iterator pointing to the first element of the container.                                                                                                                      |
| <code>crbegin()</code>  | A reverse <code>const_iterator</code> pointing to the first element of the container.                                                                                                   |
| <code>rend()</code>     | A reverse iterator pointing to the end of the container.                                                                                                                                |
| <code>crend()</code>    | A reverse <code>const_iterator</code> pointing to the end of the container. (C++11).                                                                                                    |
| <code>size()</code>     | The number of elements in the container, from <code>begin()</code> to <code>end()</code> .                                                                                              |
| <code>length()</code>   | The same as <code>size()</code> .                                                                                                                                                       |
| <code>capacity()</code> | The allocated memory space, which is greater than or equal to the <code>size()</code> . It can be increased or decreased, but it cannot be allocated.                                   |
| <code>max_size()</code> | The maximum number of elements that can be stored in the container.                                                                                                                     |
| <code>data()</code>     | A pointer of type <code>T*</code> pointing to the first element of an array. If the container is a string, the pointer should point to the first character of the string object itself. |



```
const Allocator& a = Allocator();
basic_string(initializer_list<charT>, const
~basic_string();
```

Some of the increase in constructors comes from the example, C++98 had this copy constructor:

```
basic_string(const basic_string& str, size_type npos,
size_type n = npos, const Allocator& a = Allocator())
```

C++11 replaces it with three constructors preceding list. This allows the most common use more efficiently. The really new additions are references, as discussed in Chapter 18, “Visiting a constructor with the `initializer_list` parameter.

Note that most of the constructors have a default `Allocator` parameter: `const Allocator& a = Allocator()`

Recall that the term `Allocator` is the term for a class to manage memory. The term `Allocator` is used. Thus, the constructors, by default, use the default. They give you the option of using some other. The following sections examine the constructors in

charT values:

```
basic_string(const charT* s, const Allocator& a, const traits_type& t,
```

To determine how many characters to copy, the `traits::length()` method to the array pointer (or null pointer.) For example, the following state is indicated character string:

```
string toast("Here's looking at you, kid.");
```

The `traits::length()` method for type `charT` returns how many characters to copy.

The following relationships hold after the construction:

- The `data()` method returns a pointer to the first character of the string.
- The `size()` method returns a value equal to the number of characters in the string.
- The `capacity()` method returns a value equal to the number of characters that can be stored in the string.

## Constructors That Use Part of a C-Style String

Constructors that use part of a C-style string (or a C-style string; more generally, they let you construct a string from an array of `charT` values:

```
basic_string(const charT* s, size_type n, const Allocator& a, const traits_type& t,
```

```
string ida(mel);
```

Here, `ida` would get a copy of the string `mel`.

The next constructor additionally requires a `pos`:

```
basic_string(const basic_string& str, const
```

The following relationships hold after either of these:

- The `data()` method returns a pointer to the first element pointed to by `str.data()`.
- The `size()` method returns the value of `str.size()`.
- The `capacity()` method returns a value of `str.capacity()`.

Moving along, the next constructor lets you specify the size:

```
basic_string(const basic_string& str, size_t n,
 const Allocator& a = Allocator());
```

The second argument `pos` specifies a location to begin the copying:

```
string att("Telephone home.");
string et(att, 4);
```

Position numbers begin with 0, so position 4 is the character `h` in `home`.  
to `"phone home."`.

## Constructors That Use an Rvalue

C++11 adds move semantics to the `string` class. This involves adding a move constructor, which uses a reference:

```
basic_string(basic_string&& str) noexcept
```

This constructor is invoked when the actual

```
string one("din"); // C-style string constructor
string two(one); // copy constructor
string three(one+two); // move constructor
```

As discussed in Chapter 18, the intent is that the object constructed by `operator+` rather than the original be destroyed.

The second rvalue constructor additionally

```
basic_string(const basic_string&& str, const
```

The following relationships hold after either

- The `data()` method returns a pointer to the first element is pointed to by `str.data()`.
- The `size()` method returns the value of `str.size()`.
- The `capacity()` method returns a value of `str.capacity()`.

The begin iterator points to the element  
end points to one past the last location to be

You can use this form with arrays, strings,

```
char cole[40] = "Old King Cole was a merry
string title(cole + 4, cole + 8);
vector<char> input;
char ch;
while (cin.get(ch) && ch != '\n')
 input.push_back(ch);
string str_input(input.begin(), input.end());
```

In the first use, InputIterator is evaluated  
InputIterator is evaluated to type vector<

The following relationships hold after the

- The `data()` method returns a pointer  
copying elements from the range `[begin, end)`.
- The `size()` method returns the distance  
measured in units equal to the size of  
dereferenced.)
- The `capacity()` method returns a value



```
void resize(size_type n)
```

```
void resize(size_type n, charT c)
```

```
void reserve(size_type res_arg = 0)
```

```
void shrink_to_fit()
```

```
void clear() noexcept
```

```
bool empty() const noexcept
```

---

```
cout << word.at(0); // display the t
```

The difference (besides the syntax differences) is that `at()` does bounds checking and throw an `out_of_range` exception if `pos` is type `size_type`, which is unsigned; the `operator[]()` methods don't do bounds checking. The `at()` methods don't do bounds checking if `pos >= size()`, except that the `const` version of `at()` if `pos == size()`.

Thus, you get a choice between safety (using `at()`) and execution speed (using array notation).

There is also a function that returns a new `basic_string` object: `substr()`.  
`basic_string substr(size_type pos = 0, size_type n = npos);`

It returns a string that's a copy of the string starting at `pos` characters or to the end of the string, whichever is less. For example, if `pet` initializes `pet` to "donkey":

```
string message("Maybe the donkey will leave");
string pet(message.substr(10, 6));
```

C++11 adds these four access methods:

```
const charT& front() const;
charT& front();
const charT& back() const;
charT& back();
```

```
source = "x",
join = name + source; // now with move s
awkward = {'C','l','o','u','s','e','a','u'}
```

## String Searching

The `string` class provides six search functions. The following sections describe them briefly.

### The `find()` Family

Here are the `find()` prototypes as provided by the C++ standard:

```
size_type find (const basic_string& str, size_type pos, size_type npos)
size_type find (const charT* s, size_type pos, size_type npos)
size_type find (const charT* s, size_type pos, size_type npos, const locale& loc)
size_type find (charT c, size_type pos = 0, size_type npos = string::npos, const locale& loc)
```

The first member returns the beginning position of the first occurrence of `str` in the invoking object, with the search beginning at `pos`. If no match is found, the method returns `npos`.

Here's code for finding the location of the first occurrence of "hat" in a string:

```
string longer("That is a funny hat.");
string shorter("hat");
size_type pos = longer.find(shorter);
if (pos != string::npos) {
 cout << "Found 'hat' at position " << pos << endl;
}
```

## The `rfind()` Family

The `rfind()` methods have these prototypes:

```
size_type rfind(const basic_string& str,
 size_type pos = npos) const;
size_type rfind(const charT* s, size_type pos) const;
size_type rfind(const charT* s, size_type pos, size_type n) const;
size_type rfind(charT c, size_type pos = npos) const;
```

These methods work like the analogous `find()` methods. If an occurrence of a string or character that starts at `pos` is found, the method returns the location of the first character. If not found, the method returns `npos`.

Here's code for finding the location of the last occurrence of "hat" at the end of the longer string:

```
string longer("That is a funny hat.");
string shorter("hat");
size_type loc1 = longer.rfind(shorter);
size_type loc2 = longer.rfind(shorter, loc1);
```

```

size_type find_last_of (const basic_string& s,
 size_type pos = npos,
 const charT* s2) const;
size_type find_last_of (const charT* s,
 size_type pos,
 const charT* s2) const;
size_type find_last_of (charT c, size_type pos,
 const charT* s) const;

```

These methods work like the corresponding `find` methods, but are looking for a match of the entire substring, the character, or the character in the substring.

Here's code for finding the location of the last occurrence of the letters in "fluke" in a longer string:

```

string longer("That is a funny hat.");
string shorter("hat");
size_type loc1 = longer.find_last_of(shorter);
size_type loc2 = longer.find_last_of("any");

```

The last occurrence of any of the three letters 'a', 'n', or 'y' is the last occurrence of any of the three characters 'h', 'a', or 't'.

## **The `find_first_not_of()` Family**

The `find_first_not_of()` methods have the following signatures:

```

size_type find_first_not_of(const basic_string& s,
 size_type pos = 0,
 const charT* s2) const;
size_type find_first_not_of(const charT* s,
 size_type pos,
 const charT* s2) const;
size_type find_first_not_of(charT c, size_type pos,
 const charT* s) const;

```

```

size_type n)
size_type find_last_not_of (const charT*
size_type pos
size_type find_last_not_of (charT c, size

```

These methods work like the corresponding ones for `string` and `wstring`; they search for the last occurrence of any character in the range `[first, last)`.

Here's code for finding the location of the last space character not in "This" in a longer string:

```

string longer("That is a funny hat.");
string shorter("That.");
size_type loc1 = longer.find_last_not_of(' ');
size_type loc2 = longer.find_last_not_of(' ', loc1);

```

The last space in `longer` is the last character in `shorter`. The `f` in the `longer` string is the last character not in `shorter` through position 10.

## Comparison Methods and Functions

The `string` class offers methods and functions for comparing strings. Here are the method prototypes:

```
int a13 = s1.compare(s3); // a13 is < 0
int a12 = s1.compare(s2); // a12 is > 0
```

The second method is like the first, except that it compares the substring starting at position `pos1` in the first string for the comparison.

The following example compares the first four characters of the first string with the first four characters of the second string:

```
string s1("bellflower");
string s2("bell");
int a2 = s1.compare(0, 4, s2); // a2 is 0
```

The third method is like the first, except that it compares the substring starting at position `pos1` in the first string, and `n2` characters of the second string, for the comparison. For example, the following code compares the first four characters of the first string with the first four characters of the second string:

```
string st1("stout boar");
string st2("mad about ewe");
int a3 = st1.compare(2, 3, st2, 6, 3); // a3 is 0
```

The fourth method is like the first, except that it compares the substring starting at position `pos1` in the first string, and `n2` characters of the second string, for the comparison. For example, the following code compares the first four characters of the first string with the first four characters of the second string:

The fifth and sixth methods essentially are the same as the fourth method, except that they take a character array instead of a string object for the second string.

`append()` method. All throw a `length_error` if the new length would be greater than the maximum string size. The `+=` operator can be used to append a string to an array, or an individual character to another string.

```
basic_string& operator+=(const basic_string& s);
basic_string& operator+=(const charT* s);
basic_string& operator+=(charT c);
```

The `append()` methods also let you append an individual character to another string. In addition, you can append a string to a string object by specifying an initial position and a final position, or by specifying a range. You can append part of a string to another string by specifying a range. You can append part of the string to use. The version for appending a character to a string appends a specified number of instances of that character to copy. Here are the methods:

```
basic_string& append(const basic_string& s);
basic_string& append(const basic_string& s, const_iterator first, const_iterator last);
basic_string& append(const basic_string& s, size_type n);
template<class InputIterator>
basic_string& append(InputIterator first, InputIterator last);
basic_string& append(const charT* s);
basic_string& append(const charT* s, size_type n);
basic_string& append(size_type n, charT c);
void push_back(charT c);
```



ods, which allow you to assign a whole string of characters to a string object. Here are the methods:

```
basic_string& assign(const basic_string& s);
basic_string& assign(basic_string&& str);
basic_string& assign(const basic_string& s,
 size_type n);
basic_string& assign(const charT* s, size_type n);
basic_string& assign(const charT* s);
basic_string& assign(size_type n, charT c);
template<class InputIterator>
 basic_string& assign(InputIterator first, InputIterator last);
basic_string& assign(initializer_list<charT> ilist);
```

Here are a couple examples:

```
string test;
string stuff("set tubs clones ducks");
test.assign(stuff, 1, 5); // test is "et"
test.assign(6, '#'); // test is "####"
```

```

iterator insert(const_iterator p, const char* s)
iterator insert(const_iterator p, size_type n, const char* s)
template<class InputIterator>
 void insert(iterator p, InputIterator first, InputIterator last)
iterator insert(const_iterator p, initial_size_type n, const char* s)

```

For example, the following code inserts the string "former " before the string "banker.":

```

string st3("The banker.");
st3.insert(4, "former ");

```

Then the following code inserts the string "waltzed!" (which would be the ninth character) just before the string "banker.":

```

st3.insert(st3.size() - 1, " waltzed!", 8);

```

## Erase Methods

The erase() methods remove characters from a string.

```

basic_string& erase(size_type pos = 0, size_type n = npos)
iterator erase(const_iterator position);
iterator erase(const_iterator first, const_iterator last);
void pop_back();

```

```

 size_type n2);
basic_string& replace(size_type pos, size_type n, const basic_string& s);
basic_string& replace(size_type pos, size_type n, const charT* s, size_type n2);
basic_string& replace(const_iterator i1, const_iterator i2, const basic_string& s);
basic_string& replace(const_iterator i1, const_iterator i2, const charT* s, size_type n2);
basic_string& replace(const_iterator i1, const_iterator i2, const charT* s);
basic_string& replace(const_iterator i1, const_iterator i2, const basic_string& s,
 size_type n, charT c);
template<class InputIterator>
 basic_string& replace(const_iterator i1, const_iterator i2, InputIterator j1,
 InputIterator j2);
basic_string& replace(const_iterator i1, const_iterator i2, const charT* s,
 const basic_string& initializer);
list<charT> replace(const_iterator i1, const_iterator i2, const charT* s,
 const list<charT>& initializer);

```

Here is an example:

```

string test("Take a right turn at Main Street");
test.replace(7,5,"left"); // replace right

```





chapter assumes that you know about iterators.

## The STL and C++11

Just as the changes brought by C++11 to the STL are covered completely in this book, so are the changes to the STL in this appendix. However, we can summarize the changes.

C++11 brings several new elements to the STL. First, it adds new containers. Second, it adds a few new features to the old containers. Third, it adds new functions to its family of algorithms. All these changes are covered in this appendix, but you may find an overview of the first three changes in the next section.

### New Containers

C++11 adds the following containers: `array`, `unordered_map`, `unordered_multimap`, `unordered_set`, and `unordered_multiset`.

An `array` container, once declared, is fixed in size and stored in statically allocated memory rather than dynamically allocated memory. It is similar to `vector`, but it is of a fixed array type; it's more limited than `vector`, but it is faster.

The `list` container is a bidirectional linked list. It is similar to `vector`, but it is not stored in a contiguous memory block. It has two pointers, one at each end, linked to the item before it and the one after it.

emplacement is a means to increase efficiency

```
class Items
{
 double x;
 double y;
 int m;
public:
 Items();
 Items (double xx, double yy, int mm);
 ...
};
...
vector<Items> vt(10);
...
vt.push_back(Items(8.2, 2.8, 3)); //
```

The call to `insert()` causes the memory of the old object at the end of `vt`. Next, the `Items()` constructor of this object is copied to a location at the front of the vector. The old object is deleted. With C++11, you can do this

```
vi.emplace_back(8.2, 2.8, 3);
```

|                                 |                                                                                                                                               |
|---------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------|
| <code>X::value_type</code>      | T, the element type                                                                                                                           |
| <code>X::reference</code>       | T &.                                                                                                                                          |
| <code>X::const_reference</code> | const T &.                                                                                                                                    |
| <code>X::iterator</code>        | Iterator type pointing to the first element.                                                                                                  |
| <code>X::const_iterator</code>  | Iterator type pointing to the first element.                                                                                                  |
| <code>X::difference_type</code> | Signed integral type used to compute the difference between two iterators (the number of elements between two iterators to another iterator). |
| <code>X::size_type</code>       | Unsigned integral type used to count the number of objects, number of elements.                                                               |

---

The class definition uses a typedef to define a type that can be used to declare suitable variables. For example, the code below replaces the first occurrence of "bonus" in a vector with "prize". In order to show how you can use member type

```
using namespace std;
vector<string> input;
string temp;
while (cin >> temp && temp != "quit")
 input.push_back(temp);
vector<string>::iterator want=
```





## Operation

## Description

|                                 |                                                                |
|---------------------------------|----------------------------------------------------------------|
| <code>X u;</code>               | Constructs an empty object                                     |
| <code>X()</code>                | Constructs an empty object                                     |
| <code>X(a)</code>               | Constructs a copy of object <code>a</code>                     |
| <code>X u(a)</code>             | <code>u</code> is a copy of <code>a</code> (copy constructor)  |
| <code>X u = a;</code>           | <code>u</code> is a copy of <code>a</code> (copy constructor)  |
| <code>r = a</code>              | <code>r</code> equals the value of <code>a</code> (assignment) |
| <code>X u(rv)</code>            | <code>u</code> equals the value that <code>rv</code> points to |
| <code>X u = rv;</code>          | <code>u</code> equals the value that <code>rv</code> points to |
| <code>a = rv</code>             | <code>a</code> equals the value that <code>rv</code> points to |
| <code>(&amp;a) -&gt;~X()</code> | Destructor applied to every object                             |
| <code>begin()</code>            | Returns an iterator to the beginning                           |
| <code>end()</code>              | Returns an iterator to past the end                            |
| <code>cbegin()</code>           | Returns a <code>const</code> iterator to the beginning         |
| <code>cend()</code>             | Returns a <code>const</code> iterator to past the end          |
| <code>size()</code>             | Returns the number of elements                                 |
| <code>maxsize()</code>          | Returns the size of the largest container                      |
| <code>empty()</code>            | Returns <code>true</code> if the container is empty            |

The unordered set and unordered map containers support the following optional container operations in Table G.4, but not the other operations.

Table G.4    **Optional Container Operations**

| Operation          | Description                                                                                                       |
|--------------------|-------------------------------------------------------------------------------------------------------------------|
| <code>&lt;</code>  | <code>a &lt; b</code> returns <code>true</code> if <code>a</code> is lexicographically less than <code>b</code> . |
| <code>&gt;</code>  | <code>a &gt; b</code> returns <code>b &lt; a</code> .                                                             |
| <code>&lt;=</code> | <code>a &lt;= b</code> returns <code>!(a &gt; b)</code> .                                                         |
| <code>&gt;=</code> | <code>a &gt;= b</code> returns <code>!(a &lt; b)</code> .                                                         |

The `>` operator for a container assumes that the container is ordered. A lexicographic comparison is a generalization of the comparison of two containers, element-by-element, until it differs. If the containers are identical through the end, the containers are considered to be in the same order as the first container. For example, if two containers are identical through the end, the first container is less than the second.

valid dereferenceable const iterator to a. [q1]  
 an integer of `X::size_type`. Args is a template  
 parameter pack with the pattern `Args&&`.

**Table G.5 Additional Operations Defined for**

| <b>Operation</b>                 | <b>Description</b>                                                                                          |
|----------------------------------|-------------------------------------------------------------------------------------------------------------|
| <code>X(n, t)</code>             | Constructs a sequence of <code>n</code> elements of type <code>t</code> .                                   |
| <code>X a(n, t)</code>           | Constructs a sequence of <code>n</code> elements of type <code>t</code> .                                   |
| <code>X(i, j)</code>             | Constructs a sequence of <code>j - i + 1</code> elements of type <code>t</code> .                           |
| <code>X a(i, j)</code>           | Constructs a sequence of <code>j - i + 1</code> elements of type <code>t</code> .                           |
| <code>X(il);</code>              | Constructs a sequence of <code>il</code> elements of type <code>t</code> .                                  |
| <code>a = il;</code>             | Copies the values of <code>il</code> into <code>a</code> .                                                  |
| <code>a.emplace(p, args);</code> | Inserts an object of type <code>t</code> at position <code>p</code> using the arguments <code>args</code> . |
| <code>a.insert(p, t)</code>      | Inserts a copy of <code>t</code> at position <code>p</code> .                                               |
| <code>a.insert(p, rv)</code>     | Inserts a copy of <code>rv</code> at position <code>p</code> .                                              |

|                              |                                          |
|------------------------------|------------------------------------------|
|                              | ment that had follo                      |
| <code>a.erase(q1, q2)</code> | Erases the elemen<br>pointing to the ele |
| <code>a.clear()</code>       | Does the same thi                        |
| <code>a.front()</code>       | Returns <code>*a.begin</code>            |

---

Table G.6 lists methods common to some `forward_list`, `list`, and `deque`).

**Table G.6 Operations Defined for Some Seq**

| <b>Operation</b>             | <b>Description</b>                                                    |
|------------------------------|-----------------------------------------------------------------------|
| <code>a.back()</code>        | Returns <code>*--a.end()</code>                                       |
| <code>a.push_back(t)</code>  | Inserts <code>t</code> before <code>a.end()</code>                    |
| <code>a.push_back(rv)</code> | Inserts <code>t</code> before <code>a.end()</code><br>move semantics. |
| <code>a.pop_back()</code>    | Erases the last elem                                                  |

The vector template additionally has the container and `n` is an integer of `X::size_type`

Table G.7 Additional Operations for Vectors

| Operation                 | Description                                                                                                                                                                                                                                                                                                                             |
|---------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>a.capacity()</code> | Returns the total number of elements the vector can hold without requiring reallocation.                                                                                                                                                                                                                                                |
| <code>a.reserve(n)</code> | Alerts object <code>a</code> that <code>n</code> more elements may be added without the need for reallocation. After the method call, the number of elements in <code>a</code> may increase. Reallocation occurs if the new capacity is less than <code>n</code> or greater than <code>a.max_size()</code> , which throws an exception. |

|                                       |                                                           |
|---------------------------------------|-----------------------------------------------------------|
| <code>a.merge(b)</code>               | Merges the<br>tor defined<br>lent to an e<br>List b is em |
| <code>a.merge(b, Compare comp)</code> | Merges the<br>function or<br>to an eleme<br>b is empty    |
| <code>a.sort()</code>                 | Sorts list a                                              |
| <code>a.sort(Compare comp)</code>     | Sorts list a                                              |
| <code>a.reverse()</code>              | Reverses th                                               |

---

The forward-list operations are similar. The forward-list class iterator can't go backwards, some methods like `erase()`, and `splice()` methods are replaced by `splice_before()` and `splice_after()` methods, all of which operate on a position rather than preceding it.

Associative containers provide the method `compare`. The comparison object need not require that values with *equivalent keys* means that two values, which may be equal, are not. In the table, `x` is a container class, and `a` is an object of type `x` (is, is set or map), `a_uniq` is an object of type `x` (is, is set or multimap), `a_eq` is an object of type `x`. As to elements of `value_type`, `[i, j)` is a valid range of dereferenceable iterators to `a`, `[q1, q2)` is a valid range of dereferenceable iterators to `a_uniq` (which may be a pair), and `k` is a value of `x`: `initializer_list<value_type>` object.



`a_eq.insert(t)`

component points  
key of `t`.

`a.insert(p,t)`

Inserts `t` and returns  
Inserts `t`, using `p` as a hint for  
search. If `a` is a `multiset`,  
if and only if `a` does not contain  
key; otherwise, insertion  
takes place, the mapped value  
location with an empty value.

`a.insert(i,j)`

Inserts elements `t` in the range `[i,j)`.

`a.insert(il)`

Inserts elements `t` in the range `[il,il+1)`.

`a_uniq.emplace(args)`

Like `a_uniq.insert(t)`, but  
parameter list may be a  
pack.

`a_eq.emplace(args)`

Like `a_eq.insert(t)`, but  
parameter list matches the  
signature of `operator()`.

`a.emplace_hint(args)`

Like `a.insert(p,t)`, but  
parameter list matches the  
signature of `operator()`.

`a.erase(k)`

Erases all elements with  
key `k` and returns the number  
of elements erased.

# Unordered Associative C

As mentioned earlier, the unordered associative containers, `unordered_multiset`, `unordered_map`, and `unordered_multimap`, use buckets and buckets tables to provide rapid access to data. Let's take `unordered_map` as an example. A `hash` function converts a key to an index value. For example, a `hash` function could sum the numeric codes for the characters in the key and then take the modulus 13, thus giving an index in the range 0 to 12. This gives us 13 buckets to store strings. Any string with a hash value of 4 goes into bucket 4. If you wished to search the container for a string, you would calculate the hash value from the key and just search the bucket with that index. You would have enough buckets that each one would have a small number of strings.

The C++11 library provides a `hash<Key>` interface for the unordered containers use by default. Specializations are provided for integral types, pointer types, point types, for pointers, and for some templates.

Table G.11 lists types used for these containers.

The interface for the unordered associative containers is defined in `<unordered_map>`. In particular, Table G.10 also lists the exceptions with the following exceptions: The `lower_bound` and `upper_bound` methods are not required, nor is the `x(i, j, c)` constructor. The ordered containers are ordered allows them to use a comparison function.

container class, `a` is an object of type `X`, `b` is a  
 an object of type `unordered_set` or `unordered`  
`unordered_multiset` or `unordered_multima`  
 of type `key_equal`, `n` is a value of type `size_`  
 before, `i` and `j` are input iterators referring to  
 range, `p` and `q2` are iterators to `a`, `q` and `q1` are  
 valid range, `t` is a value of `X::value_type` (w  
`X::key_type`. Also `il` is an `initializer_lis`

**Table G.12 Additional Operations Defined for  
Multimaps**

| Operation                   | Description                                                                              |
|-----------------------------|------------------------------------------------------------------------------------------|
| <code>X(n, hf, eq)</code>   | Constructs a<br>using <code>hf</code> as<br>predicate. If<br>key equality<br>used as the |
| <code>X a(n, hf, eq)</code> | Constructs a<br>using <code>hf</code> as<br>predicate. If<br>key equality<br>used as the |

|                                   |                                                                                                                                                                    |
|-----------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>b.max_bucket_count()</code> | Returns an integer value representing the number of buckets that the bucket container contains.                                                                    |
| <code>b.bucket(k)</code>          | Returns the bucket index of the element with a key <code>k</code> .                                                                                                |
| <code>b.bucket_size(n)</code>     | Returns the number of elements in the bucket with index <code>n</code> .                                                                                           |
| <code>b.begin(n)</code>           | Returns an iterator to the first element in the bucket with index <code>n</code> .                                                                                 |
| <code>b.end(n)</code>             | Returns an iterator to the end of the bucket with index <code>n</code> .                                                                                           |
| <code>b.cbegin(n)</code>          | Returns a constant iterator to the first element in the bucket with index <code>n</code> .                                                                         |
| <code>b.cbegin(n)</code>          | Returns a constant iterator to the first element in the bucket with index <code>n</code> .                                                                         |
| <code>b.load_factor()</code>      | Returns the load factor of the bucket container.                                                                                                                   |
| <code>b.max_load_factor()</code>  | Returns the maximum load factor of the bucket container; the container will automatically resize the buckets when the load factor exceeds the maximum load factor. |
| <code>b.max_load_factor(z)</code> | May change the maximum load factor of the bucket container to <code>z</code> .                                                                                     |

related operators, and numeric operations. (C moved the STL to the numeric library, but that doesn't affect this *operation* indicates that the function takes a pair or sequence, to be operated on. The term *mutator* indicates the container.

## **Nonmodifying Sequence Operations**

Table G.13 summarizes the nonmodifying sequence operations shown, and overloaded functions are listed just below the prototypes, follows the table. Thus, you can see what an operation does and then look up the details if you need them.

|                              |                                                                                         |
|------------------------------|-----------------------------------------------------------------------------------------|
| <code>for_each()</code>      | Applies a nonmodifying                                                                  |
| <code>find()</code>          | Finds the first occurrence                                                              |
| <code>find_if()</code>       | Finds the first value that                                                              |
| <code>find_if_not()</code>   | Finds the first value that<br>range. (C++11)                                            |
| <code>find_end()</code>      | Finds the last occurrence<br>values of a second sequence<br>applying a binary predicate |
| <code>find_first_of()</code> | Finds the first occurrence<br>matches a value in the<br>may be evaluated with           |
| <code>adjacent_find()</code> | Finds the first element<br>ing it. Matching may be<br>predicate.                        |
| <code>count()</code>         | Returns the number of                                                                   |
| <code>count_if()</code>      | Returns the number of<br>with a match determined                                        |

## **all\_of() (C++11)**

```
template<class InputIterator, class Predicate>
bool all_of(InputIterator first, InputIterator last,
 Predicate pred);
```

The `all_of()` function returns true if `pred` returns true for all elements in the range `[first, last)` or if the range is empty. Otherwise, it returns false.

## **any\_of() (C++11)**

```
template<class InputIterator, class Predicate>
bool any_of(InputIterator first, InputIterator last,
 Predicate pred);
```

The `any_of()` function returns false if `pred` returns false for all elements in the range `[first, last)` or if the range is empty. Otherwise, it returns true.

## **none\_of() (C++11)**

```
template<class InputIterator, class Predicate>
bool none_of(InputIterator first, InputIterator last,
 Predicate pred);
```

The `none_of()` function returns true if `pred` returns false for all elements in the range `[first, last)` or if the range is empty. Otherwise, it returns false.

last) for which the function object can predicate. If no such item is found, last is not found.

### **find\_if\_not()**

```
template<class InputIterator, class Predicate>
InputIterator find_if_not(InputIterator first, InputIterator last,
 Predicate pred);
```

The `find_if_not()` function returns an iterator [first, last) for which the function object returns false. If no such item is not found.

### **find\_end()**

```
template<class ForwardIterator1, class ForwardIterator2>
ForwardIterator1 find_end(ForwardIterator1 first1, ForwardIterator1 last1,
 ForwardIterator2 first2, ForwardIterator2 last2);
```

```
template<class ForwardIterator1, class ForwardIterator2,
 class BinaryPredicate>
ForwardIterator1 find_end(ForwardIterator1 first1, ForwardIterator1 last1,
 ForwardIterator2 first2, ForwardIterator2 last2,
 BinaryPredicate pred);
```



uses the == operator for the value type to compare elements. The first version takes a binary predicate function object pred to compare elements. Elements found by it1 and it2 match if pred(\*it1, \*it2) is true. If no such pair is found, not found.

## **adjacent\_find()**

```
template<class ForwardIterator>
ForwardIterator adjacent_find(ForwardIterator first1,
 ForwardIterator last);
```

```
template<class ForwardIterator, class BinaryPredicate>
ForwardIterator adjacent_find(ForwardIterator first1,
 ForwardIterator last, BinaryPredicate pred);
```

The adjacent\_find() function returns an iterator [first1, last1) such that the element matches the element at [first2, last2). If no such pair is found, returns last. The first version uses the == operator to compare elements. The second version uses pred to compare elements. That is, elements p1 and p2 match if pred(\*it1, \*it2) is true.

```

 mismatch(InputIterator1 first1,
 InputIterator1 last1, InputIterator2 first2,
 InputIterator2 last2, const BinaryPredicate& pred) const;

template<class InputIterator1, class InputIterator2, class BinaryPredicate>
pair<InputIterator1, InputIterator2>
mismatch(InputIterator1 first1,
 InputIterator1 last1, InputIterator2 first2,
 InputIterator2 last2, BinaryPredicate pred);

```

Each of the `mismatch()` functions finds the first element that doesn't match the corresponding element. If no mismatch is found, the return value is `pair<last1, first2>`. If a mismatch is found, the return value is `pair<last1, first2>` where `last1` is the iterator to the first mismatch found, the return value is `pair<last1, first2>`. The second function uses the `==` operator to test matching. The second function uses the `==` operator to test matching. The second function uses the `==` operator to test matching. The second function uses the `==` operator to test matching. That is, the elements match if `pred(*it1, *it2)` is false.

## **equal()**

```

template<class InputIterator1, class InputIterator2>
bool equal(InputIterator1 first1, InputIterator1 last1,
 InputIterator2 first2);

```

last1) matches the corresponding element in sequence beginning at first2; it returns false. The second version of the == operator for the value type to compare elements uses a predicate function object pred to compare elements. It returns true if \*it1 and \*it2 match if pred(\*it1, \*it2) is true.

## **search()**

```
template<class ForwardIterator1, class ForwardIterator2>
ForwardIterator1 search(
 ForwardIterator1 first1, ForwardIterator1 last1,
 ForwardIterator2 first2, ForwardIterator2 last2)

template<class ForwardIterator1, class ForwardIterator2,
 class BinaryPredicate>
ForwardIterator1 search(
 ForwardIterator1 first1, ForwardIterator1 last1,
 ForwardIterator2 first2, ForwardIterator2 last2,
 BinaryPredicate pred);
```

The search() function finds the first occurrence in the second sequence that matches the corresponding sequence found in the first sequence. It returns last1 if no such sequence is found. The first version is for the value type to compare elements. The second version

overloaded functions are listed just once. A function follows the table. Thus, you can scan the table and then look up the details if you find the function.

Now let's take a more detailed look at the `copy()` function, the discussion shows the prototype. As we saw earlier, pairs of iterators indicate ranges, with the first indicating the type of iterator. As usual, a range is first up to, but not including, last. Functions which can be function pointers or objects for Chapter 16, a predicate is a Boolean function. `copy()` is a Boolean function with two arguments. (5) It returns 0 as they return 0 for false and a nonzero value for true. A function object is one that takes a single argument and that takes two arguments.

## **copy()**

```
template<class InputIterator, class OutputIterator>
OutputIterator copy(InputIterator first,
 InputIterator last,
 OutputIterator result)
```

The `copy()` function copies the elements in the range `[first, last)` to the range `[result, result + (last - first))`. It returns `result + (last - first)`.

|                                |                                                                                                                                            |
|--------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------|
| <code>swap()</code>            | Exchanges two values                                                                                                                       |
| <code>swap_ranges()</code>     | Exchanges corresponding                                                                                                                    |
| <code>iter_swap()</code>       | Exchanges two values                                                                                                                       |
| <code>transform()</code>       | Applies a function object to each element in a pair of ranges, storing the result at the corresponding location of another range           |
| <code>replace()</code>         | Replaces each occurrence of a value with another value                                                                                     |
| <code>replace_if()</code>      | Replaces each occurrence of a value with another value if a predicate function object returns true                                         |
| <code>replace_copy()</code>    | Copies one range to another, replacing each occurrence of a specified value with another value                                             |
| <code>replace_copy_if()</code> | Copies one range to another, replacing each occurrence of a specified value with another value if a predicate function object returns true |
| <code>fill()</code>            | Sets each value in a range to a specified value                                                                                            |
| <code>fill_n()</code>          | Sets <i>n</i> consecutive elements in a range to a specified value                                                                         |
| <code>generate()</code>        | Sets each value in a range to the result of a function object that takes the current index as an argument                                  |
| <code>generate_n()</code>      | Sets the first <i>n</i> values in a range to the result of a function object that takes the current index as an argument                   |

|                                 |                                                                                                                                             |
|---------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------|
| <code>rotate()</code>           | Creates a range as a sub-range of the original range.                                                                                       |
| <code>rotate_copy()</code>      | Copies one range to another, rotating the process.                                                                                          |
| <code>random_shuffle()</code>   | Randomly rearranges the elements in a range.                                                                                                |
| <code>shuffle()</code>          | Randomly rearranges the elements in a range of type satisfying C++11 (C++11)                                                                |
| <code>is_partitioned()</code>   | Returns true if a range is partitioned.                                                                                                     |
| <code>partition()</code>        | Places all the elements in a range that don't satisfy a predicate at the beginning of the range.                                            |
| <code>stable_partition()</code> | Places all the elements in a range that don't satisfy a predicate at the beginning of the range, preserving the relative order of elements. |
| <code>partition_copy()</code>   | Copies all the elements in a range that satisfy a predicate to the output range and the rest to the output range. (C++11)                   |
| <code>partition_point()</code>  | For a range partitioned by <code>partition()</code> , returns the first element not satisfying the predicate.                               |

---

## **copy\_backward()**

```
template<class BidirectionalIterator1,
 class BidirectionalIterator2>
BidirectionalIterator2 copy_backward(BidirectionalIterator1 first,
BidirectionalIterator1 last, BidirectionalIterator2 result)
```

The `copy_backward()` function copies the elements in the range `[first, last)` into the range `[result, result + (last - first))`. The elements are being copied to location `result - 1` and processed in reverse order. The function returns `result - (last - first)`—that is, an iterator pointing to the location before the first element copied. The function requires that `result` be a bidirectional iterator. However, because copying is done backward, it is possible to use a reverse iterator.

## **move() (C++11)**

```
template<class InputIterator, class OutputIterator>
OutputIterator copy(InputIterator first, InputIterator last,
OutputIterator result)
```

The `move()` function uses `std::move()` to move the elements in the range `[first, last)` into the range `[result, result + (last - first))`. The elements are being copied to location `result - 1` and processed in reverse order. The function returns `result - (last - first)`—that is, an iterator pointing one past the last element copied.

The `swap()` function exchanges values stored in memory. In C++11, `swap()` is overloaded. (C++ moves this function to the utility header `<utility>`.)

### **swap\_ranges()**

```
template<class ForwardIterator1, class ForwardIterator2, class ForwardIterator3>
ForwardIterator2 swap_ranges(ForwardIterator1 first1, ForwardIterator1 last1,
 ForwardIterator2 first2, ForwardIterator3 last2)
```

The `swap_ranges()` function exchanges values in the range `[first1, last1)` with the corresponding values in the range `[first2, last2)` beginning at `first1`.

### **iter\_swap()**

```
template<class ForwardIterator1, class ForwardIterator2>
void iter_swap(ForwardIterator1 a, ForwardIterator2 b)
```

The `iter_swap()` function exchanges values stored in memory at `a` and `b`.

### **transform()**

```
template<class InputIterator, class OutputIterator, class UnaryOperation>
OutputIterator transform(InputIterator first, InputIterator last,
 OutputIterator result, UnaryOperation op)
```



[first, last) with the value new\_value.

### **replace\_if()**

```
template<class ForwardIterator, class Predicate>
void replace_if(ForwardIterator first, ForwardIterator last,
 Predicate pred, const T& new_value)
```

The `replace()_if` function replaces each value `old` in the range `[first, last)` which `pred(old)` is true with the value `new_value`.

### **replace\_copy()**

```
template<class InputIterator, class OutputIterator>
OutputIterator replace_copy(InputIterator first, InputIterator last,
 OutputIterator result, const T& old_value,
 const T& new_value)
```

The `replace_copy()` function copies the elements in the range `[first, last)` beginning at `result` but substituting `new_value` for `old_value`. It returns `result + (last - first)`, the past-the-end iterator.

### **replace\_copy\_if()**

```
template<class InputIterator, class OutputIterator, class Predicate>
OutputIterator replace_copy_if(InputIterator first, InputIterator last,
 OutputIterator result, Predicate pred,
 const T& new_value)
```

The `generate()` function sets each element of `first` to `gen()`. `gen` is a generator function object—that is, `gen` can be a pointer to `rand()`.

### **`generate_n()`**

```
template<class OutputIterator, class Size>
void generate_n(OutputIterator first, Size n, const Generator& gen);
```

The `generate_n()` function sets each of the `n` elements of `first` to `gen()`, where `gen` is a generator function object. `gen` can be a pointer to `rand()`. For example, `gen` can be a pointer to `rand()`.

### **`remove()`**

```
template<class ForwardIterator, class T>
ForwardIterator remove(ForwardIterator first, ForwardIterator last,
 const T& value);
```

The `remove()` function removes all occurrences of `value` from the range `[first, last)` and returns a past-the-end iterator for the resulting range. The order of the unremoved elements is preserved.

The `remove_copy()` function copies values from `first` to `last` to `result`, beginning at `result`, skipping instances of `val`. The function returns an iterator for the resulting range. The function does not move elements; the range of moved elements is unaltered.

### **`remove_copy_if()`**

```
template<class InputIterator, class OutputIterator>
OutputIterator remove_copy_if(InputIterator first, InputIterator last,
 OutputIterator result)
```

The `remove_copy_if()` function copies values from `first` to `last` to `result`, beginning at `result`, skipping instances of `val`. It returns a past-the-end iterator for the resulting range. The function does not move elements; that the order of the unremoved elements is unaltered.

### **`unique()`**

```
template<class ForwardIterator>
ForwardIterator unique(ForwardIterator first, ForwardIterator last)
```

```
template<class ForwardIterator, class BinaryPredicate>
ForwardIterator unique(ForwardIterator first, ForwardIterator last,
 BinaryPredicate pred)
```

## **reverse()**

```
template<class BidirectionalIterator>
void reverse(BidirectionalIterator first,
```

The `reverse()` function reverses the elements in the range `[first, last)`, and so on.

## **reverse\_copy()**

```
template<class BidirectionalIterator, class OutputIterator>
OutputIterator reverse_copy(BidirectionalIterator first,
 BidirectionalIterator last,
 OutputIterator result)
```

The `reverse_copy()` function copies the elements in the range `[first, last)` to the range `[result, result + (last - first))` in reverse order. The original range is not modified.

## **rotate()**

```
template<class ForwardIterator>
void rotate(ForwardIterator first, ForwardIterator middle,
 ForwardIterator last);
```

The `rotate()` function performs a left rotation of the elements in the range `[first, last)`. The element at `middle` is moved to the first position, and the elements after `middle` are shifted one position to the right.

This version of the `random_shuffle()` function takes two iterators `[first, last)`. The function object `random` determines the distribution. The expression `random(n)` should return a value in the range `[0, n)`. In C++11, the argument was an lvalue reference; in C++11, it is a constant reference.

## **shuffle()**

```
template<class RandomAccessIterator, class RandomNumberGenerator>
void shuffle(RandomAccessIterator first, RandomAccessIterator last,
 RandomNumberGenerator& gen);
```

This version of the `random_shuffle()` function takes two iterators `[first, last)`. The type of the function object `random` determines the distribution. The expression `random(n)` should return a value in the range `[0, n)`. In C++11, the argument was an lvalue reference; in C++11, it is a constant reference.

## **is\_partitioned() (C++11)**

```
template<class InputIterator, class Predicate>
bool is_partitioned(InputIterator first, InputIterator last, Predicate pred);
```

relative ordering within each of the two groups. The following algorithm moves the elements of the first group, following the last position holding a value for w

## **partition\_copy() (C++11)**

```
template<class InputIterator, class OutputIterator, class Predicate>
OutputIterator partition_copy(InputIterator first, InputIterator last,
 OutputIterator out_true, OutputIterator out_false,
 Predicate pred)
```

The `partition_copy()` function copies each element from the range `[first, last)` to the range beginning with `out_true` if `pred(val)` is true to the range beginning with `out_false`. It returns the range beginning with `out_true` and the range beginning with `out_false`. It returns the end of the range that begins with `out_true` and the end of the range that begins with `out_false`.

## **partition\_point() (C++11)**

```
template<class ForwardIterator, class Predicate>
ForwardIterator partition_point(ForwardIterator first, ForwardIterator last,
 Predicate pred)
```

full sort.

`partial_sort_copy()`

Copies a p

`is_sorted()`

Returns tr

`is_sorted_until()`

Returns th

(C++11)

`nth_element()`

Given an it

would be t

element th

`lower_bound()`

Given a va

before whi

the orderi

`upper_bound()`

Given a va

before whi

the orderi

`equal_range()`

Given a va

range such

ment in th

`binary_search()`

Returns tr

equivalent

`sort_heap()`

Sorts a h

`is_heap()`

Returns t

`is_heap_until()`

Returns t

(C++11)

`min()`

Returns t

an initi

`max()`

Returns t

initial

`minmax()`

Returns a

ments in

largest ite

`min_element()`

Finds the

range.

`max_element()`

Finds the

`minmax_element()`

Returns a

occurrenc

tor to the

(C++11)



- If *a* is equivalent to *b* and *b* is equivalent to *c*, then *a* is equivalent to *c* (equivalency is a transitive relationship).

If you think of applying `<` to integers, then `<` doesn't have to hold for more general cases. For example, with several members describing a mailing address, a function that orders the structures according to zip code would be equivalent but not equal.

Now let's take a more detailed look at the `sort` function, the discussion shows the prototype(`sort`), the function is divided into several subsections. As you saw, the function ranges, with the chosen template parameter `bool`, a range in the form `[first, last)` goes from `first` to `last`. Functions passed as arguments are function objects, the `()` operation is defined. As you learned in Chapter 1, a function with one argument, and a binary predicate with two arguments. (The functions need not be type `bool`, any nonzero value for `true`.) Also as in Chapter 1, a function with a single argument, and a binary function object.



```
Compare comp);
```

The `is_sorted()` function returns `true` if the range is sorted, and `false` otherwise. The first version uses `<`, and the second version uses a custom `comp` to determine the order.

### **`is_sorted_until()` (C++11)**

```
template<class ForwardIterator>
ForwardIterator is_sorted_until(ForwardIterator first,
 ForwardIterator last) {
 return first;
}

template<class ForwardIterator, class Compare>
ForwardIterator is_sorted_until(ForwardIterator first,
 ForwardIterator last,
 Compare comp);
```

The `is_sorted_until()` function returns the iterator to the first element that is not sorted. Otherwise, it returns the last iterator. The first version uses `<`, and the second version uses a custom `comp` to determine the order.

### **`nth_element()`**

```
template<class RandomAccessIterator>
void nth_element(RandomAccessIterator first,
 RandomAccessIterator nth,
 RandomAccessIterator last);
```

The `lower_bound()` function finds the first position in front of which `value` can be inserted without changing the order that points to this position. The first version uses a comparison object `comp` to determine the order.

## upper bound()

```
template<class ForwardIterator, class T>
ForwardIterator upper_bound(ForwardIterator first,
 ForwardIterator last,
 const T& value)
```

```
template<class ForwardIterator, class T,
 ForwardIterator upper_bound(ForwardIterator
const T& value, Compare comp);
```

The `upper_bound()` function finds the last position in the range `[first, last)` from the front of which `value` can be inserted without changing the relative order of the elements that points to this position. The first version uses the `<` comparison object `comp` to determine the order.

```
equal range()
```

```
template<class ForwardIterator, class T>
pair<ForwardIterator, ForwardIterator> eq
 ForwardIterator first, ForwardIterator second,
```

## Note

Recall that if  $<$  is used for ordering, the values  $a$  and  $b$  are `false`. For ordinary numbers, equivalent structures sorted on the basis of just one member function where a new value can be inserted and sorted. The comparison object `comp` is used for ordering, equivalent to `comp(b, a)` are `false`. (This is a generalization of the fact that  $a < b$  is equivalent to  $b > a$ .)

## Merging

The merging functions assume that ranges are sorted.

### `merge()`

```
template<class InputIterator1, class InputIterator2,
 class OutputIterator>
OutputIterator merge(InputIterator1 first1, InputIterator2 first2,
 OutputIterator result)

template<class InputIterator1, class InputIterator2,
 class OutputIterator,
 class Compare>
```

and [middle, last)—into a single sorted sequence. Elements from the first range precede equivalent elements from the second range. The first version uses <, and the second version uses <= to determine the order.

## Working with Sets

Set operations work with all sorted sequences. Multisets are containers that hold more than one instance of a value. They are ordered and unique. A union of two multisets contains the largest number of instances of an element, and an intersection contains the smallest number of instances. For example, suppose Multiset A contains the string "apple" four times. Multiset B contains the same string four times. Then the intersection of A and B contains four instances of "apple", and the intersection with Multiset C contains the same string four times.

### **includes()**

```
template<class InputIterator1, class InputIterator2>
bool includes(InputIterator1 first1, InputIterator1 last1,
 InputIterator2 first2, InputIterator2 last2)
```

```
template<class InputIterator1, class InputIterator2>
bool includes(InputIterator1 first1, InputIterator1 last1,
 InputIterator2 first2, InputIterator2 last2)
```

the comparison object comp to determine the

## **set\_intersection()**

```
template<class InputIterator1, class InputIterator2,
 class OutputIterator>
OutputIterator set_intersection(InputIterator1 first1,
 InputIterator1 last1, InputIterator2 first2,
 InputIterator2 last2, OutputIterator result)

template<class InputIterator1, class InputIterator2, class OutputIterator, class Compare>
OutputIterator set_intersection(InputIterator1 first1, InputIterator1 last1,
 InputIterator2 first2, InputIterator2 last2, OutputIterator result,
 Compare comp);
```

The `set_intersection()` function constructs the intersection of the two input ranges `[first1, last1)` and `[first2, last2)` and stores the result in the range `[result, result + (last1 - first1) + (last2 - first2) - 1)`. The resulting range should not be modified. The function returns a past-the-end iterator for the intersection. The first version of the function returns a past-the-end iterator for the intersection. The second version uses the comparison object `comp` to determine the

```

OutputIterator set_symmetric_difference(
 InputIterator1 first1, InputIterator1 last1,
 InputIterator2 first2, InputIterator2 last2,
 OutputIterator result);

template<class InputIterator1, class InputIterator2,
 class OutputIterator, class Compare = std::less<typename InputIterator1::value_type> >
OutputIterator set_symmetric_difference(
 InputIterator1 first1, InputIterator1 last1,
 InputIterator2 first2, InputIterator2 last2,
 OutputIterator result,
 Compare comp)

```

The `set_symmetric_difference()` function computes the symmetric difference between the ranges `[first1, last1)` and `[first2, last2)` and stores the result in the range `[result, result + (last1 - first1 + last2 - first2))`. The result is the set of elements found in either of the original ranges. The function returns a past-the-end iterator. The symmetric difference is the set that contains elements found in either of the two sets but not in the second and the elements found in the first set but not in the second. The symmetric difference is as the difference between the union and the intersection of the two sets. The second version uses the comparison object `comp` to compare the elements.



```
void push_heap(RandomAccessIterator first,
 Compare comp);
```

The `push_heap()` function assumes that the range `[first, last)` is a valid heap. It adds the value at location `last - 1` (that is, the last element) to the heap, making `[first, last)` a valid heap. The first version uses `<` to determine the ordering, and the second version uses `comp`.

## **pop\_heap()**

```
template<class RandomAccessIterator>
void pop_heap(RandomAccessIterator first,
 RandomAccessIterator last,
 Compare comp);
```

The `pop_heap()` function assumes that the range `[first, last)` is a valid heap. It removes the value at location `last - 1` with the value at location `first` and then re-heapifies the range `[first, last - 1)` to be a valid heap. The first version uses `<` to determine the ordering, and the second version uses the `comp` comparison object.

comp to determine the order.

## **is\_heap\_until() (C++11)**

```
template<class RandomAccessIterator>
RandomAccessIterator is_heap_until(RandomAccessIterator first,
 RandomAccessIterator last,
 Compare comp) {
 return first;
}

template<class RandomAccessIterator, class Compare>
RandomAccessIterator is_heap_until(
 RandomAccessIterator first,
 RandomAccessIterator last,
 Compare comp);
```

The `is_heap_until()` function returns last element of the heap and the element following it. Otherwise, it returns the last element of the sequence. The first version uses `<`, and the second version uses `comp` to determine the order.

## **Finding Minimum and Maximum Values**

The `min` and `max` functions return the minimum and maximum values of pairs of values and of sequences of values.

```
template<class T, class Compare>
const T& max(const T& a, const T& b, Compare comp)
```

These versions of the `max()` function return the first value if the values are equivalent, they return the first value. The first version uses the `less` comparison object, and the second version uses the `comp` comparison object.

```
template<class T> T max(initializer_list<T> t)
```

```
template<class T, class Compare>
T max(initializer_list<T> t, Compare comp)
```

These versions of `max()` function (C++11) return the first value in the list `t`. If the two or more values are equivalent, they return the first item having that value. The first version uses the `less` comparison object, and the second version uses the `comp` comparison object.

## **minmax() (C++11)**

```
template<class T>
pair<const T&,const T&> minmax(const T& a, const T& b)
```

```
template<class T, class Compare>
pair<const T&,const T&> minmax(const T& a, const T& b,
 Compare comp)
```

The `min_element()` function returns the first element in the range `[first, last)` such that no element in the range is less than `*it`. The first version uses the `<` operator for ordering, and the second version uses the compare function `comp`.

### **`max_element()`**

```
template<class ForwardIterator>
ForwardIterator max_element(ForwardIterator first,
 ForwardIterator last) {
 return min_element(first, last, greater());
}

template<class ForwardIterator, class Compare>
ForwardIterator max_element(ForwardIterator first,
 ForwardIterator last,
 Compare comp);
```

The `max_element()` function returns the first element in the range `[first, last)` such that there is no element that `*it` is less than. The first version uses the `<` operator for ordering, and the second version uses the compare function `comp`.

### **`minmax_element()` (C++11)**

```
template<class ForwardIterator>
pair<ForwardIterator, ForwardIterator>
minmax_element(ForwardIterator first,
 ForwardIterator last) {
 return minmax_element(first, last, greater());
}

template<class ForwardIterator, class Compare>
pair<ForwardIterator, ForwardIterator>
minmax_element(ForwardIterator first,
 ForwardIterator last,
 Compare comp);
```

to `*first2`. If `*first1` is less than `*first2`, then `*first1` is less than `*first2`, the function returns `false`. If `*first1` is greater than `*first2`, the function returns `true`. If `*first1` is equal to `*first2`, the function returns to the next element in each sequence. This process continues until the elements are not equivalent or until the end of one of the sequences is reached. If the sequences are equivalent until the end of one is reached, the function returns `true`. If the sequences are not equivalent until the end of one is reached, the function returns `false`. The first version of the function uses `<` for comparison. The second version uses the `comp` comparison object. The third version uses the `lex` comparison object. The fourth version uses the `lex` comparison object of an alphabetic comparison.

## Working with Permutations

A *permutation* of a sequence is a reordering of the elements of the sequence. A sequence of three elements has six possible orderings because there are three choices for the first element, two choices for the second element, and one choice for the third element. For example, the permutations of the sequence 1, 2, 3 are as follows:

123 132 213 232 312 321

In general, a sequence of  $n$  elements has  $n!$  permutations.

The permutation functions assume that the elements of the sequence are arranged in lexicographic order, as in the previous example.

```
template<class BidirectionalIterator>
bool prev_permutation(BidirectionalIterator
 BidirectionalIterator
```

```
template<class BidirectionalIterator, class
bool prev_permutation(BidirectionalIterator
 BidirectionalIterator
```

The `previous_permutation()` function transforms the permutation (the last permutation in lexicographic order) to the previous permutation in lexicographic order. If it doesn't exist, the function returns `true`. If it does exist, the function returns `false`. The first version of the function uses the `<` comparison function object for ordering elements. The second version uses the `<` ordering, and the second version uses the `<` ordering.

## Numeric Operations

Table G.16 summarizes the numeric operations. The first column shows the header file. Arguments are not shown, and only the first version of the function is shown. Each function has a version that uses `<` for ordering. The second column shows the comparison function object for ordering elements. A checkmark indicates that the function follows the table. Thus, you can scan the table to find the function you need, then look up the details if you find the function.

```
template <class InputIterator, class T, class BinaryOperation>
T accumulate(InputIterator first, InputIterator last, T init,
 BinaryOperation binary_op);
```

The `accumulate()` function initializes a value `acc` to `init` and then computes `acc = acc + *i` (first version) or `acc = binary_op(acc, *i)` (second version) for each iterator `i` in the range `[first, last)`, in order.

## **`inner_product()`**

```
template <class InputIterator1, class InputIterator2, class T,
 class BinaryOperation>
T inner_product(InputIterator1 first1, InputIterator1 last1,
 InputIterator2 first2, T init, BinaryOperation binary_op);
```

```
template <class InputIterator1, class InputIterator2, class T,
 class BinaryOperation1, class BinaryOperation2>
T inner_product(InputIterator1 first1, InputIterator1 last1,
 InputIterator2 first2, T init, BinaryOperation1 binary_op1,
 BinaryOperation2 binary_op2);
```

The `inner_product()` function initializes a value `acc` to `init` and then computes `acc = *i * *j` (first version) or `acc = binary_op1(acc, *i, *j)` (second version) for each iterator `i` in the range `[first1, last1)`, in order.

copied over the original sequence, if desired.

## **adjacent\_difference()**

```
template <class InputIterator, class OutputIterator>
OutputIterator adjacent_difference(InputIterator first, InputIterator last,
 OutputIterator result)
```

```
template <class InputIterator, class OutputIterator, class BinaryOperation>
OutputIterator adjacent_difference(InputIterator first, InputIterator last,
 OutputIterator result, BinaryOperation binary_op)
```

The `adjacent_difference()` function assigns `*result = *first`. Subsequent locations in the target range (`result + 1`, `result + 2`, ..., `result + last - first`) are assigned the value of `binary_op(*first, *(first + 1))`, `binary_op(*(first + 1), *first)` (second argument is the past-the-end iterator for the result. The algorithm returns the past-the-end iterator for the result. The algorithm allows the result to be copied over the original sequence, if desired.





## Selected Readings

- Becker, Pete. *The C++ Standard Library*. Addison-Wesley, 2007.

This book discusses the first C++ Library, an optional library for C++98, but more for C++11. Topics include the unordered container library, the random number generator, and the atomic library.

- Booch, Grady, Robert A. Maksimchuk. *Object-Oriented Analysis and Design, Third Edition*. Addison-Wesley, 2007.

This book presents the concepts behind Object-Oriented Design. It discusses OOP methods, and presents sample code.

- Cline, Marshall, Greg Lomow, and Mike Lomow. *C++ Programming*. Reading, MA: Addison-Wesley, 1998.

This book addresses a great number of C++ language features.

- Josuttis, Nicolai M. *The C++ Standard Library*. Reading, MA: Addison-Wesley, 1999.

This book describes the Standard Template Library features, such as complex number, string, and vector.

- Musser, David R, Gillmer J. Derge, and  
*C++ Programming with the Standard Temp*  
Addison-Wesley, 2001.

The STL merits a complete book to de  
such a book.

- Stroustrup, Bjarne. *The C++ Programmi*  
Addison-Wesley, 1997.

Stroustrup created C++, so this is the d  
digested if you already have some know  
guage, but it also provides many exampl  
OOP methodology. Successive editions  
and this edition includes a discussion of  
and strings.

- Stroustrup, Bjarne. *The Design and Evolu*  
1994.

If you're interested in learning how C+  
this book.

- Vandevoorde, David and Nicolai M. Jp  
Reading, MA: Addison-Wesley, 2003.

A lot can be said about templates, as thi





# Use Alternatives for Some Directives

The C/C++ preprocessor provides an array of directives that you can use those directives that are designed to manage program files. For example, using directives as a substitute for code. For example, you can use the `#include` directive as a component for managing program files. Other directives, such as `#if`, `#ifdef`, and `#ifndef`, you control whether particular blocks of code are compiled. You can also use `#pragma` to control compiler-specific compilation options, such as optimization, warnings, and so on. You should exercise caution, however, when using these directives.

## Use `const` Instead of `#define`

Symbolic constants make code more readable and maintainable. They indicate its meaning, and if you need to change a value, you only need to change it once, in the definition, and then recompile. C++ provides a way to define names for a constant:

```
#define MAX_LENGTH 100
```

The preprocessor then does a text substitution of `MAX_LENGTH` with `100` prior to compilation.

The C++ approach is to apply the `const` keyword to a variable. For example, `const int MAX_LENGTH = 100;`

The preprocessor will replace

```
int n;
```

with

```
int 5;
```

and induce a compilation error. The `dz` declaration is a constant variable. Also if necessary, `fizzle()` can use the constant as `::dz`.

C++ has borrowed the `const` keyword from C. For example, the C++ version has internal linkage for the default external linkage used by variables. A file in a program using a `const` needs that `const` definition. It sounds like extra work, but, in fact, it makes life easier. `const` definitions in a header file used by various files cause a linker error for external linkage but not for internal linkage. If defined in the file that uses it (being in a header file), you can use `const` values as array size

```
const int MAX_LENGTH = 100;
```

```
...
```

```
double loads[MAX_LENGTH];
```

```
enum {LEVEL1 = 1, LEVEL2 = 2, LEVEL3 = 4,
```

## Use inline Instead of #define

The traditional C way to create the near-equivalent  
#define macro definition:

```
#define Cube(X) X*X*X
```

This leads the preprocessor to do text substitution  
substituting argument to Cube():

```
y = Cube(x); // replaced with y = x*x*x
y = Cube(x + z++); // replaced with x + z++ * x + z++ * x + z++
```

Because the preprocessor uses text substitution,  
using such macros can lead to unexpected and  
undesired results, especially when  
by using lots of parentheses in the macro to ensure

```
#define Cube(X) ((X)*(X)*(X))
```

Even this, however, doesn't deal with cases like



serves as its own prototype.

You should use `const` in function prototypes. Similarly, you should use `const` with pointer parameters for data that is not to be altered. Not only does this prevent changing data, it also makes a function more general. A pointer or reference can process both `const` and non-`const` data. It fails to use `const` with a pointer or reference

## Use Type Casts

One of Stroustrup's pet peeves about C is its need for type casts. Type casts are often necessary, but the standard type cast is `(type) expression`. Consider the following code:

```
struct Doof
{
 double feeb;
 double steeb;
 char sgif[10];
};

Doof leam;

short * ps = (short *) & leam; // old syntax
int * pi = int * (&leam); // new syntax
```

If you've been using `malloc()` and `free()`, you should use `new` and `delete` instead. If you've been using `setjmp()` and `longjmp()`, you should use `try`, `throw`, and `catch` instead. You should try to use `true` and `false`.

## Use the New Header Organization

The C++ Standard specifies new names for the standard library headers. This is called “Setting Out to C++.” If you've been using the old names, you should move over to using the new-style names. This is not a breaking change, because older versions sometimes add new features. For example, C++11 added support for wide-character input and output. It also added `boolalpha` and `fixed` (as described in Chapter 10). The new headers offer a simpler interface than using `setf()` or `unsetf()` for formatting options. If you do use `setf()`, you should use `ios::fixed` for specifying constants; that is, you should use `ios::fixed` instead of `ios::fixed`. Also the new header files incorporate namespacing.

## Use Namespaces

Namespaces help organize identifiers used in a program. The standard library uses namespaces. Because the standard library, as implemented in the C++ Standard Library,

expression that uses the operator:

```
cout << std::fixed << x << endl; //using
```

This could get wearisome, but you could do this in a header file:

```
// mynames - a header file
using std::cin; // a u
using std::cout;
using std::endl;
```

Going a step further, you could collect usi

```
// mynames - a header file
#include <iostream>
```

```
namespace io
{
 using std::cin;
 using std::cout;
 using std::endl;
}
```

```
namespace formats
{
```

assignment operator to assign a string to a character array. You can't use `strcpy()` or `strncpy()`. You can't use the `sizeof` operator on strings; instead, you must remember to use `strlen()`. If you use the `>` operator, you don't get a syntax error; instead, you get a warning about comparing pointers instead of string contents.)

The `string` class (see Chapter 16, “The Standard Library,” and Appendix F, “The `string` Template Library”) uses dynamic memory management to represent strings. Assignment operators and the `+` operator (for concatenation) are all defined. It provides automatic memory management so that you normally don't have to worry about entering a string that either overruns an array or is not properly terminated.

The `string` class provides many convenient methods for manipulating one `string` object to another, but you can also convert a C-style string value to a `string` object. For functions that need to return a string, you can use the `c_str()` method to return a suitable C-style string.

Not only does the `string` class provide a convenient interface for string-related tasks, such as finding substrings and concatenating strings, but it is also compatible with the Standard Template Library (STL) `string` objects.

your part. And the iterator concept used to in  
algorithms aren't limited to being used with S  
applied to traditional arrays, too.

final compilation.

3. It makes definitions made in the `std` namespace available.

4. `cout << "Hello, world\n";`

or

`cout << "Hello, world" << endl;`

5. `int cheeses;`

6. `cheeses = 32;`

7. `cin >> cheeses;`

8. `cout << "We have " << cheeses << endl;`

9. The function `froop()` expects to be called with a `double`, and that the function will return a `double`. It can be used as follows:

```
int gval = froop(3.14159);
```

The function `rattle()` has no return value. It can be used as follows:

```
rattle(37);
```

Note. Don't count on the being large  
system supports universal list-initialization

```
short rbis = {80}; // = is
unsigned int q {42110}; // could
long long ants {3000000000};
```

3. C++ provides no automatic safeguards  
you can use the `climits` header file to
4. The constant `33L` is type `long`, whereas
5. The two statements are not really equivalent  
some systems. Most importantly, the first  
only on a system using the ASCII code,  
other codes. Second, `65` is a type `int` co
6. Here are four ways:

```
char c = 88;
cout << c << endl; // char

cout.put(char(88)); // put()

cout << char(88) << endl; // new-

cout << (char)88 << endl; // old-
```

- 10.
- a. `int`
  - b. `float`
  - c. `char`
  - d. `char32_t`
  - e. `double`

## Answers to Chapter Review

- 1.
- a. `char actors[30];`
  - b. `short betsie[100];`
  - c. `float chuck[13];`
  - d. `long double dipsea[64];`
- 2.
- a. `array<char, 30> actors;`
  - b. `array<short, 100> betsie;`
  - c. `array<float, 13> chuck;`
  - d. `array<long double, 64> dips`



```

 "trout",
 12,
 26.25
 };

```

10. `enum Response {No, Yes, Maybe};`

11. `double * pd = &ted;`  
`cout << *pd << "\n";`

12. `float * pf = treacle; // or = &treacle;`  
`cout << pf[0] << " " << pf[9] << "\n";`  
`// or use *pf and *(pf + 9);`

13. This assumes that the `iostream` and `vector` headers are included, and that there is a `using` directive:

```

unsigned int size;
cout << "Enter a positive integer: ";
cin >> size;
int * dyn = new int [size];
vector<int> dv(size);

```

```

17. #include <string>
 #include <vector>
 #include <array>
 const int Str_num {10}; // or = 10
 ...
 std::vector<std::string> vstr(Str_num);
 std::array<std::string, Str_num> as;

```

## Answers to Chapter Review

1. An entry-condition loop evaluates a test condition before the loop body is executed. If the condition is initially false, the loop body is not executed. An exit-condition loop evaluates a test expression after the loop body is executed. Thus, the loop body is executed once, even if the condition is false. The for and while loops are entry-condition loops. The do-while loop is an exit-condition loop.

2. It would print the following:

```
01234
```

Note that `cout << endl;` is not part of the loop body (it is outside the braces).

to `x`. The second statement is also valid.  
be evaluated as follows:

```
(y = 1) , 024;
```

That is, the left expression sets `y` to 1, and the right expression, which isn't used, is 024, or 20.

9. The `cin >> ch` form skips over spaces, tabs, and newlines. The other two forms read those characters as well.

## Answers to Chapter Review Questions

1. Both versions give the same answers, but Version 1 is more efficient. Consider what happens, for example, when reading a line of input containing spaces, tests whether the character is a space, and so on. In Version 1, the program has already established that `ch` is a character, so it doesn't need to check line. Version 2, in the same situation, skips the first test.
2. Both `++ch` and `ch + 1` have the same memory address. `++ch` prints as a character, while `ch + 1`, because it's an integer, prints as a number.

6.  $(x < 0) ? -x : x$

or

$(x \geq 0) ? x : -x;$

7. `switch (ch)`

```
{
 case 'A': a_grade++;
 break;
 case 'B': b_grade++;
 break;
 case 'C': c_grade++;
 break;
 case 'D': d_grade++;
 break;
 default: f_grade++;
 break;
}
```

8. If you use integer labels and the user types a non-integer, the program hangs because integer input can't be processed. If the user types an integer such as **5**, the program works. Then the default part of the switch can be used.

```
g. char * plot(map *pmap);
```

3. 

```
void set_array(int arr[], int size,
{
 for (int i = 0; i < size; i++)
 arr[i] = value;
}
```
4. 

```
void set_array(int * begin, int * end,
{
 for (int * pt = begin; pt != end; pt++)
 *pt = value;
}
```
5. 

```
double biggest (const double foot[],
{
 double max;
 if (size < 1)
 {
 cout << "Invalid array size\n";
 cout << "Returning a value of 0\n";
 return 0;
 }
}
```

```

 while (*str) // while not null
 {
 if (*str == c1)
 {
 *str = c2;
 count++;
 }
 str++; // advance to next char
 }
 return count;
 }
}

```

9. Because C++ interprets "pizza" as the address of the first element, the address operator yields the value of that first element. If you use "taco", C++ interprets "taco" as the address of the second element. If you use "taco", the value of the element two positions ahead of the first element. In other words, the string constant acts as the address of the first element.
10. To pass it by value, you just pass the string constant. To pass it by reference, use the address operator &glitz. Passing by reference passes the original data, but it takes time and memory but doesn't protect the original data. Passing by value passes a function parameter. Also passing by value.

```

13. typedef void (*p_f1)(applicant *);
 p_f1 p1 = f1;
 typedef const char * (*p_f2)(const a
 p_f2 p2 = f2;
 p_f1 ap[5];
 p_f2 (*pa)[10];

```

## Answers to Chapter Review

- Short, nonrecursive functions that can f  
for inline status.
- void song(const char \* name
  - None. Only prototypes contain th
  - Yes, provided that you retain the c

```
void song(char * name = "O, My
```
- You can use either the string "\"" or th  
The following functions show both me  

```
#include <iostream.h>
void iquote(int n)
```

```
 cout << "Volume = " << crate.volume << endl;
 }
}
```

```
b. void set_volume(box & crate)
{
 crate.volume = crate.height * crate.width * crate.depth;
}
```

5. First, change the prototypes to the following:

```
// function to modify array object
void fill(std::array<double, Season::length> & expenses, Season season);
// function that uses array object
void show(const std::array<double, Season::length> & expenses, Season season);
```

Note that `show()` should use `const` to avoid modifying the array object.  
Next, within `main()`, change the `fill` call to:  
`fill(expenses);`

There's no change to the `show()` call.



- b. You can't use a default for the repeat function because you can't have default values from right to left. You can write:  

```
void repeat(int times, const char * str);
```

```
void repeat(const char * str, int times);
```
- c. You can use function overloading:  

```
int average(int a, int b);
```

```
double average(double x, double y);
```
- d. You can't do this because both versions of repeat would be ambiguous.

- 7. 

```
template<class T>
T max(T t1, T t2) // or T max(const T& t1, const T& t2)
{
 return t1 > t2? t1 : t2;
}
```
- 8. 

```
template<> box max(box b1, box b2)
{
 return b1.volume > b2.volume? b1 : b2;
}
```

3. `#include <iostream>`

```
int main()
```

```
{
```

```
 double x;
```

```
 std::cout << "Enter value: ";
```

```
 while (! (std::cin >> x))
```

```
 {
```

```
 std::cout << "Bad input. Pl
```

```
 std::cin.clear();
```

```
 while (std::cin.get() != '\n'
```

```
 continue;
```

```
 }
```

```
 std::cout << "Value = " << x <<
```

```
 return 0;
```

```
}
```

4. Here is the revised code:

```
#include <iostream>
```

```
int main()
```

```
{
```

```
 using std::cin;
```

```
 using std::cout;
```

```
 using std::endl;
```

4, 1, 2

2

2

4, 1, 2

2

## Answers to Chapter Review

1. A class is a definition of a user-defined data structure to be stored, and it specifies the methods and functions to access and manipulate that data.
2. A class represents the operations you can perform on the interface of class methods; this is abstraction. A class (by default) for data members, meaning that data is hidden from member functions; this is data hiding. Data and member representation and method code, are hidden.
3. A class defines a type, including how it is stored in memory. Another data object, such as that produced by a program, according to the class definition. The relationship between the same as that between a standard type and its data.

```
void deposit(double cash);
void withdraw(double cash);
};
```

6. A class constructor is called when you explicitly call the constructor. A class d
7. These are two possible solutions (note in order to use `strncpy()` or else you class):

```
BankAccount::BankAccount(const char
{
 strncpy(name, client, 39);
 name[39] = '\\0';
 strncpy(acctnum, num, 24);
 acctnum[24] = '\\0';
 balance = bal;
}
```

or

```
BankAccount::BankAccount(const std:
 const std:
{
 name = client;
```

```

 void set_tot() { total_val = sha
public:
 Stock(); // default c
 Stock(const std::string & co, lo
 ~Stock() {} // do-nothin
 void buy(long num, double price)
 void sell(long num, double price)
 void update(double price);
 void show() const;
 const Stock & topval(const Stock
 int numshares() const { return s
 double shareval() const { return
 double totalval() const { return
 const string & co_name() const {
};

```

10. The `this` pointer is available to class methods to invoke the method. Thus, `this` is the address of the object itself.

access public members.

4. Here's a prototype for the class definition methods file:

```
// prototype
friend Stonewt operator*(double mult, const Stonewt &s)

// definition - let constructor do the work
Stonewt operator*(double mult, const Stonewt &s)
{
 return Stonewt(mult * s.pounds)
}
```

5. The following five operators cannot be overloaded with `sizeof`

```
.
.*
::
? :
```

6. These operators must be defined by user-defined functions

1  
ory allocated by new in the constructor deletes such memory, it might end up transferring one such object to another. That's because one object to another copies pointer values and this produces two pointers to the same memory. The constructor that causes initialization to be done on one object to another can produce the same data. The solution is to overload the data, not the pointers.

3. C++ automatically provides the following:
  - A default constructor if you don't define one
  - A copy constructor if you don't define one
  - An assignment operator if you don't define one
  - A default destructor if you don't define one
  - An address operator if you don't define one

The default constructor does nothing, but initializes objects. The default copy constructor uses memberwise assignment. The default address operator returns the address of the object (this pointer).

```

 talents = 0;
 }

 nifty::nifty(const char * s)
 {
 strcpy(personality, s);
 talents = 0;
 }

 ostream & operator<<(ostream & os,
 {
 os << n.personality << '\n';
 os << n.talent << '\n';
 return os;
 }
}

```

Here is another possible solution:

```

#include <iostream>
#include <cstring>
using namespace std;
class nifty
{
private: // optional
 char * personality; // creat

```



```

ostream & operator<<(ostream & os, c
{
 os << n.personality << '\n';
 os << n.talent << '\n';
 return os;
}

```

5.
  - a.
 

```

Golfer nancy; // default constructor
Golfer lulu("Little Lulu"); // Golfer with name
Golfer roy("Roy Hobbs", 12); // Golfer with name and talent
Golfer * par = new Golfer; // default constructor
Golfer next = lulu; // Golfer with name and talent
Golfer hazard = "Weed Thwacker"; // Golfer with name
*par = nancy; // default assignment operator
nancy = "Nancy Putter"; // Golfer with name and talent
// the default constructor

```

Note that some compilers additionally require a default constructor for Statements 5 and 6.

- b. The class should define an assignment operator for addresses.

4. Constructors are called in the order of inheritance. The constructor of the base class is called first. Destructors are called in the reverse order.
5. Yes, every class requires its own constructor. If a class has no members, the constructor can have an empty body.
6. Only the derived-class method is called. The base-class method is called only if the derived class does not have a method with the same name. If you use the scope-resolution operator to call a function, then any functions that will be redefined are called.
7. The derived class should define an assignment operator. Constructors use the `new` or `new []` operator to allocate memory for that class. More generally, the derived class should define an assignment operator if the default assignment is incorrect for the class.
8. Yes, you can assign the address of an object of a derived class to a pointer of a base class. You can assign the address of an object of a base class (downcasting) only by making an explicit cast. It is not safe to use such a pointer.
9. Yes, you can assign an object of a derived class to a pointer of a base class. The data members that are new to the derived class are not accessed, however. The program uses the base-class version of the data members.

```
PublicCorporation::head().
```

14. First, the situation does not fit the *is-a* model. Second, the definition of `area()` in `Circle` is wrong because the two methods have different

## Answers to Chapter Review

1. 

```
class Bear
class PolarBear
class Kitchen
class Home
class Person
class Program
class Person
class HorseAndJockey
class Person,
class Automobile
class Driver
```

```

private:
 enum {MAX = 10}; // constant
 Worker * items[MAX]; // hold items
 int top; // index of top element
public:
 Stack();
 Boolean isempty();
 Boolean isfull();
 Boolean push(const Worker * & item);
 Boolean pop(Worker * & item);
};

```

5. `ArrayTP<string> sa;`  
`StackTP< ArrayTP<double> > stck_arr;`  
`ArrayTP< StackTP<Worker *> > arr_st;`

Listing 14.18 generates four templates:  
`ArrayTP<int,5>`, and `Array< ArrayTP<int,5>`

6. If two lines of inheritance for a class share a common ancestor, the compiler will generate two copies of the ancestor's member functions. To avoid this, the base class should be made a friend of its immediate descendants.

```

 class cuff {
 public:
 void snip(muff &) { ... }
 ...
 };
 class muff {
 friend void cuff::snip(muff &) { ... }
 };
 };

```

2. No. For Class A to have a friend that's a function, the function declaration must precede the A declaration. A friend function would tell A that B is a class, but it wouldn't know B is a friend. Similarly, if B has a friend that's a member function, the member function must precede the B declaration. These two rules are necessary to ensure that the compiler can find the definition of the friend function or member function.
3. The only access to a class is through its public interface. The only thing you can do with a Sauce object is call its public methods. The other members (soy and sugar) are private.
4. Suppose the function f1() calls the function f2(). If f2() throws an exception, it causes program execution to resume at the point in f1() where the function call in function f1(). A throw statement in a function can cause an exception to be thrown, which then propagates up the call stack until it is caught by a catch statement.

```

class RQ1
{
private:
 string st; // a string object
public:
 RQ1() : st("") {}
 RQ1(const char * s) : st(s) {}
 ~RQ1() {}
 // more stuff
};

```

The explicit copy constructor, destructor, and assignment operator are needed because the string object provides dynamic memory management.

2. You can assign one string object to another without dynamic memory management so that you normally do not exceed the capacity of its holder.
3. 

```
#include <string>
#include <cctype>
using namespace std;
void ToUpper(string & str)
```

ordinary arrays as well as with iterators and algorithms, thus providing generality.

9. You can assign one vector object to another, so you can insert items into a vector and then use the `at()` method, you can get automatic memory management, and so on.
10. The two `sort()` functions and the `random_shuffle()` function provide access iterators, whereas a `list` object just provides a `begin()` and `end()` access iterator. The `list` template class `sort()` member function is called “Methods and Functions”) instead of the `sort()` member function. Sorting, but there is no member function `sort()` in the `list` class. If you could copy the list to a vector, shuffle the vector, and copy the vector back to the list.

## Answers to Chapter Review Questions

1. The `iostream` file defines the classes, `istream` and `ostream`, for input and output. These objects manage the input and output streams. The file also creates standard objects (`cin`, `cout`, and `clog`) character equivalents) used to handle the input and output connected to every program.

```
6. //rq17-6.cpp
#include <iostream>
#include <iomanip>

int main()
{
 using namespace std;
 cout << "Enter an integer: ";
 int n;
 cin >> n;
 cout << setw(15) << "base ten" << endl;
 << "base sixteen" << setw(15) << endl;
 cout.setf(ios::showbase); //
 cout << setw(15) << n << hex << endl;
 << oct << setw(15) << n << endl;

 return 0;
}
```



```

cout << "First format:\n";
 cout << setw(30) << name << ": $"
 << setw(10) << hourly << ": $"
 << setw(5) << hours << "\n";
cout << "Second format:\n";
cout.setf(ios::left, ios::adjustleft);
cout << setw(30) << name << ": $"
 << setw(10) << hourly << ": $"
 << setw(5) << hours << "\n";

 return 0;
}

```

8. Here is the output:

```
ct1 = 5; ct2 = 9
```

The first part of the program ignores spaces and tabs, and the second part doesn't. Note that the second part of the program uses the `ios::left` character following the first `q`, and it controls the alignment of the output total.

9. The `ignore()` form filters if the input stream skips only the first 80 characters.

or  $w+1$ .

$r1(w)$  is valid, and the argument is the return value of  $w$ .

In general, if an lvalue is passed to a const reference parameter, the parameter is initialized to the lvalue. If an rvalue reference parameter refers to a temporary,

$r2(w)$  is valid, and the argument  $rx$  refers

$r2(w+1)$  is an error because  $w+1$  is an rvalue

$r2(w)$  is an error because the return value

In general, if an lvalue is passed to a non-const reference parameter, the parameter is initialized to the lvalue. But a non-const reference parameter can't accept an rvalue function argument.

$r3(w)$  is an error because an rvalue reference

$r3(w+1)$  is valid, and  $rx$  refers to the temporary

$r3(w)$  is valid, and  $rx$  refers to the temporary

3.     a. `double & rx`  
       `const double & rx`  
       `const double & rx`

are special because the compiler can do these functions, depending on the context.

5. A move constructor can be used when instead of copying it, but there is no move standard array. If the `Fizzle` class used a then one can transfer ownership by raw pointer.

6. 

```
#include <iostream>
#include <algorithm>
template<typename T>
 void show2(double x, T fp) {std::cout << fp(x) << "\n";}
int main()
{
 show2(18.0, [] (double x){return x;});
 return 0;
}
```





sayings1.cpp, 656

string1.cpp, 653-656

string1.h, 652-653

potential problems, 645

strings, 133-134

structures, 145-146

when to use, 644

**& (bitwise AND operator), 1239-1240**

**~ (bitwise negation operator), 1237**

**| (bitwise OR operator), 1237-1238**

**^ (bitwise XOR operator), 1238**

**{ } (braces), 258**

**[ ] (brackets), 649-651**

**, (comma operator), 214-217**

example, 214-216

precedence, 217

**/\*...\*/ comment notation, 33**

**// comment notation, 27, 32**

**+ (concatenation operator), strings, 133-134**

**?: (conditional operator), 273-274**

**-- (decrement operator), 207-208**

pointers, 210-211

postfixing, 209-210

prefixing, 209-210

- acquire() function, 516**
- actual arguments, 314**
- adaptable binary functions, 1035**
- adaptable binary predicate, 1035**
- adaptable functors, 1032**
- adaptable generators, 1035**
- adaptable predicate, 1035**
- adaptable unary functions, 1035**
- adapters, 1002**
- adding vectors, 590**
- addition operator (+), overloading, 569-572**
- addpntrs.cpp, 167**
- address.cpp, 154**
- addresses**
  - addresses of functions, obtaining, 362
  - of arrays, 170
  - structure addresses, passing, 351-353
- adjacent\_difference() function, 1321-1322**
- adjacent\_find() function, 1287, 1290**
- ADTs (abstract data types), 552-557**
- algorithms, 1035**
  - copying, 1036
  - groups, 1035-1036
  - in-place, 1036
  - properties, 1036-1037

**arguments (functions), 49, 53**

**arith.cpp, 98**

**arithmetic, pointers, 167-172**

**arithmetic operators, 97-99**

associativity, 99-100

division (/), 100-101

functor equivalents, 1031-1032

modulus (%), 101-102

order of precedence, 99-100

overloading

addition operator (+), 569-572

multiplication operator (\*),  
574-578, 600

subtraction operator (-), 574-578

vector class, 599-600

**array notation, 173**

**array objects, 355**

fill function, 357

versus arrays, 188-189

versus vector objects, 188-189

**array template class, 187**

**ArrayDb class, 791**

**arraynew.cpp, 166**

**arrayone.cpp, 117**



- string1.h, 652-653
- potential problems, 645
- strings, 133-134
- structures, 145-146
- when to use, 644
- assignment operators, combination**
  - assignment operators, 211-212**
- assignment statements, 43-44**
- assignments, 1172-1173**
- associative containers, 1018, 1026**
  - methods, 1281-1284
  - multimap, 1023-1025
  - set, 1019-1022
- associativity, arithmetic operators, 99-100**
- associativity of operators, 1231**
  - examples, 1234
  - table of, 1232-1234
  - asterisk (\*), dereferencing operator
    - (\*)155, 159
    - pointers171-172
- at() method, 1259**
- atan() function, 348**
- atan2() function, 348**

dominance, 828-829

methods, 818-828

**base-class functions, 777**

**begin() method, 981-984, 1251-1252, 1275**

**best matches, 432-434**

**bidirectional iterators, 998**

**Big Endian, 1218**

**bigstep.cpp, 205**

**binary files, 1127-1133**

**binary functions, 1027-1030**

**binary numbers, 1217**

hexadecimal equivalents, 1217-1218

**binary operators, 601, 1234**

**binary predicates, 1027, 1030**

**binary search operations**

binary\_search() function, 1304, 1310

equal\_range() function, 1304, 1309

lower\_bound() function, 1304, 1309

upper\_bound() function, 1304, 1309

**binary searching**

binary\_search() function, 1304, 1310

equal\_range() function, 1304, 1309

lower\_bound() function, 1304, 1309

upper\_bound() function, 1304, 1309

- flushing, 1063
- Build All option (compiler), 24**
- buildstr() function, 341**
- buy() function, 516**
- bytes, 69**

## C

---

### **C language**

- ANSI C, 17
- classic C, 17
- development history, 11
- programming philosophy, 11-13

**C++, Macintosh, 25**

**C++ FAQ Lite, 1325**

**C++ FAQs, Second Edition, 1323**

***The C++ Programming Language, Third Edition, 1324***

***The C++ Standard Library: A Tutorial and Reference, 1323-1324***

***C++ Templates: The Complete Guide, 1324***

### **C++11**

- arrays, initializing, 120
- auto declarations, 109

- chartype.cpp, 81**
- CHAR\_BIT constant, 72**
- CHAR\_MAX constant, 72**
- char\_type type, 1250**
- check\_it.cpp, 1096**
- cheers() function, 307-309**
- choices.cpp, 188**
- choosing integer types, 76-77**
- cin, cin.get() function, 235-237, 241-244**
- cin object, 1067, 1093-1095**
  - get() function, 128-130**
  - getline() function, 126-127**
  - loops, 234-235**
  - operator overloading, 1095-1097**
  - stream states, 1097-1098**
    - effects, 1100-1102**
    - exceptions, 1099-1100**
    - setting, 1098**
- cin statement, 46**
- cin.get() function, 235-237, 241-244, 317**
- cin.get() member function, 1103-1105**
- cin.get(ch) function, 317**
- cinexcp.cpp, 1099**
- cinfish.cpp, 283-284**

- calling, 526-527
- conversion, 769-770
- copy constructors, 639-644, 767
- declaring, 525-526
- default constructors, 527-528,  
638-639, 766-767
- defining, 525-526
- delegating, 1180-1181
- inheriting, 1181-1183
- new operator, 659-661, 677-678
- converting class type, 677
- Customer, 690-691, 694
- data hiding, 511-513, 523
- data types, 507-508
- declarations, 509-511, 522
- defaulted and deleted methods,  
1179-1180
- defined, 36, 47, 508
- defining, 47
- definition of, 13
- derived classes, 405
- destructors, 524, 528-529,  
538-539, 768
- encapsulation, 512, 523

- base-class methods, accessing, 800-801
- base-class objects, accessing, 801
- Student class example, 798
- protected classes, 745-746, 775
- Queue
  - class declaration, 691-694
  - design, 679
  - implementation, 680-682
  - methods, 682-690
  - public interface, 679-680
- Sales
  - sales.cpp, 924
  - sales.h, 922
  - use\_sales.cpp, 925-927
- sample program, 518-520
- special member functions, 1178-1179
- Stack, 831-836
  - pointers, 837-843
- static class members, 628-637
- stdexcept exception classes, 918-920
- Stock, 511
- streambuf, 1065
- string, 131-133, 353-354, 647, 952, 960, 965-966, 1333

- 605-606
- removing ranges of, 982
- shuffling elements in, 987
- sorting, 987
- state members, 597-599
- vect1.cpp example, 980-981
- vect2.cpp sample program, 984-986
- vect3.cpp sample program, 988-991
- vector template class, 186-187
- virtual methods
  - final, 1183-1184
  - override, 1183-1184
- Worker, 810-814

## **classic C, 17**

**classifying data types, 97**

**clear() method, 1258, 1278, 1283**

**clear() stream state method, 1098-1102**

**clearing bits, 1086**

**client files, creating, 533-536**

**client/server model, 520**

**climits header file, 71-73**

**clock() function, 229**

**clog object, 1067**

**close() method, 292**

firstref.cpp, 383  
floatnum.cpp, 95  
fltadd.cpp, 96  
forloop.cpp, 196  
formore.cpp, 203–204  
forstr1.cpp, 206  
forstr2.cpp, 215  
fowl.cpp, 973  
frnd2tmp.cpp, 860–861  
fun ptr.cpp, 364, 368  
funadap.cpp, 1034–1035  
function overloading, leftover.cpp, 418  
functions, tempover.cpp, 434–437  
functions, arguments, twoarg.cpp, 316  
functions, recursion, recur.cpp, 359  
functor.cpp, 1028  
funtemp.cpp, 420  
getinfo.cpp, 45  
get\_fun.cpp, 1107  
hangman.cpp, 962–965  
hexoct2.cpp, 79  
if/cpp, 255  
ifelse.cpp, 257  
ifelseif.cpp, 259  
ilist.cpp, 1053



setf.cpp, 1085  
setf2.cpp, 1088  
setops.cpp, 1021-1022  
showpt.cpp, 1084  
somedefs.h, 1192  
sqrt.cpp, 51  
stack.cpp, 554-555  
stack.h, 553-554  
stacker.cpp, 555-557  
stacktem.cpp, 835-836  
stacktp.h, 833-834  
static.cpp, 470-471  
stcktp1.cpp, 841-842  
stcktp1.h, 839-840  
stdmove.cpp, 1174  
stock00.h, 510  
stock1.cpp, 531  
stock1.h, 530  
stock2.cpp, 543  
stock2.h, 543  
stocks.cpp, class member functions, 515  
stone1.cpp, 615  
stone.cpp, 610  
stonewt.cpp, 608

use\_stui.cpp, 804–805  
use\_tv.cpp, 882  
usestok0.cpp, 519  
valvect.cpp, 1048  
variadic1.cpp, 1199  
vect.cpp, 593  
vect.h, 591  
vect1.cpp, 980  
vect2.cpp, 984–985  
vect3.cpp, 988  
vegnews.cpp, 634  
vslice.cpp, 1049–1050  
waiting.cpp, 229  
while.cpp, 225  
width.cpp, 1080  
Worker0.cpp, 811–812  
Worker0.h, 810–811  
workermi.cpp, 823–825  
workermi.h, 821–822  
workmi.cpp, 826–827  
worktest.cpp, 813  
write.cpp, 1073–1074

**code style, 40**

**colon (), scope-resolution operator (::), 514**

**combination assignment operators, 211–212**

- functions, 1039-1041
- properties, 1008-1010
- sequence requirements, 1011-1012
- functors, 1027-1030
- iterators, models, 1000-1001
- concurrent programming, 1202-1203**
- condit.cpp, 273**
- conditional operator (?::), 273-274**
- const, reference returns, 400**
- const keyword, 90-92, 473-474, 771-772**
  - arrays, 327-328
  - pointers, 334-336
  - reference variables, 401
  - temporary variables, 392-394
- const member functions, 537**
- const modifier as alternative to #define, 1327-1329**
- const objects, returning references to, 662-665**
- constant time, 1009**
- constant time complexity, 1009-1010**
- constants, 78-80. See also strings**
  - char constants. *See* char data type
  - const keyword, 90-92

C++11, 1271-1273

vectors, methods, 1278-1280

**containment, 785**

**continue statement, 280-282**

**conversion constructors, 769-770**

**conversion operators, explicit, 1159-1160**

**convert.cpp, 57**

**converting**

class type, 677

rectangular coordinates to polar  
coordinates, 348-351

to standard C++, 1327

auto\_ptr template, 1333

C++ features, 1331

const instead of #define,  
1327-1329

function prototypes, 1330

header files, 1331

inline instead of #define,  
1329-1330

namespaces, 1331-1333

STL (Standard Template  
Library), 1334

string class, 1333

type casts, 1330-1331

**data objects, pointers, 161**

**data types, 507-508**

- ADTs (abstract data types), 552-557

- aliases, creating, 230

- bool, 90

- classifying, 97

- compound types, 116

- double, 50

- floating-point numbers, 92

  - advantages/disadvantages, 96-97

  - constants, 96

  - decimal-point notation, 92

  - double, 94-96

  - E notation, 92-93

  - float, 94-96

  - long double, 94-96

- integers, 68

  - char, 80-89

  - choosing integer types, 76-77

  - climits header file, 71-73

  - constants, 78-80, 90-92

  - initializing, 73

  - int, 68-70

  - long, 68-70

- defining declarations, 463
- definitions, 463
- delegating constructors, 1180-1181
- delete operator, 163-164, 180-183, 400, 454, 476-477, 668
- delete.cpp, 181
- deleted methods, classes, 1179-1180
- deque class templates, 1013
- deque containers, 1013
- deques, methods, 1278-1280
- dequeue() method, 689
- dereferencing (\*) operator, 155-159
- dereferencing operators, 1242-1246
- derived classes, 405
  - constructors, 713-715
  - creating, 711-712
  - header files, 716
  - method definitions, 716
  - objects, creating, 717-718
  - relationships with base classes, 718-720
- derived types, 116
- design
  - bottom-up, 13
  - top-down, 12

enumerators, 150-151

scoped, 1158

C++11, 551-552

value ranges, 153

values, setting, 152

**enumerators, 150-151**

as labels, 278-280

**EOF (end-of-file) conditions, 237-241**

**eof() function, 238**

**eof() method, 296**

**eof() stream state methods, 1097-1102**

**eofbit stream state, 1097-1102**

**equal sign (=)**

assignment operator (=), 644,

767-768, 772-775

custom definitions, 645-646

enumerator values, setting, 152

overloading, 652-658

potential problems, 645

strings, 133-134

structures, 145-146

when to use, 644

equality operator (==), compared to

assignment operator, 218-220

**equal() function, 1288-1292**

- declaration-statement expressions, 202
- logical AND (&&), 262
  - alternative representations, 270
  - example, 263-265
  - precedence, 269-270
  - ranges, 265-267
- logical NOT (!), 267-269
  - alternative representations, 270
  - precedence, 269
- logical OR (| |), 260-262
  - alternative representations, 270
  - example, 261-262
  - precedence, 269
- relational operators, 217-218, 220
  - C-style strings, comparing, 220-223
  - equality operator (==), 218-220
  - string class strings, comparing, 223-224
  - table of, 217
- sequence points, 208-209
- side effects, 201, 208-209
- type conversions, 105-106
- extern keyword, 467-472**
  - functions, 474



- `find_first_of()` method, 961, 1262
- `find_if()` function, 1287-1289
- `find_last_not_of()` method, 1263
- `find_last_of()` method, 961, 1262
- `firstref.cpp`, 383
- fixed manipulator, 1091
- flags, 1084
- flags, setting, 1083
- float data type, 94-96
- floating points, display precision, 1082-1087, 1090
- floating-point data types, default behavior, 1076
- floating-point numbers, 92
  - advantages/disadvantages, 96-97
  - constants, 96
  - decimal-point notation, 92
  - double data type, 94-96
  - E notation, 92-93
  - float data type, 94-96
  - long double data type, 94-96
- `floatnum.cpp`, 95
- `fltadd.cpp`, 96

**free store, 454**

**free store (memory), 182-183**

**freeing memory, delete operator, 163-164**

**friend classes, 578-580, 877-888**

base-class friends, accessing, 801-804

compared to class member

functions, 886

templates, 858

bound template friend functions,  
861-864

non-template friend functions,  
858-861

unbound template friend functions,  
864-865

Tv class example, 878-879, 883

tv.cpp, 880-882

tv.h, 879-880

tvfm.h, 885-886

use\_tv.cpp, 882

**friend functions, 578-580**

creating, 579-580

type conversion, 618-621

**friend keyword, 579-580**

- private, 513
- properties, 777-778
- public, 513
- qualified names, 514
- this pointer, 539-546
- unqualified names, 514
- compared to container methods, 1039-1041
- conversion functions, 677
- defining, 306-309
- definitions, 29
- formatted input, 1094
- friend functions, 578-580
  - creating, 579-580
  - type conversion, 618-621
- function prototypes, 1330
- headers, 29-31
- inline functions, 379-382
  - compared to macros, 382
  - square(), 381-382
- input, unformatted, 1102
- lambda functions, 1184
- language linking, 475-476

**generators, 1027**

**generic programming, 14, 419, 951,  
978, 992**

- associative containers, 1018-1026
  - multimap, 1023-1025
  - set, 1019-1022

- container concepts, 1007
  - container methods compared to functions, 1039-1041
  - properties, 1008-1010
  - sequence requirements, 1011-1012

- container types
  - deque, 1013
  - list, 1014-1017
  - priority\_queue, 1017-1018
  - queue, 1017
  - stack, 1018
  - vector, 1012-1013

- iterators, 992-997
  - back insert, 1005-1007
  - bidirectional, 998
  - concepts, 1000-1001
  - copy() function, 1001-1002
  - forward, 998
  - front insert, 1005-1007

**hexoct2.cpp, 79**  
**hierarchy, iterators, 999-1000**  
**high-level languages, 11**  
**history of C++, 10-15**

    C language

        development history, 11

        programming philosophy, 11-13

    generic programming, 14

    OOP, 13-14

**hmean() function, 898-905**

|

---

**I/O (input/output), 1062, 1270**

    buffers, 1063-1067

    redirecting, 1067-1068

    streams, 1063-1067

    text files, 287-288

        reading, 292-298

        writing to, 288-292

**identifiers, special meanings, 1223**

**IDEs (integrated development  
environments), 19**

- virtual base classes, 815-829
- Worker class example, 810-814
- polymorphic public inheritance, 722-723
  - base-class functions, 777
  - Brass class declaration, 723-726
  - Brass class implementation, 727-731
  - Brass class objects, 732-733
  - BrassPlus class declaration, 723-726
  - BrassPlus class implementation, 727-731
  - BrassPlus class objects, 732-733
  - constructors, 742
  - dynamic binding, 737-740
  - pointer compatibility, 737-739
  - reference type compatibility, 737-739
  - static binding, 737-740
  - virtual destructors, 737, 742-743, 776
  - virtual functions, 734-736, 739-745, 775-776

- instr1.cpp, 125**
- instr2.cpp, 127**
- instr3.cpp, 129**
- int data type, 68-70**
- int main() function header, 30-31**
- integer values, displaying with cout, 44-45**
- integers, 68**
  - bool, 90
  - char, 80-87
    - escape sequences, 84-87
    - signed char, 88-89
    - universal character names, 87-88
    - unsigned char, 88-89
    - wchar\_t, 89
  - choosing integer types, 76-77
  - climits header file, 71-73
  - constants, 78-80
    - const keyword, 90-92
    - symbolic names, 90-92
  - initializing, 73
  - int, 68-70
  - long, 68-70
  - pointers, 160
  - short, 68-70

**jump.cpp, 280-281**

**K&R (Kernighan and Ritchie) C standard, 17**  
**keywords, 56. See also statements**

auto, 472

catch, 900

class, 831

const, 90-92, 473-474, 771-772,  
1327-1329

arrays, 327-328

pointers, 334-336

reference variables, 401

temporary variables, 392-394

decltype, 439

explicit, 610

extern, 467-472

functions, 474

friend, 579-580

implicit, 610

inline, 517, 1329-1330

mutable, 472-473

namespace, 483-486



- linked lists, 680
- linking multiple libraries, 453
- Linux, g++ compiler, 22
- list class templates, 1014-1017
  - member functions, 1014-1016
- list containers, 1014-1017
  - member functions, 1014-1016
- list initialization, C++11, 537
- list.cpp, 1015-1016
- listrmv.cpp, 1039-1040
- lists
  - linked lists, 680
  - methods, 1278-1280
- literal operators, 1204
- Little Endian, 1218
- local scope, 454
  - variables, 455-457
- local variables, 314-315
  - compared to global variables, 467
- logical AND operator (&&), 262
  - alternative representations, 270
  - example, 263-265
  - precedence, 269-270
  - ranges, 265-267

---

machine language, definition of, 18  
 Macintosh, C++, 25  
 macros, compared to inline functions, 382  
 magval() method, 602  
 main() function, 29-30  
     calling, 30  
     importance of, 32  
     int main() header, 30-31  
 make\_heap() function, 1305, 1314  
 malloc() function, 160  
 mangling names, 418  
 manip.cpp, 1079  
 manipulators, 38, 1090-1091  
     endl, 37-38  
     iomanip header file, 1091  
     number base display, 1078-1079  
 mantissas, 93  
 manyfrnd.cpp, 865  
 mapped\_type type, 1281, 1284  
 maps, methods, 1281-1284  
 math operators, 97. *See also* arithmetic operators  
 max() function, 1305, 1316

- unwinding, 909-914
- static storage, 183
- storage class specifiers, 472-473
- storage duration, 453-454
  - automatic variables, 455-459
  - scope and linkage, 454
  - static variables, 459-463, 466-472
- storage methods, 182
- memory allocation, dynamic (auto\_ptr class), 969, 973-975**
- memory leaks, 163**
- merge() function, 1305, 1310-1311**
- merge() method, 1016-1017, 1280**
- merging**
  - inplace\_merge() function, 1305, 1311
  - merge() function, 1305, 1310-1311
- methods. See also specific methods**
  - base-class methods, accessing, 800-801
  - defaulted and deleted methods, classes, 1179-1180
  - end(), 984
  - inheritance, multiple, 826
  - insert(), 1015-1016
  - STL, 1161
  - virtual base classes, 818-828

- replace() function, 1294, 1298, 1302
- replace\_copy() function, 1294, 1298
- replace\_copy\_if() function, 1294, 1298
- replace\_if() function, 1294, 1298
- reverse() function, 1295
- reverse\_copy() function, 1295, 1301
- rotate() function, 1295, 1301
- rotate\_copy() function, 1295, 1302
- stable\_partition() function, 1295, 1303
- swap() function, 1294, 1297
- swap\_ranges() function, 1294, 1297
- transform() function, 1294, 1297
- unique() function, 1295, 1300
- unique\_copy() function, 1295, 1301

## **myfirst.cpp program, 27-29**

- comments, 32-33
- header filenames, 34
- iostream file, 33-34
- main() function, 29-30
  - calling, 30
  - importance of, 32
  - int main() header, 30-31
- namespaces, 35-36

- free store, 183
- memory allocation, 160-162
- placement new, 671-676
- reference variables, 400
- newline character (\n), 38-39**
- newstrct.cpp, 179-180**
- next\_permutation() algorithm, 1039**
- next\_permutation() function, 1306, 1319**
- noboolalpha manipulator, 1090**
- noexcept, C++11, 1248**
- non-const objects, returning**
  - references to, 663
- non-member functions, 986-991**
- non-type arguments (arrays), 843-845**
- nonexpressions, for loops, 202**
- nonmodifying sequence operations, 1286**
  - adjacent\_find() function, 1287, 1290
  - count() function, 1287, 1291
  - count\_if() function, 1287, 1291
  - equal() function, 1288, 1291-1292
  - find() function, 1287-1289
  - find\_end() function, 1287-1290
  - find\_first\_of() function, 1287, 1290
  - find\_if() function, 1287-1289

- declaring, 546
- example, 547-549
- fill function, 357
- initializing, 546
- as exceptions, 903-908
- assignable, 1008
- associating with files, 289
- base-class objects, accessing, 801
- cerr, 1067
- cin, 1067, 1093-1095
  - cin.get() function, 235-237, 241-244
  - get() function, 128-130
  - getline() function, 126-127
  - loops, 234-235
  - operator overloading, 1095-1097
  - stream states, 1097-1102
- class, 788
- clog, 1067
- contained objects
  - compared to private inheritance, 806
  - initializing, 791
  - interfaces, 792-795

- operator\*(), 574-578
- operator+(), 569-572
- operator-(), 574-578
- operator<<(), 581-585, 587
- restrictions, 573-574
- subtraction operator (-), 574-578
- vector class, 588-590, 600
  - adding vectors, 590
  - declaring, 591-592
  - displacement vectors, 589
  - implementation comments, 602
  - member functions, 592, 597
  - multiple representations, 599
  - overloaded arithmetic operators, 599-600
  - overloading overloaded operators, 601
  - Random Walk sample program, 602, 605-606
  - state members, 597-599
  - with classes, string, 965
- operator\*() function, 574-578**
- operator+() function, 569-572**
- operator+() method, 1266**
- operator-() function, 574-578**

- operator\*(), 574-578
  - operator+(), 569-572
  - operator-(), 574-578
  - operator<<(), 581-587
- overloading, 101, 564-565
  - addition operator (+), 569-572
  - assignment operator, 652-658
  - example, 565-569
  - left shift operator (<<), 581-587, 676
  - member versus nonmember functions, 587-588
  - multiplication operator (\*), 574-578
  - operator functions, 565
  - overloading overloaded operators, 601
  - restrictions, 573-574
  - subtraction operator (-), 574-578
- precedence, 1231
  - examples, 1234
  - table of, 1232-1234
- reference operator (&), 383-386
- reinterpret\_cast, 946



left shift operator (<<),

581-587, 676

member versus nonmember  
functions, 587-588

multiplication operator (\*),  
574-578

operator functions, 565

restrictions, 573-574

subtraction operator (-), 574-578

vector class, 588-590

reference parameters, 415

templates, 422-424

overload resolution, 431-438

**override, 1183-1184**

**ownership, 973**

## P

---

**pairs.cpp, 848**

**palindromes, 1057**

**pam() function, 362-363**

**parameter lists, 30**

**parameterized types, 419**

- new operator, 160-162
- passing variables, 386-390
- pointer arithmetic, 167-172
- pointer notation, 173
- pointers to objects, 665-670
- pointers to pointers, 335
- stacks of pointers, 837-843
- strings, 173-178
- this, 539-546

**polar coordinates, 347**

- converting rectangular coordinates to, 348-351

**polymorphic public inheritance, 722-723**

- base-class functions, 777
- Brass class declaration, 723-726
- Brass class implementation, 727-731
- Brass class objects, 732-733
- BrassPlus class declaration, 723-726
- BrassPlus class implementation, 727-731
- BrassPlus class objects, 732-733
- constructors, 742
- dynamic binding, 737-740
- pointer compatibility, 737-739

- list, 1014-1017
- priority\_queue, 1017-1018
- queue, 1017
- stack, 1018
- vector, 1012-1013
- iterators, 992, 997-1005
  - back insert iterators, 1005-1007
  - bidirectional iterators, 998
  - concepts, 1000
  - copy() function, 1001-1002
  - forward iterators, 998
  - front insert iterators, 1005-1007
  - hierarchy, 999-1000
  - importance of, 992-996
  - input iterators, 997-998
  - insert iterators, 1005-1007
  - istream iterator template, 1003
  - models, 1001
  - ostream iterator template, 1002
  - output iterators, 998
  - pointers, 1001
  - random access iterators, 999
  - reverse iterators, 1003-1005
  - types, 997

BrassPlus class implementation, 727,  
730-731

BrassPlus class objects, 732-733  
constructors, 742

dynamic binding, 737-740

pointer compatibility, 737-739

reference type compatibility, 737-739

static binding, 737, 740

virtual destructors, 737, 742-743, 776

virtual functions, 739-742, 775-776

    behavior, 734-736

    friends, 743, 776

    memory and execution speed, 742

    redefinition, 743-745

    virtual function tables, 740

**public interfaces, 509**

**public keyword, 511-513**

**public member functions, 513**

**pure virtual functions, 748**

**push\_back() function, 1041**

**push back() method, 982-984**

**push\_heap() function, 1305**

**pushing values onto heap, 1314**

**put() method, 1071-1075**

**putback() member function, 1109-1114**

- real numbers, 50**
- recommended reading, 1323-1324**
- rect\_to\_polar() function, 348-349**
- rectangular coordinates, 346**
  - converting to polar coordinates, 348-351
- recur.cpp, 355, 358**
- recurs() function, 357-359**
- recursion, 357**
  - multiple recursive calls, 359-361
  - single recursive call, 358-359
  - variadic template functions, 1199-1202
- recursive use of templates, 846-847**
- redefining virtual functions, 743-745**
- redirecting I/O, 1067-1068**
- redirection, 238**
- refcube() function, 391-393**
- reference, passing by, 343, 386, 389-390, 770**
- reference arguments, 392-394, 408-409**
- reference counting, 973**
- reference operator (&), 383-386**
- reference parameters, overloading, 415**

associative, multimap, 1023-1025  
list, 1014, 1017  
    member functions, 1014-1016  
vector, 1012-1013

**review questions**

chapter 2, 62  
chapter 3, 110-111  
chapter 4, 191-192  
chapter 5, 250  
chapter 6, 298-300  
chapter 7, 372-373  
chapter 8, 443-444  
chapter 9, 498-501  
chapter 10, 558-559  
chapter 11, 623  
chapter 12, 700-702  
chapter 13, 779-780  
chapter 14, 869-870  
chapter 15, 947-949  
chapter 16, 1056-1057  
chapter 17, 1146-1147  
chapter 18, 1209-1212

**rewrite rule, 517**

**rfind() method, 961, 1261**

**right manipulator, 1091**

## sequences

mutating sequence operations

copy\_backward() function,  
1294-1297

copy() function, 1293-1296

fill() function, 1294, 1299

fill\_n() function, 1294, 1299

generate() function, 1294, 1299

generate\_n() function, 1294, 1299

iter\_swap() function, 1294

partition() function, 1295, 1302-1303

random\_shuffle() function,  
1295, 1302

remove\_copy() function, 1295, 1300

remove\_copy\_if() function,  
1295, 1300

remove\_if() function, 1295, 1300

remove() function, 1295, 1299

replace\_copy() function, 1294, 1298

replace\_copy\_if() function,  
1294, 1298

replace() function, 1294, 1298, 1302

replace\_if() function, 1294, 1298

reverse\_copy() function, 1295, 1301

- `show()`, array objects, 357
- `show_array()` function, 327-328
- `Show()` function, 818-820
- `show()` method, 514, 537
- `show_polar()` function, 347, 351
- `show_time()` function, 344-345
- showbase manipulator, 1090
- showperks() function, 744
- showpoint manipulator, 1091
- showpos manipulator, 1091
- showpt.cpp, 1084
- SHRT\_MAX constant, 72
- SHRT\_MIN constant, 72
- side effects of expressions, 201, 208-209
- signatures (functions), 413
- signed char data type, 88-89
- singly linked lists, 680
- size, string size
  - finding, 960
  - automatic sizing feature, 966-967
- `size()` function, 136, 960
- `size()` method, 509, 787, 981, 984, 1251-1252, 1275
- `size_type` constant, 1251
- `size_type` type, 1250, 1273



- concatenated output, 46-47
- cout.put() function, 83-84
- endl manipulator, 37-38
- integer values, displaying, 44-45
- \n newline character, 38-39
- declaration statements, 41-43
- defined, 29-30
- enum, 278-280
- examples, 41, 45
- for, declaration-statement expressions, 203
- if, 254
  - bug prevention, 260
  - example, 255
  - syntax, 254
- if else, 255
  - example, 256-257
  - formatting, 257-258
  - if else if else construction, 258-260
  - syntax, 255
- return statements, 30
- switch, 274-278
  - enumerators as labels, 278-280
  - example, 275-278
  - syntax, 275

- remove(), 1295, 1299
- remove\_copy(), 1295, 1300
- remove\_copy\_if(), 1295, 1300
- remove\_if(), 1295, 1300
- replace(), 1294, 1298, 1302
- replace\_copy(), 1294, 1298
- replace\_copy\_if(), 1294, 1298
- replace\_if(), 1294, 1298
- reverse(), 1295
- reverse\_copy(), 1295, 1301
- rotate(), 1295, 1301
- rotate\_copy(), 1295, 1302
- search(), 1288, 1292-1293
- search\_n(), 1288, 1293
- stable\_partition(), 1295, 1303
- swap(), 1294, 1297
- swap\_ranges(), 1294, 1297
- transform(), 1294, 1297
- unique(), 1295, 1300
- unique\_copy(), 1295, 1301
- functors, 1026-1027
  - adaptable, 1032
  - concepts, 1027-1028, 1030
  - predefined, 1030-1032
- generic programming, 992

- scope and linkage, 454
- static variables, 459-462
  - external linkage, 463, 466-467
  - internal linkage, 467-470
  - no linkage, 470-472

**str1.cpp, 953**

**str2.cpp, 966-967, 971**

**strcat() function, 136**

**strcmp() function, 221-222, 648**

**strcpy() function, 177-178, 633**

**strctfun.cpp, 348**

**strctptr.cpp, 352-353**

**stream objects, 1067**

**stream states, 1097-1098**

- effects, 1100-1102

- exceptions, 1099-1100

- file I/O, 1118-1119

- get() and getline() input effects, 1108

- setting, 1098

**streambuf class, 1065**

**streams, 1063-1064, 1067**

- istream class, 47

- ostream class, 47

**strfile.cpp, 958-959**

**strgfun.cpp, 340**

- accessing, 1259
- accessing with for loops, 206-207
- appending, 1265-1266
- assigning, 1266
- C-style, 120-122
  - combining with numeric input, 130-131
  - concatenating, 122
  - empty lines, 130
  - failbits, 130
  - in arrays, 123-124
  - null characters, 121
  - passing as arguments, 339-341
  - pointers, 173-178
  - returning from functions, 341-343
  - string input, entering, 124-126
  - string input, reading with `get()`, 127-130
  - string input, reading with `getline()`, 126-127
- comparing, 1263-1265
  - C-style strings, 220-223
  - string class strings, 223-224
- concatenating, 128, 1266

- `strquote.cpp`, 402-403
- `strtref.cpp`, 395
- `strtype1.cpp`, 132
- `strtype2.cpp`, 134
- `strtype4.cpp`, 137
- `struct` keyword, 140
- `structur.cpp`, 142
- structure initialization, C++11, 144
- structure members, 141
- structured programming, 12
- structures, 140-142, 343, 346
  - addresses, passing, 351-353
  - arrays, 147-148
  - assignment, 145-146
  - bit fields, 148
  - compared to classes, 514
  - dynamic structures, new operator, 178-180
  - example, 142-144
  - nested structures, 682
  - passing/returning structures, 344-345, 348-351
  - polar coordinates, 347
  - rectangular coordinates, 346

- auto\_ptr, 969, 973-975
- complex, 1045
- deque, 1013
- explicit instantiations, 850
- explicit specializations, 850-851
- implicit instantiations, 850
- list, 1014-1017
  - member functions, 1014-1016
- members, 854-855
- partial specializations, 851-852
- pointers, stacks of pointers, 837-843
- priority\_queue, 1017-1018
- queue, 1017
- stack, 1018
- valarray, 1045-1046, 1049-1051
- vector, 979-991, 1012-1013, 1045-1046, 1049-1051
  - adding elements to, 982-983
  - past-the-end iterators, 981-982
  - removing ranges of, 982
  - shuffling elements in, 987
  - sorting, 987
  - vect1.cpp example, 980-981
  - vect2.cpp sample program, 984-986
  - vect3.cpp sample program, 988-991

- cin object, 234-235
- end-of-file conditions, 237-241
- sentinel characters, 234
- text files, 287-288, 1129**
  - reading, 292-298
  - writing to, 288-292
- textin1.cpp, 234**
- textin2.cpp, 236**
- textin3.cpp, 239**
- textin4.cpp, 242**
- third-generation specialization, 425-426**
- this pointer, 539**
- throw keyword, 900**
- throwing exceptions, 900, 915-916**
- tilde, 529, 1237**
- time, 1009**
- time-delay loops, 229-230**
- tmp2tmp.cpp, 862-864**
- toggling bits, 1241**
- tokens, 39**
  - alternative tokens, table of, 1222
- tolower() function, 273, 1041**
- top-down design, 12**
- top-down programming, 331**
- topfive.cpp, 353**

- ULONG\_MAX constant, 72
- UML (Unified Modeling Language), 08
- unary functions, 1027, 1030
- unary minus operator, 601
- unary operators, 601, 1234
- unbound template friend functions, 864-865
- uncaught exceptions, 928-931
- underscore (`_`), 1222
- unexpected exceptions, 928-931
- unexpected() function, 929
- unformatted input functions, 1102
- Unicode, 88
- Unified Modeling Language (UML), 1208
- Unified Modeling Language User Guide*, 1323
- uniform initialization, 1154
  - initializer\_list, 1155
  - narrowing, 1154
- unions, 149
  - anonymous unions, 150
  - declaring, 149
- unique\_copy() function, 1295, 1301
- unique() function, 1041, 1295, 1300
- unique() method, 1016-1017, 1280
- unique\_ptr versus auto\_ptr, 975-977



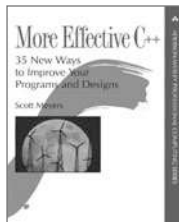
- enum, 150-152
  - enumerators, 150-151
  - value ranges, 153
  - values, setting, 152
- floating-point numbers, 92
  - advantages/disadvantages, 96-97
  - constants, 96
  - decimal-point notation, 92
  - double data type, 94-96
  - E notation, 92-93
  - float data type, 94-96
  - long double data type, 94-96
- global, compared to local
  - variables, 467
- indeterminate values, 73
- initializing, 52, 73
- integers, 68
  - bool, 90
  - char, 80-89
  - choosing integer types, 76-77
  - climits header file, 71-73
  - constants, 78-80, 90-92
  - initializing, 73

- past-the-end iterators, 981-982
- Random Walk sample program, 602, 605-606
- removing ranges of, 982
- shuffling elements in, 987
- sorting, 987
- state members, 597-599
- vect1.cpp example, 980-981
- vect2.cpp sample program, 984-986
- vect3.cpp sample program, 988-991
- vector class templates, 1012-1013**
- vector containers, 1012-1013**
- vector objects versus arrays, 188-189**
- vector template class, 120, 186-187**
- vectors, methods, 1278-1280**
- vegnews.cpp, 634**
- versatility of templates, 845-846**
- version1() function, 403**
- version2() function, 404**
- version3() function, 405**
- ViewAcct() function, 725-726, 730**
- virtual base classes, 815-817**
  - combining with nonvirtual base classes, 828
  - constructors, 817-818





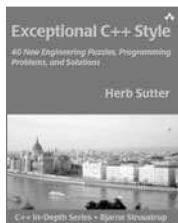
978-0-321-33487-9



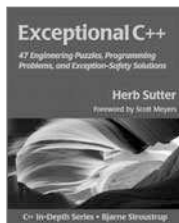
978-0-201-63371-9



978-



978-0-201-76042-2



978-0-201-61562-3



978-

  
Addison  
Wesley

  
PRENTICE  
HALL

**informIT**  
**Safari**   
Books Online

**T**he **Developer's Library Series** from practicing programmers with unique tutorials on the latest programming language use in their daily work. All books in the D expert technology practitioners who are and presenting information in a way that

Developer's Library books cover a wide source programming languages and data Microsoft, and Java, to Web development Mac/iPhone programming, and Android

